



## Hardware Opportunities in the Machine Learning Lifecycle

Joseph E. Gonzalez  
Co-director of the RISE Lab  
[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)

1

Get the latest slides and links to literature



<https://tinyurl.com/isscc-lifecycle>

2

1

2

### About Me

- Co-director of the RISE Lab
- Co-founder of Turi Inc.
- Member of the Apache Spark PMC
- Research
  - Artificial Intelligence
  - Data Science
  - Distributed Data Systems
  - Graph Processing Systems
- I don't study processor **architecture**
  - But I probably should ...



3

### Outline

- History and the **Co-evolution** of Hardware and AI
  - The **Feedback Cycle** driving the 3<sup>rd</sup> wave of AI
- Machine Learning is **not a single workload**
  - Stages of the **Machine Learning Lifecycle**
- Security** and machine learning

*Along the way, I will talk about some of the research in my group addressing interesting aspects of the lifecycle.*

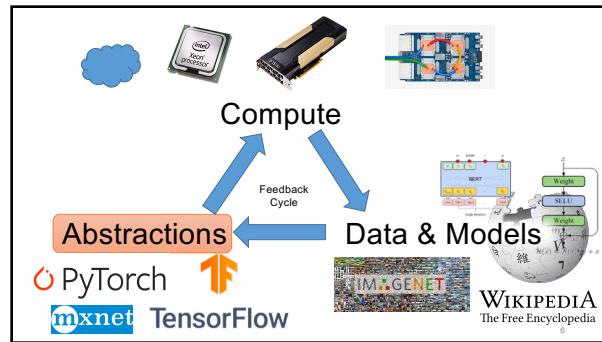
4

4

### Hardware and the History of AI

- 1950 to 1974: Birth of AI**
  - 1951 Marvin Minsky builds first **neural network hardware** (SNARC)
- 1974 to 1980: First AI Winter**
  - Limited processing power and data
- 1980 to 1987: Second Wave of AI**
  - XCON (AI for **hardware configuration**) for DEC → boom in AI hardware companies
- 1987 to 1993: Second AI Winter**
  - Brittle AI and the collapse of the **AI Hardware Market**
- 1993 to 2011: AI → Machine Learning**
  - Confluence of ideas + **Compute** + **Big Data** → AI starts to really work
- 2011 to 2019: Third Wave (Deep Learning)**
  - **Compute** + **data** + **abstractions** → feedback cycle

5



6

## Abstractions are Enabling Innovation

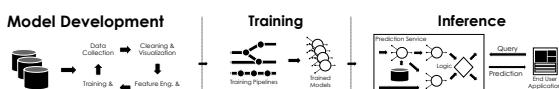
- ❑ Much of machine learning before 2010
    - Research focused on machine learning **algorithms**
    - Programs written using **high-level imperative languages**
      - Matlab/R/C++/Java
    - **Big abstractions:** linear algebra, map-reduce, graph systems
  - ❑ Today:
    - Research focused on **model design**
    - **Models** written in high-level **DSLs**
      - TensorFlow/Pytorch
    - Big abstractions: **tensor operations, loss minimization**, linear algebra, ...
  - ❑ Models written in **TensorFlow** can now run on **hardware** that didn't exist when the models were created.

8

# How do we make hardware for Machine learning?

# Machine learning is not a single application.

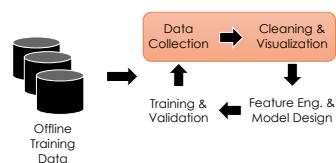
Machine learning is  
**multiple applications** with different requirements.



## Machine Learning Lifecycle

10

## Model Development



## **Identifying** potential sources of data

## Joining data from multiple sources

## Addressing missing values and outliers

## Plotting trends to identify anomalies

## Model Development



- Building informative features functions
  - Designing new model architectures
  - Tuning training algos.
  - Validating prediction accuracy

13

12

Model Development *Frameworks*



13

## Model Development → Hardware

- Need to test **multiple designs and hyperparameters** quickly
  - May be better to run many parallel experiments than one experiment faster
- **Debug heavy** → sources of error → data, hyperparams., & model
  - **System should not be a source of error**
  - **Avoid** cutting corners (e.g., quantization, async) for increased performance
    - Unless you can make a case for stable convergence ...
- **Data preparation** is often a bottleneck
  - Opportunity for **data tooling**
  - Accelerate data transformation and augmentation
- Emerging Trends
  - **Attention Models** and **Graph Neural Networks**: reduced locality, sparsity
  - **Dynamic Networks**: gating, cascades, mixtures, ...
  - Increased emphasis on **DNN features** and **Fine-Tuning**
    - Reuse of common architectures and weights

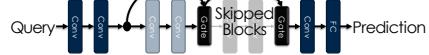
14

## Dynamic Networks for **fast** and **accurate** inference

**IDK Cascades:** Using the fastest model possible [UAI'18]

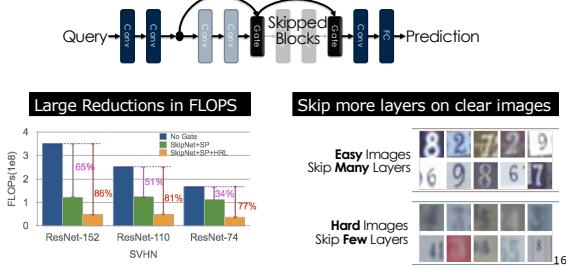


**SkipNet:** dynamic execution within a model [ECCV'18]



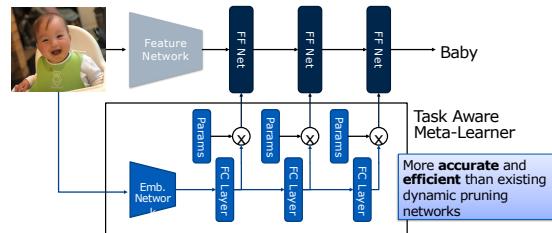
15

## SkipNet: dynamic execution within a model [ECCV'18]



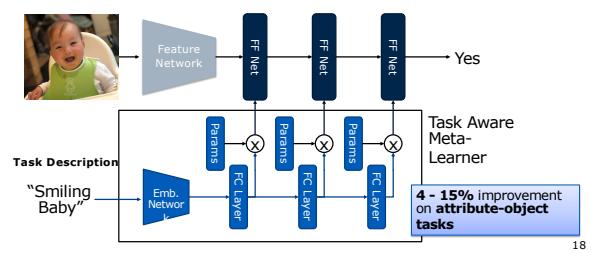
16

## Task Aware Feature Embeddings [CVPR'19]



17

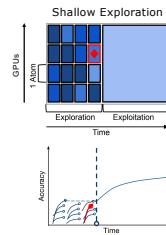
## Task Aware Feature Embeddings [CVPR'19]



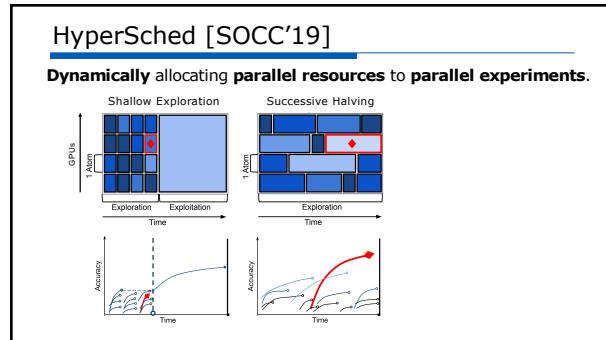
18

## HyperSched [SOCC'19]

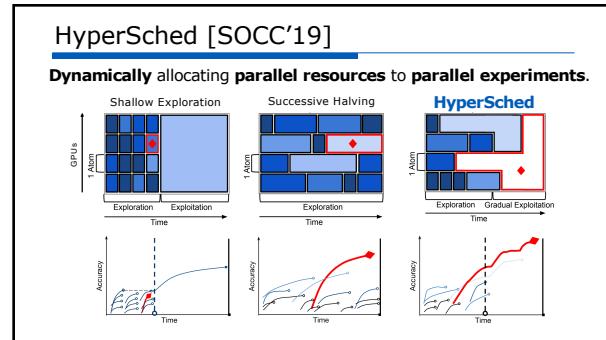
**Dynamically** allocating **parallel resources** to **parallel experiments**.



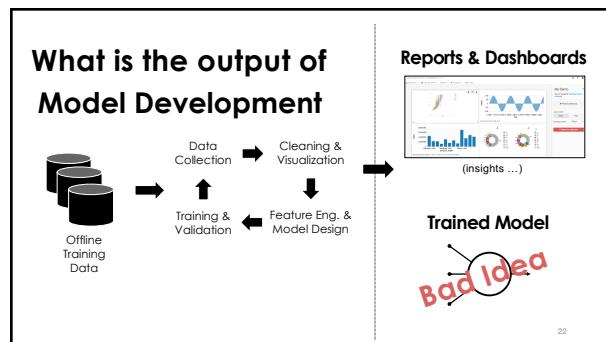
19



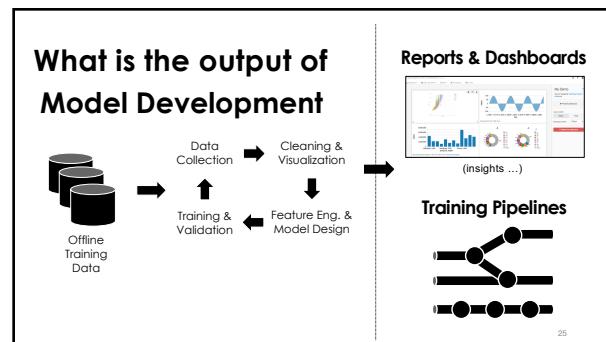
20



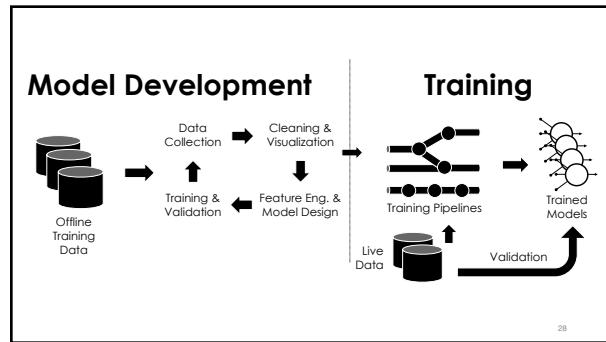
21



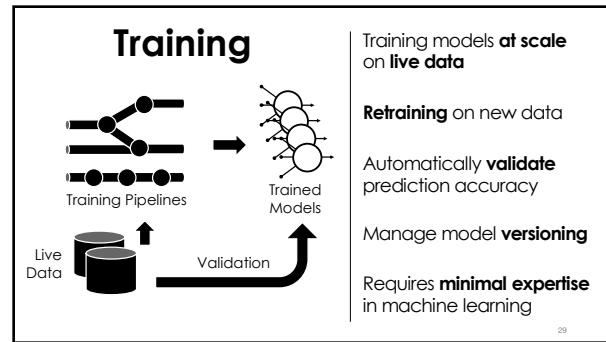
22



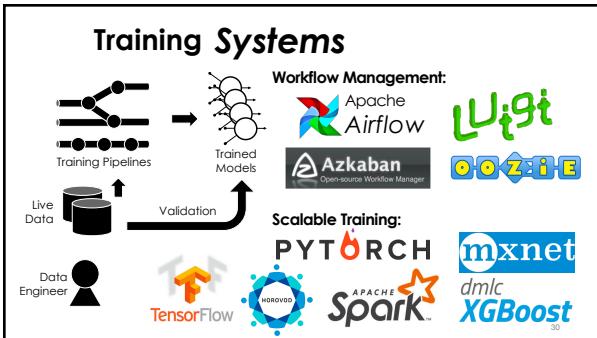
25



28



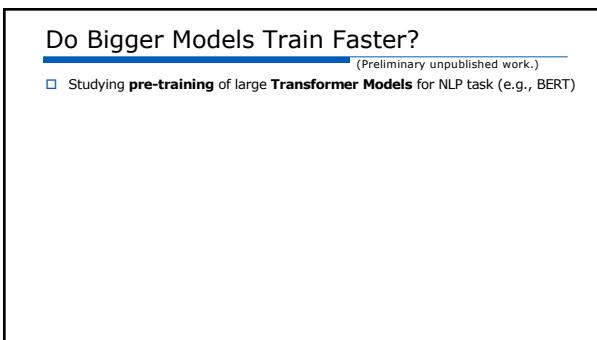
29



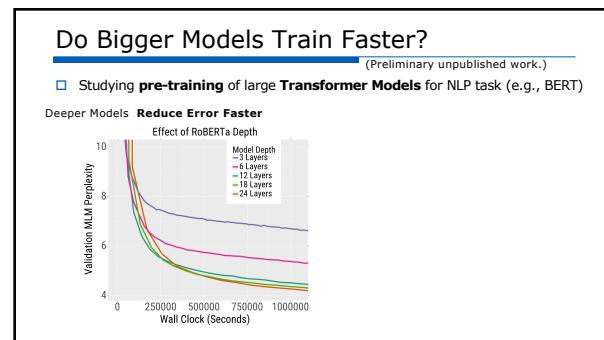
30

- ### Model Training → Hardware
- Fewer models to train → need **distributed training of individual models**
    - Often train with **more data**
  - **Larger models and mini-batch sizes**
    - Need larger on-device memory
    - Counter trends → reversible networks, **optimal checkpointing**, ...
  - Models and hyperparameters are vetted → focus on **system optimizations**
    - Can tolerate some **system error** (quantization and async.)
    - Need adequate stability to meet deadlines
  - **Data preparation** is still potentially an issue (as with model dev.)
  - Need to deal with **composition** of multiple models

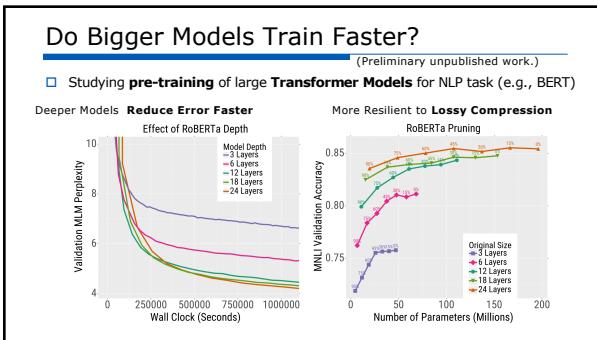
31



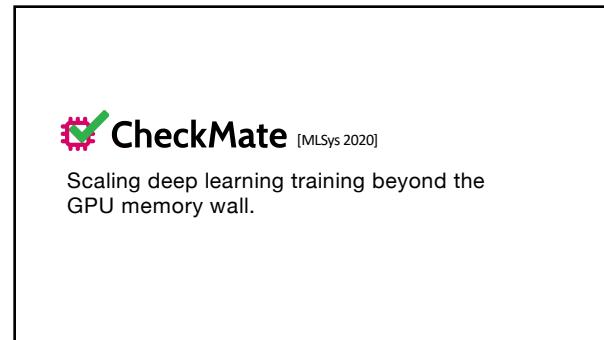
32



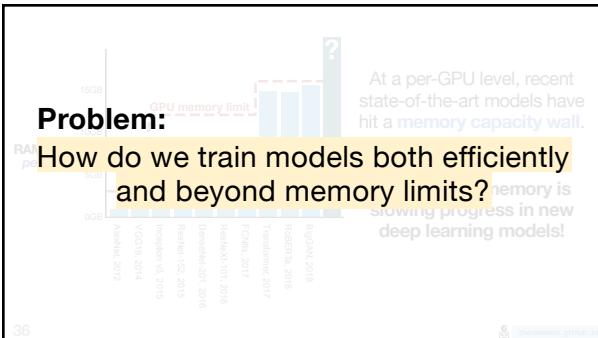
33



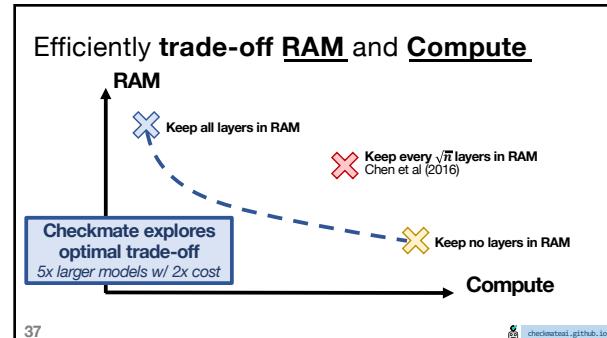
34



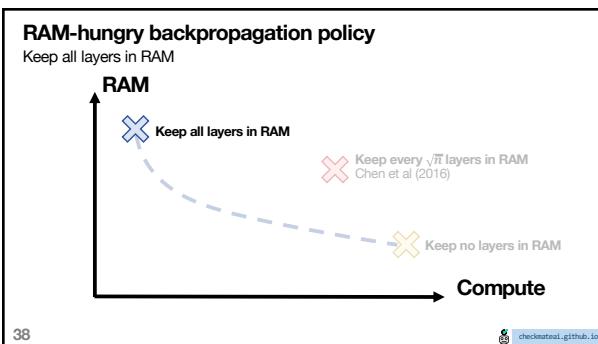
35



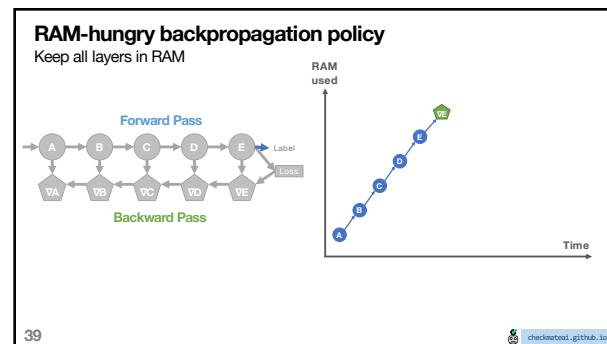
36



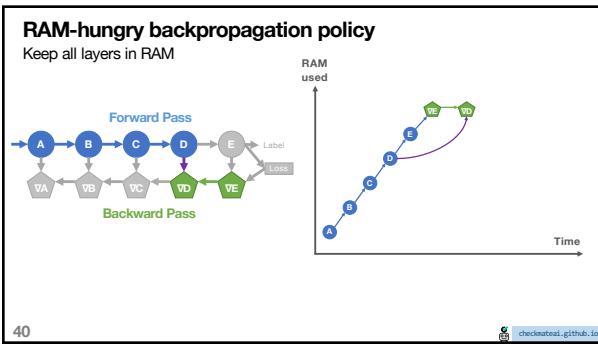
37



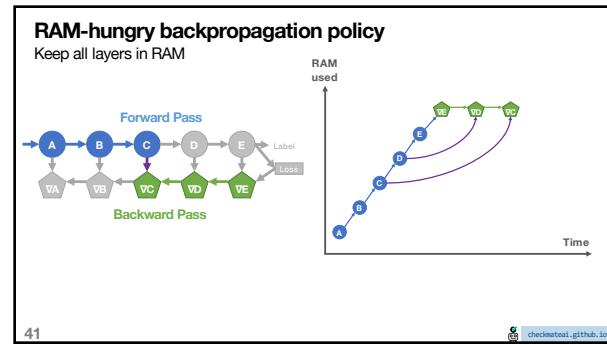
38



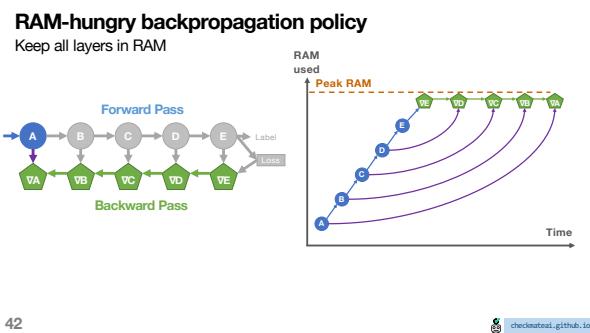
39



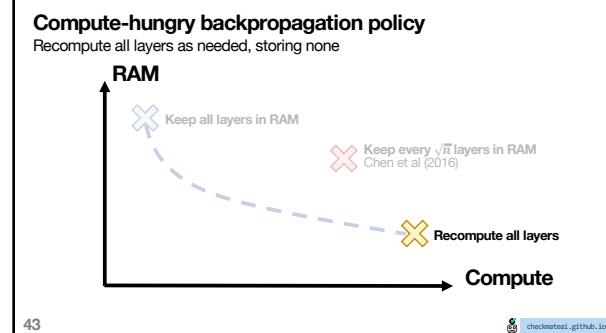
40



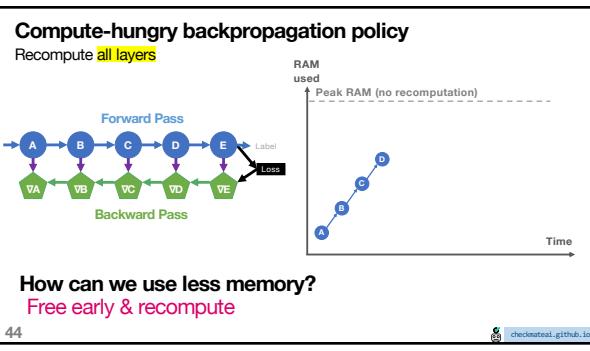
41



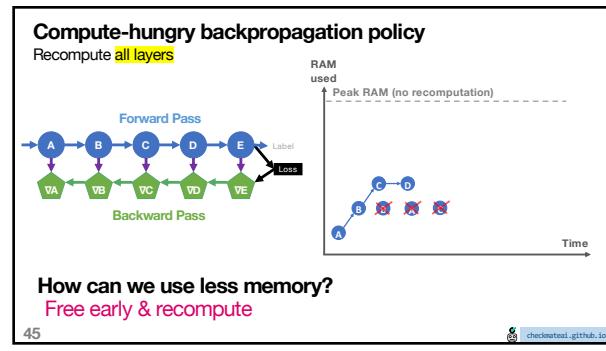
42



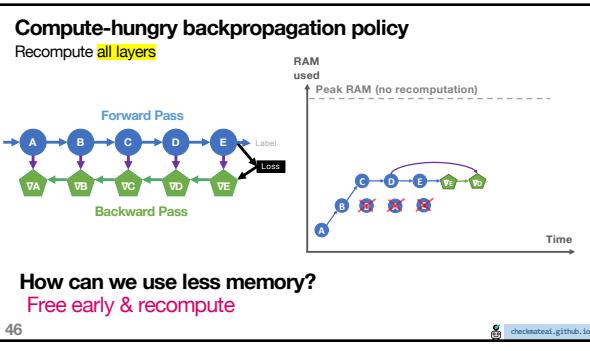
43



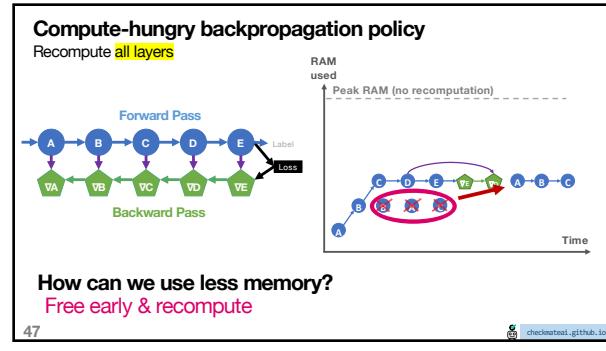
44



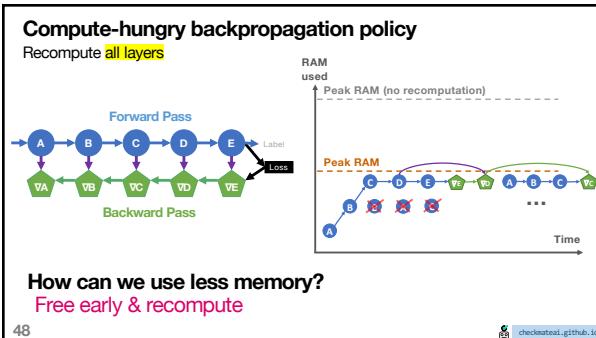
45



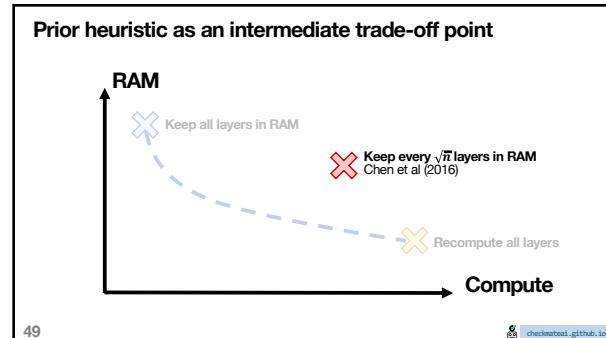
46



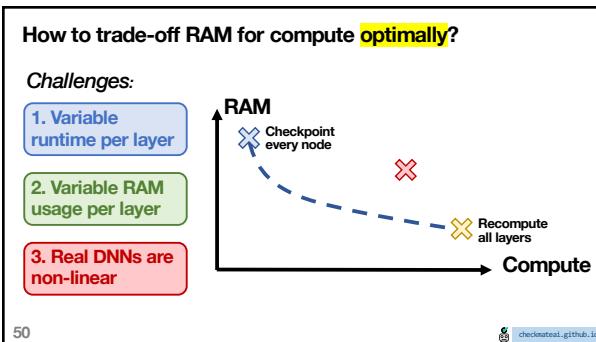
47



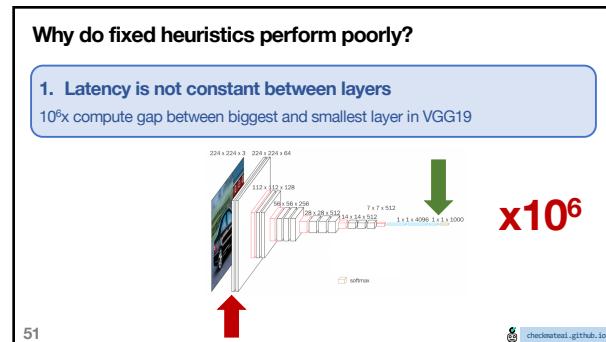
48



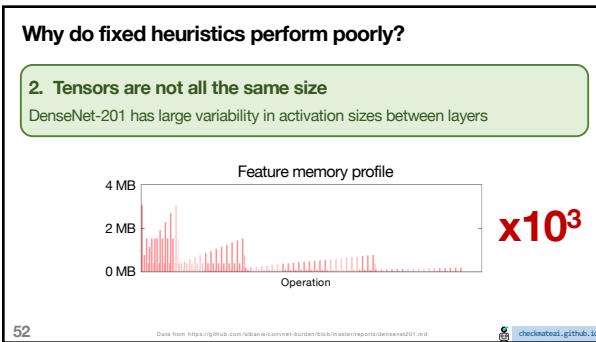
49



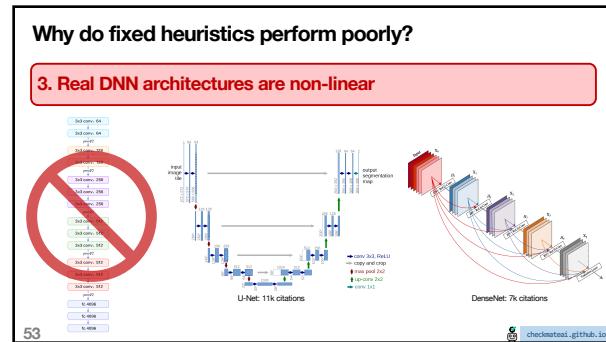
50



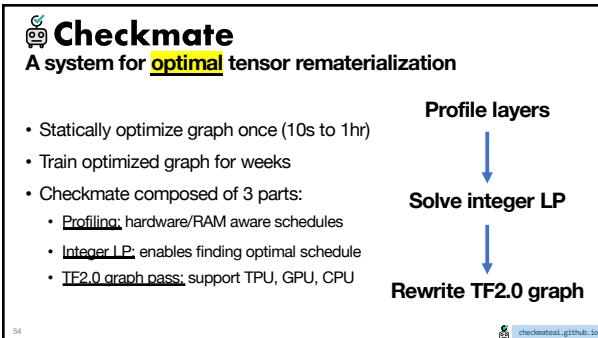
51



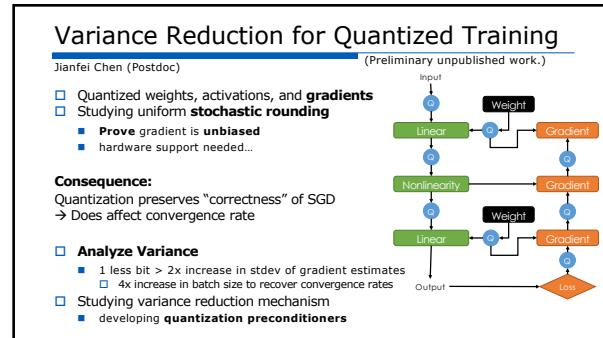
52



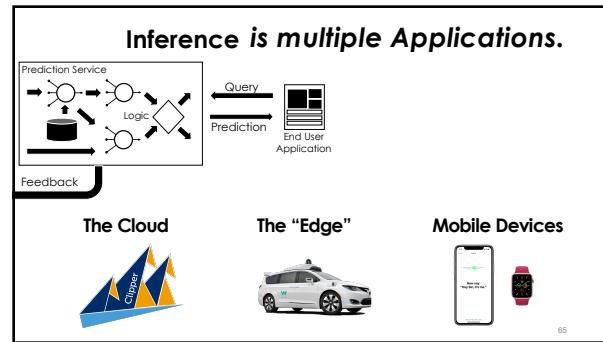
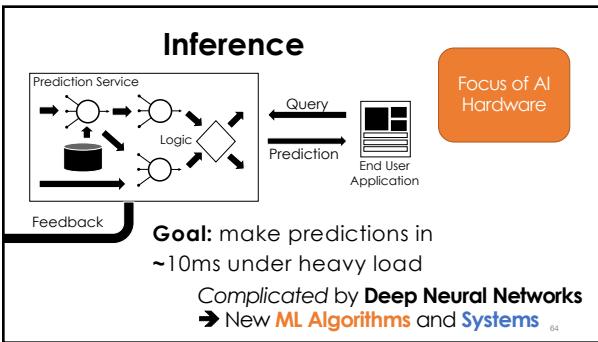
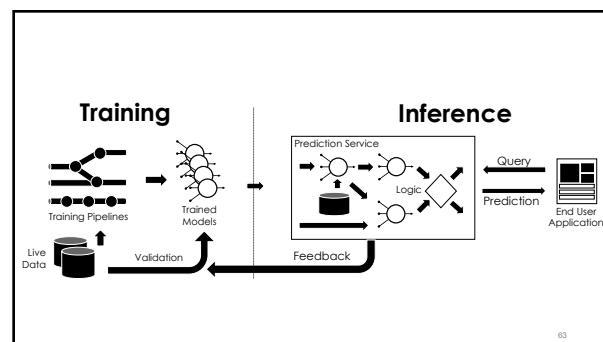
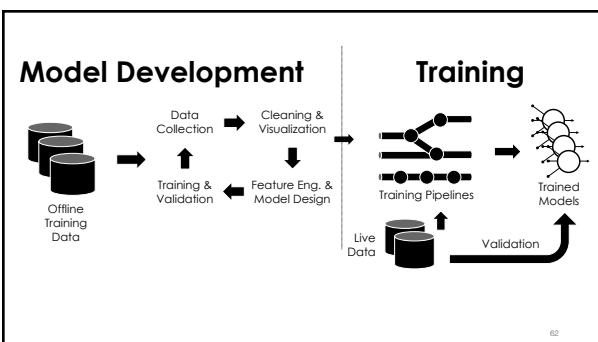
53

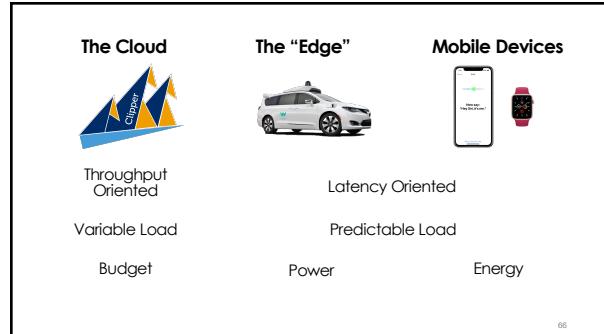


54

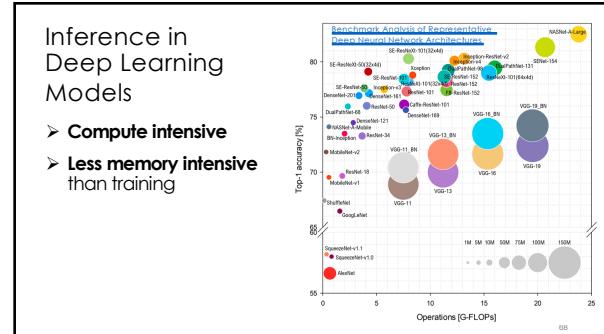


55





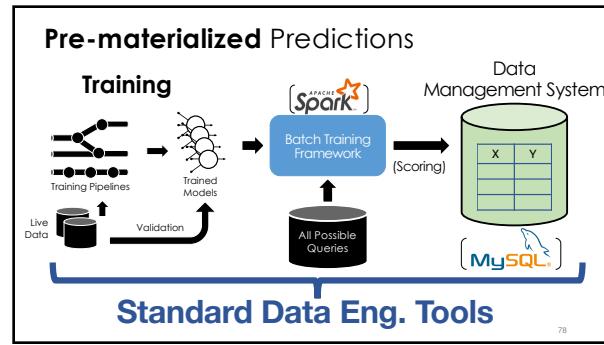
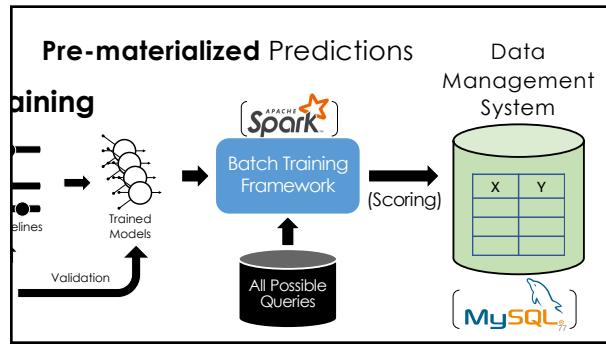
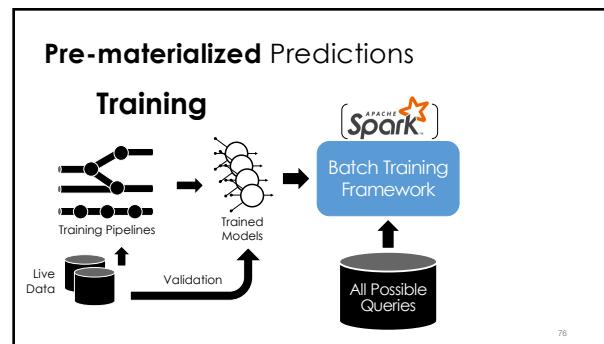
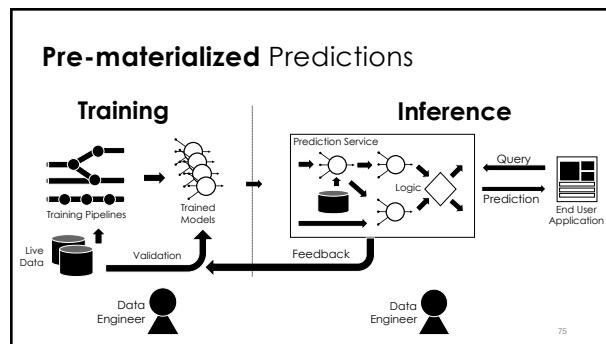
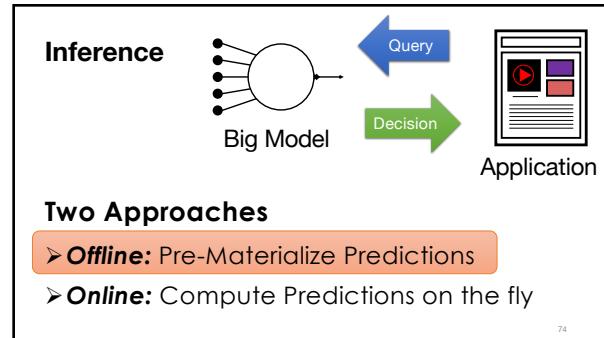
66

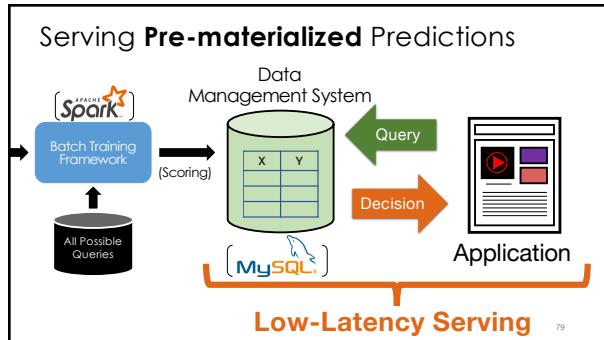


**Other Challenges?**

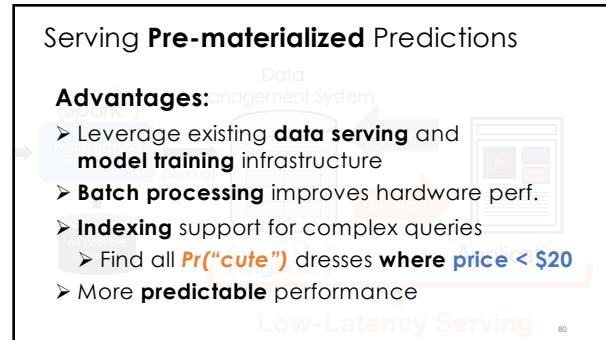
- **Bursty load** →
  - overprovision resources →
    - expensive
  - TPU reports 28% utilization of vector units in production
  - Solutions
    - statistical multiplexing → hardware not designed for multitenancy
    - could try to predict arrival process → generally difficult to predict
- **Versioning and testing models**
- **Prediction pipelines** → more on this soon

73

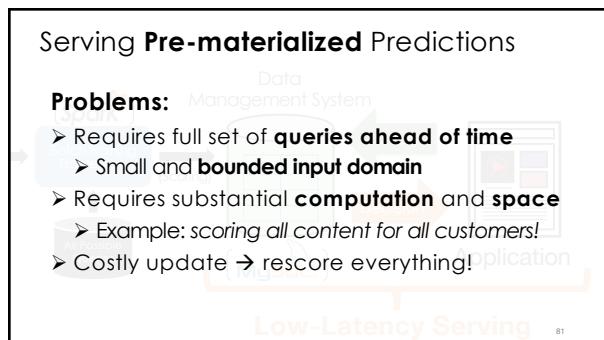




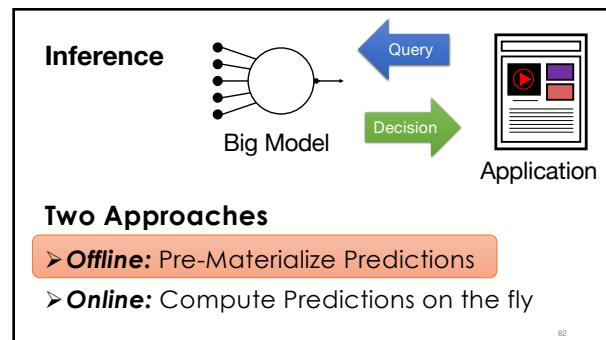
79



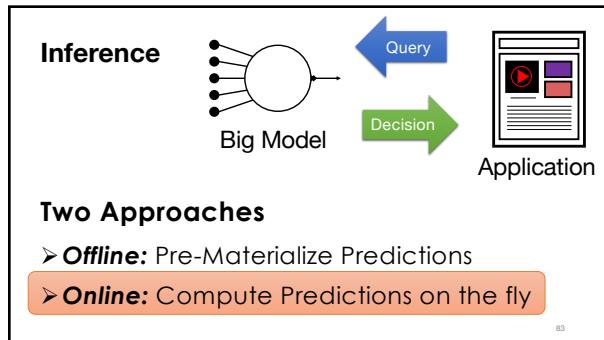
80



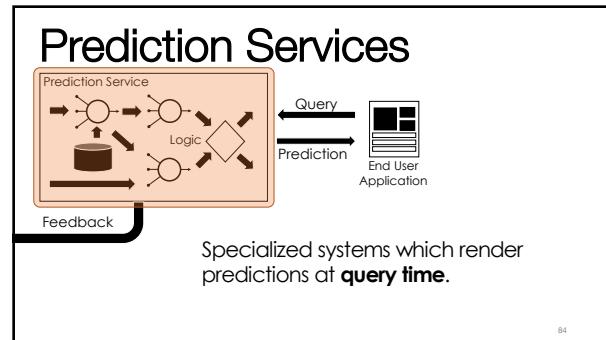
81



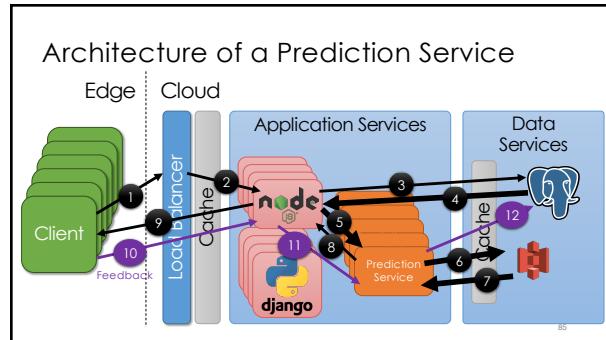
82



83



84

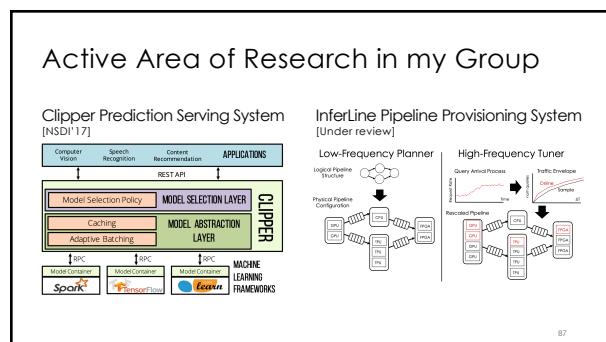


85

## Online: Compute Predictions at Query Time

- **Examples**
  - Signals processing: speech recognition & image tagging
  - Ad-targeting based on search terms, available ads, user features
- **Advantages**
  - Compute only necessary queries
  - Enables models to be changed rapidly (e.g., bandit exploration)
  - Queries do not need to be from small ground set
- **Disadvantages**
  - Increases complexity and computation overhead of serving system
  - Requires low and predictable latency from models

86

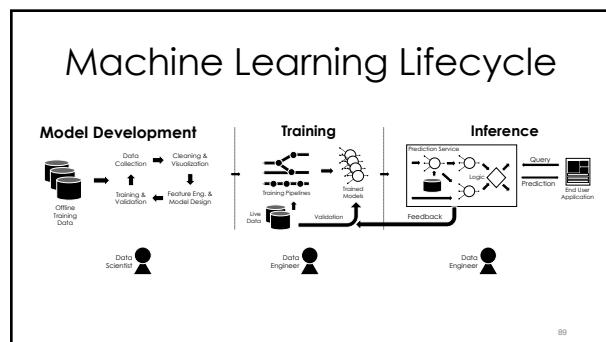


87

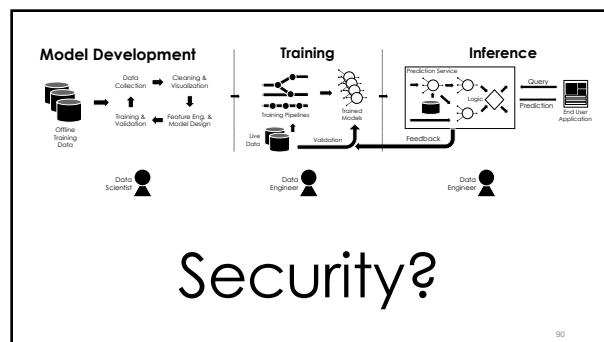
## Prediction Serving → Hardware

- Inference requires less memory → **focus on compute**
- Greater emphasis on **latency** instead of throughput
  - Focus on **small batch** inference (batch size = 1)
  - Opportunity to exploit **pipeline parallelism**
  - Need **high availability** → esp. in mission critical settings
- Often runs multiple **concurrent prediction tasks**
  - Cloud → Multitenancy → Performance isolation
  - Edge → supporting multiple data streams
- Tolerate **model compression and quantization**
  - As low as 4-bit activations and weights
- **Bursty load**
  - Statistical multiplexing
  - Use inference hardware for background training?

88



89



90

## Protect the data, the model, and the query

### High-Value Data is Sensitive

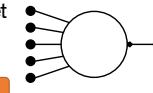
- Medical Info.
- Home video
- Finance

Biggest opportunity for hardware in ML.



### Models capture value in data

- Core Asset
- “Contain” the data



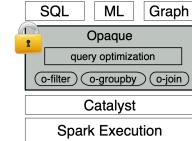
Queries can be as sensitive as the data



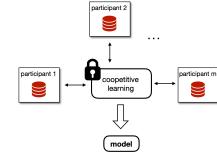
91

## Our recent work in secure ML

### Opaque | Oblivious Spark over SGX [NSDI'17]



### Helen | Competitive Learning Using Cryptographic Primitives [S&P'19]



92

## Security and Hardware

- Improved Access to Data
  - User willing to share data with models but **not companies (people)**
  - **Differential Privacy** can increase data sharing incentives
- Better **isolation** of co-tenant models on hardware accelerators
- Cooperative Learning**: Secure multiparty computation for ML
  - Example: Competing banks collaborate to construct a shared fraud model without sharing data.
- Models have access to **more sensitive inputs**
  - Example: Alexa could see where you are when asking to turn on a light.

93

93

## Conclusion

- History:** AI and Computer Systems have Co-evolved
- Feedback Cycle:** Hardware, Abstractions, and Data
- ML is many Applications:** *Machine Learning Lifecycle*
  - Model Development: Exploration
  - Training: Scale and Composition
  - Inference: Cloud – Edge Spectrum
- Security:** Opportunity for hardware innovation in AI

Thank you!

94

94