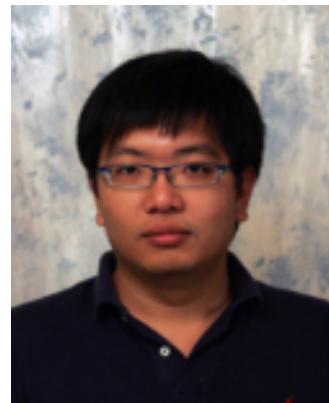


Parallel Gibbs Sampling

From Colored Fields to Thin Junction Trees



Joseph
Gonzalez



Yucheng
Low



Arthur
Gretton



Carlos
Guestrin

Sampling as an Inference Procedure

- Suppose we wanted to know the probability that coin lands “heads”

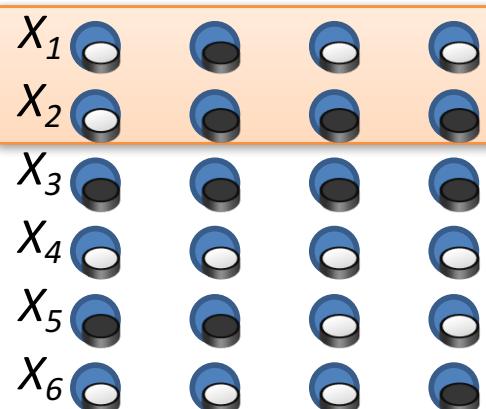
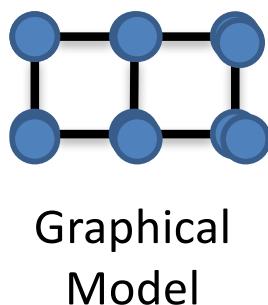
“Draw Samples”



$$P(\text{Heads}) \approx \frac{4}{10}$$

Counts
4x Heads
6x Tails

- We use the same idea for graphical model inference

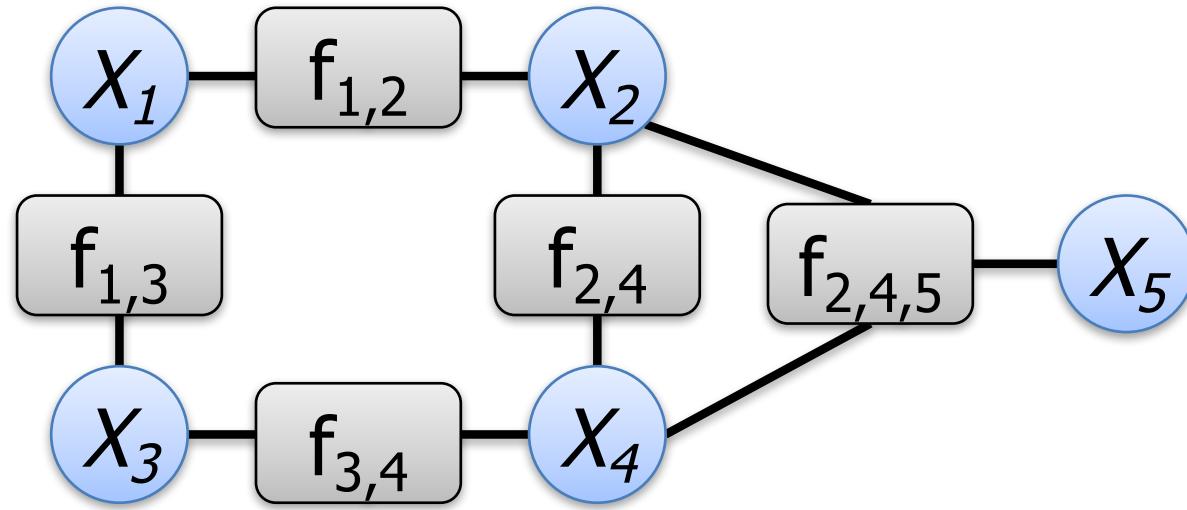


Inference:
 $P(X_1 = X_2) \approx 2/4$

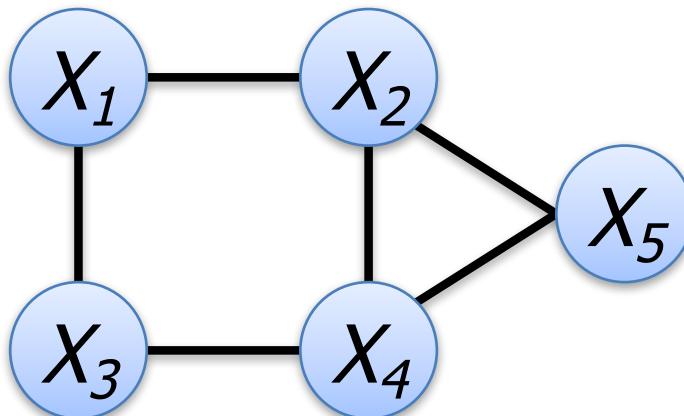
Terminology: Graphical Models

- Focus on **discrete** factorized models with **sparse structure**:

Factor
Graph



Markov
Random
Field



Terminology: Ergodicity

- The goal is to estimate:

$$\mathbf{E} [h(X_1, \dots, X_n)]$$

- Example: marginal estimation

$$h_i(x) = \mathbf{I}[x == i] \Rightarrow \mathbf{E}[h_i(X_k)] = \mathbf{P}(X_k = i)$$

- If the sampler is **ergodic** the following is true*:

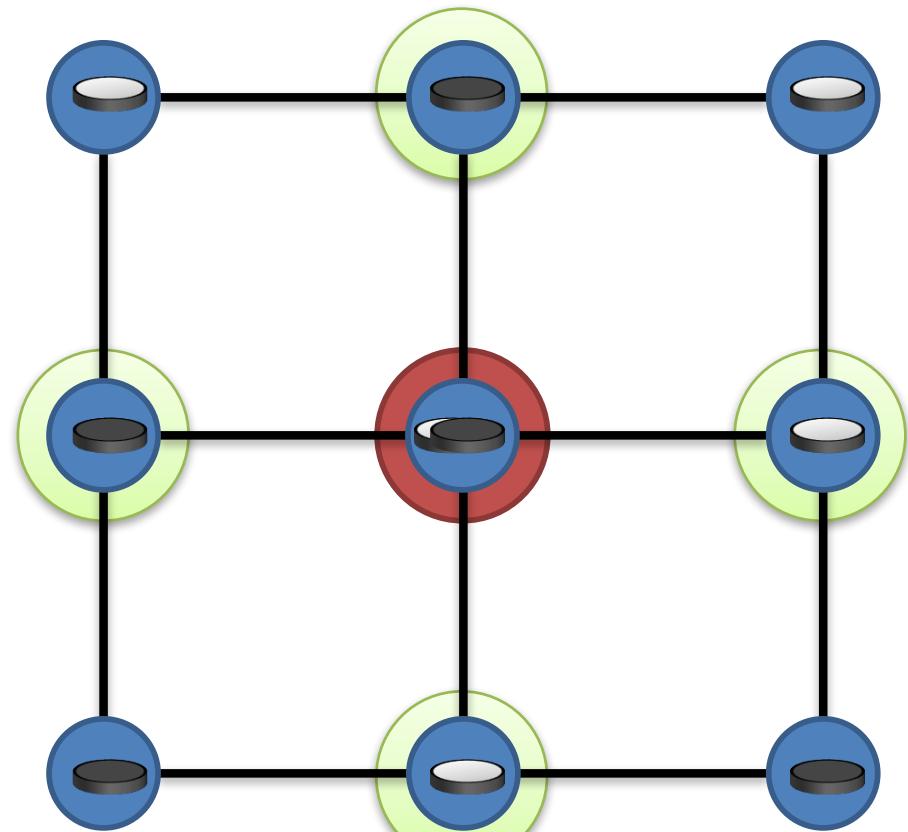
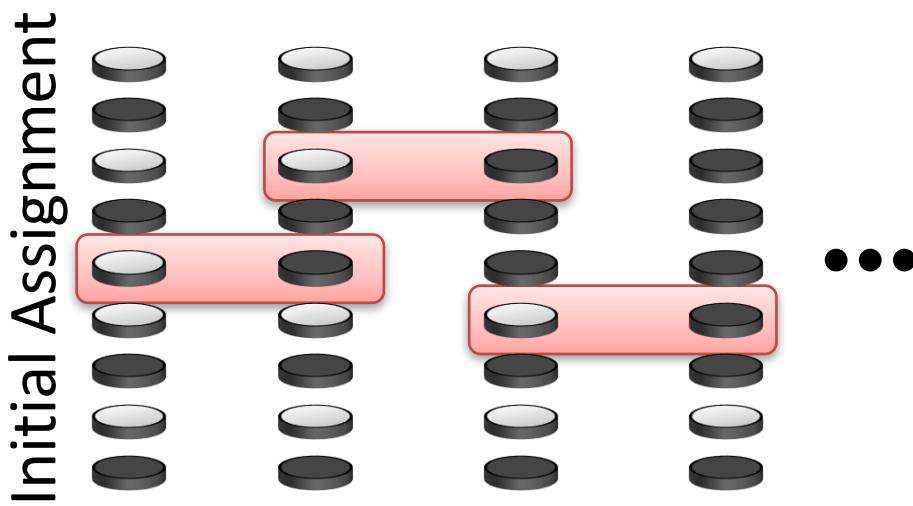
$$\lim_{m \rightarrow \infty} \frac{1}{m} \sum_{t=1}^m h(x_1^{(t)}, \dots, x_n^{(t)}) \xrightarrow{a.s.} \mathbf{E} [h(X_1, \dots, X_n)]$$

*Consult your statistician about potential risks before using.

Gibbs Sampling [Geman & Geman, 1984]

- Sequentially for each variable in the model

- Select **variable**
- Construct conditional given **adjacent assignments**
- Flip coin and update assignment to **variable**



Why Study Parallel Gibbs Sampling?

“The Gibbs sampler ... might be considered the workhorse of the MCMC world.”

–Robert and Casella

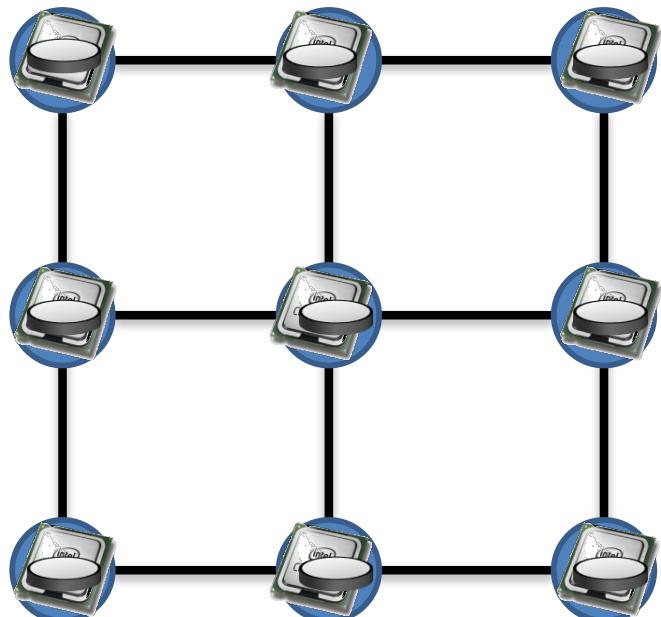
- Ergodic with geometric convergence
- Great for high-dimensional models
 - No need to tune a joint proposal
- Easy to construct algorithmically
 - WinBUGS
- Important Properties that help Parallelization:
 - **Sparse structure → factorized computation**

Is the Gibbs Sampler
trivially parallel?

From the original paper on Gibbs Sampling:

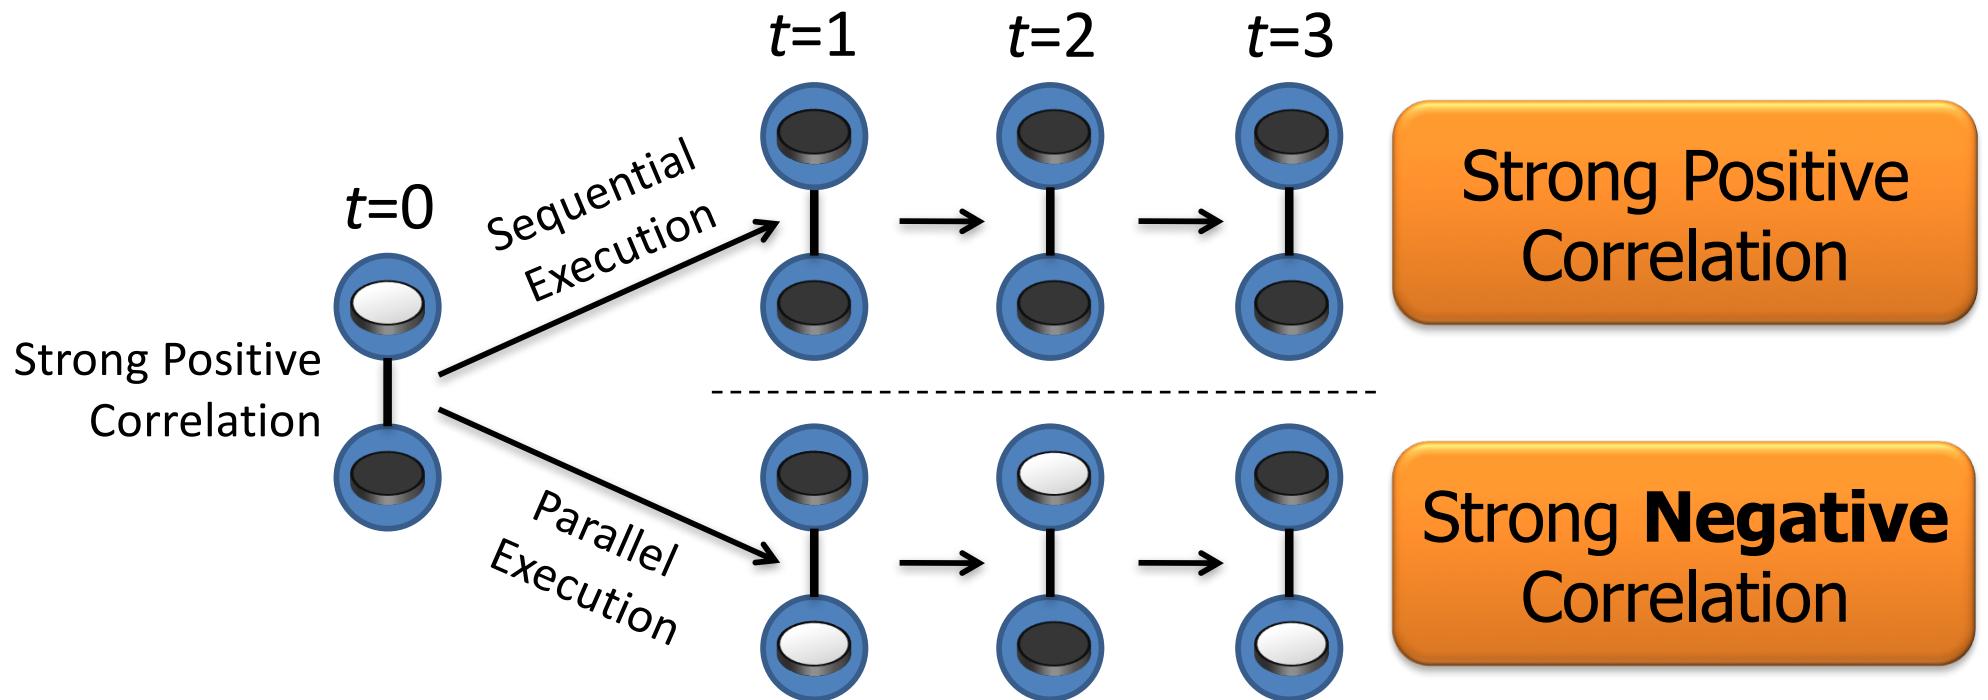
*“...the MRF can be divided into collections of [variables] with each collection assigned to an **independently** running **asynchronous processor**.”*

-- Stuart and Donald Geman, 1984.



Converges to the
wrong distribution!

The problem with Synchronous Gibbs



- *Adjacent variables **cannot** be sampled simultaneously.*

How has the machine
learning community solved
this problem?



Two Decades later

1. Newman et al., *Scalable Parallel Topic Models*. Jnl. Intelligen. Comm. R&D, 2006.
 2. Newman et al., *Distributed Inference for Latent Dirichlet Allocation*. NIPS, 2007.
 3. Asuncion et al., *Asynchronous Distributed Learning of Topic Models*. NIPS, 2008.
 4. Doshi-Velez et al., *Large Scale Nonparametric Bayesian Inference: Data Parallelization in the Indian Buffet Process*. NIPS 2009
 5. Yan et al., *Parallel Inference for Latent Dirichlet Allocation on GPUs*. NIPS, 2009.
-
- Same problem as the original Geman paper
 - Parallel version of the sampler is **not ergodic**.
 - Unlike Geman, the recent work:
 - Recognizes the issue
 - Ignores the issue
 - Propose an “approximate” solution

Two Decades Ago

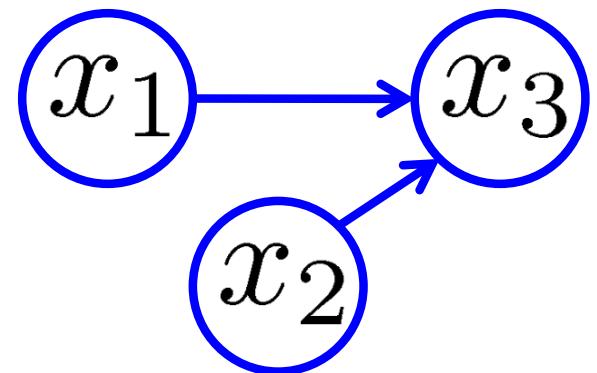
- Parallel computing community studied:

Sequential Algorithm

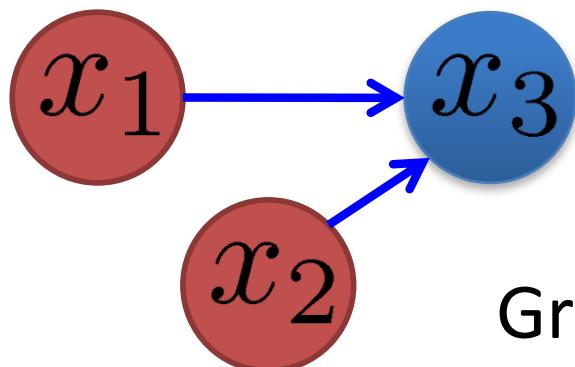
Time ↓

$$\begin{aligned}x_1^{(t+1)} &= f_1(x_3^{(t)}) \\x_2^{(t+1)} &= f_2(x_2^{(t)}, x_3^{(t)}) \\x_3^{(t+1)} &= f_3(x_1^{(t+1)}, x_1^{(t+1)})\end{aligned}$$


Directed Acyclic
Dependency Graph



Construct an **Equivalent** Parallel Algorithm

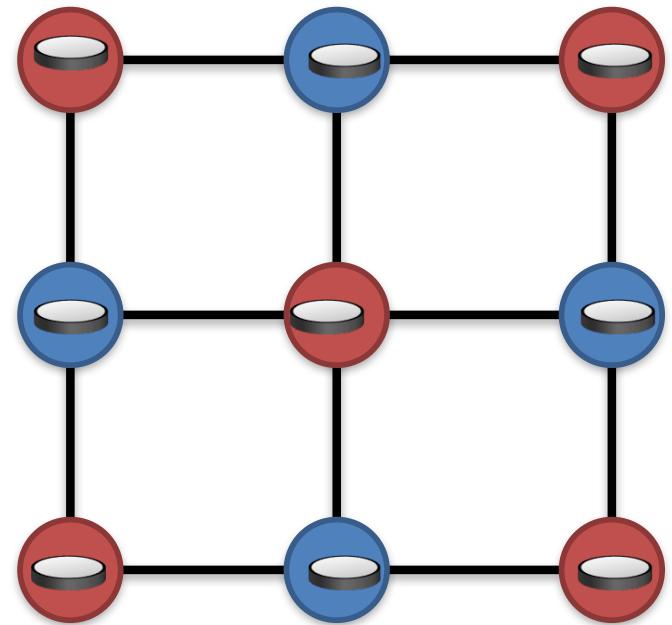


Using
Graph Coloring

$$\begin{aligned}x_1^{(t+1)} &= f_1(x_3^{(t)}) & x_2^{(t+1)} &= f_2(x_2^{(t)}, x_3^{(t)}) \\x_3^{(t+1)} &= f_3(x_1^{(t+1)}, x_1^{(t+1)})\end{aligned}$$

Chromatic Sampler

- Compute a k-coloring of the graphical model
- Sample all variables with same color in parallel
- Sequential Consistency:



Time →

Chromatic Sampler Algorithm

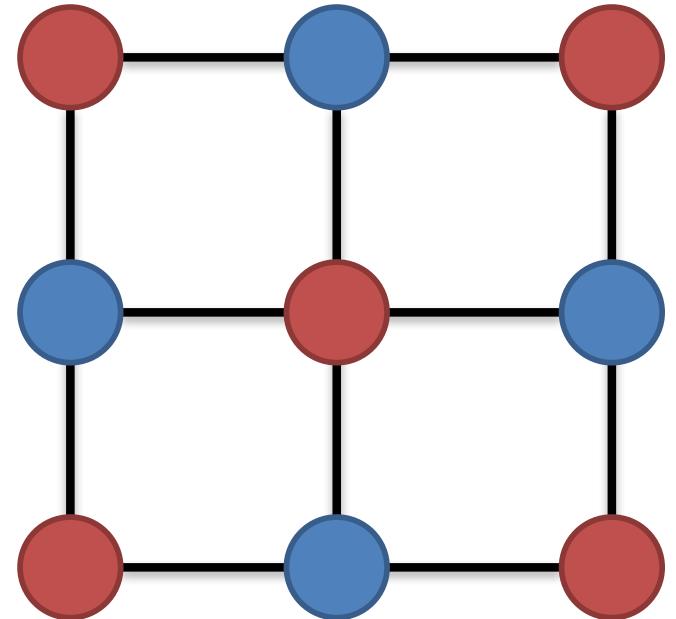
For t from 1 to T do

$$x^{(t)} \leftarrow x^{(t-1)}$$

For k from 1 to K do

Parfor i in color k :

$$x_i^{(t)} \sim \mathbf{P}(X_i \mid x_{-i}^{(t)})$$



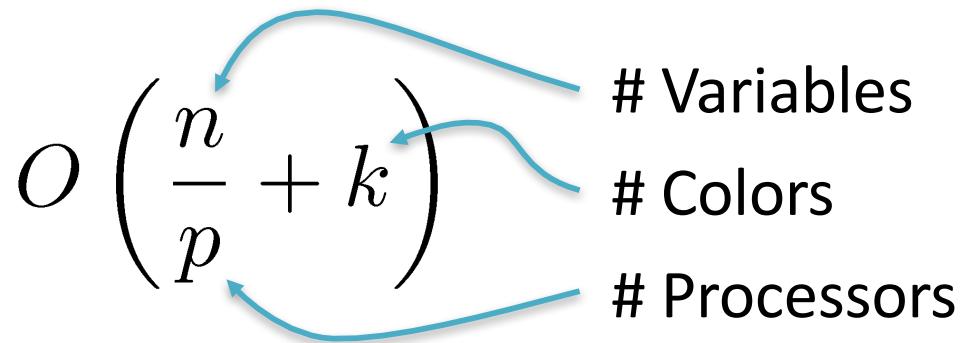
Asymptotic Properties

- Quantifiable acceleration in **mixing**

Time to update
all variables once

$$O\left(\frac{n}{p} + k\right)$$

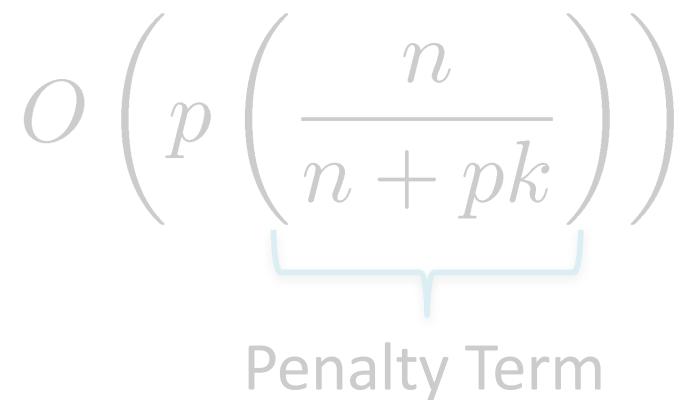
Variables
Colors
Processors



- Speedup:

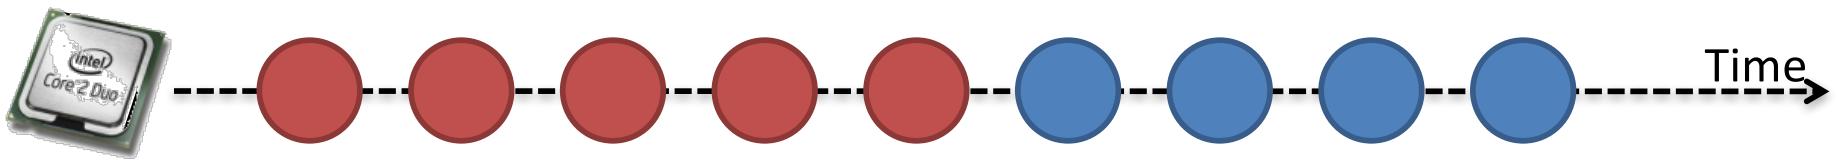
$$O\left(p\left(\frac{n}{n + pk}\right)\right)$$

Penalty Term

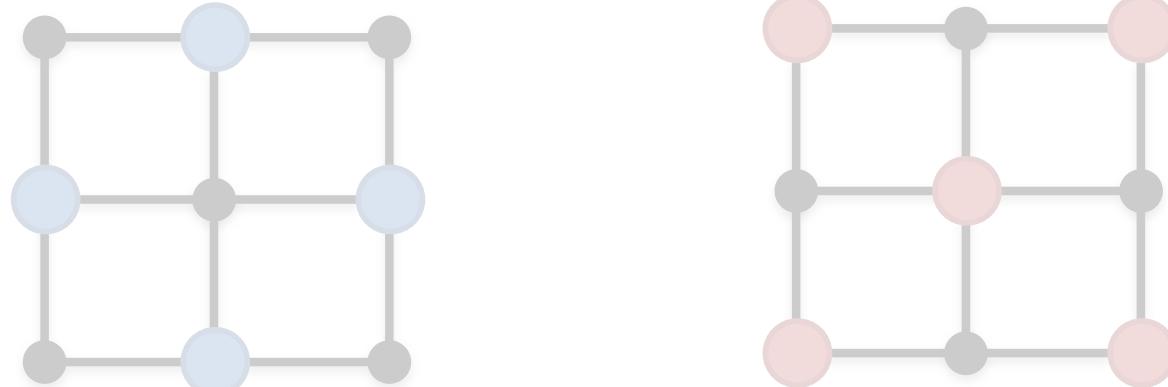


Proof of Ergodicity

- Version 1 (Sequential Consistency):
 - **Chromatic Gibbs Sampler** is *equivalent* to a **Sequential Scan** Gibbs Sampler

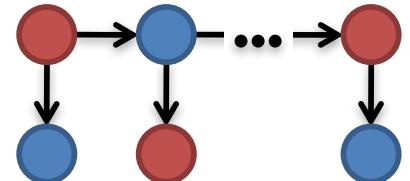
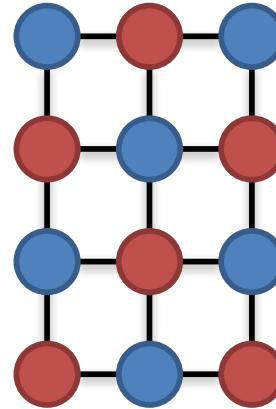
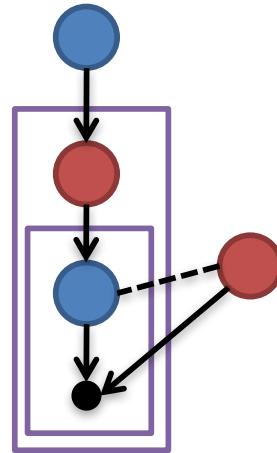
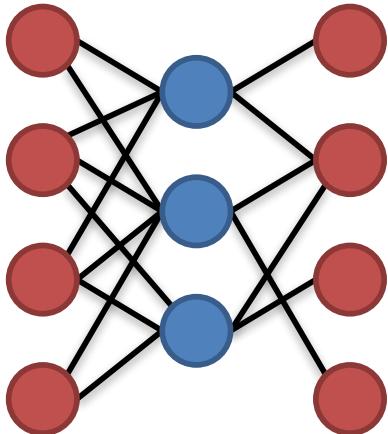


- Version 2 (Probabilistic Interpretation):
 - Variables in same color are **Conditionally Independent** →
Joint Sample is *equivalent* to Parallel Independent Samples



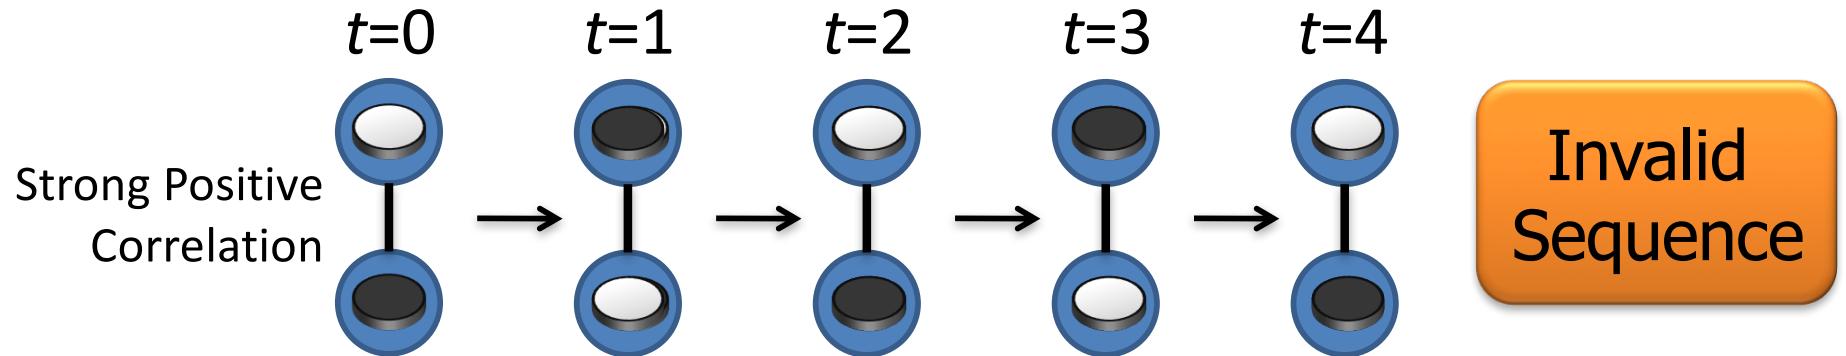
Special Properties of 2-Colorable Models

- Many common models have two colorings

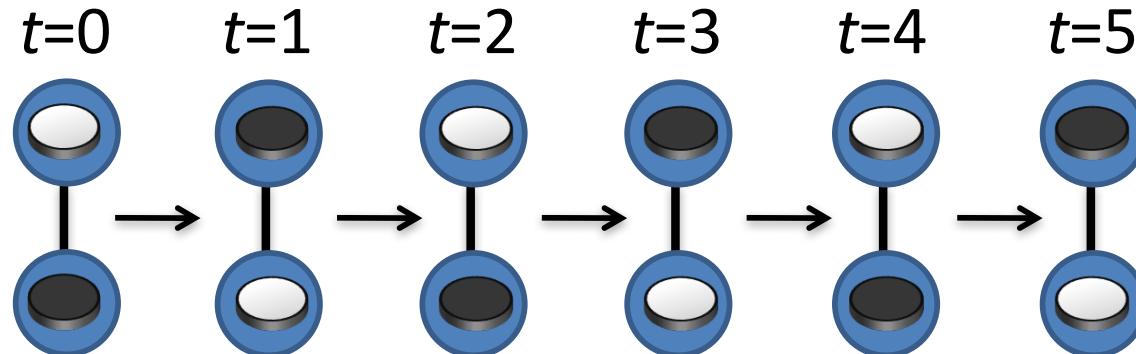


- For the [Incorrect] Synchronous Gibbs Samplers
 - Provide a method to correct the chains
 - Derive the stationary distribution

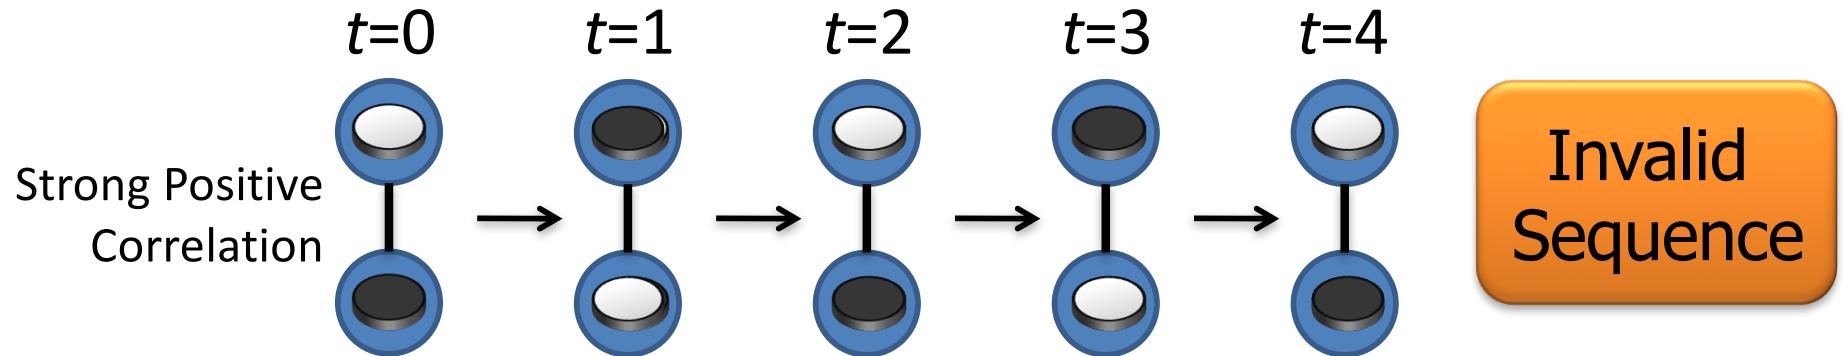
Correcting the *Synchronous* Gibbs Sampler



- We can derive two **valid** chains:

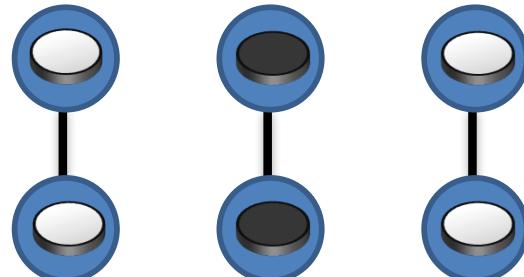


Correcting the *Synchronous* Gibbs Sampler



- We can derive two **valid** chains:

Chain 1



Converges to the
Correct Distribution

Chain 2

- Stationary distribution of **Synchronous Gibbs**:

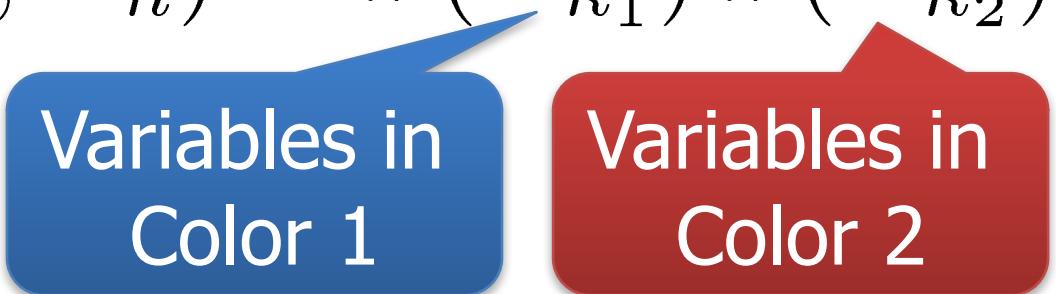
$$\begin{aligned} \mathbf{P}(x') &= \sum_x K(x' | x) \pi(x_{\kappa_1}) \pi(x_{\kappa_2}) \\ &= \sum_{x_{\kappa_1}} \sum_{x_{\kappa_2}} \pi(x'_{\kappa_1} | x_{\kappa_2}) \pi(x'_{\kappa_2} | x_{\kappa_1}) \pi(x_{\kappa_1}) \pi(x_{\kappa_2}) \\ &= \sum_{x_{\kappa_1}} \sum_{x_{\kappa_2}} \pi(x_{\kappa_1}, x'_{\kappa_2}) \pi(x'_{\kappa_1}, x_{\kappa_2}) \\ &= \pi(x'_{\kappa_1}) \pi(x'_{\kappa_2}) \end{aligned}$$

Variables in
Color 1

Variables in
Color 2

- Stationary distribution of **Synchronous Gibbs**

$$\mathbf{P}_{\text{sync}}(X_1, \dots, X_n) = \pi(X_{\kappa_1}) \pi(X_{\kappa_2})$$



Variables in
Color 1

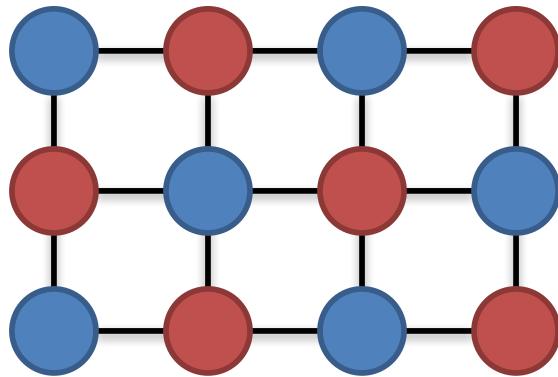
Variables in
Color 2

- **Corollary:** Synchronous Gibbs sampler is **correct** for single variable marginals.

$$\lim_{m \rightarrow \infty} \frac{1}{m} \sum_{t=1}^m h(x_i^{(t)}) \xrightarrow{a.s.} \mathbf{E}[h(X_i)]$$

From Colored Fields to Thin Junction Trees

Chromatic Gibbs Sampler



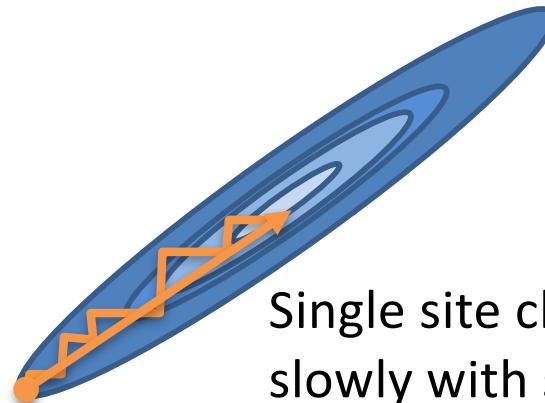
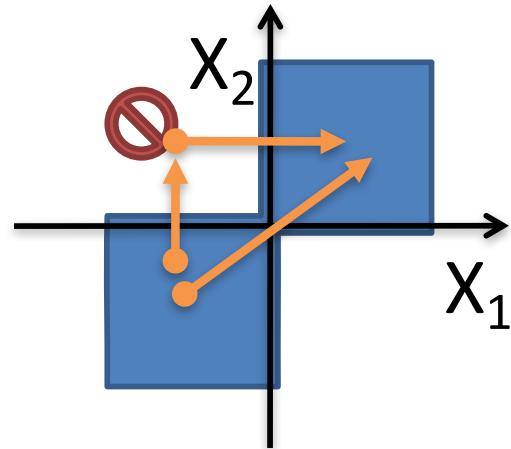
- Ideal for:
 - **Rapid mixing** models
 - Conditional structure does not admit **Splash**

Slowly Mixing Models

- Ideal for: ?
 - Slowly mixing models
 - Conditional structure admits Splash
 - Discrete models

Models With Strong Dependencies

- **Single variable** Gibbs updates tend to mix **slowly**:

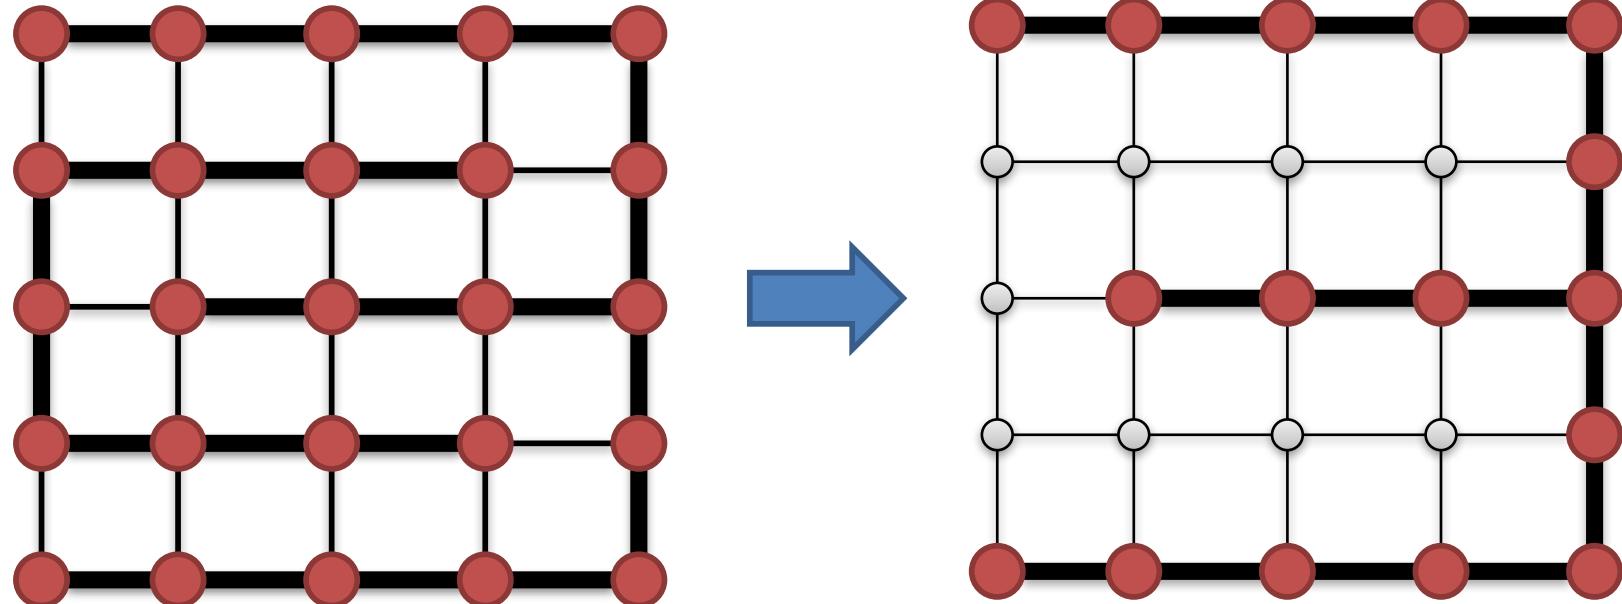


Single site changes move slowly with strong correlation.

- Ideally we would like to draw joint samples.
 - Blocking

Blocking Gibbs Sampler

- Based on the papers:
 1. Jensen et al., *Blocking Gibbs Sampling for Linkage Analysis in Large Pedigrees with Many Loops.* TR 1996
 2. Hamze et al., *From Fields to Trees.* UAI 2004.

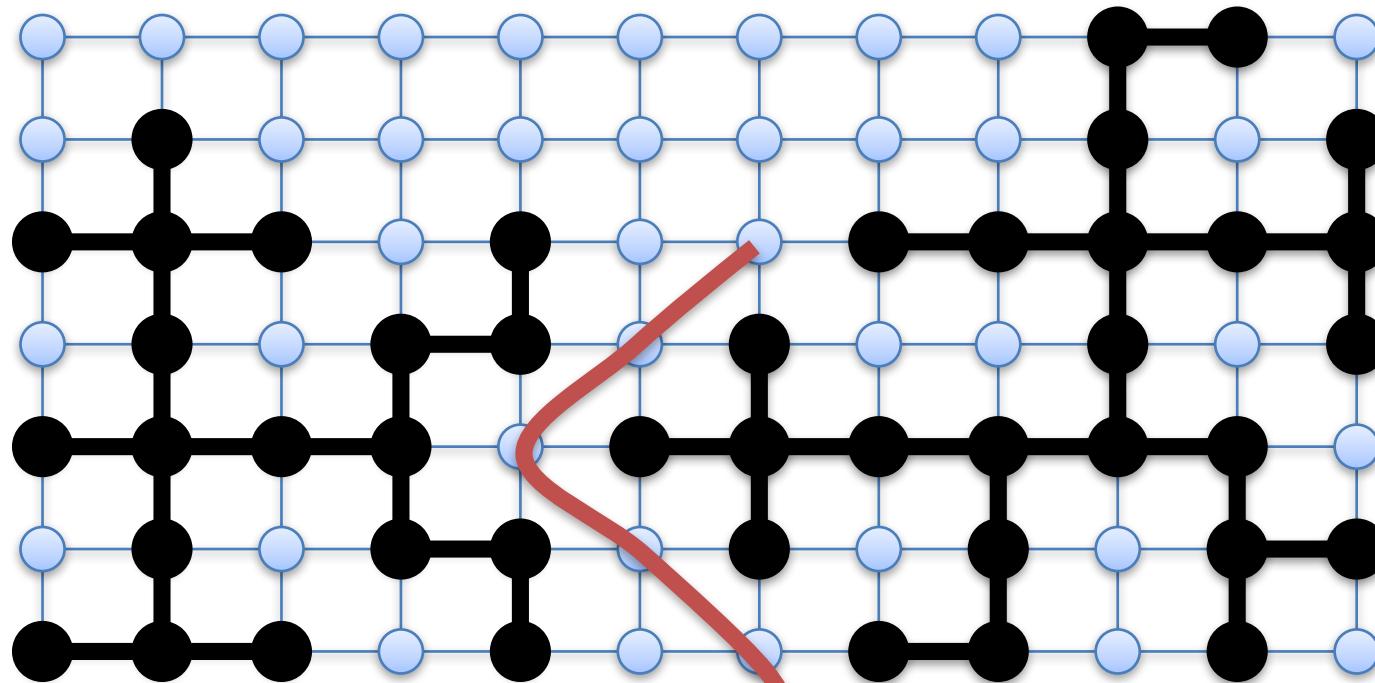


Splash Gibbs Sampler

An asynchronous Gibbs Sampler that adaptively addresses strong dependencies.

Splash Gibbs Sampler

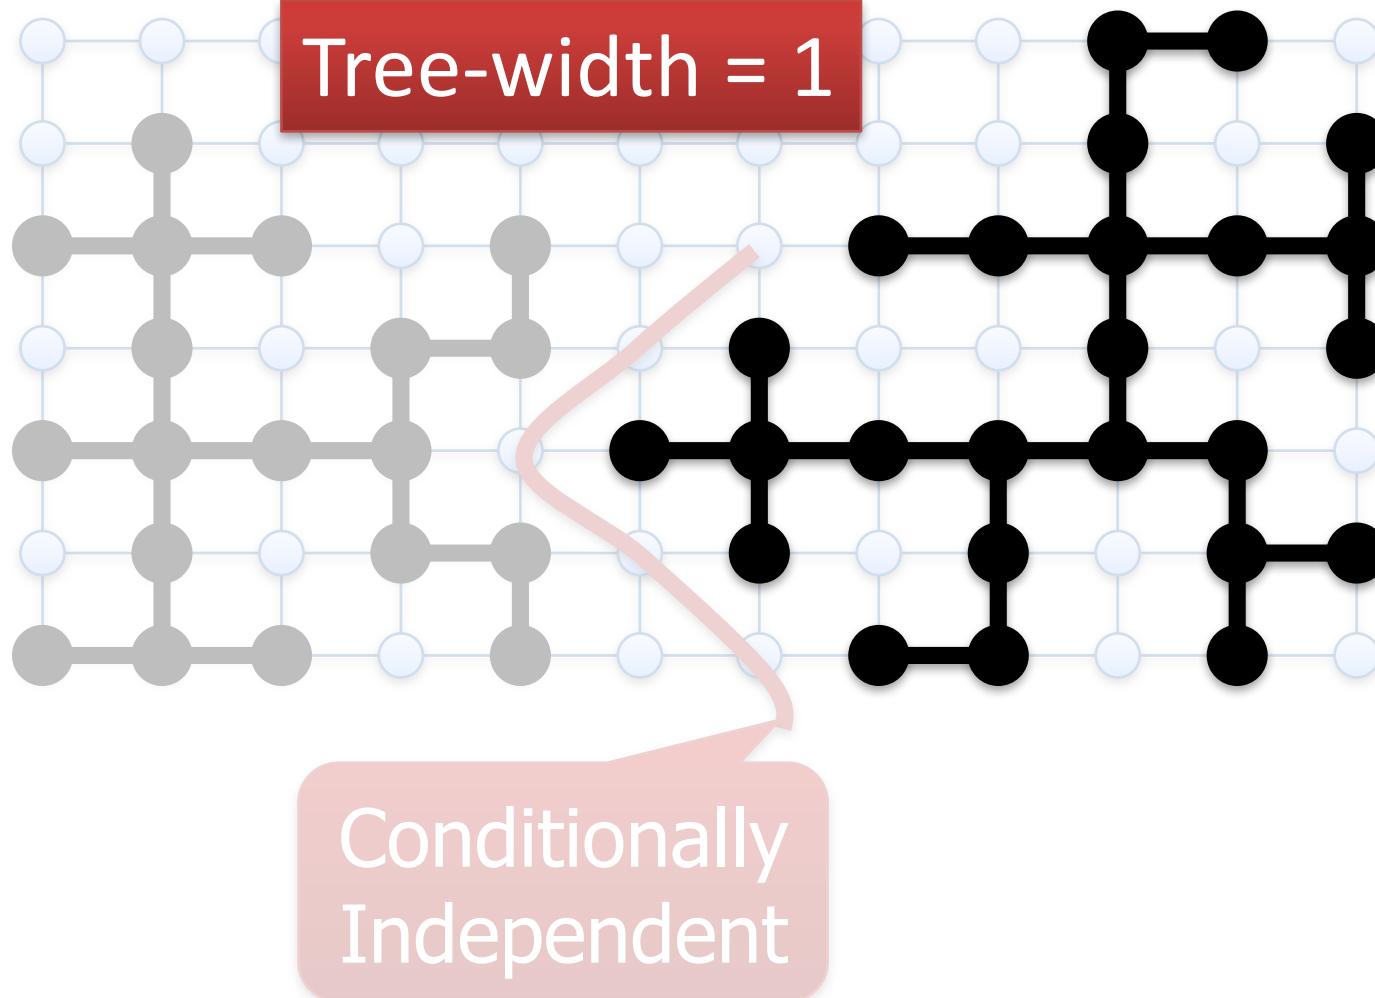
- **Step 1:** Grow multiple Splashes in parallel:



Conditionally
Independent

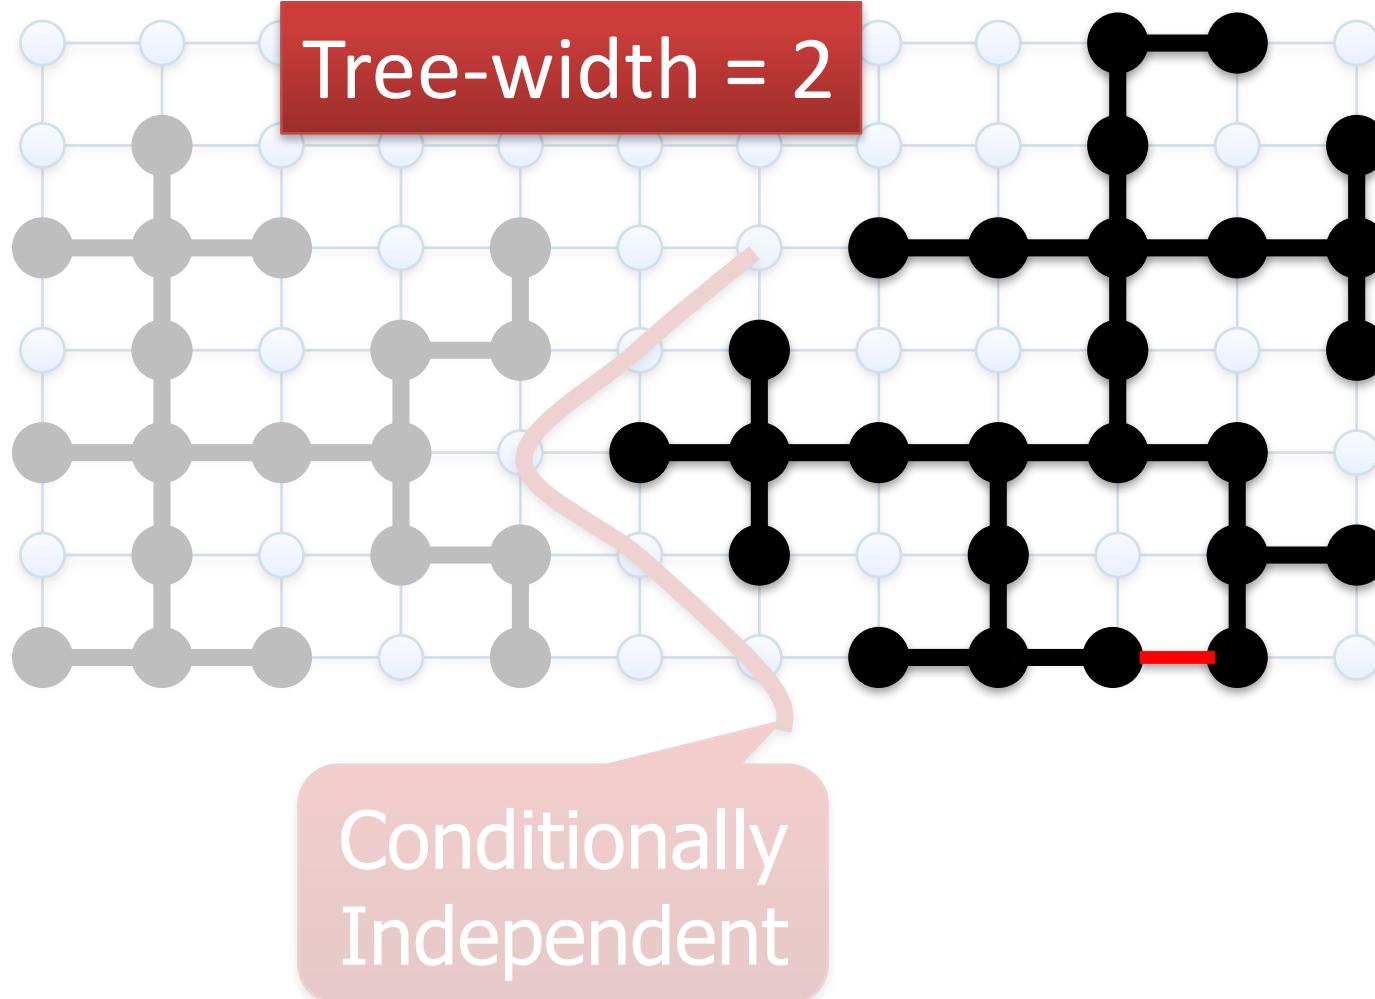
Splash Gibbs Sampler

- **Step 1:** Grow multiple Splashes in parallel:



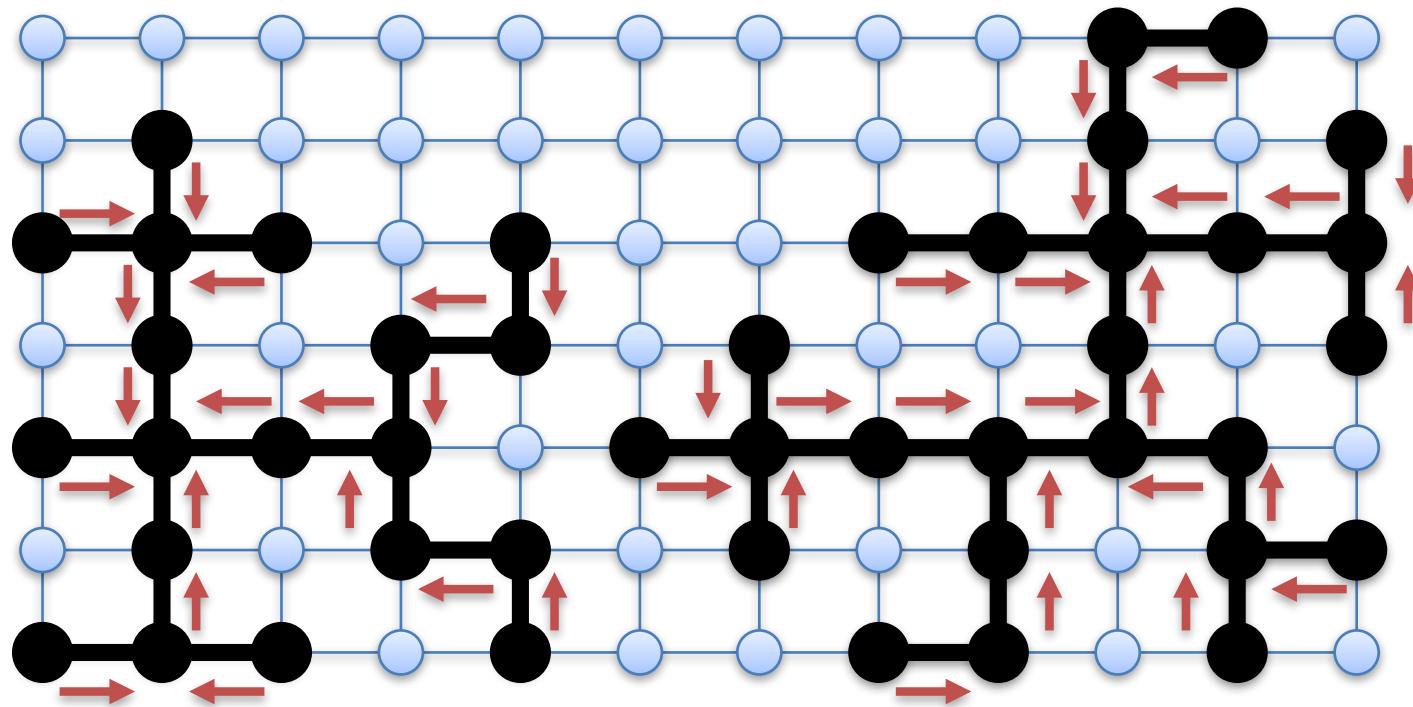
Splash Gibbs Sampler

- **Step 1:** Grow multiple Splashes in parallel:



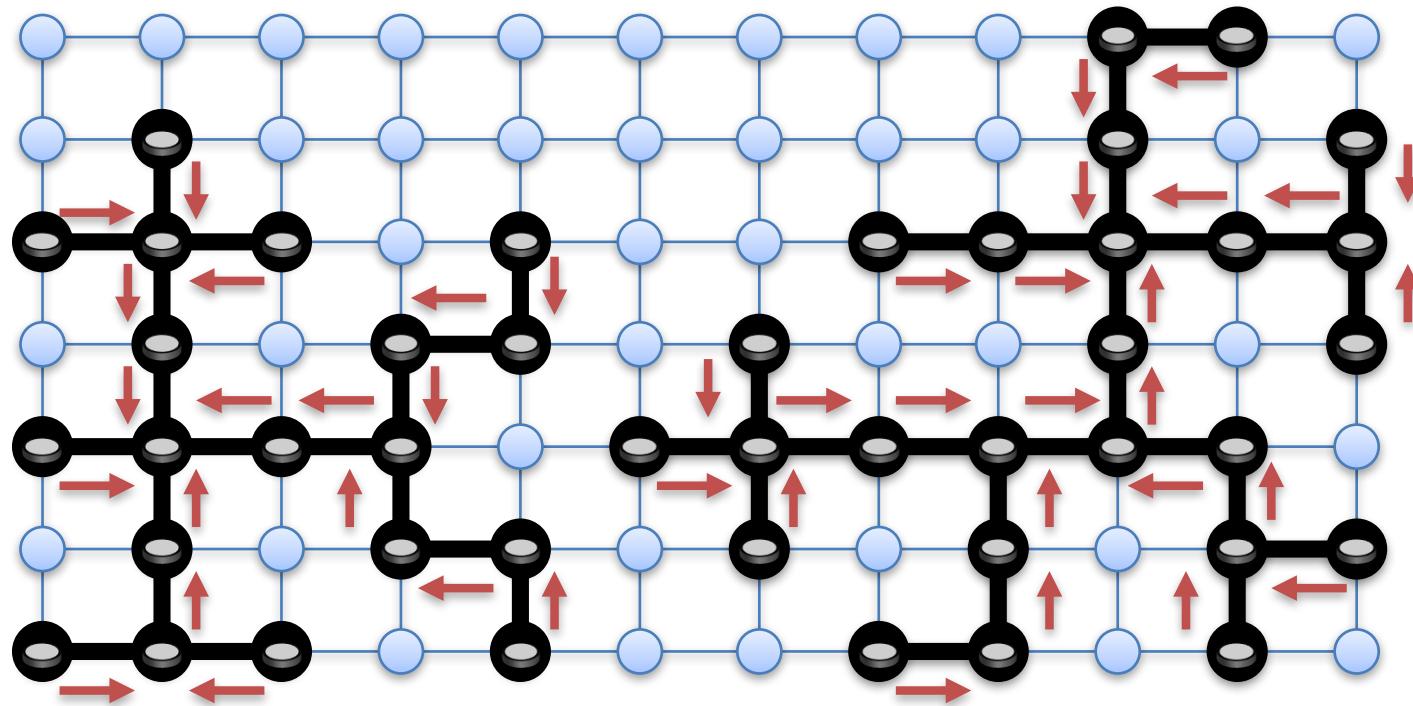
Splash Gibbs Sampler

- Step 2: Calibrate the trees in parallel



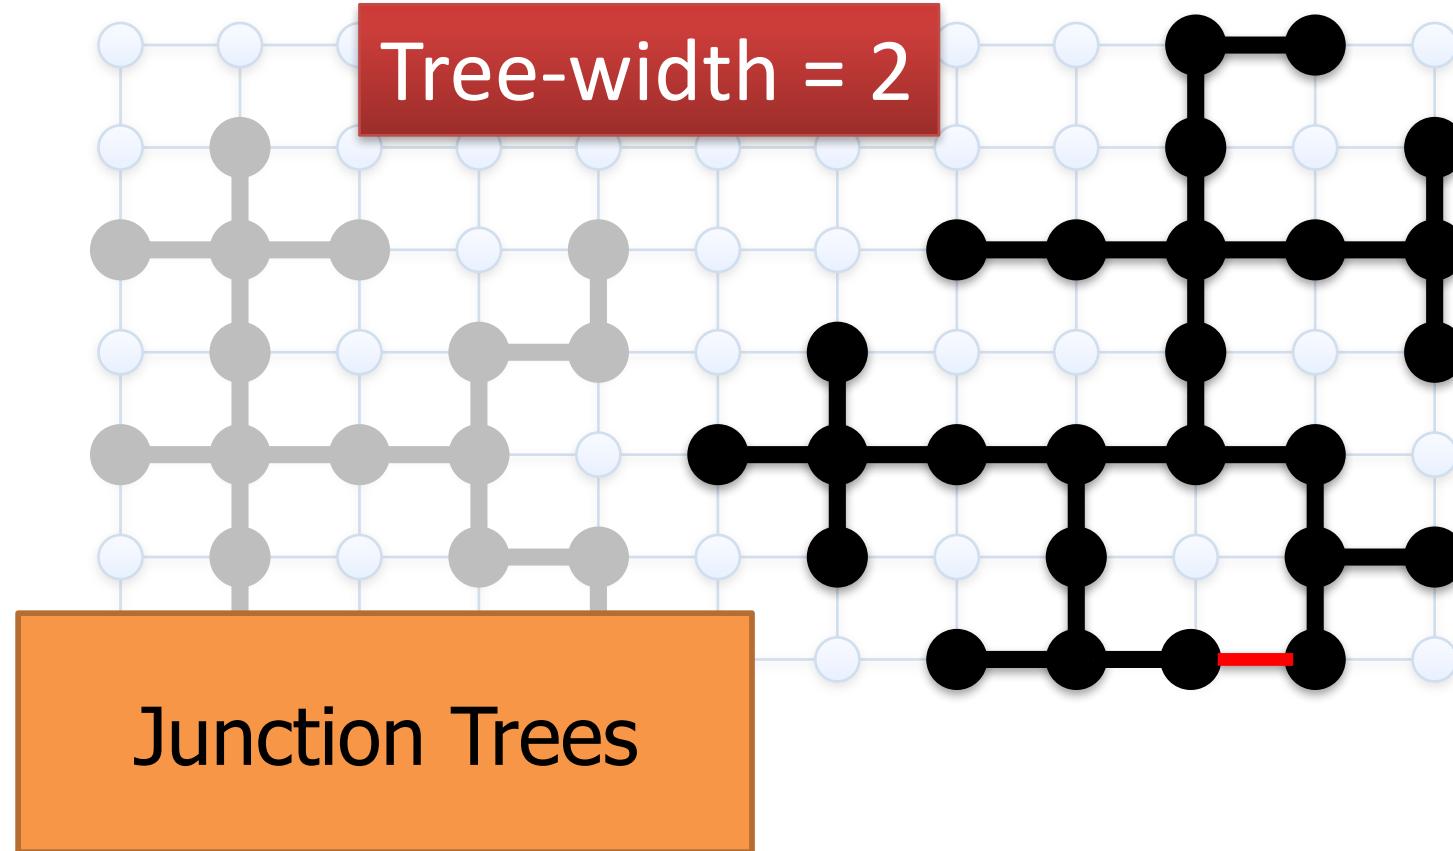
Splash Gibbs Sampler

- **Step 3:** Sample trees in parallel



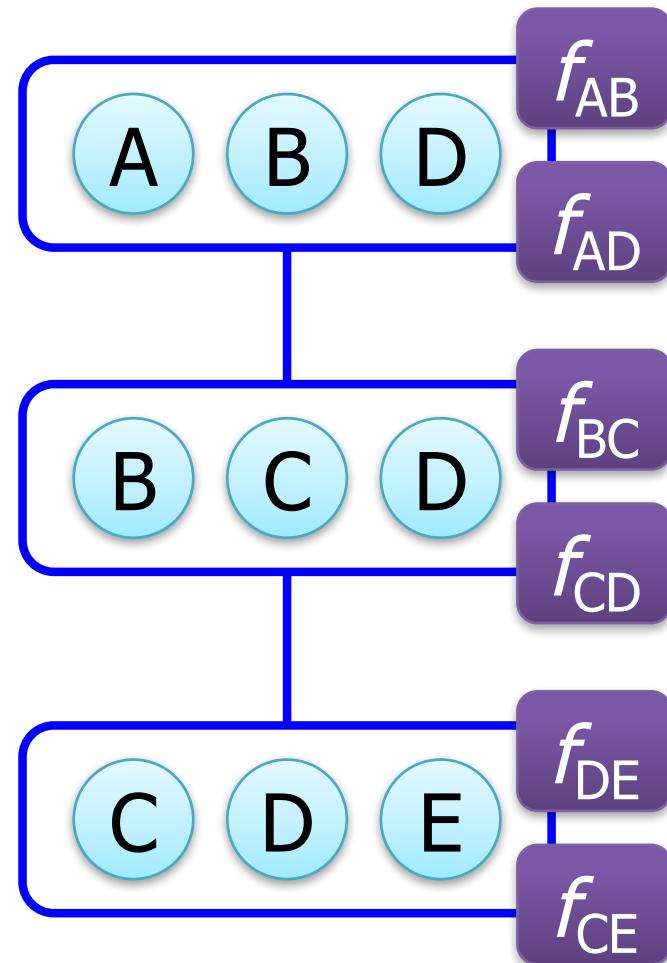
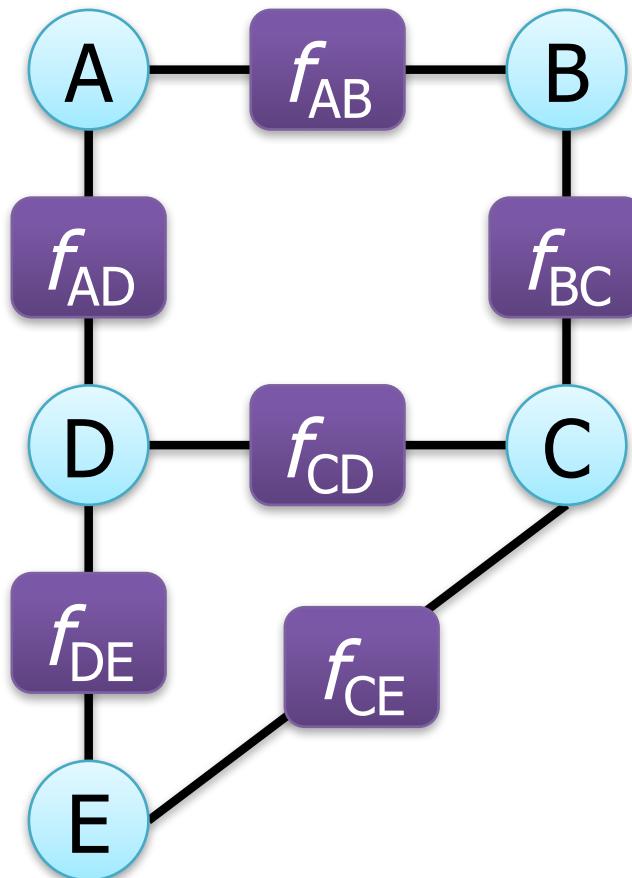
Higher Treewidth Splashes

- Recall:



Junction Trees

- Data structure used for exact inference in loopy graphical models



Tree-width = 2

Splash Thin Junction Tree

- Parallel Splash Junction Tree Algorithm

- Construct multiple conditionally independent thin (bounded treewidth) junction trees **Splashes**

- Sequential junction tree extension

- Calibrate the each thin junction tree in parallel

- Parallel belief propagation

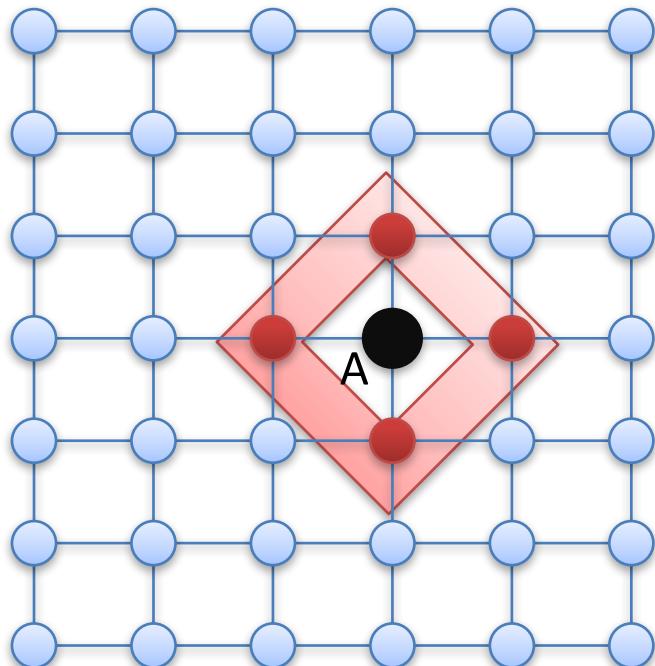
- Exact backward sampling

- Parallel exact sampling

Splash generation

- Frontier extension algorithm:

Markov Random Field



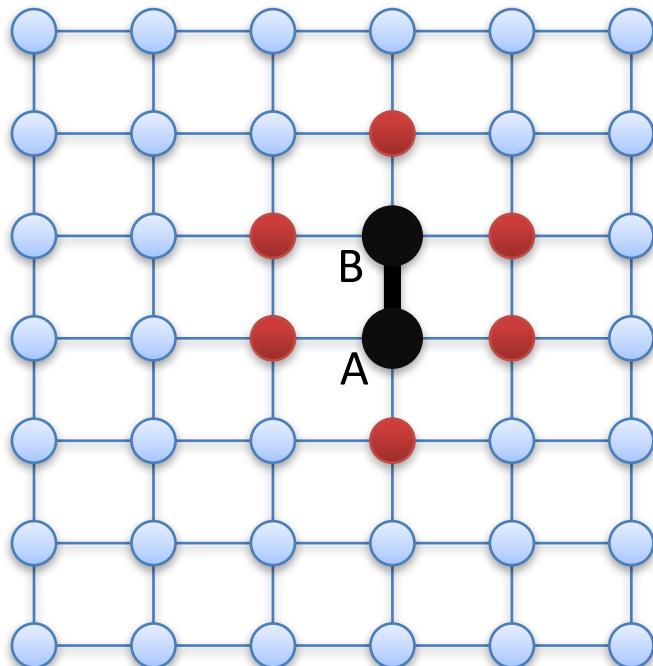
Corresponding Junction tree



Splash generation

- Frontier extension algorithm:

Markov Random Field



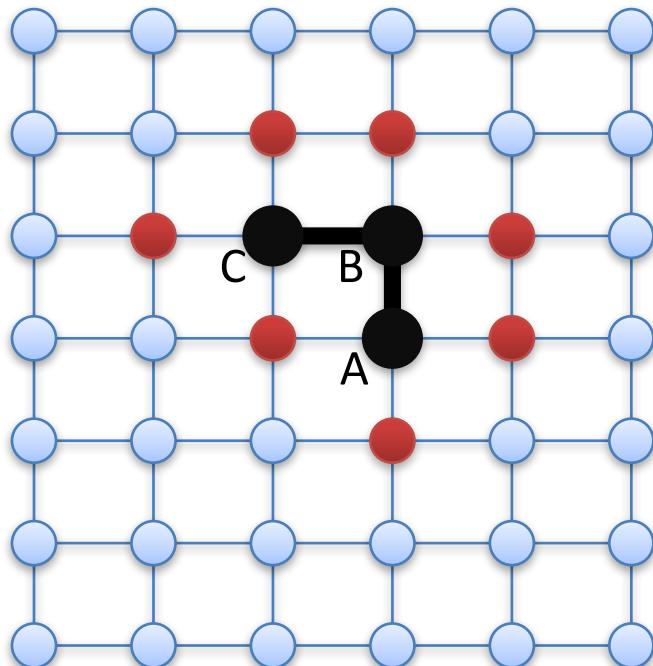
Corresponding Junction tree



Splash generation

- Frontier extension algorithm:

Markov Random Field



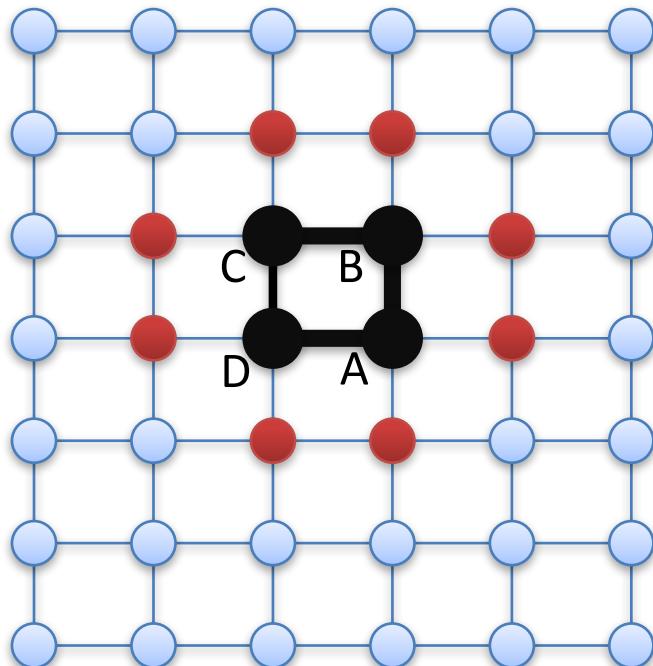
Corresponding Junction tree



Splash generation

- Frontier extension algorithm:

Markov Random Field



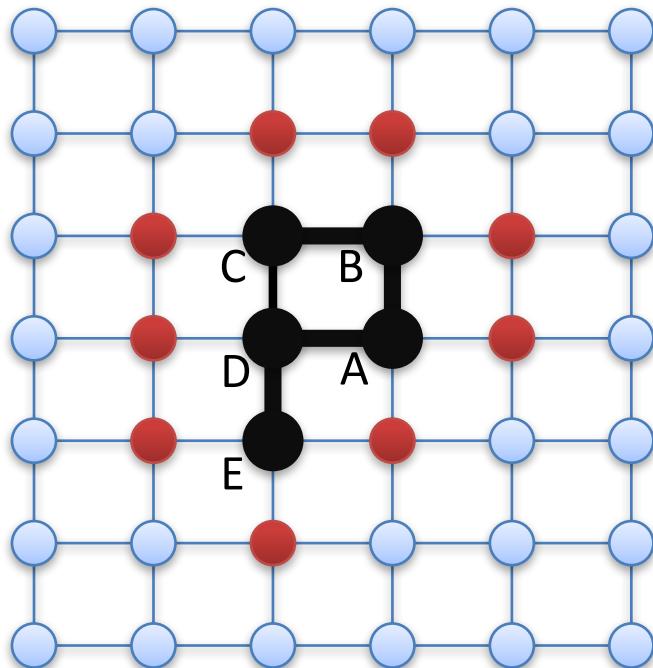
Corresponding Junction tree



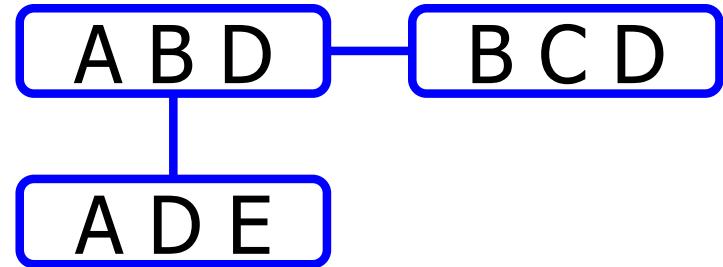
Splash generation

- Frontier extension algorithm:

Markov Random Field



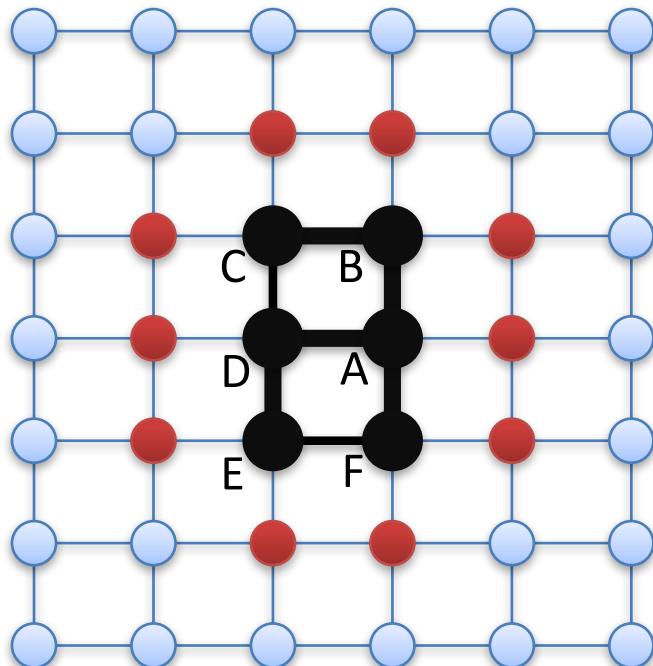
Corresponding Junction tree



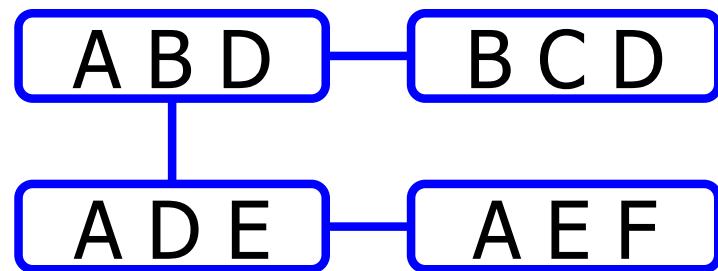
Splash generation

- Frontier extension algorithm:

Markov Random Field



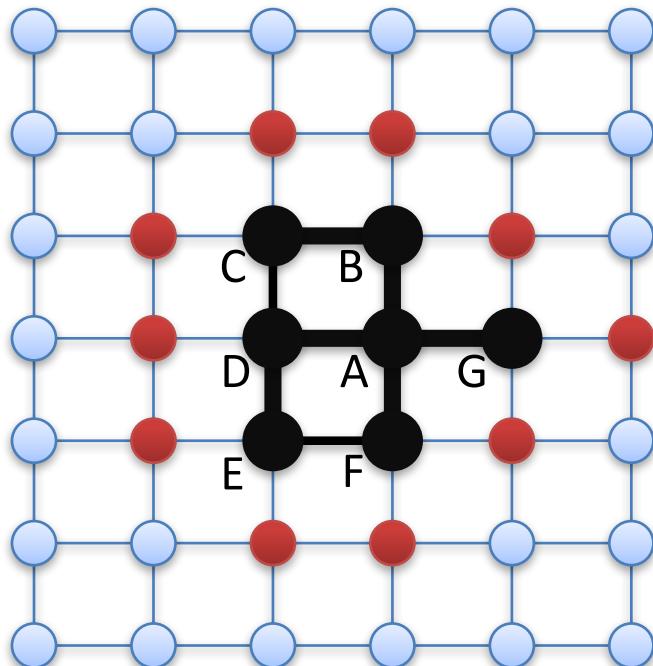
Corresponding Junction tree



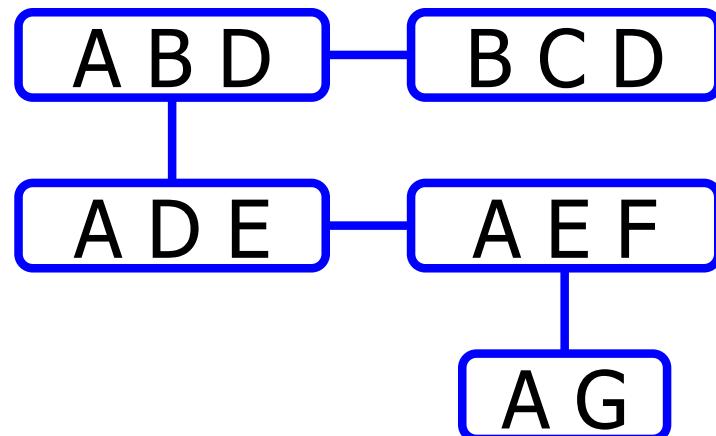
Splash generation

- Frontier extension algorithm:

Markov Random Field



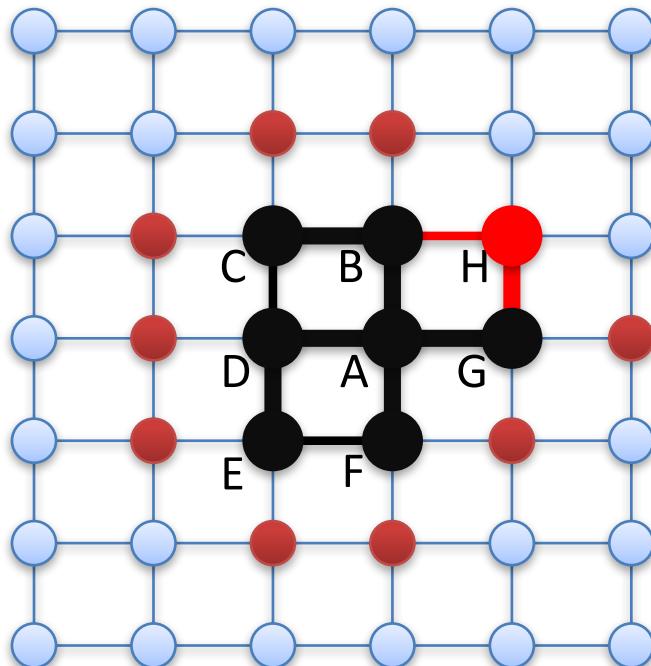
Corresponding Junction tree



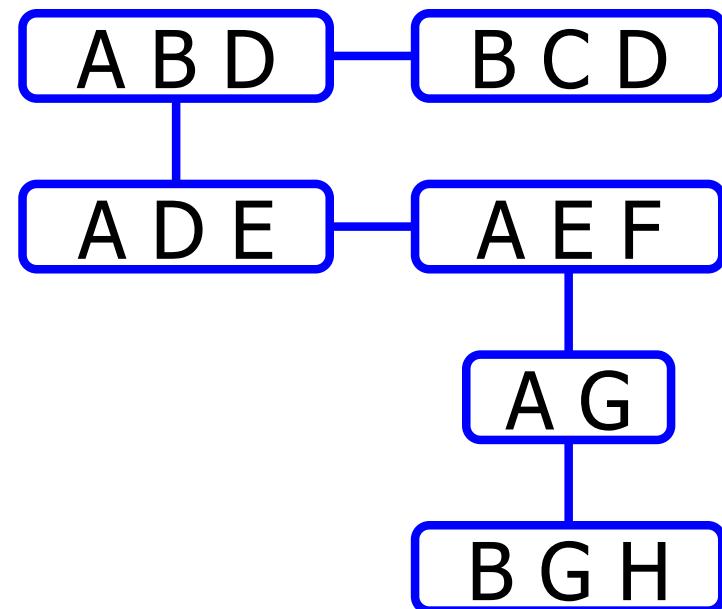
Splash generation

- Frontier extension algorithm:

Markov Random Field



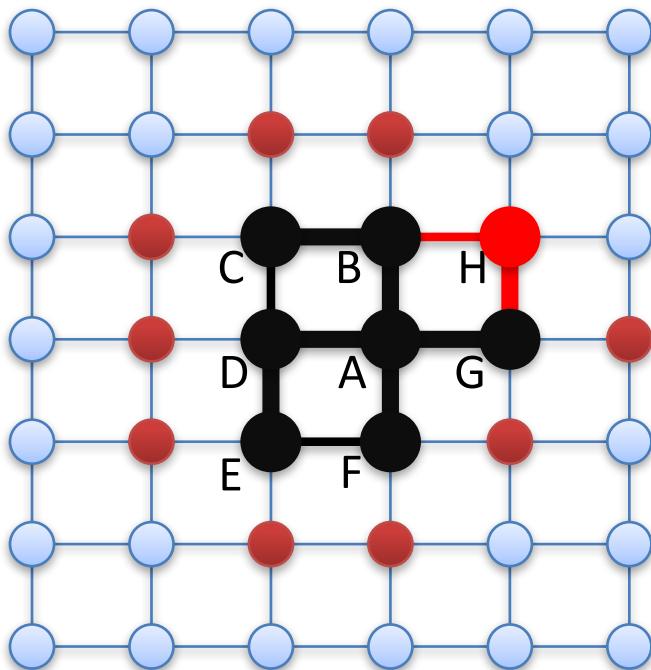
Corresponding Junction tree



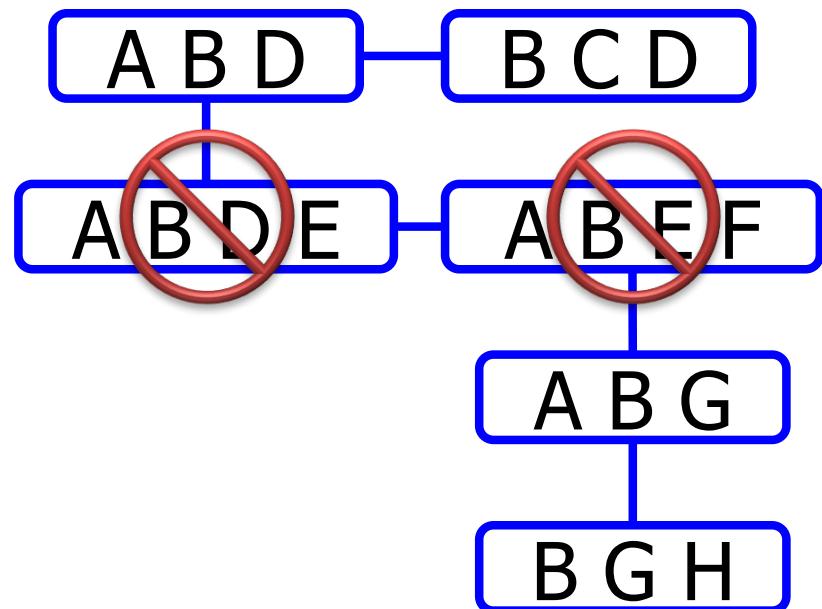
Splash generation

- Frontier extension algorithm:

Markov Random Field



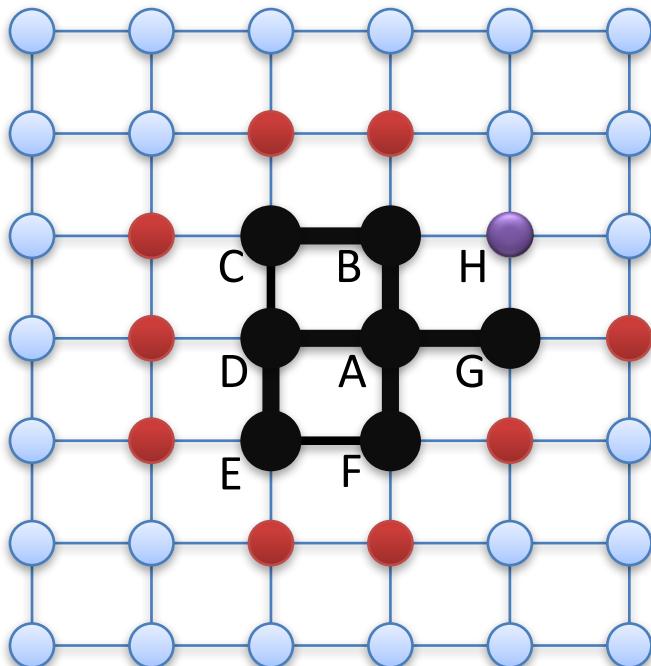
Corresponding Junction tree



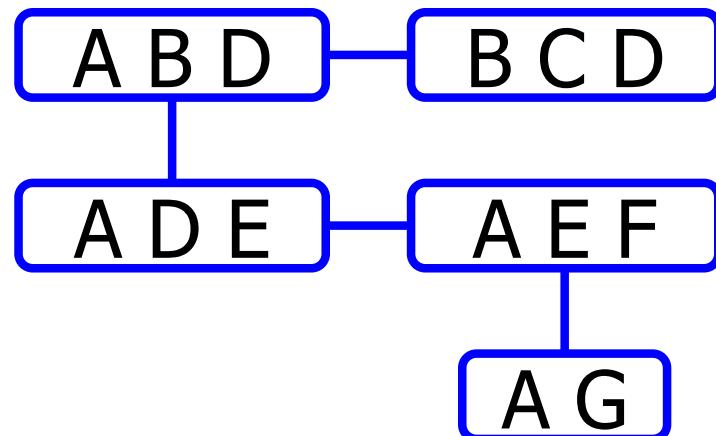
Splash generation

- Frontier extension algorithm:

Markov Random Field



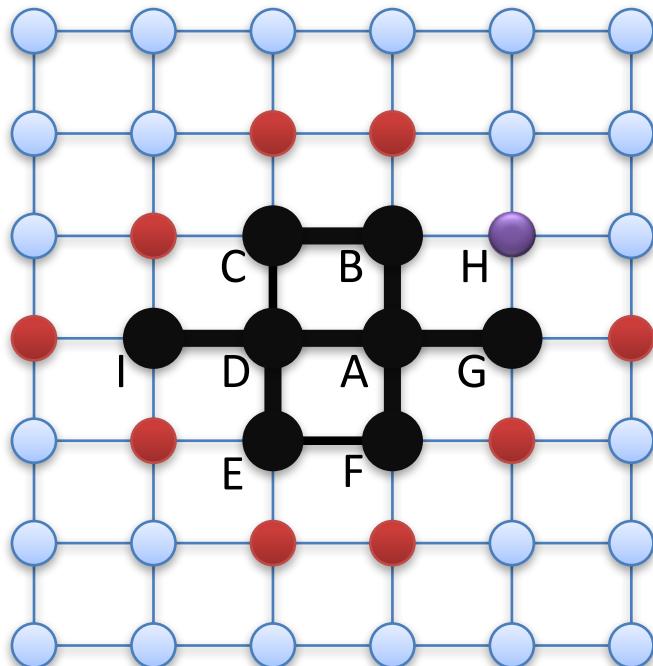
Corresponding Junction tree



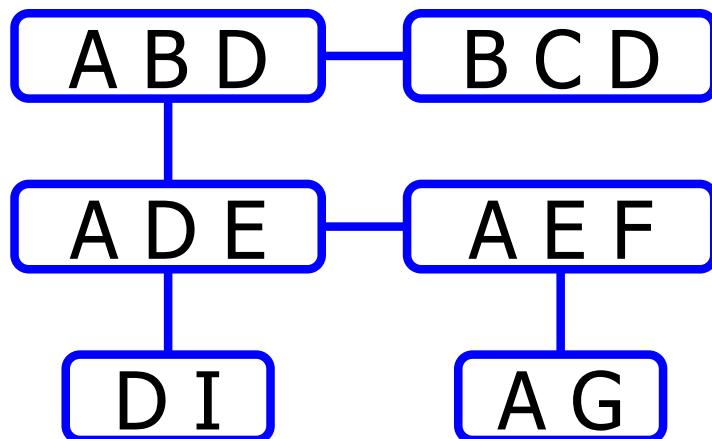
Splash generation

- Frontier extension algorithm:

Markov Random Field

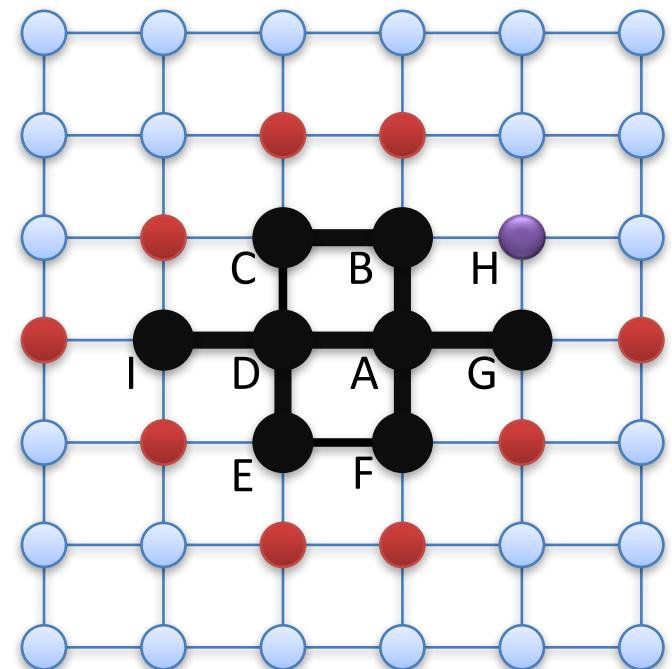


Corresponding Junction tree



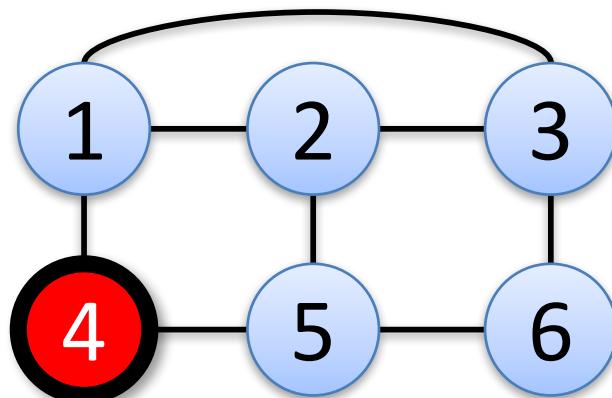
Splash generation

- Challenge:
 - Efficiently reject vertices that violate treewidth constraint
 - Efficiently extend the junction tree
 - Choosing the next vertex
- Solution Splash Junction Trees:
 - Variable elimination with reverse visit ordering
 - I,G,F,E,D,C,B,A
 - Add new clique and update RIP
 - If a clique is created which exceeds treewidth terminate extension
 - Adaptive prioritize boundary



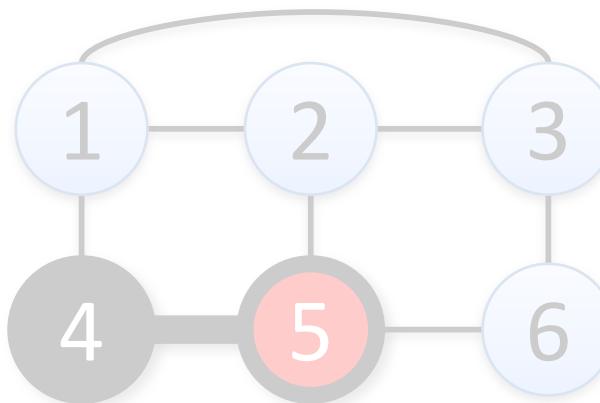
Incremental Junction Trees

- First 3 Rounds:



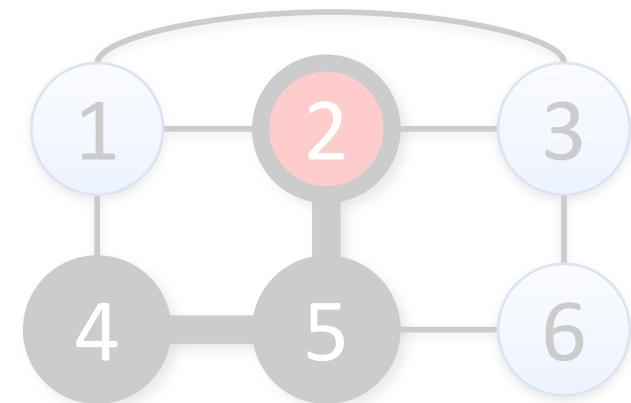
Junction Tree: $\underline{4}$

Elim. Order: {4}



$\underline{4} \leftarrow \underline{4,5}$

{5,4}

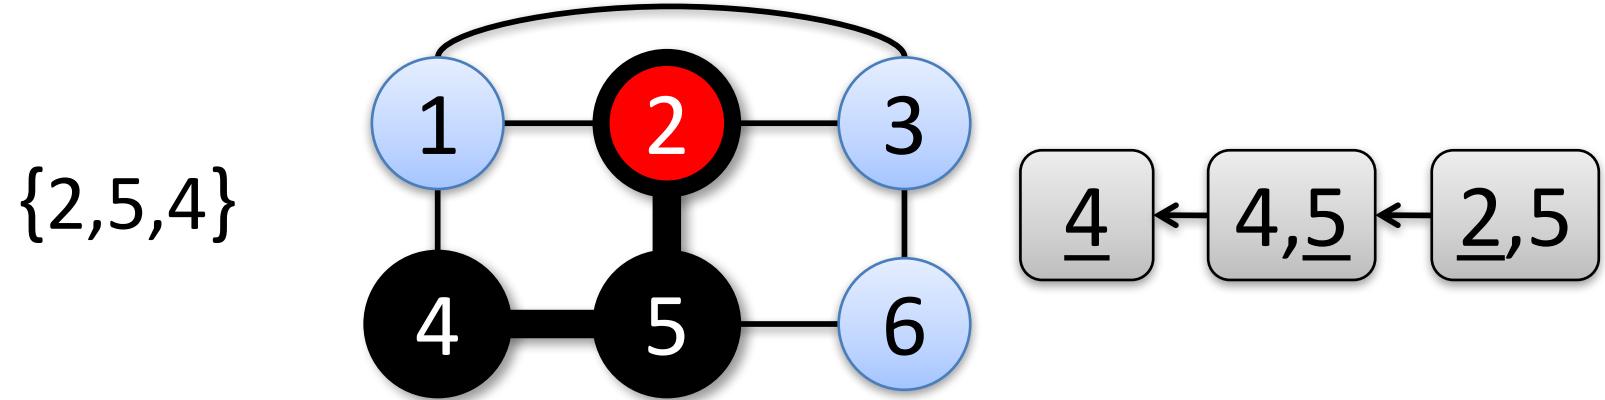


$\underline{4} \leftarrow \underline{4,5} \leftarrow \underline{2,5}$

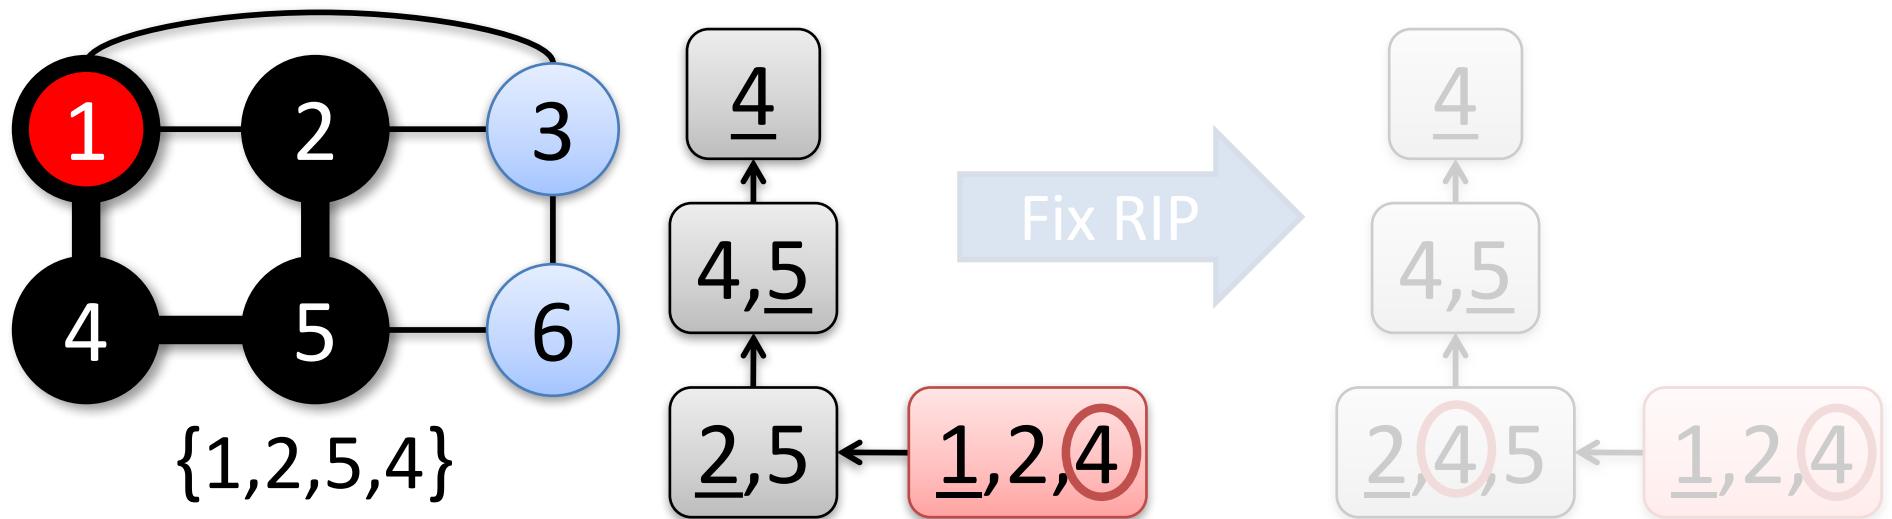
{2,5,4}

Incremental Junction Trees

- Result of third round:



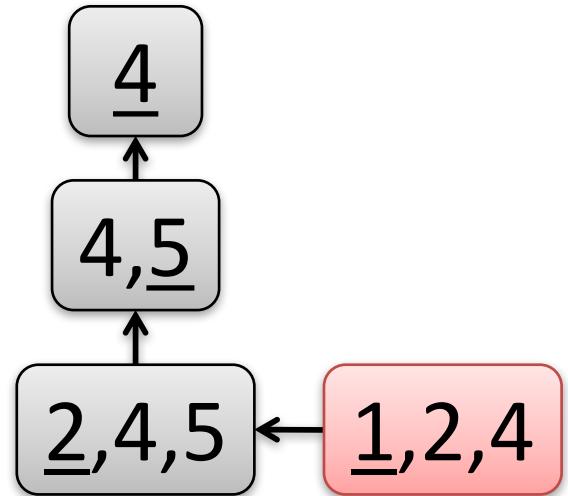
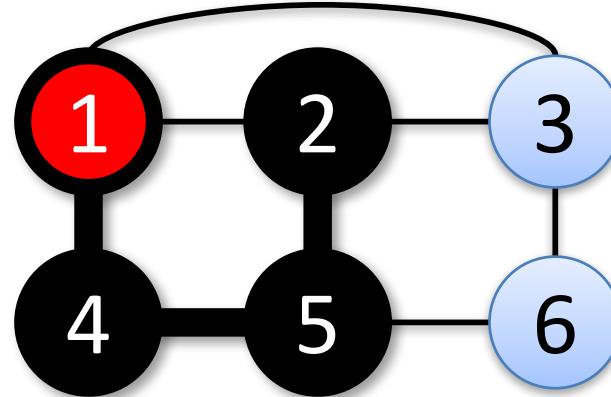
- Fourth round:



Incremental Junction Trees

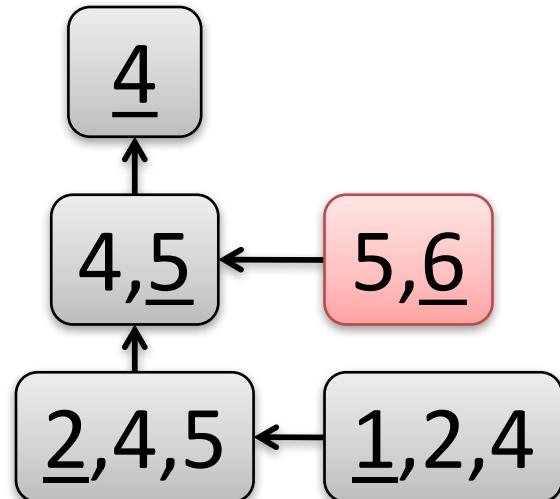
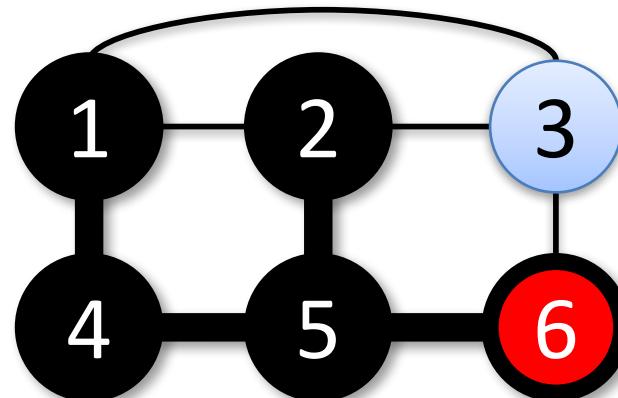
- Results from 4th round:

{1,2,5,4}



- 5th Round:

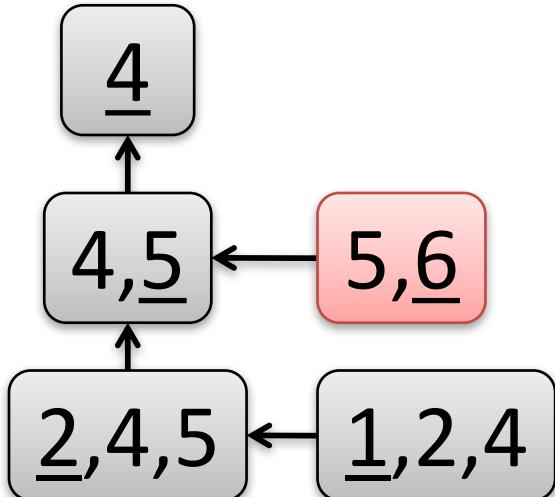
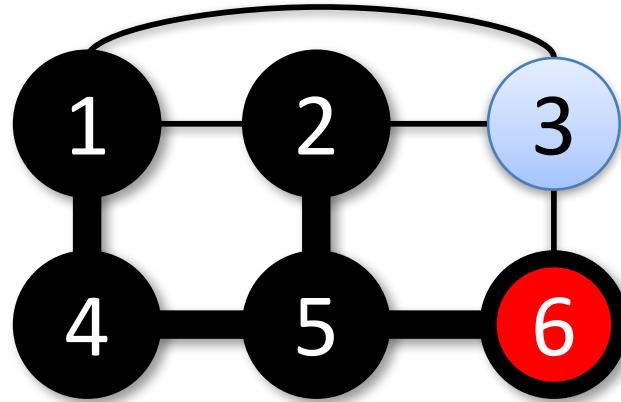
{6,1,2,5,4}



Incremental Junction Trees

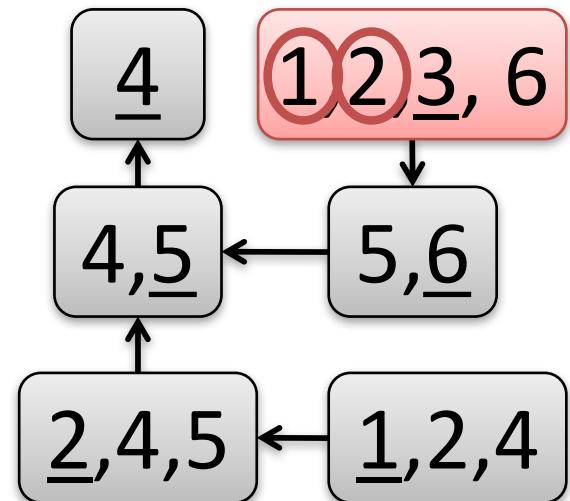
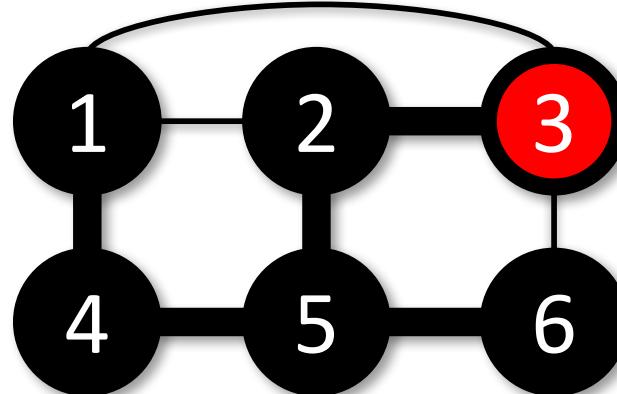
- Results from 5th round:

$\{6,1,2,5,4\}$



- 6th Round:

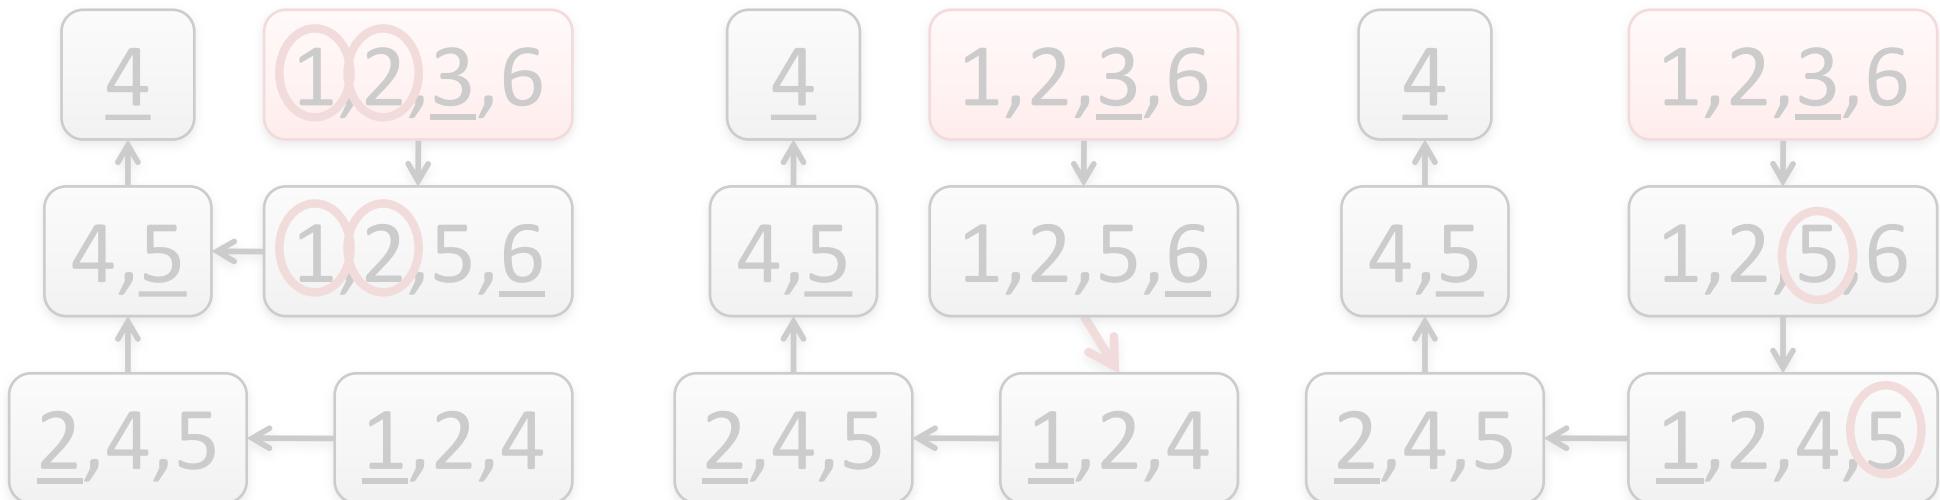
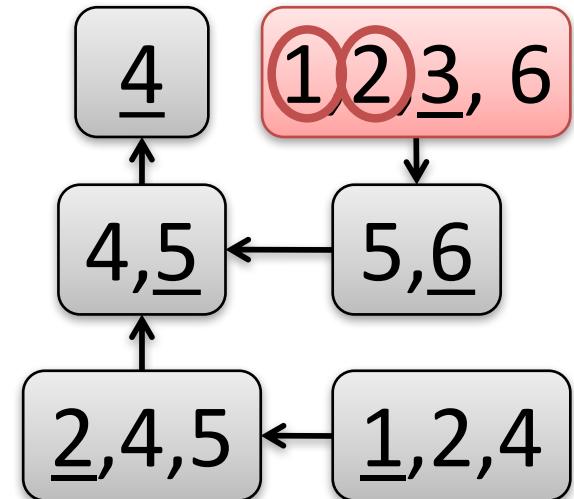
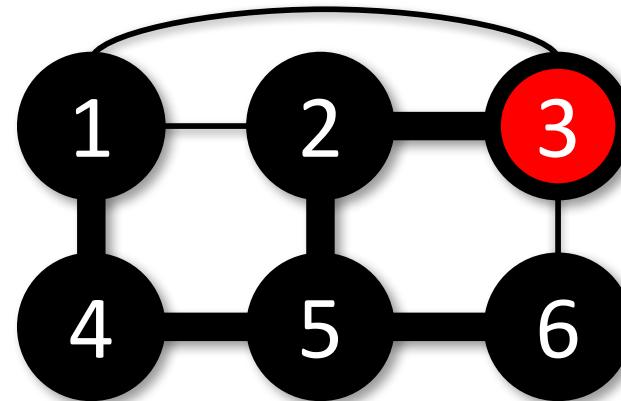
$\{3,6,1,2,5,4\}$



Incremental Junction Trees

- Finishing 6th round:

$\{3,6,1,2,5,4\}$



Algorithm Block [Skip]

Input: The original junction tree $(\mathcal{C}, E) = \mathbf{J}_{\mathcal{S}}$.

Input: The variable X_i to add to $\mathbf{J}_{\mathcal{S}}$

Output: $\mathbf{J}_{\mathcal{S}+i}$

Define : \mathbf{C}_u as the clique created by eliminating $u \in \mathcal{S}$

Define : $v[\mathbf{C}] \in \mathcal{S}$ as the variable eliminated when creating \mathbf{C}

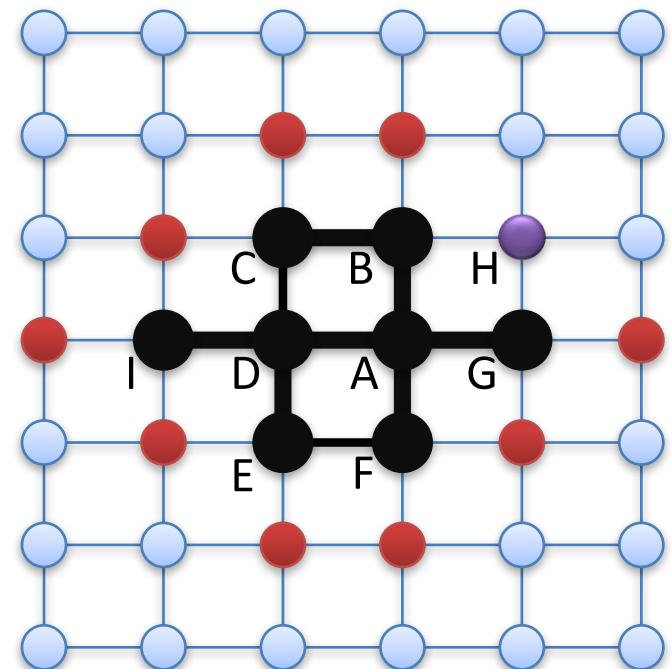
Define : $t[v]$ as the time $v \in \mathcal{S}$ was added to \mathcal{S}

Define : $P[v] \in \mathcal{N}_v \cap \mathcal{S}$ as the next neighbor of v to be eliminated.

```
1  $\mathbf{C}_i \leftarrow (\mathcal{N}_i \cap \mathcal{S}) \cup \{i\}$ 
2  $P[i] \leftarrow \arg \max_{v \in \mathbf{C}_i \setminus \{i\}} t[v]$ 
// ----- Repair RIP -----
3  $\mathcal{R} \leftarrow \mathbf{C}_i \setminus \{i\}$  // RIP Set
4  $v \leftarrow P[i]$ 
5 while  $|\mathcal{R}| > 0$  do
6    $\mathbf{C}_v \leftarrow \mathbf{C}_v \cup \mathcal{R}$  // Add variables to parent
7    $w \leftarrow \arg \max_{w \in \mathbf{C}_v \setminus \{v\}} t[w]$  // Find new parent
8   if  $w = P[v]$  then
9      $\mathcal{R} \leftarrow (\mathcal{R} \setminus \mathbf{C}_i) \setminus \{i\}$ 
10  else
11     $\mathcal{R} \leftarrow (\mathcal{R} \cup \mathbf{C}_i) \setminus \{i\}$ 
12     $P[v] \leftarrow w$  // New parent
13   $v \leftarrow P[v]$  // Move upwards
```

Splash generation

- Challenge:
 - Efficiently reject vertices that violate treewidth constraint
 - Efficiently extend the junction tree
 - Choosing the next vertex
- Solution Splash Junction Trees:
 - Variable elimination with reverse visit ordering
 - I,G,F,E,D,C,B,A
 - Add new clique and update RIP
 - If a clique is created which exceeds treewidth terminate extension
 - Adaptive prioritize boundary



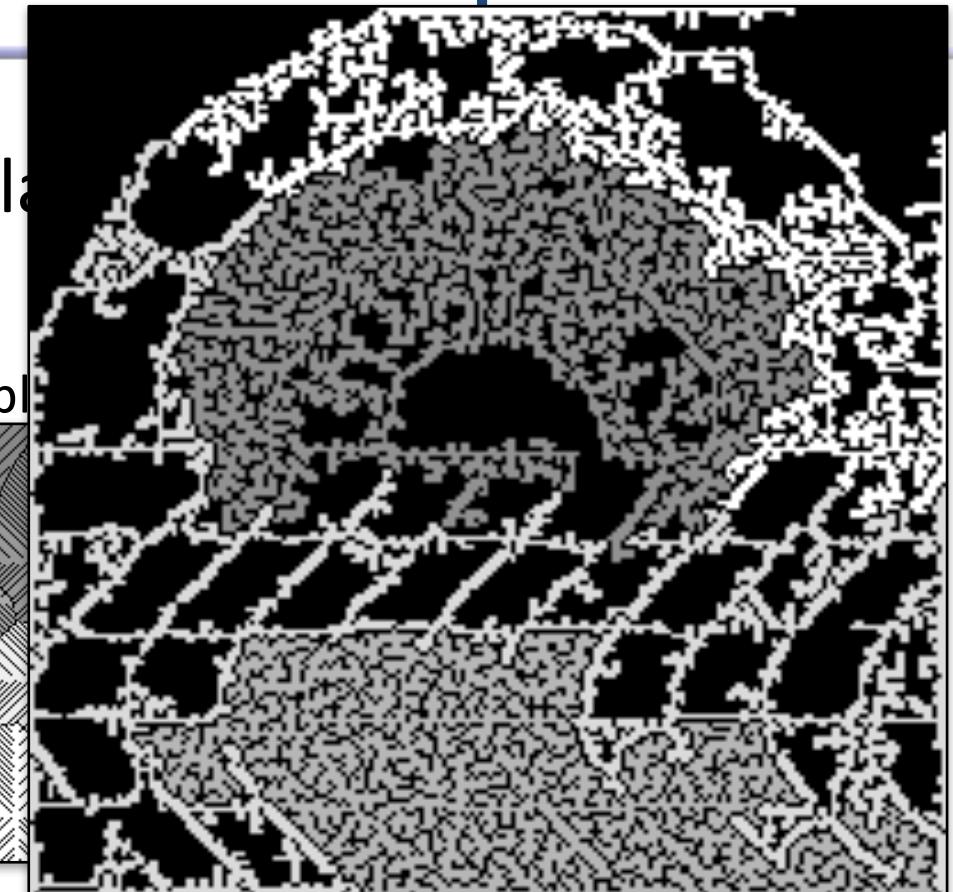
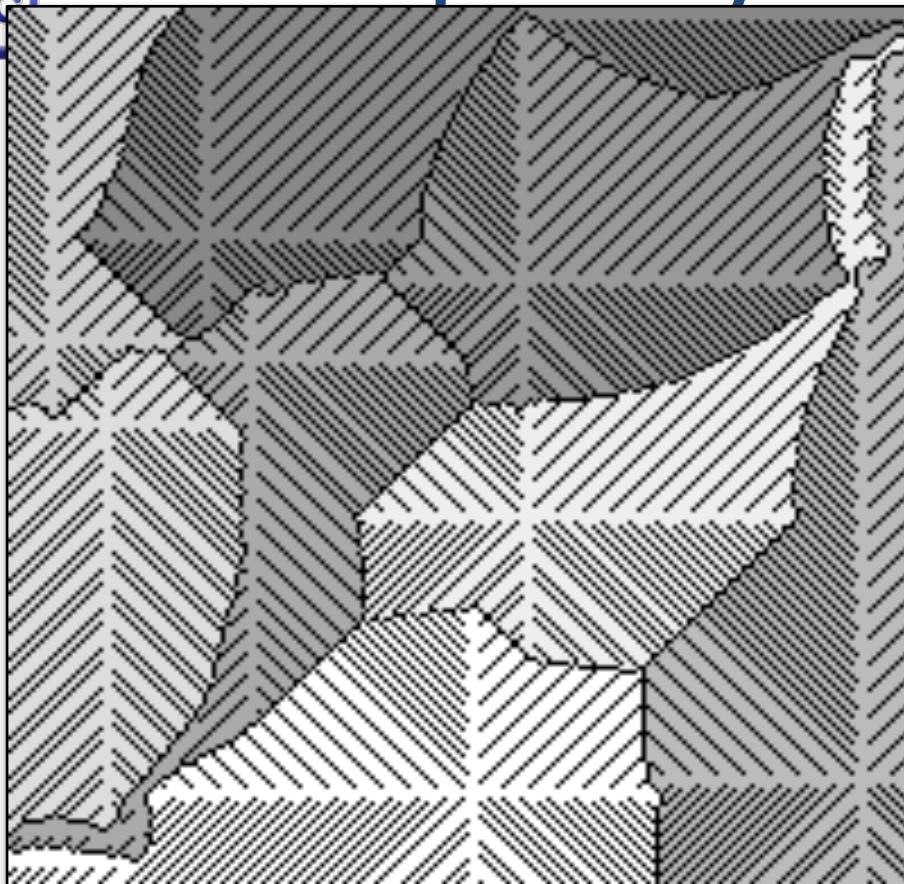
Adaptive Vertex Priorities

- Assign priorities to boundary vertices:

$$s[X_v] = \left\| \log \frac{\sum_x \pi(X_{\mathcal{S}}, X_v = x \mid X_{-\mathcal{S}} = x_{-\mathcal{S}}^{(t)})}{\pi(X_{\mathcal{S}} \mid X_v = x_v^{(t)}, X_{-\mathcal{S}} = x_{-\mathcal{S}}^{(t)})} \right\|_1$$

- Can be computed using only factors that depend on X_v
- Based on current sample
- Captures difference between marginalizing out the variable (in Splash) fixing its assignment (out of Splash)
- Exponential in treewidth
- Could consider other metrics ...

Adaptively Prioritized Splashes



- Provably converges to the correct distribution
 - Requires vanishing adaptation
 - Identify a bug in the Levine & Casella seminal work in adaptive random scan

Experiments

- Implemented using GraphLab
 - Treewidth = 1 :
 - Parallel tree construction, calibration, and sampling
 - No incremental junction trees needed
 - Treewidth > 1 :
 - Sequential tree construction (use multiple Splashes)
 - Parallel calibration and sampling
 - Requires incremental junction trees
 - Relies heavily on:
 - Edge consistency model to prove ergodicity
 - FIFO/ Prioritized scheduling to construct Splashes
- Evaluated on 32 core Nehalem Server

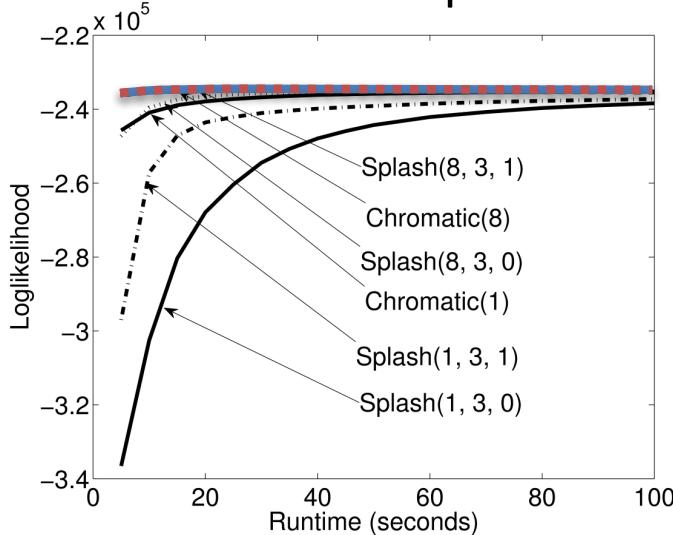
Rapidly Mixing Model

- Grid MRF with weak attractive potentials

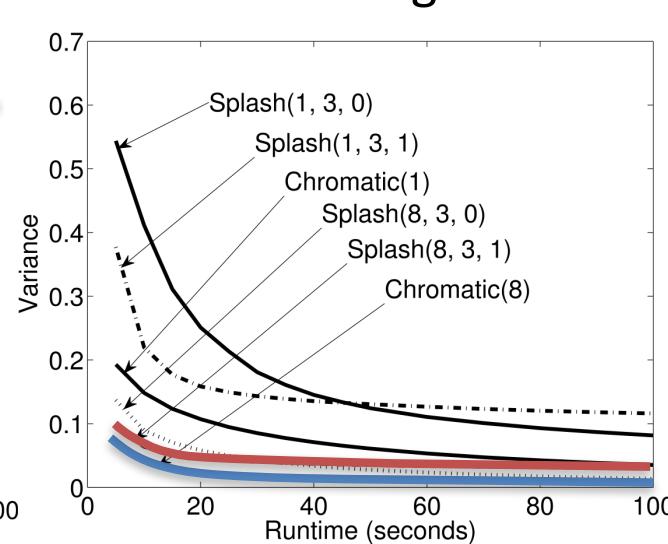
40K Variables

80K Factors

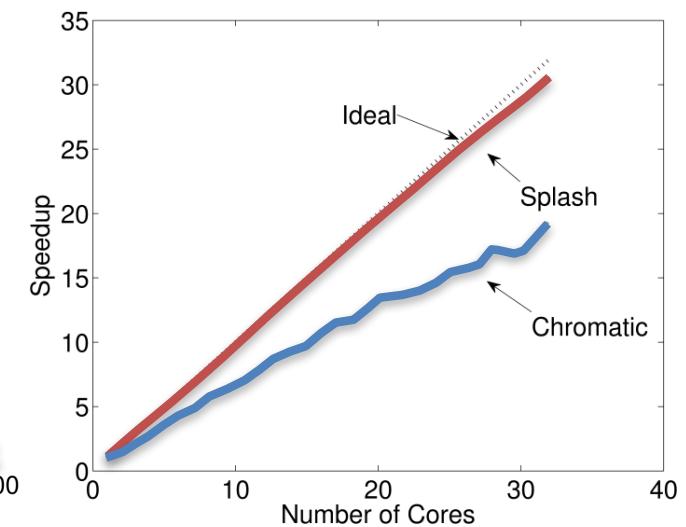
Likelihood
Final Sample



“Mixing”



Speedup



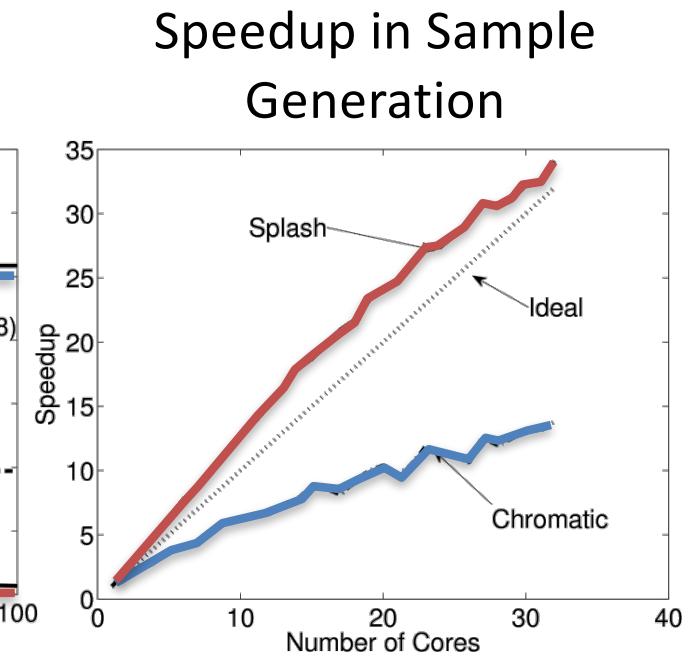
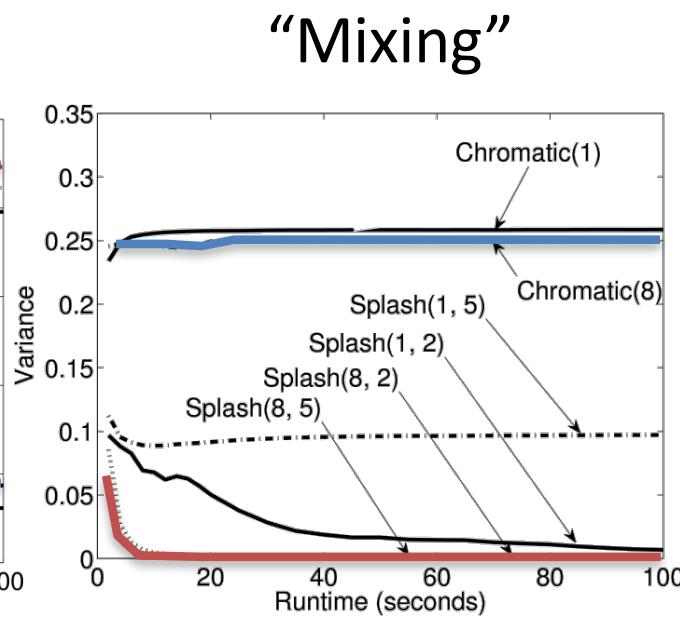
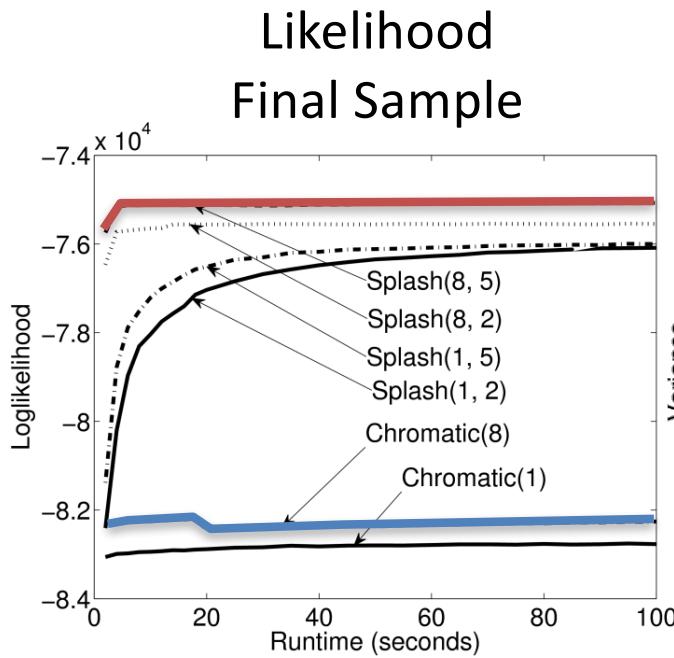
- The Chromatic sampler slightly outperforms the Splash Sampler

Slowly Mixing Model

- Markov logic network with strong dependencies

10K Variables

28K Factors



- The *Splash* sampler outperforms the *Chromatic* sampler on models with **strong** dependencies

Conclusion

- **Chromatic Gibbs** sampler for models with *weak dependencies*
 - *Converges to the correct distribution*
 - *Quantifiable improvement in mixing*
- **Theoretical analysis** of the Synchronous Gibbs sampler on *2-colorable models*
 - *Proved marginal convergence on 2-colorable models*
- **Splash Gibbs** sampler for models with *strong dependencies*
 - *Adaptive asynchronous tree construction*
 - *Experimental evaluation demonstrates an improvement in mixing*

Future Work

- Extend Splash algorithm to models with continuous variables
 - Requires continuous junction trees (Kernel BP)
- Consider “freezing” the junction tree set
 - Reduce the cost of tree generation?
- Develop better adaptation heuristics
 - Eliminate the need for vanishing adaptation?
- Challenges of Gibbs sampling in high-coloring models
 - Collapsed LDA
- High dimensional pseudorandom numbers
 - Not currently addressed in the MCMC literature