

# Deploying Interactive Machine Learning Applications with Clipper



Joseph E. Gonzalez  
Co-director of the RISE Lab  
[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)



# Managing the Machine Learning Lifecycle

Joseph E. Gonzalez  
Co-director of the RISE Lab  
[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)

# About Me

- Co-director of the RISE Lab
- Co-founder of Turi Inc.
- Member of the Apache Spark PMC

## Research

- Artificial Intelligence
- Data Science
- Distributed Data Systems
- Graph Processing Systems



Conjecture

# *Machine learning models are the next “big data”*

Evidence

1. Everyone is talking about models but few have them.
2. They have the opportunity to transform industries.
3. They are a consequence of mastering big data.
4. Today, their full value is only realized with advanced skills and technologies

Conjecture

Machine learning models  
are the next “big data”

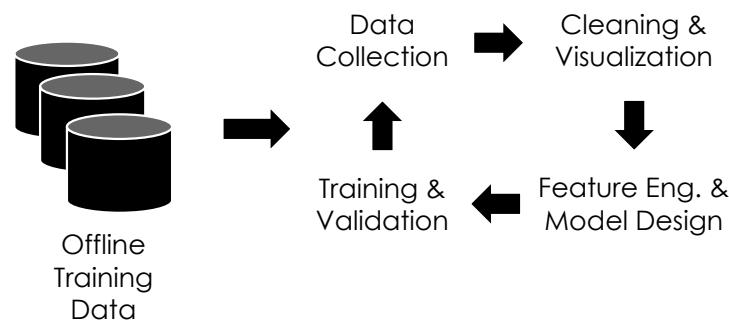
Corollaries

Data Engineers will need to manage  
data & **machine learning models**

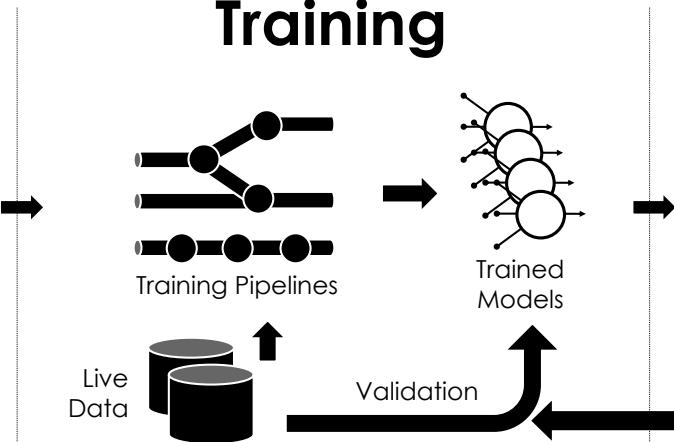
We need new technologies to manage  
the **machine learning lifecycle**

# What is the *Machine Learning Lifecycle*?

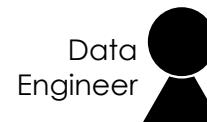
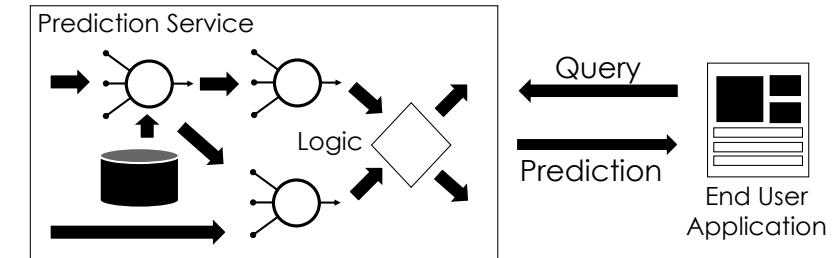
## Model Development



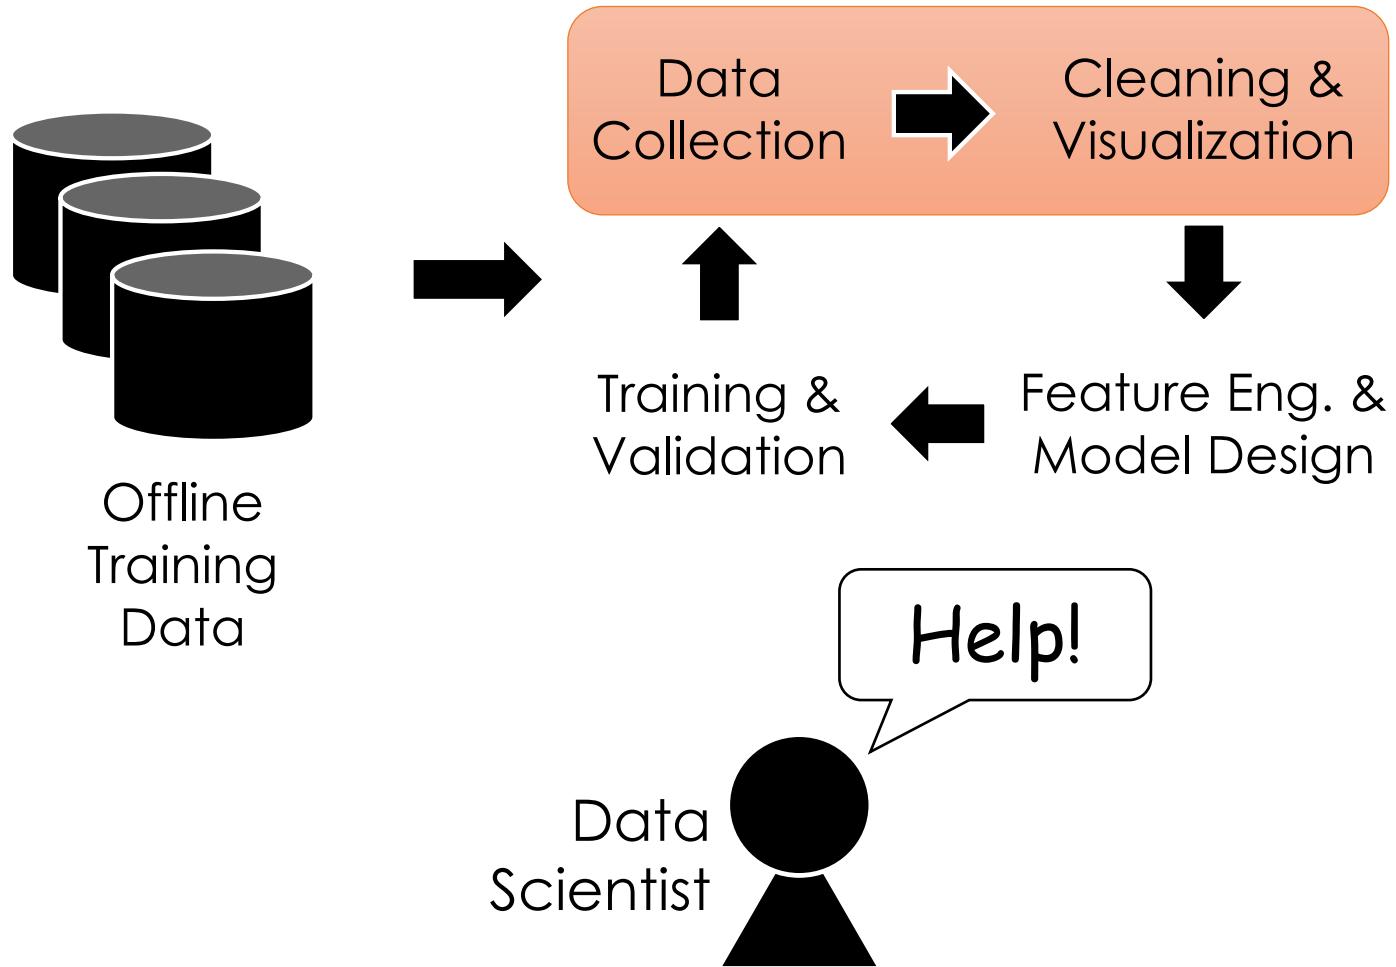
## Training



## Inference



# Model Development



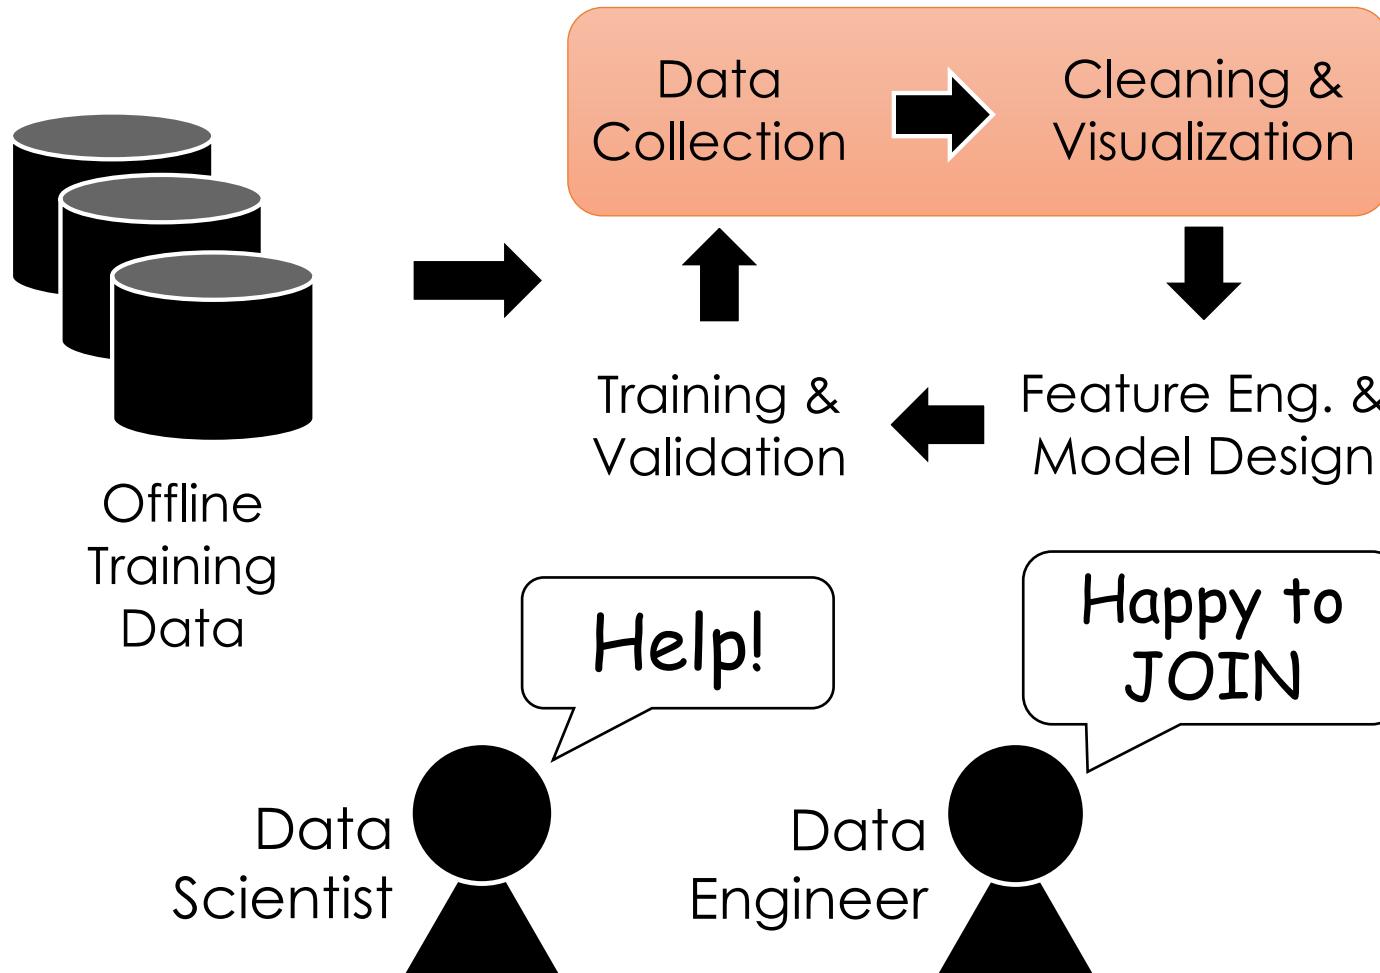
**Identifying** potential sources of data

**Joining** data from multiple sources

Addressing **missing values** and **outliers**

**Plotting** trends to identify **anomalies**

# Model Development



**Identifying** potential sources of data

**Joining** data from multiple sources

Addressing **missing values** and **outliers**

**Plotting** trends to identify **anomalies**



**Big Data Borat**

@BigDataBorat

Follow

In Data Science, 80% of time spent prepare data, 20% of time spent complain about need for prepare data.

6:47 PM - 26 Feb 2013

---

533 Retweets 330 Likes



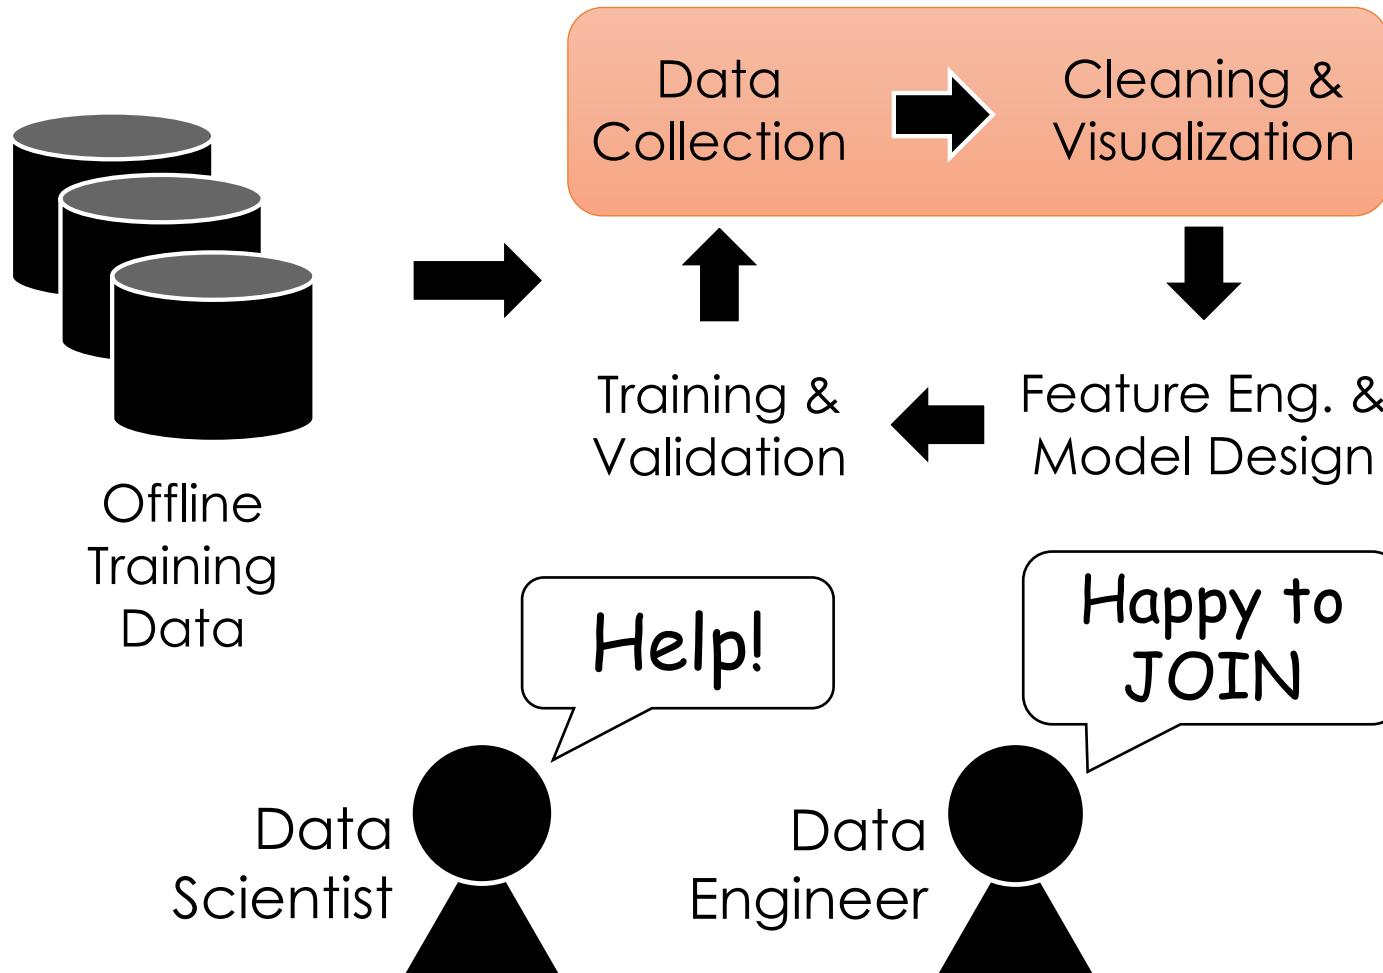
---

12

533

330

# Model Development



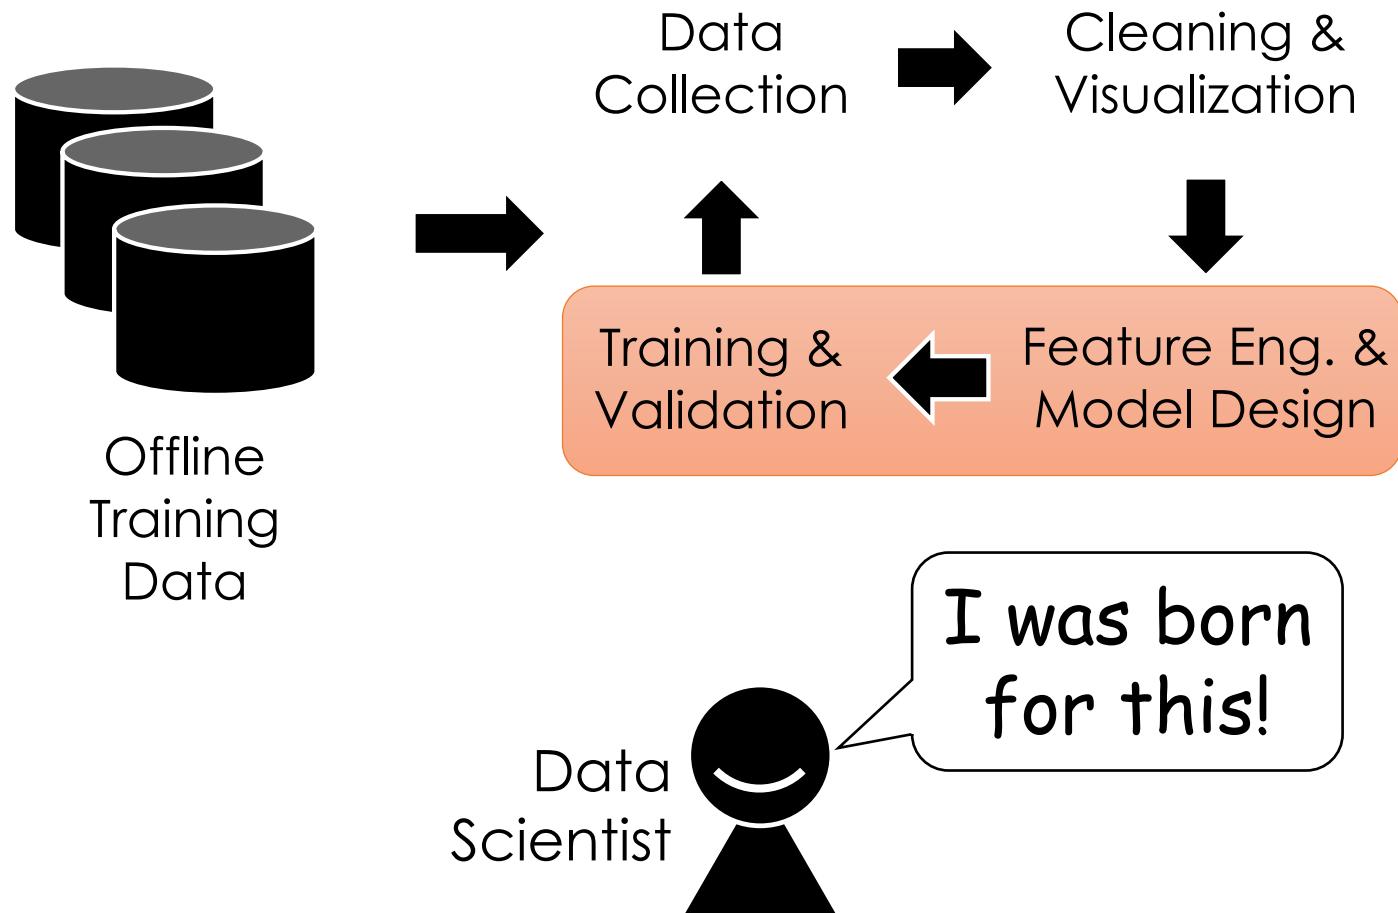
**Identifying** potential sources of data

**Joining** data from multiple sources

Addressing **missing values** and **outliers**

**Plotting** trends to identify **anomalies**

# Model Development



Building informative  
**features functions**

Designing new **model architectures**

**Tuning** training algos.

**Validating** prediction accuracy

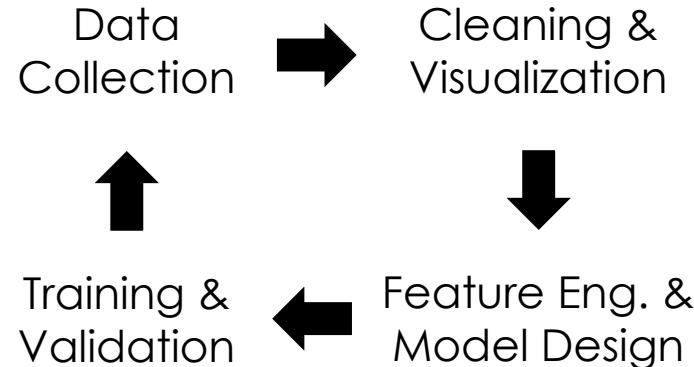
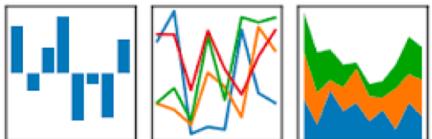
# Model Development Technologies



Offline  
Training  
Data

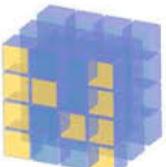
matplotlib

pandas  
 $y_t = \beta' x_t + \mu_t + \epsilon_t$



jupyter

☕ Caffe2



NumPy



scikit  
learn



TensorFlow



PYTORCH

K Keras

mxnet

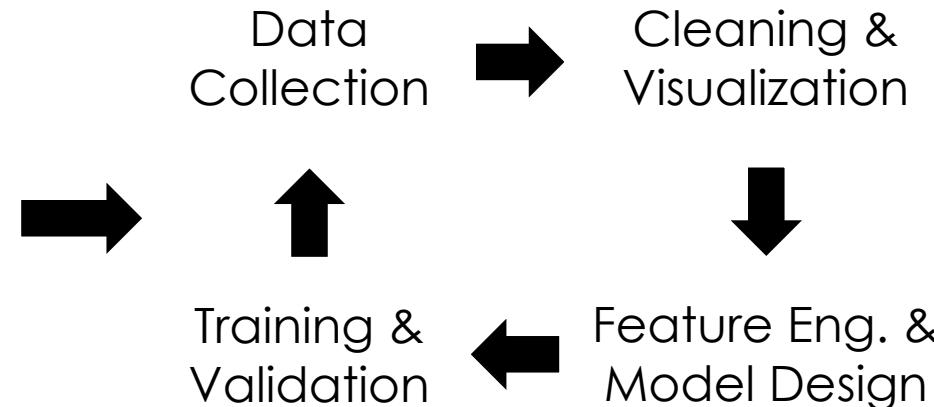
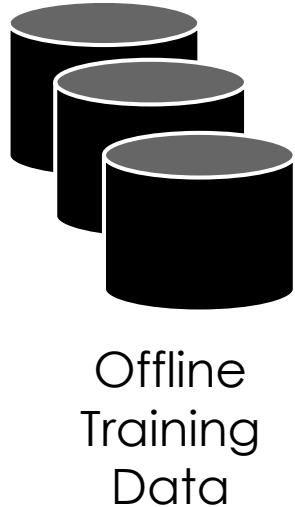
dmlc  
XGBoost

APACHE  
Spark™

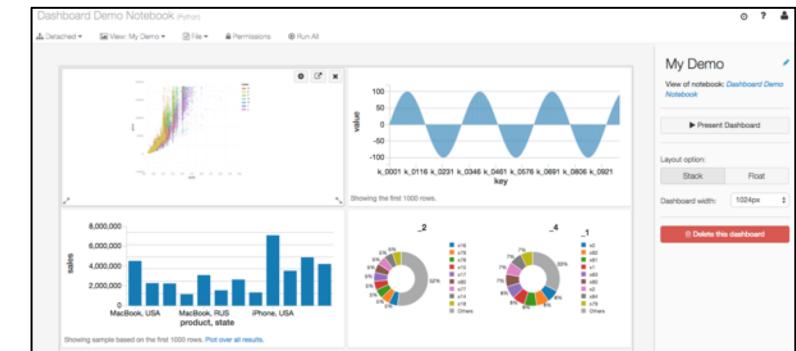
DASK

HIVE

# What is the output of Model Development

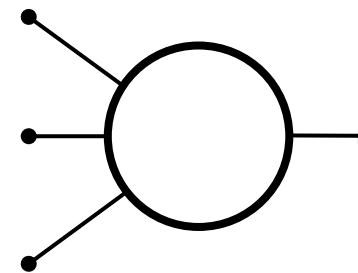


## Reports & Dashboards



(insights ...)

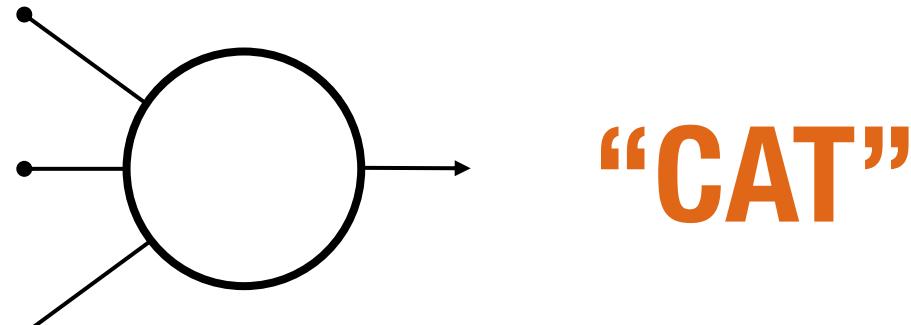
## Trained Model



# A learned function from a **query** to a **prediction**



**Trained Model**

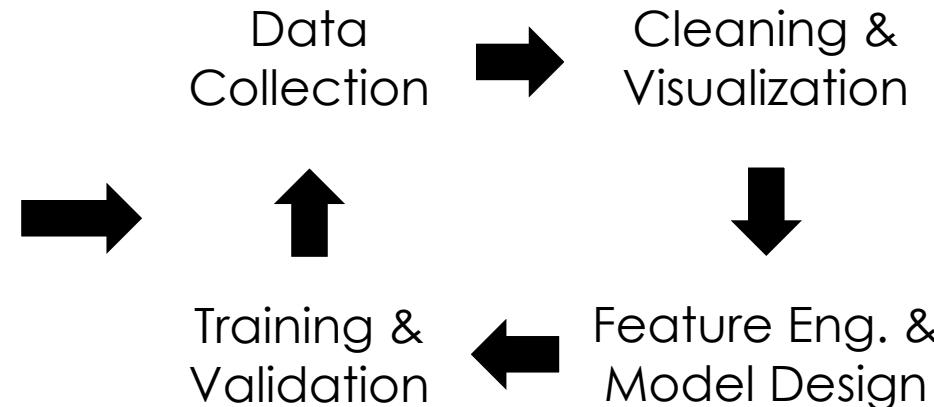
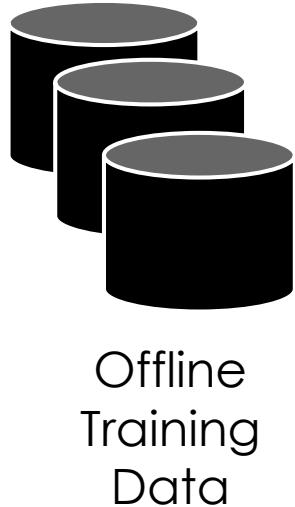


consisting of **parameters** and **model structure**.

Data  
(10B to 10GB)

How to use the  
parameters...

# What is the output of Model Development

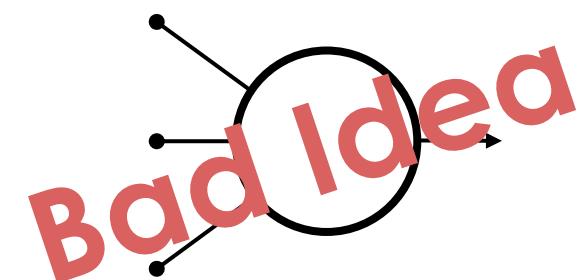


## Reports & Dashboards



(insights ...)

## Trained Model

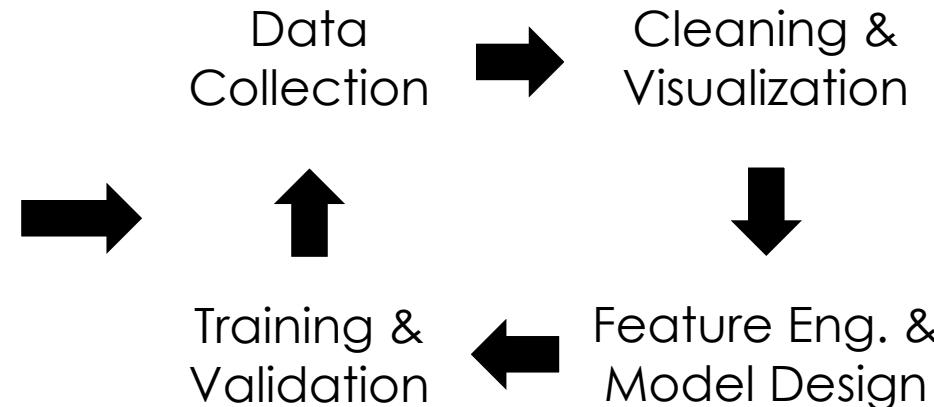
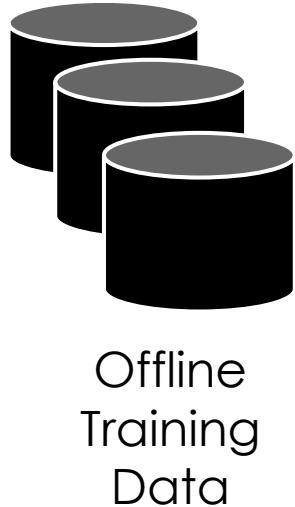


# Why is it a **Bad Idea** to directly produce trained models from model development?

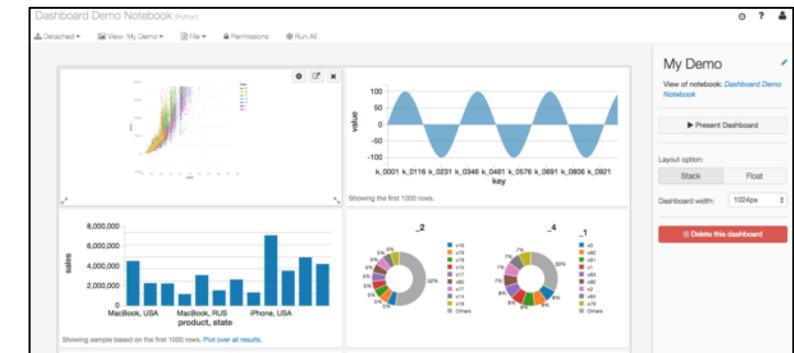
With just a trained model we are **unable to**

1. **retrain** models with new data
2. track data and code for **debugging**
3. capture **dependencies** for deployment
4. audit training for **compliance** (e.g., GDPR)

# What is the output of Model Development

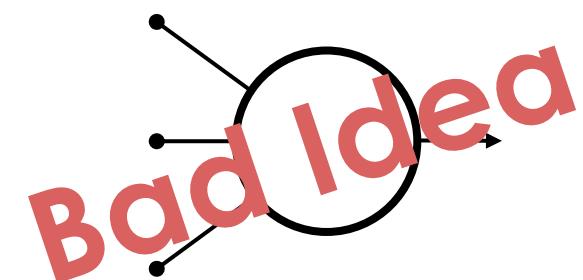


## Reports & Dashboards

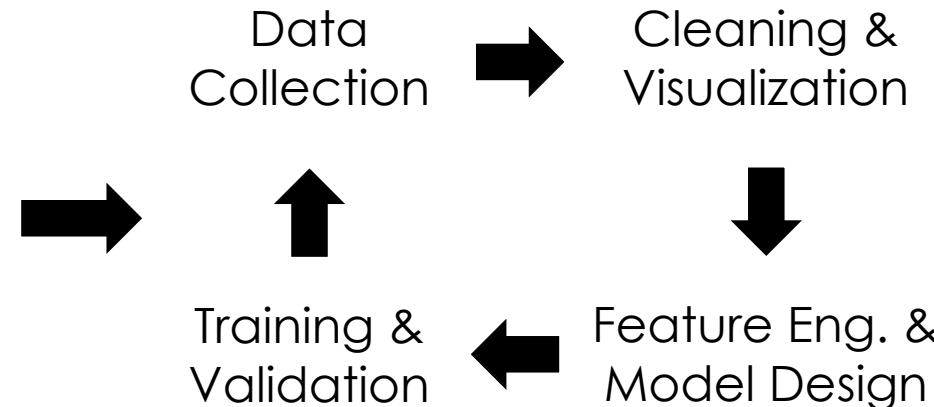
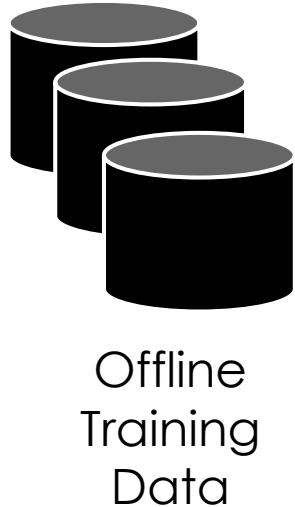


(insights ...)

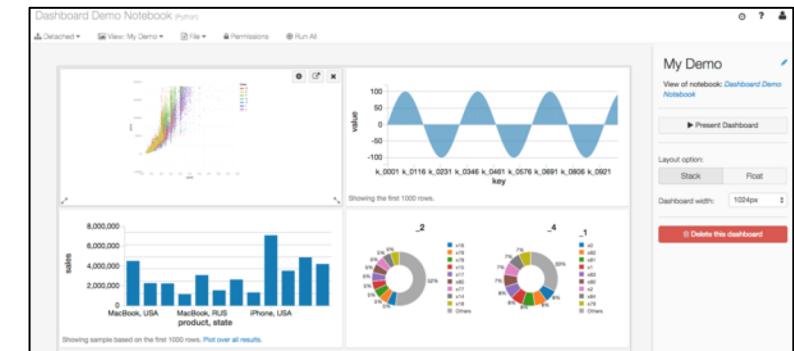
## Trained Models



# What is the output of Model Development

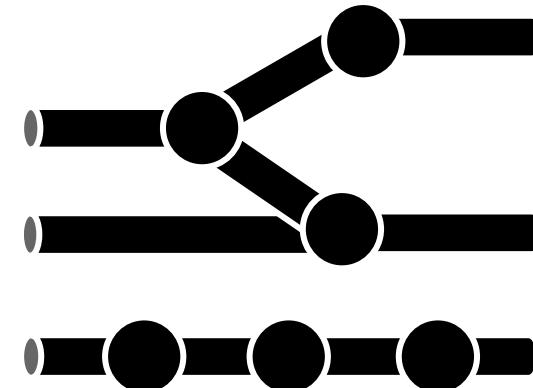


## Reports & Dashboards



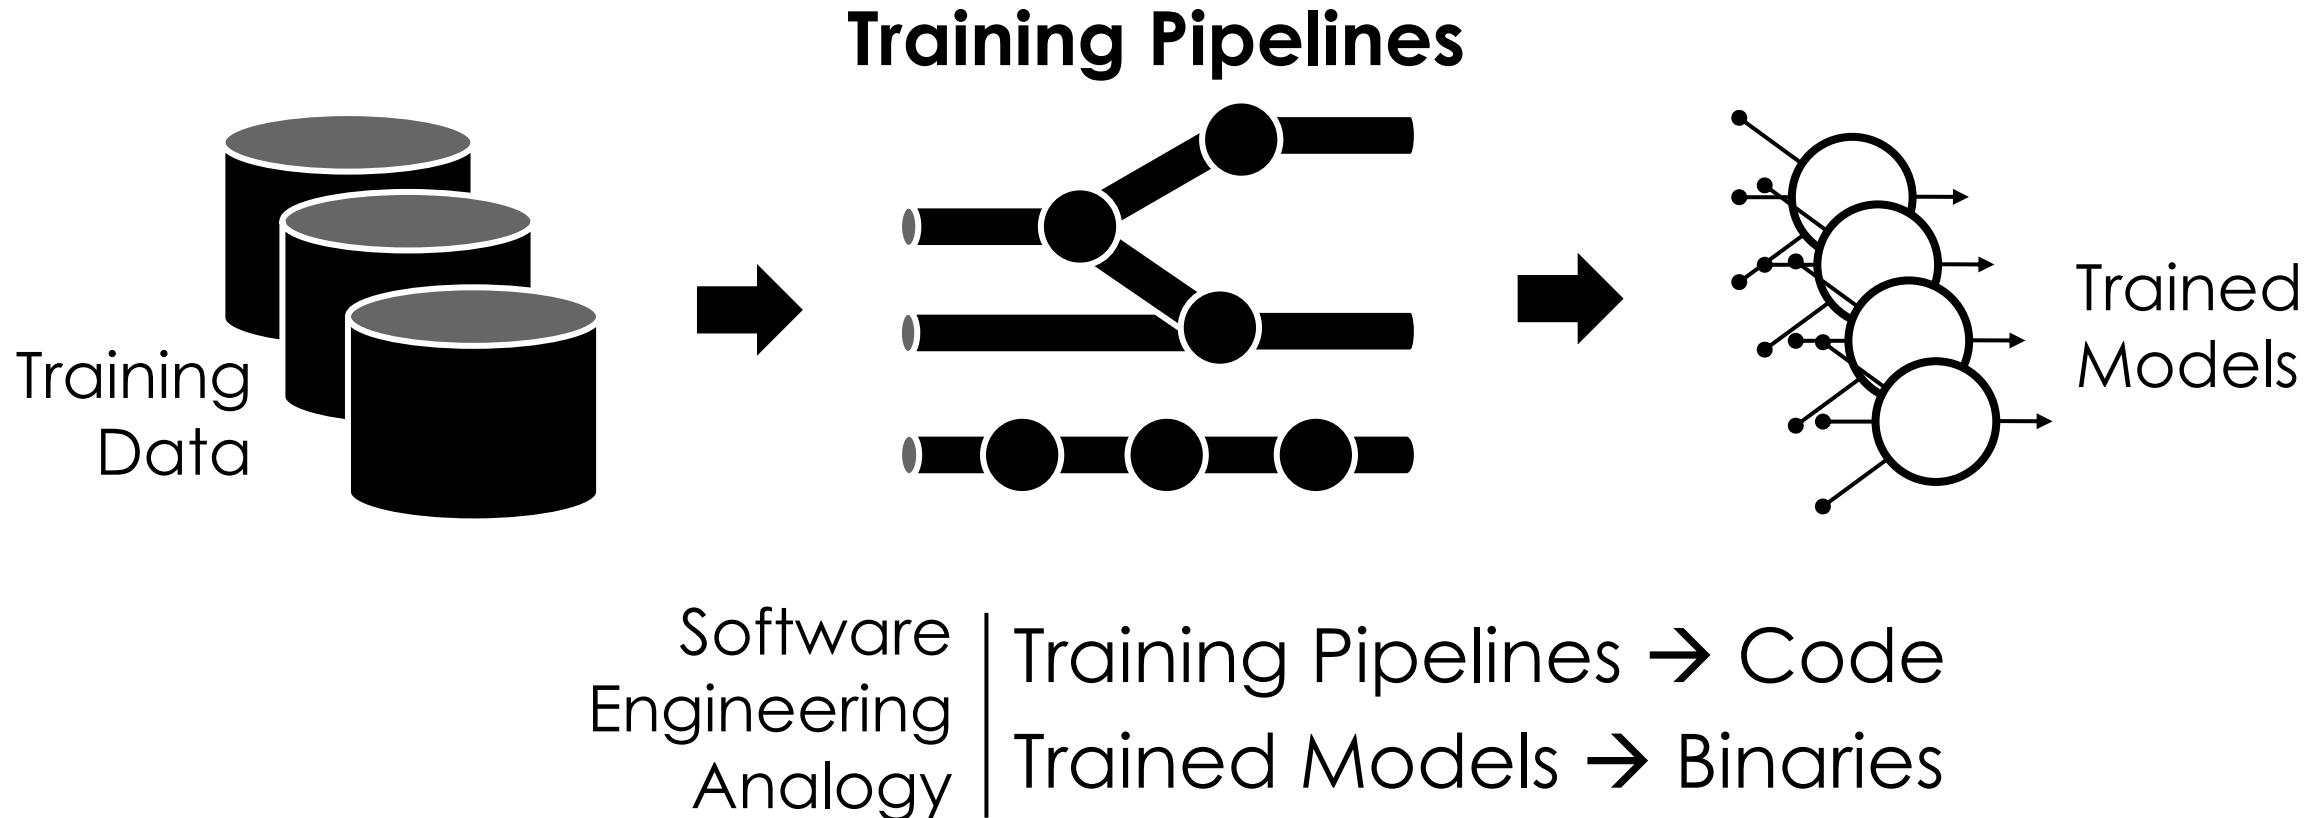
(insights ...)

## Training Pipelines

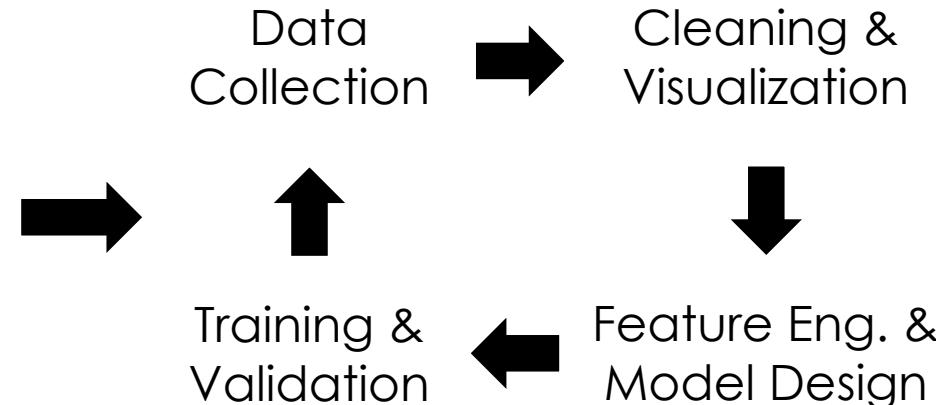
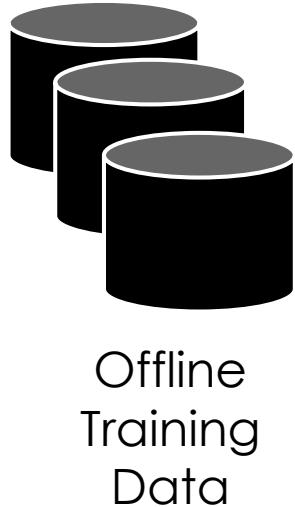


# Training Pipelines Capture the Code and Data Dependencies

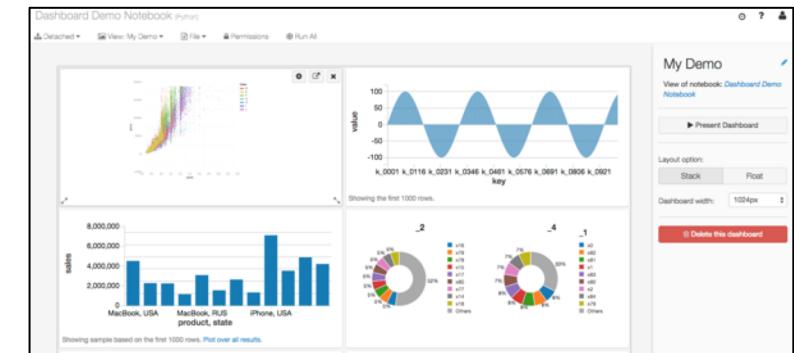
- Description of how to train the model from data sources



# What is the output of Model Development

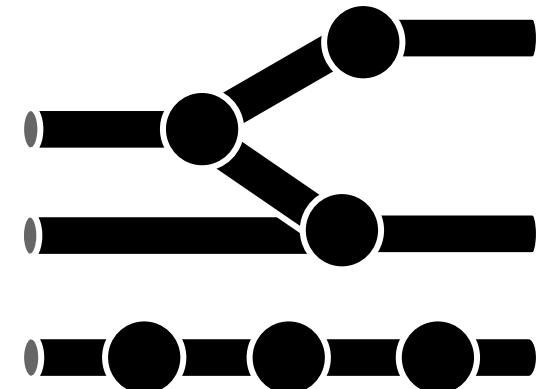


## Reports & Dashboards



(insights ...)

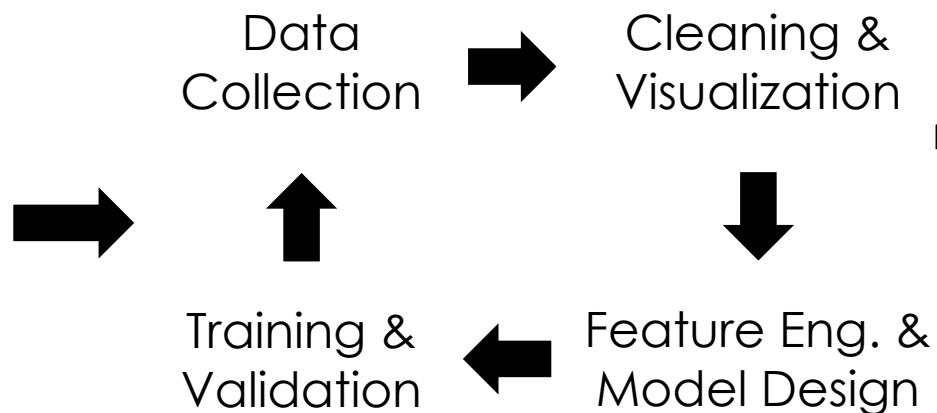
## Training Pipelines



# Model Development

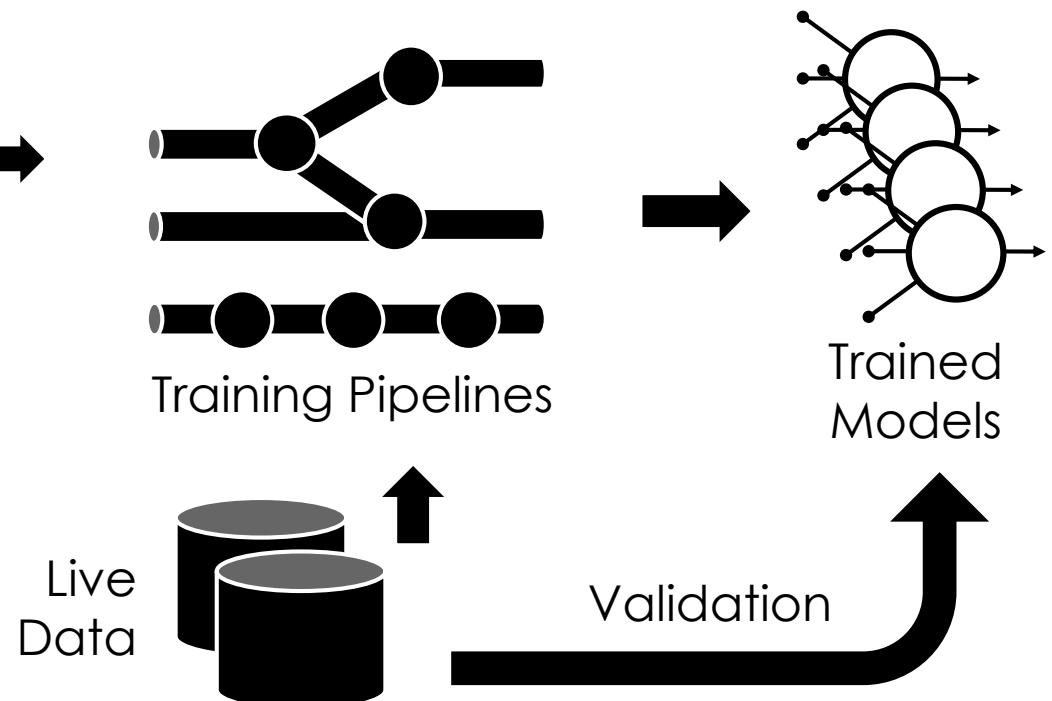


Offline  
Training  
Data



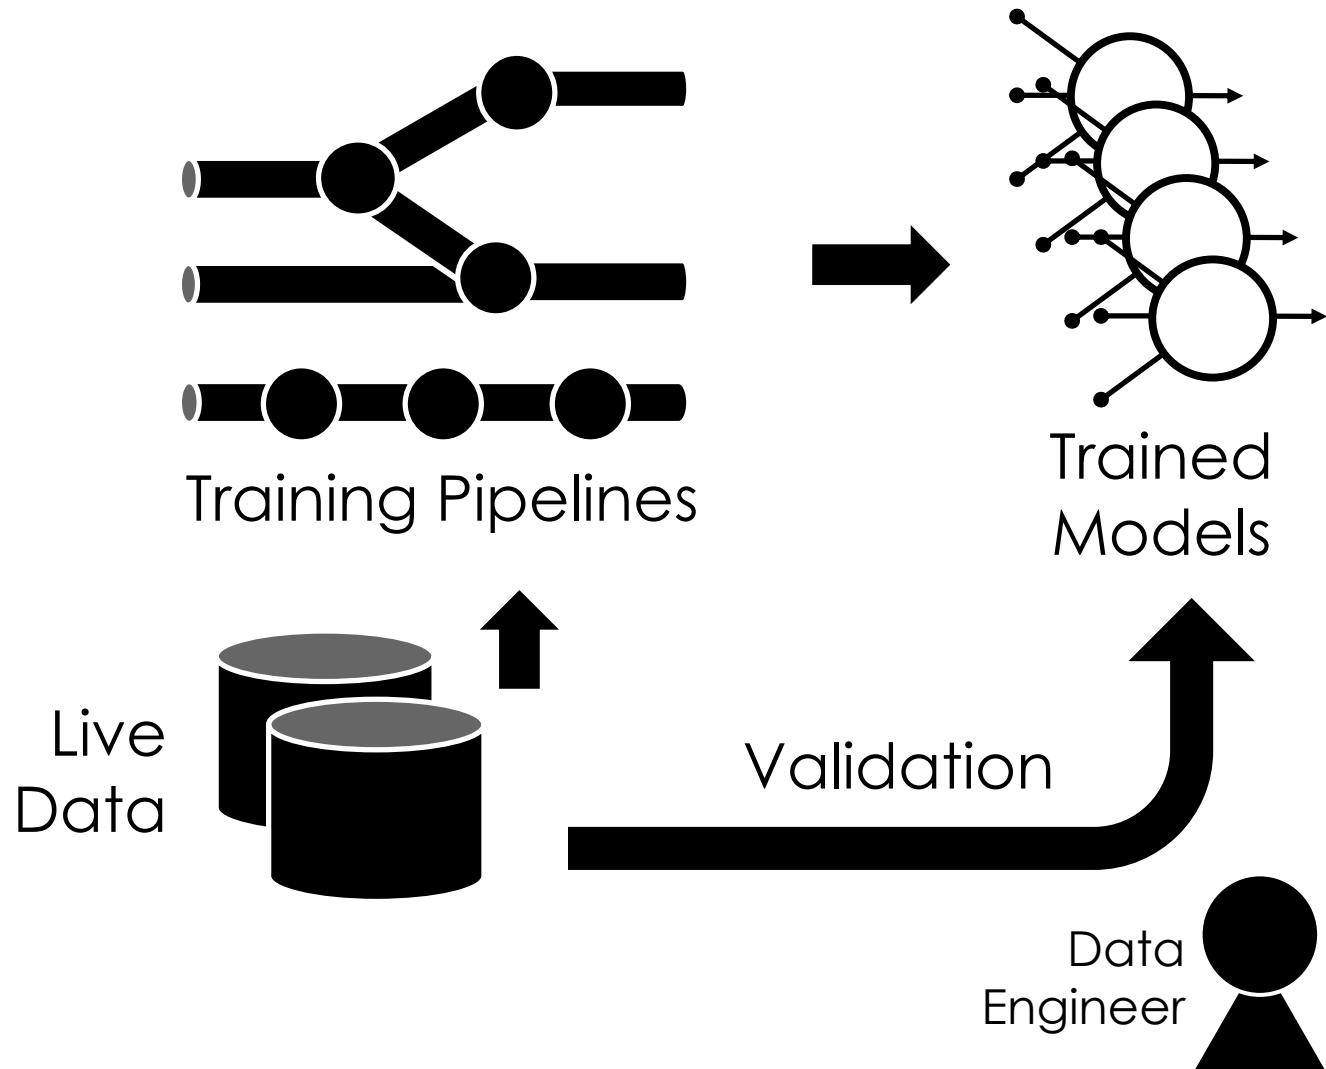
Data  
Scientist

# Training



Data  
Engineer

# Training



Training models **at scale** on **live data**

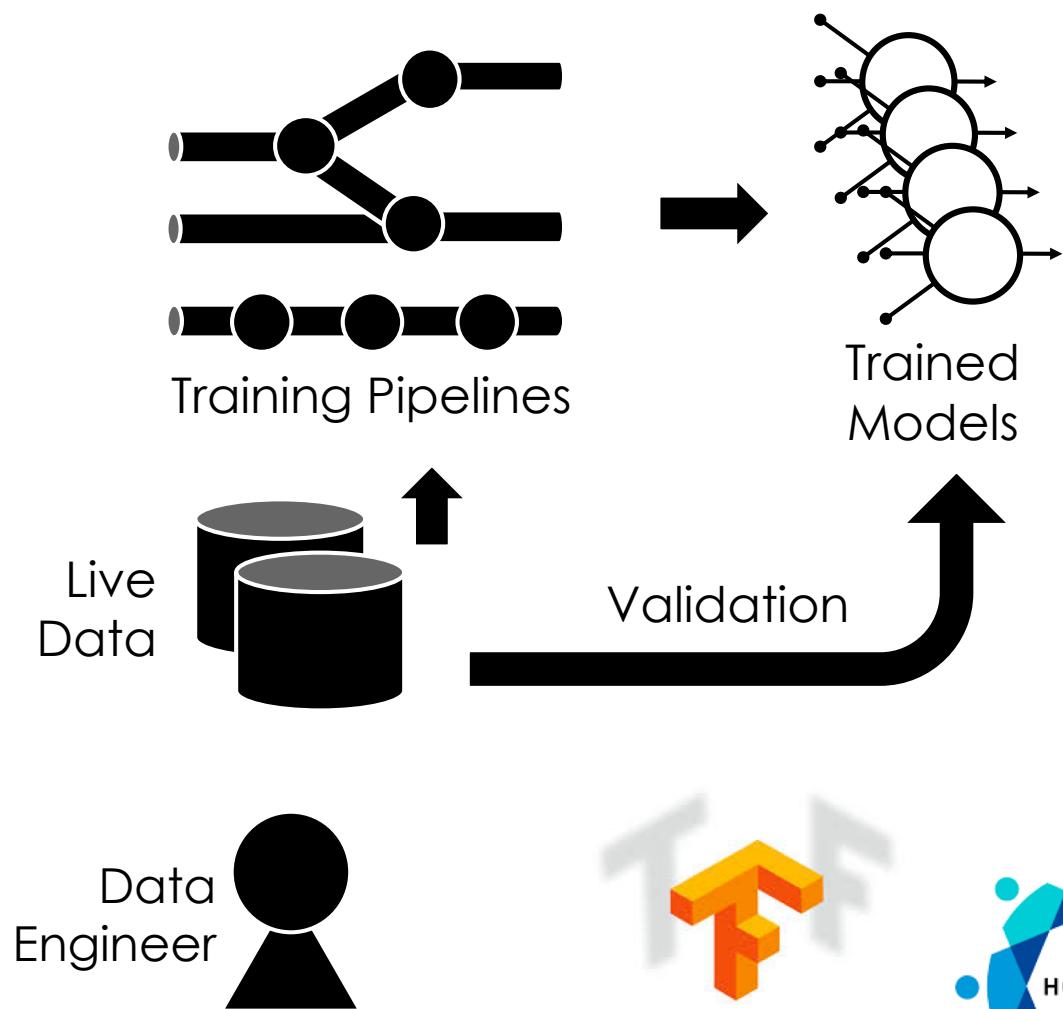
**Retraining** on new data

Automatically **validate** prediction accuracy

Manage model **versioning**

Requires **minimal expertise** in machine learning

# Training Technologies



## Workflow Management:



Apache  
Airflow



## Scalable Training:

PYTORCH



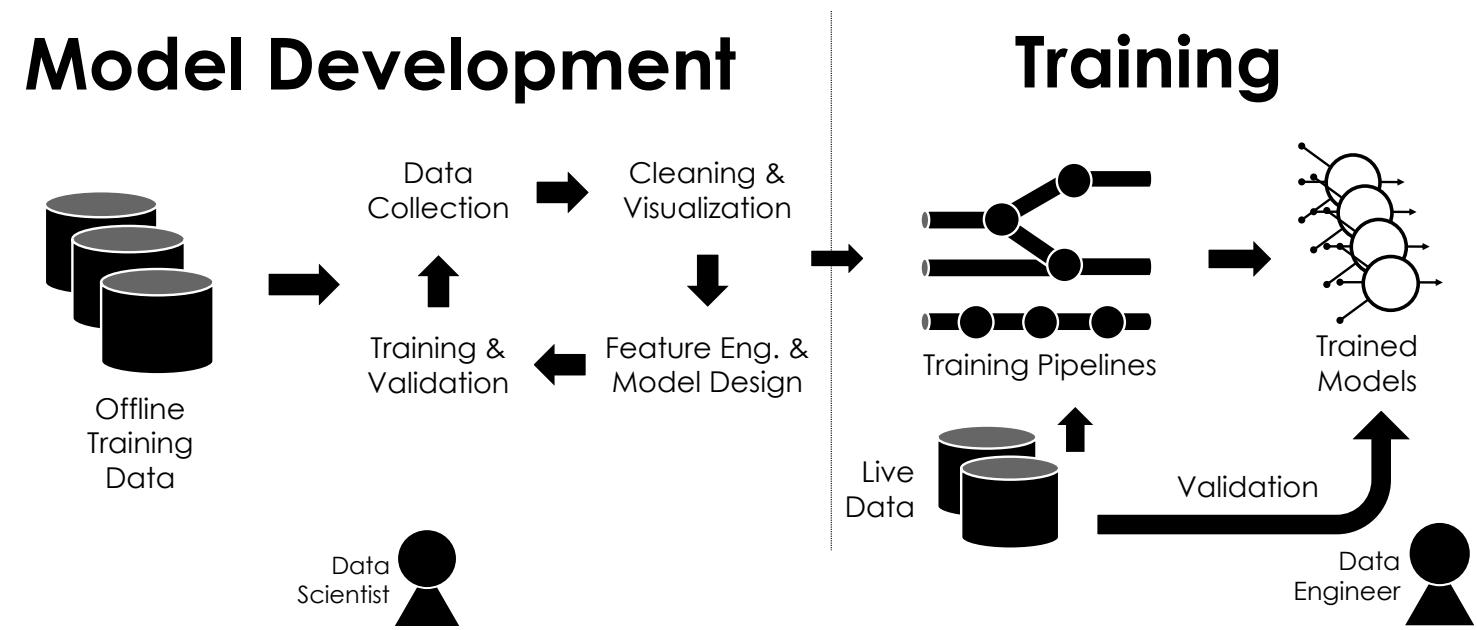
APACHE Spark™

TensorFlow



# Open Problems

Context & Composition



# Context How, What, & Who?

- **How** was the model or data created?
- **What** is the latest or best version?
- **Who** is responsible? (blame...)

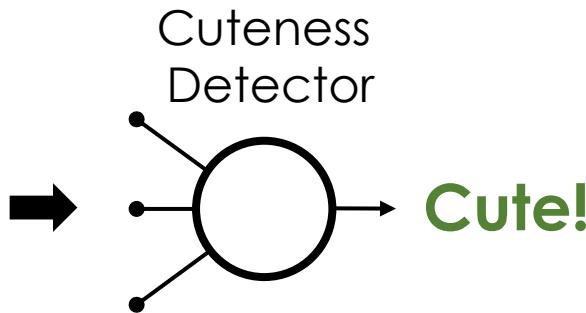


Track relationships between

1. **Code** versions ✓  **git**
2. **Model & Data** versions
3. **People** (versions?)

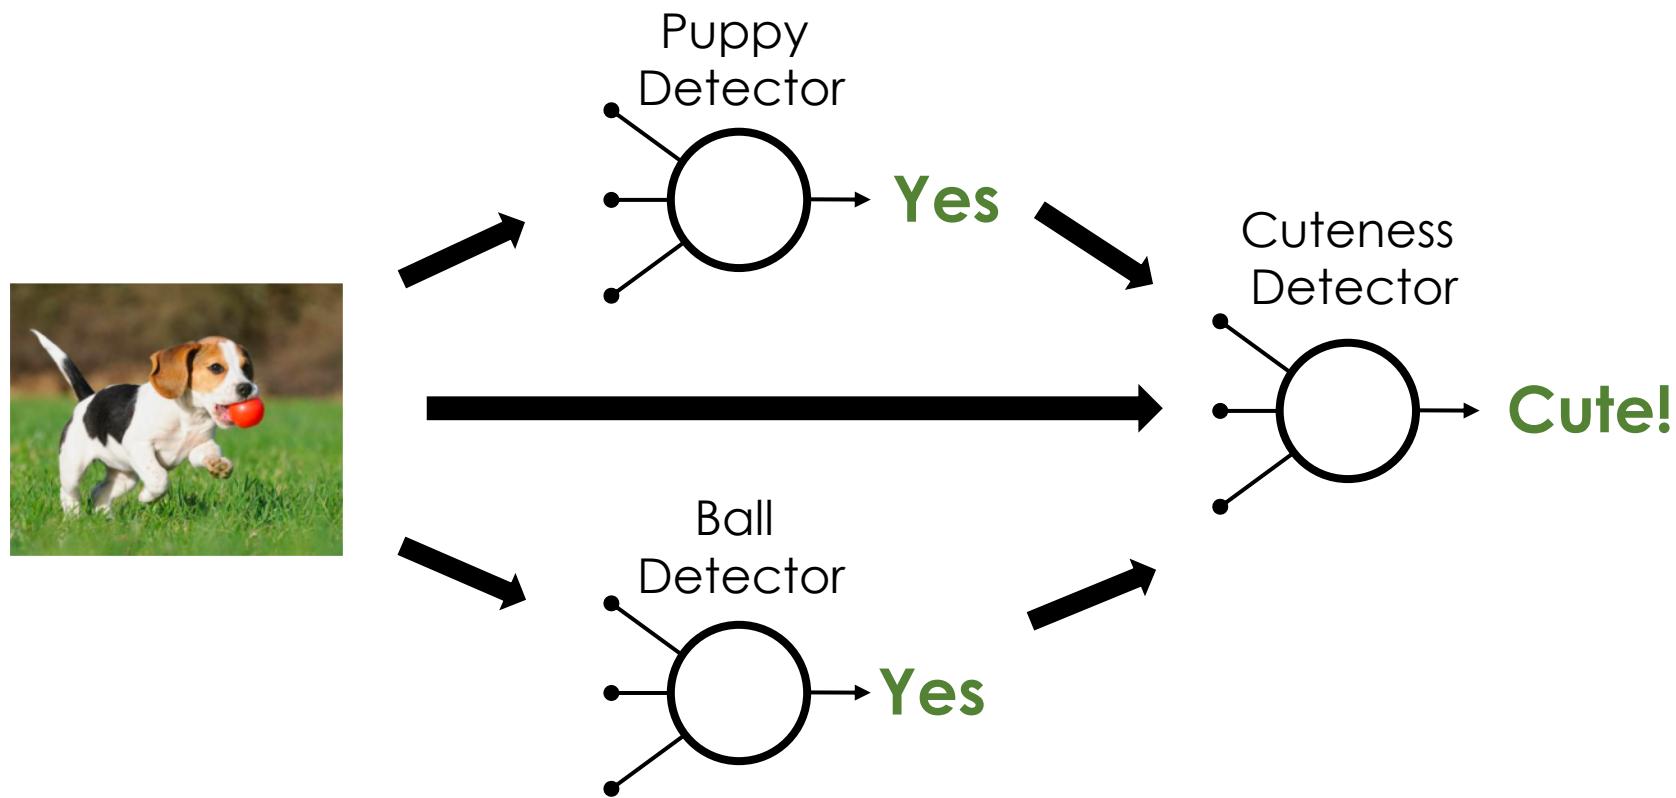
# Composition

Models are being composed to solve new problems



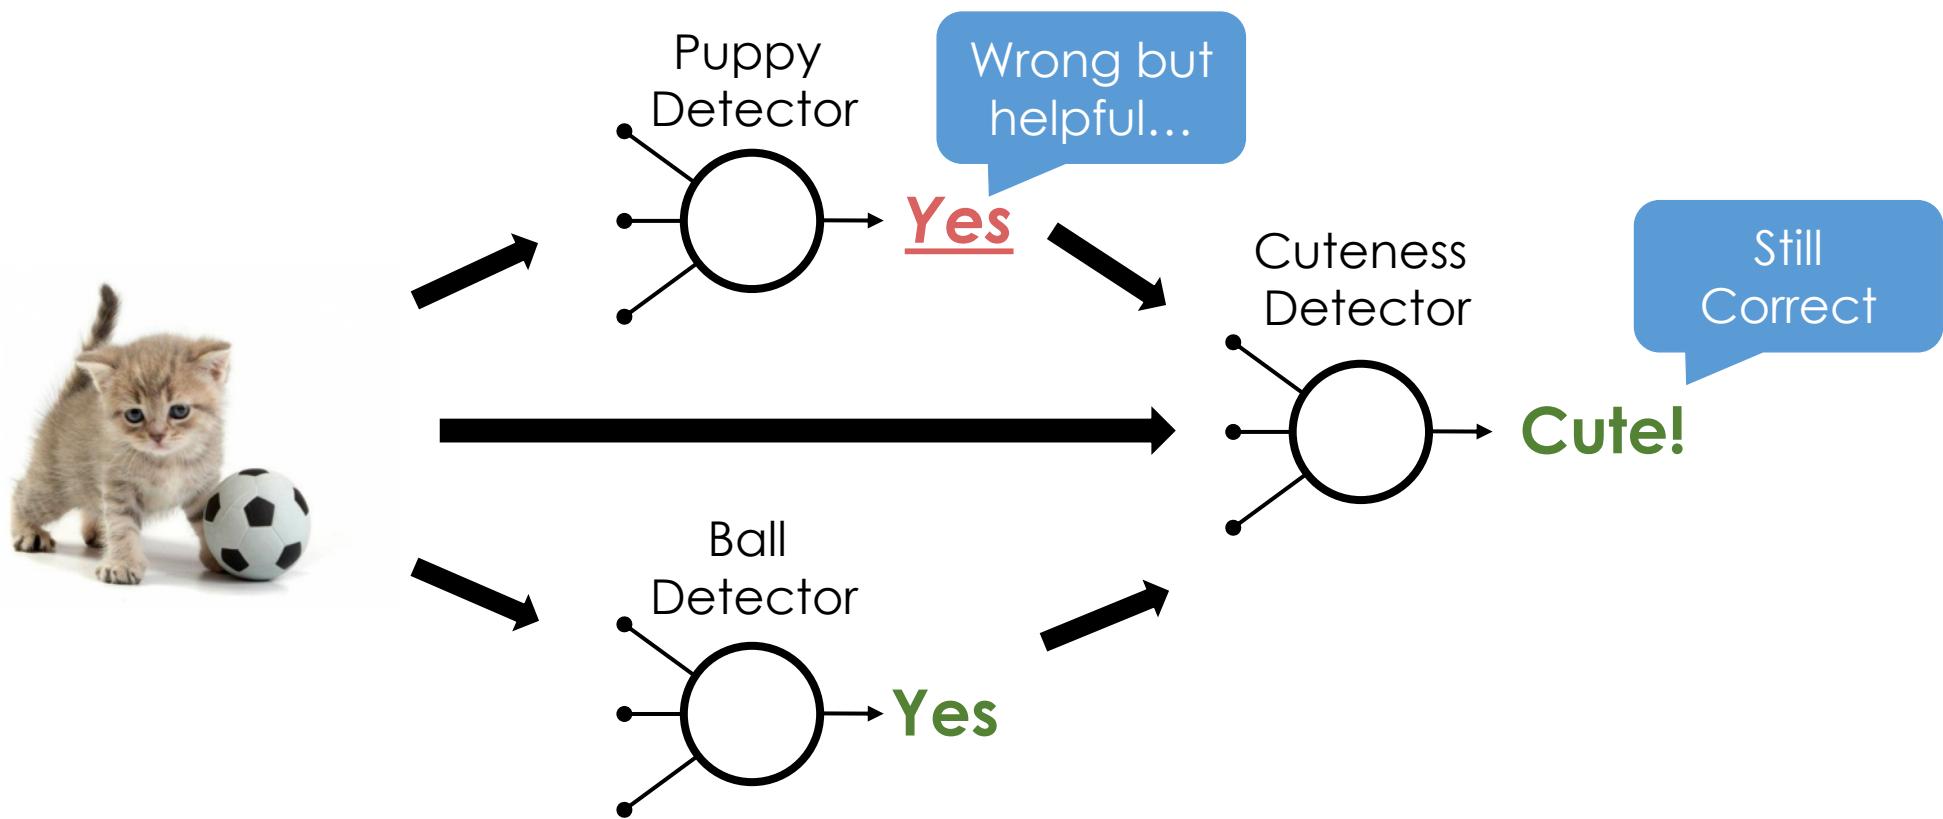
# Composition

Models are being composed to solve new problems



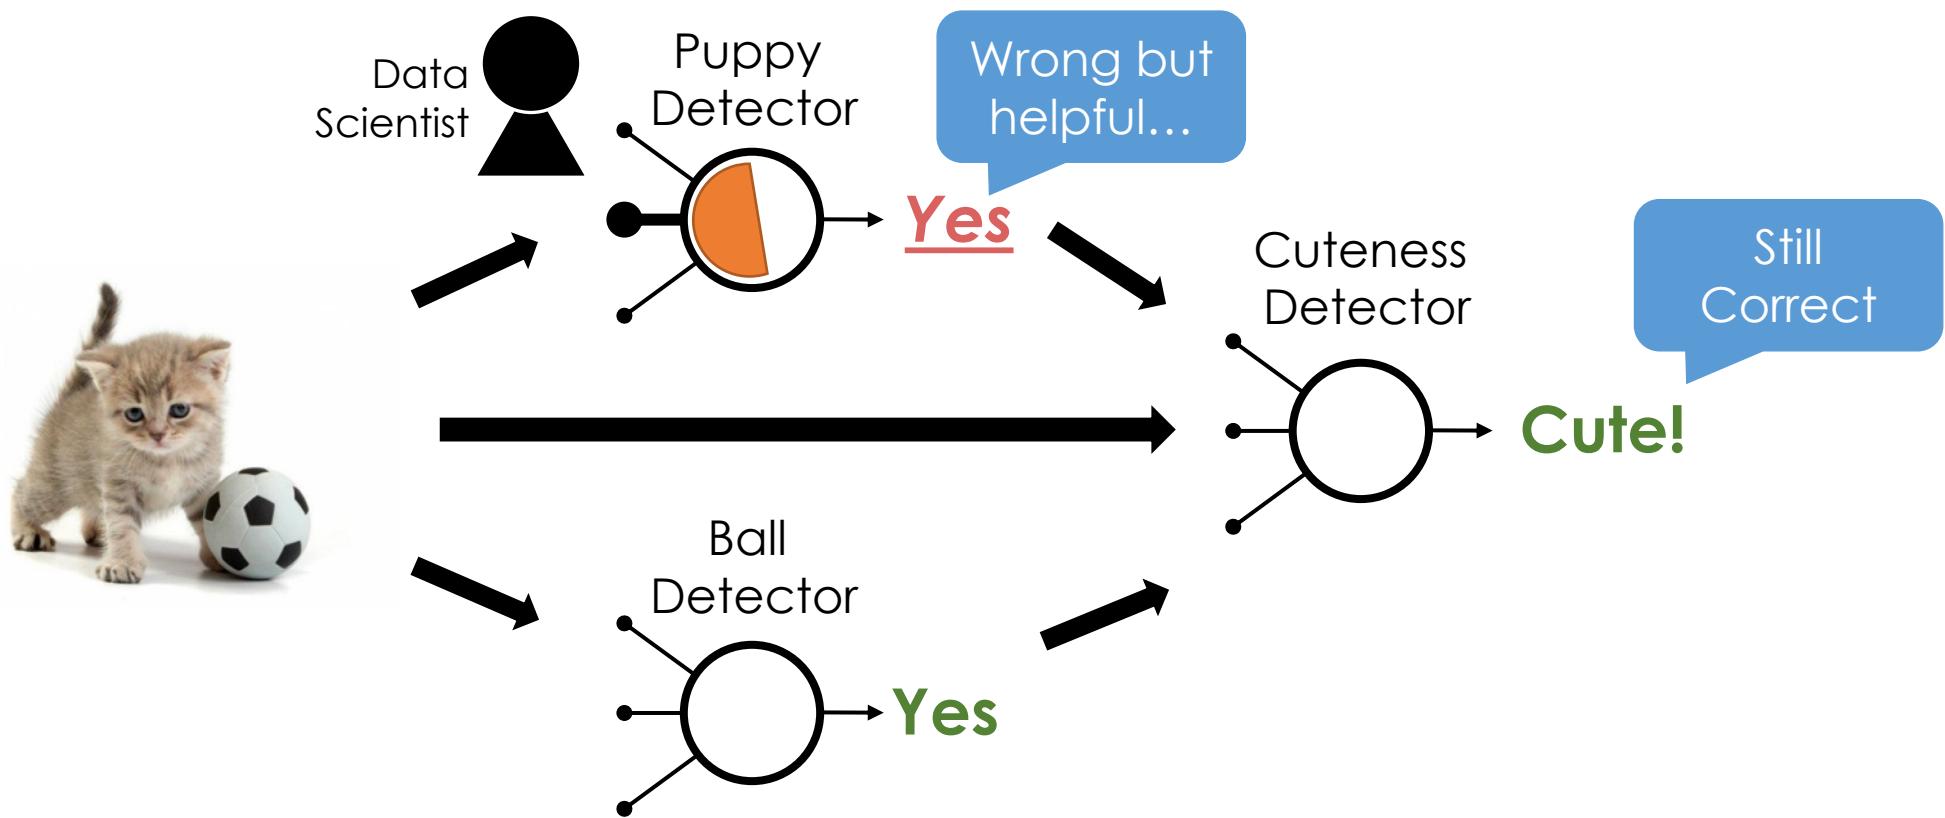
# Composition

Models are being composed to solve new problems



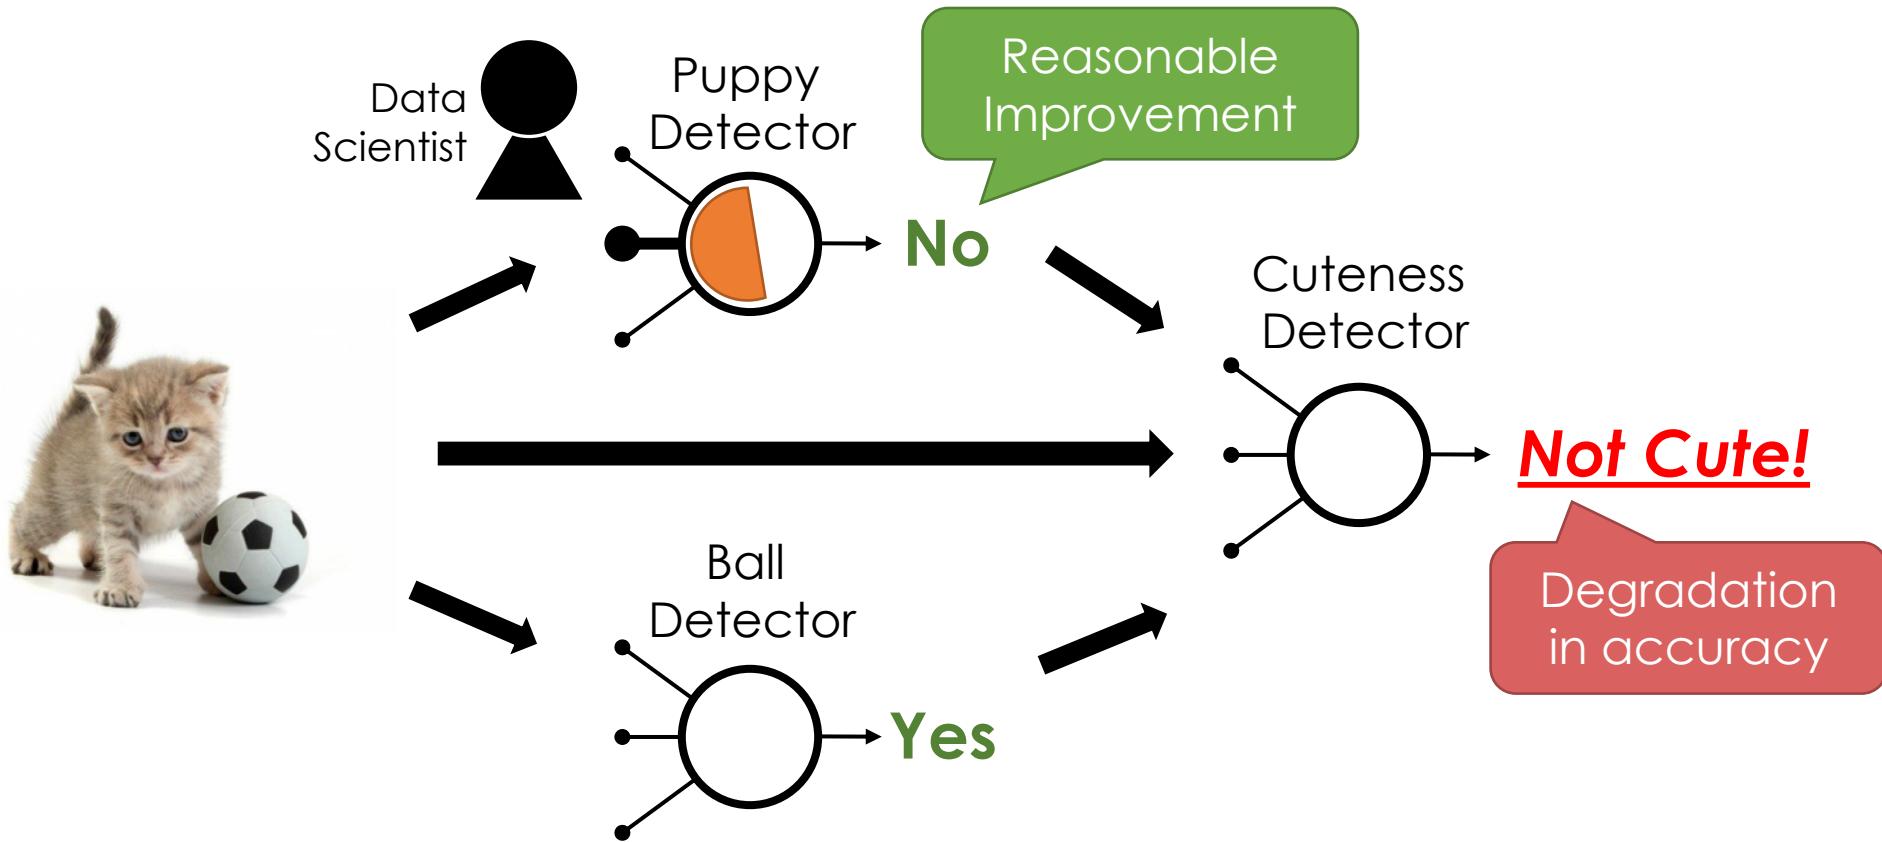
# Composition

Models are being composed to solve new problems



# Composition

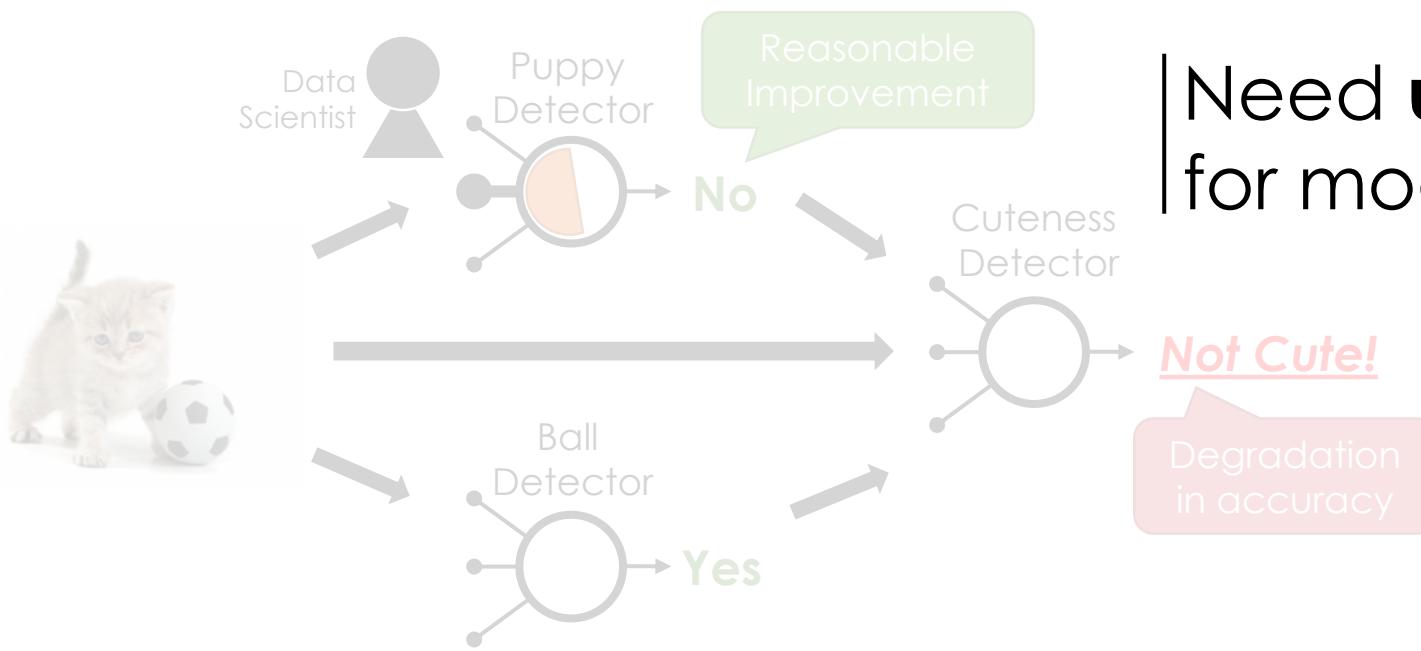
Models are being composed to solve new problems



# Composition

Models are being composed to solve new problems

| Need to track composition and validate **end-to-end accuracy**.



| Need **unit** and **integration** testing for models.

# Active Research in the UC Berkeley for Model Development and Training



A an open source  
**context management** service  
that spans multiple  
data systems

<http://www.ground-context.org/>



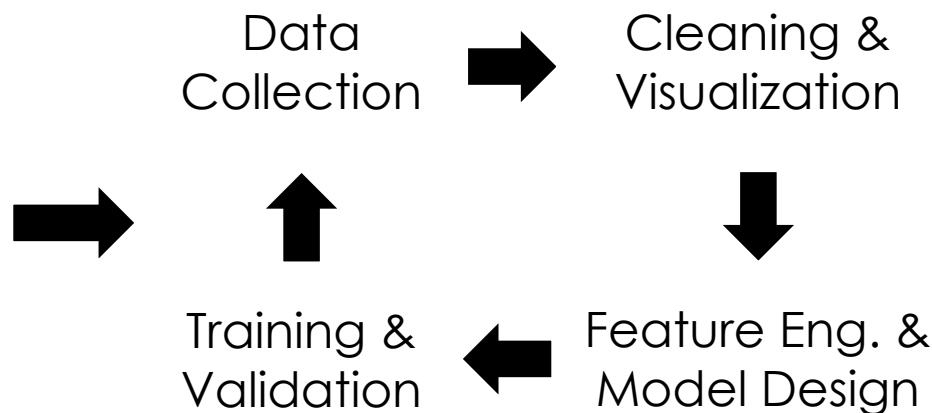
An **experiment management**  
designed to track  
data, code, and people and  
address reproducibility

<https://github.com/ucbrise/flor>

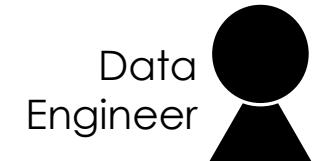
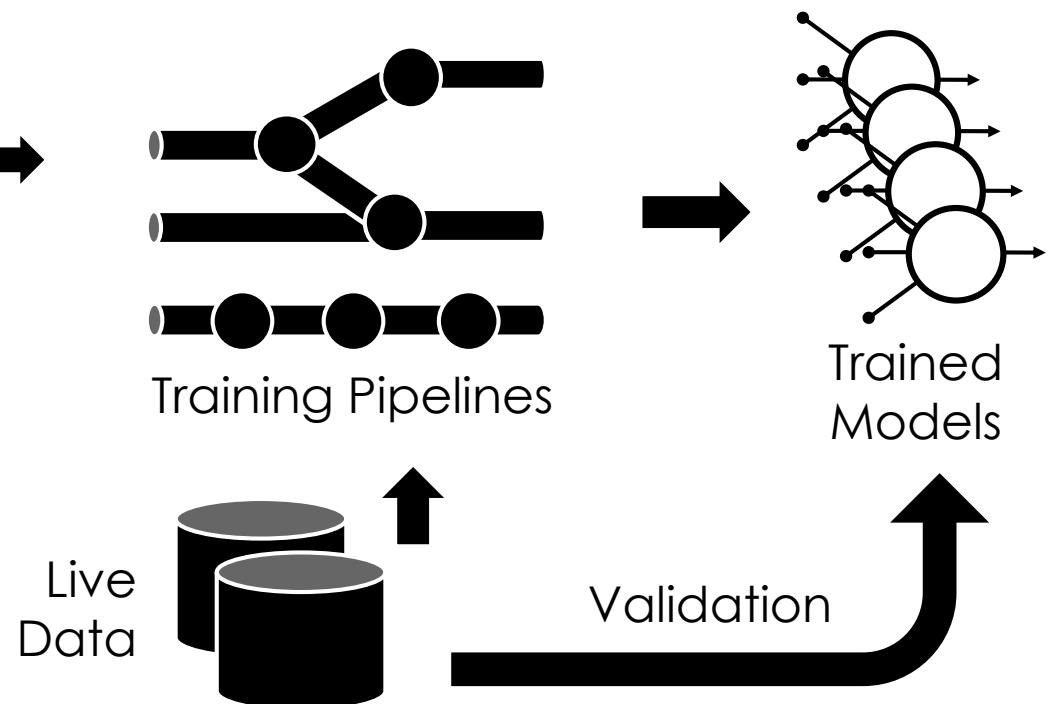
# Model Development



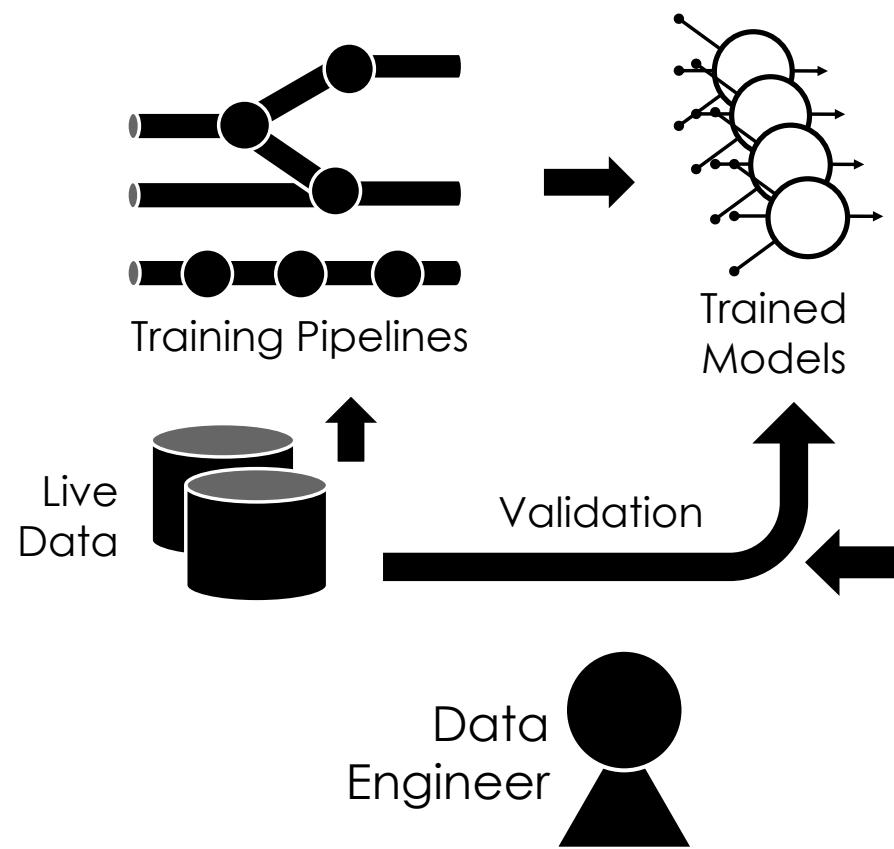
Offline  
Training  
Data



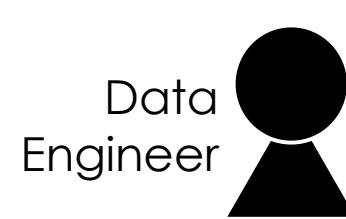
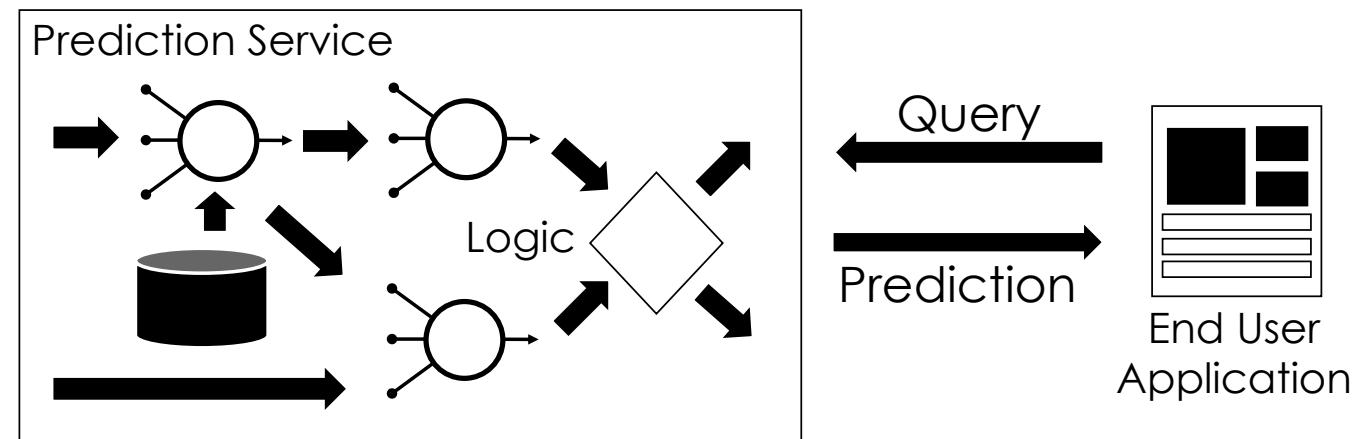
# Training



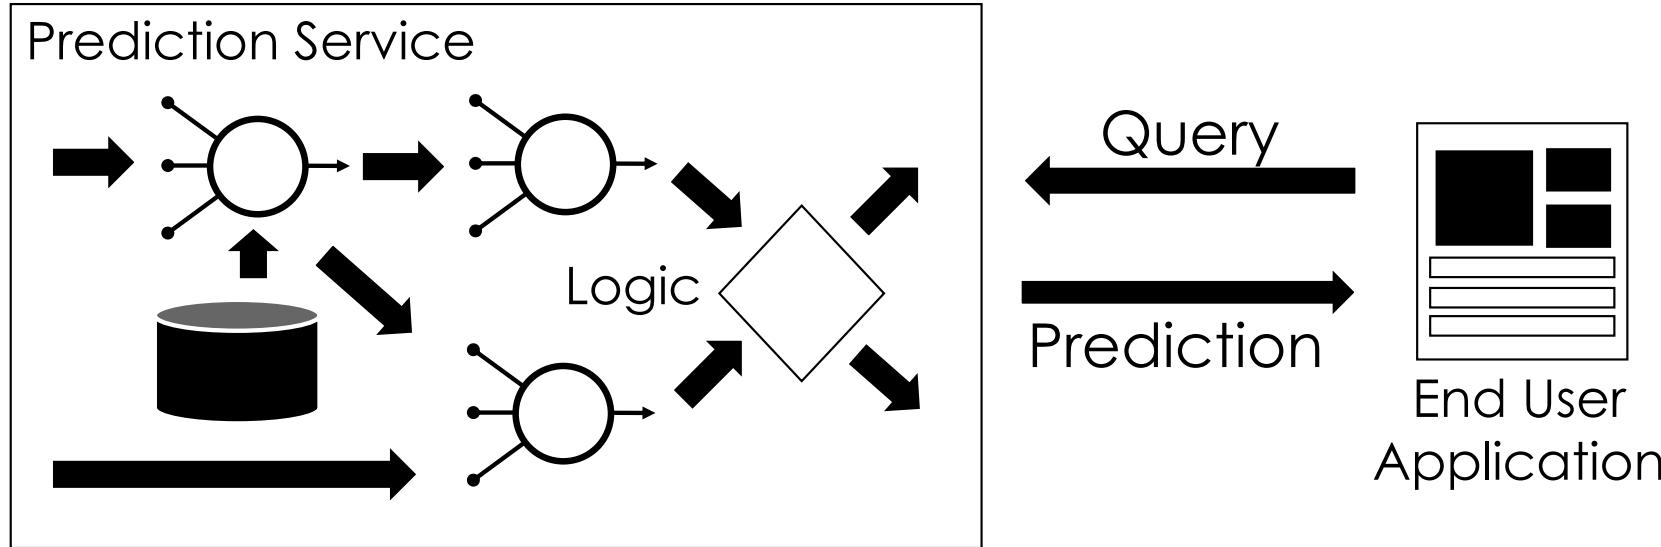
# Training



# Inference

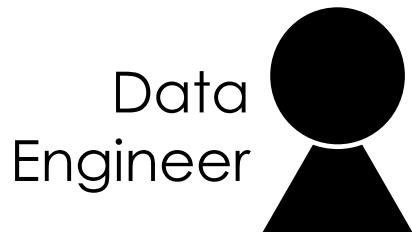


# Inference



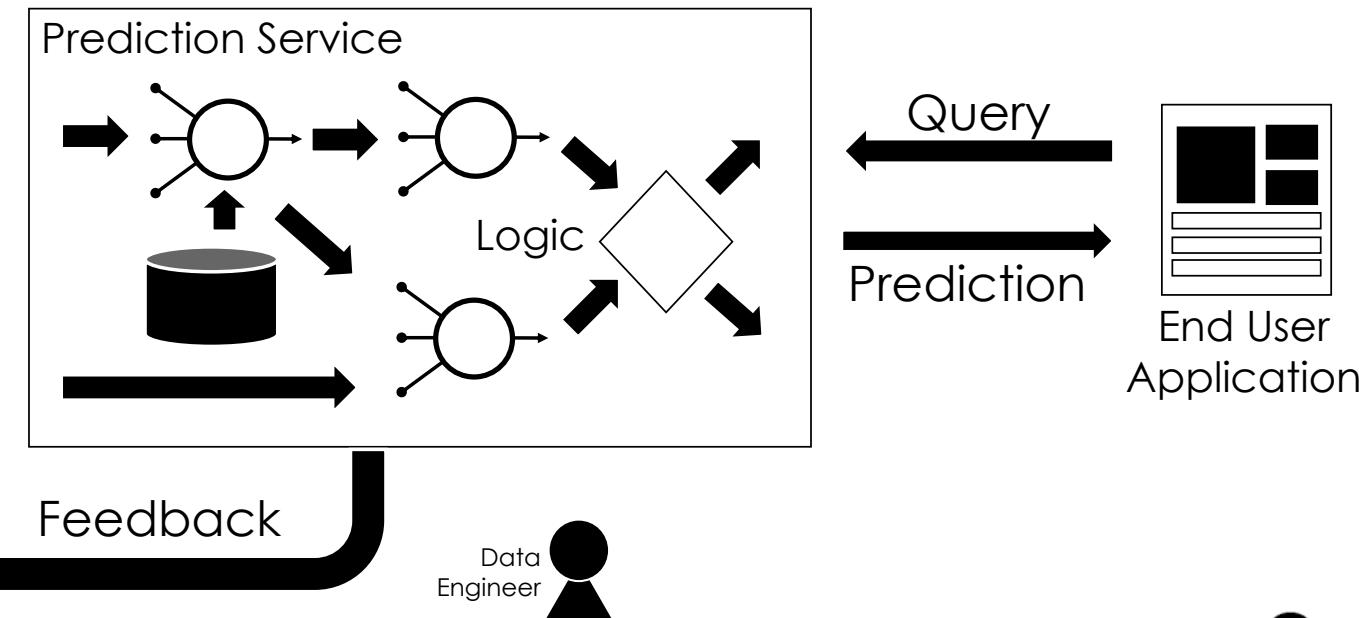
Feedback

**Goal:** make predictions in  
~10ms under heavy load

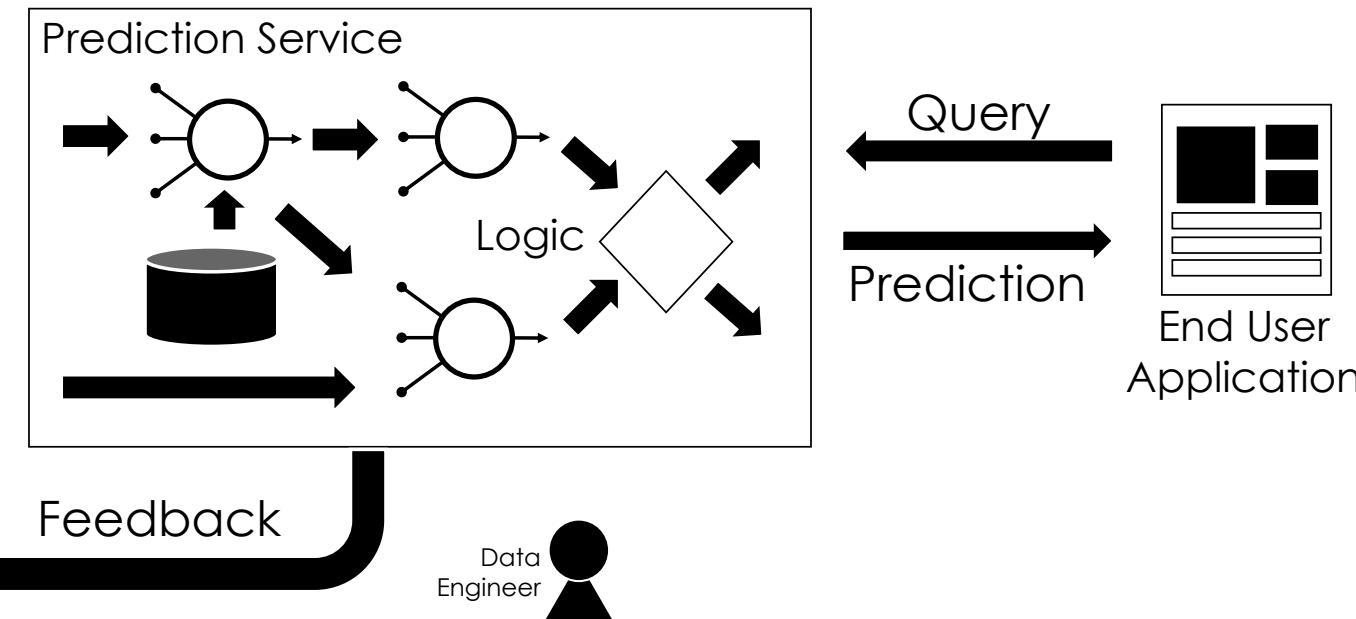


Complicated by **Deep Neural Networks**  
→ New **ML Algorithms** and **Systems**

# Inference Technologies



# Inference Technologies



Specialized in Particular Models or Frameworks



# Deploying Interactive Machine Learning Applications with Clipper

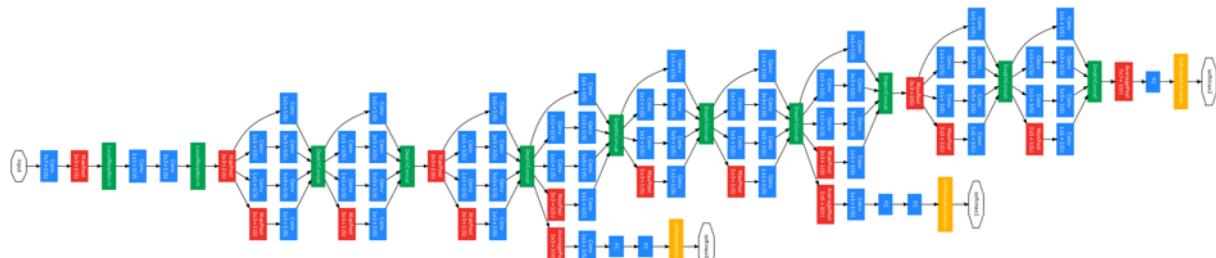


Joseph E. Gonzalez  
Co-director of the RISE Lab  
[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)

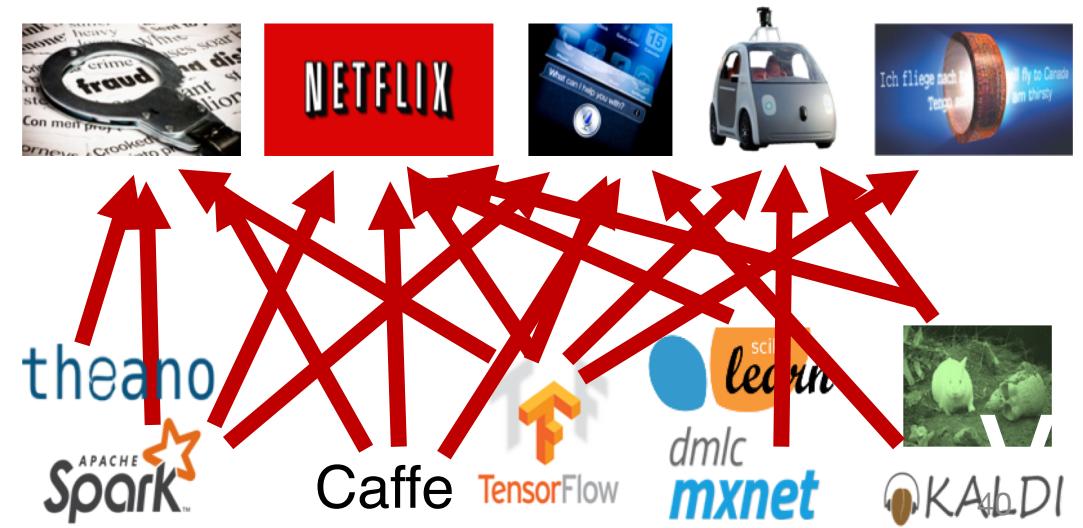
*The remainder of this talk ...*

- **Challenges** of prediction serving
- **Clipper architecture** overview
- **Open-source** system effort

# Prediction-Serving Challenges

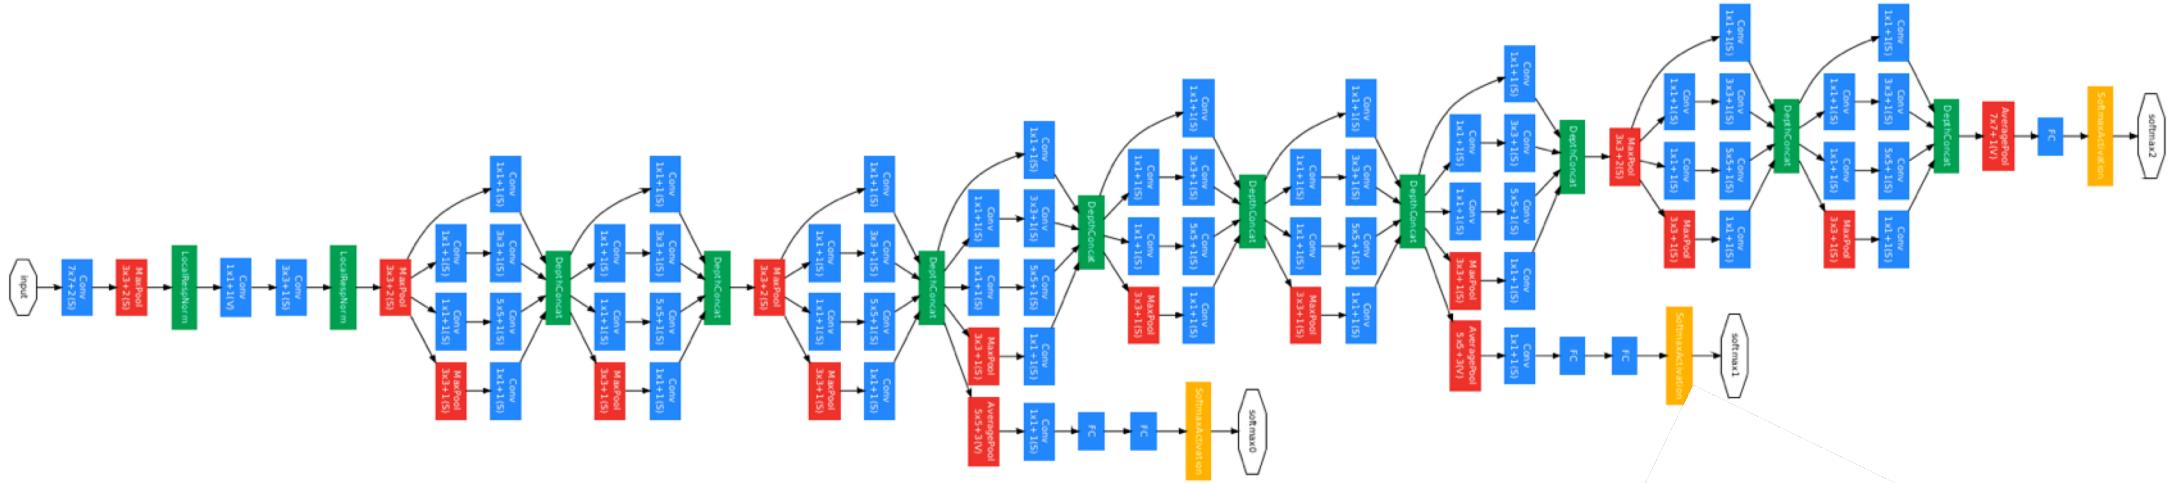


Support low-latency, high-throughput serving workloads



Large and growing ecosystem of ML models and frameworks

# Support low-latency, high-throughput serving workloads



## Models getting more complex

- 10s of GFLOPs [1]

## Deployed on critical path

- Maintain SLOs under heavy load

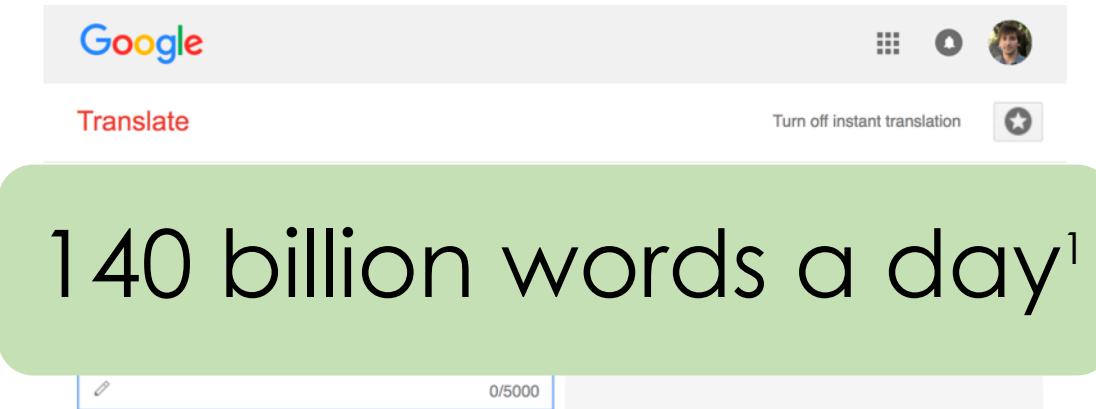


## Using specialized hardware for predictions

[1] Deep Residual Learning for Image Recognition. He et al. CVPR 2015.

# Google Translate

## Serving



82,000 GPUs  
running 24/7

Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation

Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le, Mohammad Norouzi  
`yonghui,schuster,zhifengc,qvl,mnorouzi@google.com`

Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey, Jeff Klingner, Apurva Shah, Melvin Johnson, Xiaobing Liu, Łukasz Kaiser, Stephan Gouws, Yoshikiyo Kato, Taku Kudo, Hideto Kazawa, Keith Stevens, George Kurian, Nishant Patil, Wei Wang, Cliff Young, Jason Smith, Jason Riesa, Alex Rudnick, Oriol Vinyals, Greg Corrado, Macduff Hughes, Jeffrey Dean

*"If each of the **world's Android phones** used the new Google voice search for just **three minutes a day**, these engineers realized, the company would **need twice as many data centers.**"*  
– Wired

**Designed New Hardware!  
Tensor Processing Unit (TPU)**

# Building Application Specific Systems

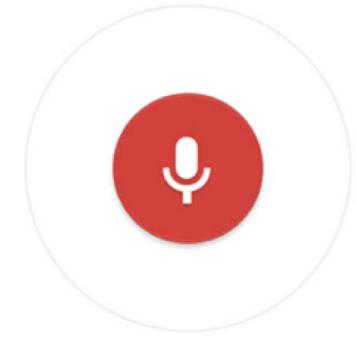


**NETFLIX**

**Google Ads**

**CONVIVA®**

OK GOOGLE



# Building Application Specific Systems

## Problems:

- **Expensive to build and maintain**
- **Require ML and systems expertise**
- **Tightly-coupled** model and application
  - Difficult to change or update application
- Only supports **single ML framework**

# Growing ecosystem of ML Frameworks

Fraud  
Detection



Content  
Rec.



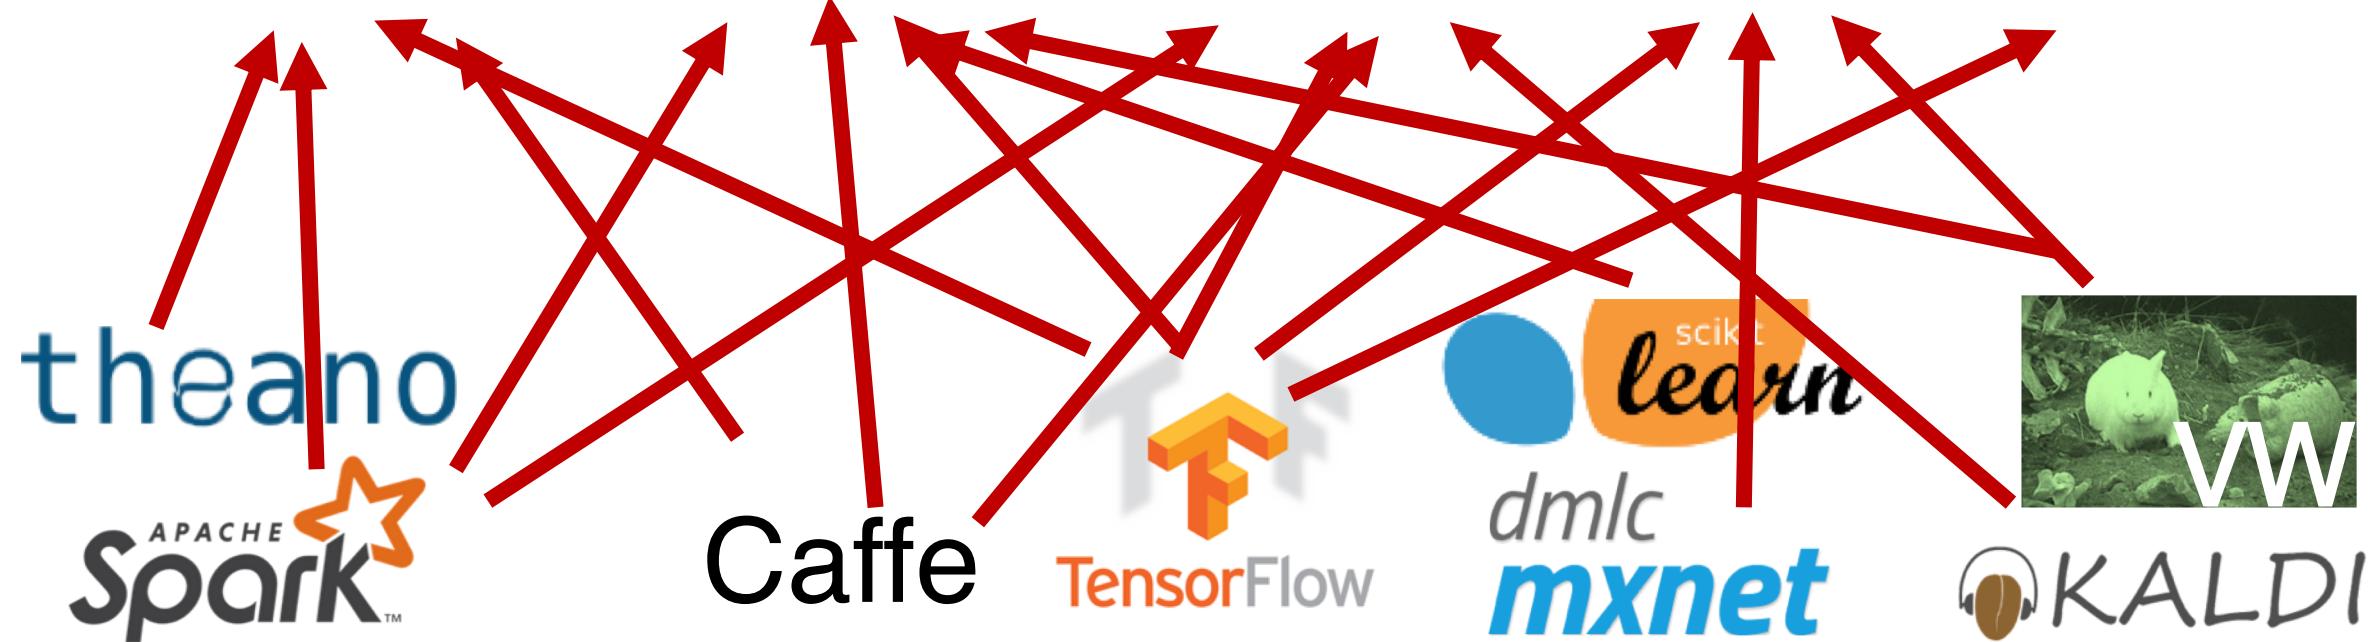
Personal  
Asst.



Robotic  
Control



Machine  
Translation

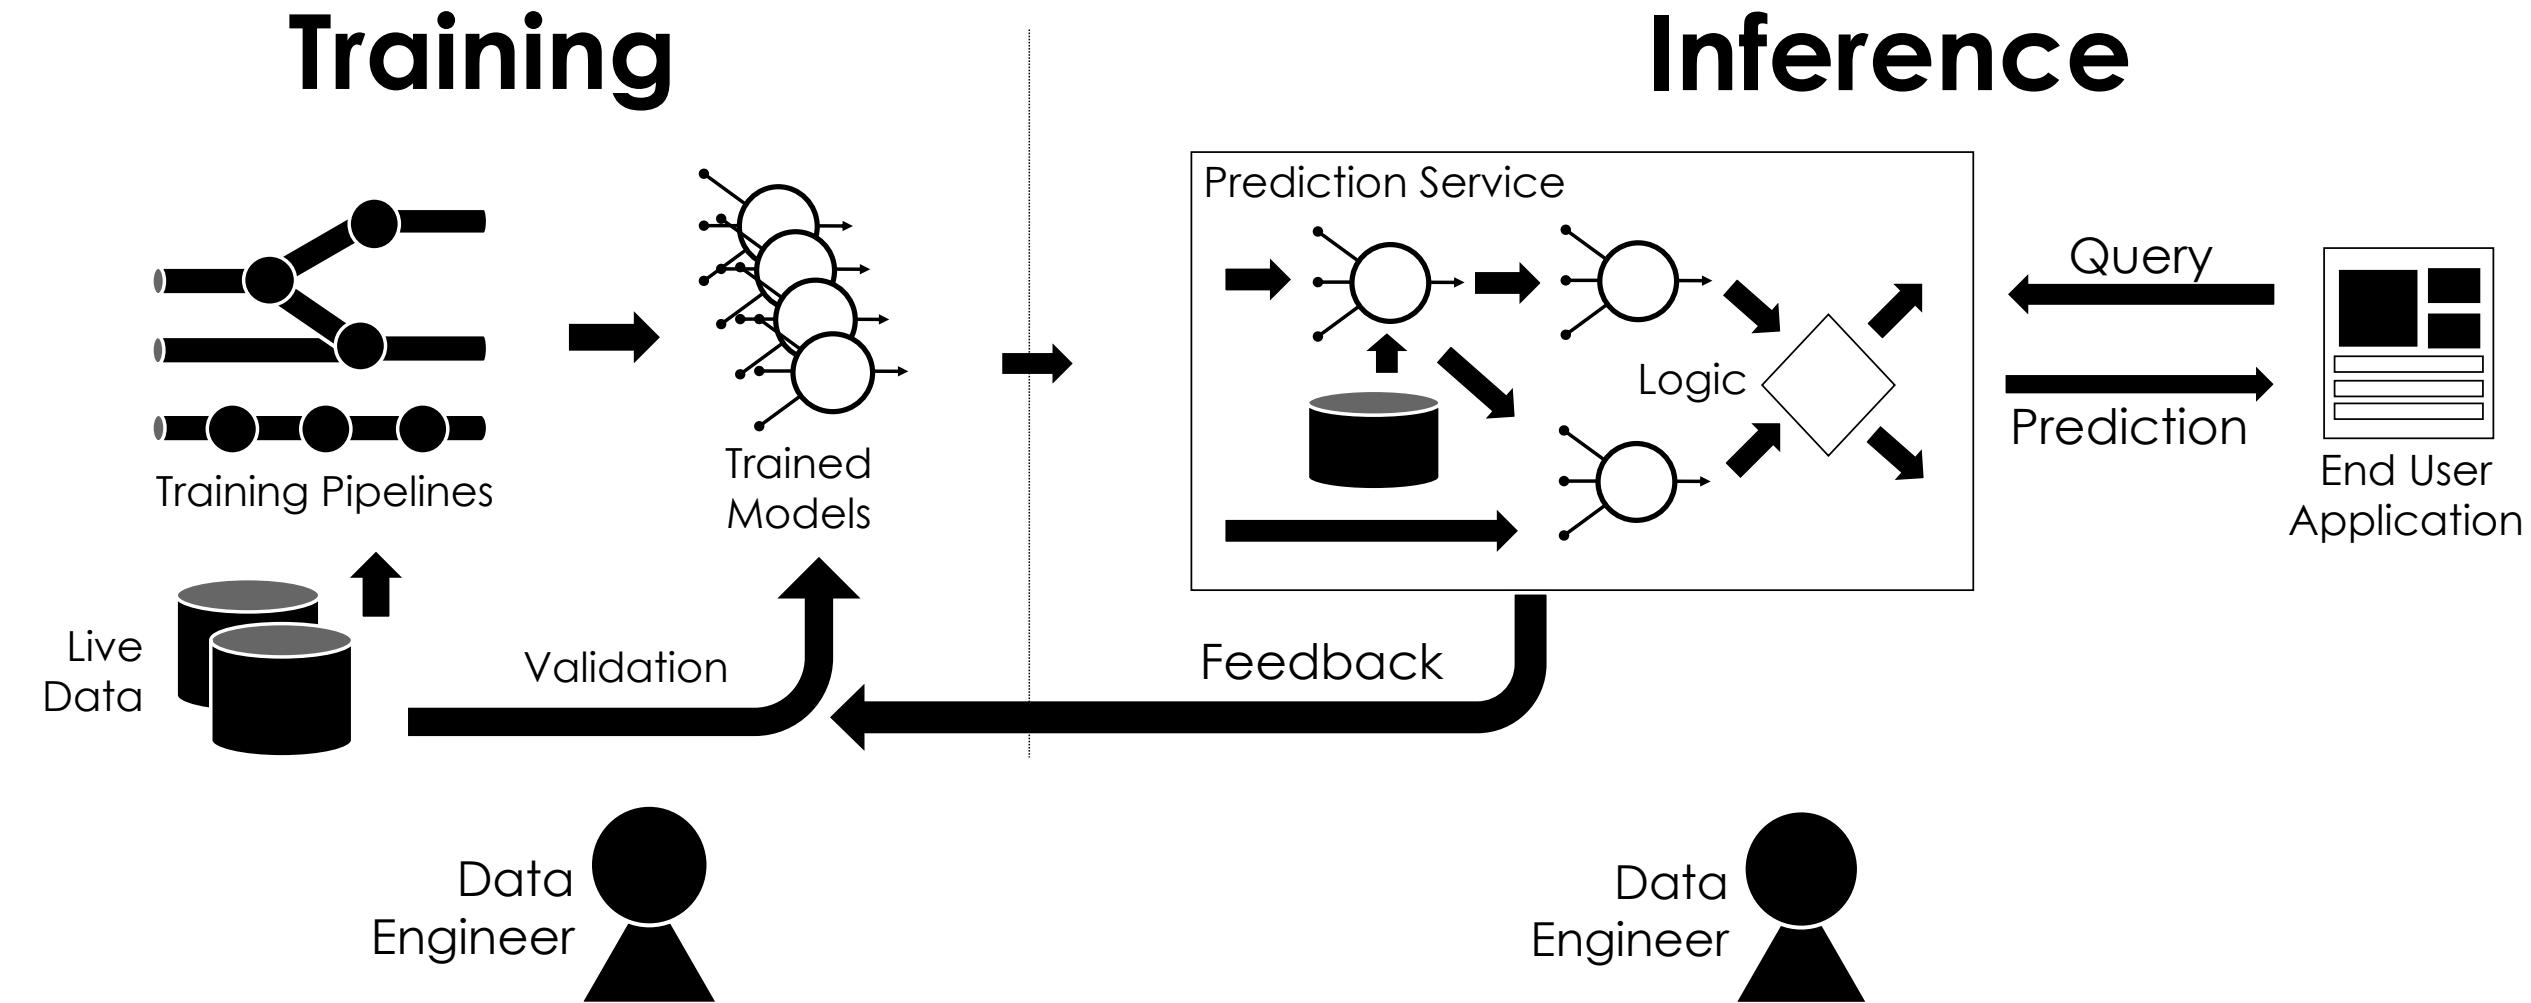


**Building & maintaining  
separate serving systems  
for each framework  
is expensive!**

## Solution

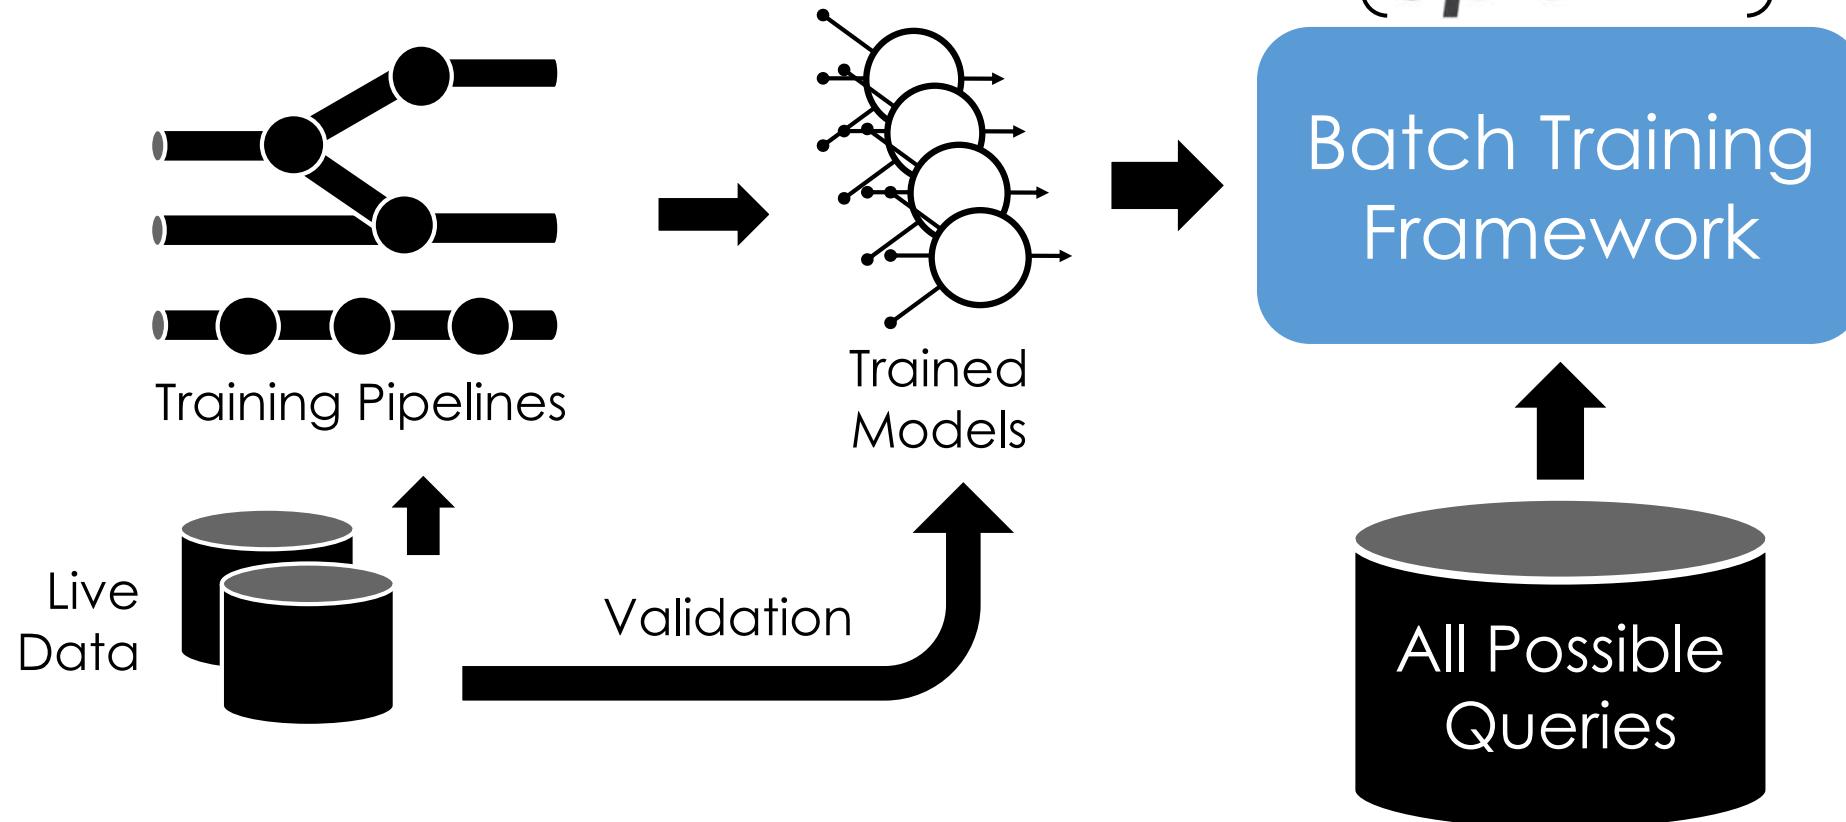
**Pre-materialize** predictions into a  
low latency **Data Management System**

# Pre-materialized Predictions



# Pre-materialized Predictions

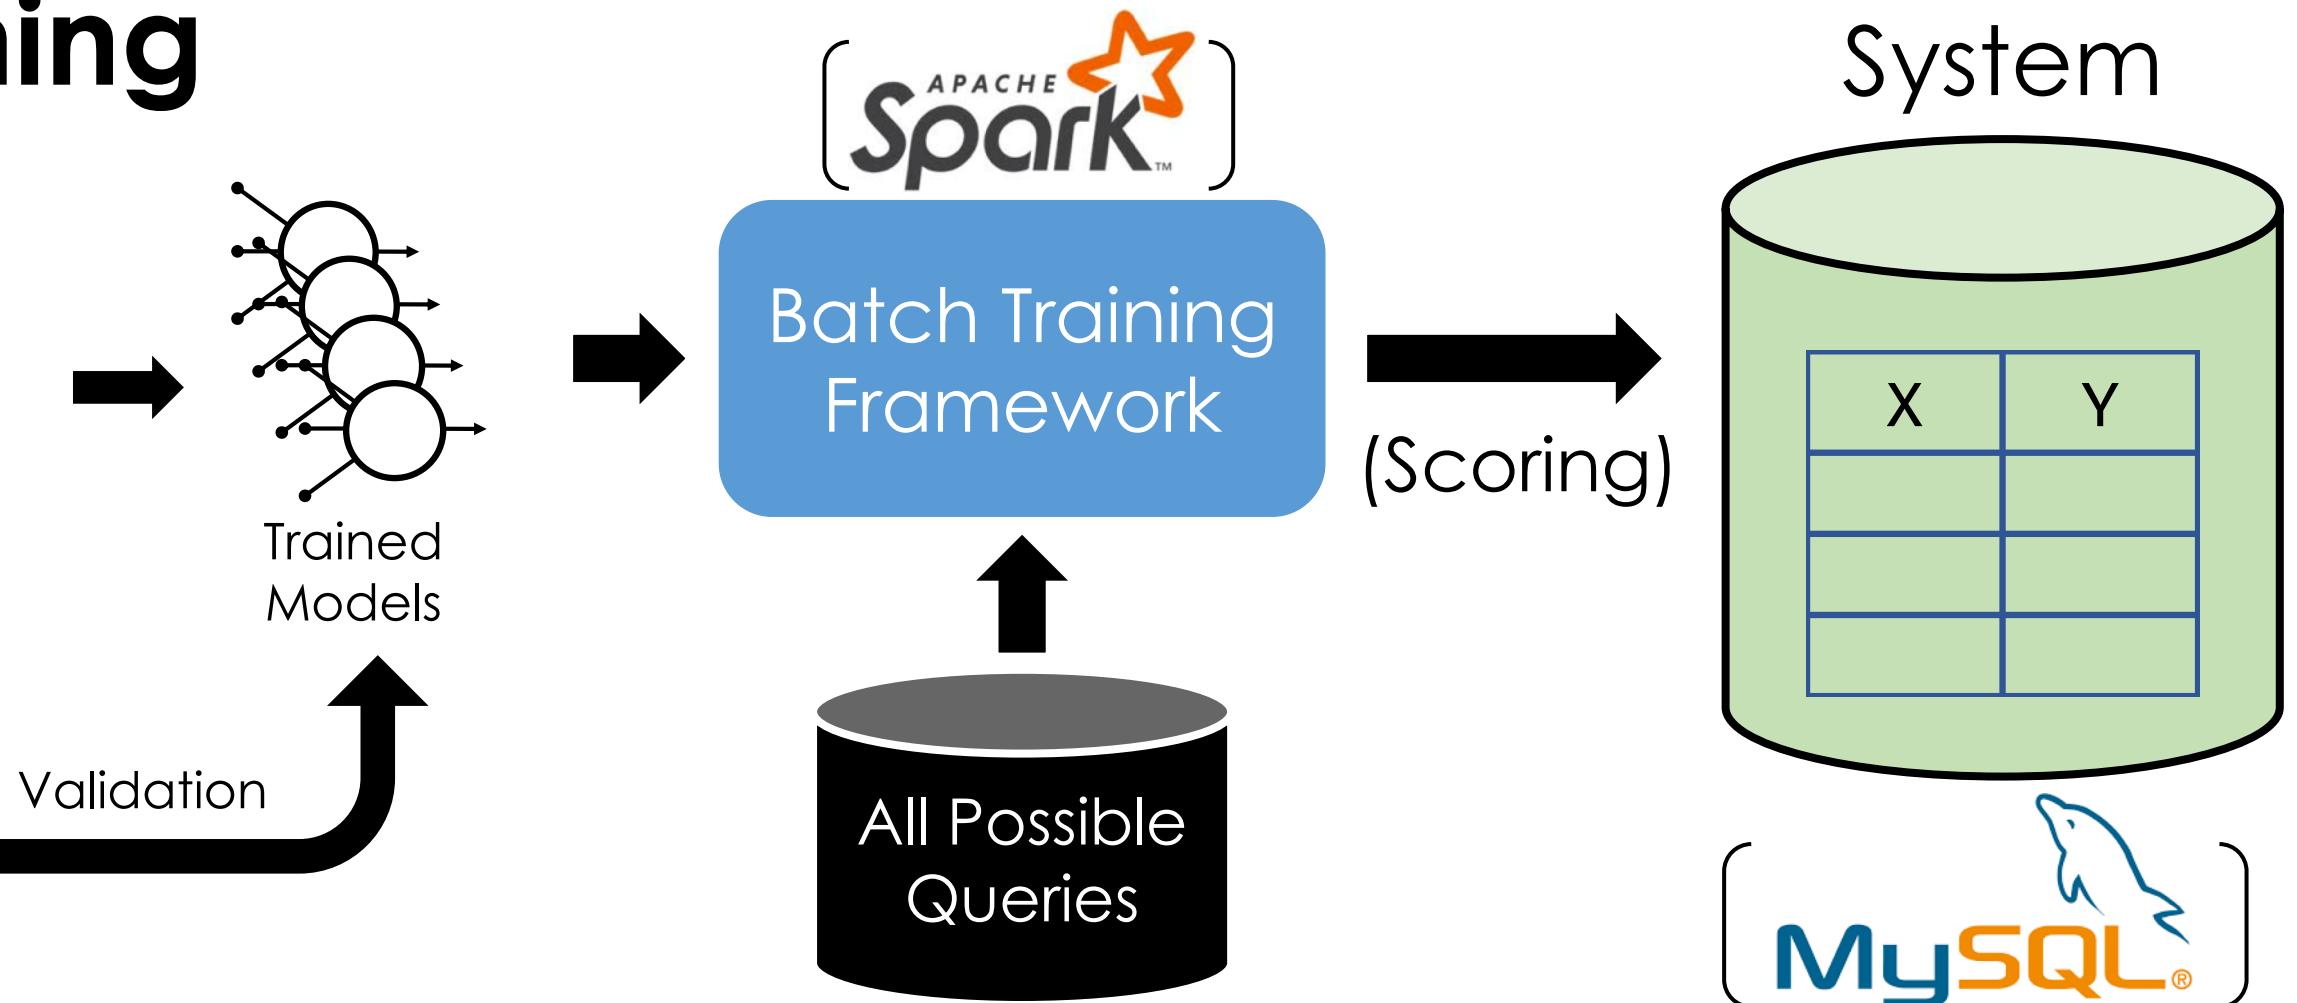
## Training



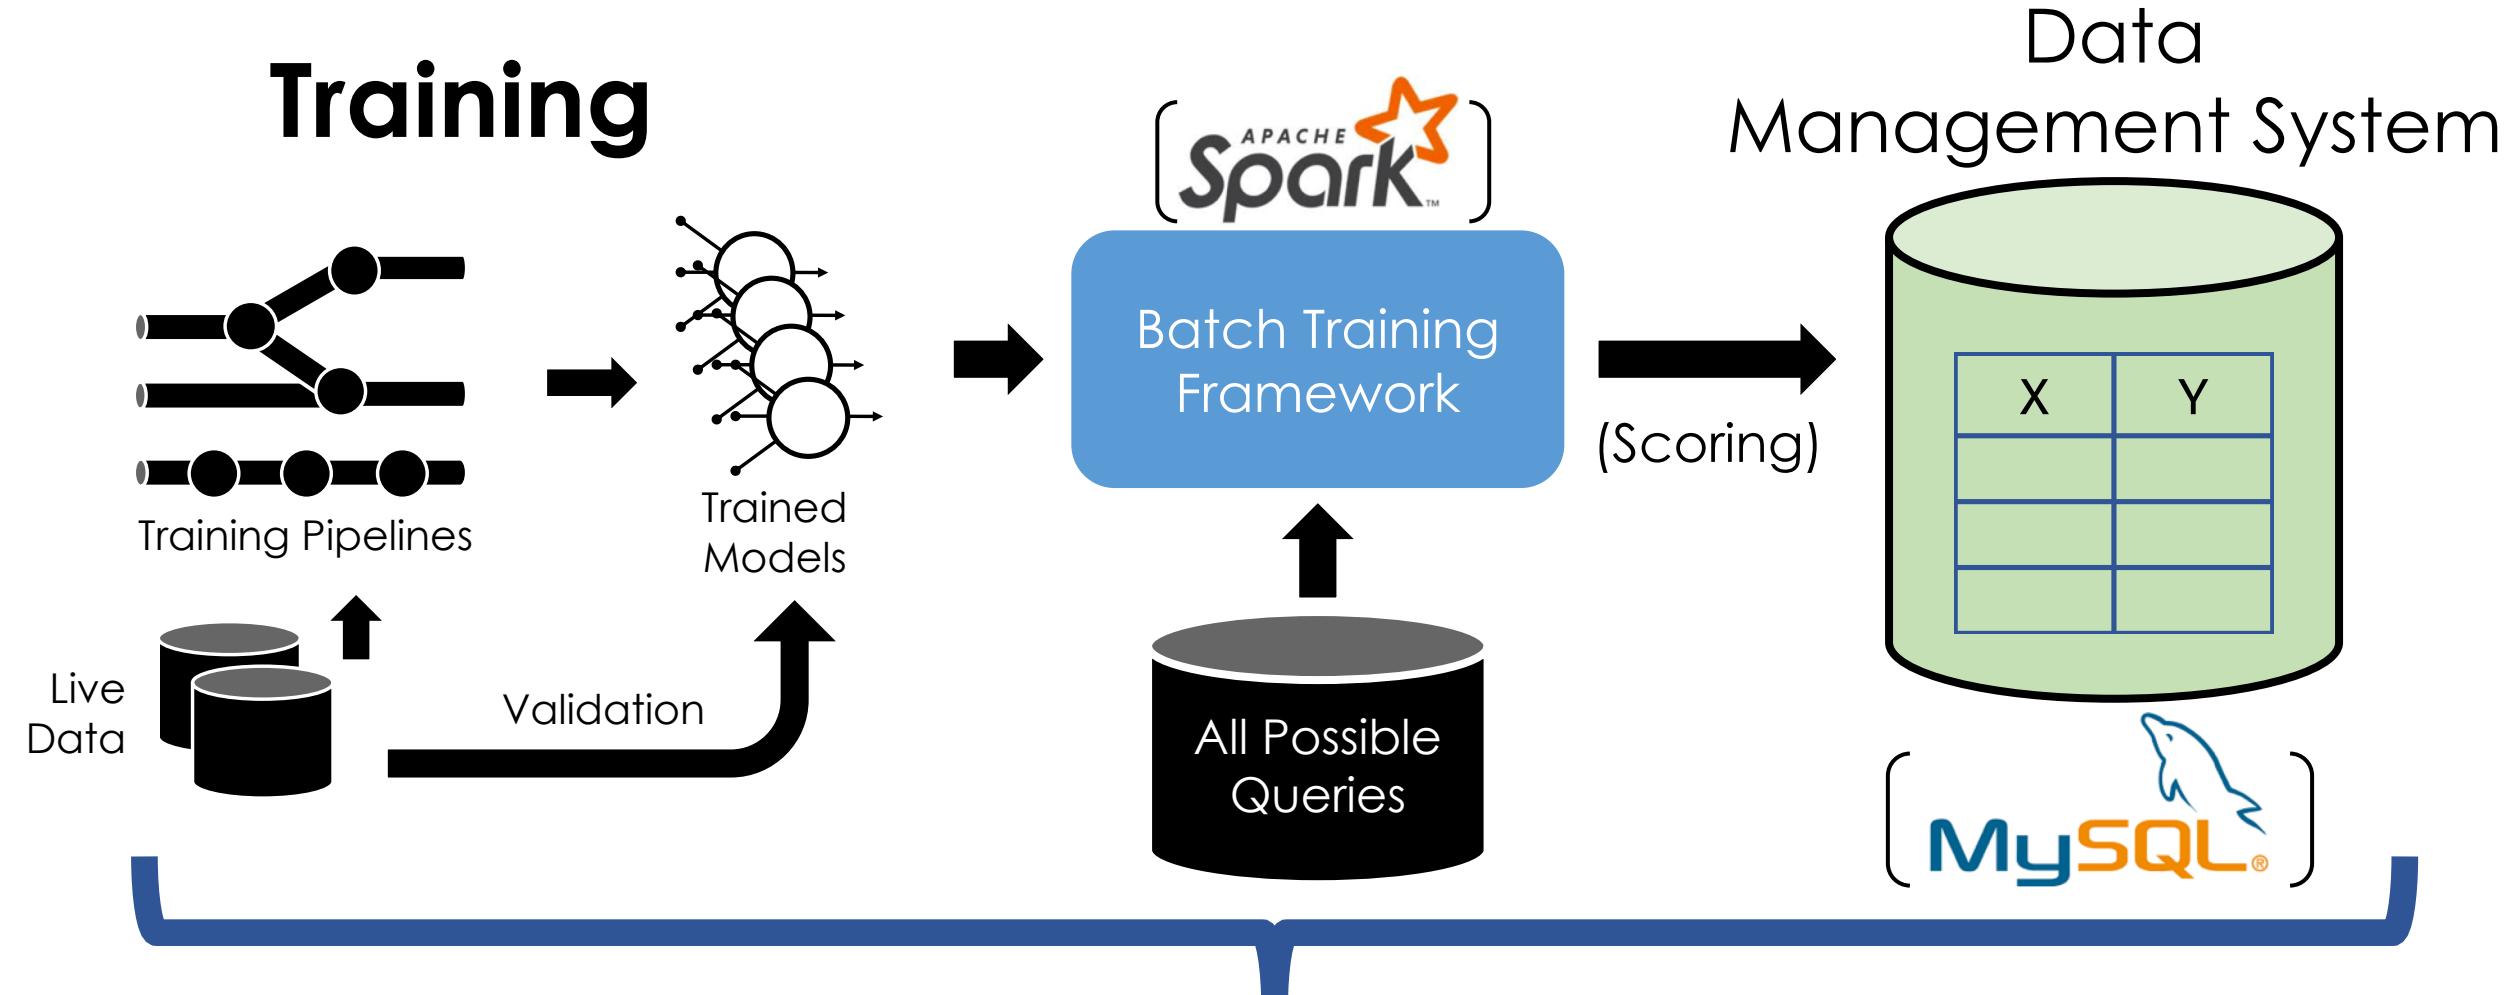
# Training

## Pre-materialized Predictions

Data Management System

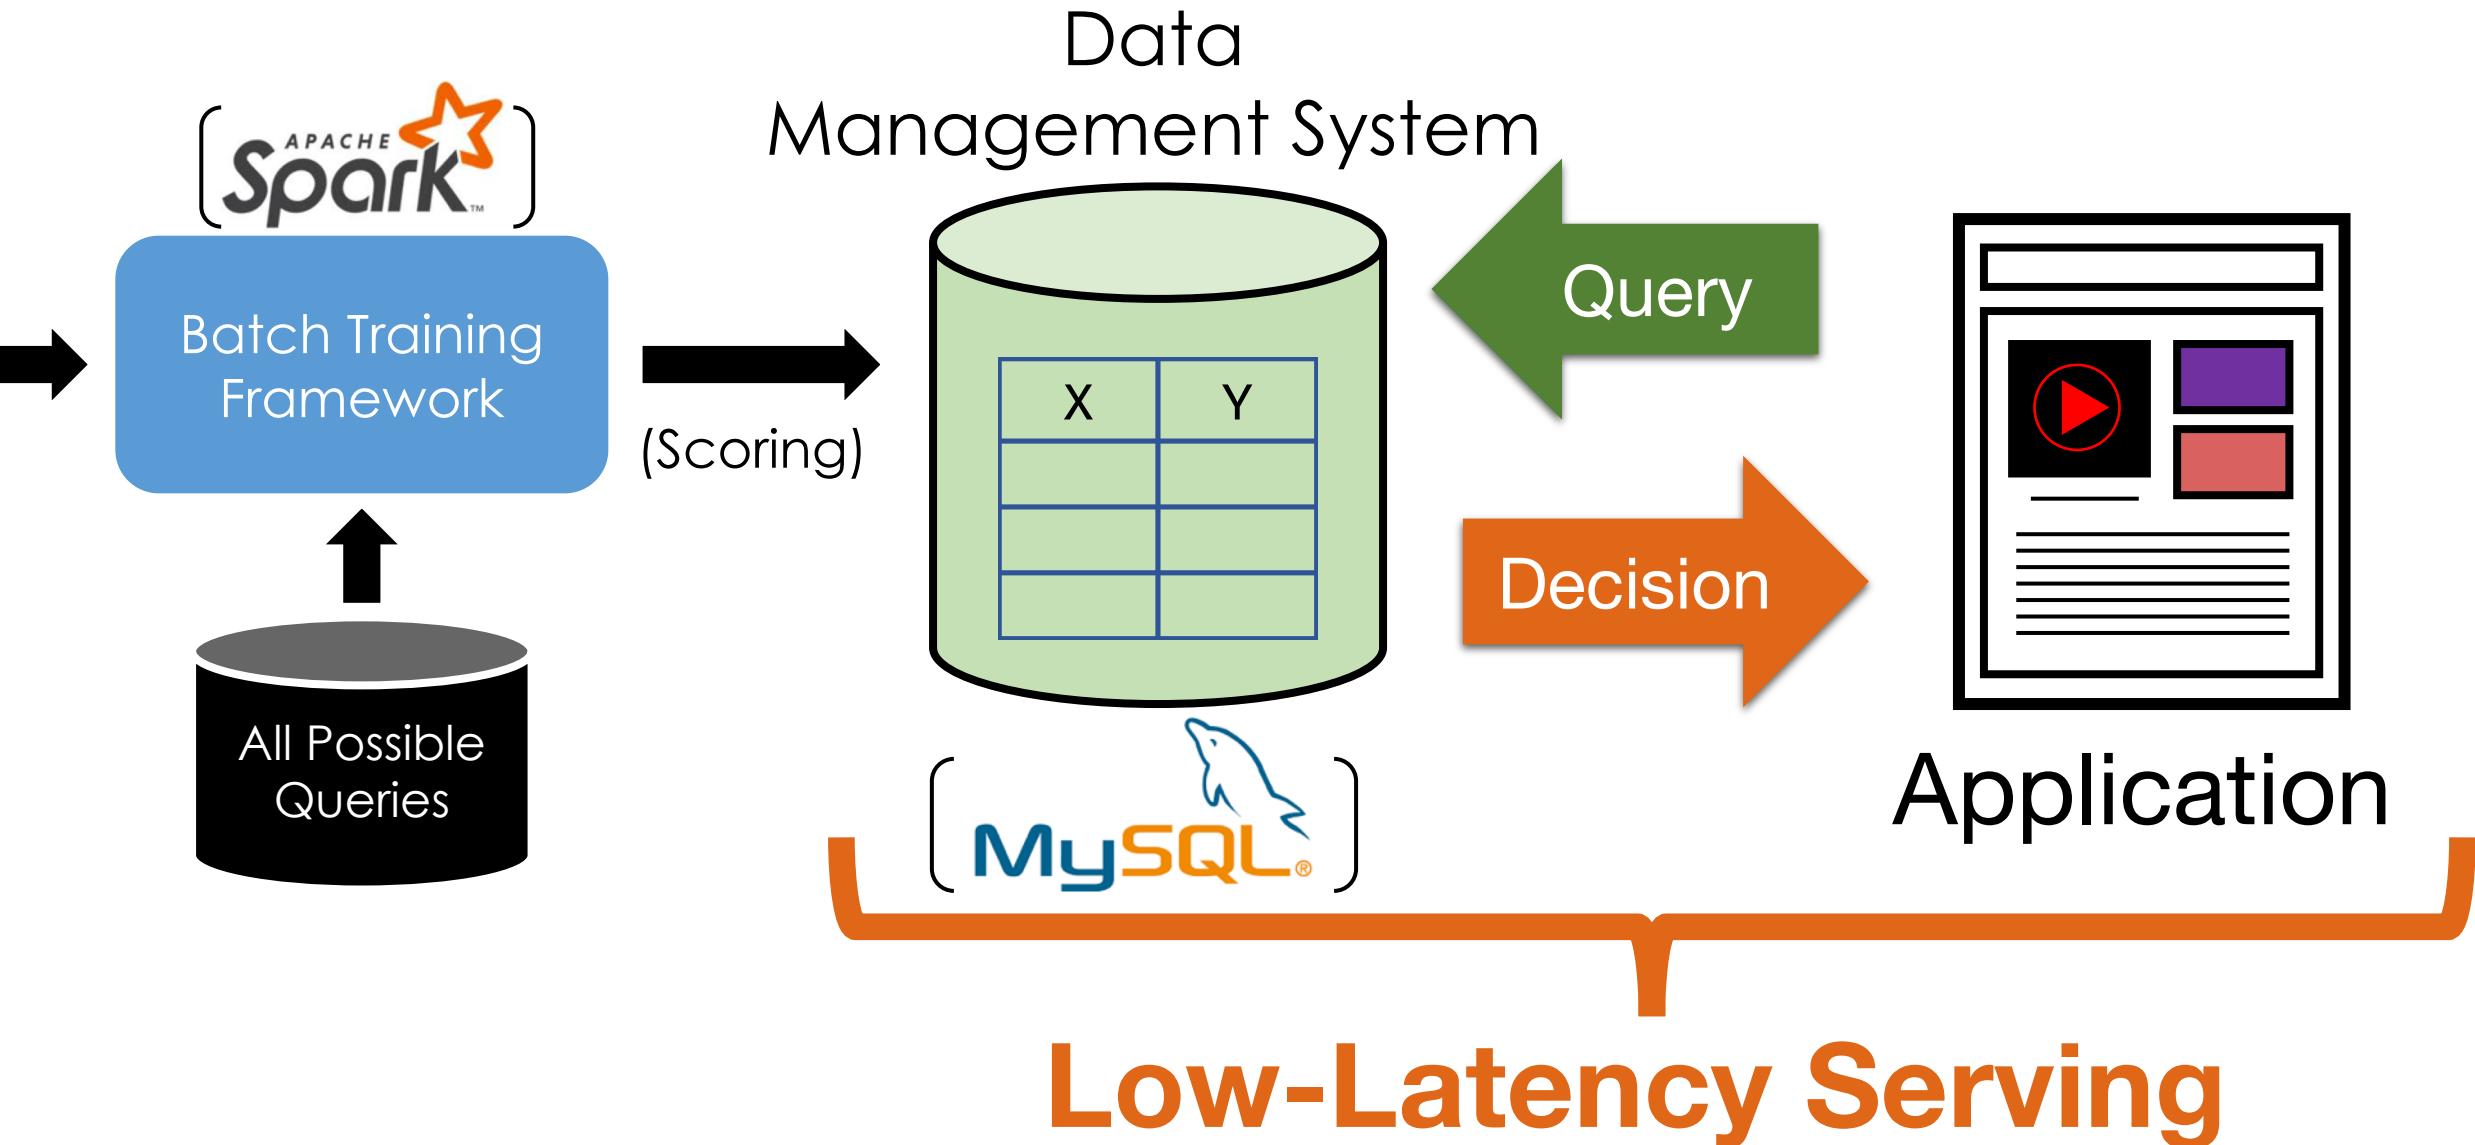


# Pre-materialized Predictions



Standard Data Eng. Tools

# Serving Pre-materialized Predictions

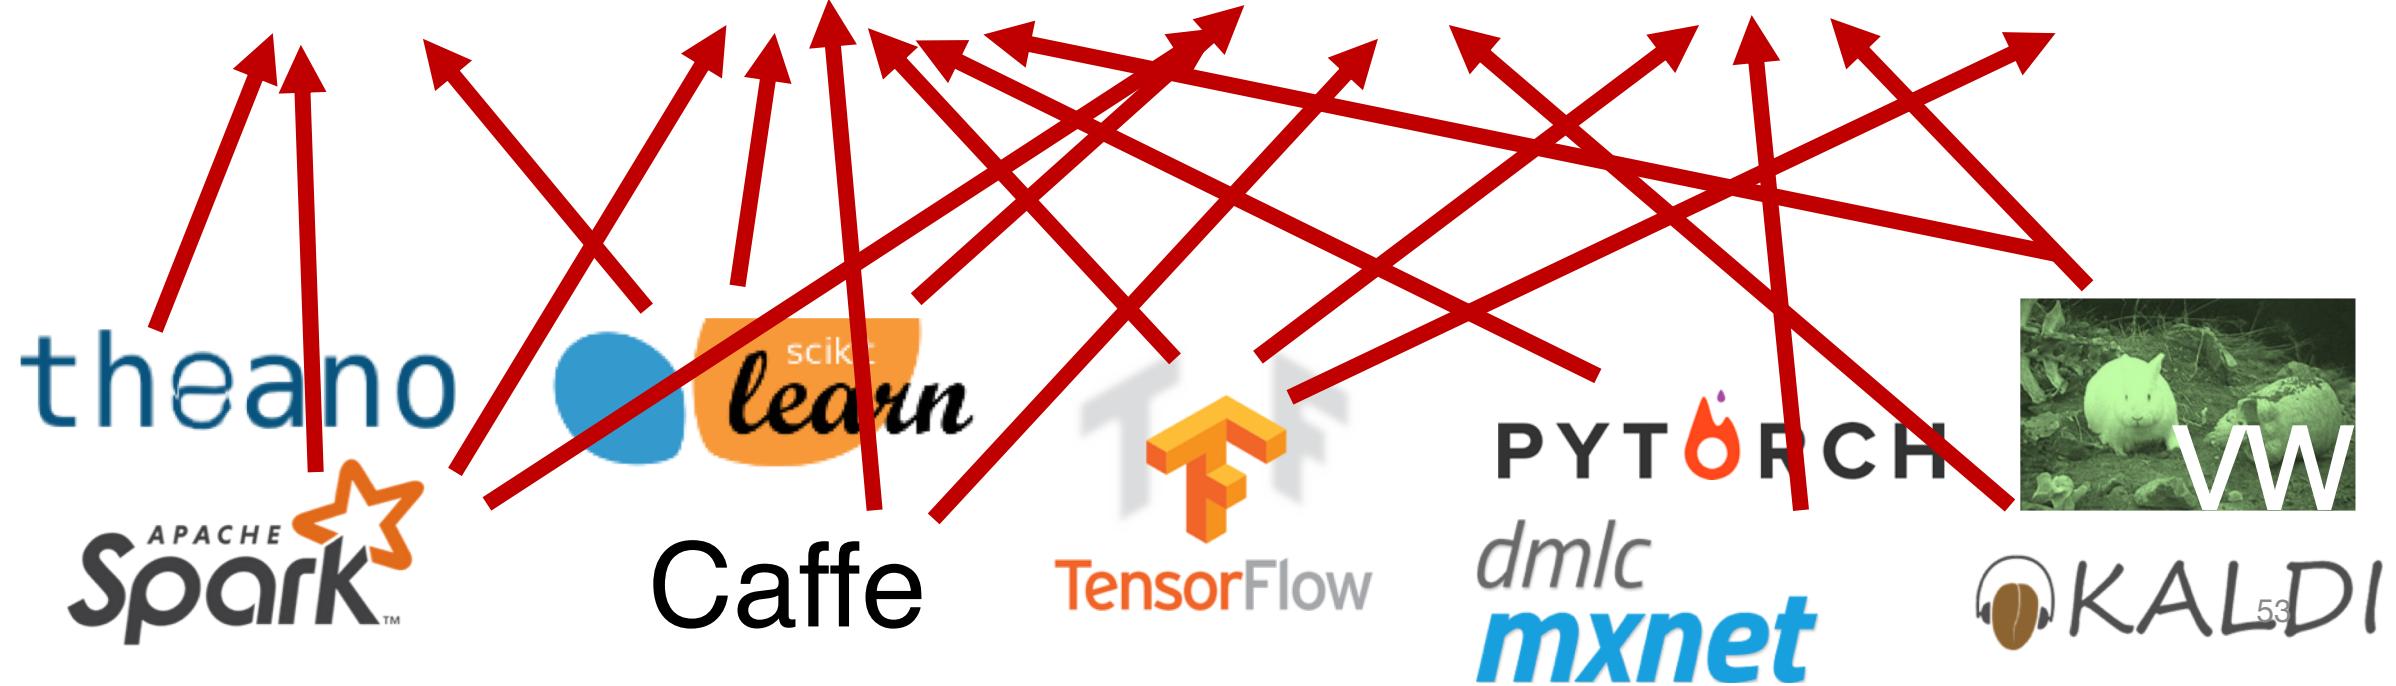


# Serving Pre-materialized Predictions

## Problems:

- Requires full set of **queries ahead of time**
  - Small and **bounded input domain**
- Requires substantial **computation** and **space**
  - Example: scoring all content for all customers!
- Costly update → rescore everything!

# Wide range of application and frameworks





---

# Middle layer for prediction serving.

Common  
Abstraction

System  
Optimizations

---

theano  
APACHE  
Spark™

 scikit-learn

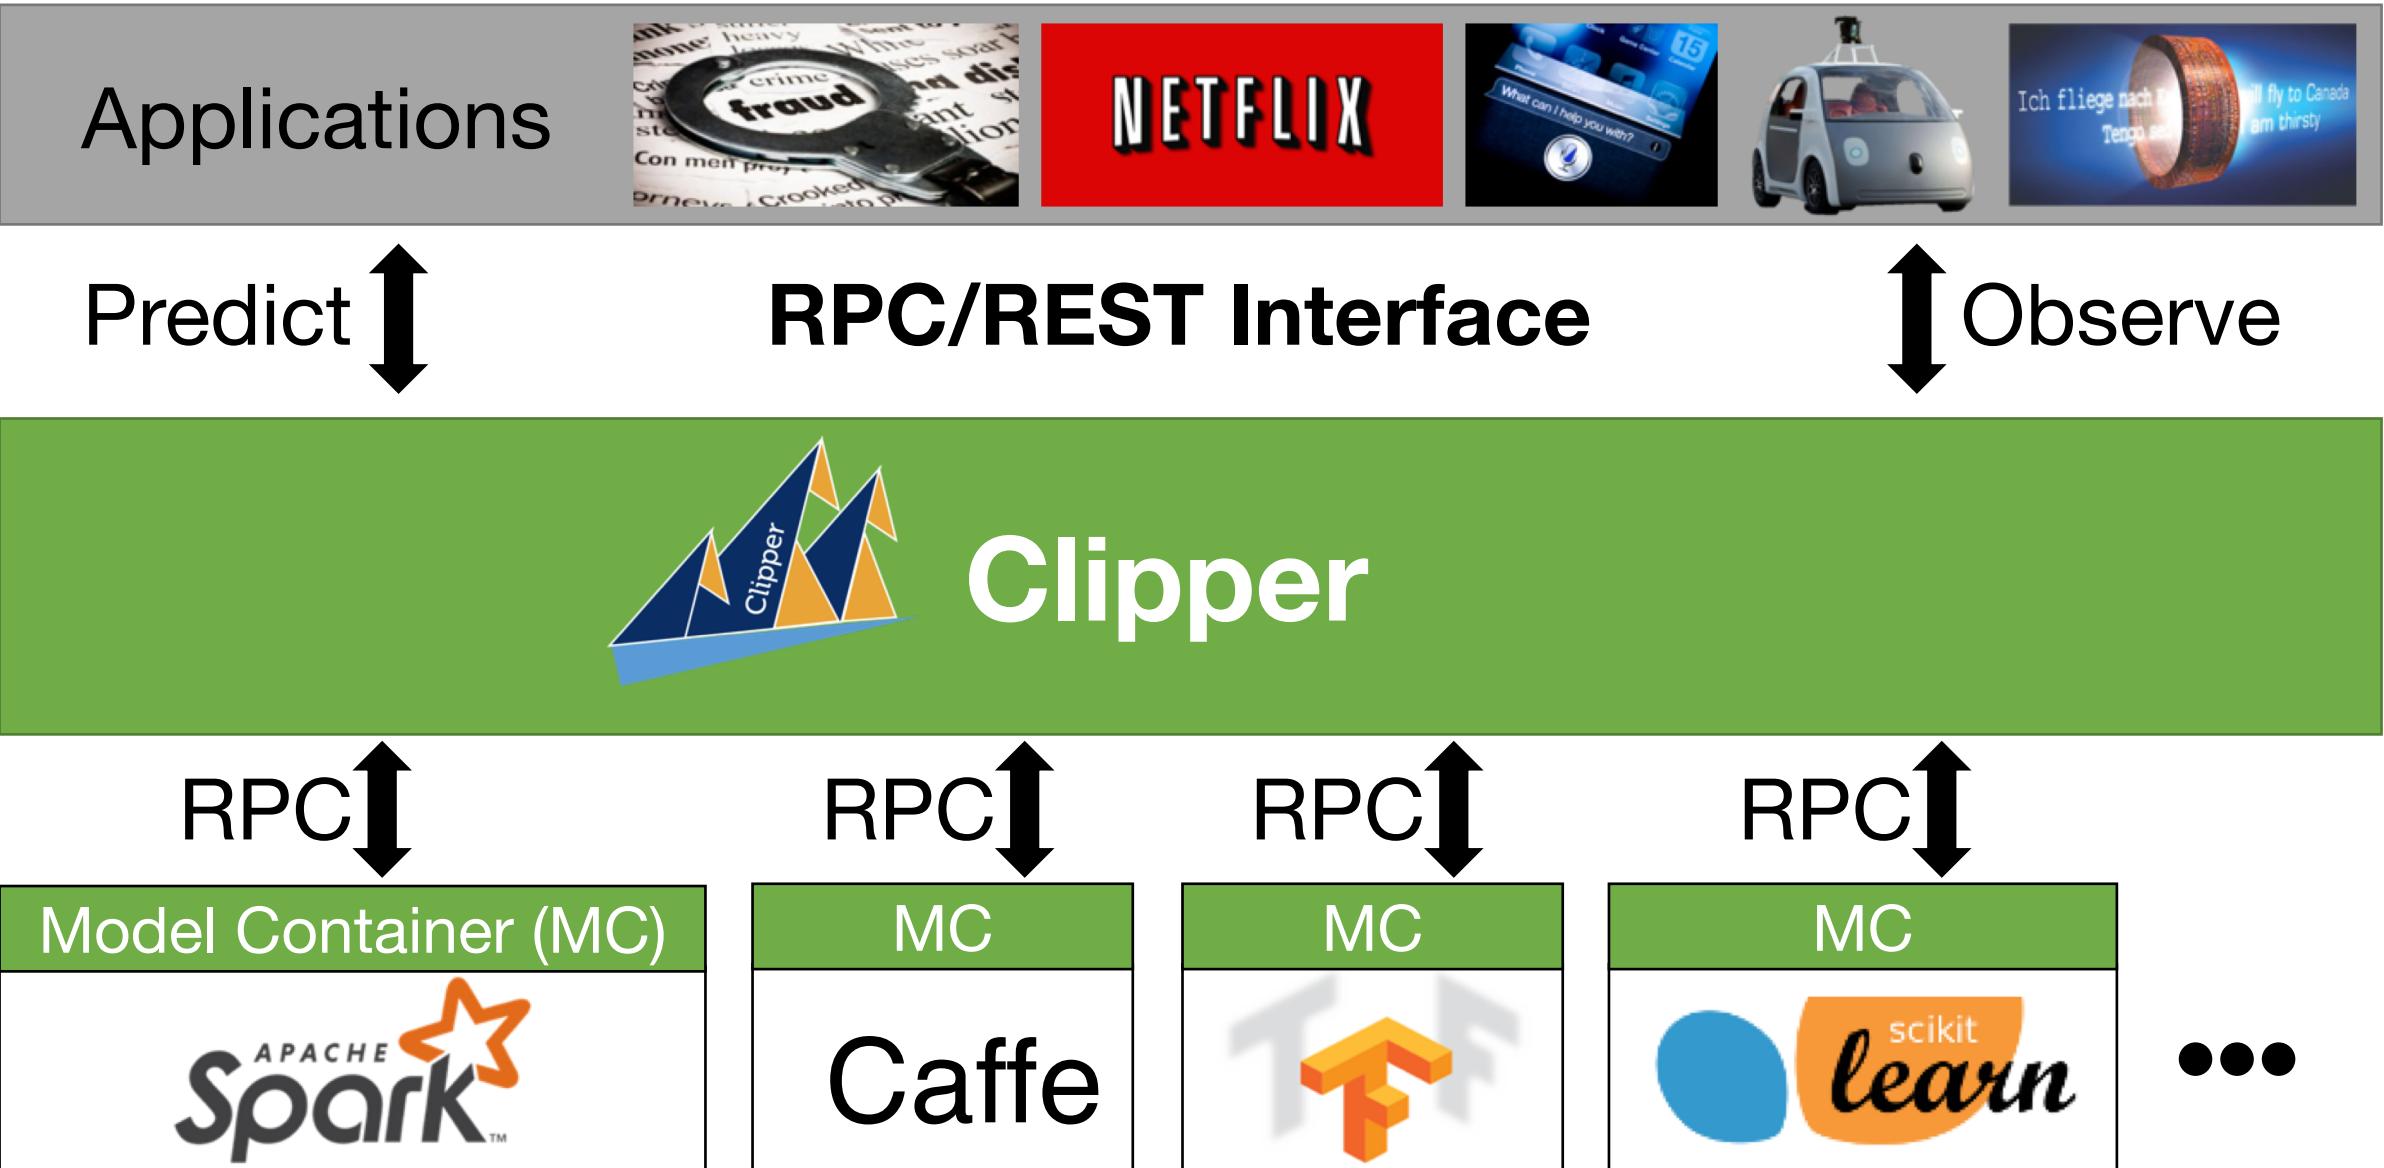
Caffe

 TensorFlow

PYTORCH  
dmlc  
mxnet

 VW KALDI

# Clipper Decouples Applications and Models





Predict 

## RPC/REST Interface

Observe 



# Clipper

- **Core system:** 10K lines of C++ and 8K lines of Python
- Open Source (**Apache License**) – <http://clipper.ai>
- Designed to support **production level** query traffic
  - Deliver low + predictable latency
  - Research goal: **study reality** ...

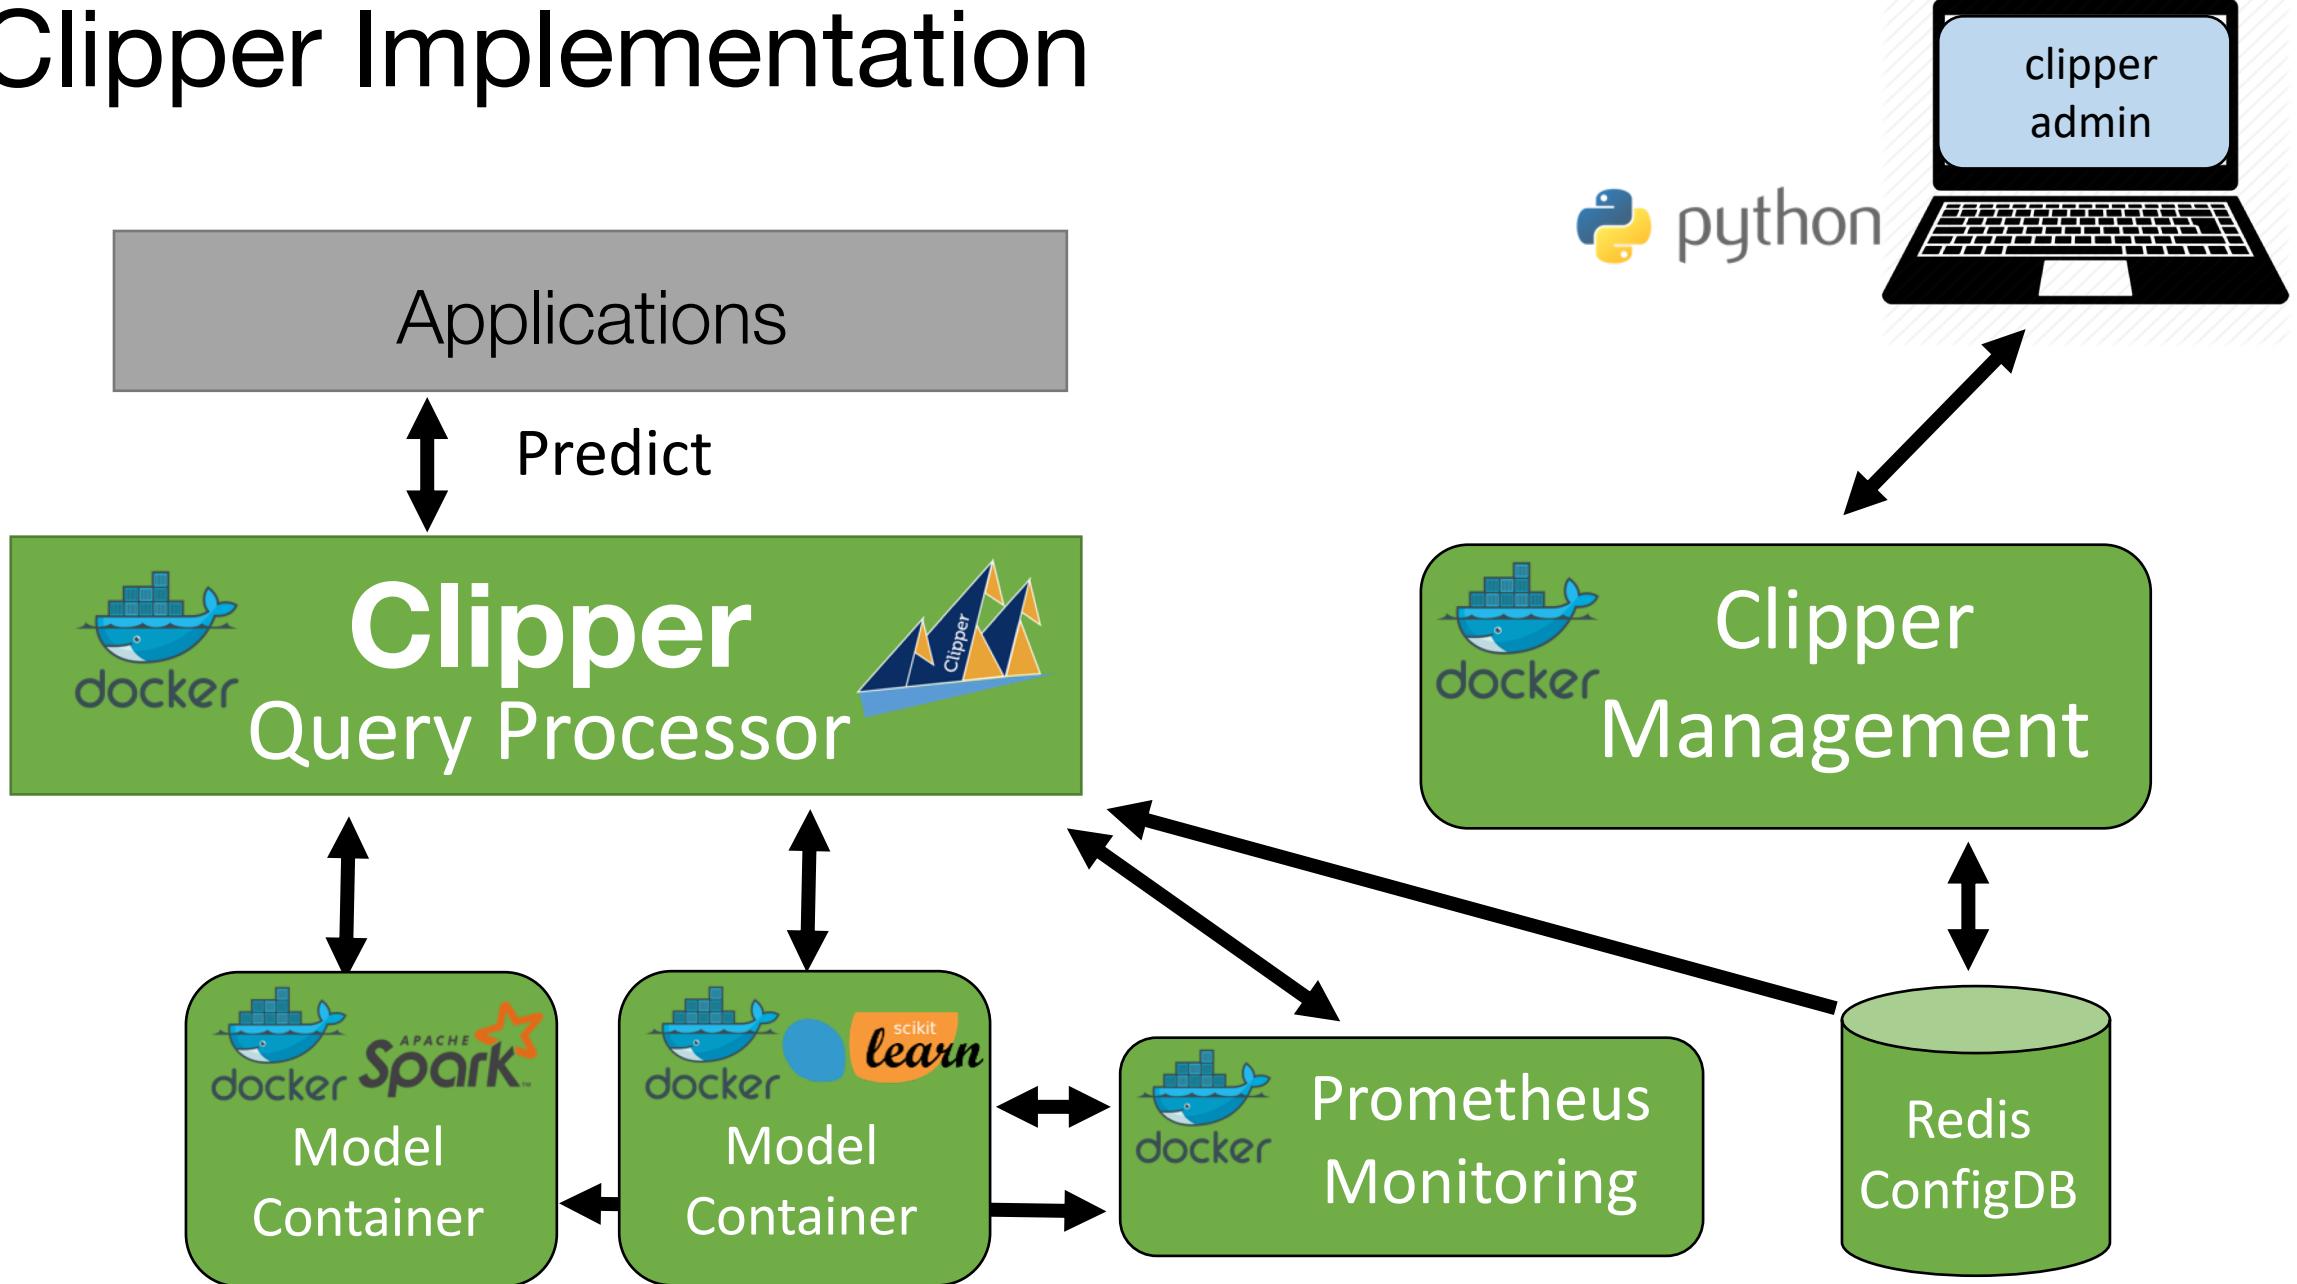
RPC 

RPC 

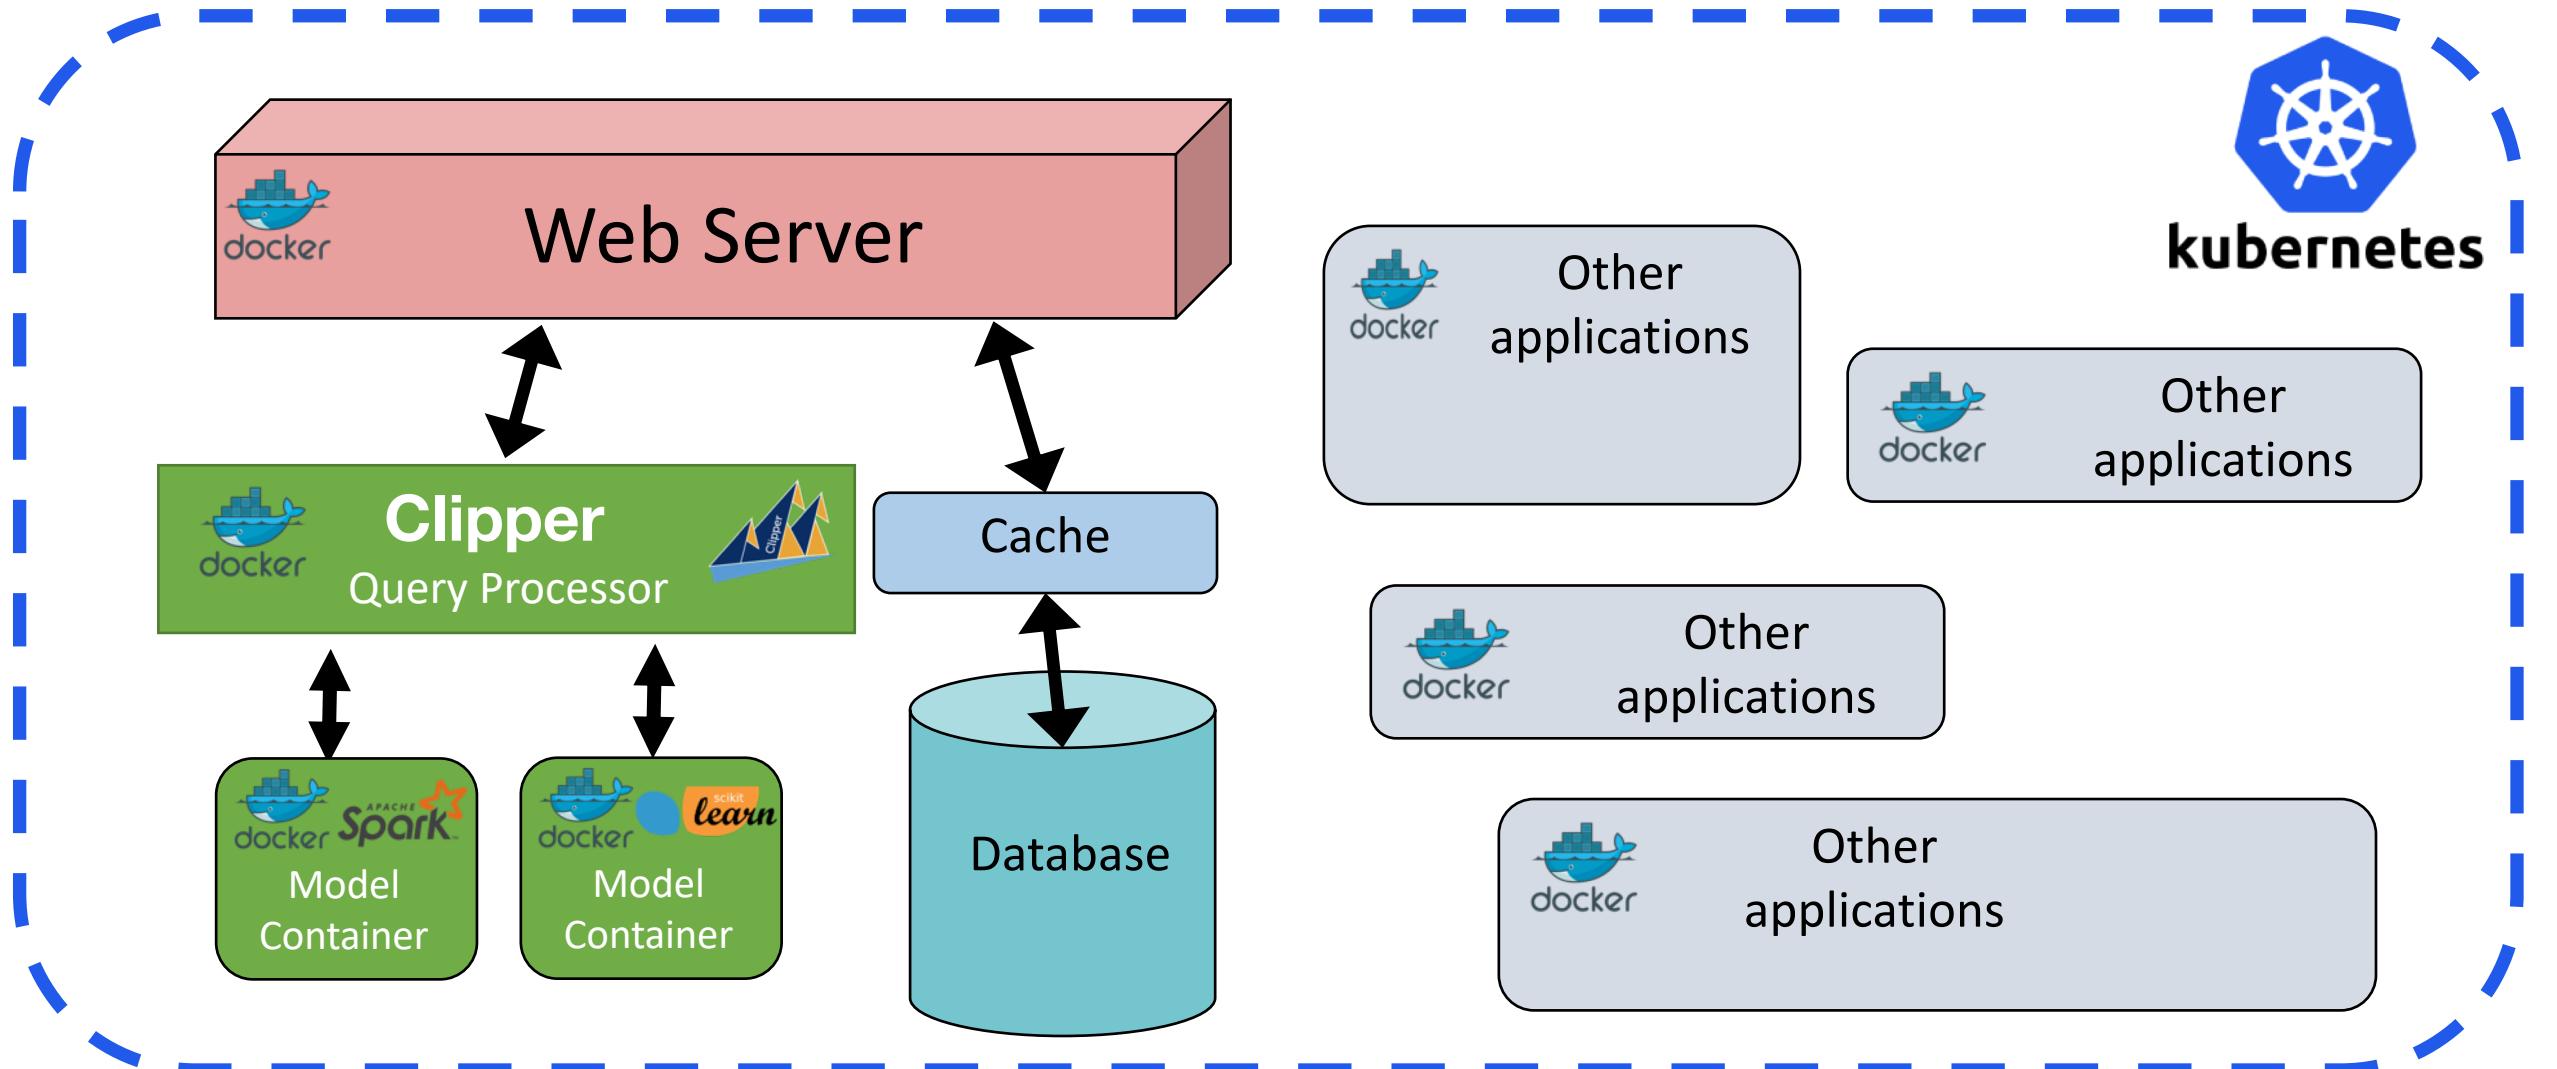
RPC 

RPC 

# Clipper Implementation



# Run alongside other applications with **Kubernetes**



# Getting Started with Clipper

**Tutorials** at <http://clipper.ai>

Docker images available **on DockerHub**

Clipper admin is distributed as **pip package**:

```
pip install clipper_admin
```

Get up and running **without compiling**

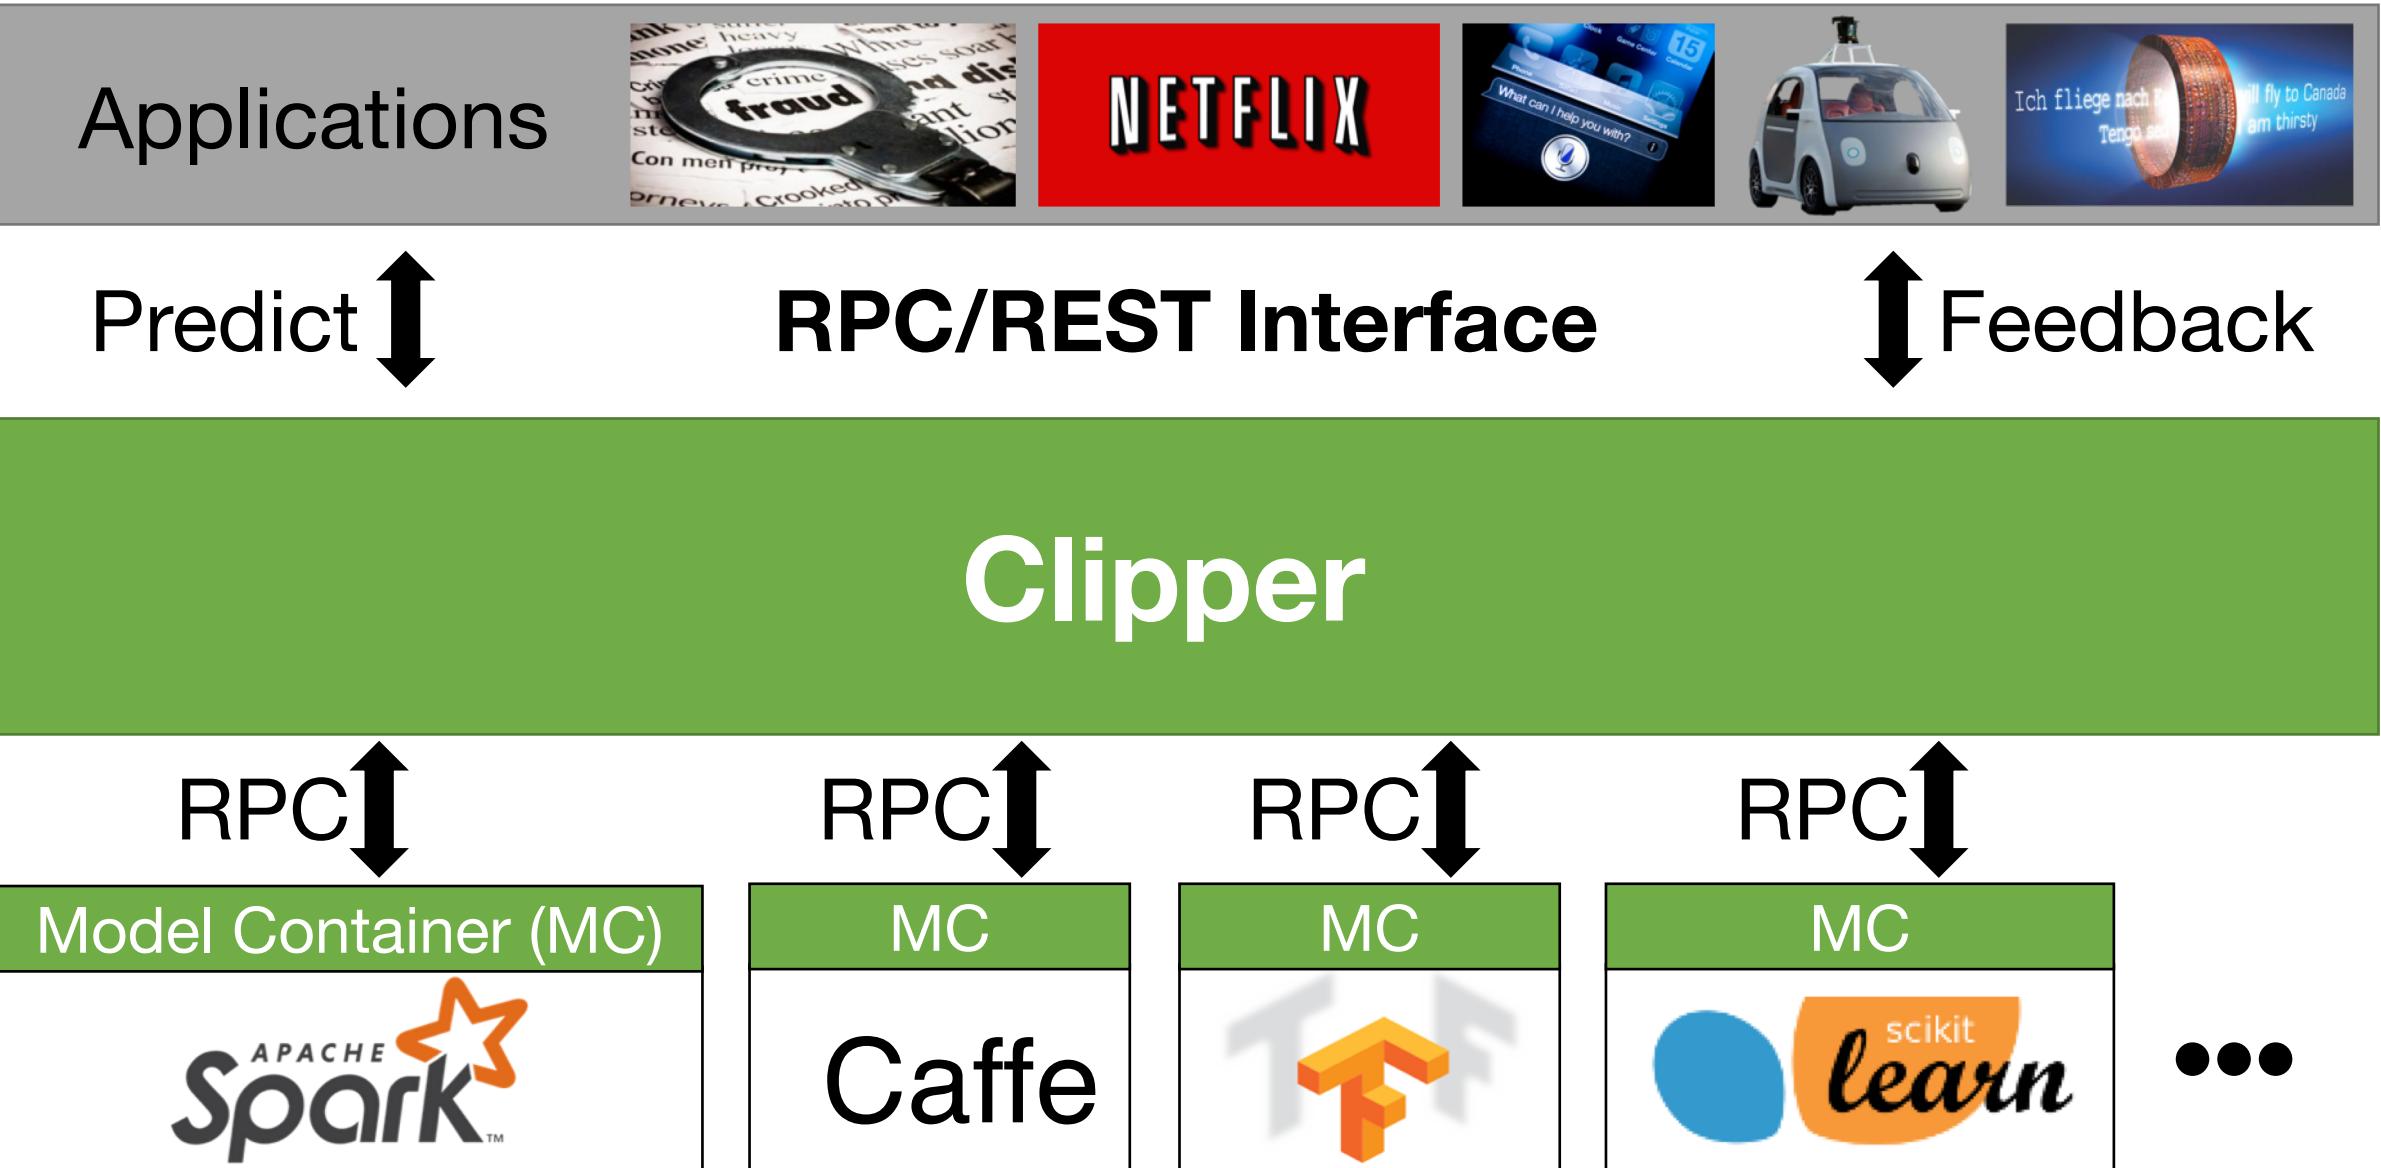
# Clipper Design Innovations

**Containerized frameworks:** unified abstraction and framework level isolation and scaling

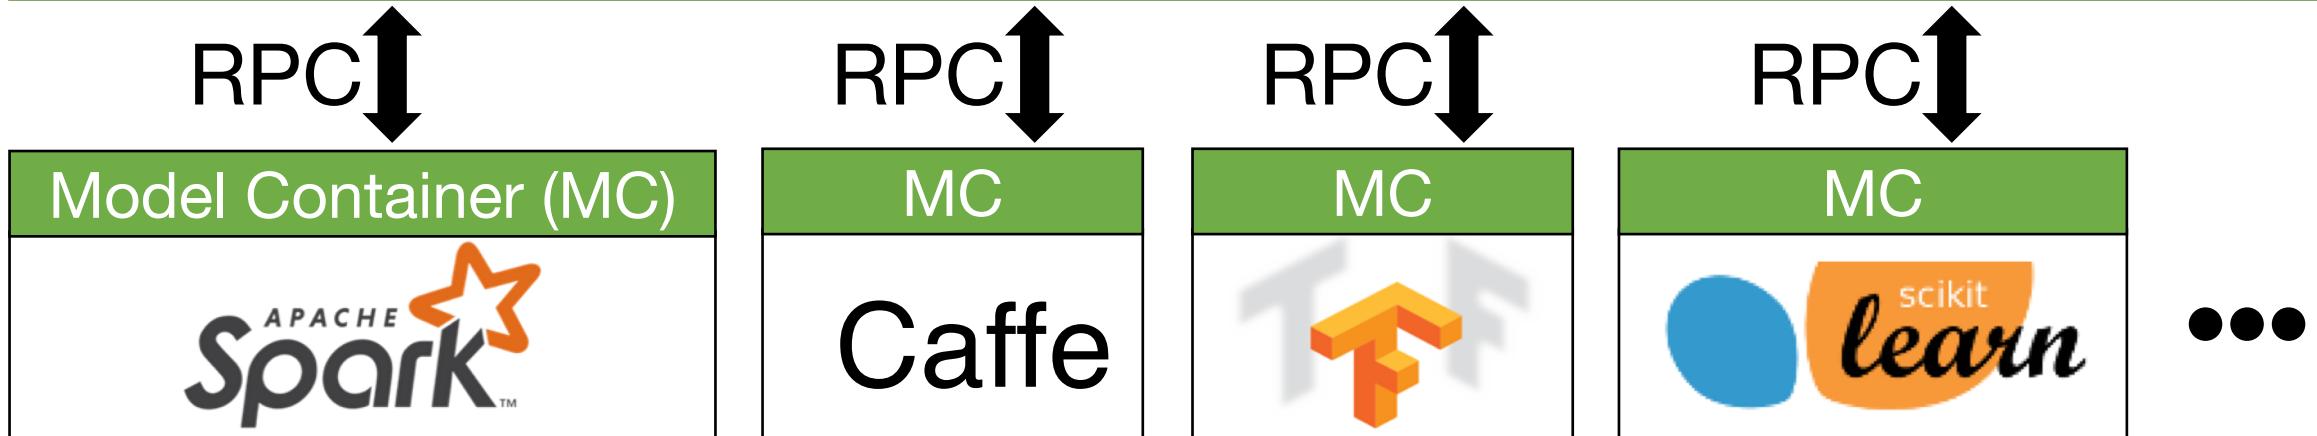
**Cross-framework caching and batching:** optimize throughput and latency

**Cross-framework model composition:** improved accuracy through ensembles and bandits

# Clipper Architecture



# Clipper



Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems

# Container-based Model Deployment

*Implement Model API:*

```
class ModelContainer:  
    def __init__(model_data)  
    def predict_batch(inputs)
```

- API support for many programming languages
  - Python
  - Java
  - C/C++
  - R

# Container-based Model Deployment

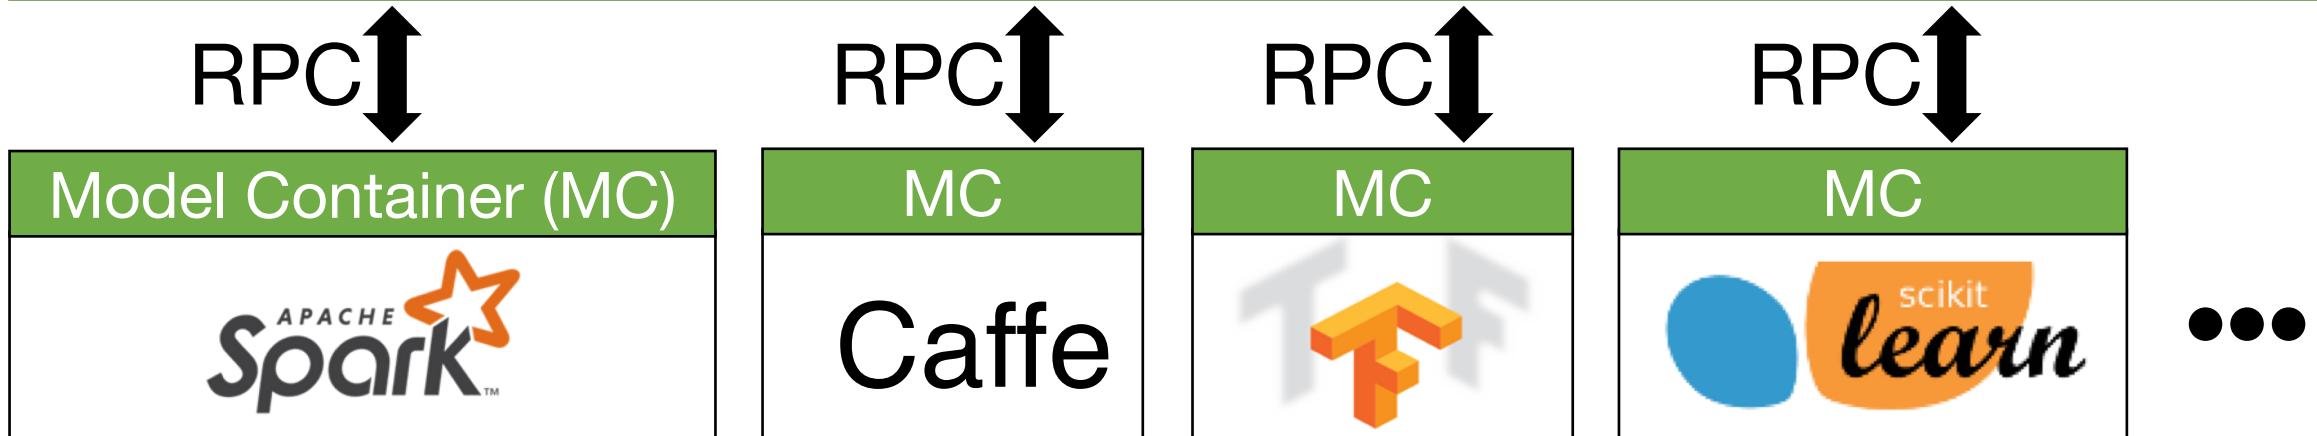
*Package model implementation and dependencies*

## Model Container (MC)

```
class ModelContainer:  
    def __init__(model_data)  
    def predict_batch(inputs)
```



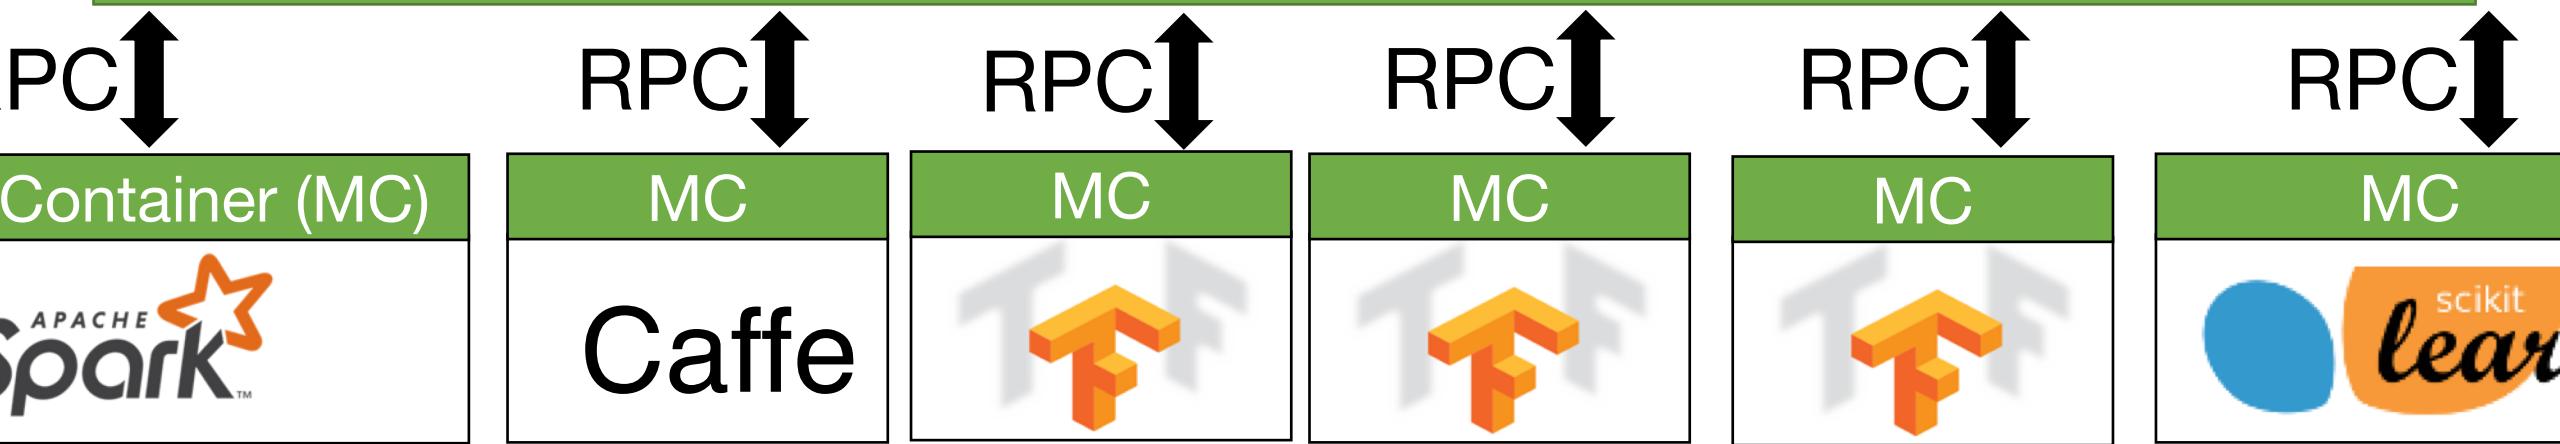
# Clipper



Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes as Docker containers
- Resource isolation: ML frameworks can be buggy

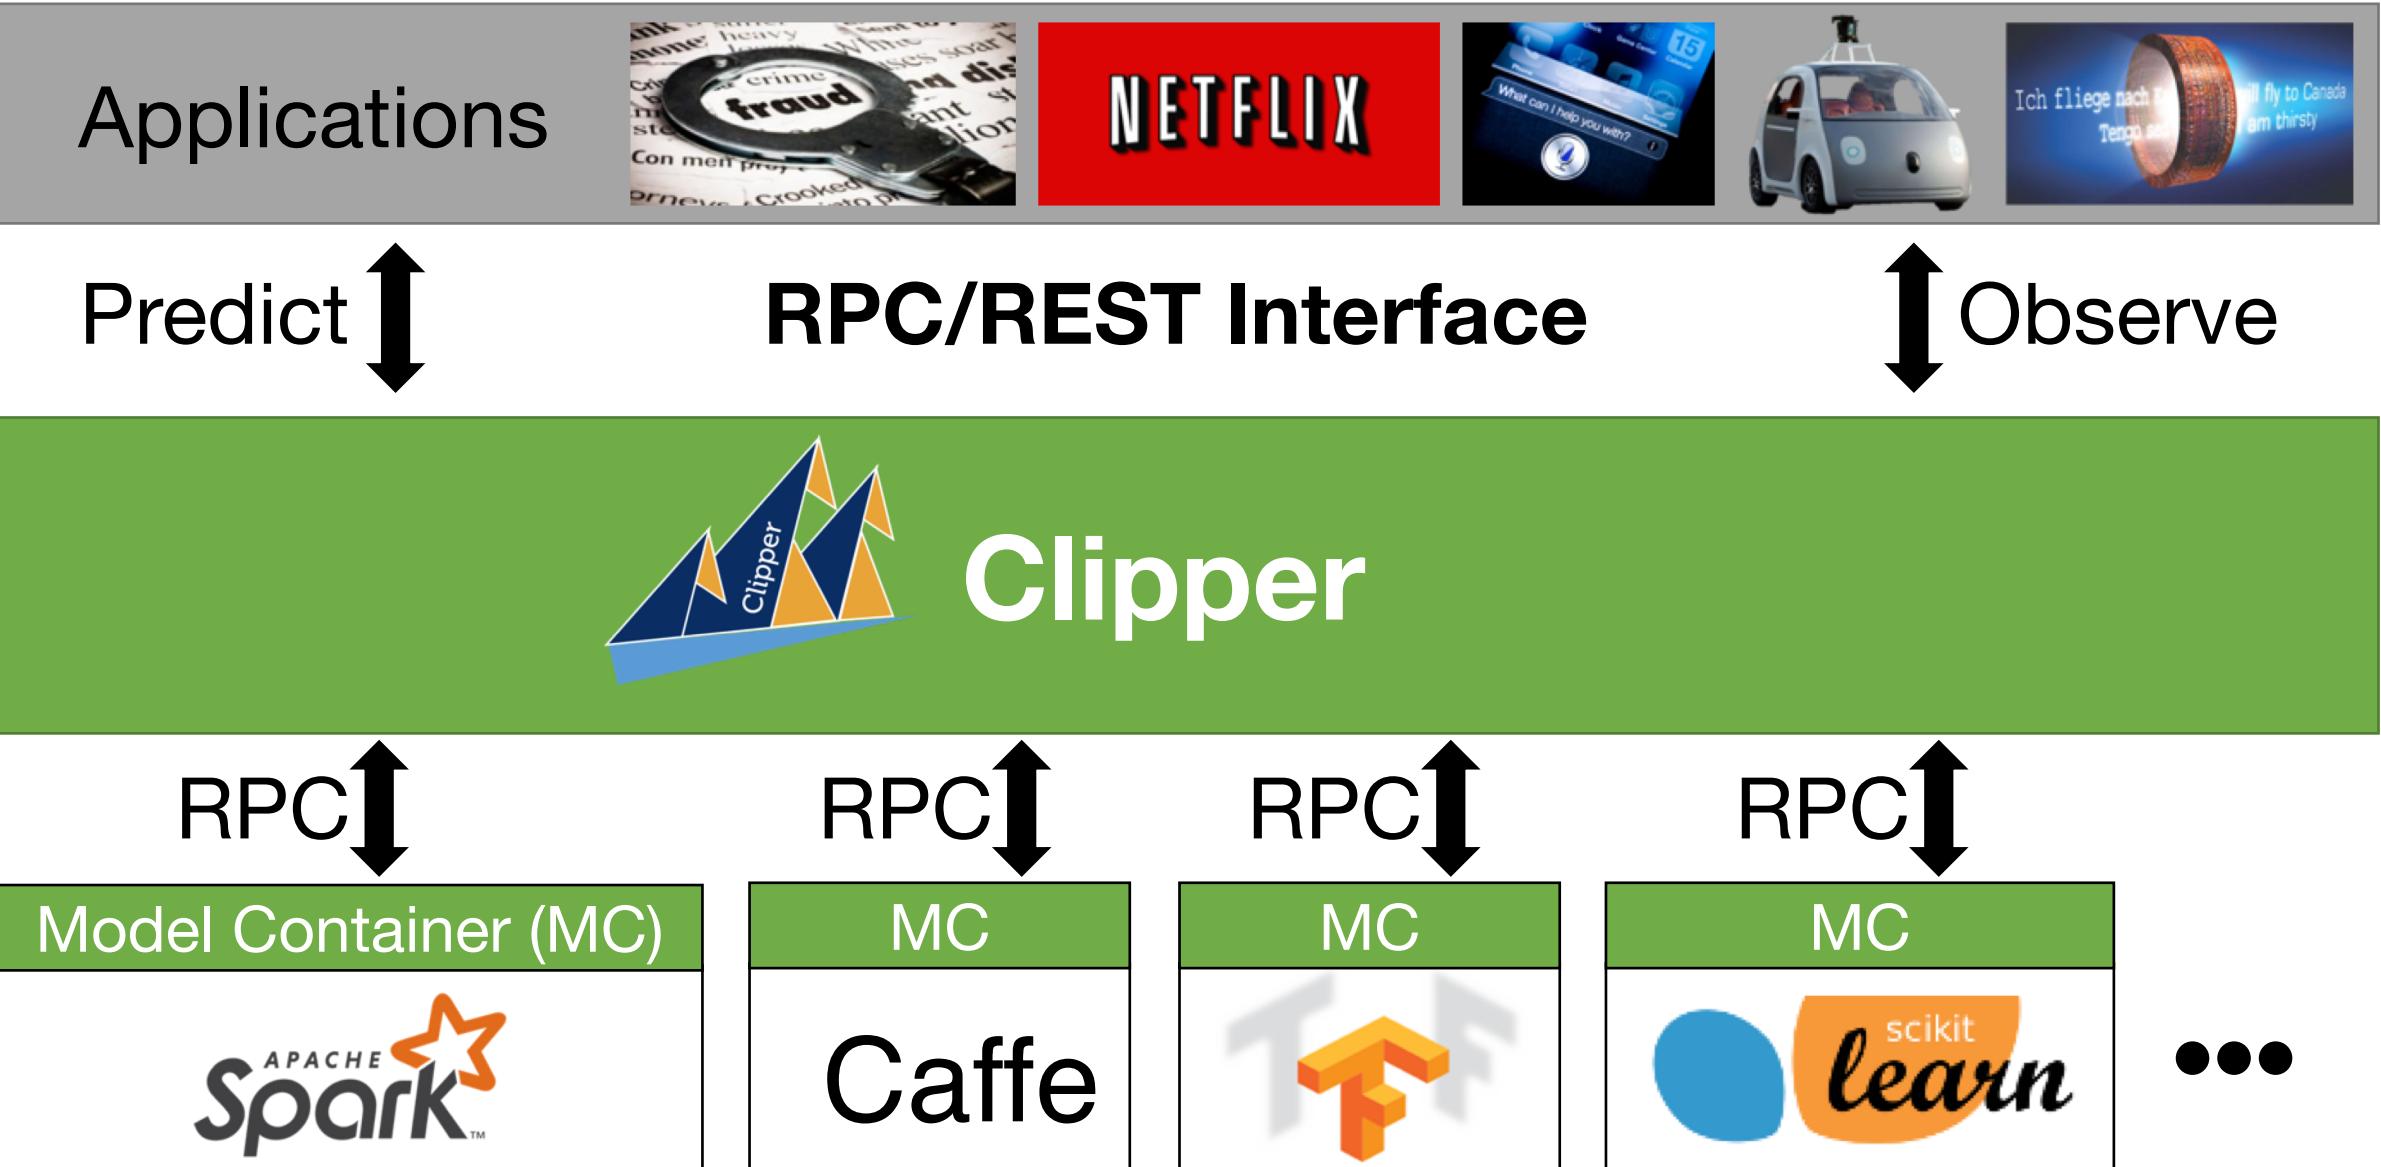
# Clipper



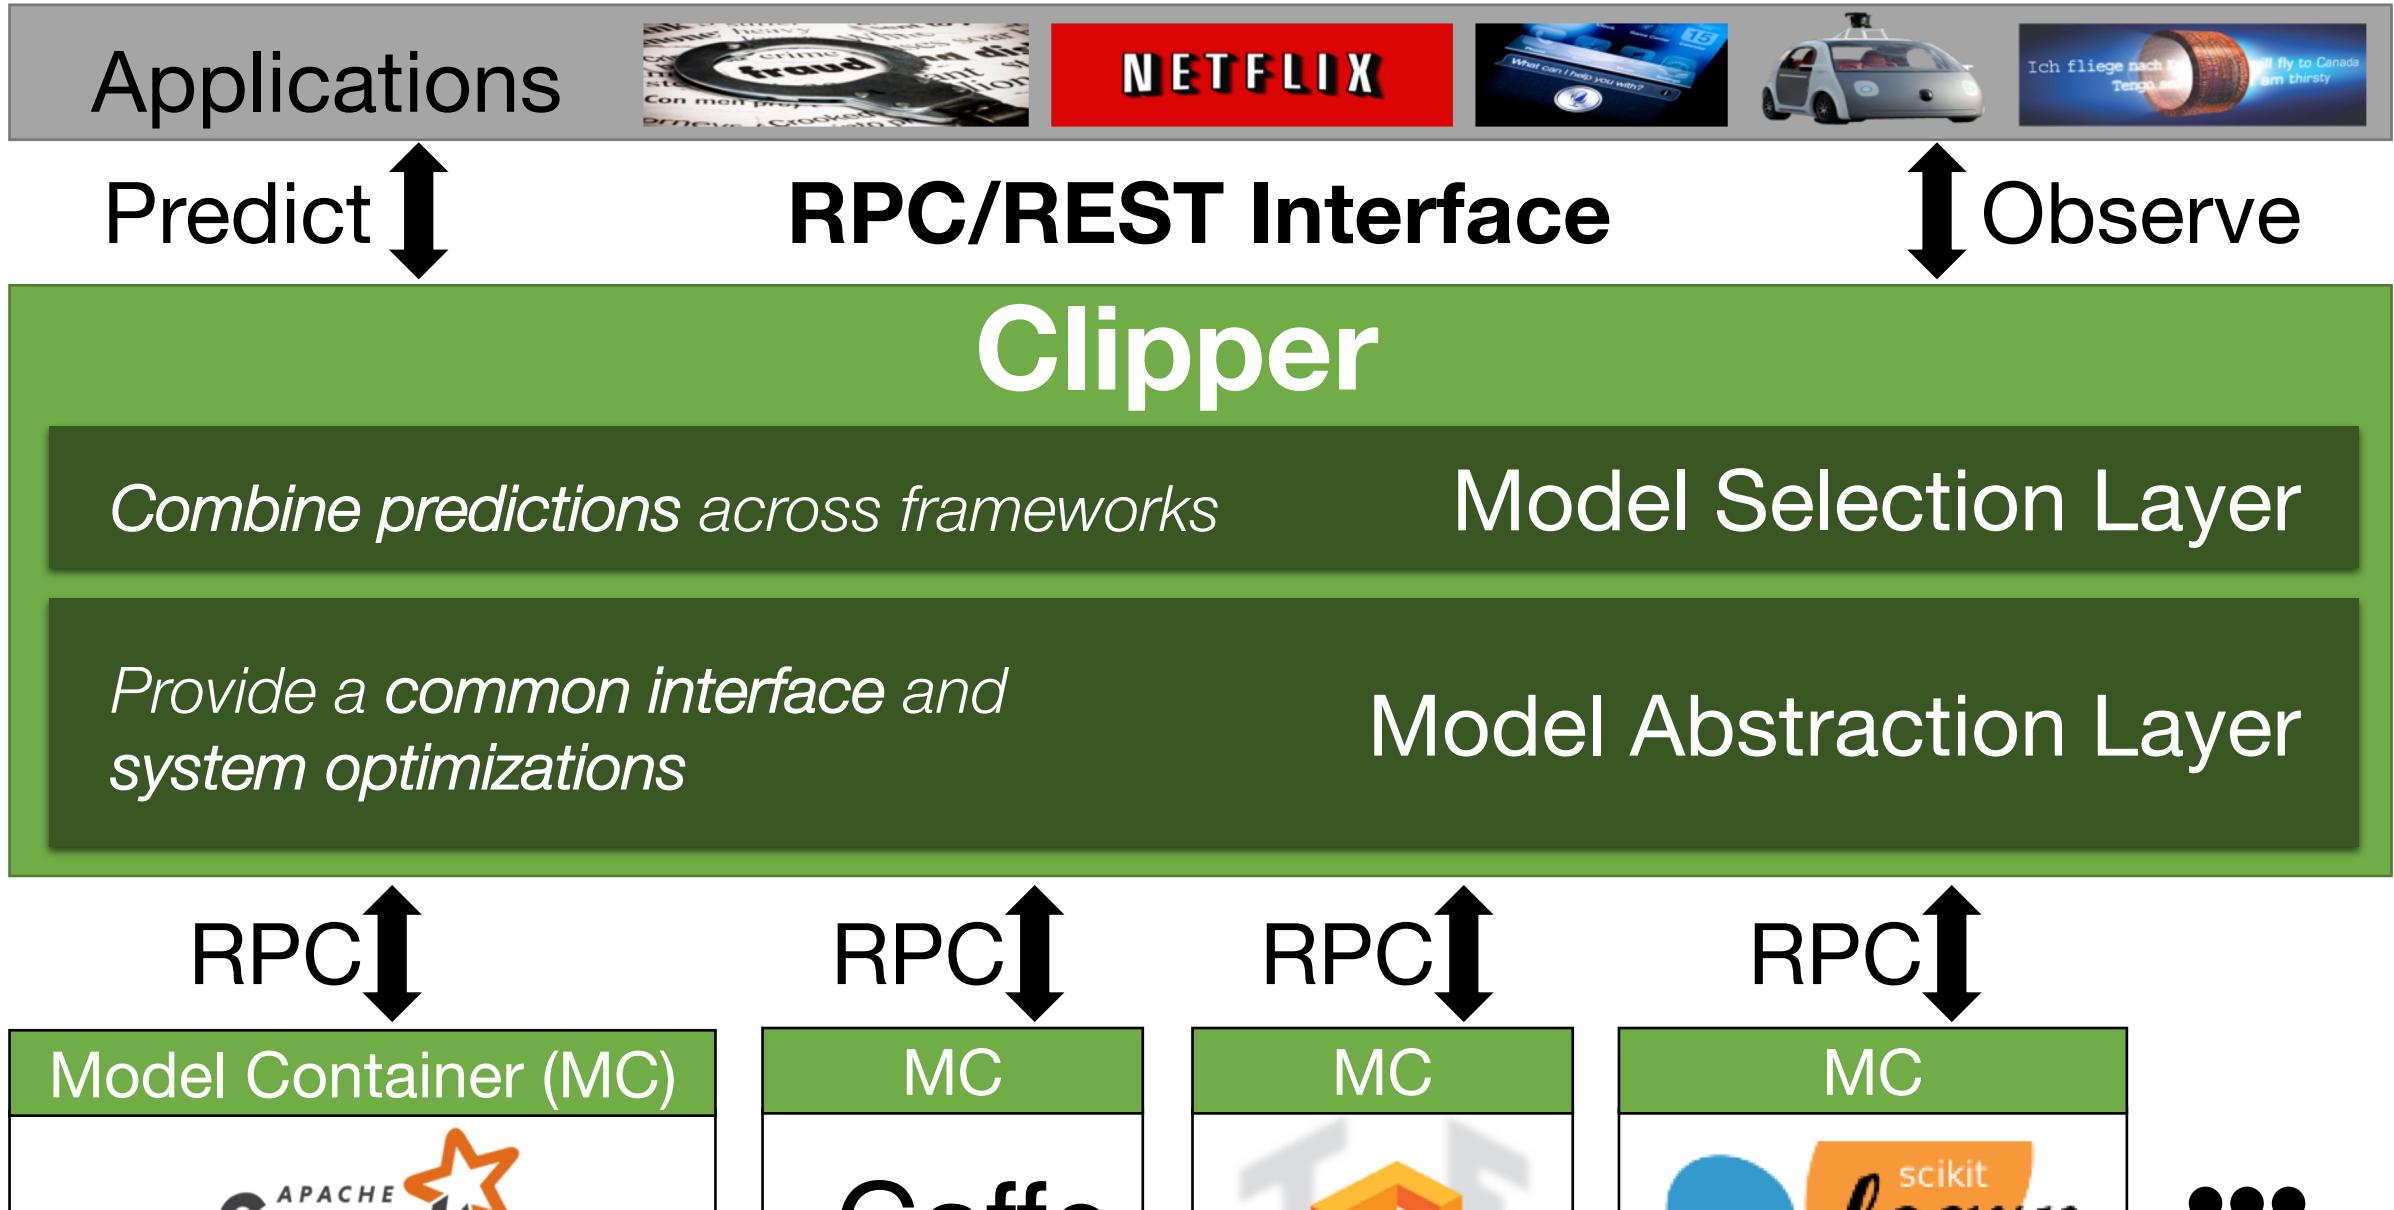
Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes as Docker containers
  - Resource isolation: ML frameworks can be buggy
  - Scale-out at the level of individual models

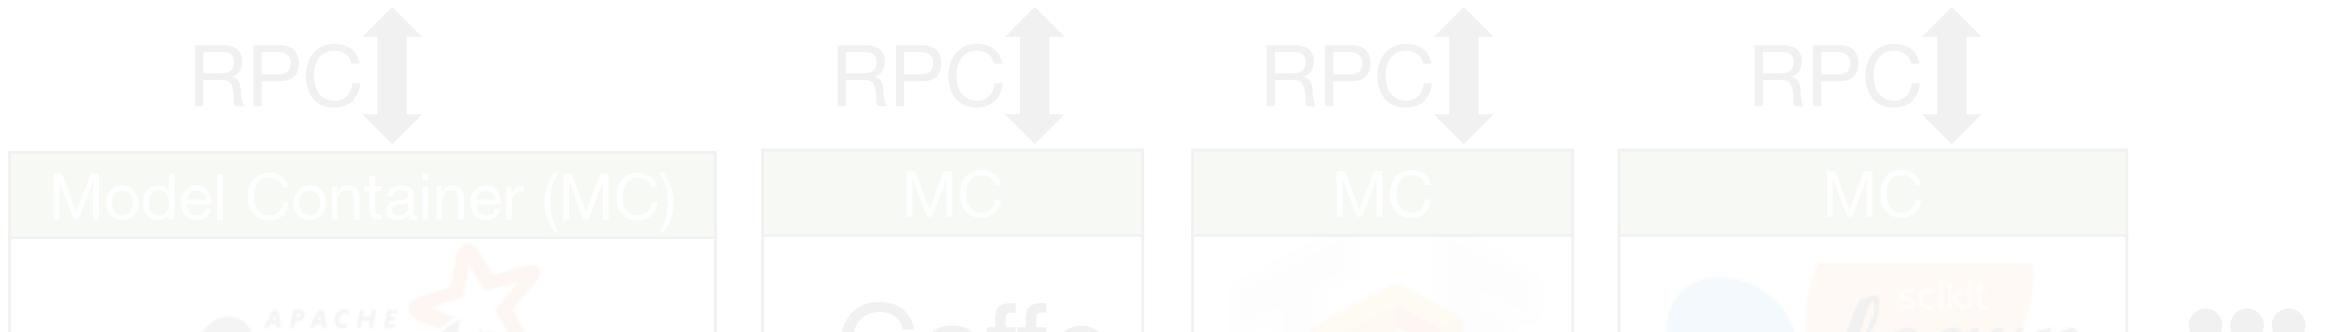
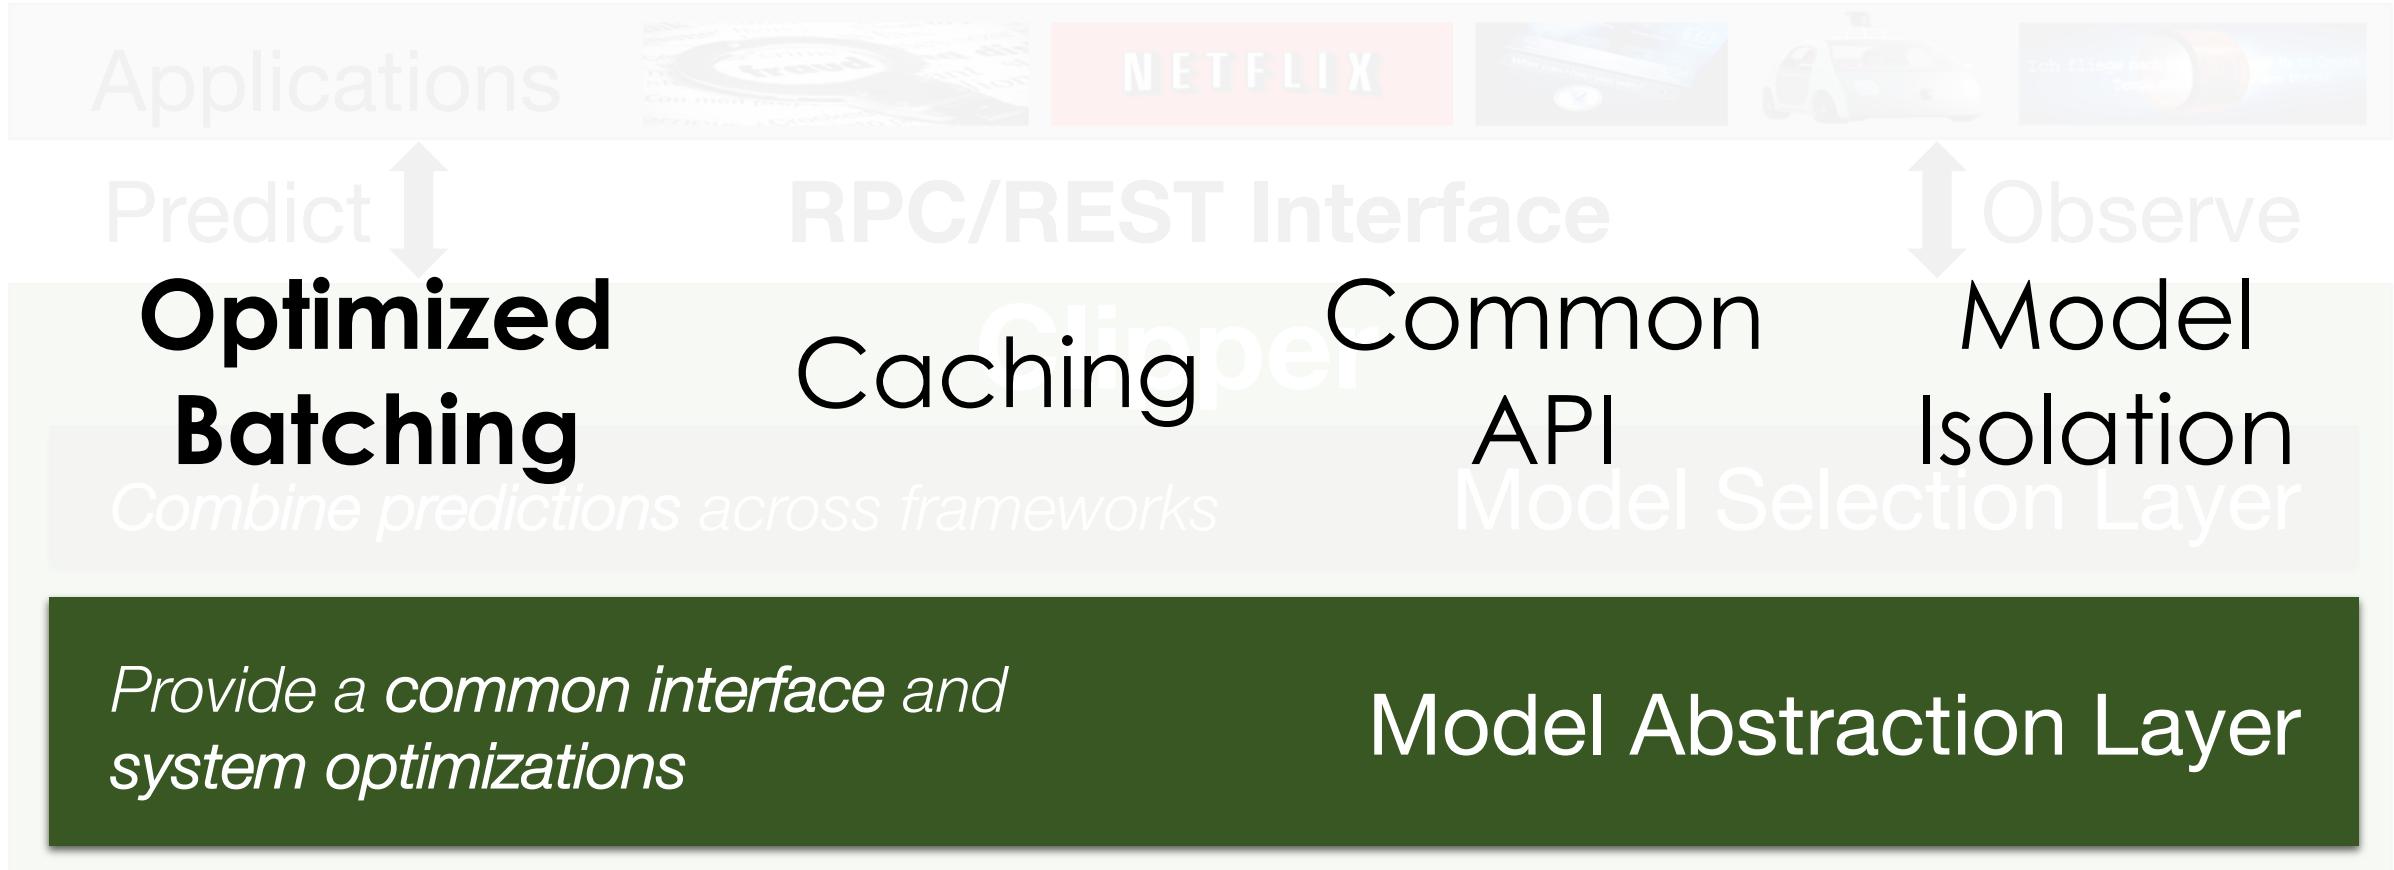
# Clipper Architecture



# Clipper Architecture

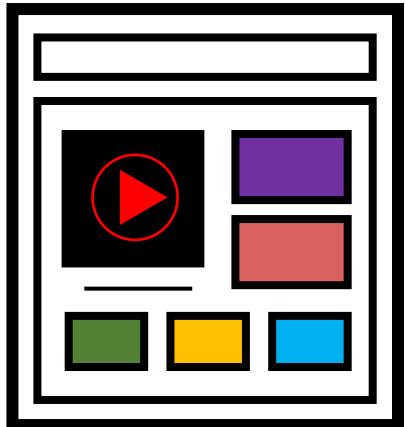


# Clipper Architecture



# Batching to Improve Throughput

- Why batching helps:

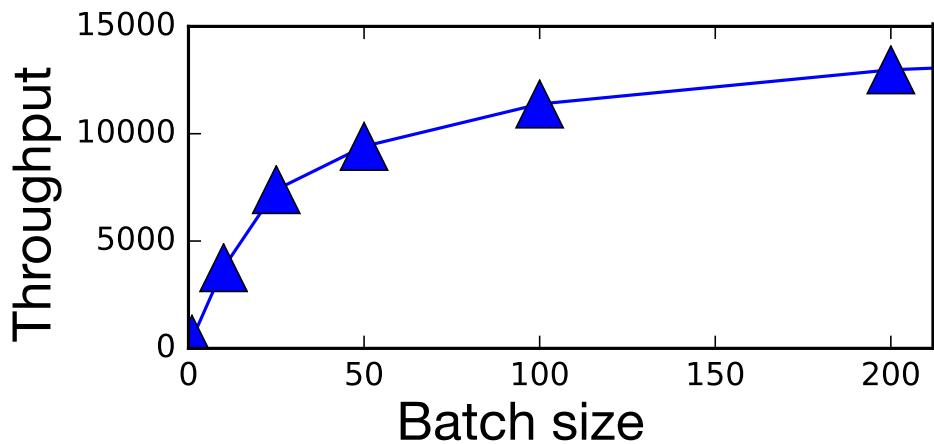


A single page load may generate many queries

- Optimal batch depends on:

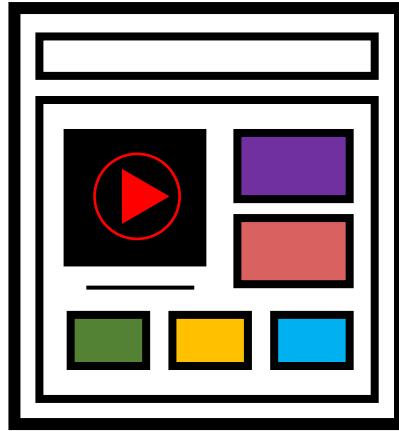
- hardware configuration
- model and framework
- system load

## Throughput-optimized frameworks



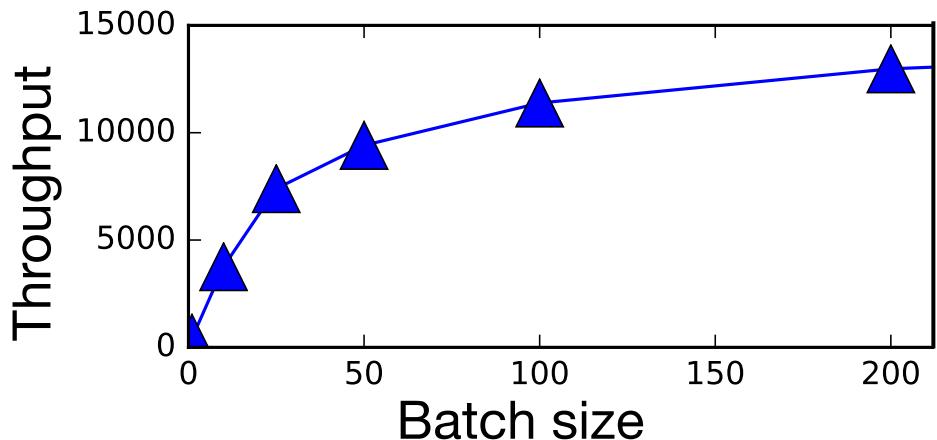
# Latency-aware Batching to Improve Throughput

- Why batching helps:



A single page load may generate many queries

Throughput-optimized frameworks



- Optimal batch depends on:

- hardware configuration
- model and framework
- system load

Clipper Solution:

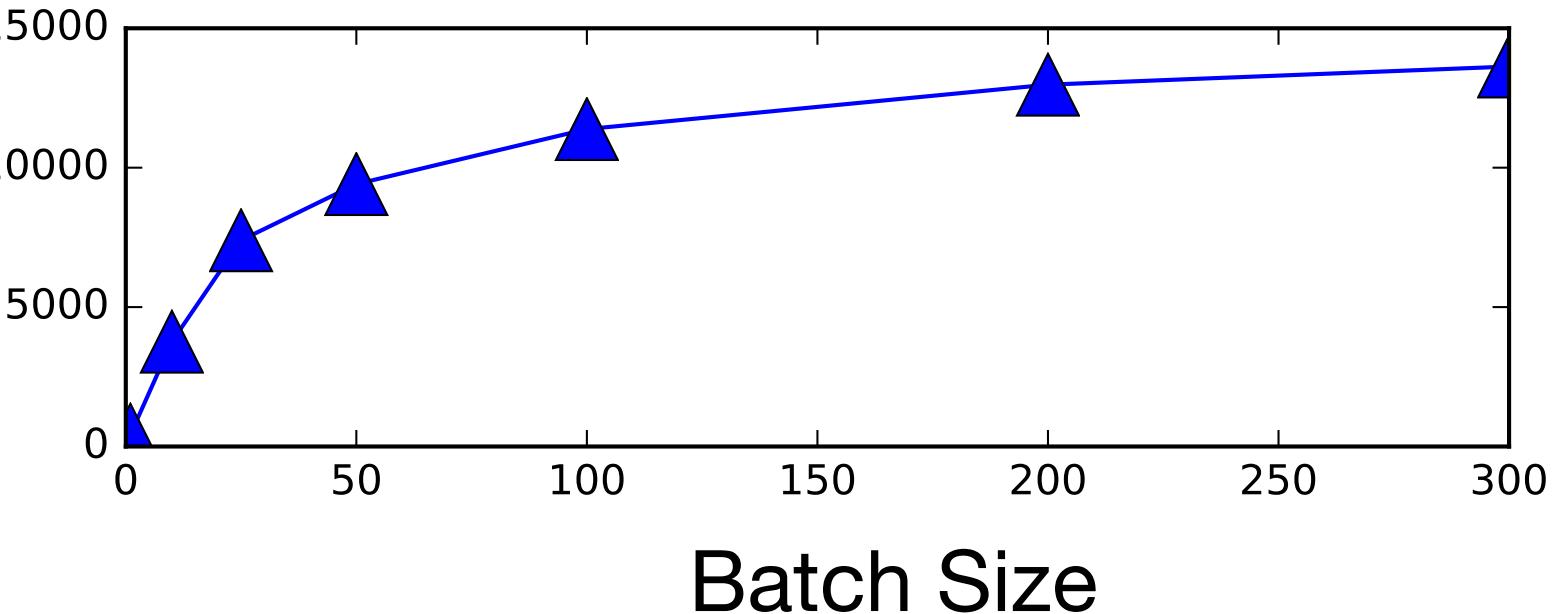
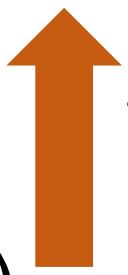
*Adaptively tradeoff latency and throughput...*

- Inc. batch size *until the latency objective is exceeded (Additive Increase)*
- If latency exceeds SLO cut batch size by a fraction (**Multiplicative Decrease**)

# Tensor Flow Conv. Net (GPU)

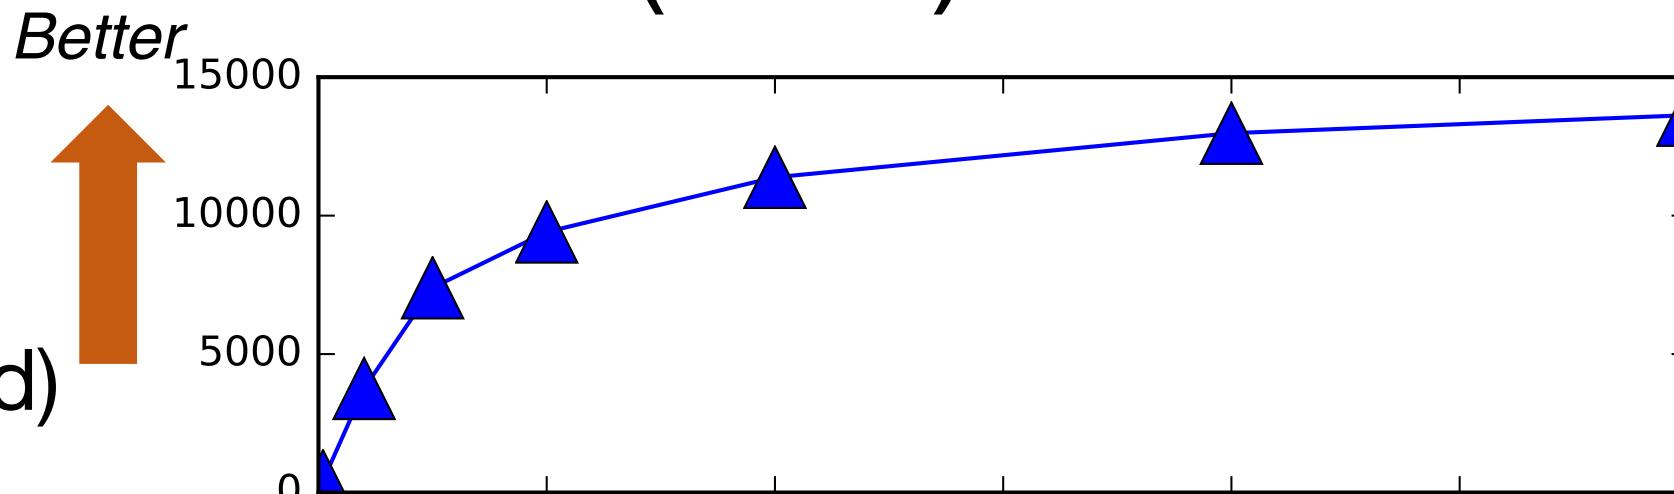
Throughput  
(Queries Per Second)

*Better*

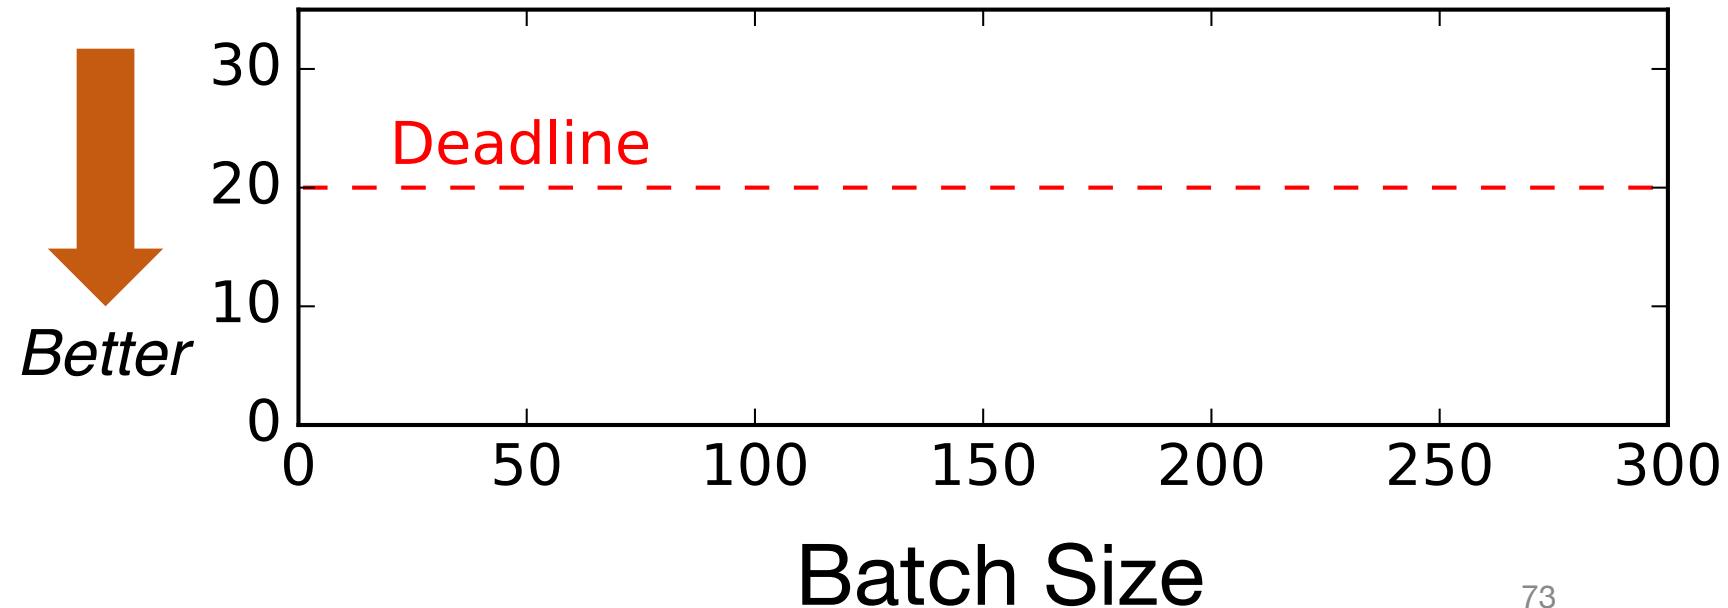


# Tensor Flow Conv. Net (GPU)

Throughput  
(Queries Per Second)



Latency (ms)



# Tensor Flow Conv. Net (GPU)

Throughput  
(Queries Per Second)

Better



15000

10000

5000

0

Optimal Batch Size

Latency (ms)

Better



30

20

10

0

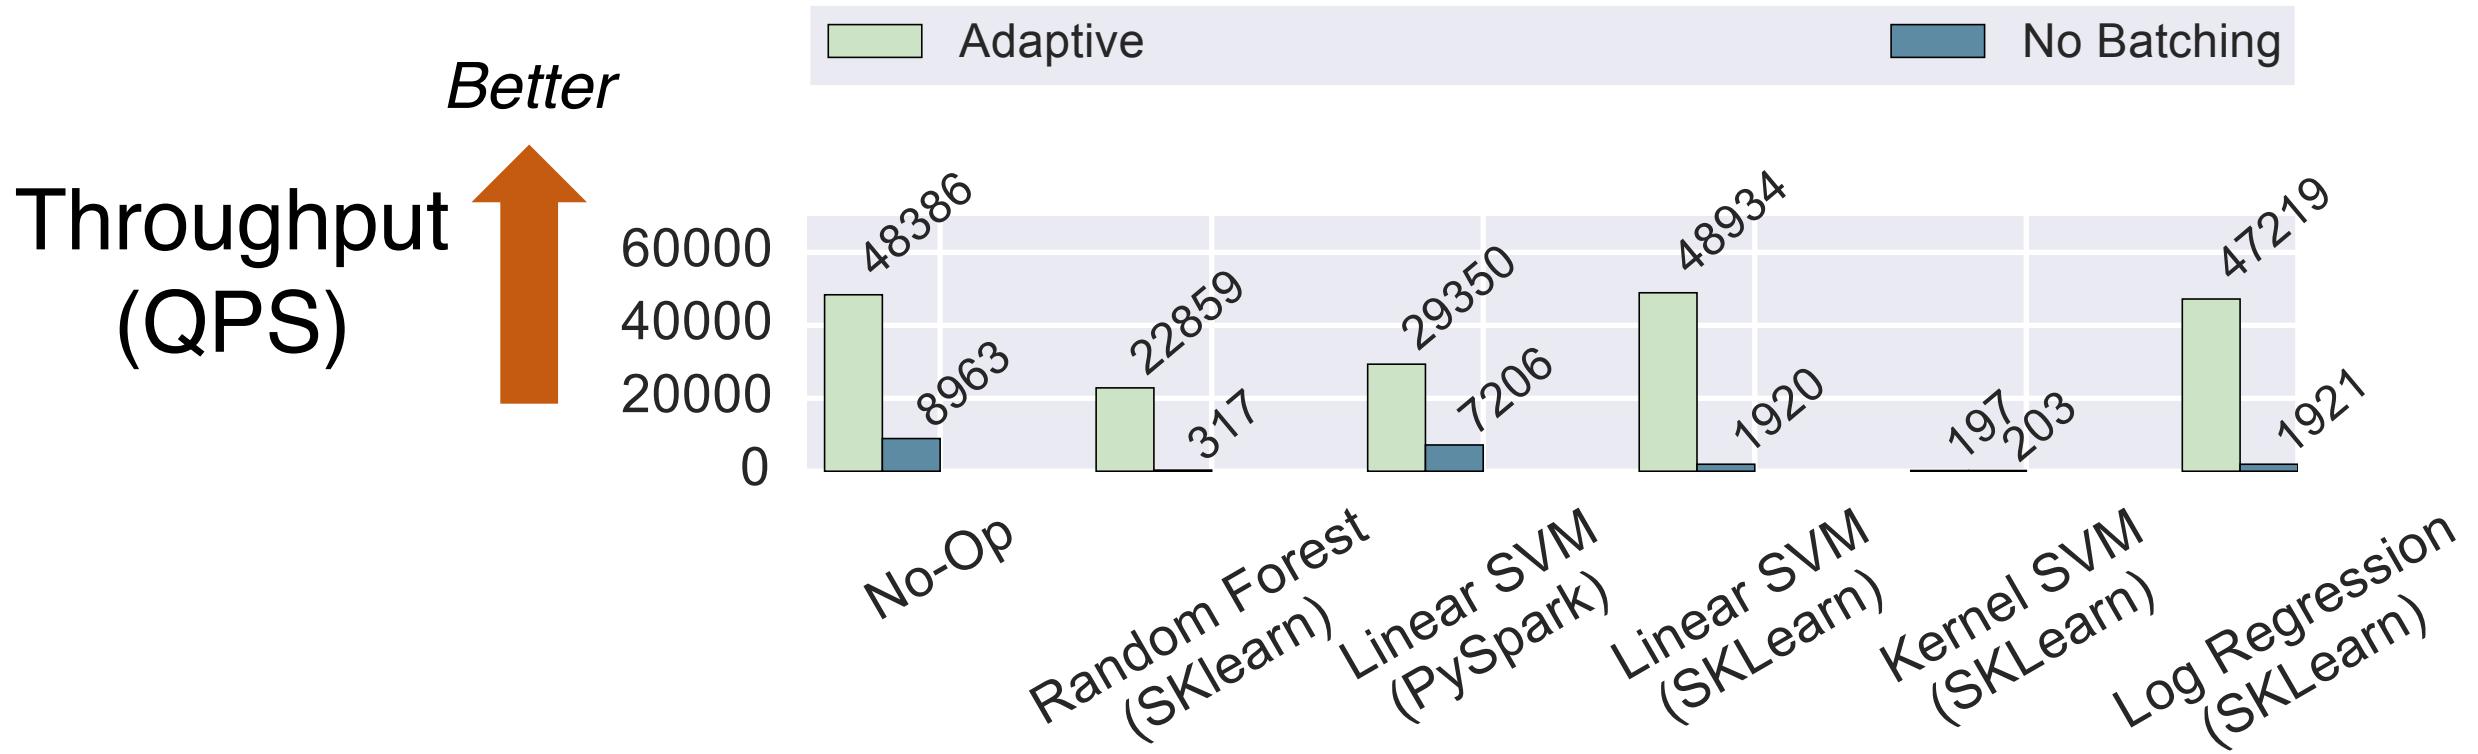
Deadline

P99

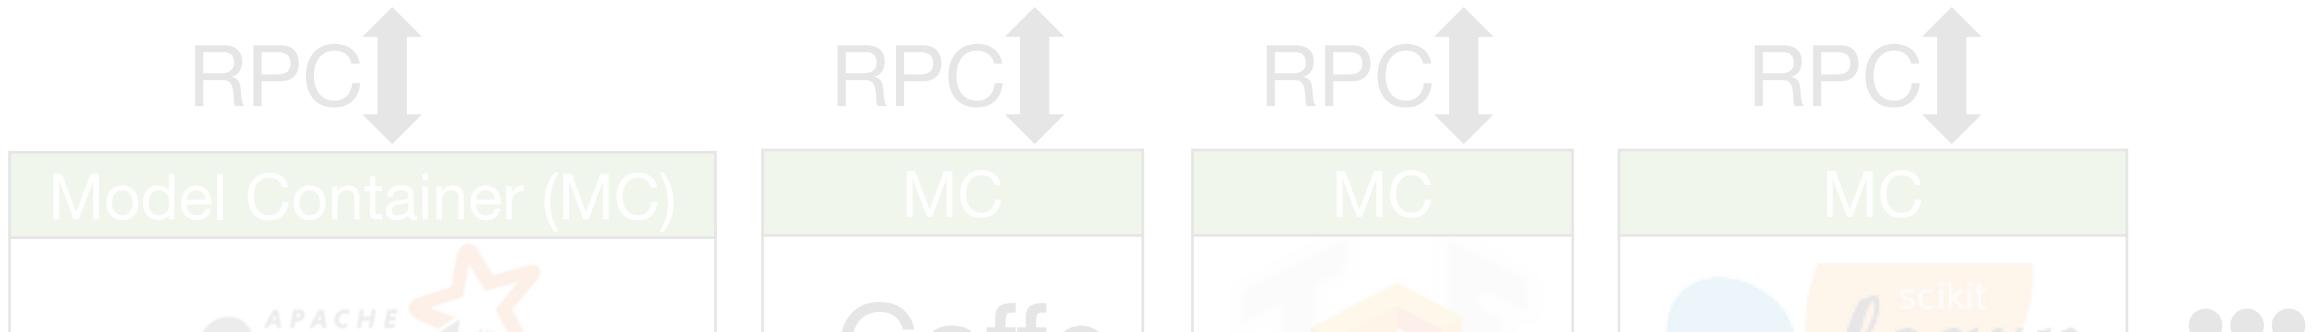
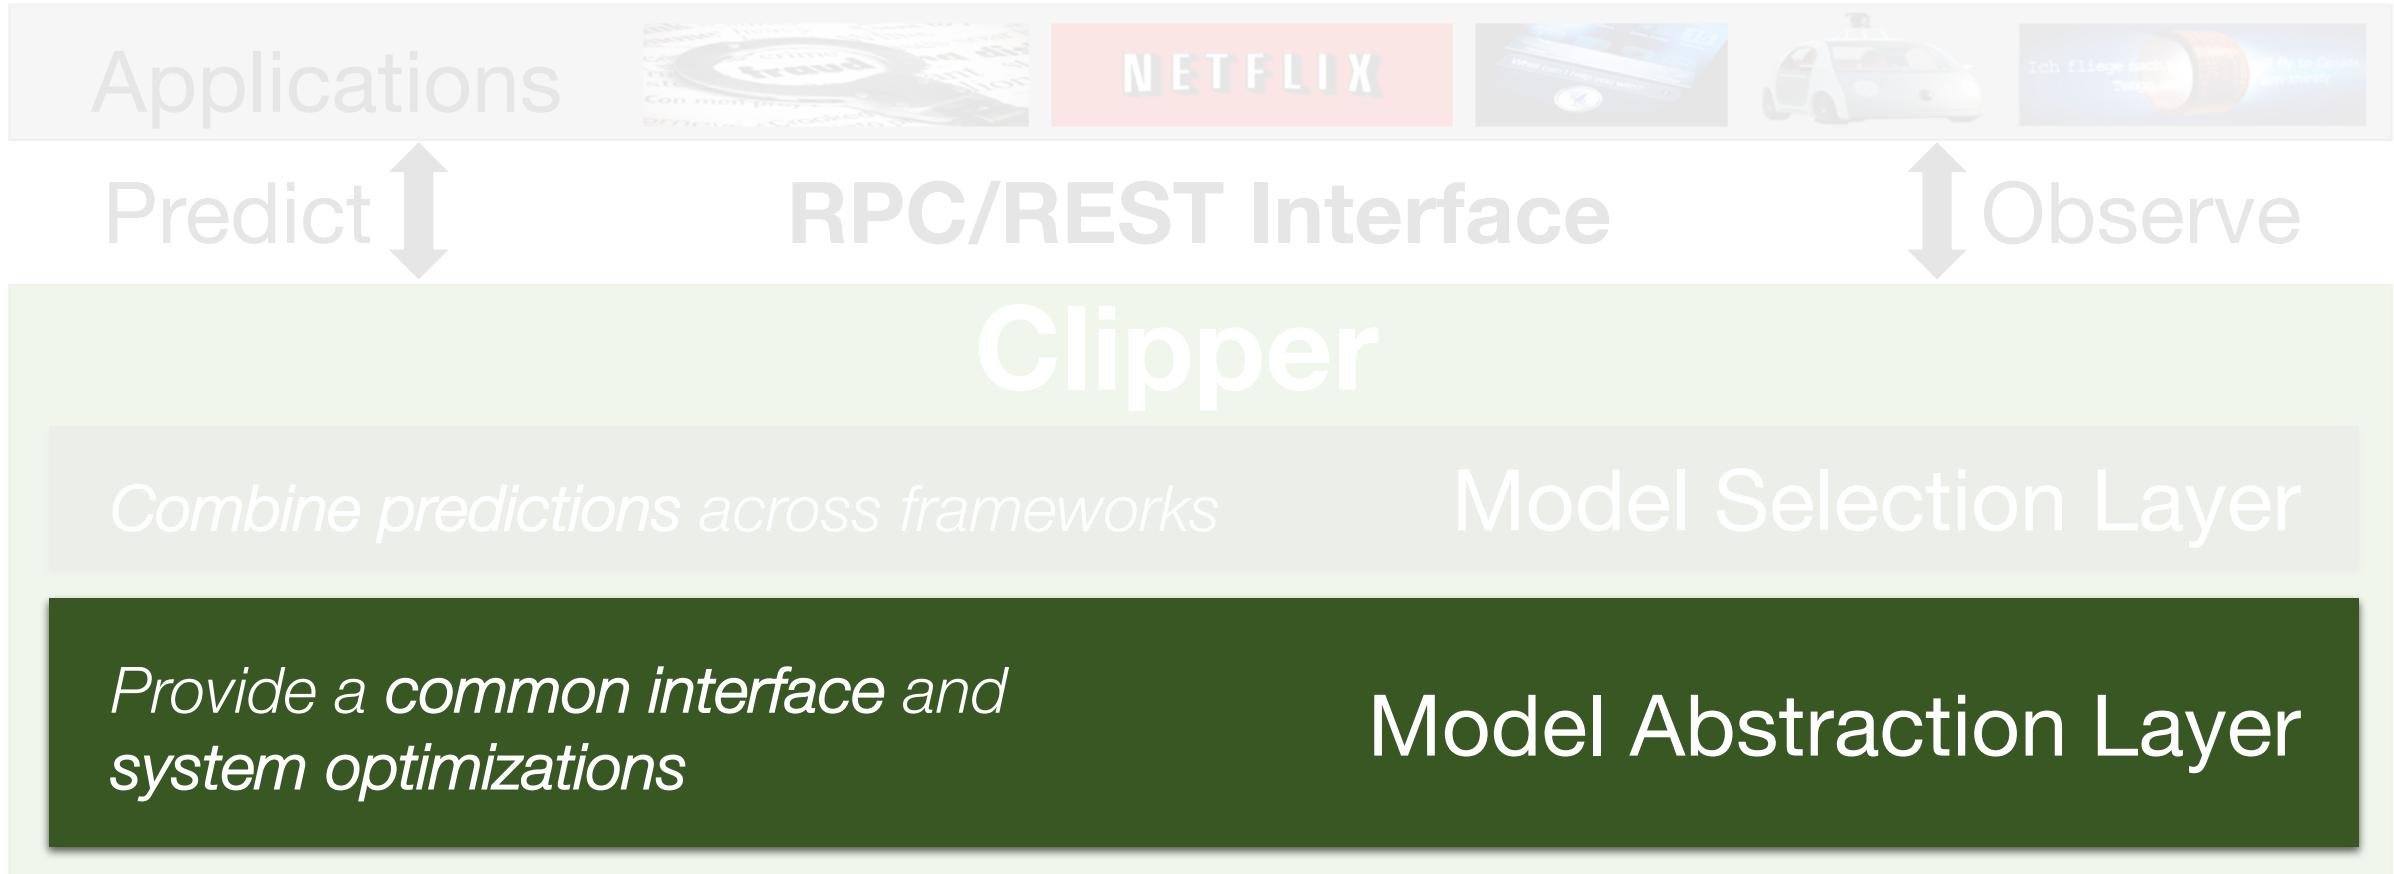
Batch Size

200

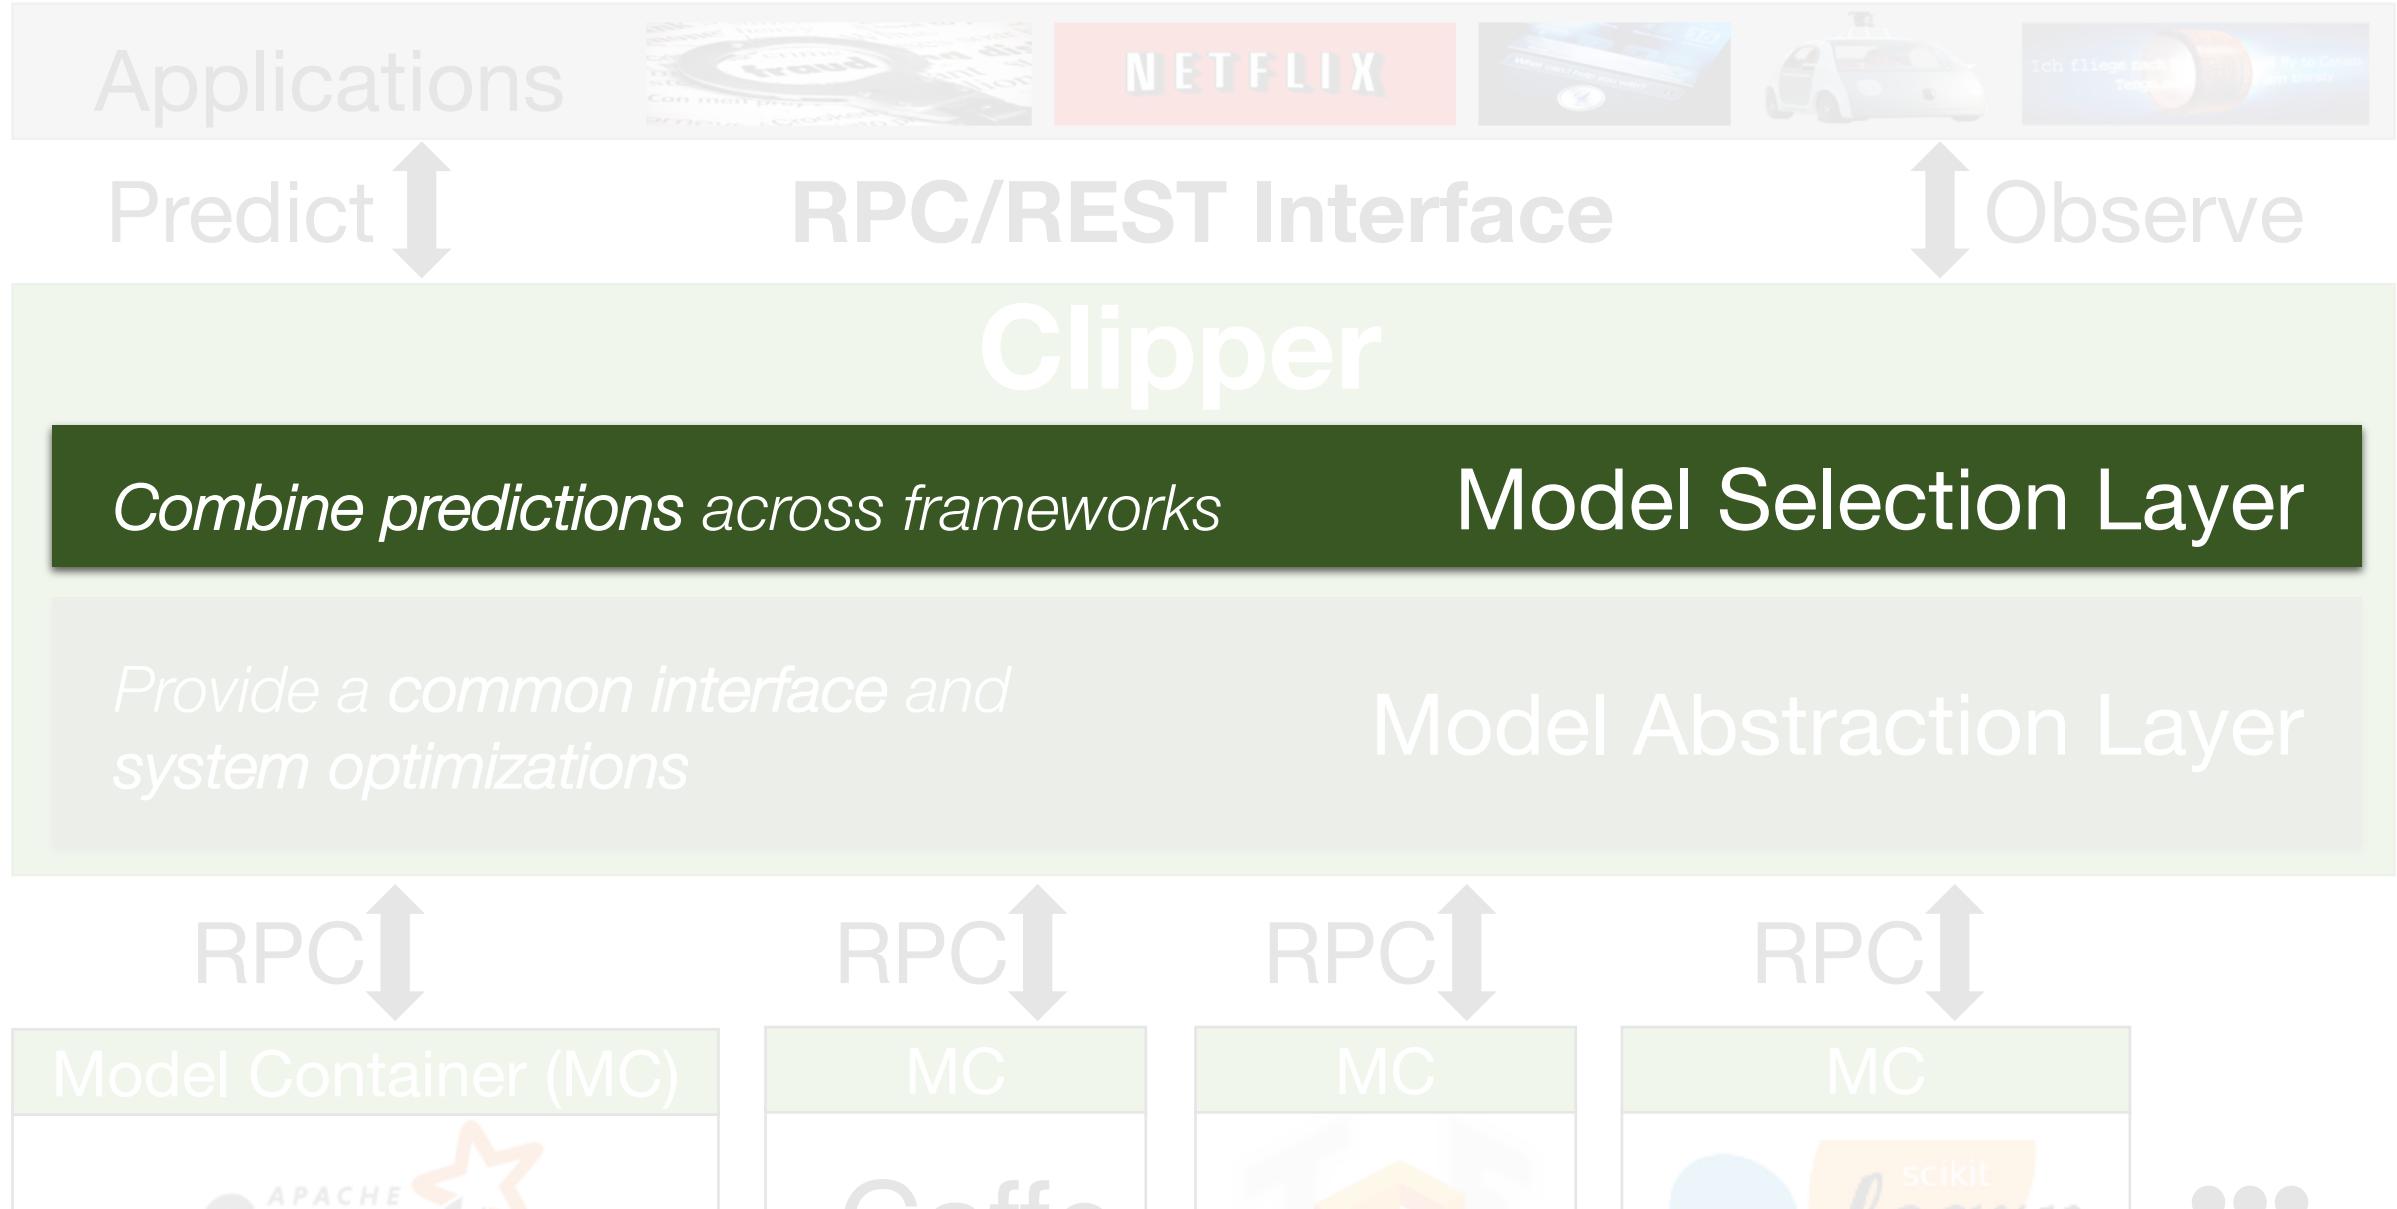
# *Latency-aware Batching to Improve Throughput*



# Clipper Architecture



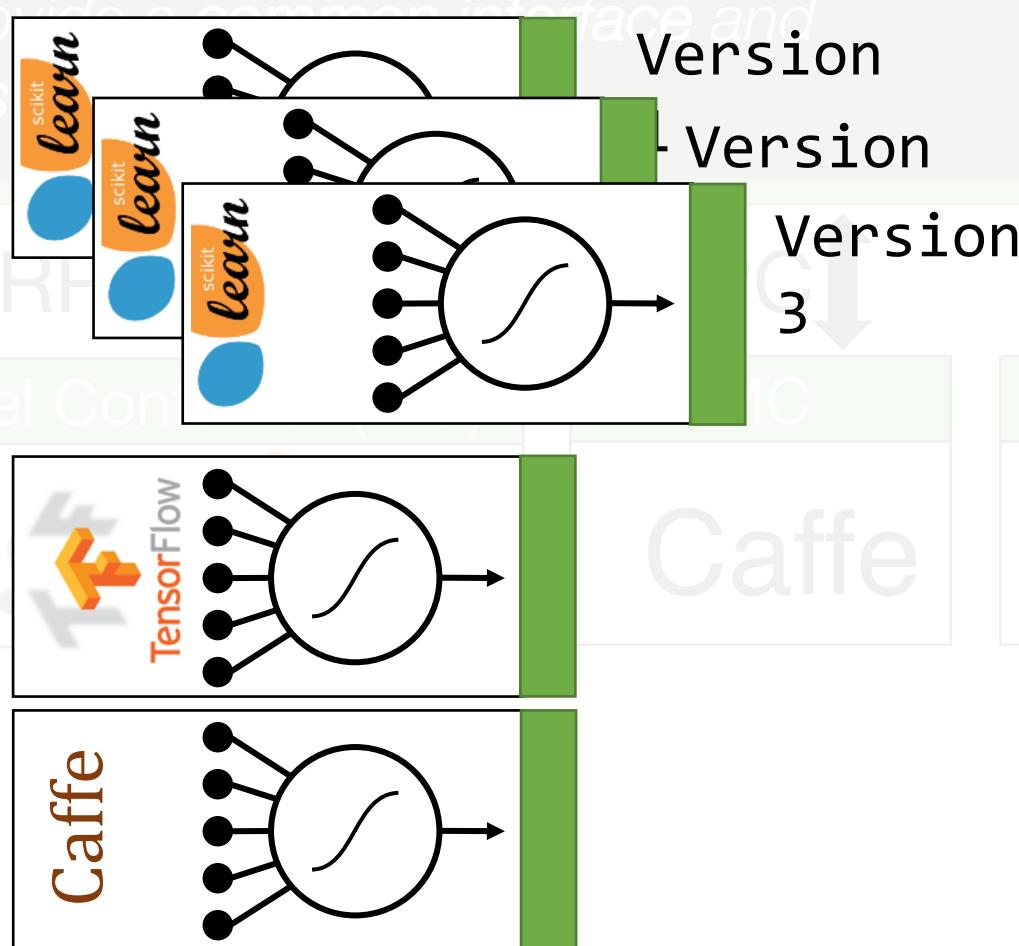
# Clipper Architecture



# Clipper

*Combine predictions across frameworks*

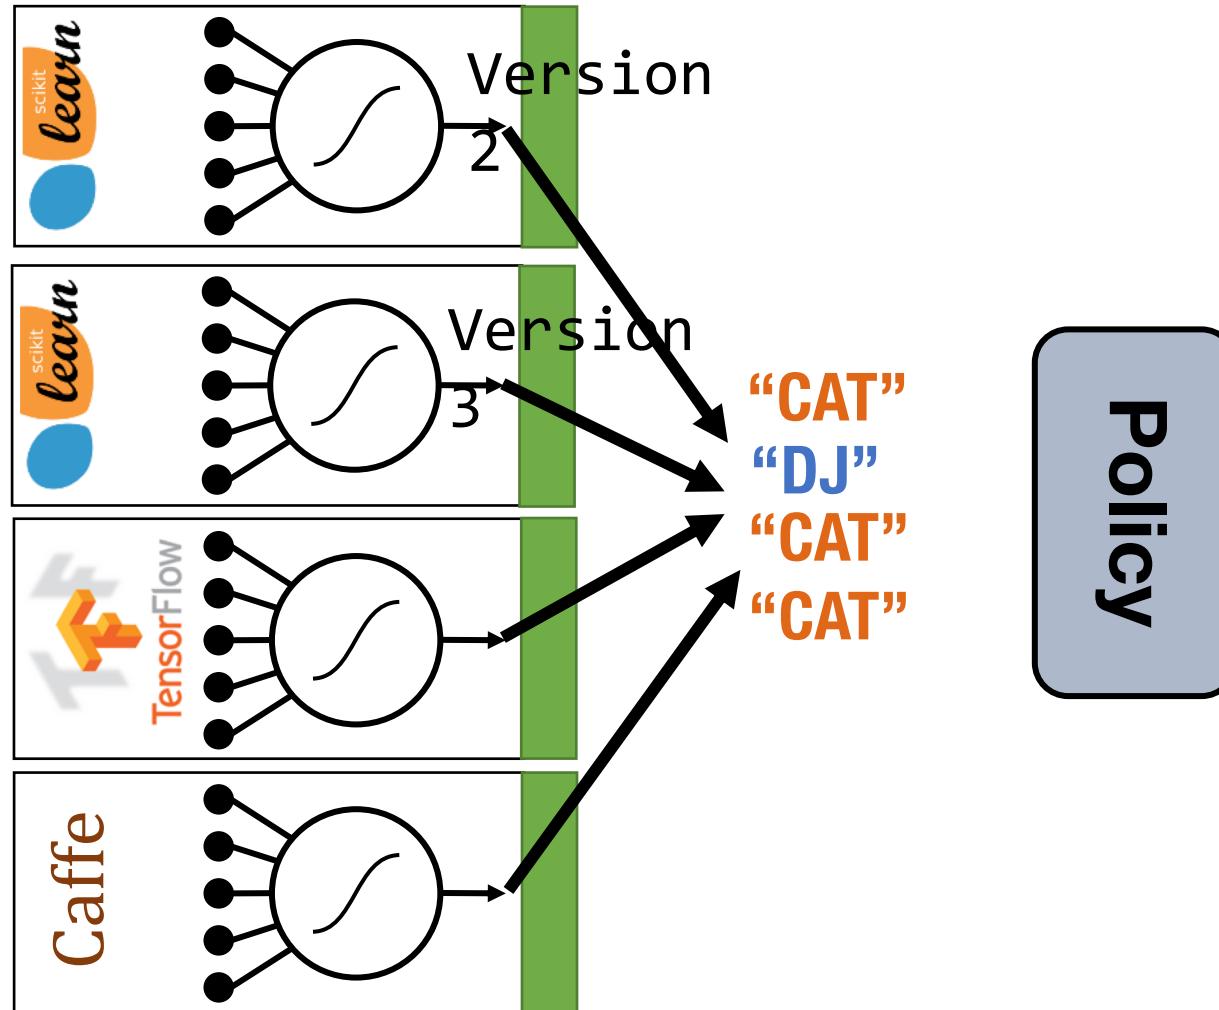
Model Selection Layer



Model Abstraction Layer  
*Periodic retraining*

*Experiment with new  
models and frameworks*

# Selection Policy can Calibrate Confidence

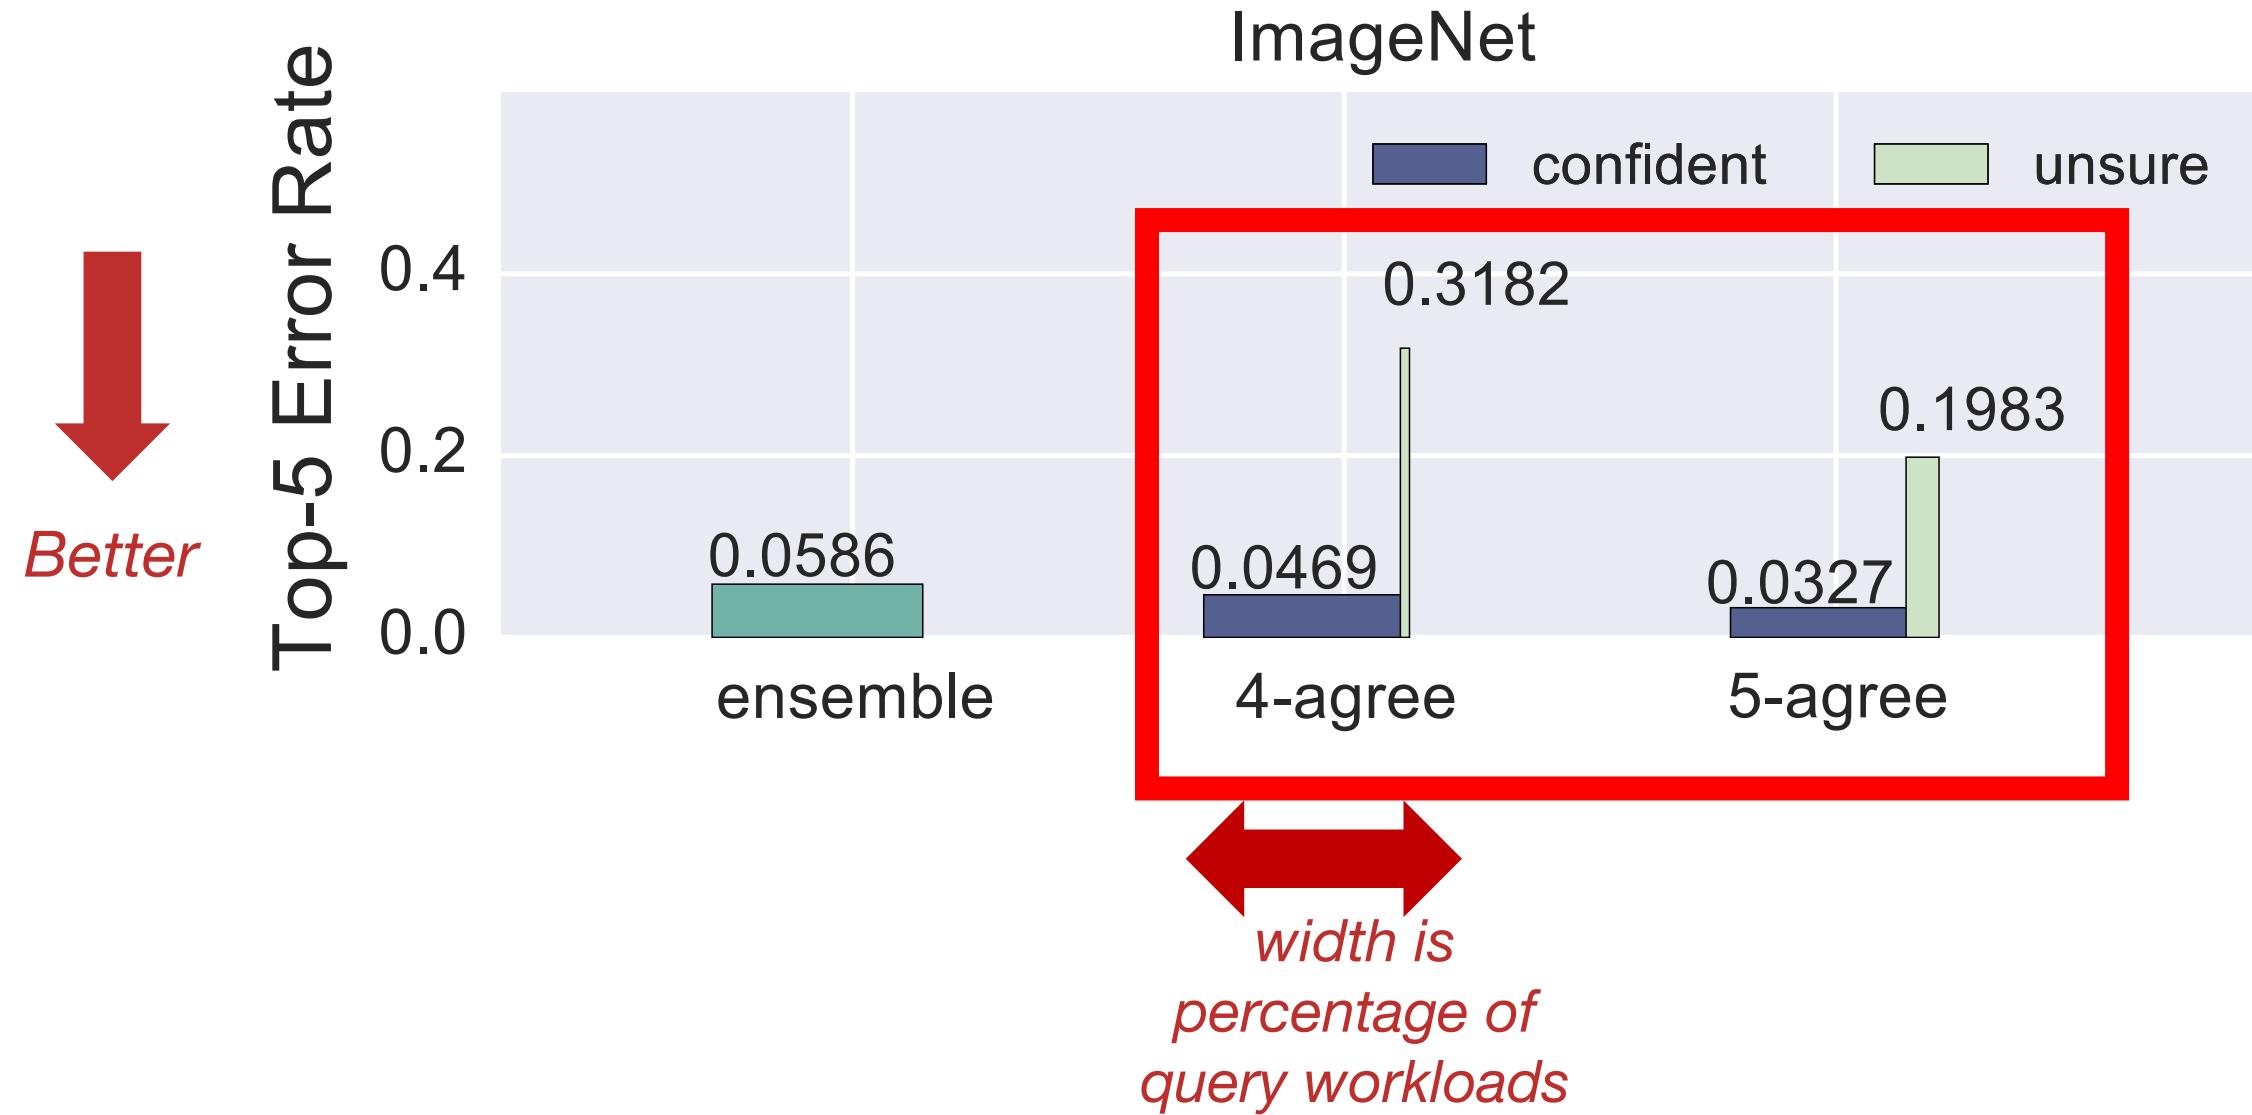


**"CAT"**  
**UNSURE**

# Selection Policy: Estimate confidence



# Selection Policy: Estimate confidence



# Project Status and Development

- Current development focus:
  - **stability** and **performance** improvements
  - **easy model deployment** for common ML frameworks: Pytorch, caffe2 (via onnx), tensorflow, xgboost, mxnet, pyspark, scikit learn
  - **metrics** and **monitoring** infrastructure using Prometheus
- **Development Team:** 22 active contributors
  - Including 8 from outside Berkeley
- Working with several organizations on **production deployments**
  - SAP, Scotia Bank, Pacific AI, ARM...

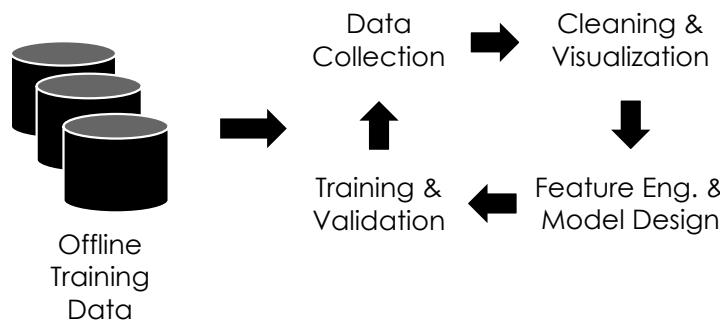
I made a case for

## **Model & Data Engineering**

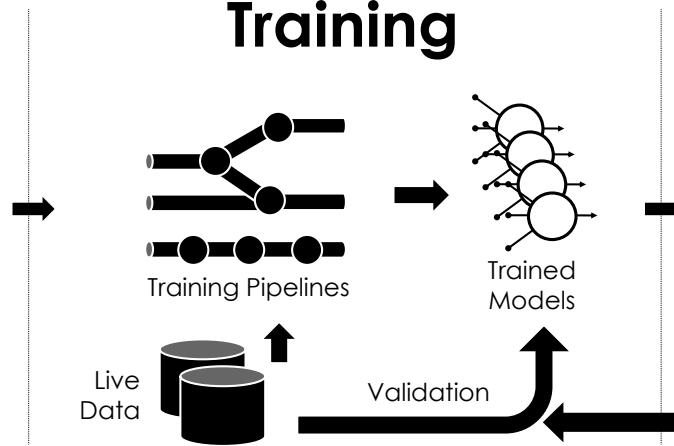
and outlined some of the  
**opportunities & challenges**

# Machine Learning Lifecycle

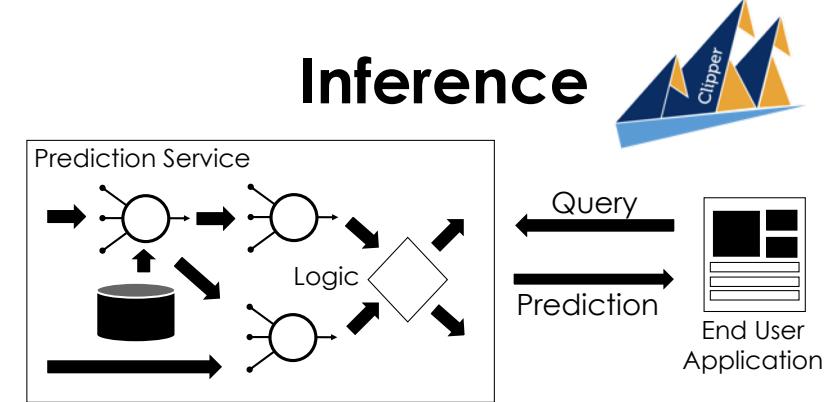
## Model Development



## Training



## Inference





Middle layer for prediction serving.

Common  
Abstraction

System  
Optimizations

theano  
APACHE  
Spark™

scikit  
learn

Caffe

TensorFlow

PYTORCH  
dmrc  
mxnet

VW  
KALDI



<http://clipper.ai>

Open-source prediction serving system for  
low-latency, high-throughput predictions across  
machine learning models and frameworks.

# Thank you!

[jegonzal@berkeley.edu](mailto:jegonzal@berkeley.edu)

## Collaborators



Daniel  
Crankshaw



Rolando  
Garcia



Joe  
Hellerstein



Yika  
Luo



Simon  
Mo



Vikram  
Sreekanti



Ion  
Stoica



Alexey  
Tumanov



Xin  
Wang



Neeraja  
Yadwadkar



Corey  
Zumar

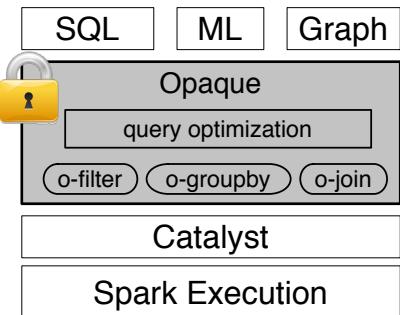
## Research Sponsors



# Bonus!

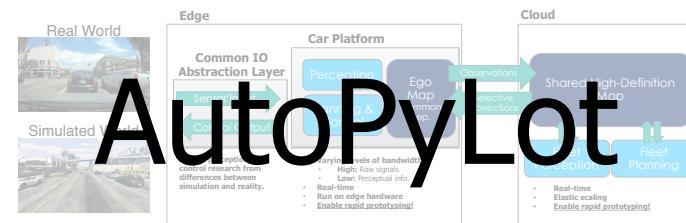
# A few of the RISE Lab projects ...

## Real-time Inference

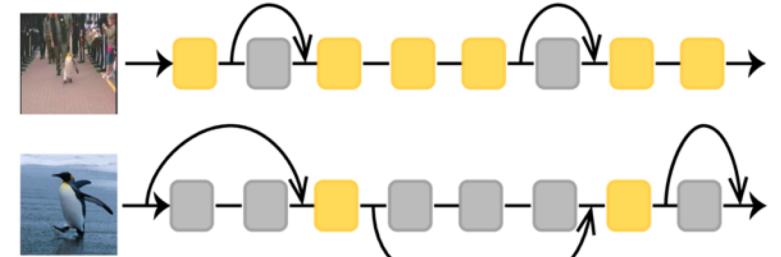


## Hardware Security for Apache Spark

**IDK** Prediction Cascades  
Teaching AI to think fast  
by  
**learning not to overthinking**



An open platform for  
**Autonomous Vehicles**



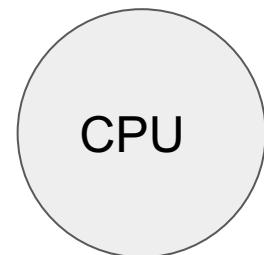
**SkipNets:** RL for Dynamic Network Design

**RAY**

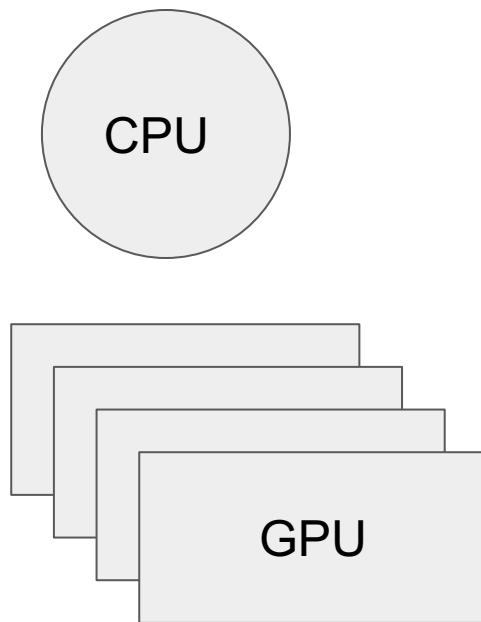
**Parallel Python** for  
Reinforcement Learning

# Ray Tune

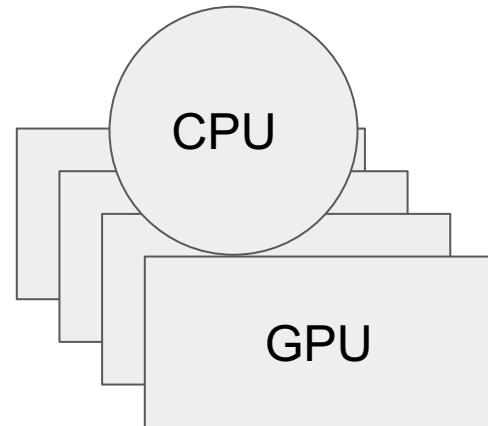
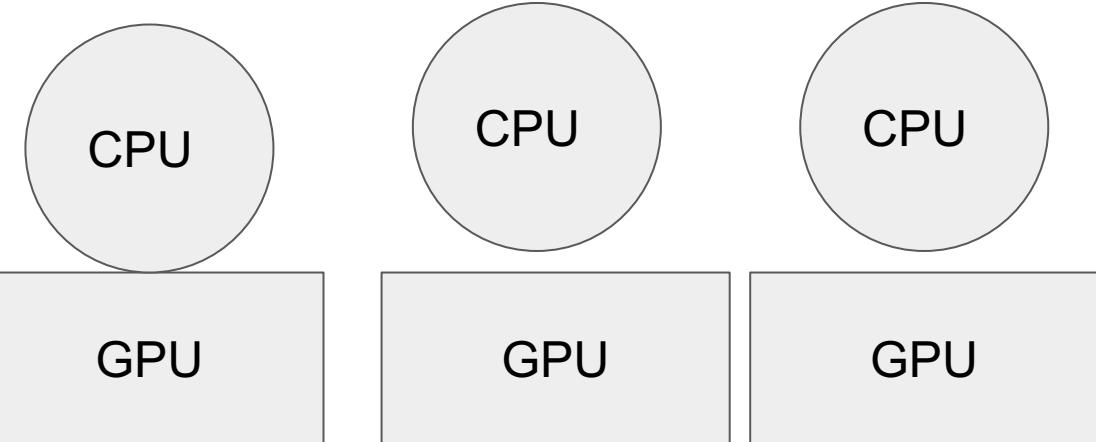
RL and Deep Learning workloads demand different resource requirements.



Logistic  
Regression

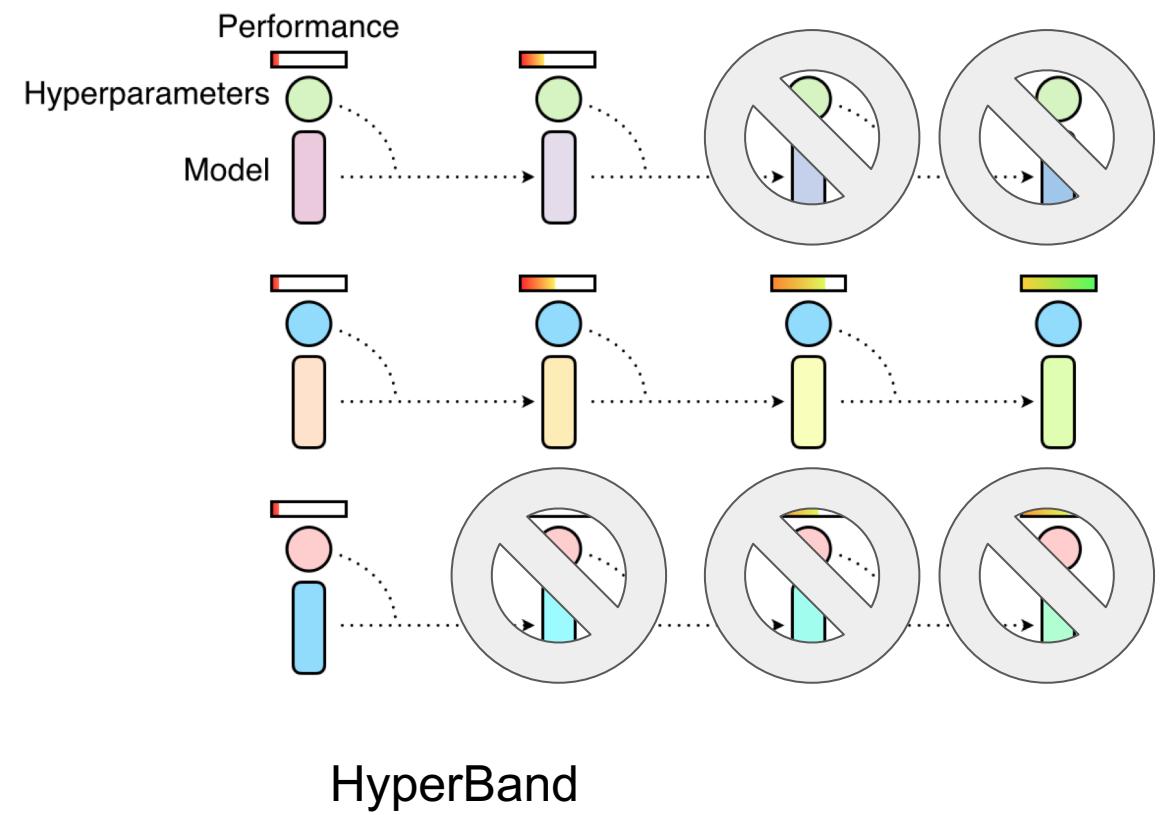
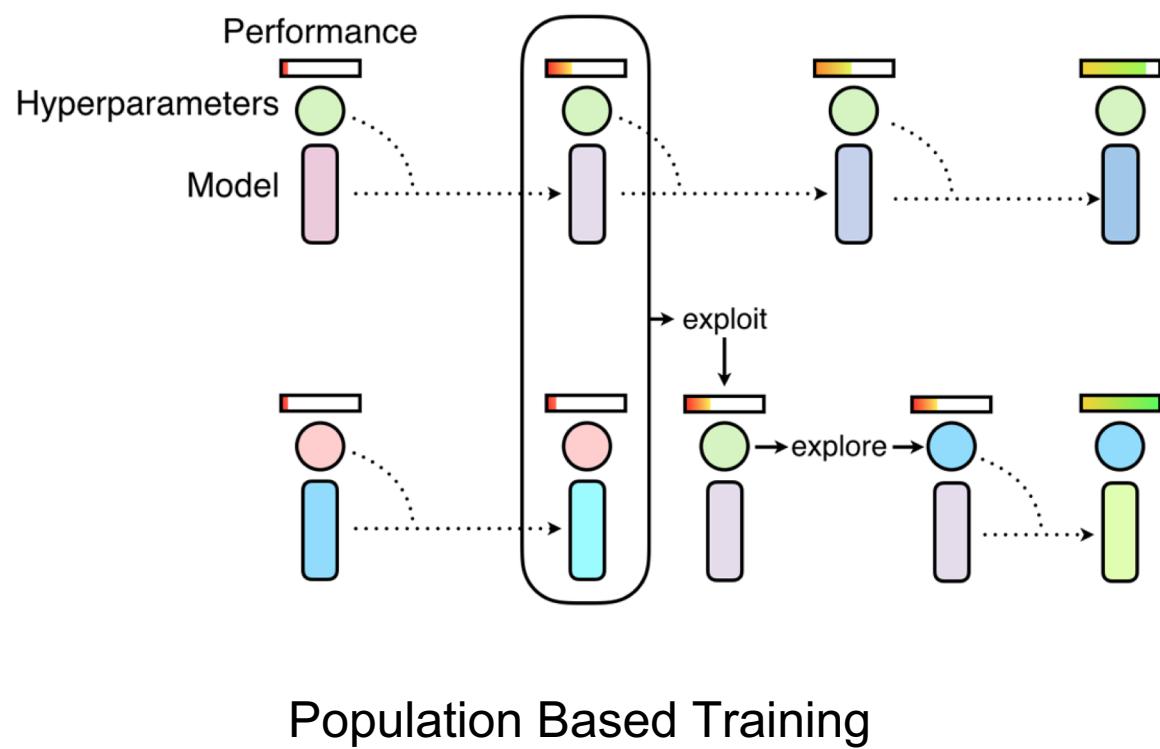


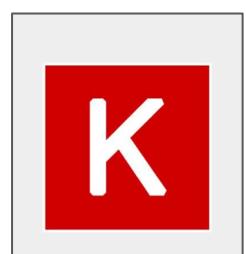
Deep Learning



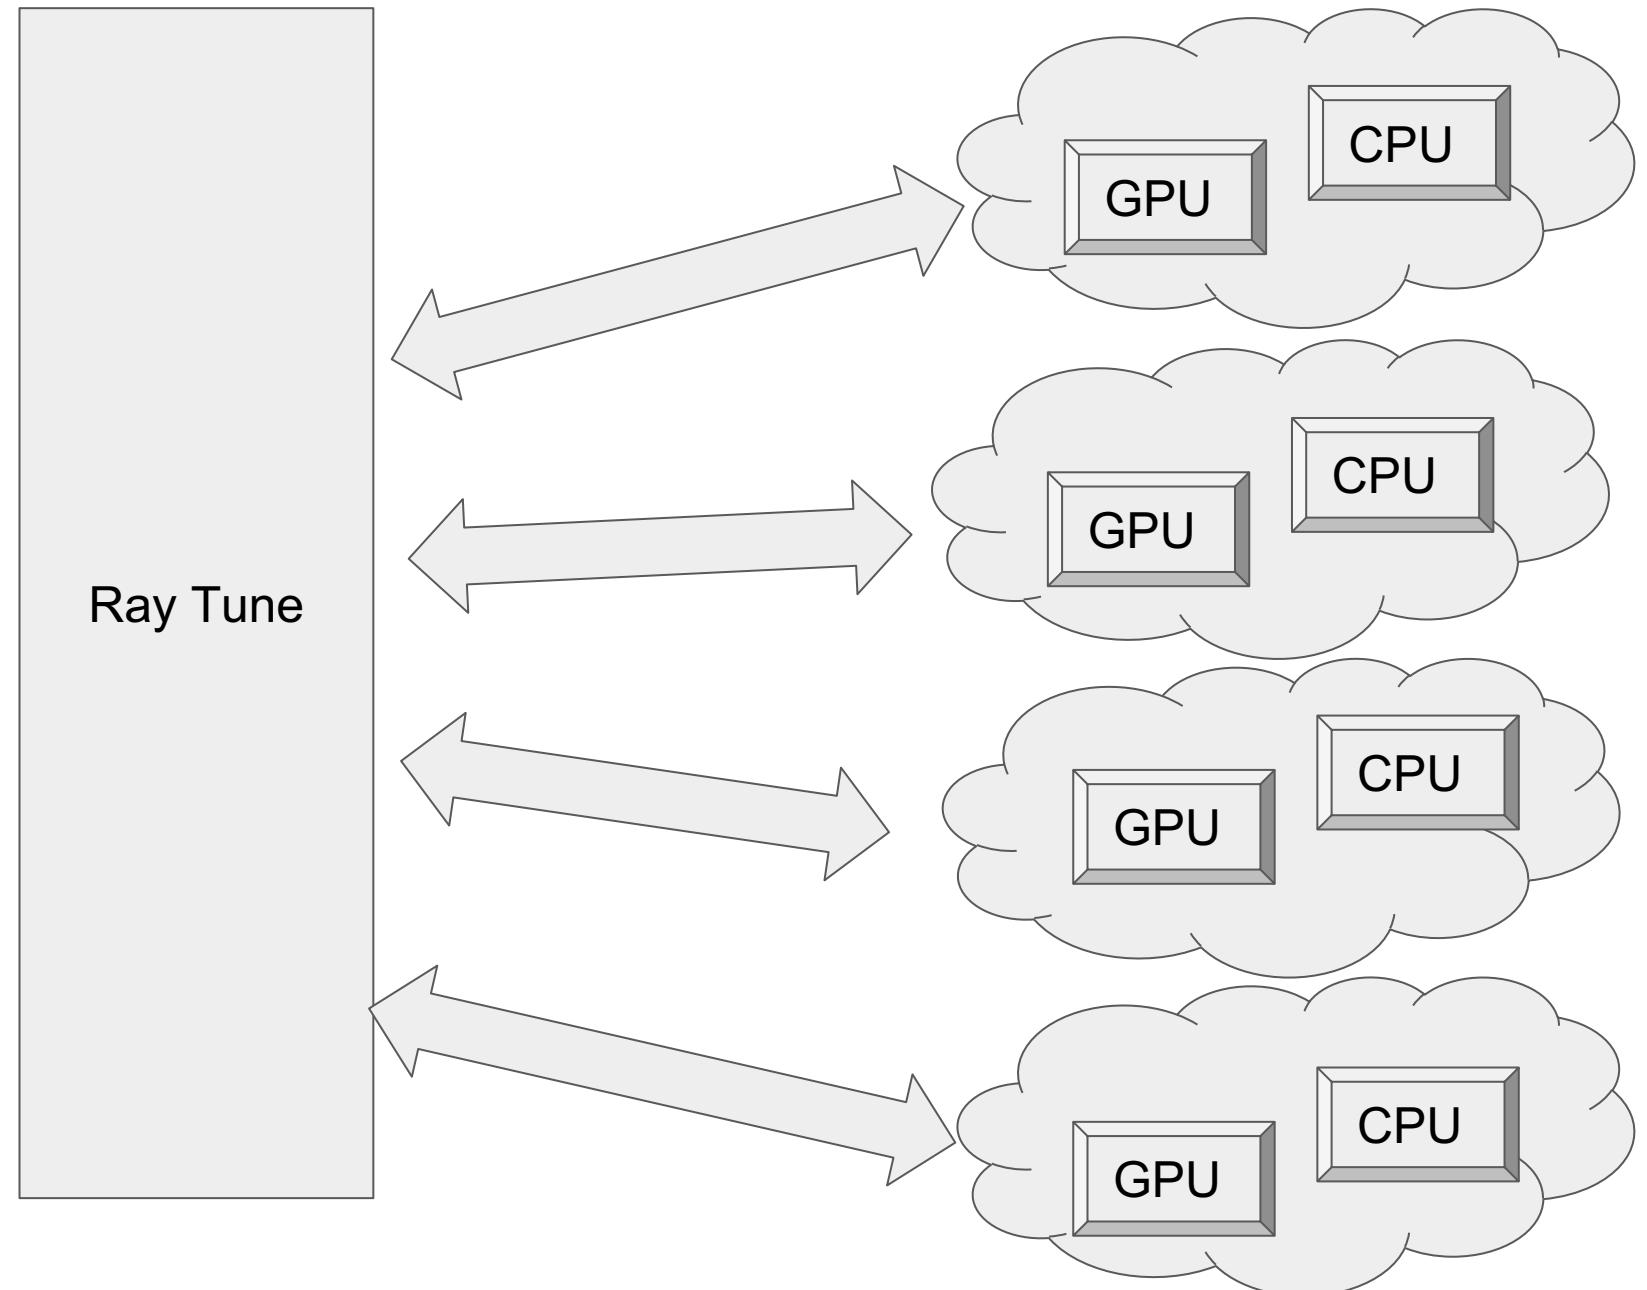
Deep Reinforcement  
Learning

# New Algorithms for hyperparameter tuning require more complicated control flows



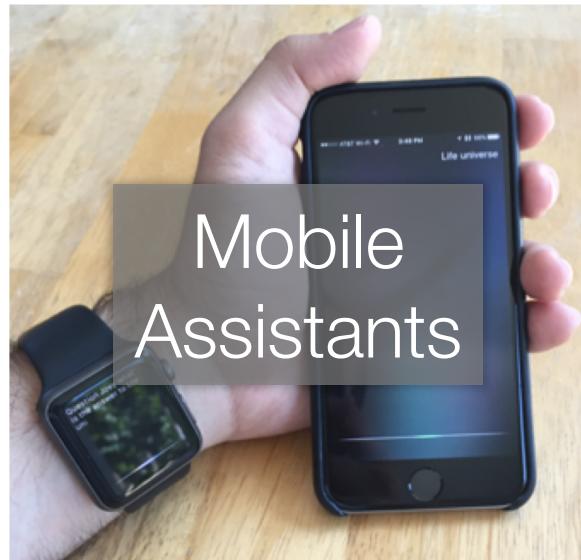
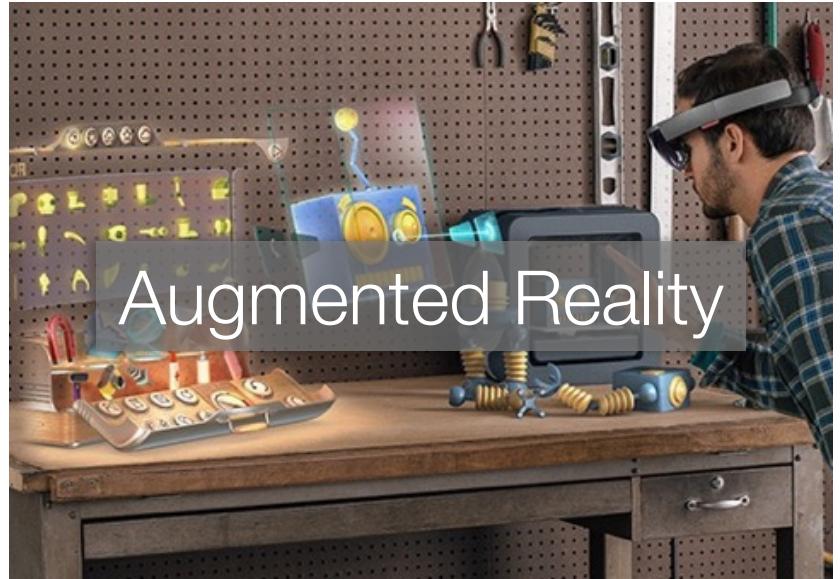


Ray Tune



# Security

# Machine Learning is on the critical path



# Machine Learning in Sensitive Contexts

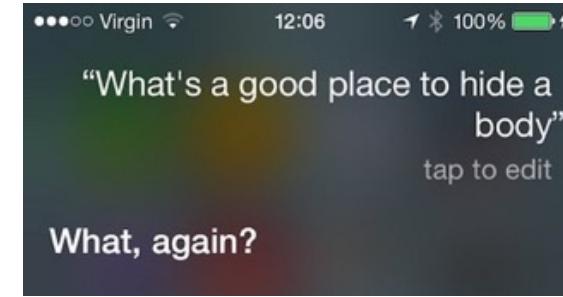
AR/VR Systems



Home Monitoring



Voice Technologies



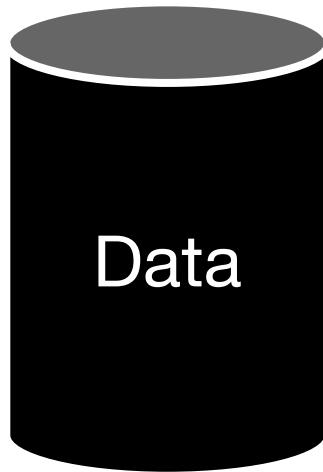
Medical Imaging



Protect the data, the model, and the query

# Protect the data, the model, and the query

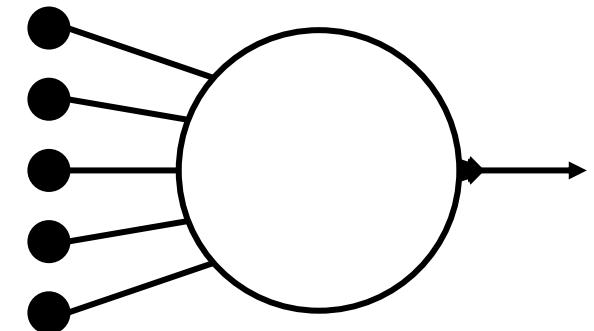
High-Value **Data is Sensitive**



- Medical Info.
- Home video
- Finance

**Models** capture **value** in data

- Core Asset
- “Contain”  
the data

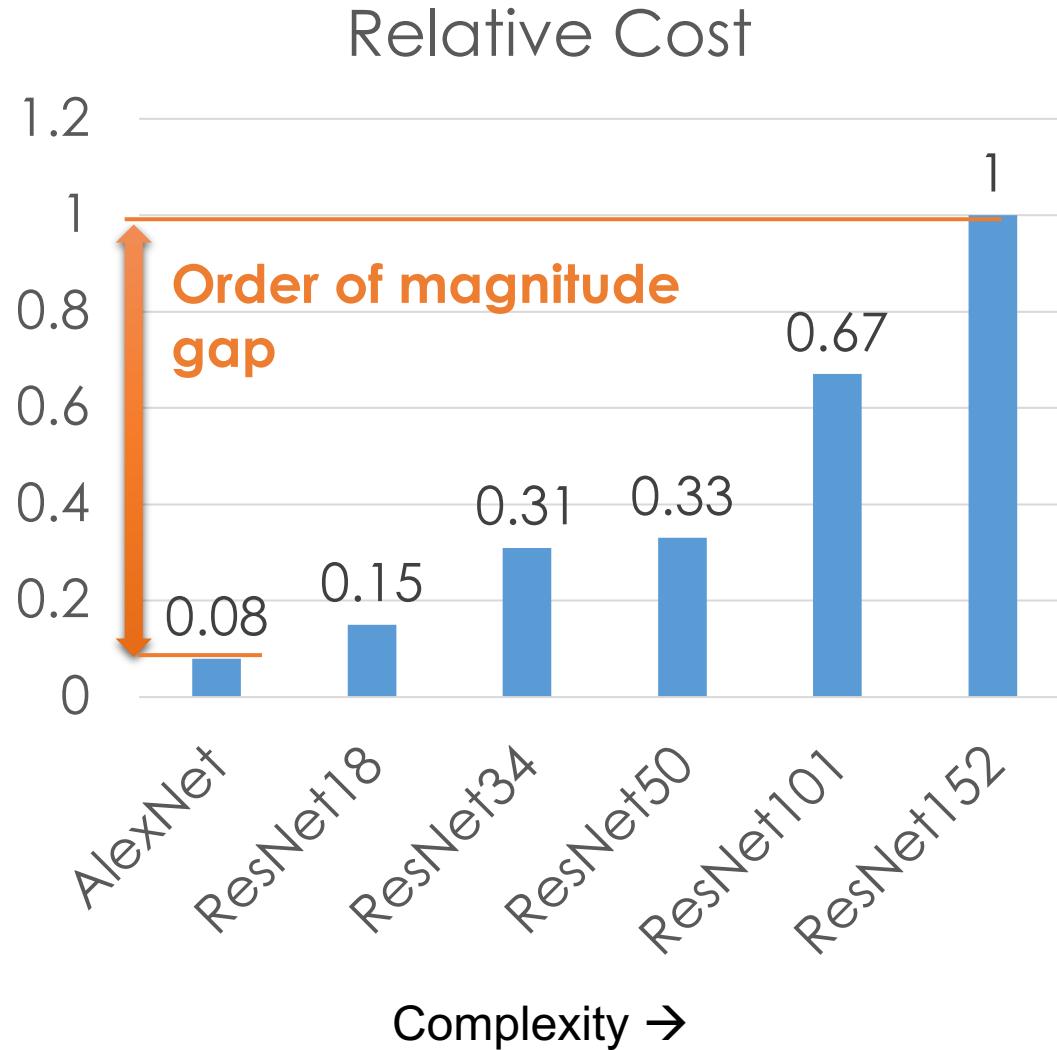
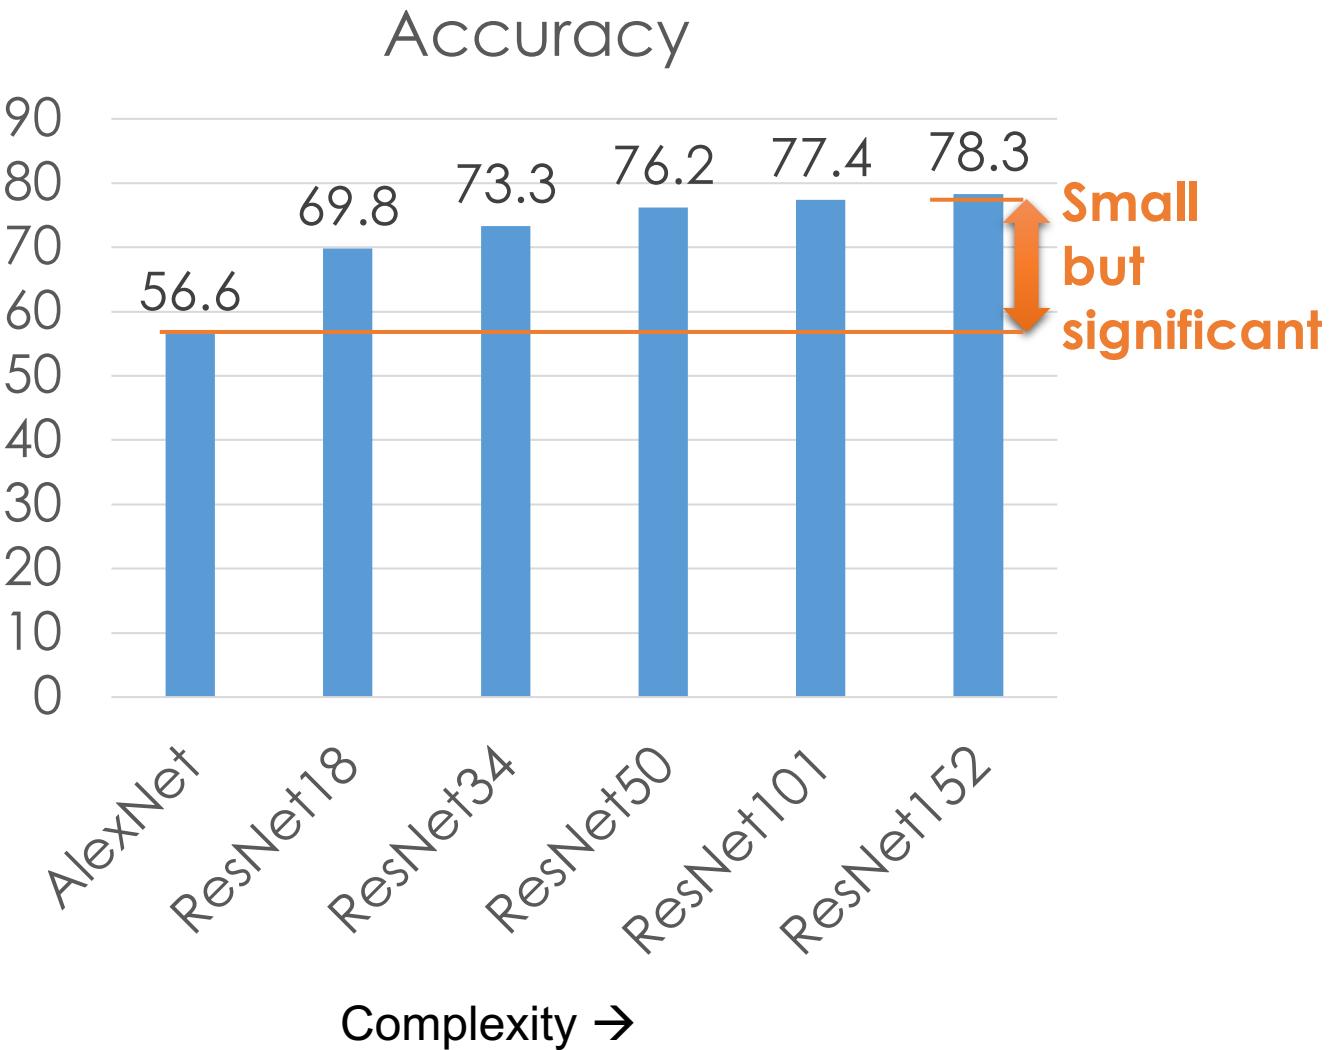


**Queries** can be as sensitive as the data



# Dynamic Inference

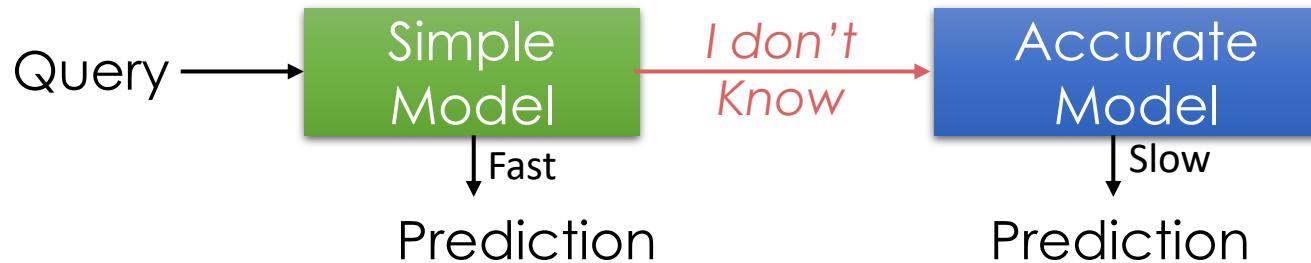
# Model **costs** are increasing much faster than gains in **accuracy**.



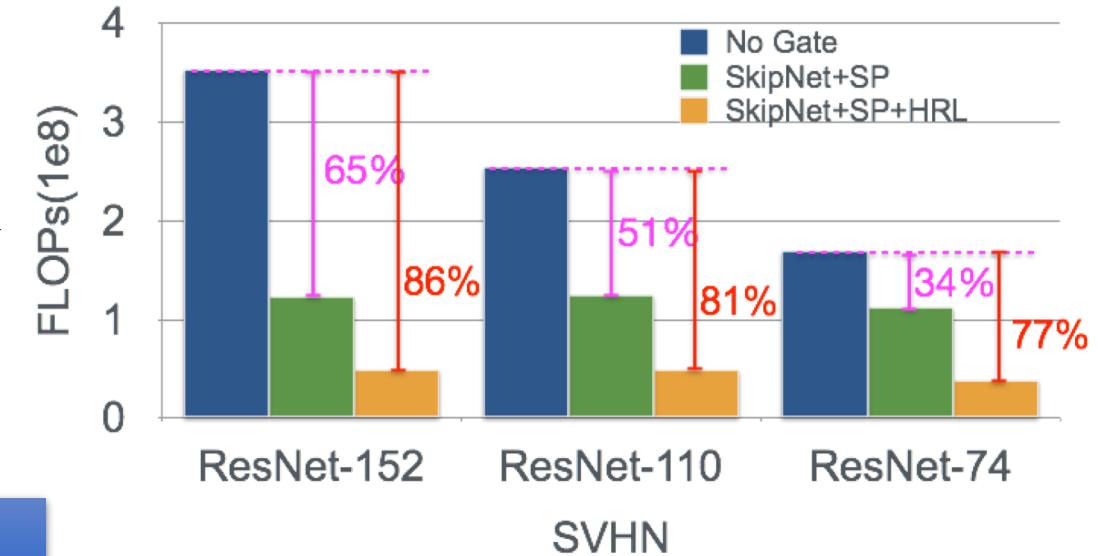
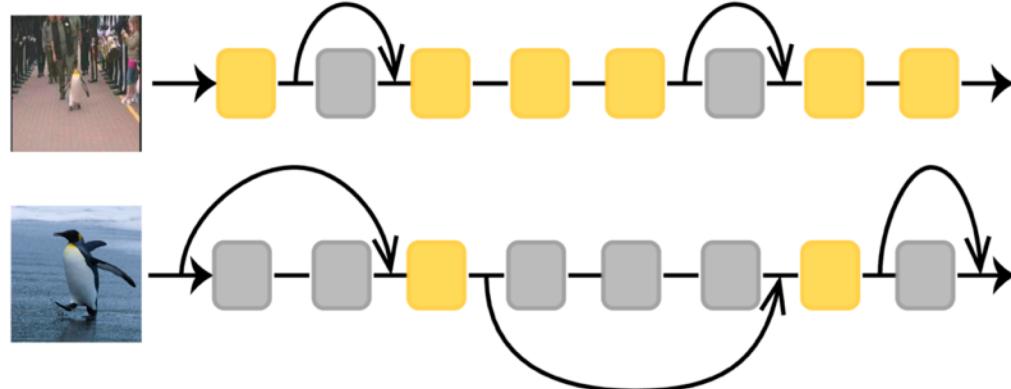
# Dynamic Networks

## Learning Not to Overthink

### IDK Prediction Cascades



### SkipNets: gated execution



**Easy** Images  
Skip **Many** Layers



**Hard** Images  
Skip **Few** Layers



# IDK Prediction Cascades

Simple models for simple tasks

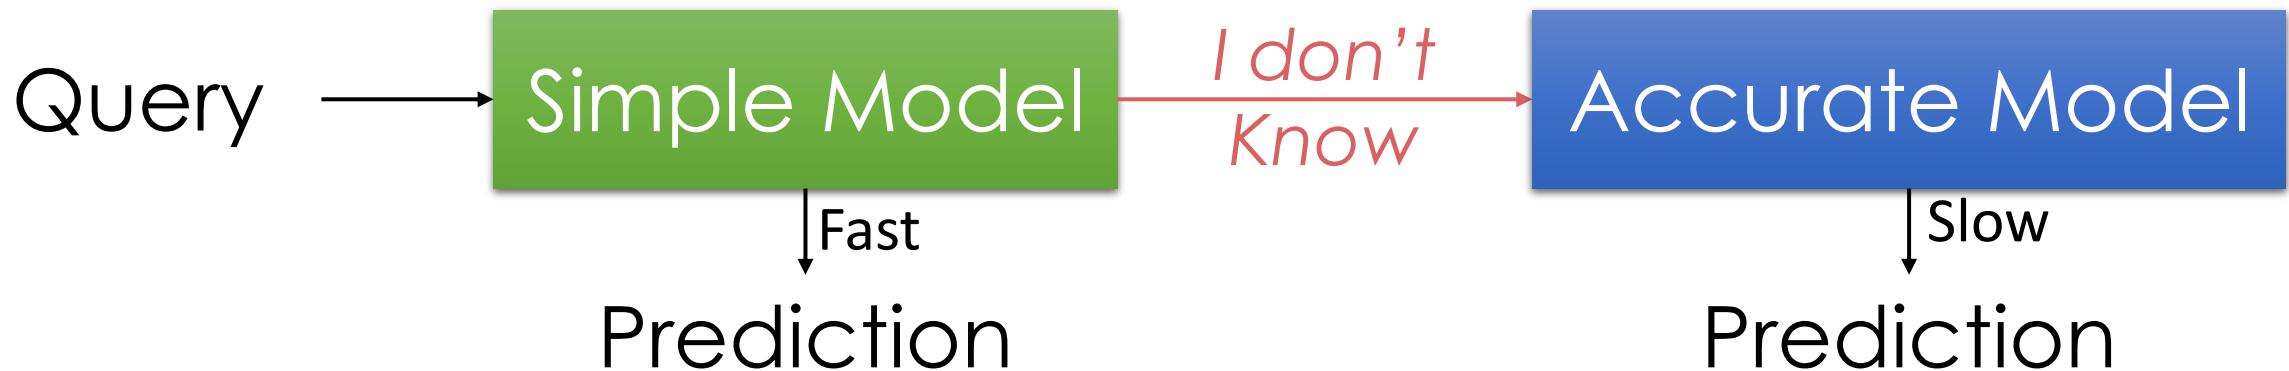


Xin  
Wang

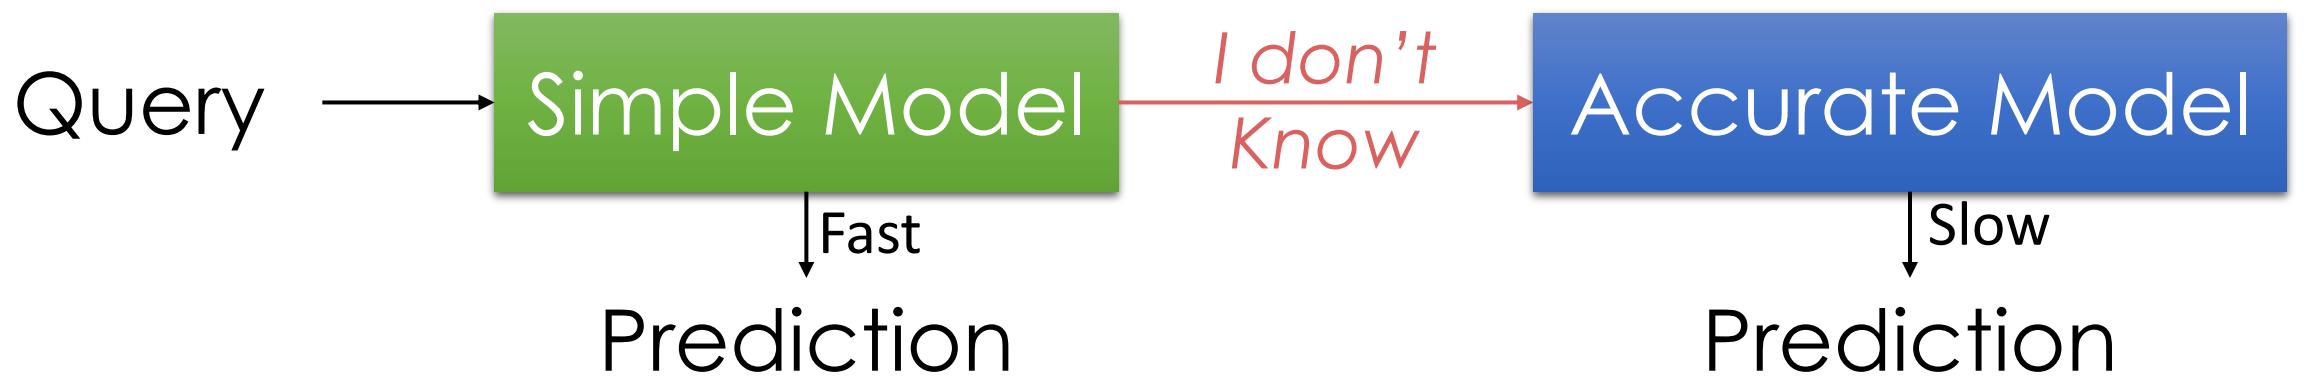
Yika  
Luo

Zi-Yi  
Duo

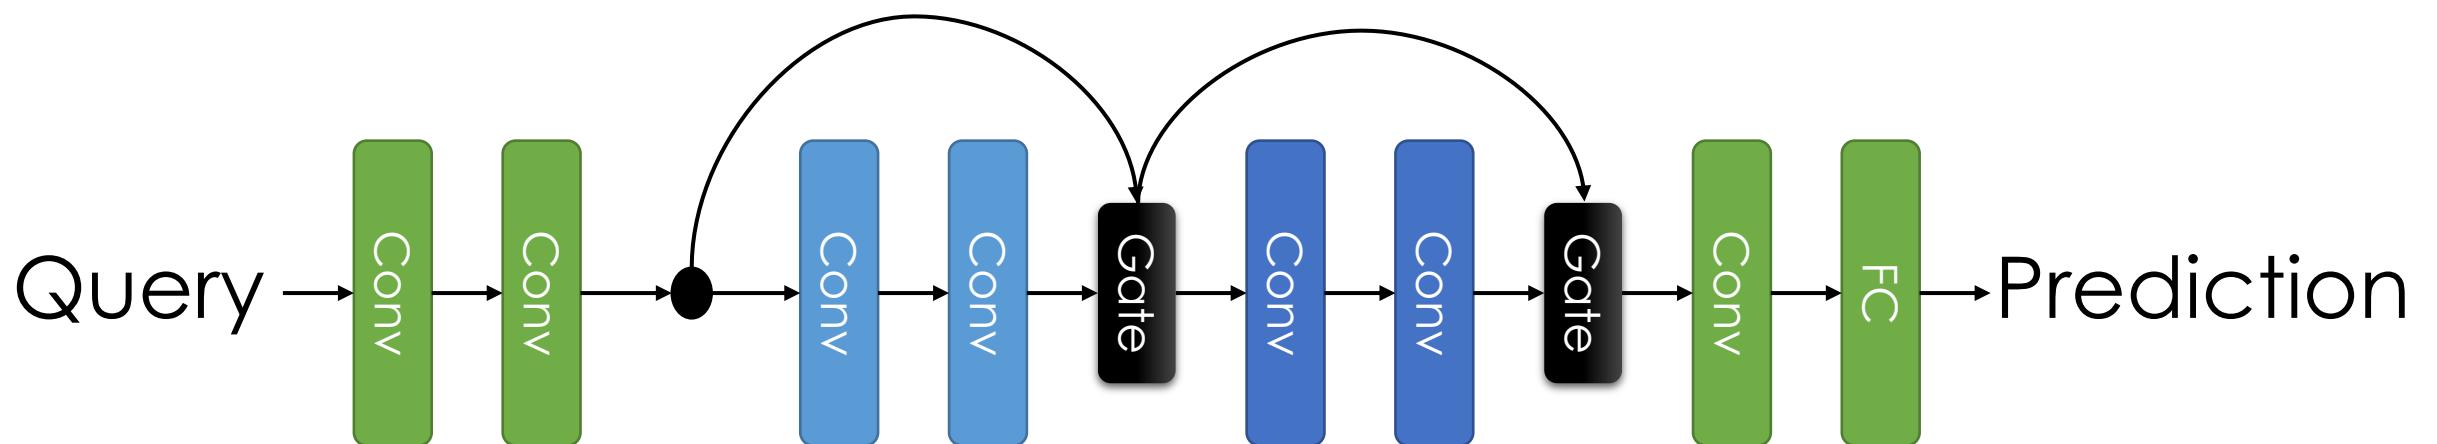
Fisher  
Yu

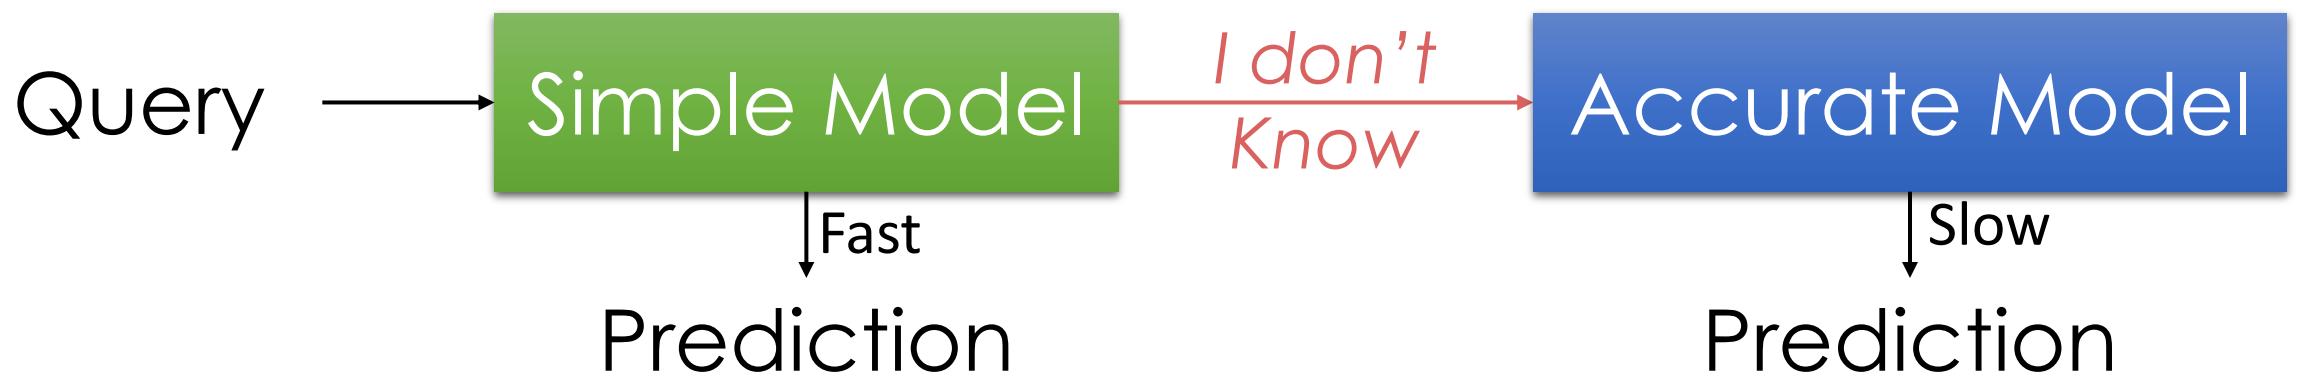


**Learn** to combine **fast (inaccurate) models** with **slow (accurate) models** to maximize accuracy while reducing computational costs.

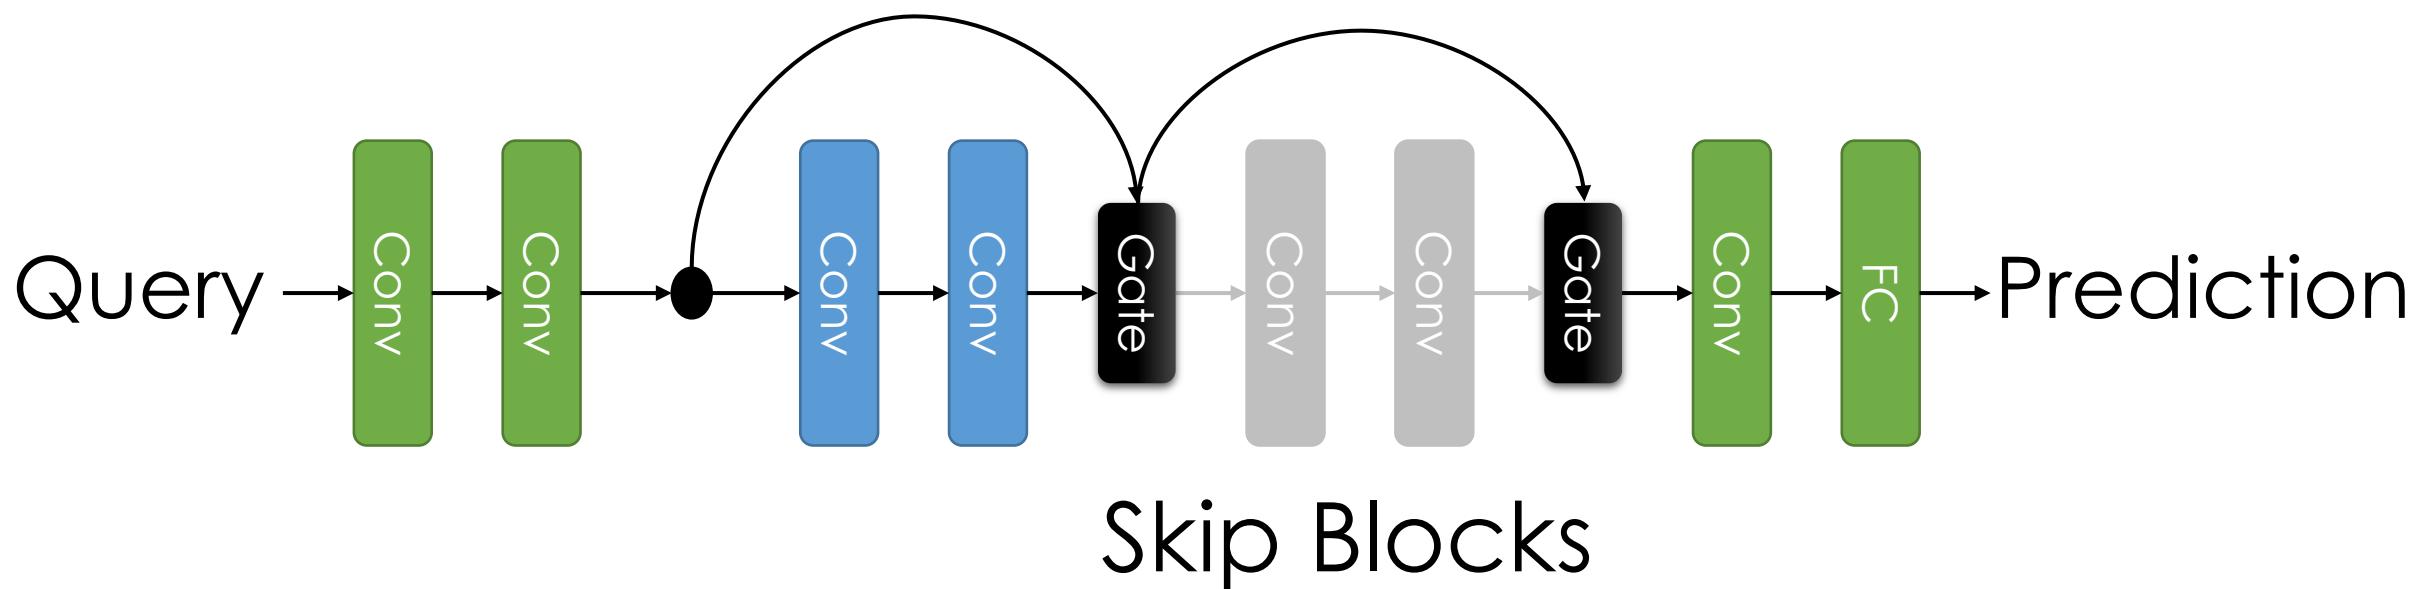


## SkipNet: dynamic execution within a model

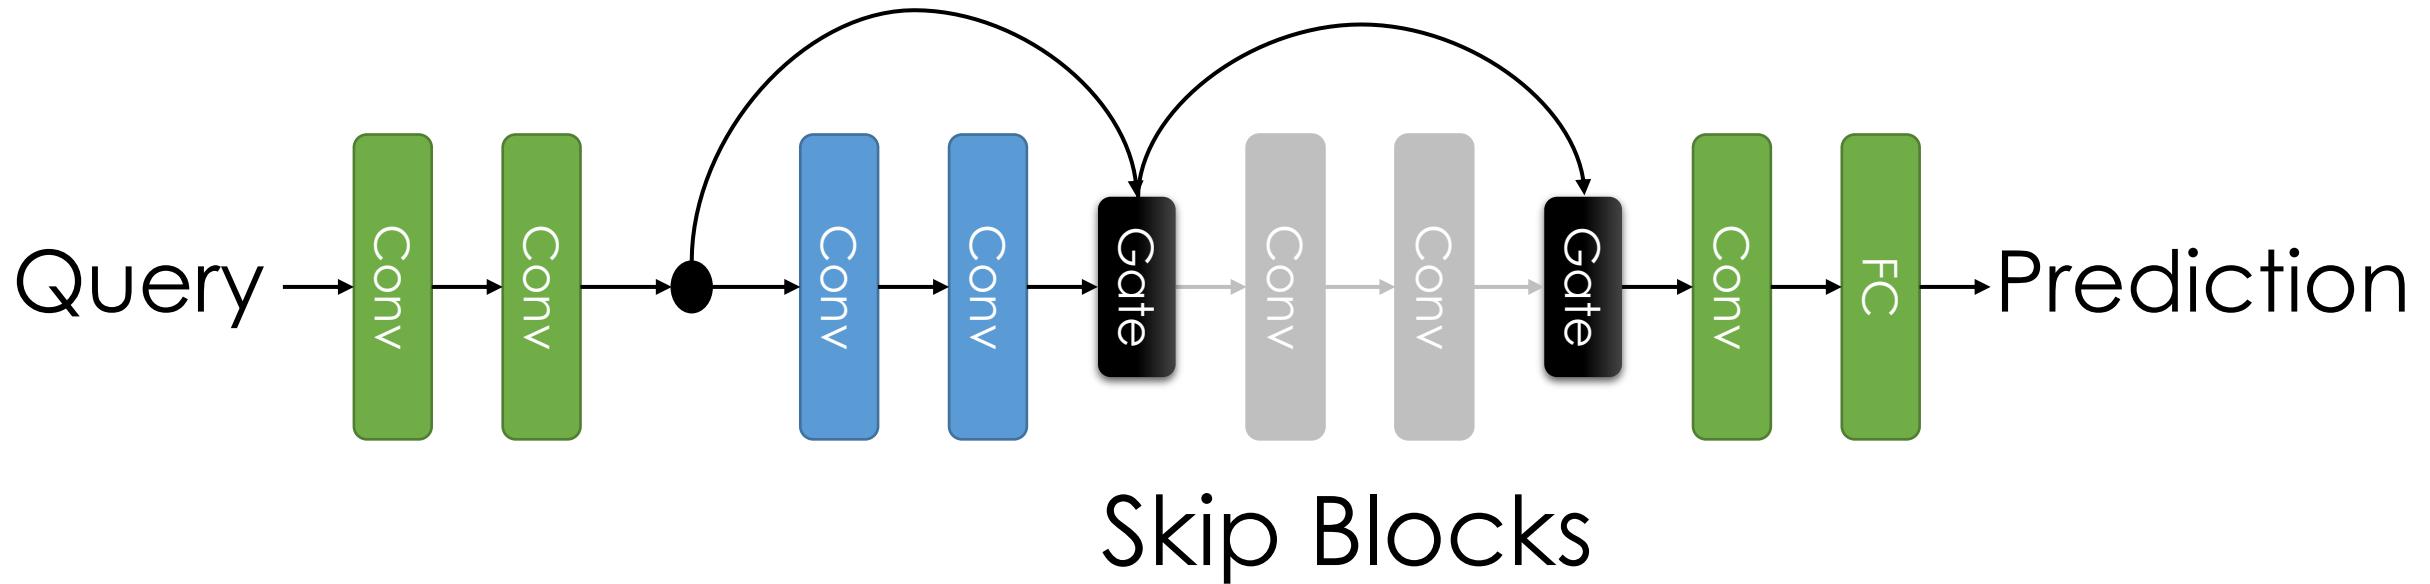




## SkipNet: dynamic execution within a model

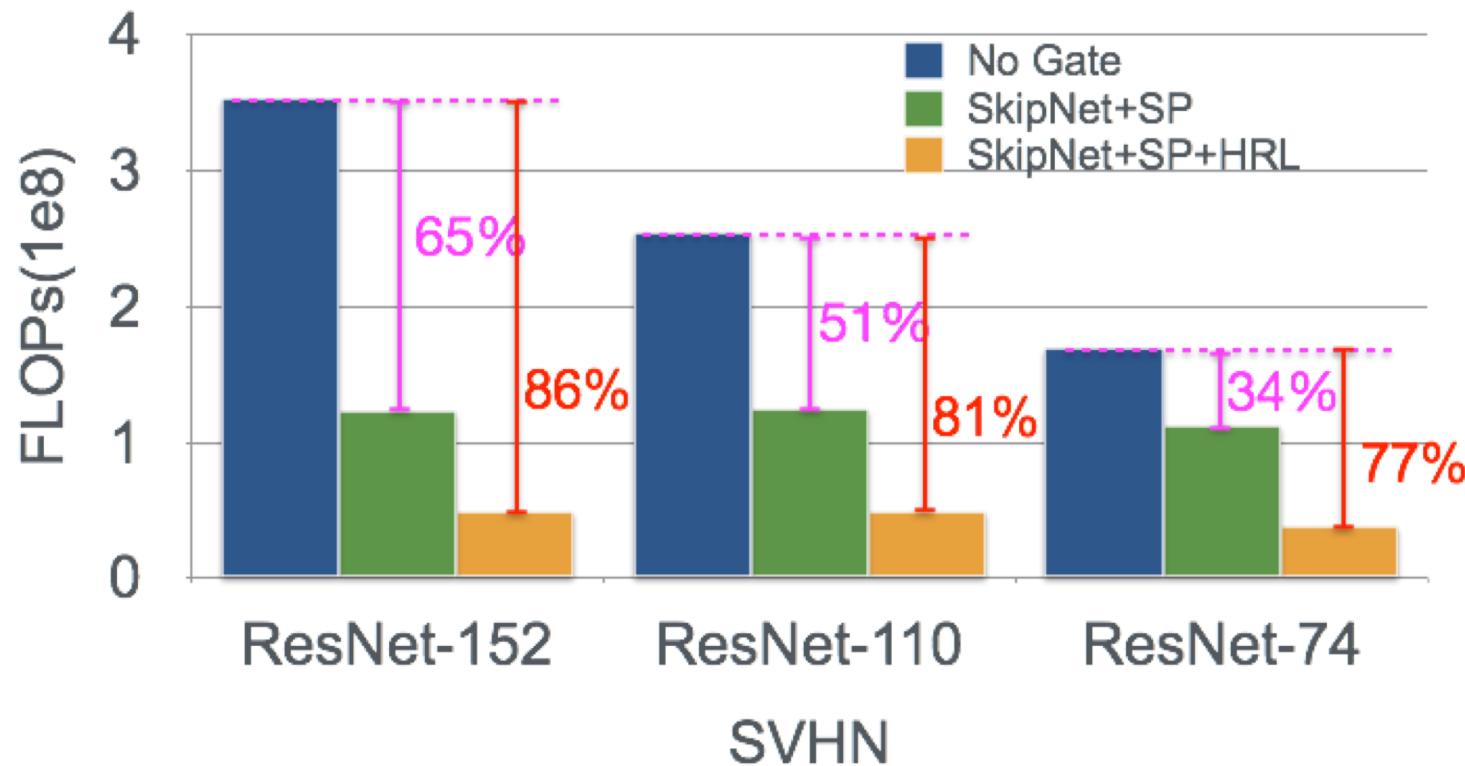


# SkipNet: dynamic execution within a model



- Combine **reinforcement learning** with **supervised pre-training** to learn a gating policy

# SkipNet Performance



**Easy Images**  
**Skip Many Layers**



**Hard Images**  
**Skip Few Layers**