

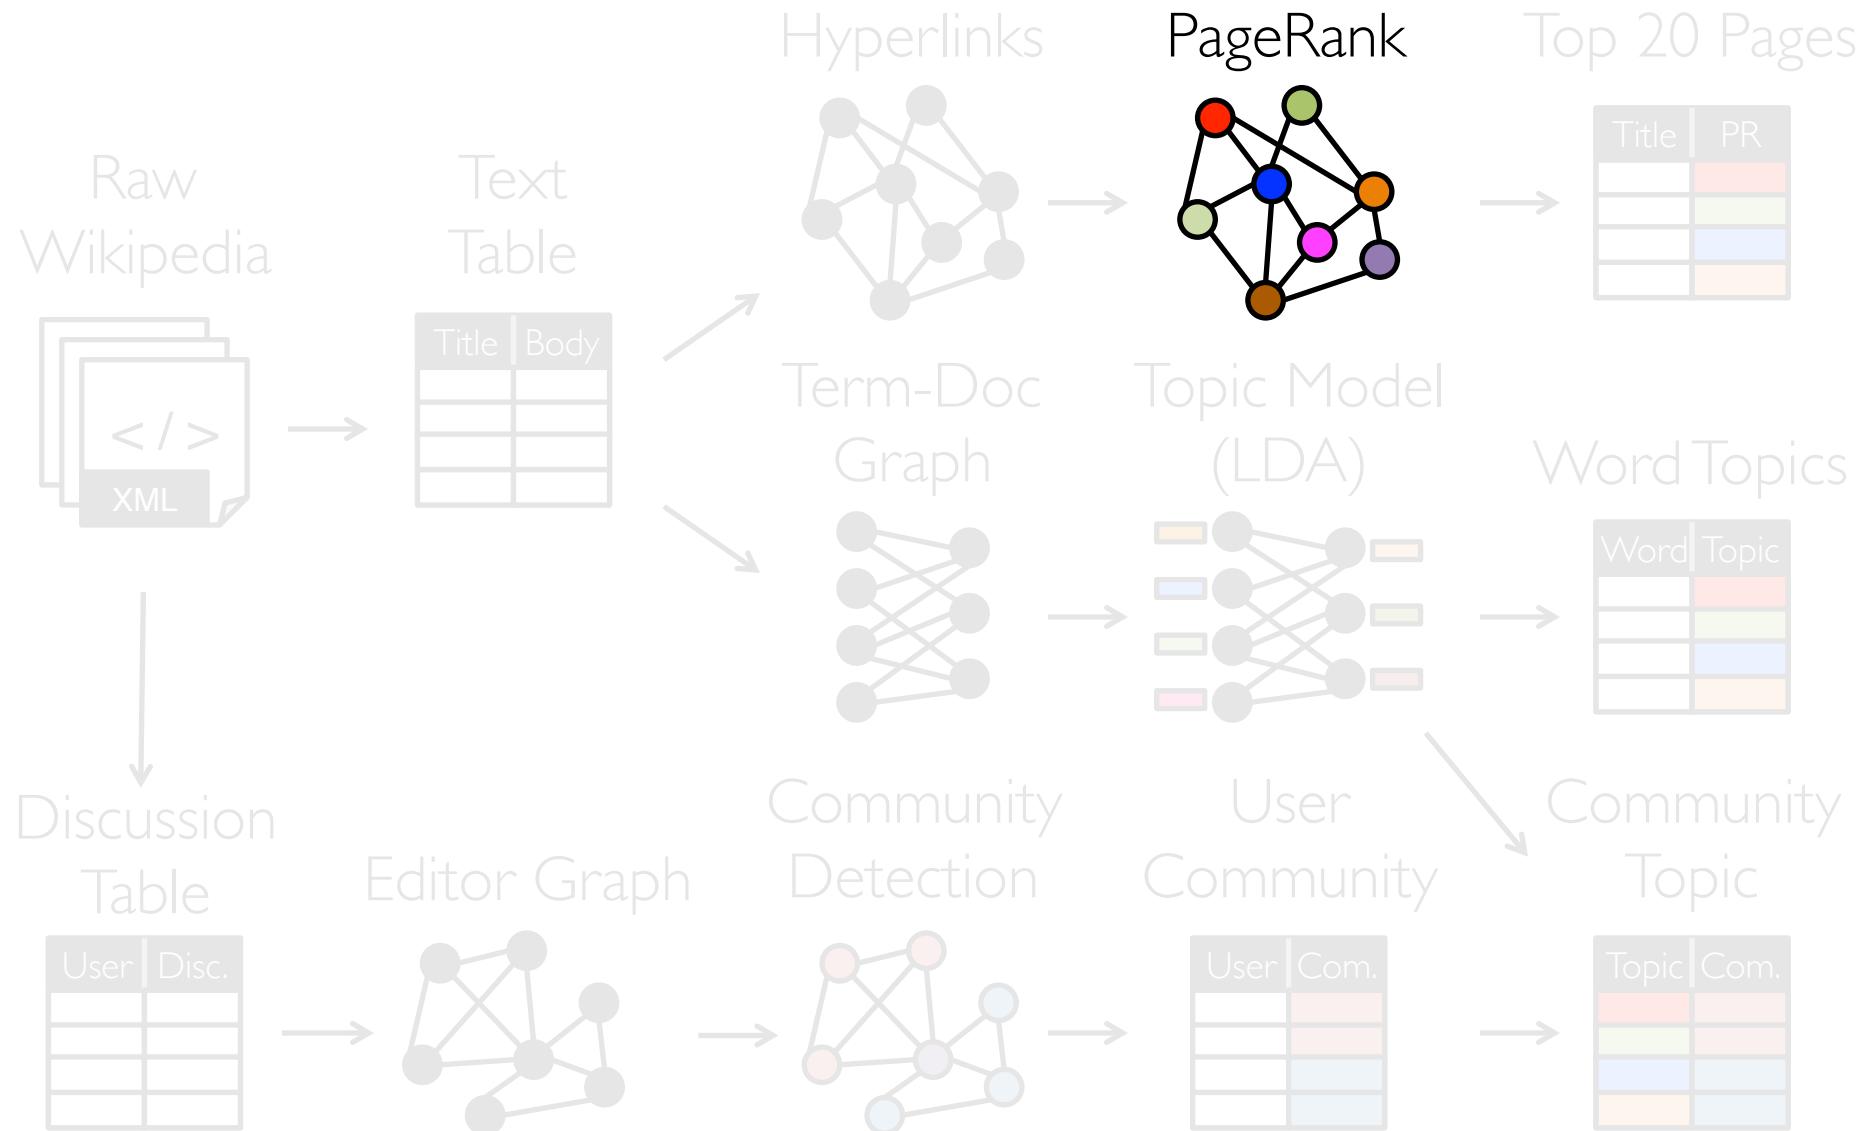
GraphX: Unifying Table and Graph Analytics

Presented by Joseph Gonzalez

Joint work with Reynold Xin, Daniel Crankshaw, Ankur Dave,
Michael Franklin, and Ion Stoica

IPDPS 2014

Graphs are Central to Analytics



PageRank: Identifying Leaders

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

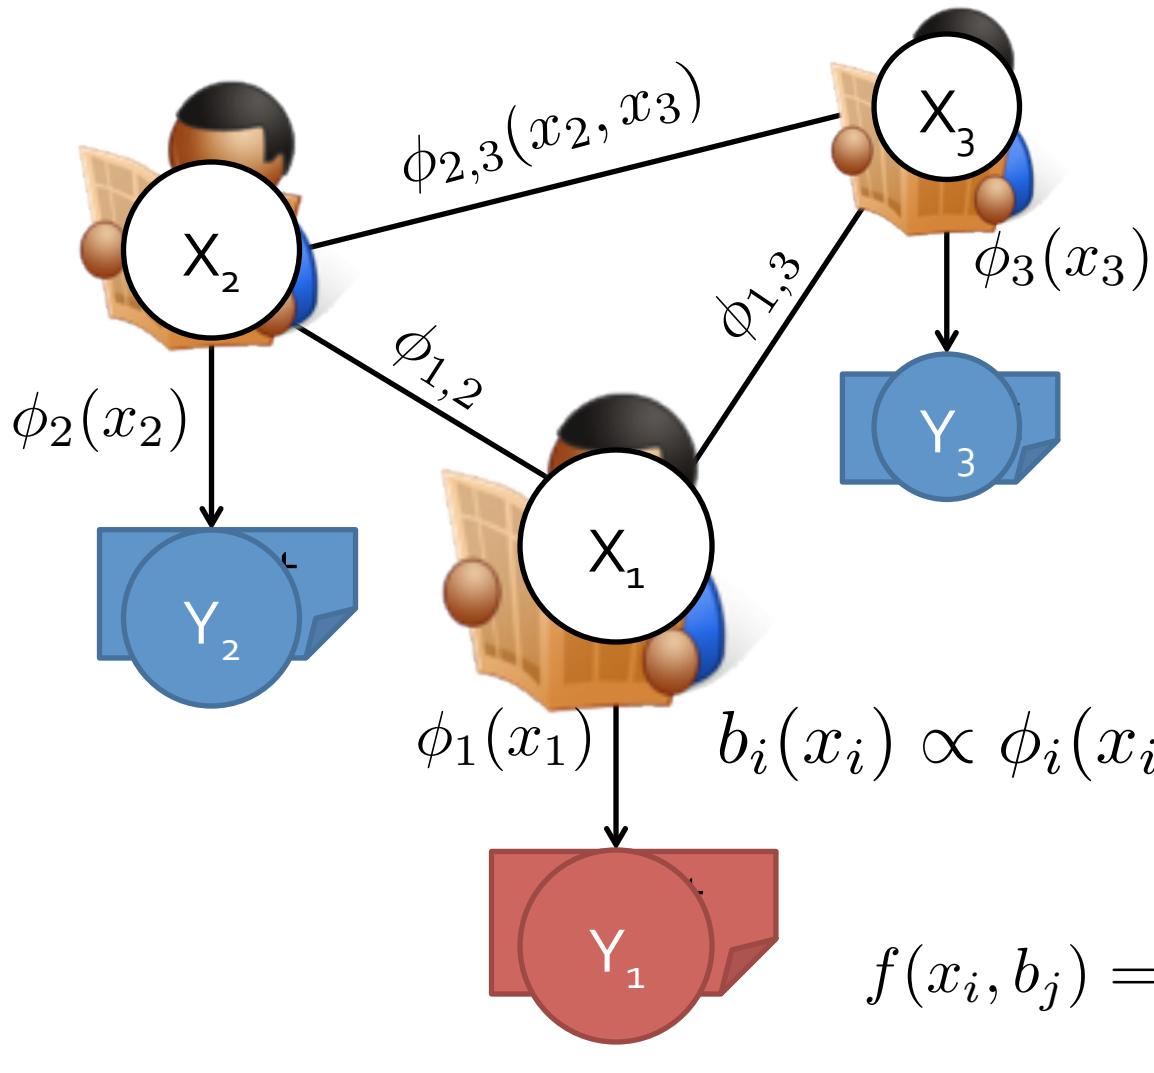
Rank of
user i

Weighted sum of
neighbors' ranks

Update ranks in parallel

Iterate until convergence

Mean Field Algorithm

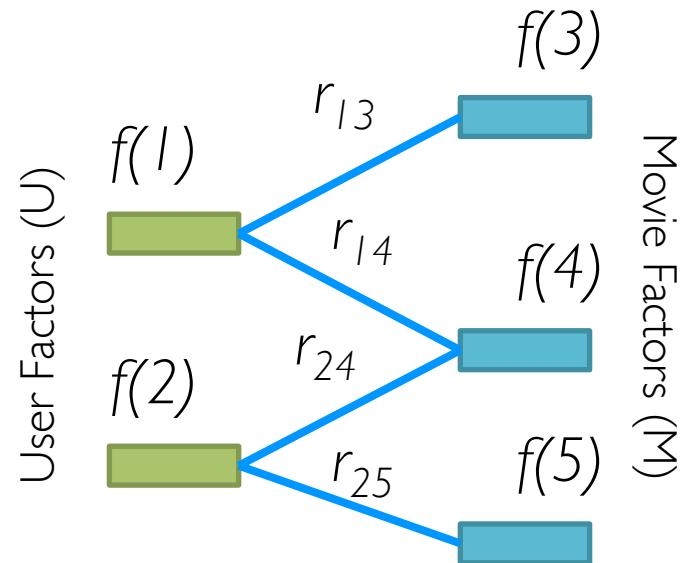
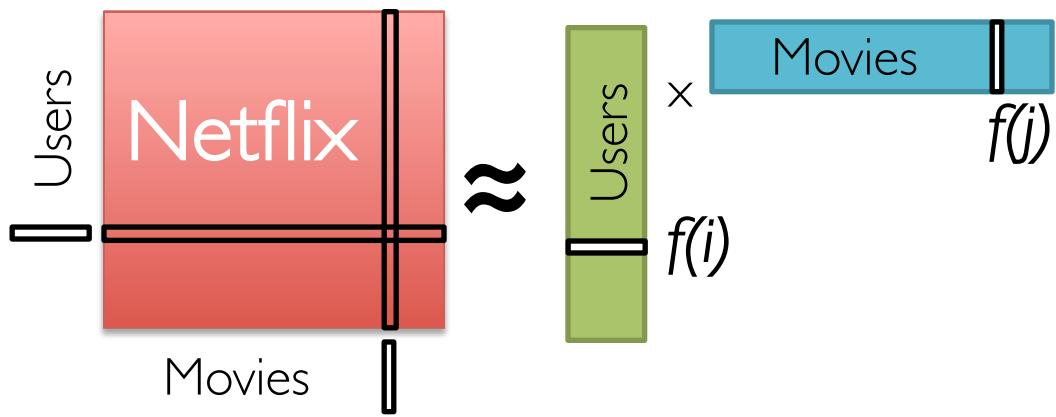


Sum over
Neighbors

$$f(x_i, b_j) = \sum_{x_j} b_j(x_j) \log \phi_{i,j}(x_i, x_j)$$
$$b_i(x_i) \propto \phi_i(x_i) \exp \left(\sum_{j \in N_i} f(x_i, b_j) \right)$$

Recommending Products

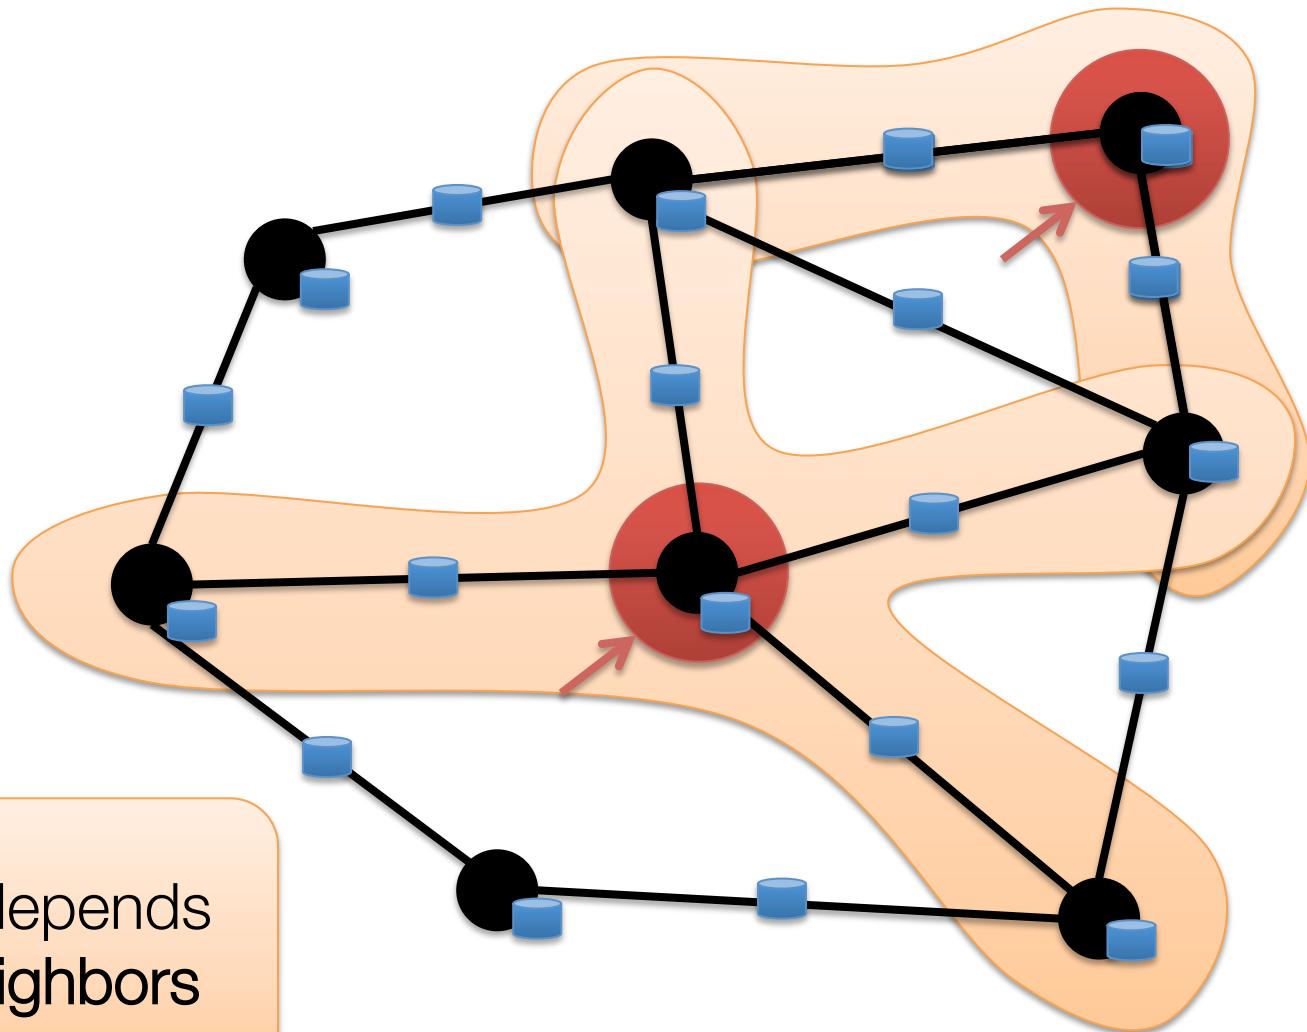
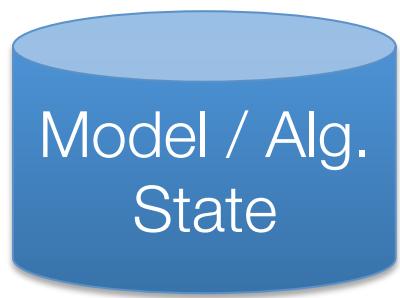
Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda \|w\|_2^2$$

The Graph-Parallel Pattern



Computation depends
only on the **neighbors**

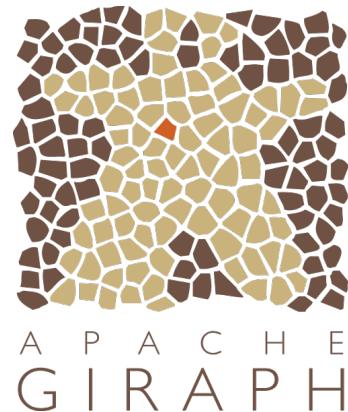
Many Graph-Parallel Algorithms

- Collaborative Filtering
 - Alternating Least Squares
 - Stochastic Gradient Descent
 - Tensor Factorization
- **MACHINE LEARNING**
 - Loopy Belief Propagation
 - Gibbs Sampling
- Semi-supervised ML
 - Graph SSL
 - CoEM

- Community Detection
 - SOCIAL NETWORK ANALYSIS**
 - K-core Decomposition
 - K-Truss
- Graph Analytics
 - PageRank
 - Personalized PageRank
 - Shortest Path
- **GRAPH ALGORITHMS**
 - Depth-First Search
 - Breadth-First Search
 - Connected Components
 - Spanning Trees
 - Minimum Spanning Tree
 - Max Flow
 - Min Cut
 - Topological Sort
 - Strongly Connected Components
 - Bellman-Ford
 - Dijkstra
 - Floyd-Warshall
 - Tarjan's Strongly Connected Components
 - Kosaraju's Strongly Connected Components
 - Depth-First Search
 - Breadth-First Search
 - Connected Components
- Classification
 - Neural Networks

Graph-Parallel Systems

Pregel
oo^{gle}



GraphLab

Expose *specialized APIs* to simplify graph
programming.

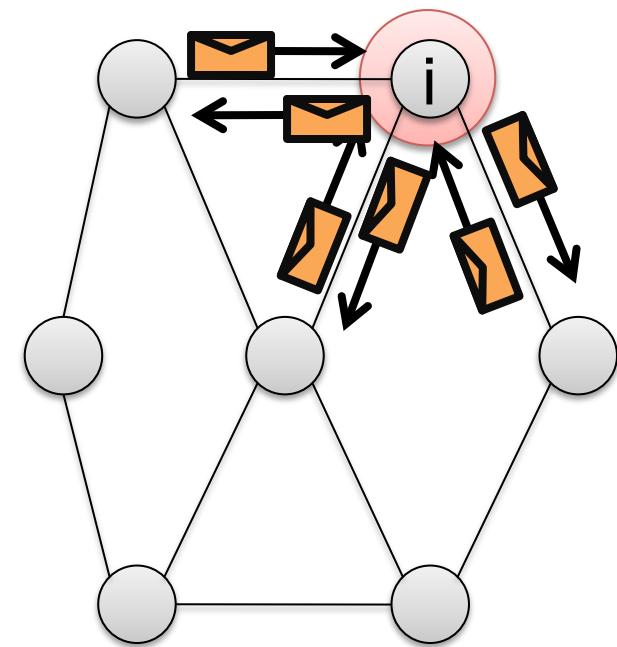
“Think like a Vertex.”

- Pregel [SIGMOD'10]

The Pregel (Push) Abstraction

Vertex-Programs interact by sending messages.

```
Pregel_PageRank(i, messages) :  
    // Receive all the messages  
    total = 0  
    foreach( msg in messages ) :  
        total = total + msg  
  
    // Update the rank of this vertex  
    R[i] = 0.15 + total  
  
    // Send new messages to neighbors  
    foreach(j in out_neighbors[i]) :  
        Send msg(R[i]) to vertex j
```



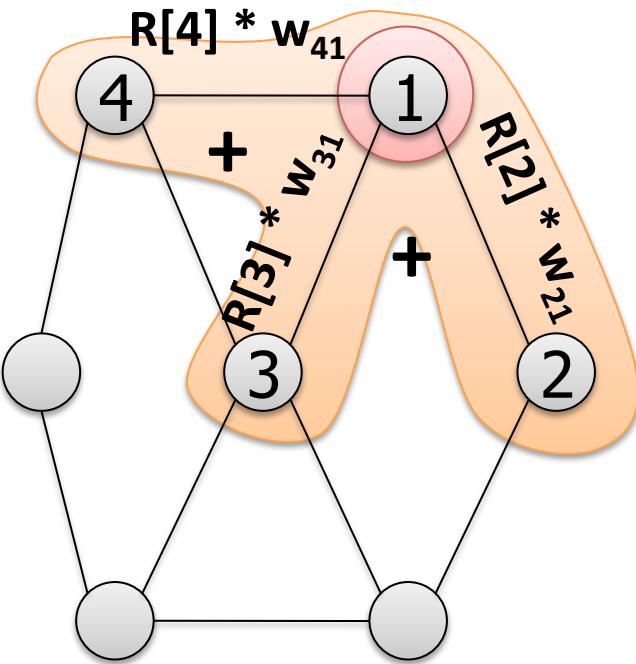
The GraphLab (Pull) Abstraction

Vertex Programs directly access adjacent vertices and edges

```
GraphLab_PageRank(i)
```

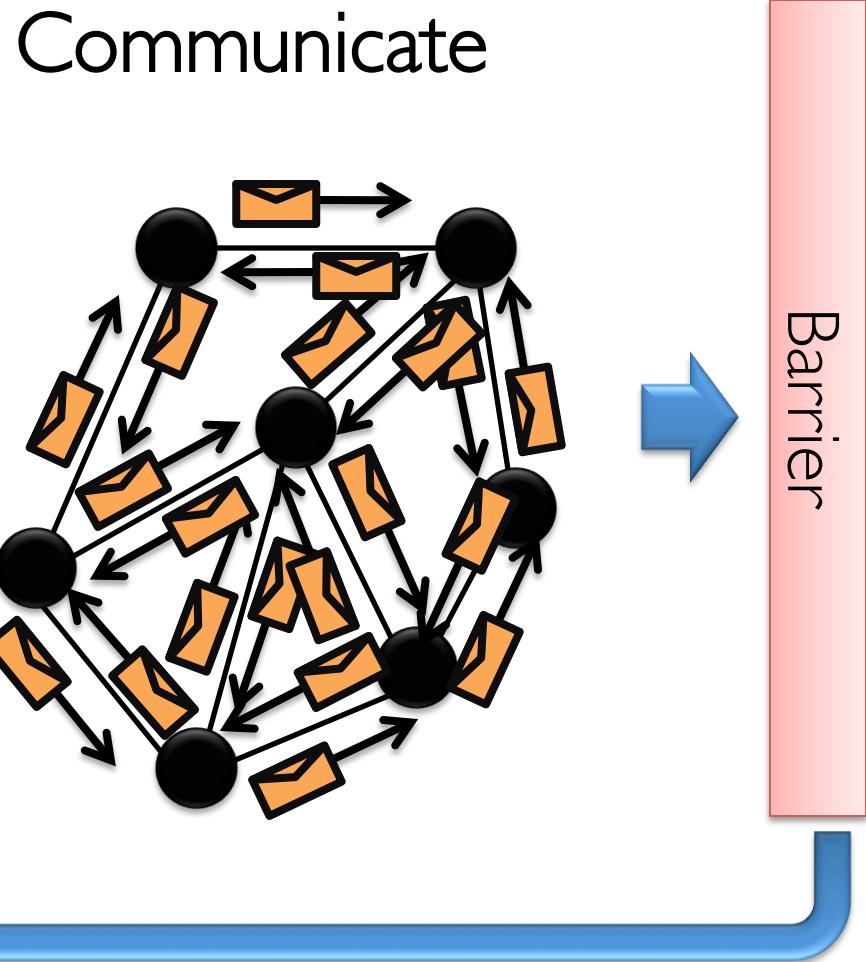
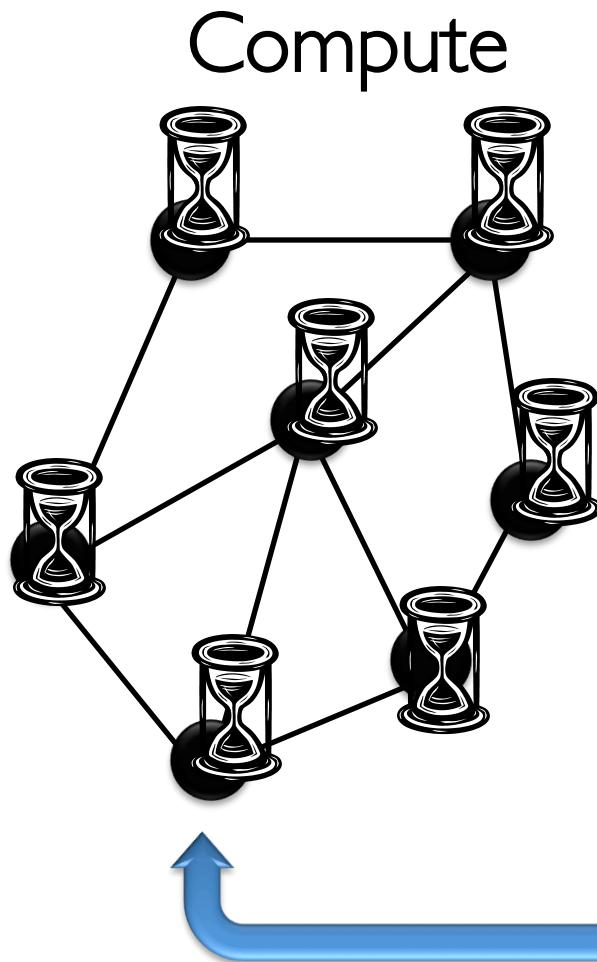
```
// Compute sum over neighbors
total = 0
foreach( j in neighbors(i)):
    total = total + R[j] * wji
```

```
// Update the PageRank
R[i] = 0.15 + total
```



Data movement is managed by the system
and not the user.

Iterative Bulk Synchronous Execution



Graph-Parallel Systems

Pregel
oo^{gle}

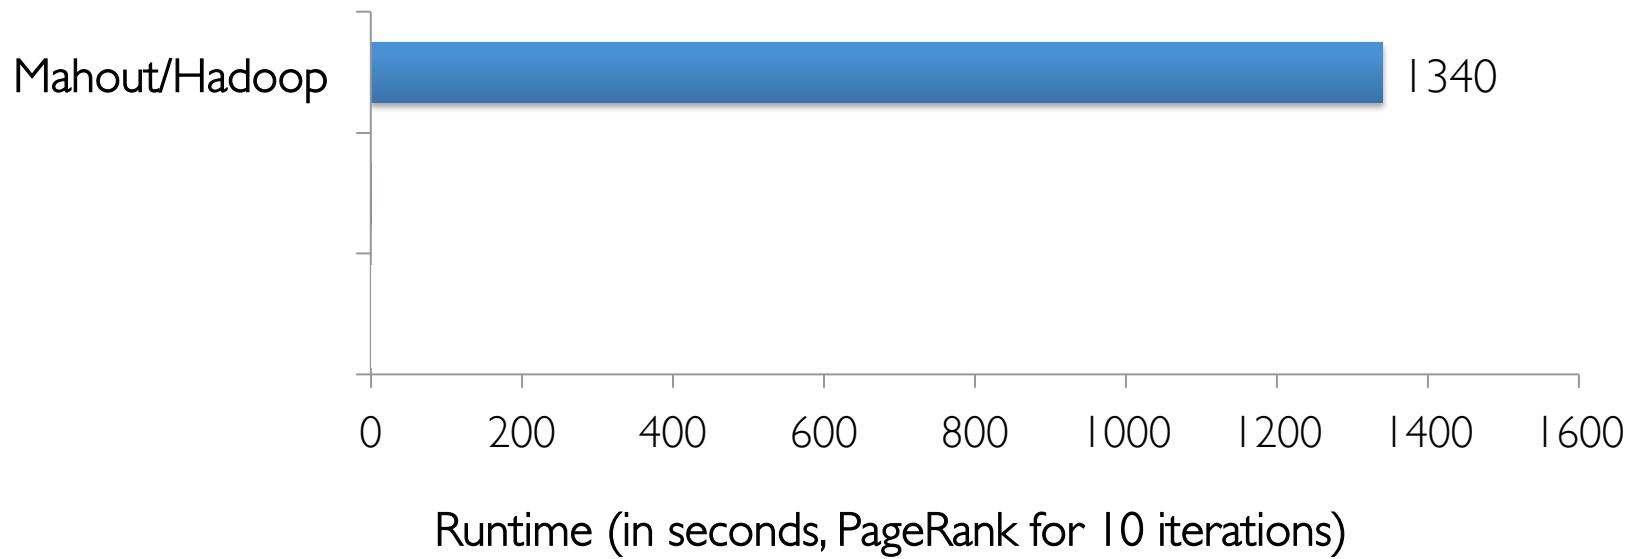


GraphLab

Expose *specialized APIs* to simplify graph programming.

Exploit graph structure to achieve *orders-of-magnitude performance gains* over more general data-parallel systems.

PageRank on the Live-Journal Graph



Spark is *4x faster* than Hadoop
GraphLab is *16x faster* than Spark

Triangle Counting on Twitter

40M Users, 1.4 Billion Links

Counted: 34.8 Billion Triangles

Hadoop
[WWW'11]

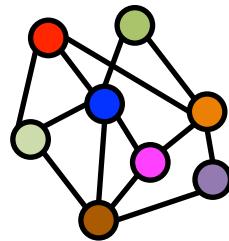
1536 Machines
423 Minutes

GraphLab

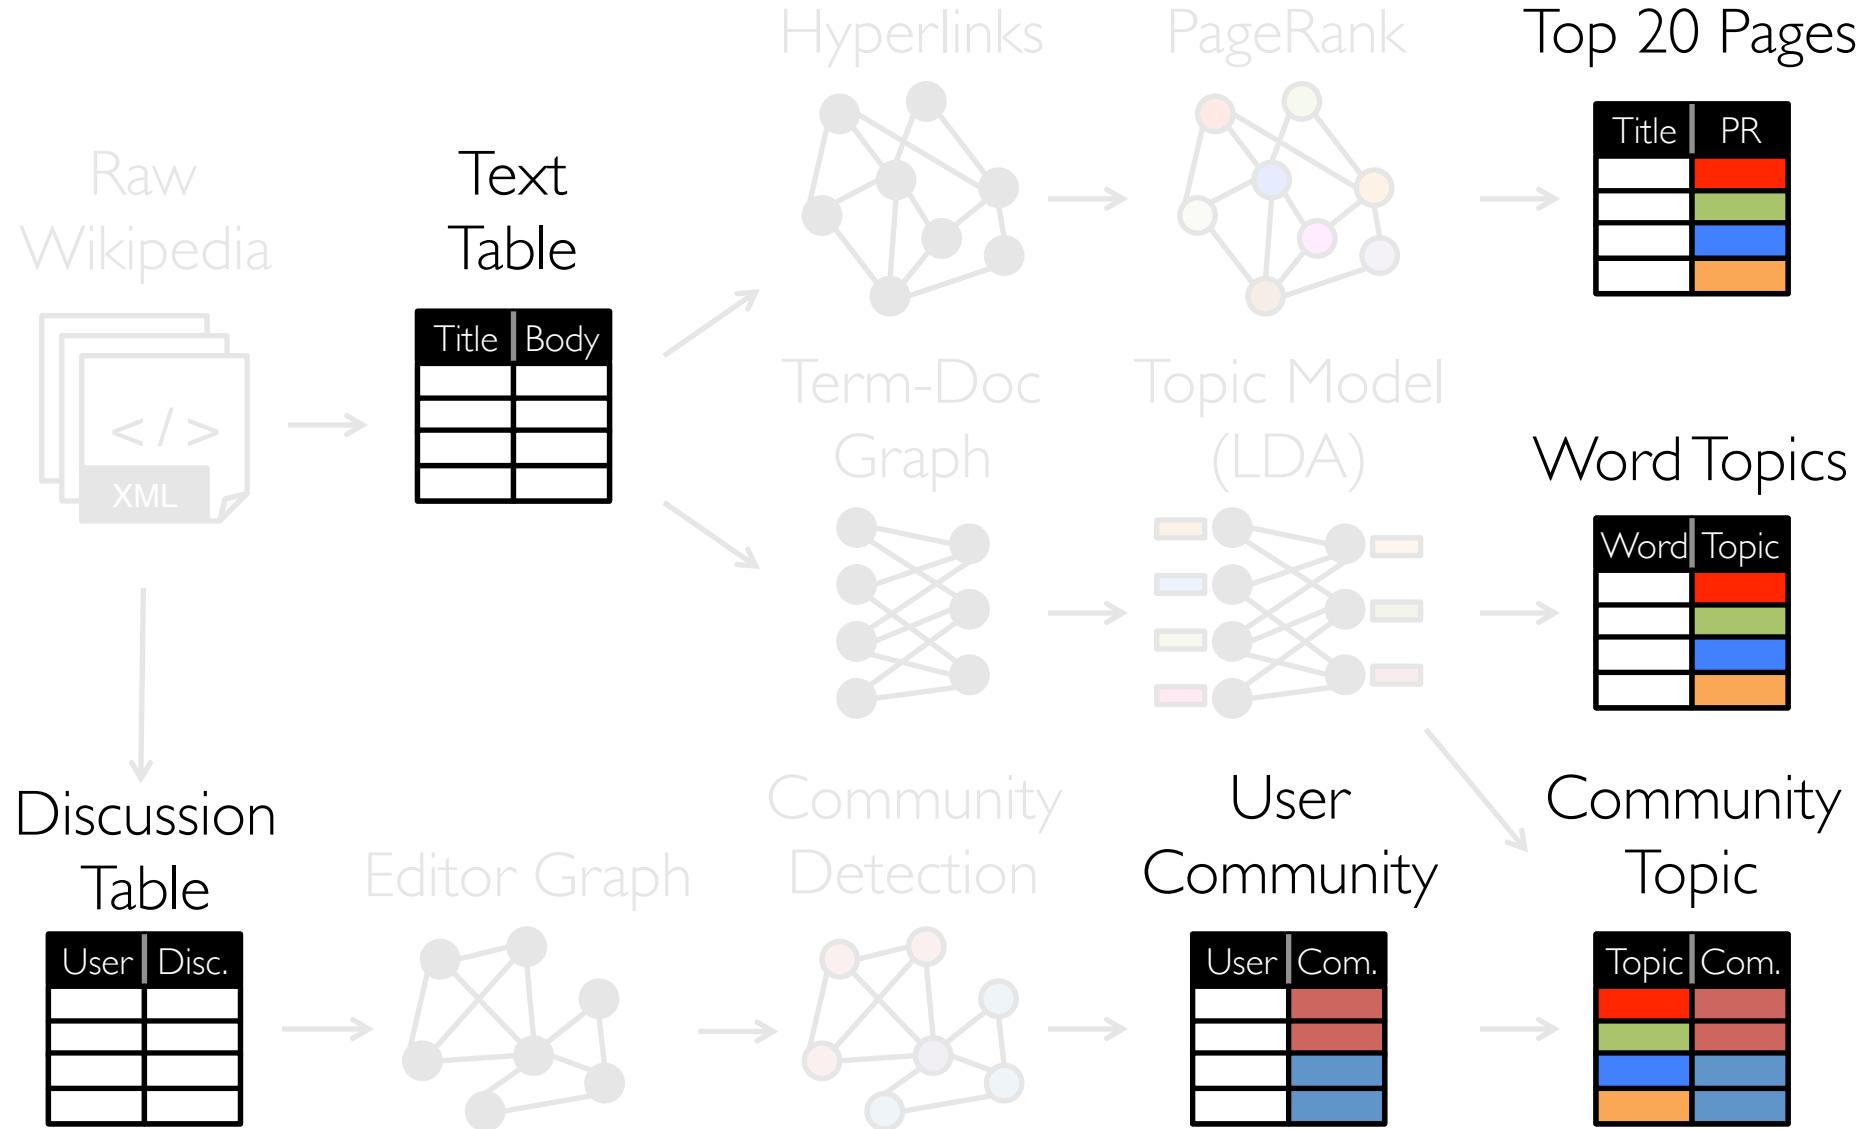
64 Machines
15 Seconds

1000 × Faster

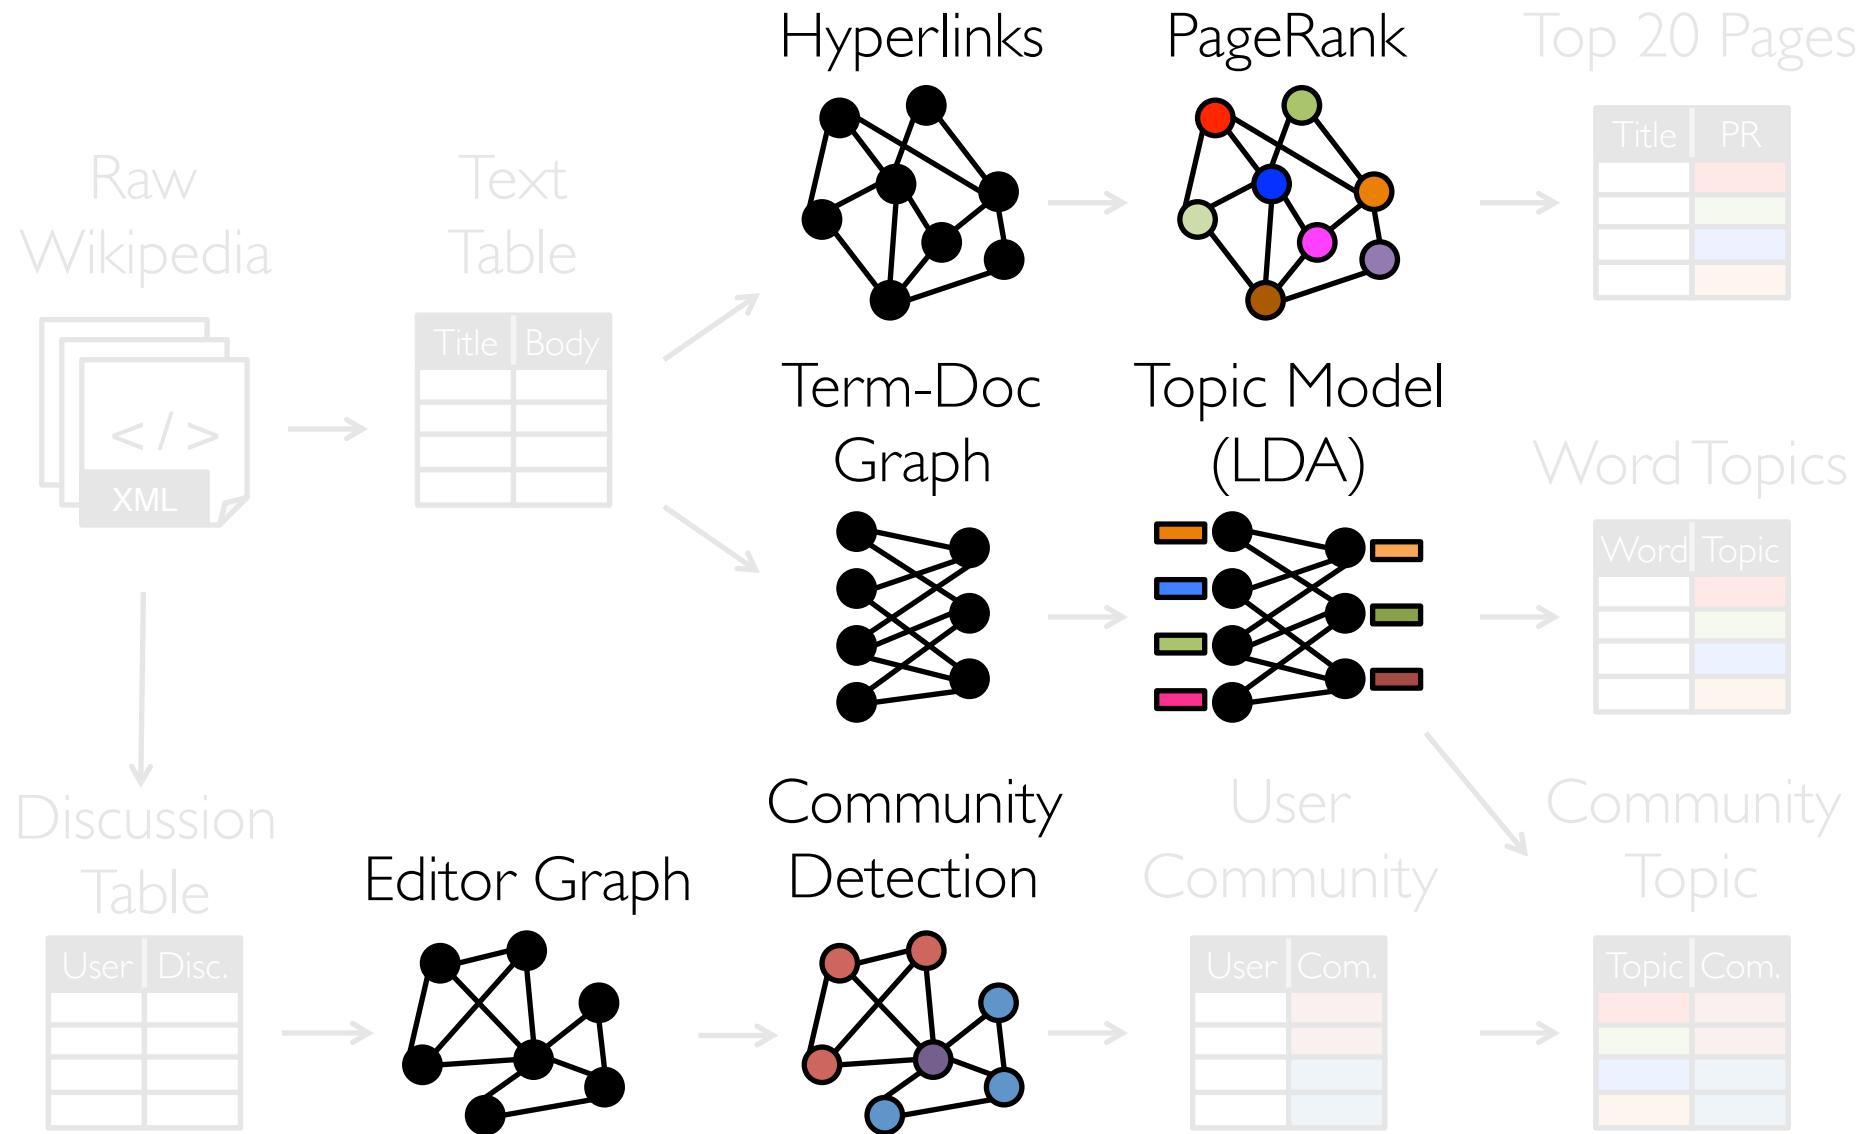
PageRank



Tables

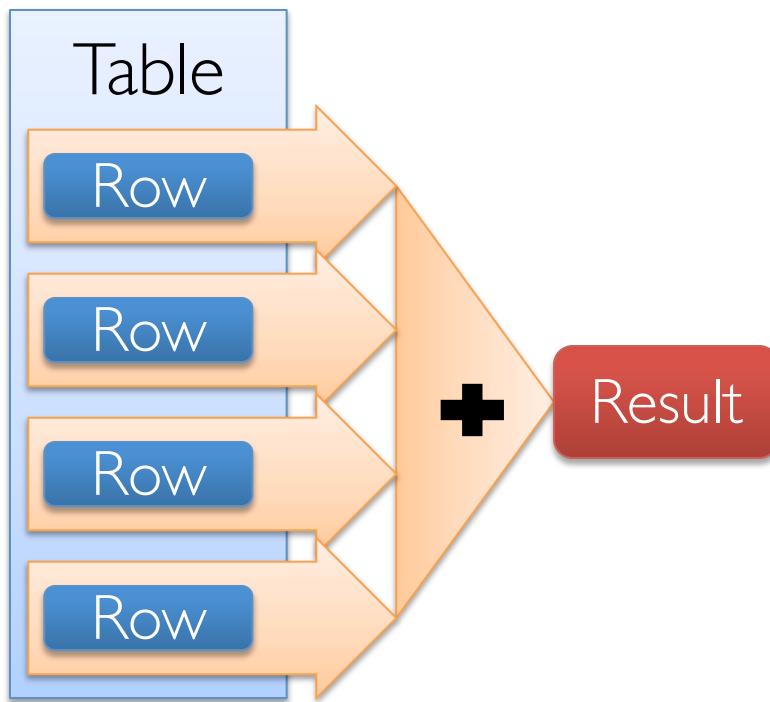


Graphs



Separate Systems to Support Each View

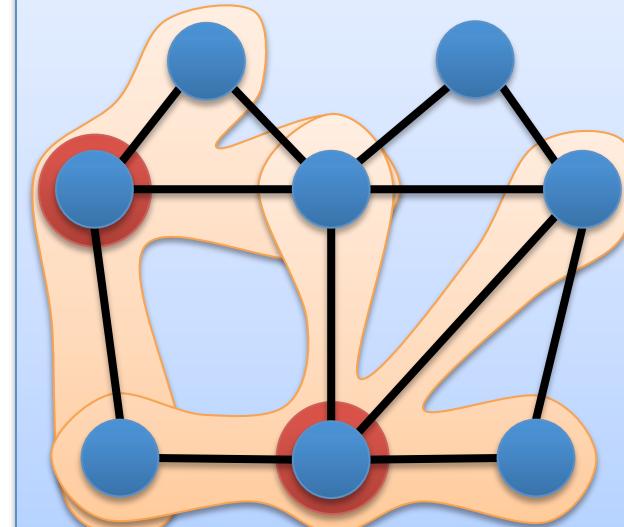
Table View



Graph View



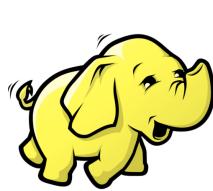
Dependency Graph



*Having separate systems
for each view is
difficult to use and inefficient*

Difficult to Program and Use

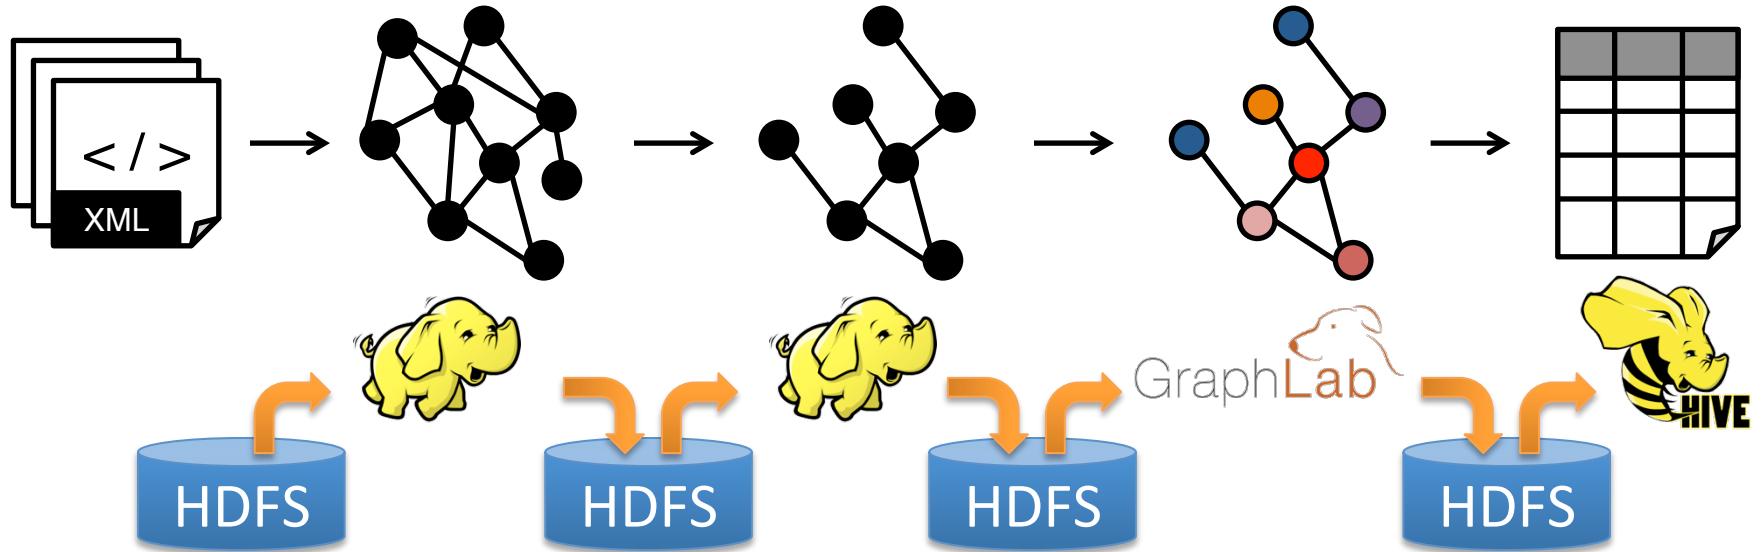
Users must *Learn*, *Deploy*, and *Manage* multiple systems



Leads to brittle and often complex interfaces

Inefficient

Extensive **data movement** and **duplication** across
the network and file system



Limited reuse internal data-structures
across stages

GraphX Solution: Tables and Graphs are *views* of the same *physical* data

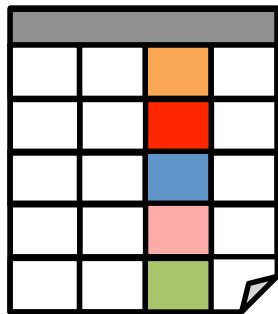
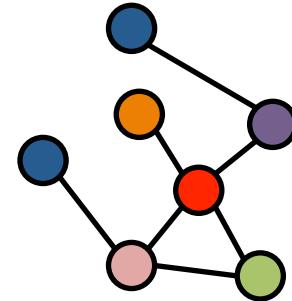
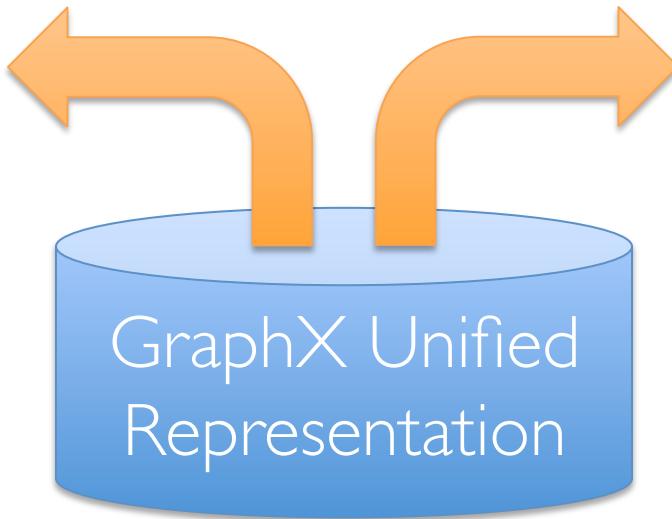


Table View



Graph View

Each view has its own *operators* that
exploit the semantics of the view
to achieve efficient execution

Graphs → Relational Algebra

1. Encode graphs as distributed tables
 2. Express graph computation in relational algebra
 3. Recast graph systems optimizations as:
 1. Distributed join optimization
 2. Incremental materialized maintenance
-

Integrate Graph and
Table data processing
systems.

Achieve performance
parity with specialized
systems.

Distributed Graphs as Distributed Tables

Property Graph

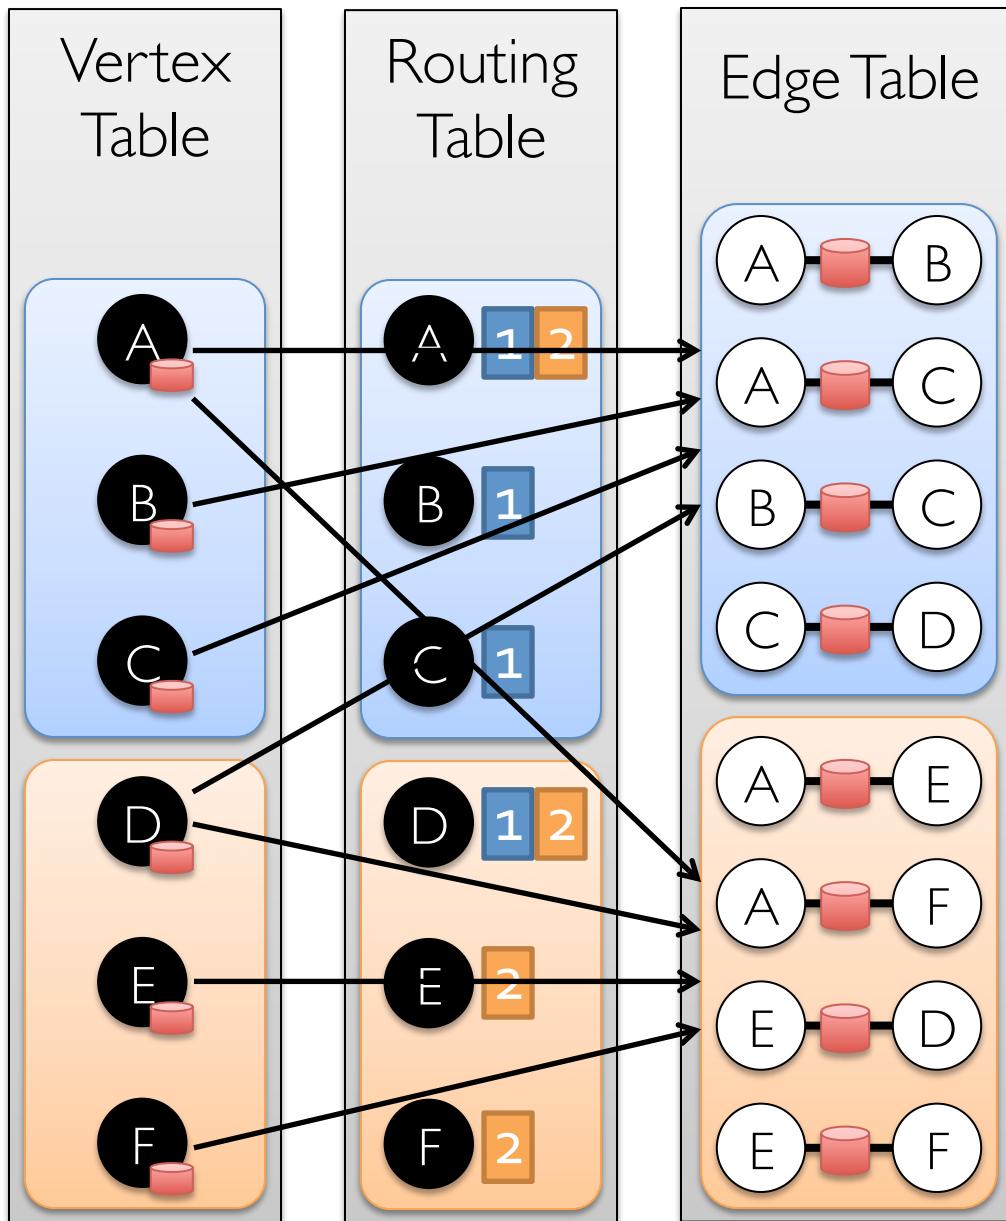
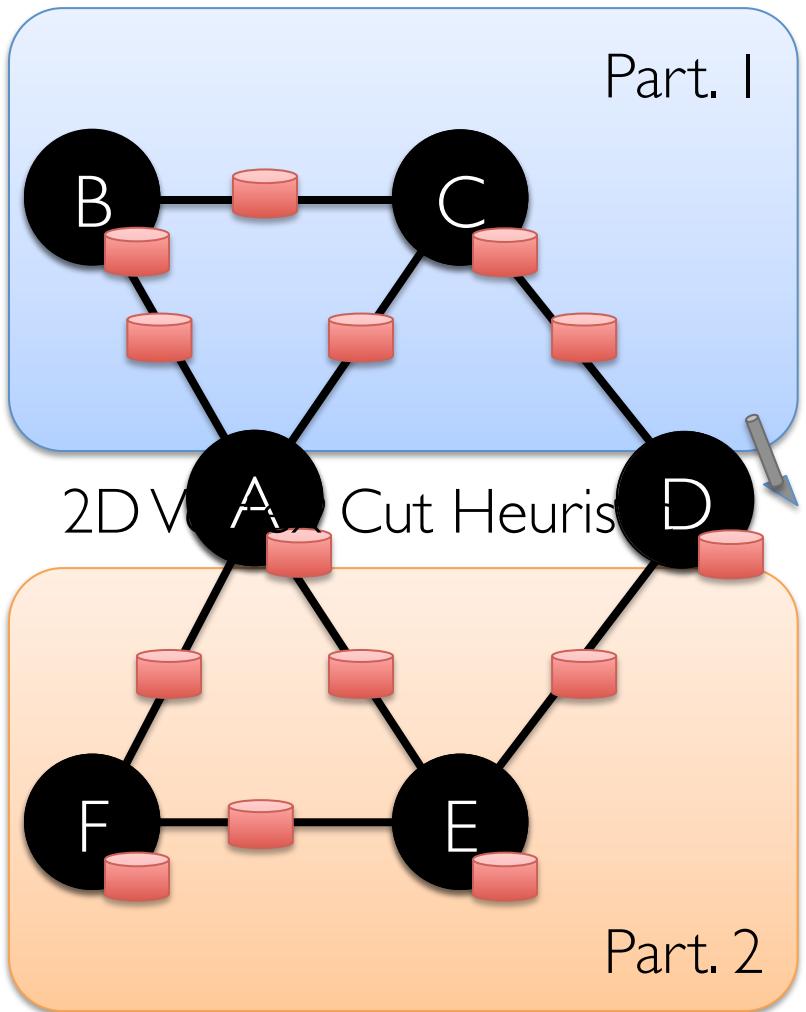


Table Operators

Table operators are inherited from Spark:

map	reduce	sample
filter	count	take
groupBy	fold	first
sort	reduceByKey	partitionBy
union	groupByKey	mapwith
join	cogroup	pipe
leftOuterJoin	cross	save
rightOuterJoin	zip	...

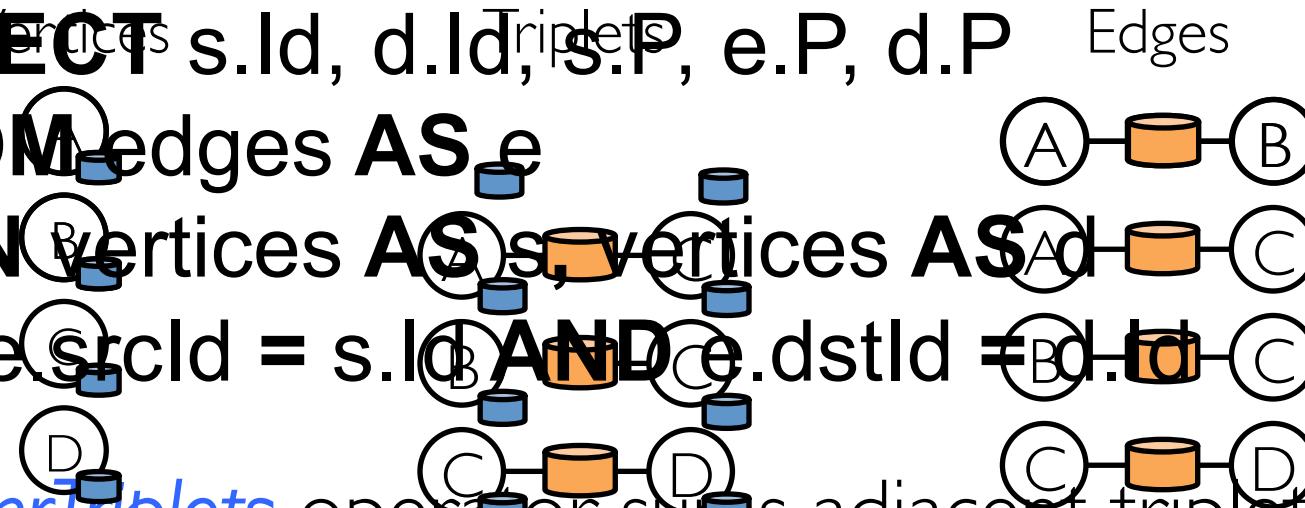
Graph Operators

```
class Graph [ V, E ] {  
    def Graph(vertices: Table[ (Id, V) ],  
              edges: Table[ (Id, Id, E) ])  
        // Table views -----  
        def vertices: Table[ (Id, V) ]  
        def edges: Table[ (Id, Id, E) ]  
        def triplets: Table [ ((Id, V), (Id, V), E) ]  
        // Transformations -----  
        def reverse: Graph[V, E]  
        def subgraph(pV: (Id, V) => Boolean,  
                    pE: Edge[V, E] => Boolean): Graph[V, E]  
        def mapV(m: (Id, V) => T ): Graph[T, E]  
        def mapE(m: Edge[V, E] => T ): Graph[V, T]  
        // Joins -----  
        def joinV(tbl: Table [(Id, T)]): Graph[(V, T), E ]  
        def joinE(tbl: Table [(Id, Id, T)]) : Graph[V, (E, T)]  
        // Computation -----  
        def mrTriplets(mapF: (Edge[V, E]) => List[(Id, T)],  
                      reduceF: (T, T) => T): Graph[T, E]  
}
```

Triplets Join Vertices and Edges

The *triplets* operator joins vertices and edges:

```
SELECT s.Id, d.Id, s.P, e.P, d.P  
FROM Edges AS e  
JOIN Vertices AS s, Vertices AS d  
ON e.srcId = s.Id AND e.dstId = d.Id
```



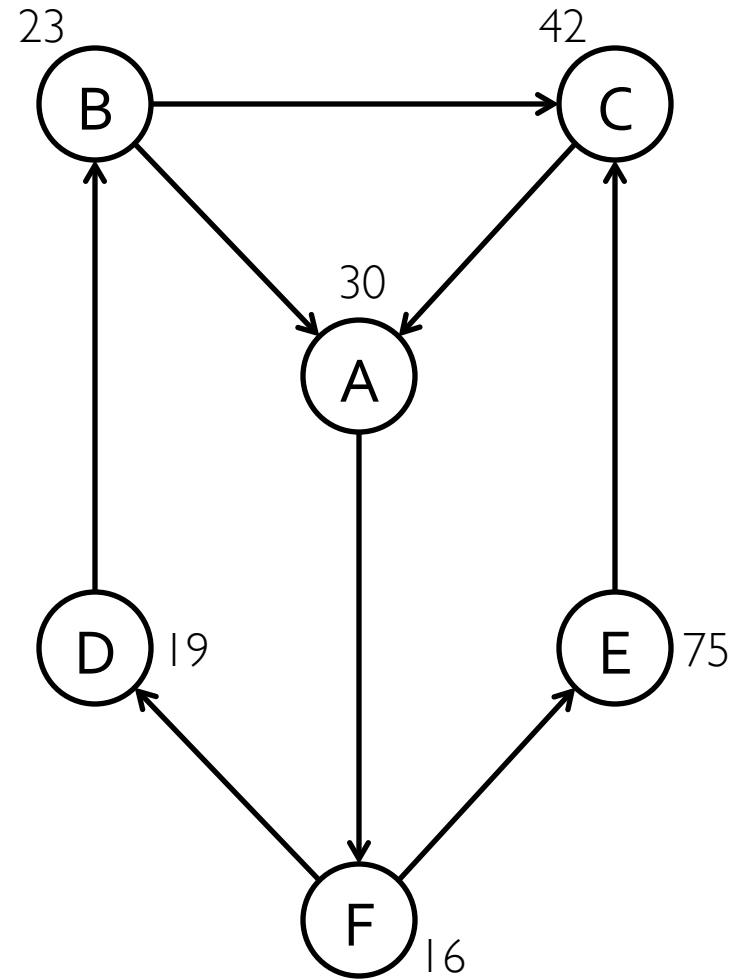
The *mrTriplets* operator sums adjacent triplets.

```
SELECT t.dstId, reduce( map(t) ) AS sum  
FROM triplets AS t GROUPBY t.dstId
```

Example: Oldest Follower

Calculate the number of older followers for each user?

```
val olderFollowerAge = graph
  .mrTriplets(
    e => // Map
      if(e.src.age < e.dst.age) {
        (e.srcId, 1)
      } else { Empty }
    ,
    (a,b) => a + b // Reduce
  )
  .vertices
```



We express *enhanced* Pregel and GraphLab abstractions using the GraphX operators in less than 50 lines of code!

Enhanced Pregel in GraphX

```
pregelPR(i, messageSum) :  
    // Receive all the messages  
    total = 0  
    foreach( msg in messageList ) :  
        total = total + msg  
  
    // Update the rank of this vertex  
    R[i] = 0.15 + total  
combineMsg(a, b) :  
    // Compute sum of two messages  
    sendMsg(i → newR[i], R[i], E[i, j]) :  
    foreach(j ← out_neighbors[i]) :  
        // Compute single message  
        Send msg(R[i]/E[i, j]) to vertex
```

Require Message Combiners

Remove Message Computation from the Vertex Program

PageRank in GraphX

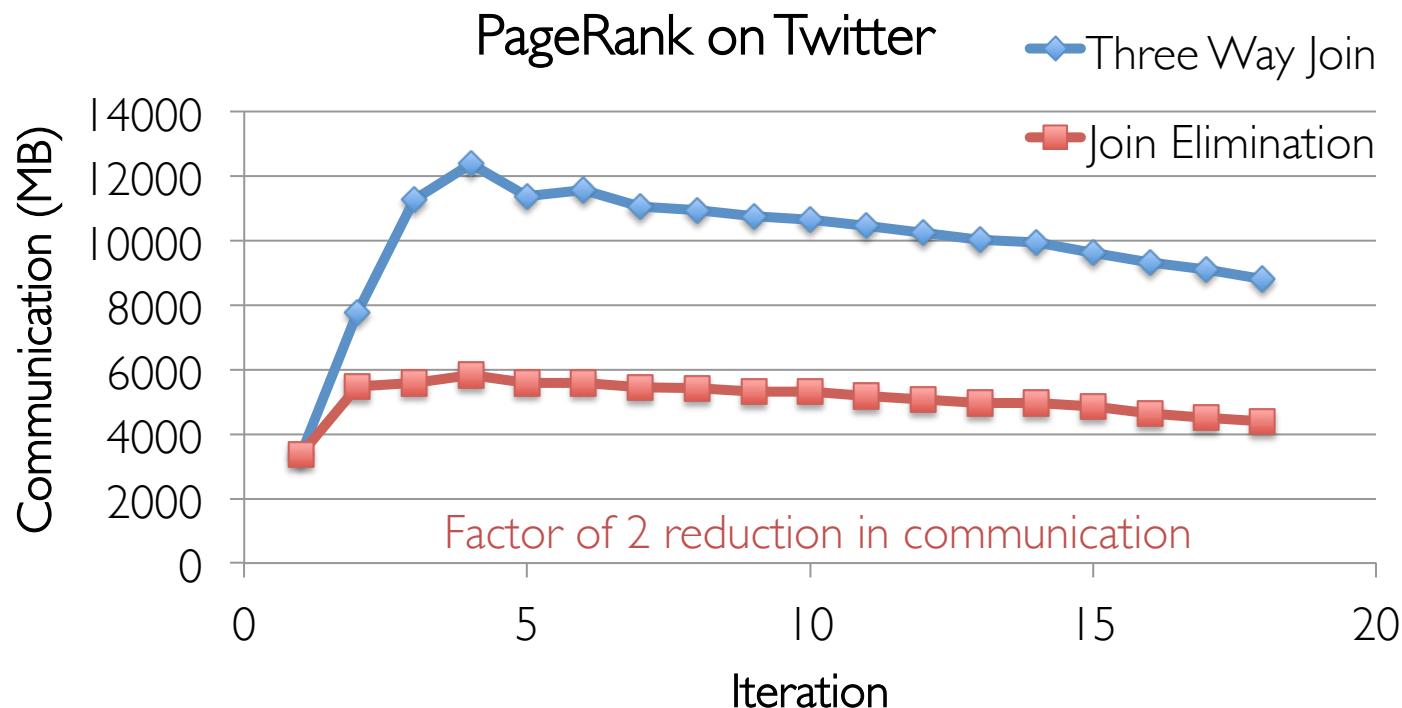
```
// Load and initialize the graph
val graph = GraphBuilder.text("hdfs://web.txt")
val prGraph = graph.joinVertices(graph.outDegrees)

// Implement and Run PageRank
val pageRank =
  prGraph.pregel(initialMessage = 0.0, iter = 10) (
    (oldV, msgSum) => 0.15 + 0.85 * msgSum,
    triplet => triplet.src.pr / triplet.src.deg,
    (msgA, msgB) => msgA + msgB)
```

Join Elimination

Identify and bypass joins for unused triplet fields

```
sendMsg(i→j, R[i], R[j], E[i,j]):  
    // Compute single message  
    return msg(R[i]/E[i,j])
```



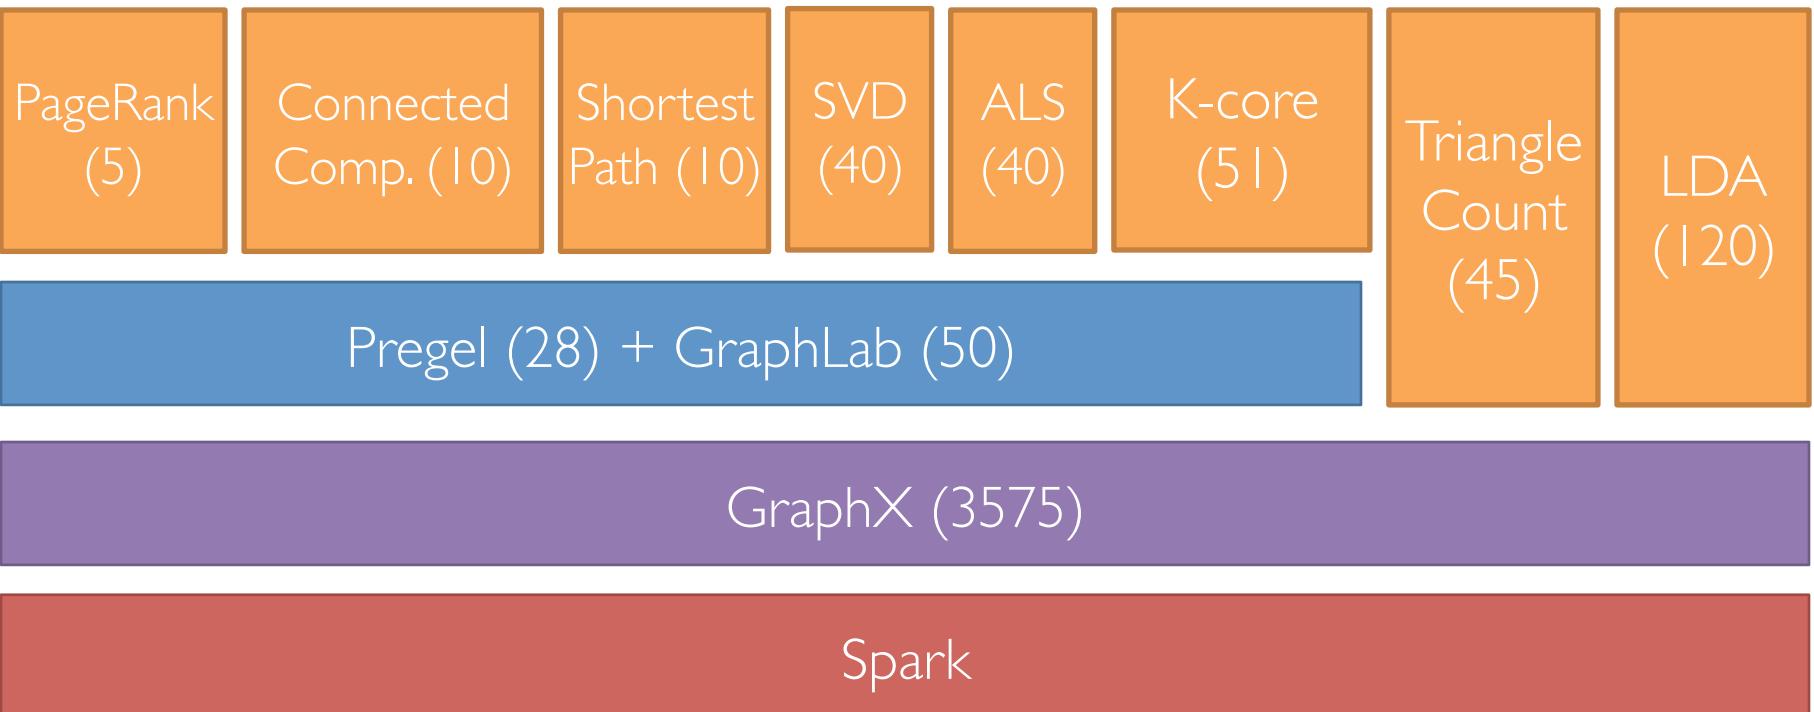
We express the Pregel and GraphLab *like* abstractions using the GraphX operators in less than 50 lines of code!

By composing these operators we can construct entire graph-analytics pipelines.

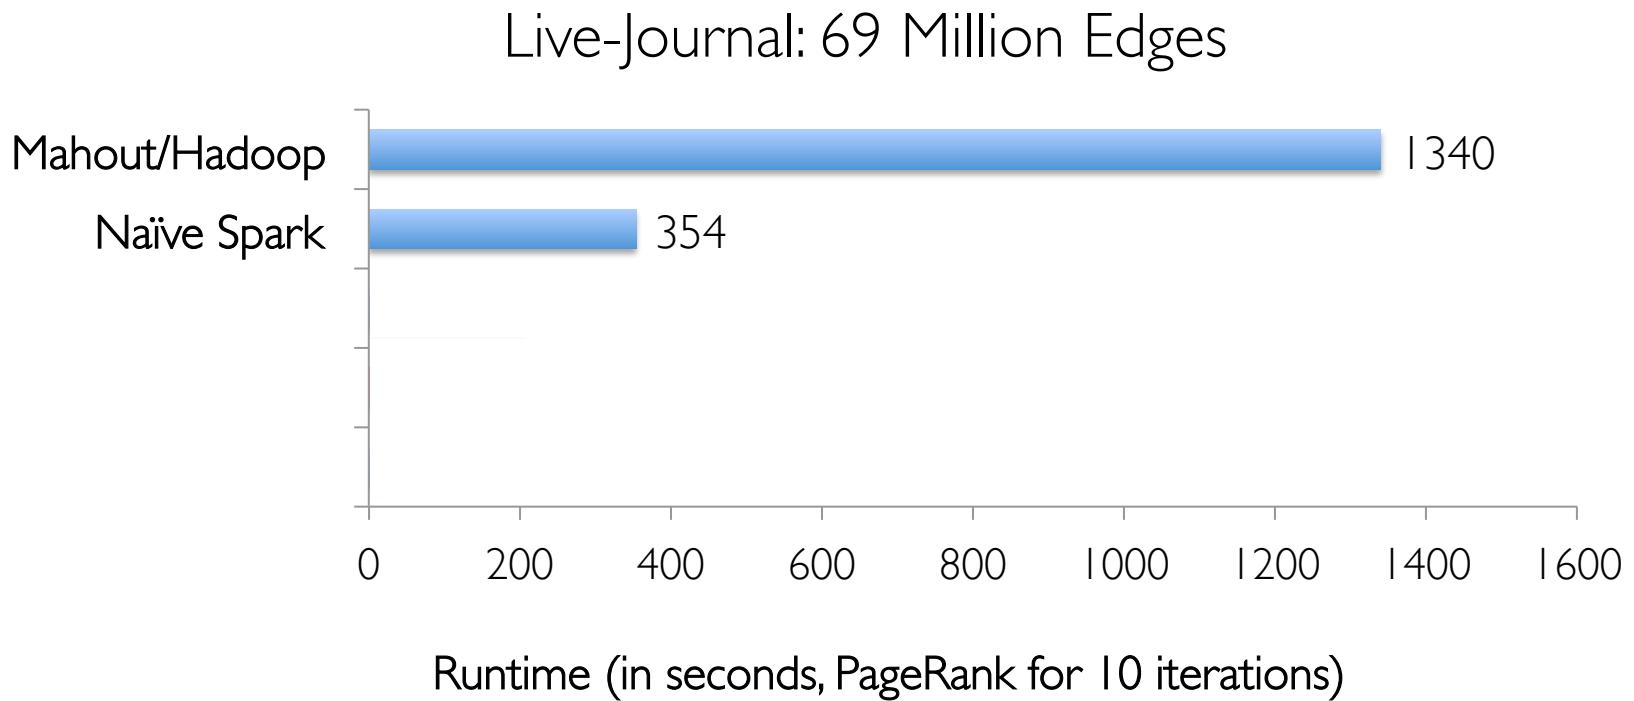
Example Analytics Pipeline

```
// Load raw data tables  
val verts = sc.textFile("hdfs://users.txt").map(parserV)  
val edges = sc.textFile("hdfs://follow.txt").map(parserE)  
// Build the graph from tables and restrict to recent links  
val graph = new Graph(verts, edges)  
val recent = graph.subgraph(edge => edge.date > LAST_MONTH)  
// Run PageRank Algorithm  
val pr = graph.PageRank(tol = 1.0e-5)  
// Extract and print the top 25 users  
val topUsers = verts.join(pr).top(25).collect  
topUsers.foreach(u => println(u.name + '\t' + u.pr))
```

The GraphX Stack (Lines of Code)

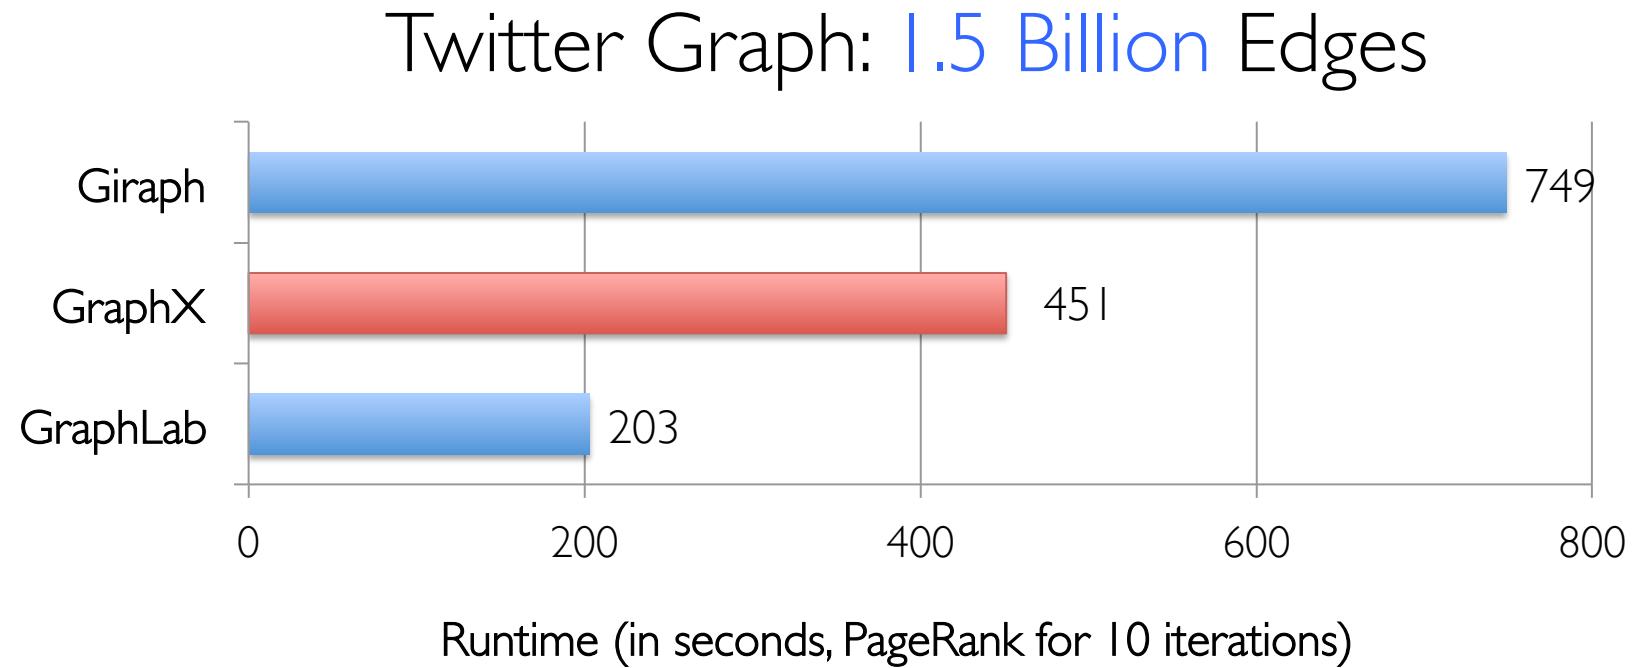


Performance Comparisons



GraphX is roughly *3x slower* than GraphLab

GraphX scales to larger graphs



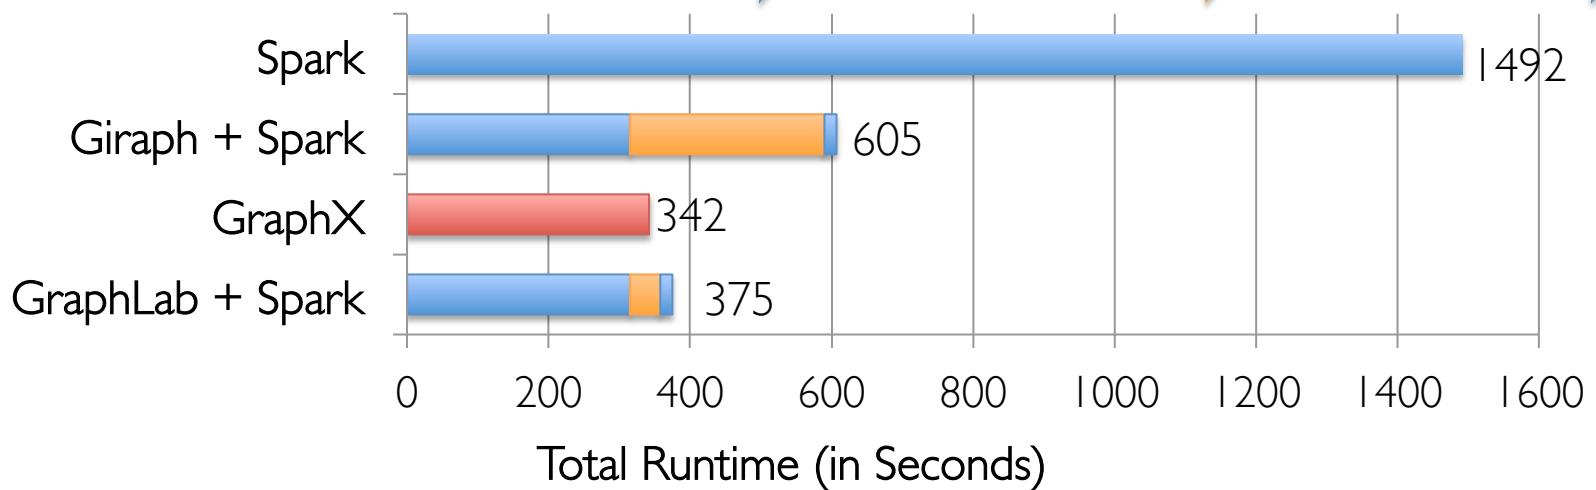
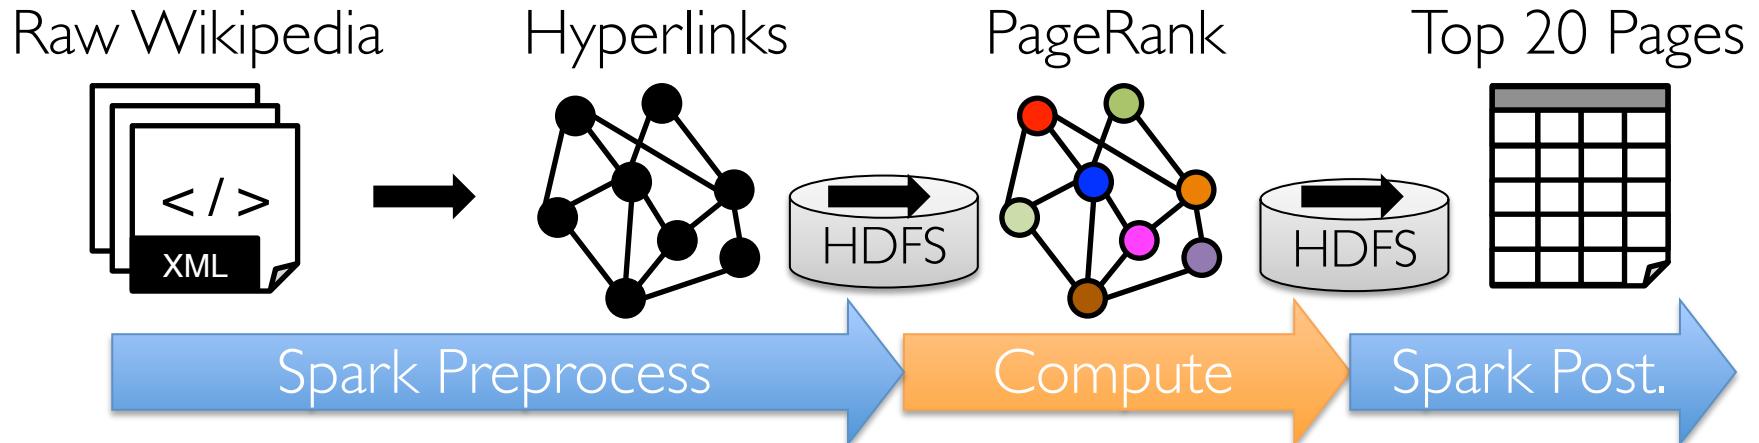
GraphX is roughly *2x slower* than GraphLab

- » Scala + Java overhead: Lambdas, GC time, ...
- » No shared memory parallelism: *2x increase* in comm.

PageRank is just one stage....

What about a pipeline?

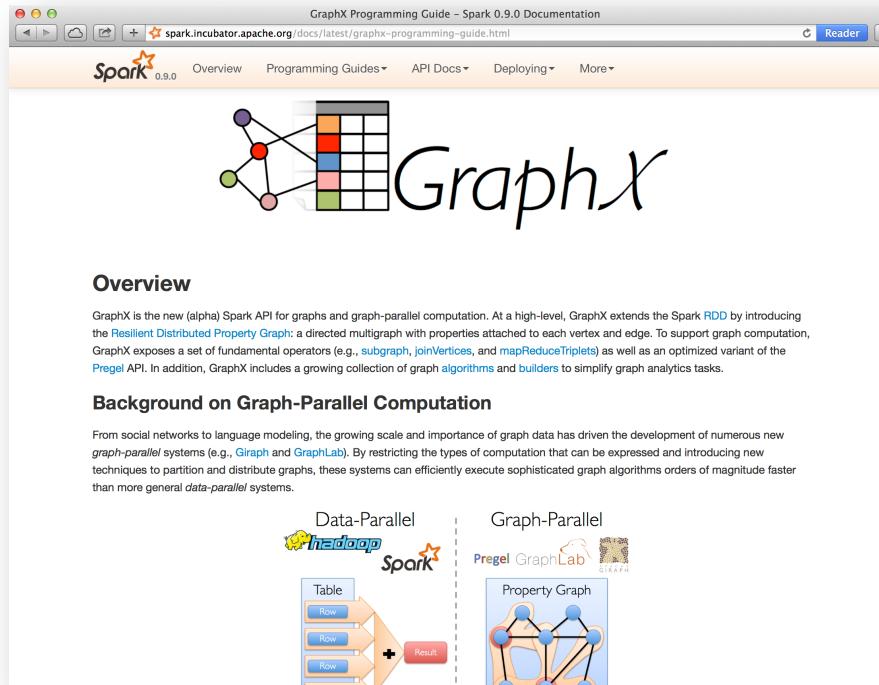
A Small Pipeline in GraphX



Timed end-to-end GraphX is *faster* than GraphLab

Status

Part of Apache Spark



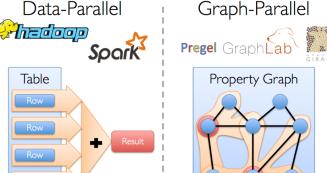
The screenshot shows a web browser displaying the "GraphX Programming Guide – Spark 0.9.0 Documentation" at spark.incubator.apache.org/docs/latest/graphx-programming-guide.html. The page features a large logo on the left with the word "GraphX" next to a grid of colored nodes and edges. The main content area includes sections like "Overview" and "Background on Graph-Parallel Computation", along with a diagram comparing Data-Parallel and Graph-Parallel processing.

Overview

GraphX is the new (alpha) Spark API for graphs and graph-parallel computation. At a high-level, GraphX extends the Spark [RDD](#) by introducing the [Resilient Distributed Property Graph](#): a directed multigraph with properties attached to each vertex and edge. To support graph computation, GraphX exposes a set of fundamental operators (e.g., `subgraph`, `joinVertices`, and `mapReduceTriplets`) as well as an optimized variant of the [Pregel API](#). In addition, GraphX includes a growing collection of graph [algorithms](#) and [builders](#) to simplify graph analytics tasks.

Background on Graph-Parallel Computation

From social networks to language modeling, the growing scale and importance of graph data has driven the development of numerous new *graph-parallel* systems (e.g., [Giraph](#) and [GraphLab](#)). By restricting the types of computation that can be expressed and introducing new techniques to partition and distribute graphs, these systems can efficiently execute sophisticated graph algorithms orders of magnitude faster than more general *data-parallel* systems.



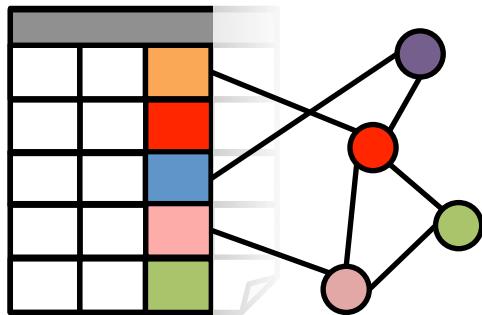
The diagram illustrates the difference between Data-Parallel and Graph-Parallel processing. On the left, under "Data-Parallel", there is a "Table" icon with four "Row" icons pointing to a central "+" sign, which then points to a "Result" icon. This represents a row-wise operation on a table. On the right, under "Graph-Parallel", there is a "Property Graph" icon showing a network of nodes and edges, with the "Pregel" and "GraphLab" logos above it. This represents a graph-based computation where data is processed based on its connections to other nodes.

In production at several large technology companies

GraphX: Unified Analytics

New API

*Blurs the distinction between
Tables and Graphs*



New System

*Combines Data-Parallel
Graph-Parallel Systems*



Enabling users to **easily** and **efficiently**
express the entire graph analytics pipeline

A Case for Algebra in Graphs

A standard algebra is essential for graph systems:

- e.g.: SQL → proliferation of relational system

By embedding graphs in *relational algebra*:

- Integration with tables and preprocessing
- Leverage advances in relational systems
- Graph opt. recast to relational systems opt.

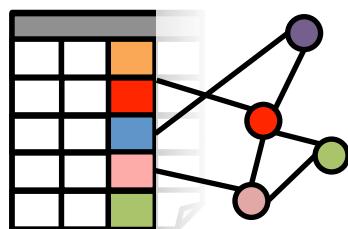
Conclusions

Composable domain specific views and operators

Single system that efficiently spans the pipeline

Graphs through the lens of database systems

- » Graph-Parallel Pattern → Triplet joins in relational alg.
- » Graph Systems → Distributed join optimizations



Joseph E. Gonzalez
jegonzal@eecs.berkeley.edu
<http://tinyurl.com/ampgraphx>

Thanks!

<http://amplab.cs.berkeley.edu/projects/graphx/>

ankurd@eecs.berkeley.edu

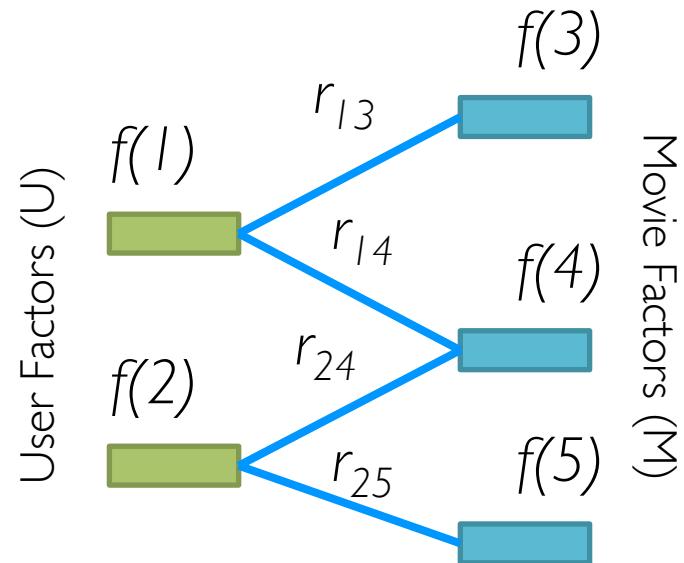
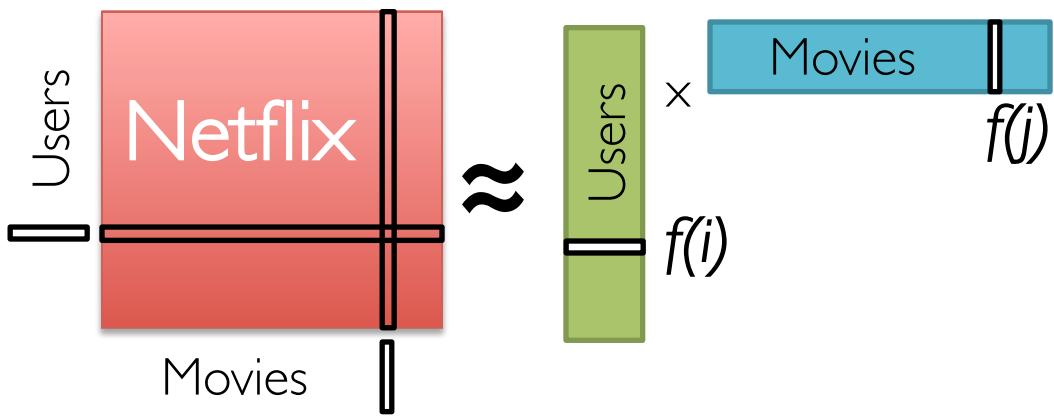
crankshaw@eecs.berkeley.edu

rxin@eecs.berkeley.edu

jegonzal@eecs.berkeley.edu

Recommending Products

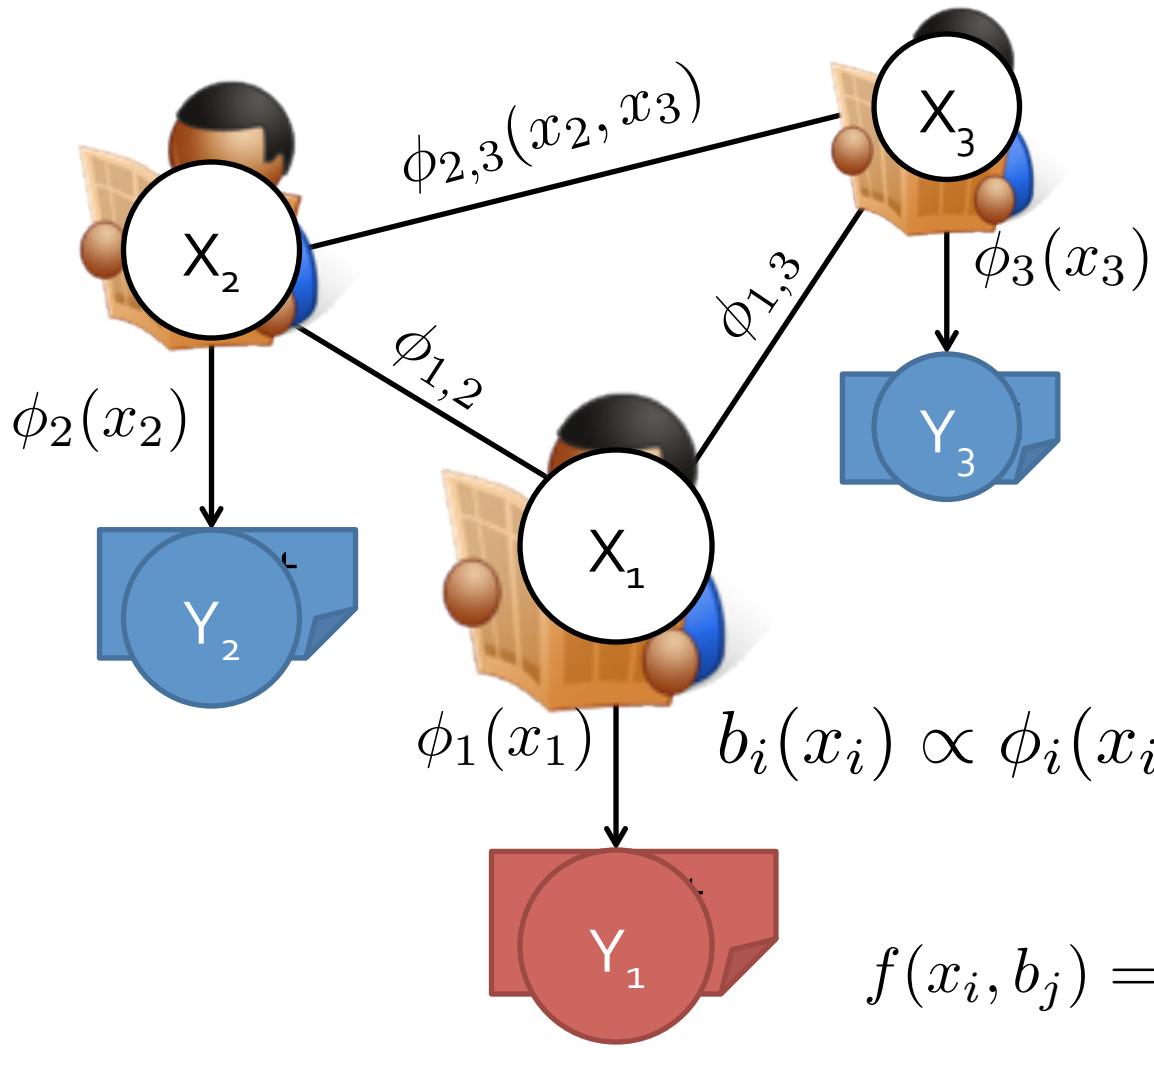
Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda \|w\|_2^2$$

Mean Field Algorithm



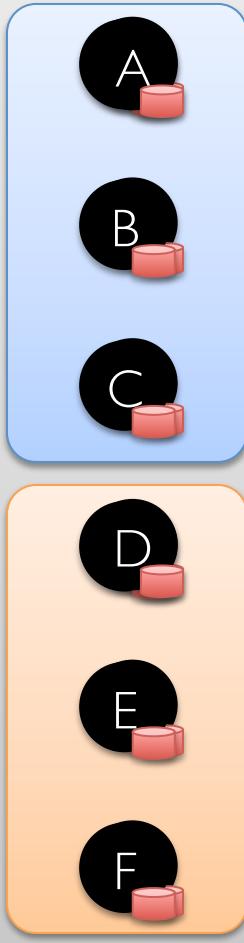
Sum over
Neighbors

$$f(x_i, b_j) = \sum_{x_j} b_j(x_j) \log \phi_{i,j}(x_i, x_j)$$
$$b_i(x_i) \propto \phi_i(x_i) \exp \left(\sum_{j \in N_i} f(x_i, b_j) \right)$$

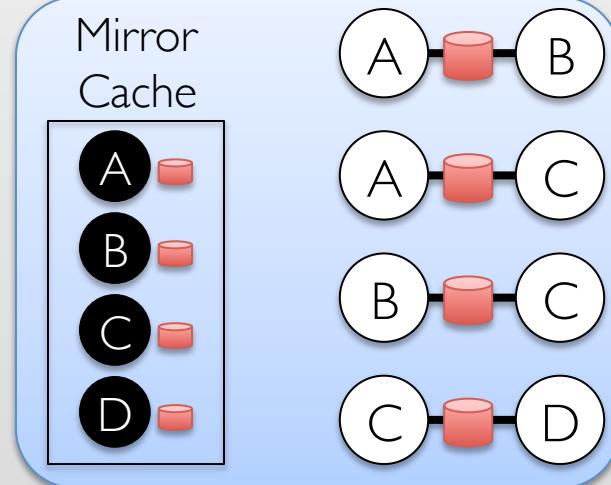
GraphX System Design

Caching for Iterative mrTriplets

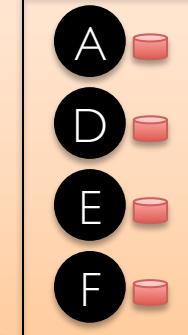
Vertex
Table
(RDD)



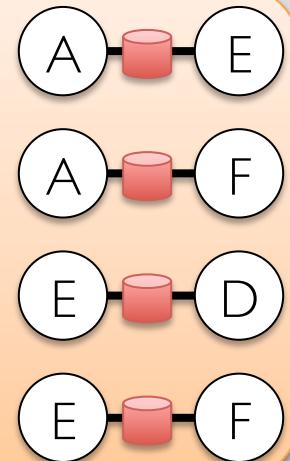
Edge Table
(RDD)



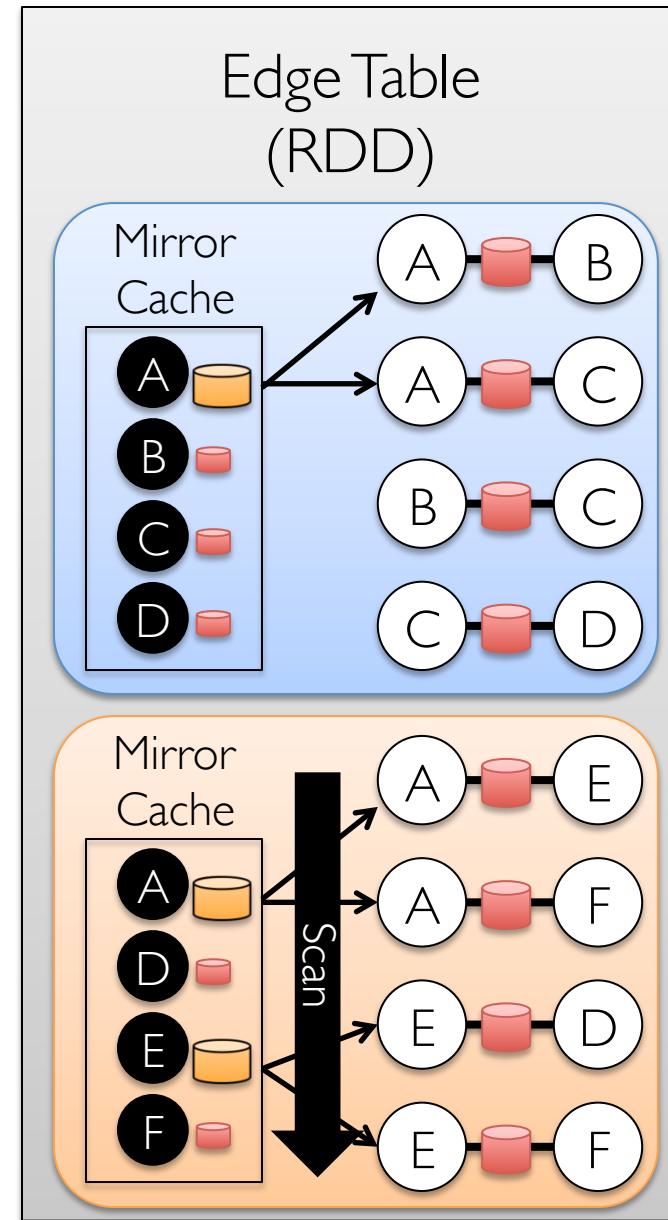
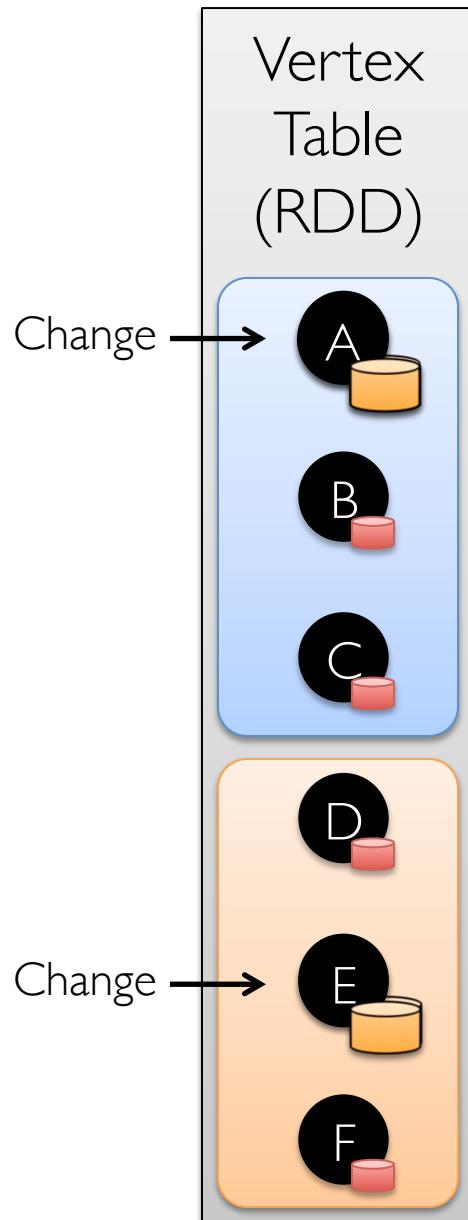
Mirror
Cache



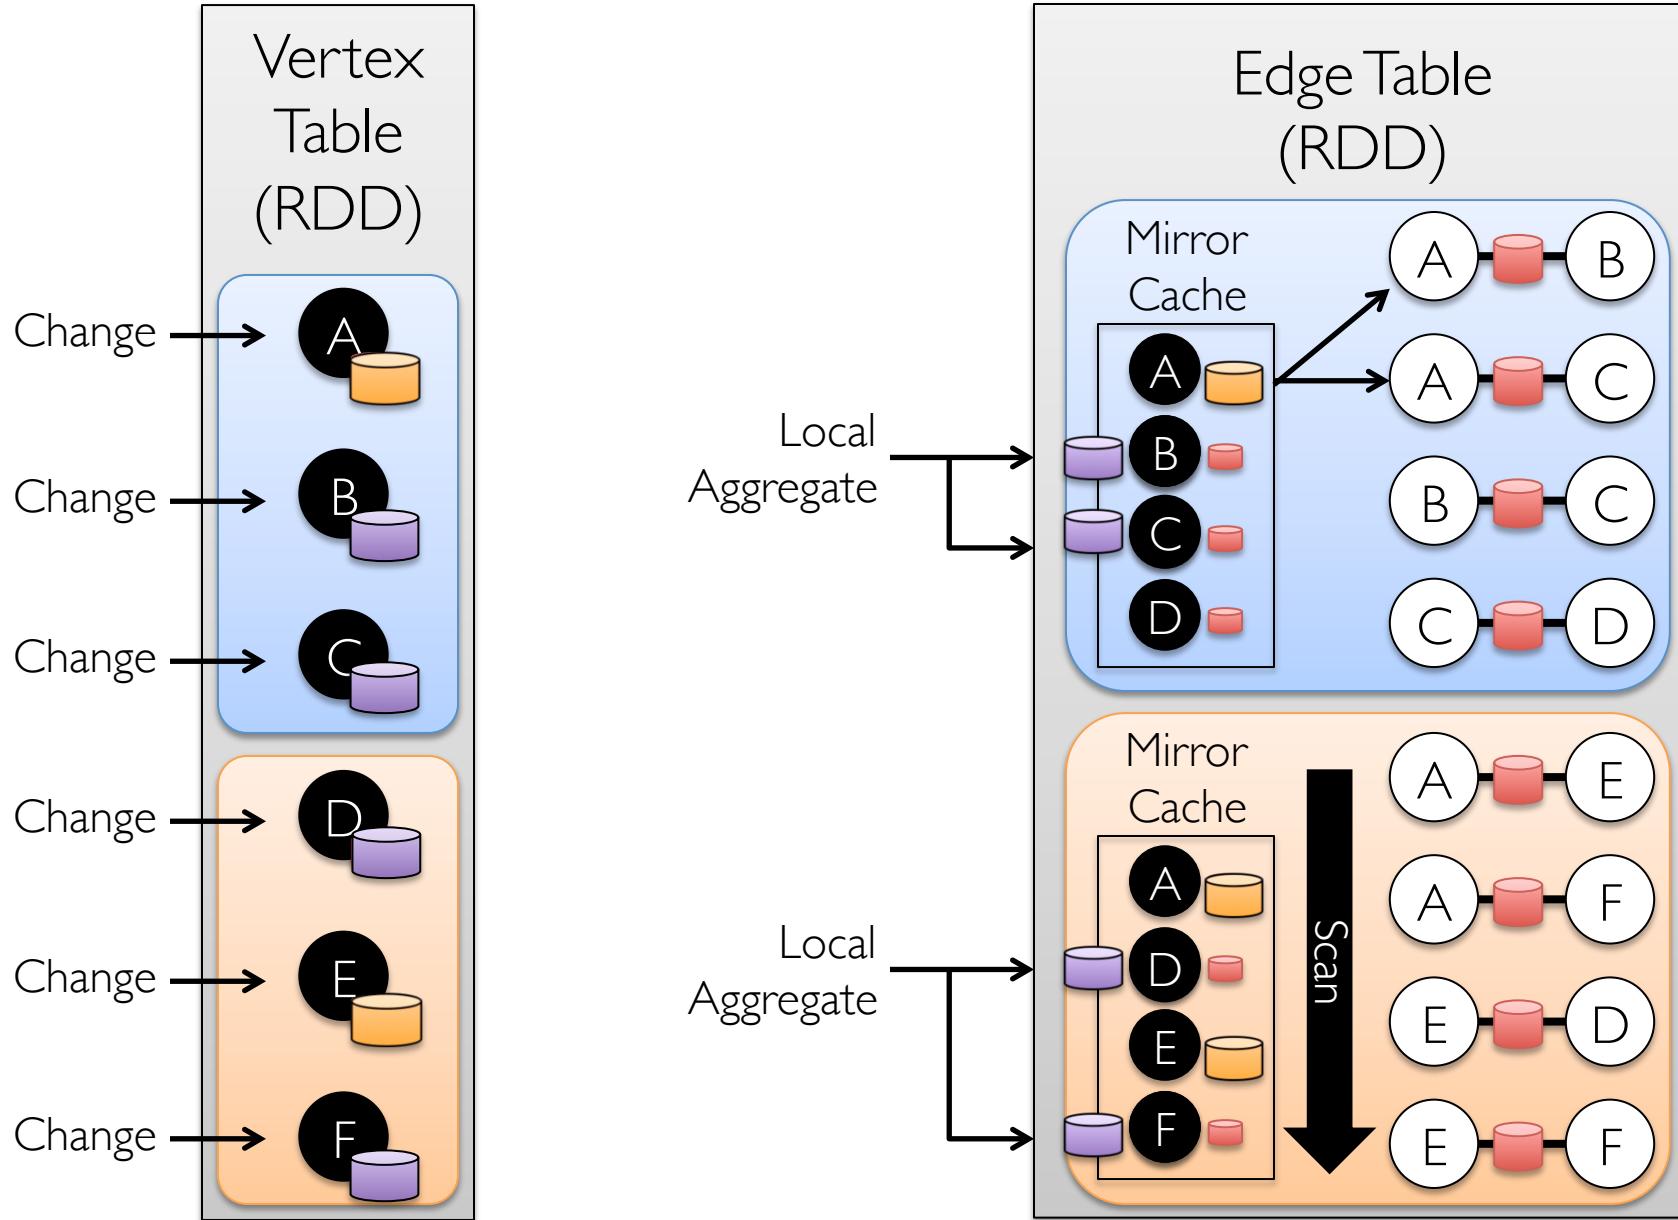
Mirror
Cache



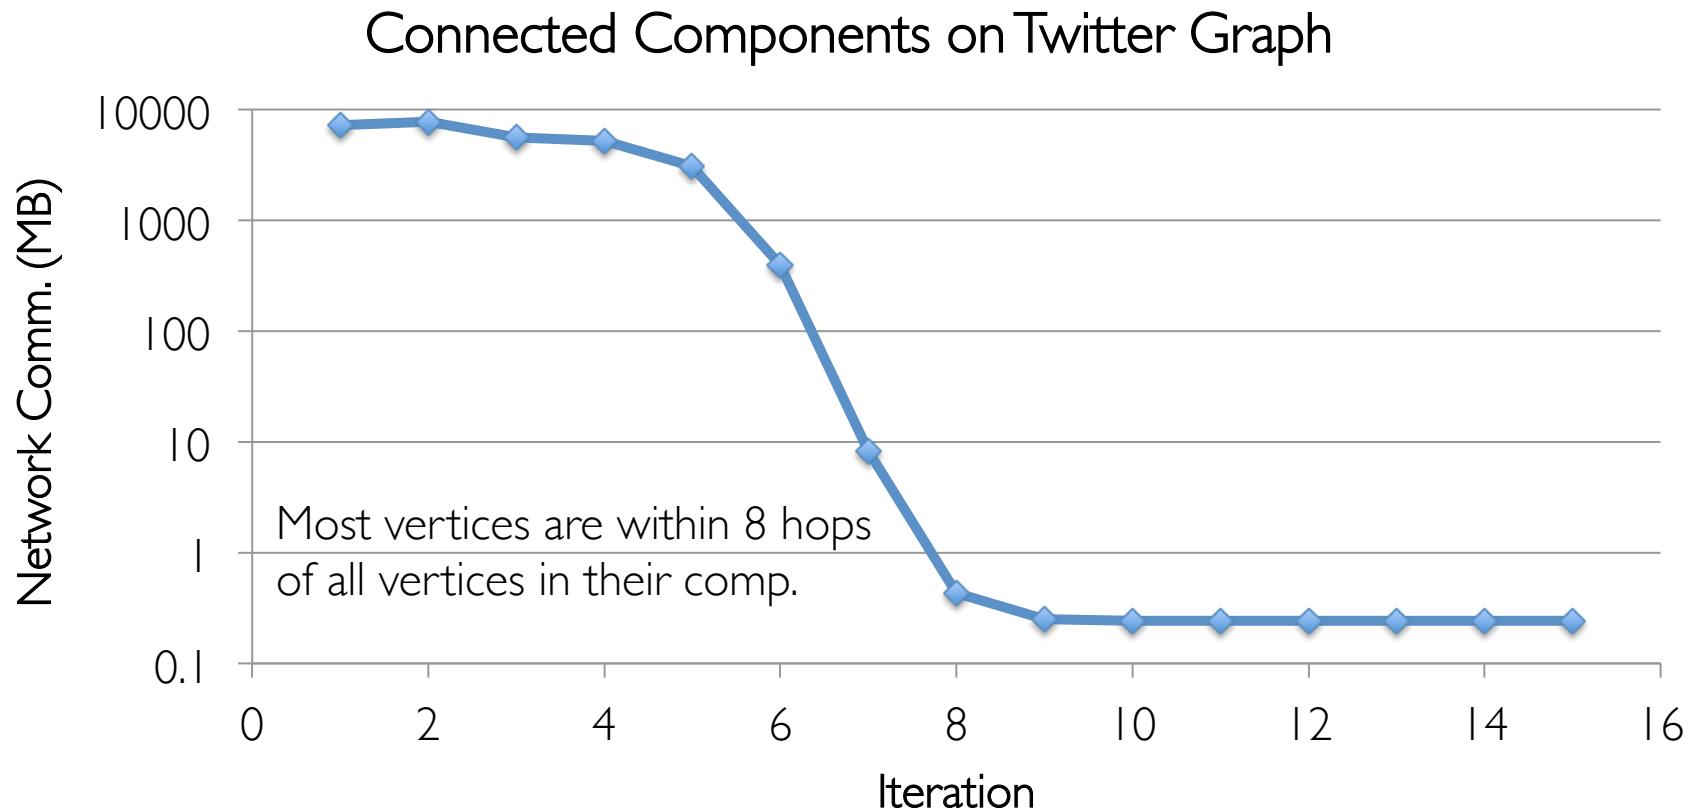
Incremental Updates for Iterative mrTriplets



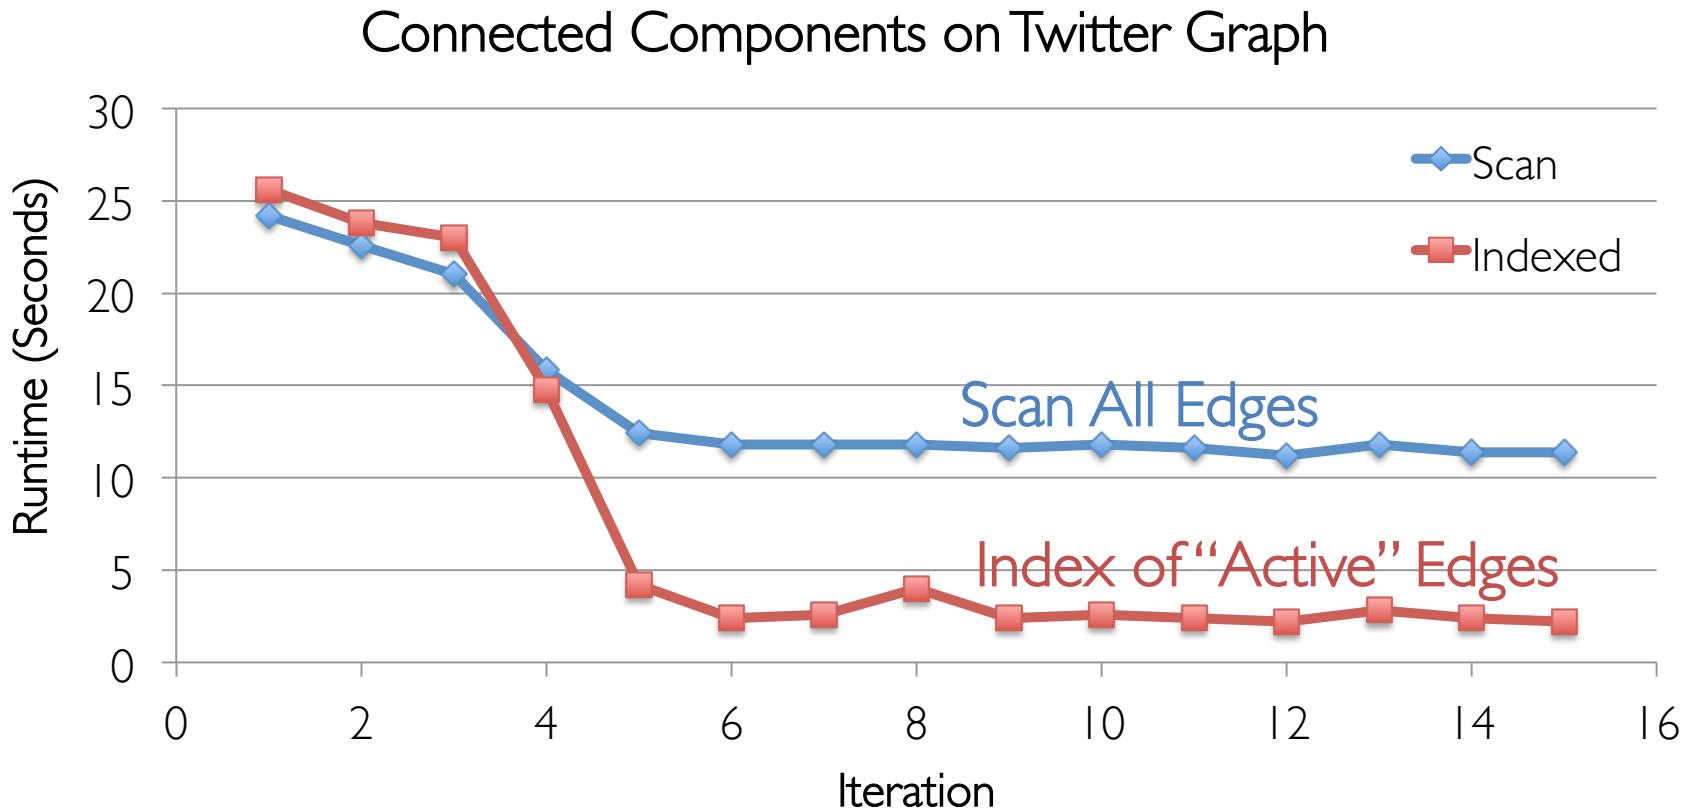
Aggregation for Iterative mrTriplets



Reduction in Communication Due to Cached Updates



Benefit of Indexing Active Edges



Additional Query Optimizations

Indexing and Bitmaps:

- » To accelerate joins across graphs
- » To efficiently construct sub-graphs

Substantial Index and Data Reuse:

- » Reuse routing tables across graphs and sub-graphs
- » Reuse edge adjacency information and indices