

Concurrency Control for Scalable Bayesian Inference

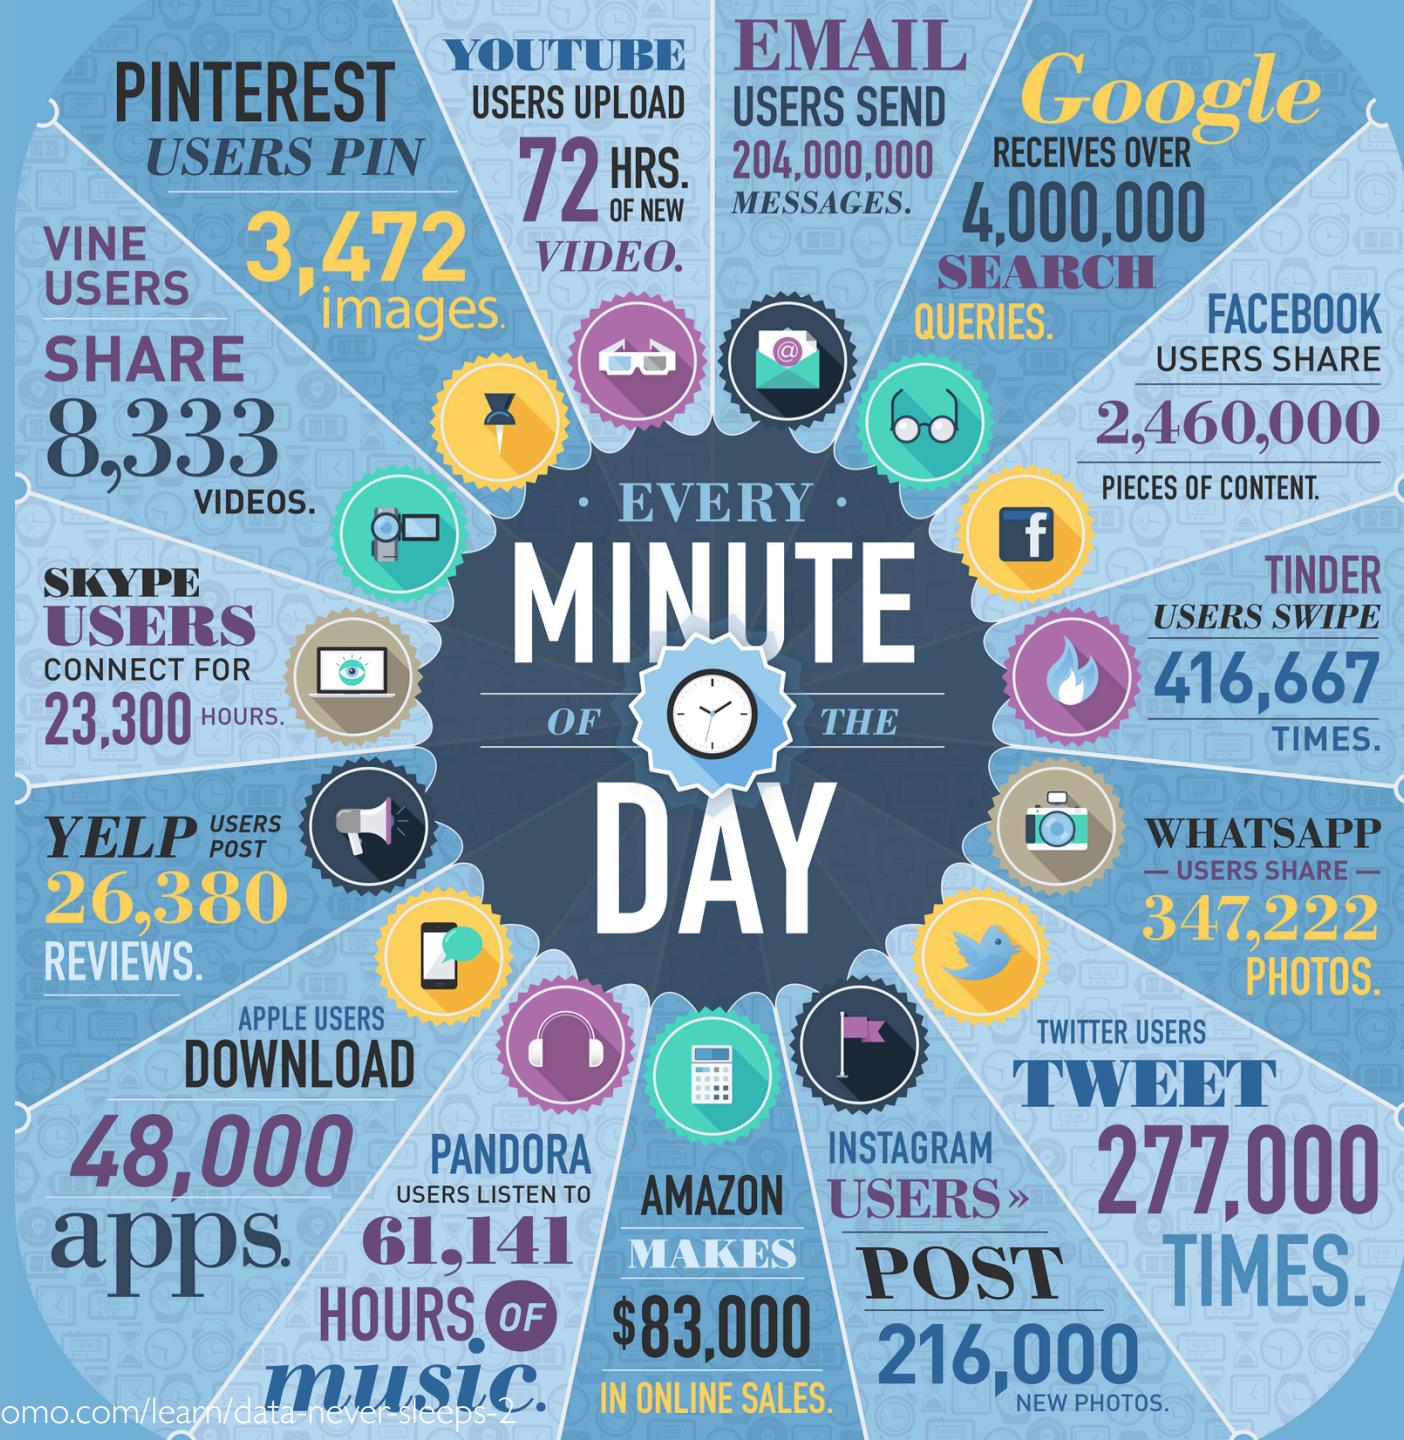
Joseph E. Gonzalez
Postdoc, UC Berkeley AMPLab
Co-founder, GraphLab Inc.
jegonzal@eecs.berkeley.edu

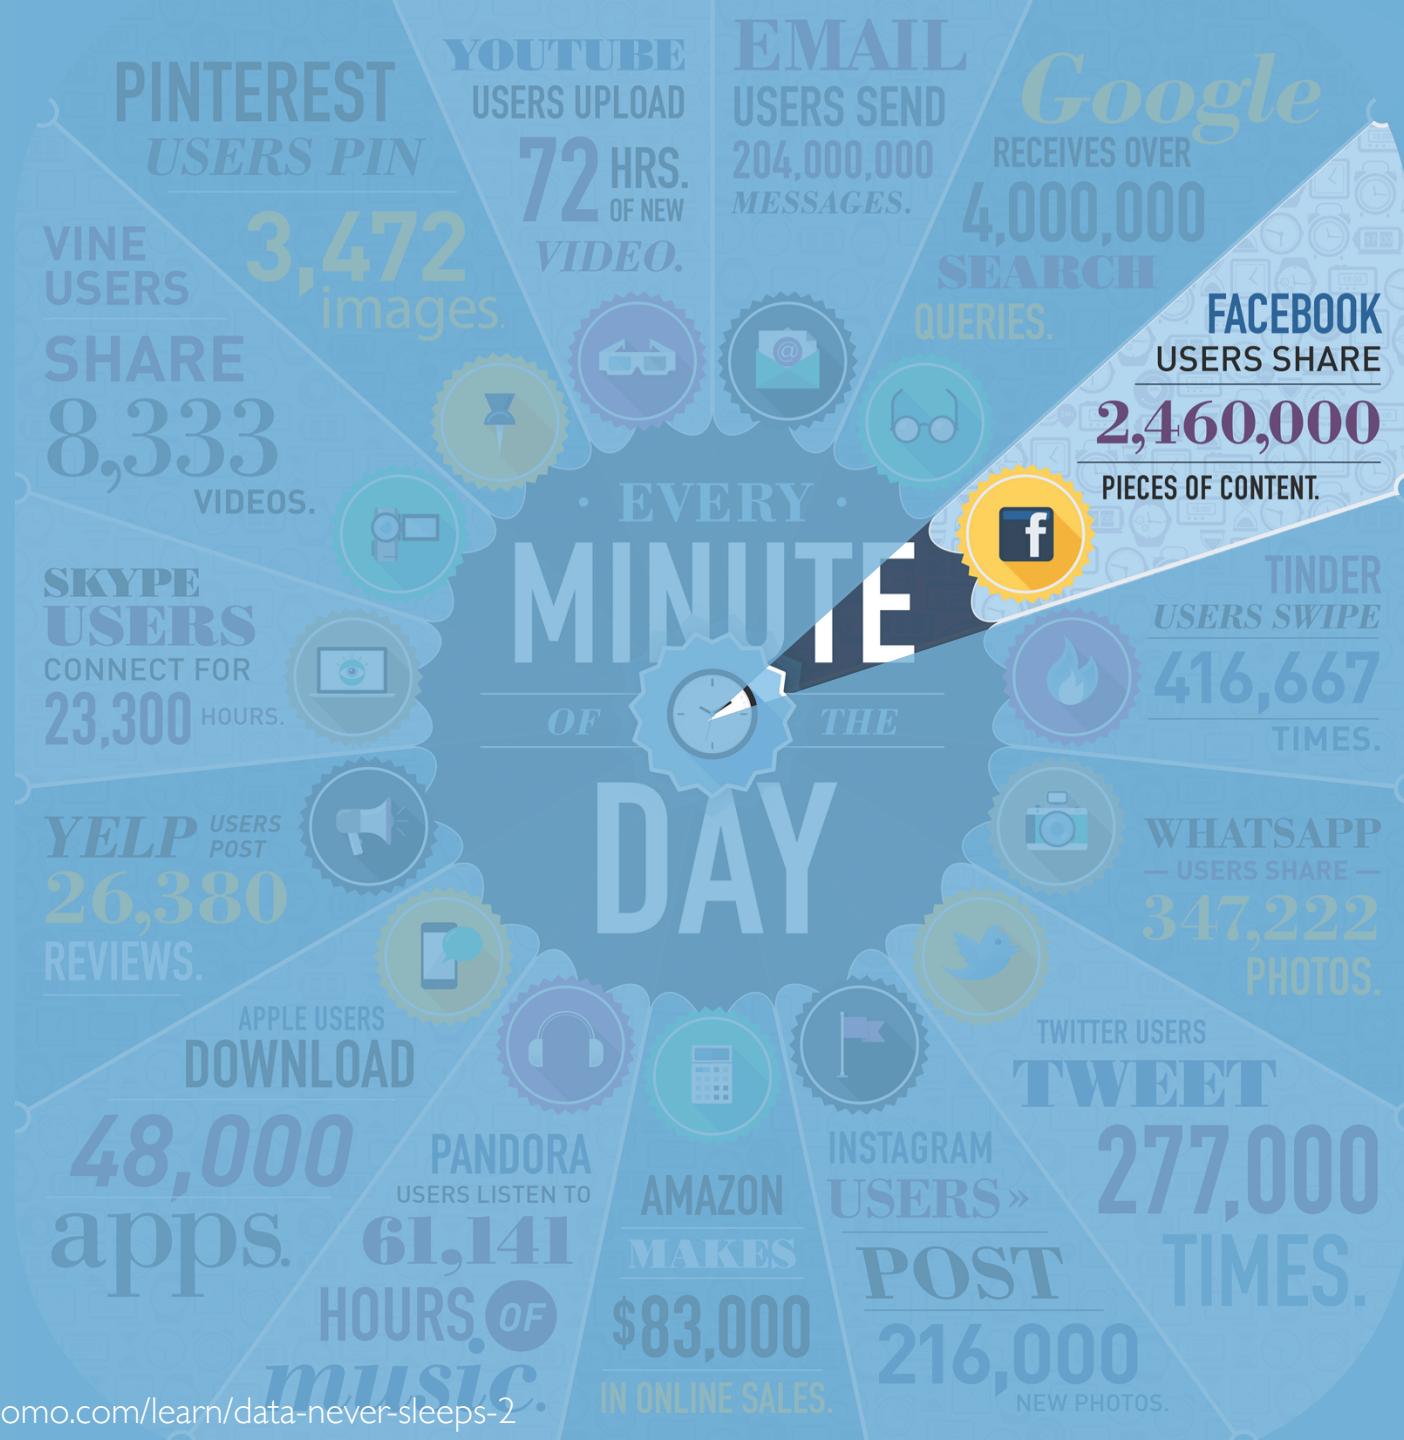
ISBA' 2014

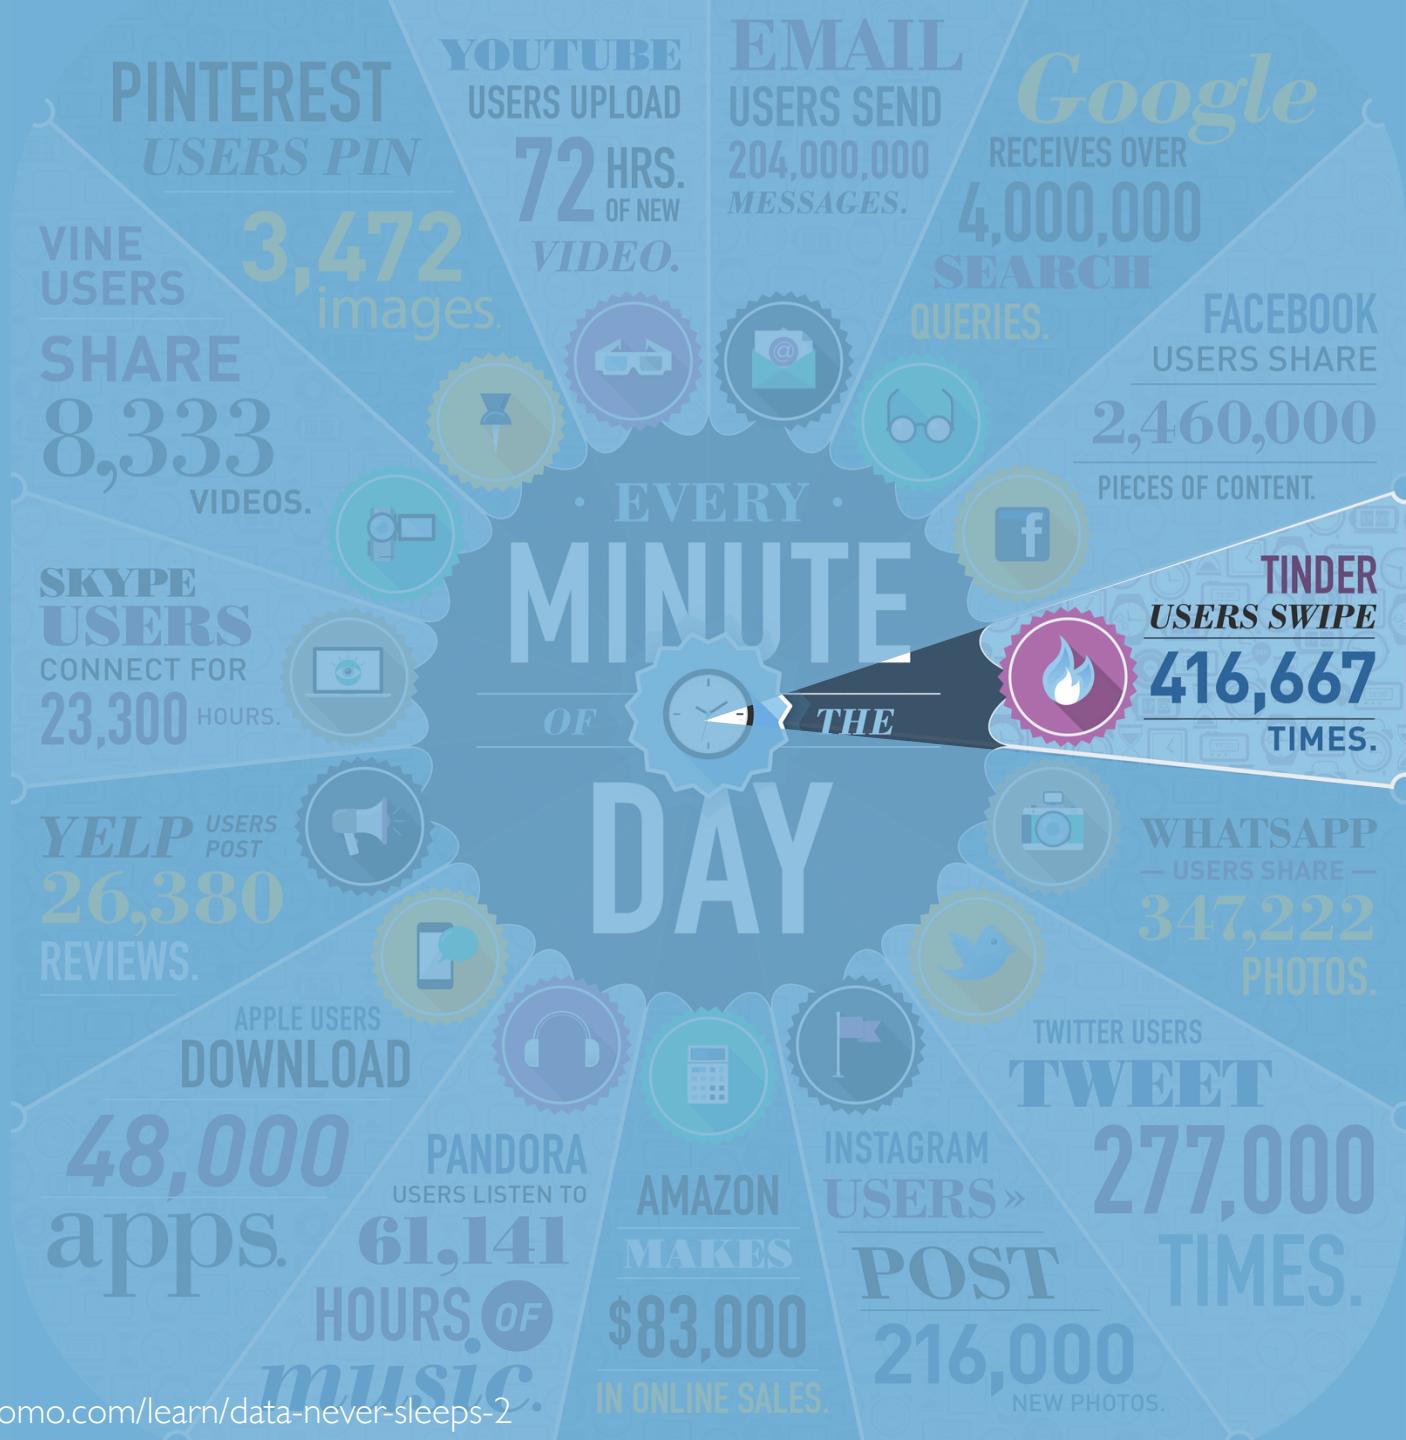
A Systems Approach to Scalable Bayesian Inference

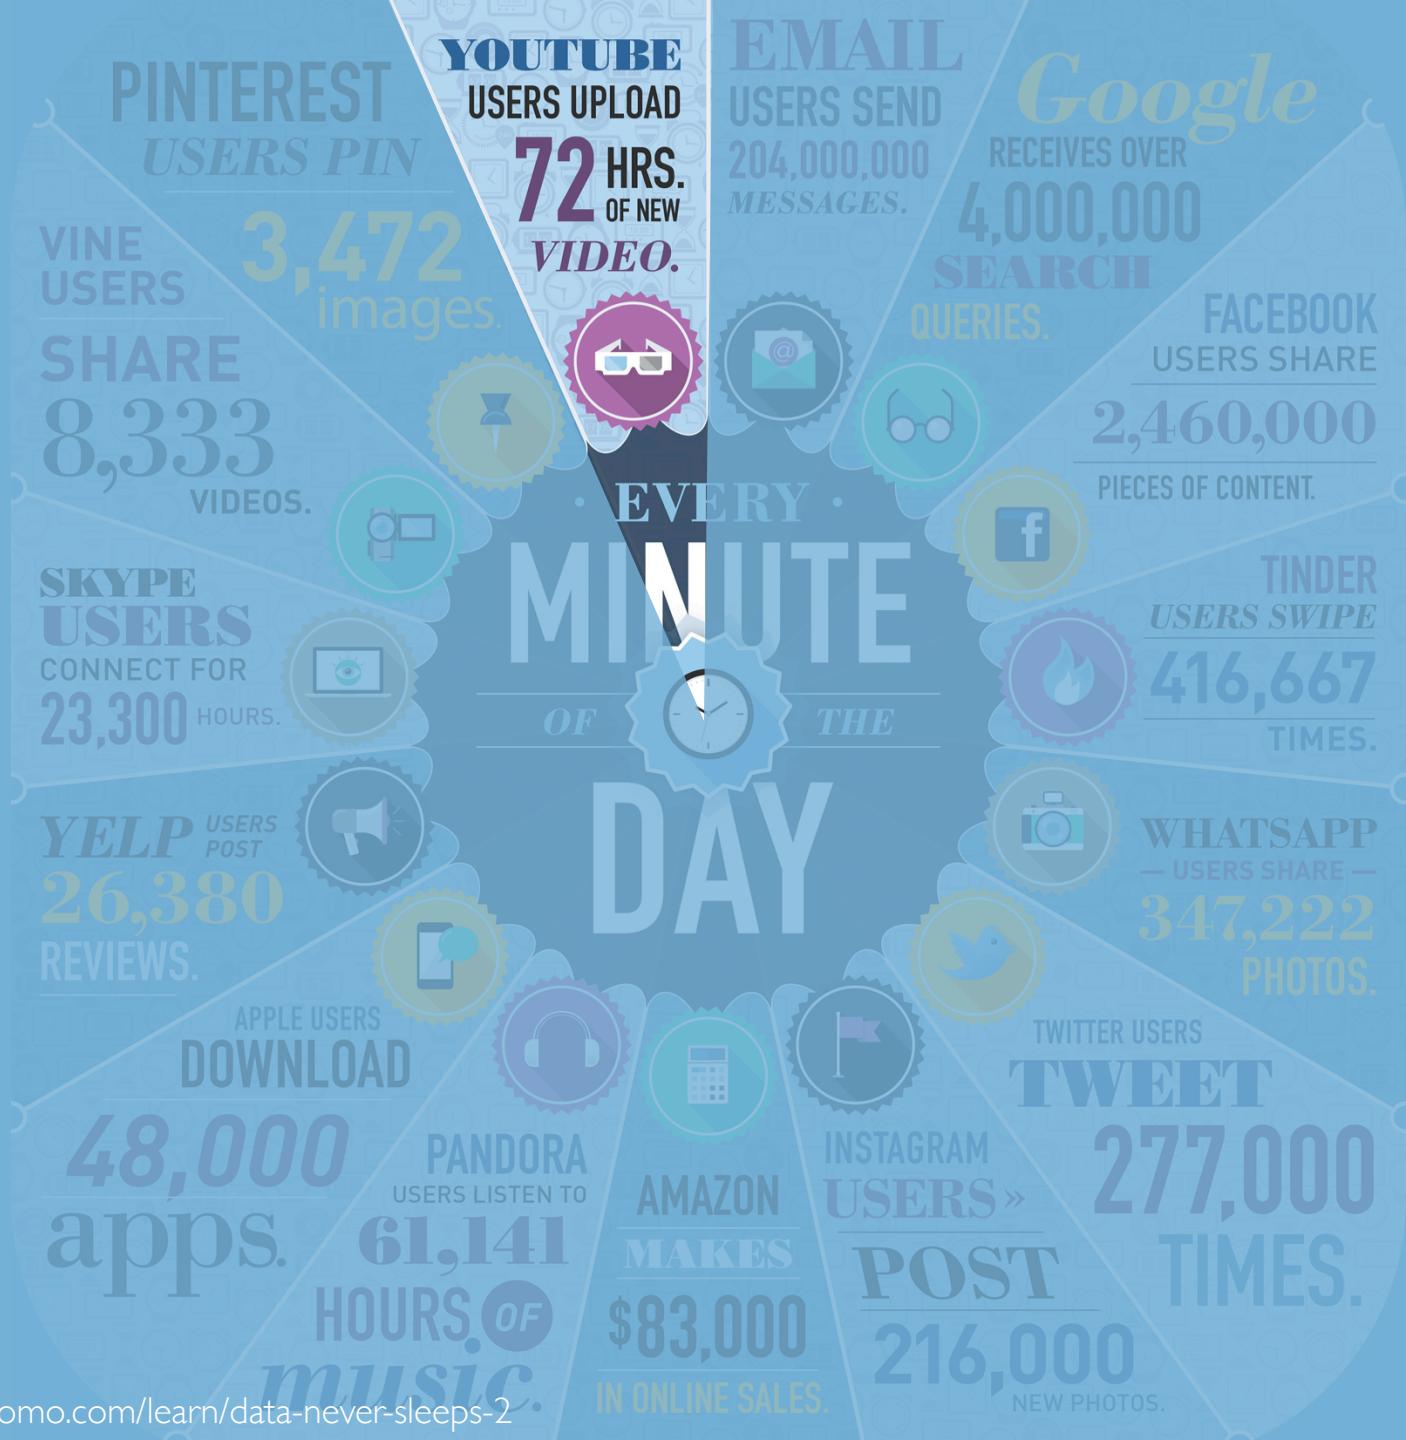
Joseph E. Gonzalez
Postdoc, UC Berkeley AMPLab
Co-founder, GraphLab Inc.
jegonzal@eecs.berkeley.edu

ISBA' 2014





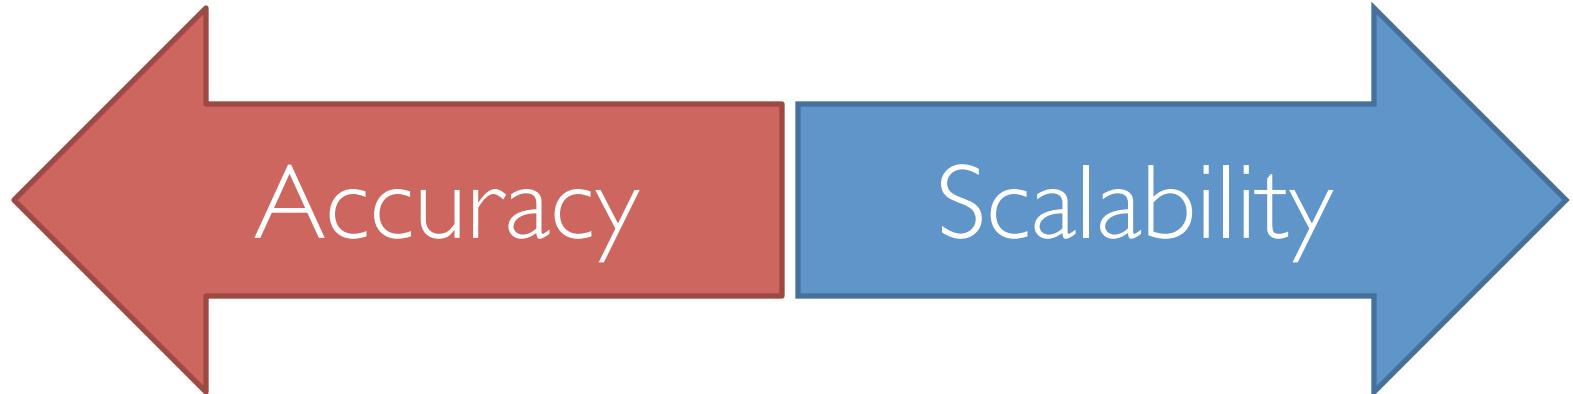




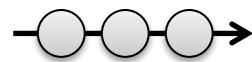
*Data Velocity
is an opportunity for
Bayesian Nonparametrics.*

How do we scale
Bayesian inference?

Opposing Forces

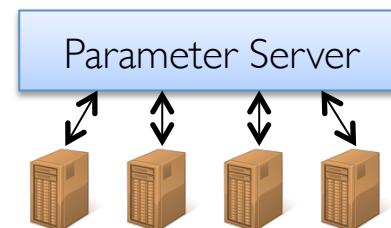


Ability to estimate the posterior distribution



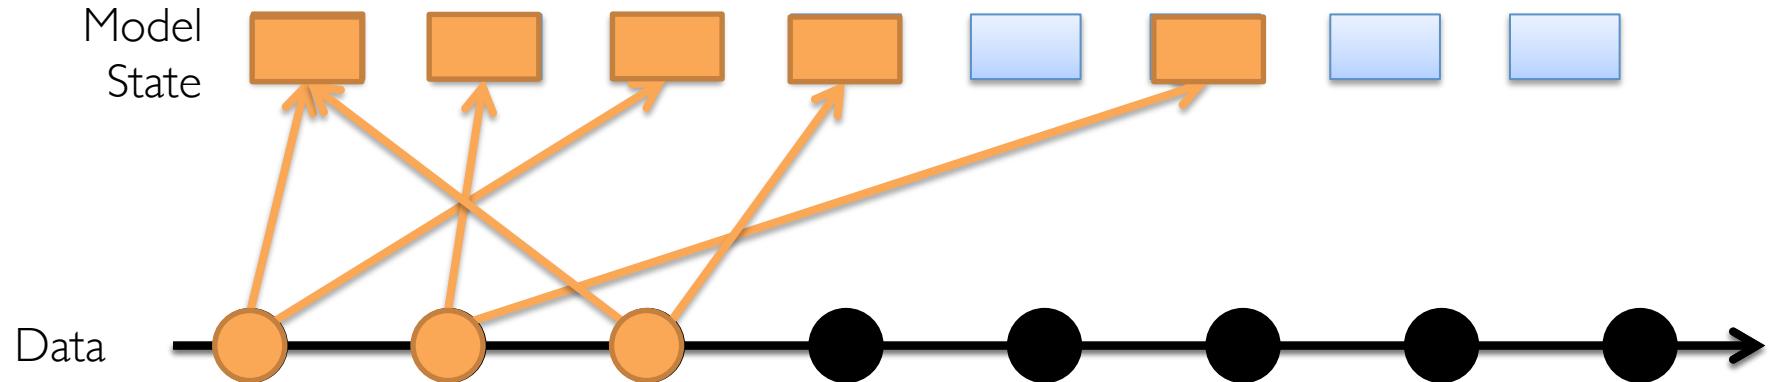
Serial
Inference

Ability to effective use parallel resources

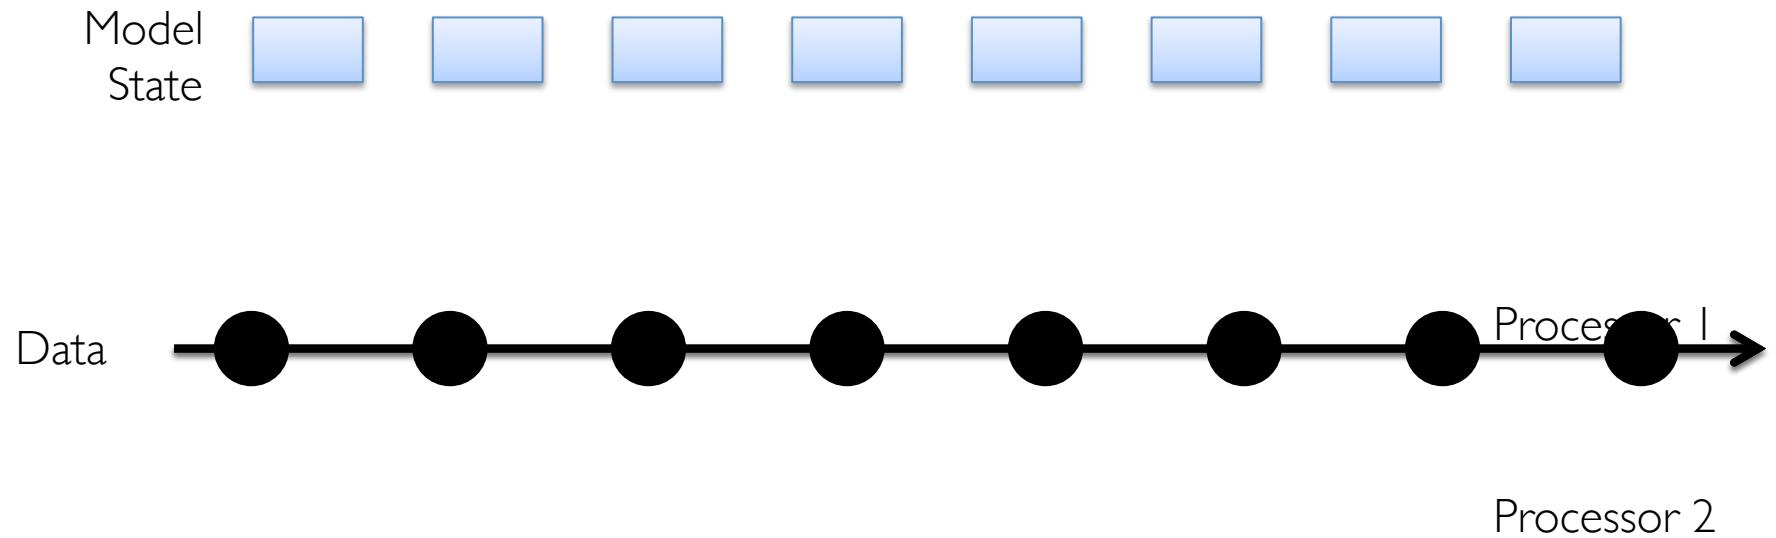


Coordination Free
Samplers

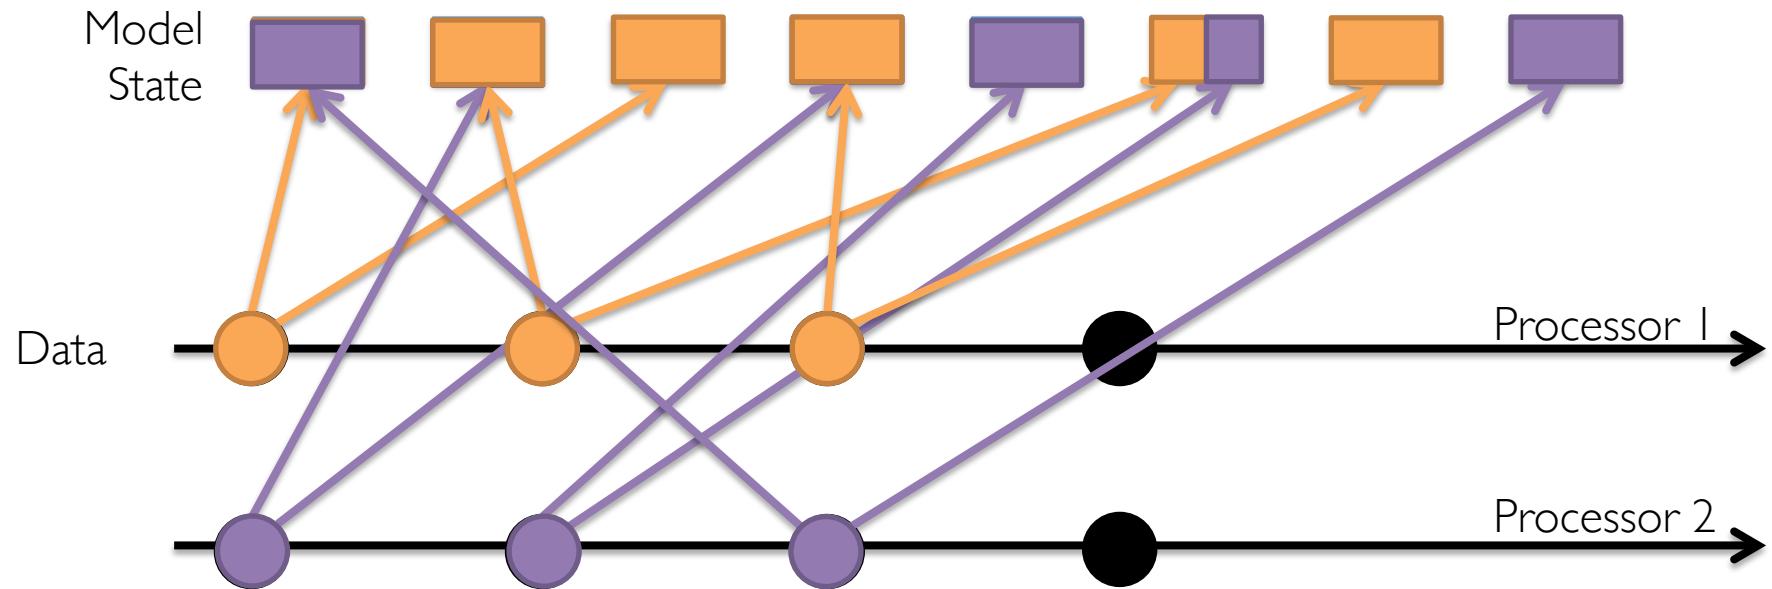
Serial Inference



Coordination Free Parallel Inference



Coordination Free Parallel Inference



Keep Calm and Carry On.

Parameter Servers

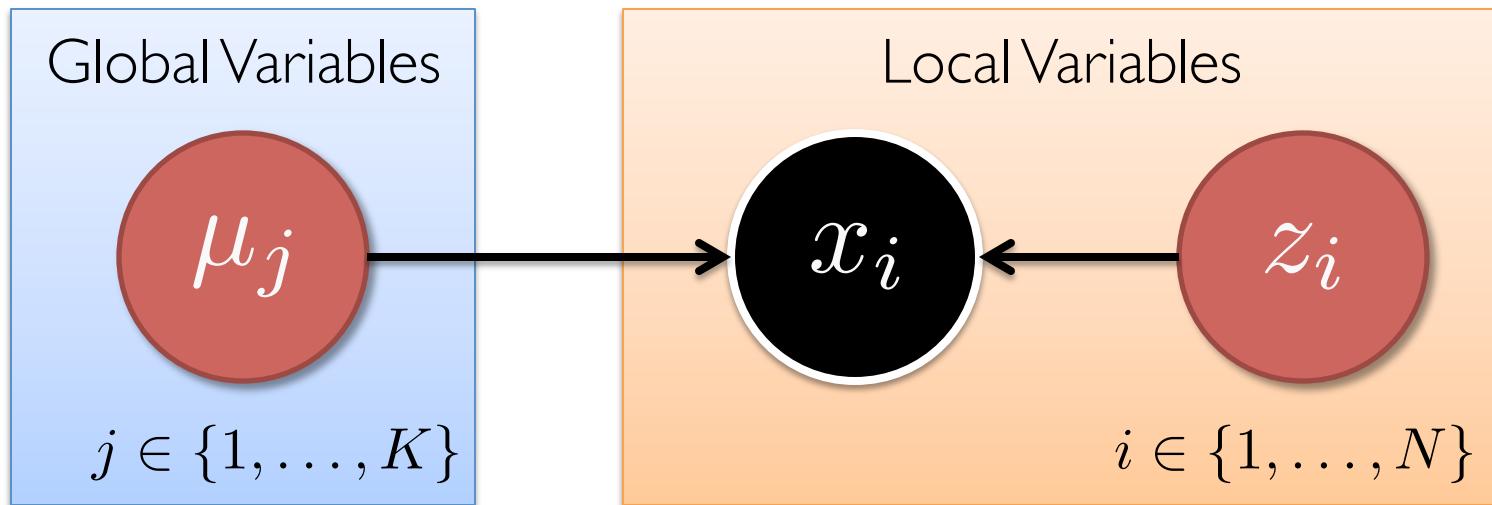
System for Coordination Free Inference

D. Newman, A. Asuncion, P. Smyth, and M. Welling. *Distributed inference for latent Dirichlet allocation*. In NIPS, 2007.

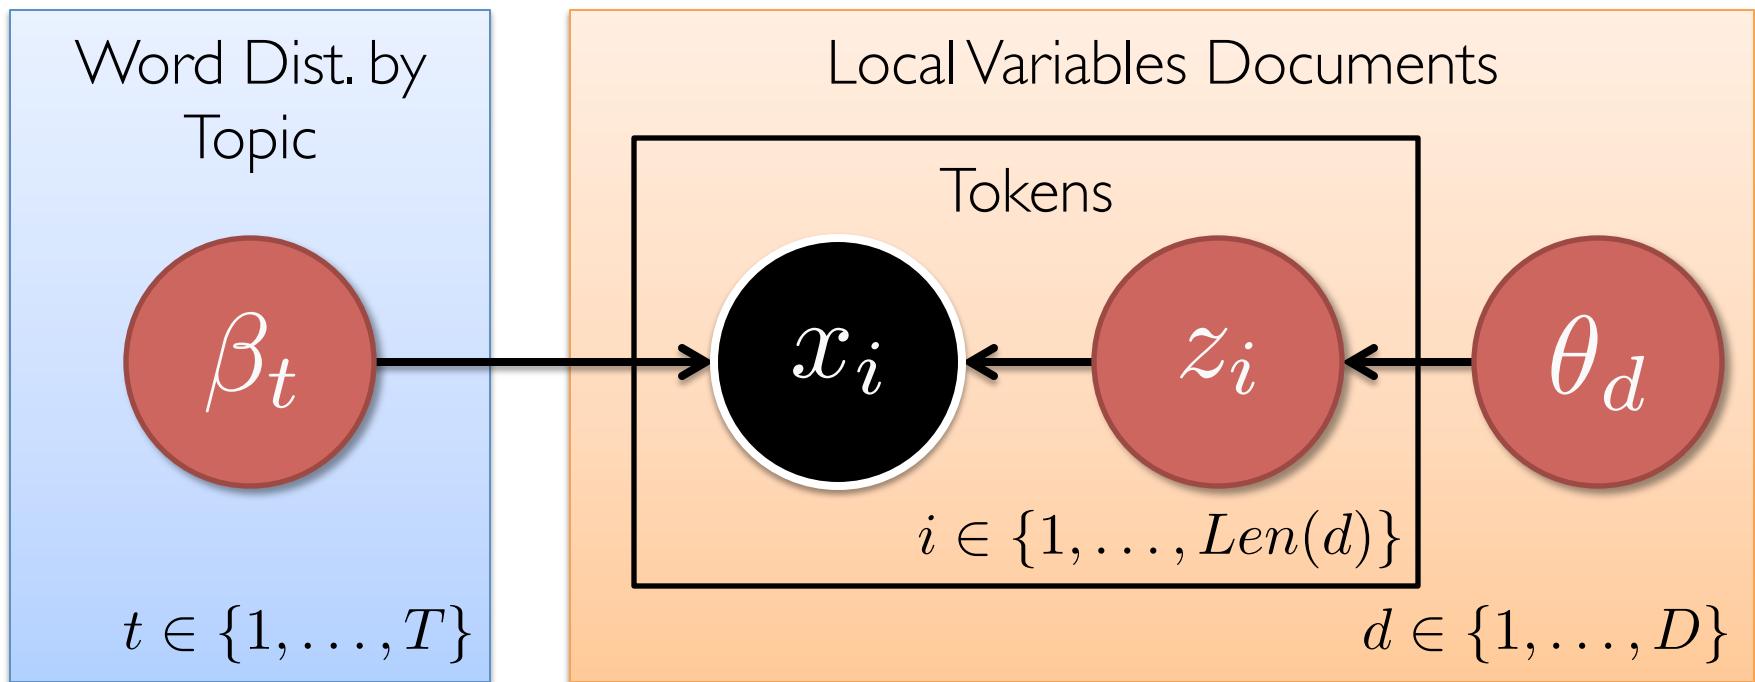
A. Smola and S. Narayananurthy. *An architecture for parallel topic models*. VLDB'10
Ahmed, M. Aly, J. Gonzalez, S. Narayananurthy, and A. J. Smola.
Scalable inference in latent variable models. WSDM '12

Ho et al. “More Effective Distributed ML via a Stale Synchronous Parallel Parameter Server” NIPS’13

Hierarchical Clustering



Example: Topic Modeling with LDA

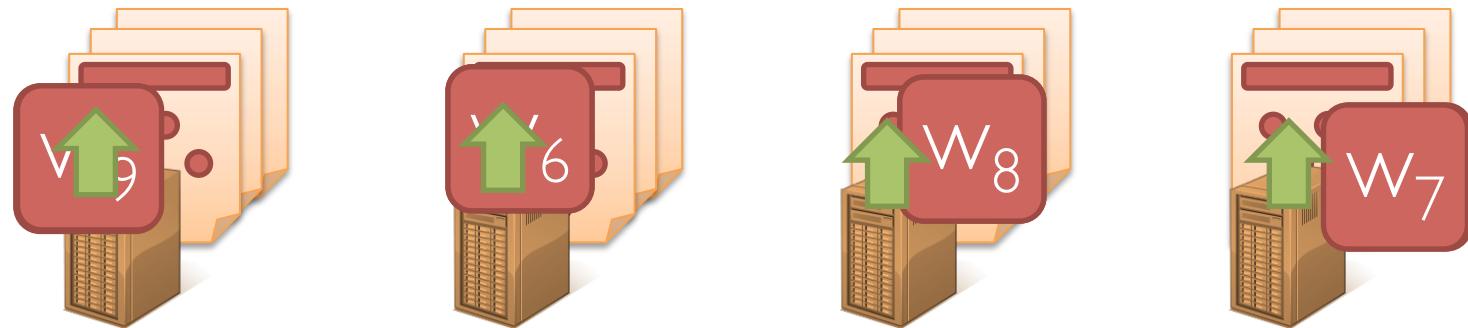
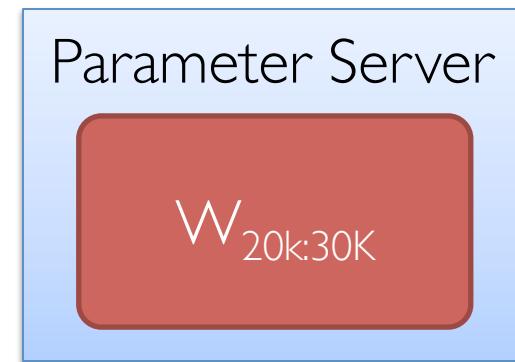
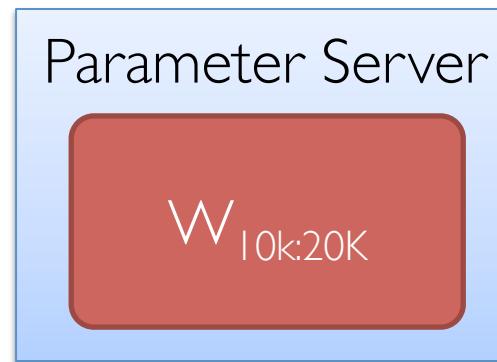
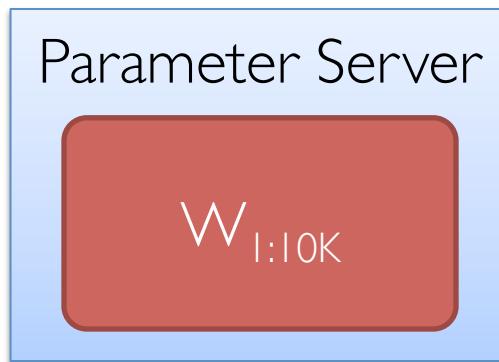


Maintained by the
Parameter Server

Maintained by the
Workers Nodes

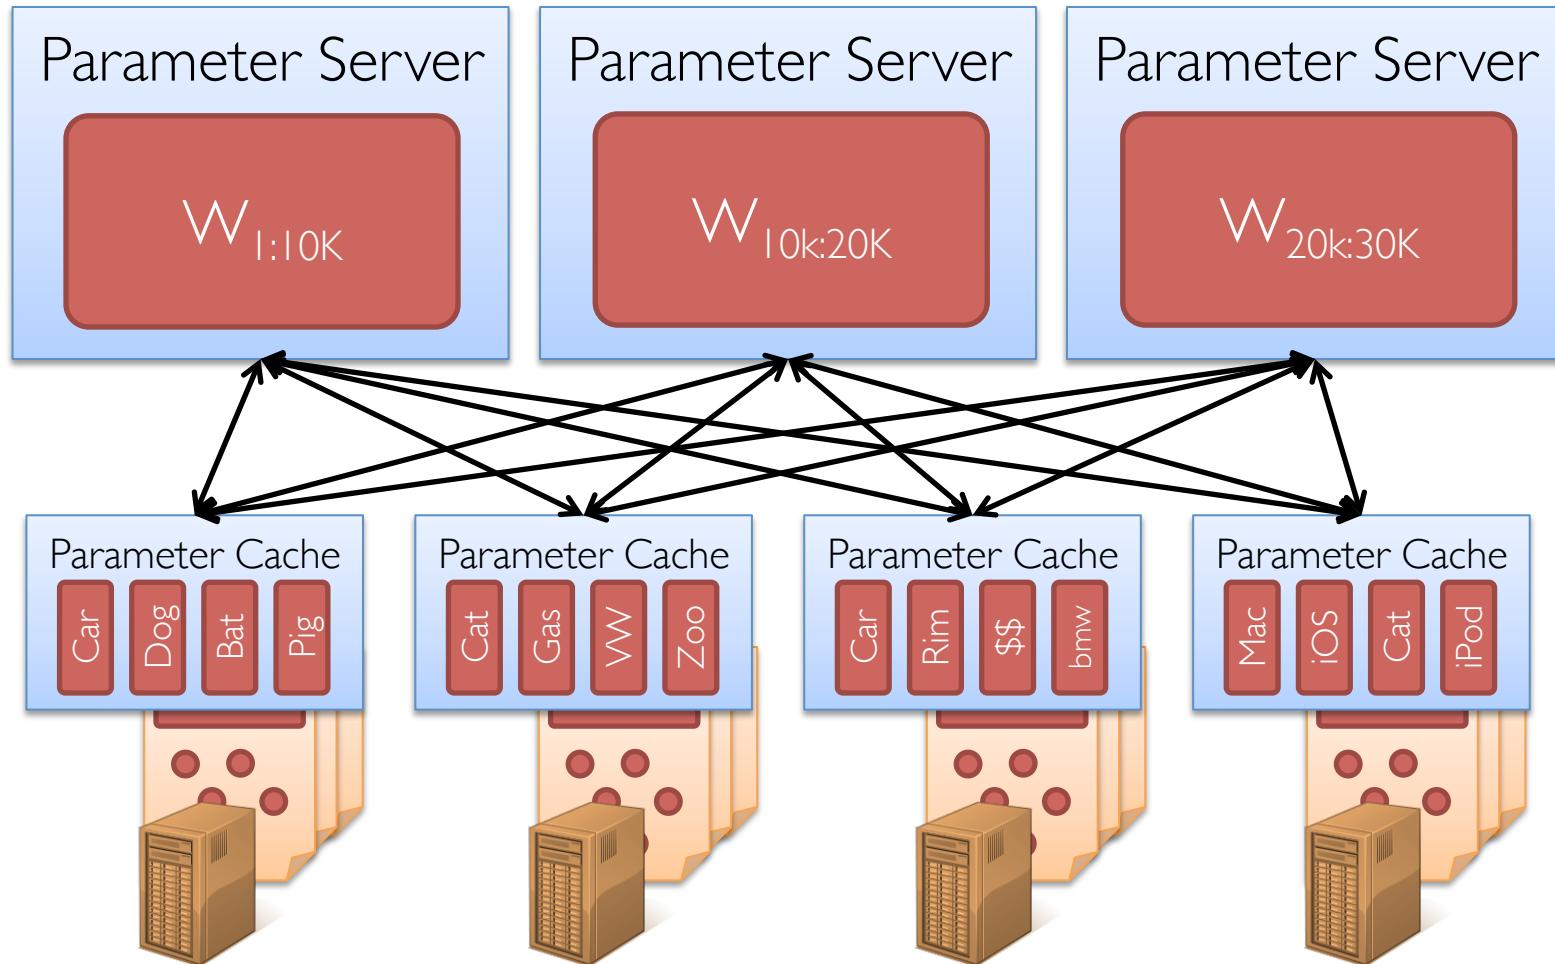
Ex: Collapsed Gibbs Sampler for LDA

Partitioning the model and data



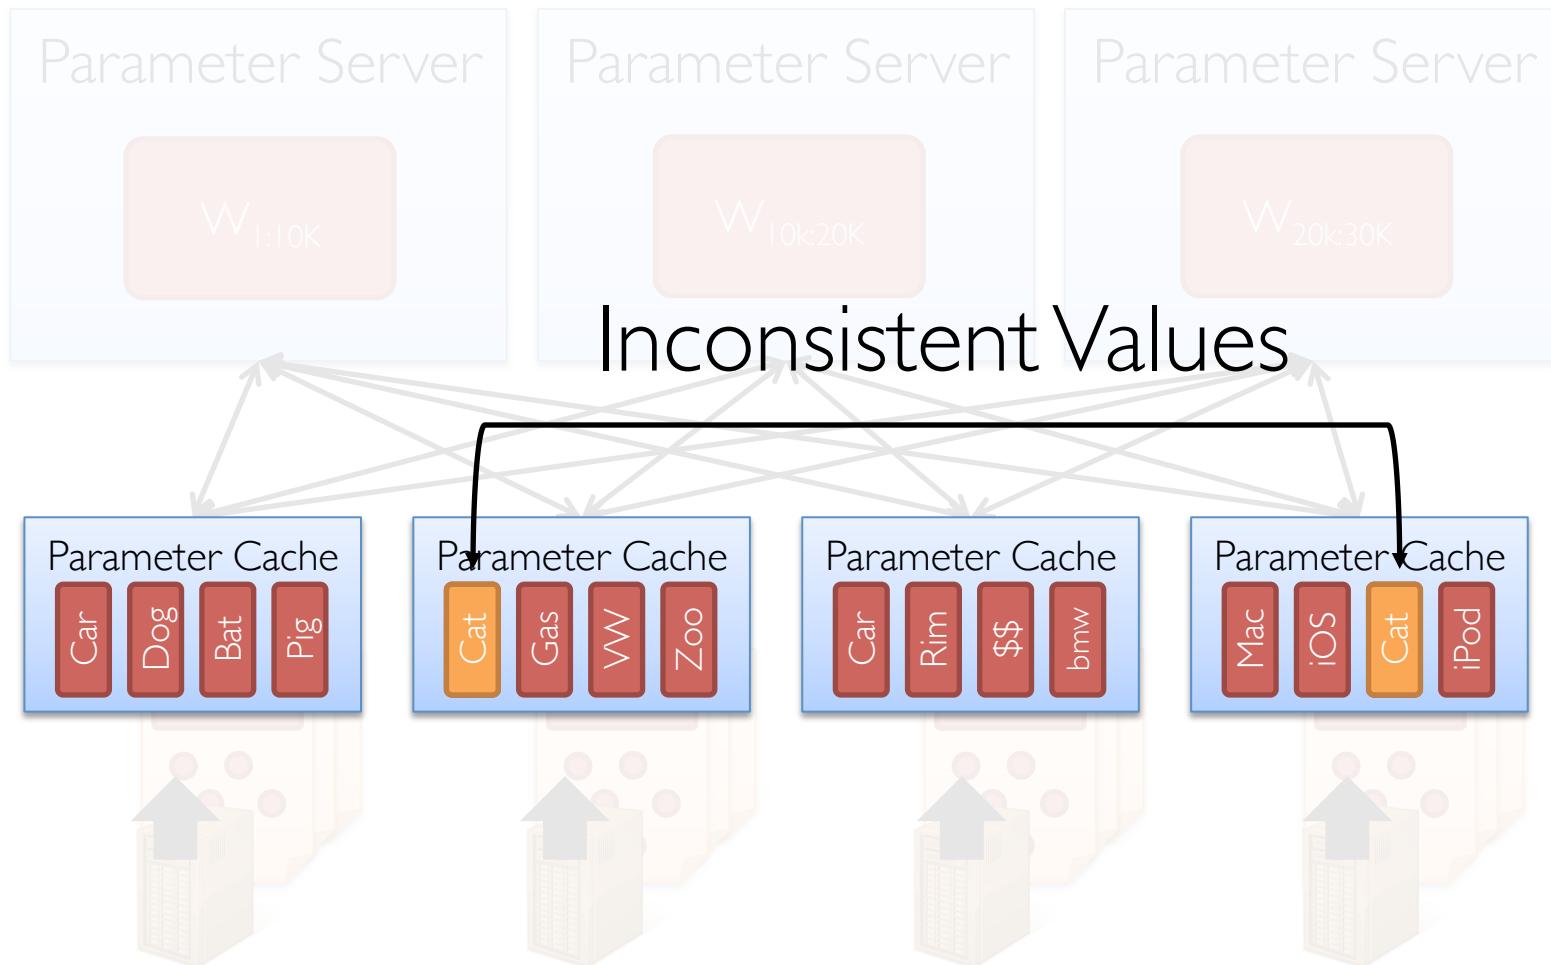
Ex: Collapsed Gibbs Sampler for LDA

Partitioning the model and data



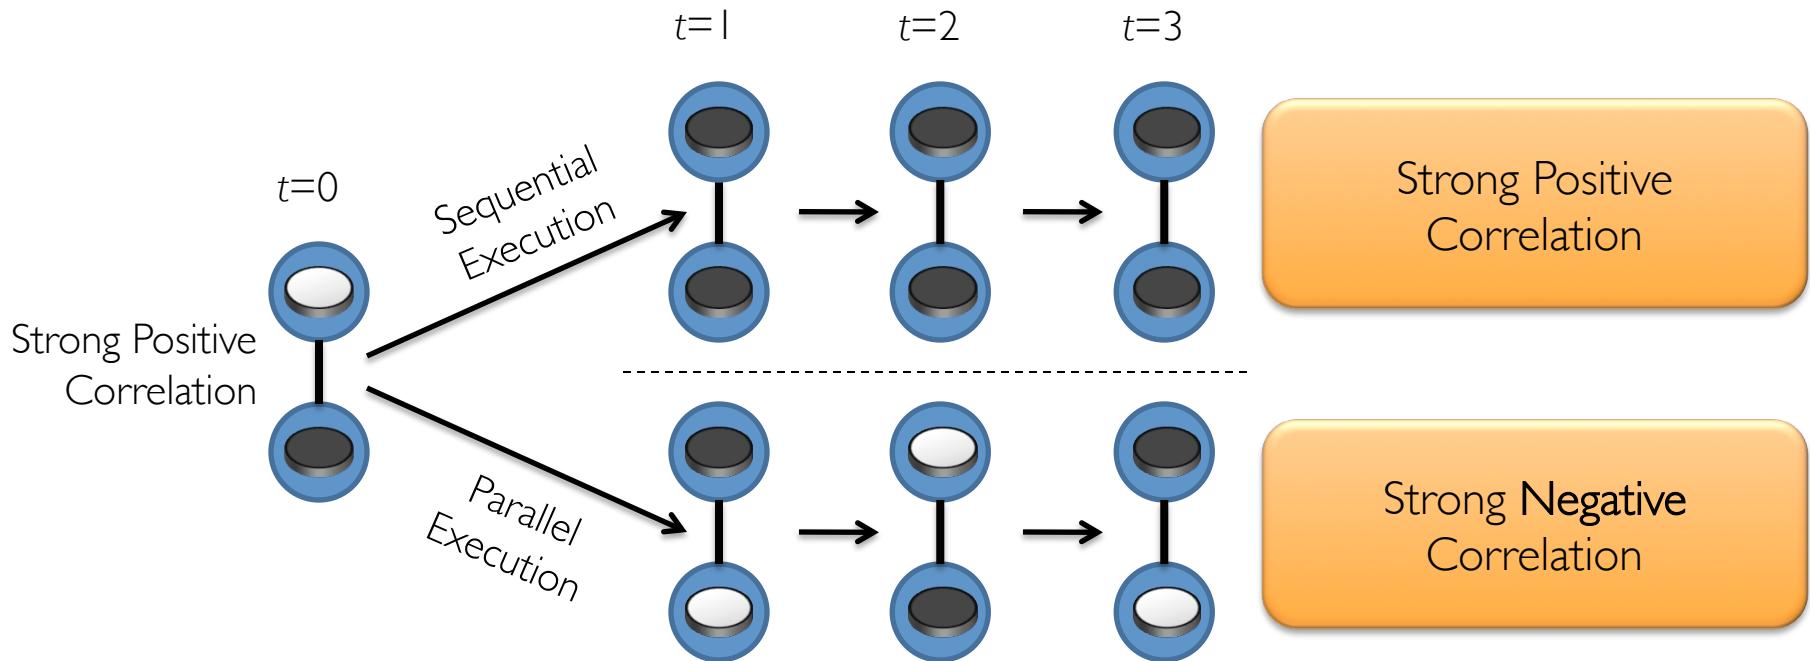
Ex: Collapsed Gibbs Sampler for LDA

Inconsistent model replicas



Parallel Gibbs Sampling

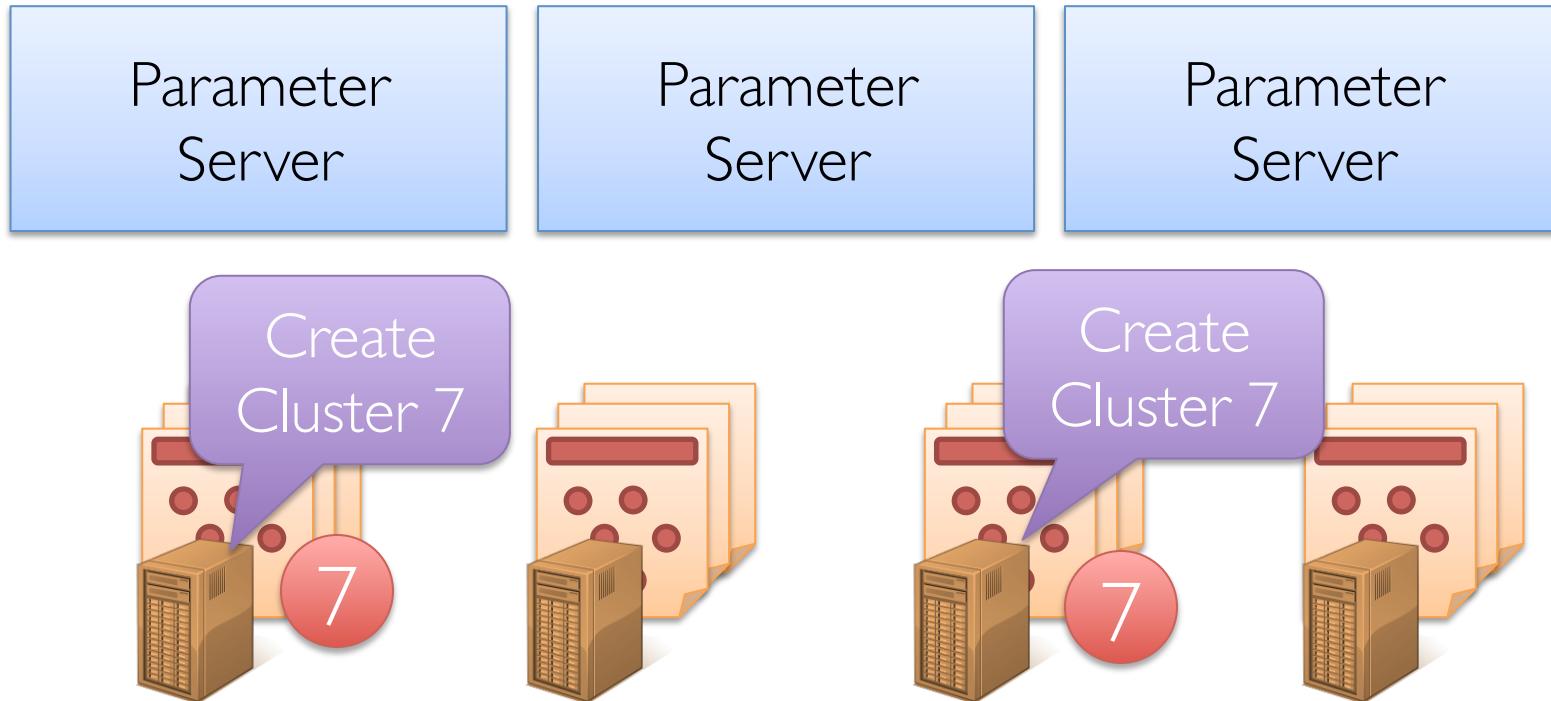
Incorrect Posterior



dependent variables cannot in general be sampled simultaneously.

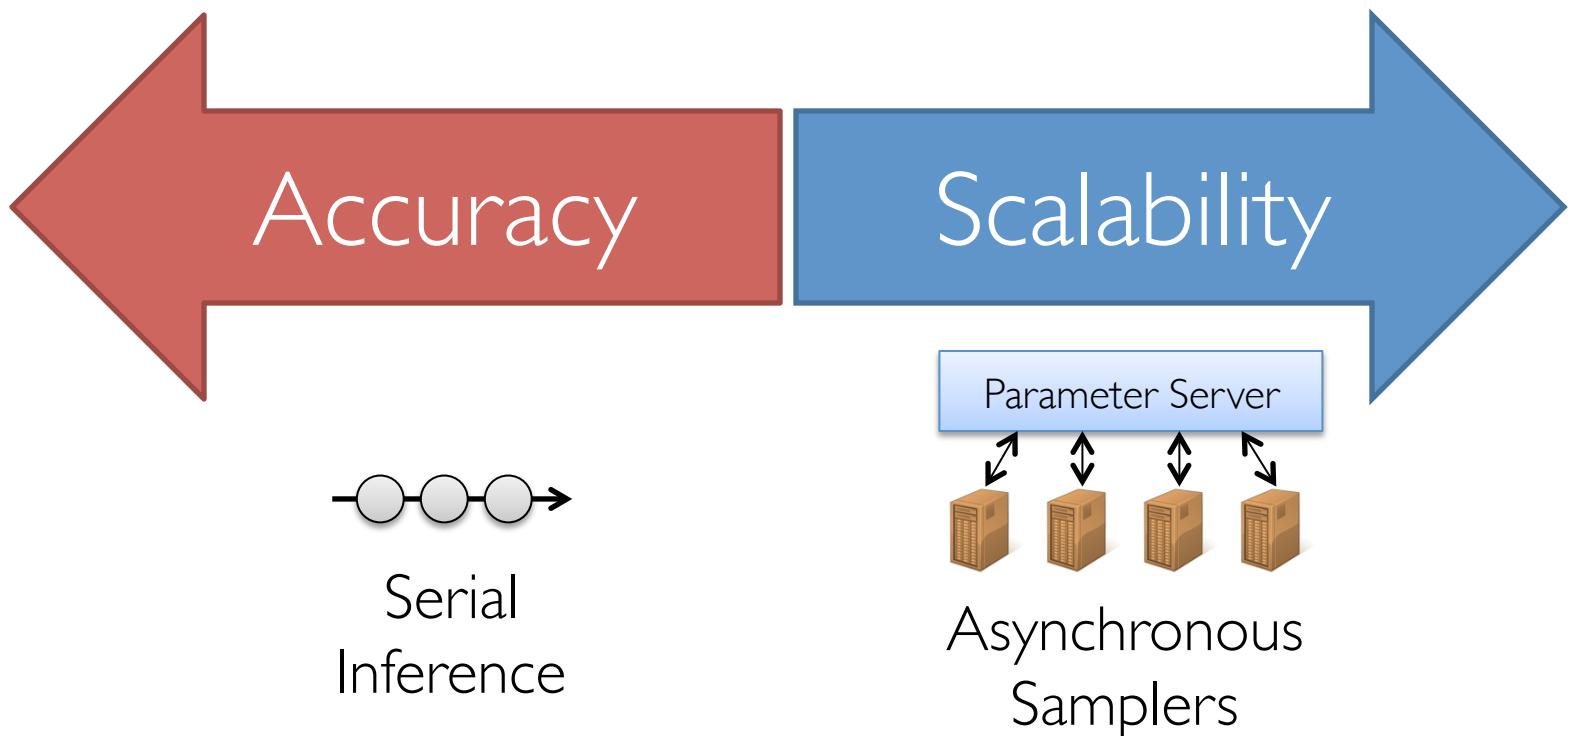
Issues with Nonparametrics

Difficult to introduce new clusters asynchronously:

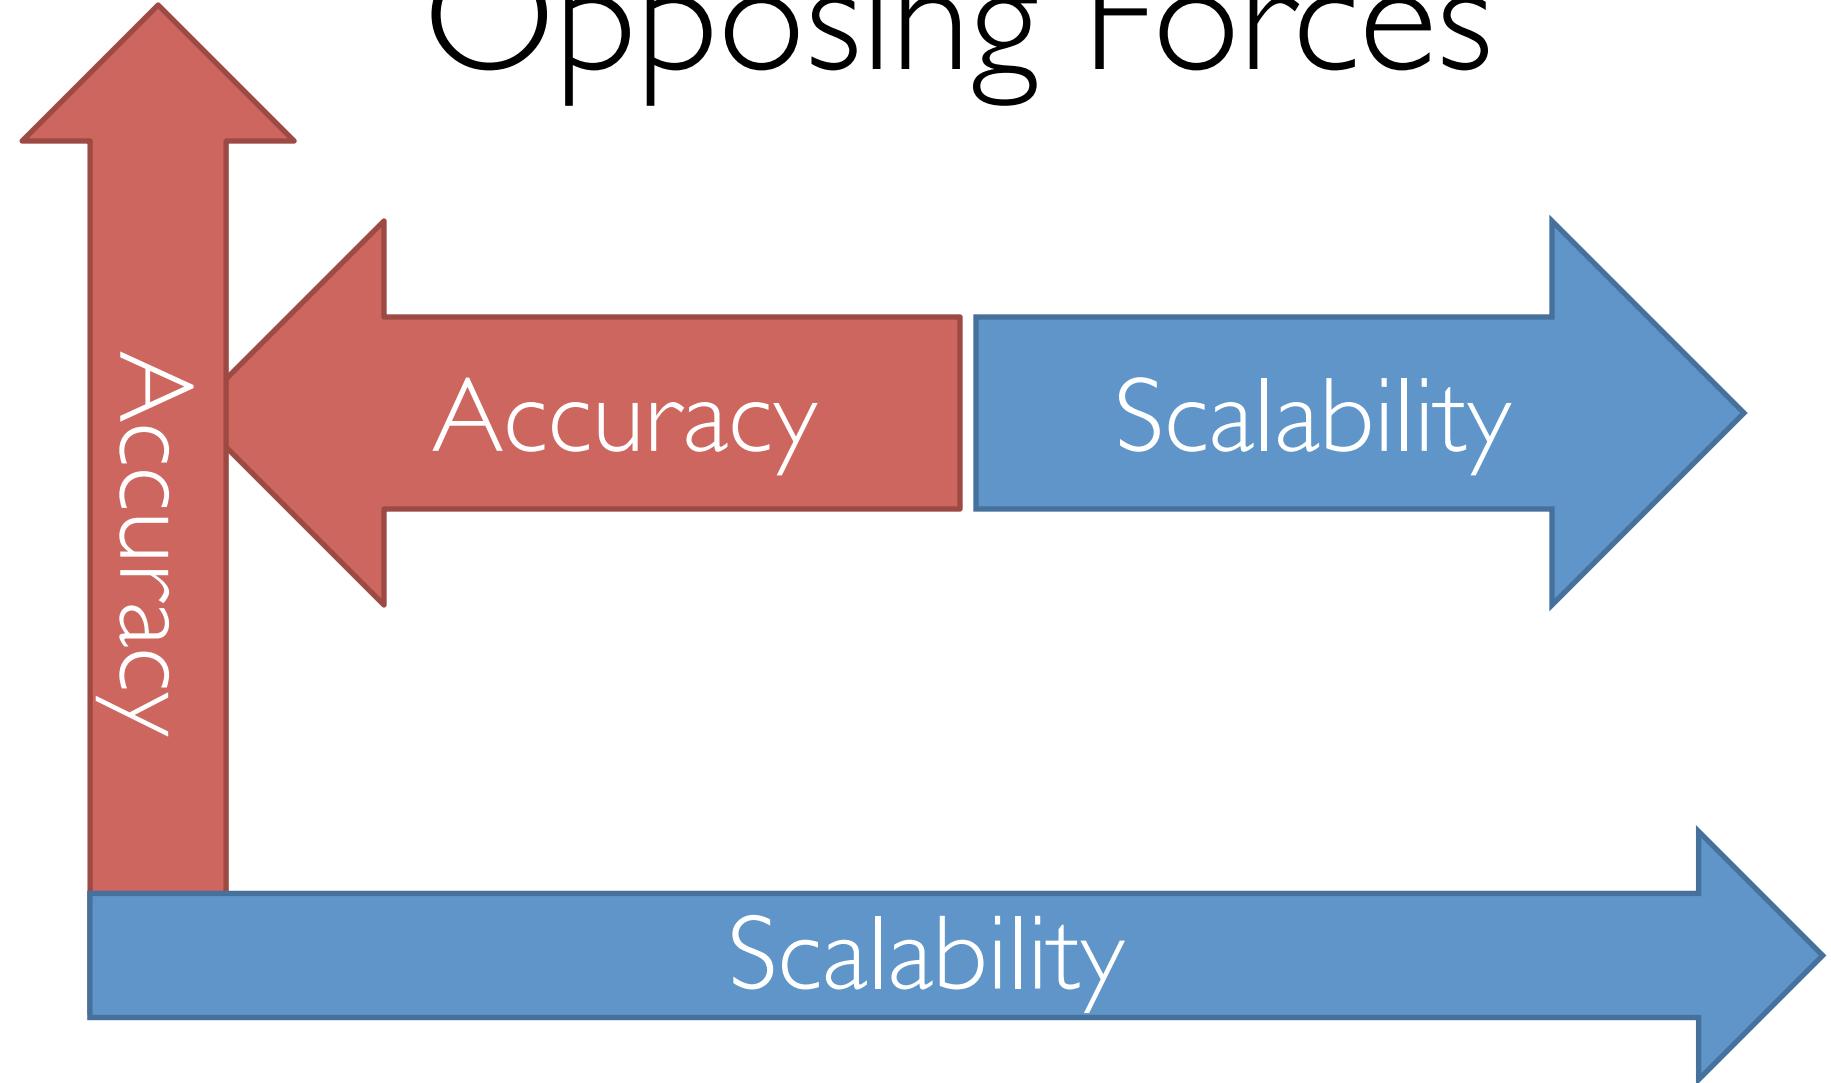


Leads to too many clusters!

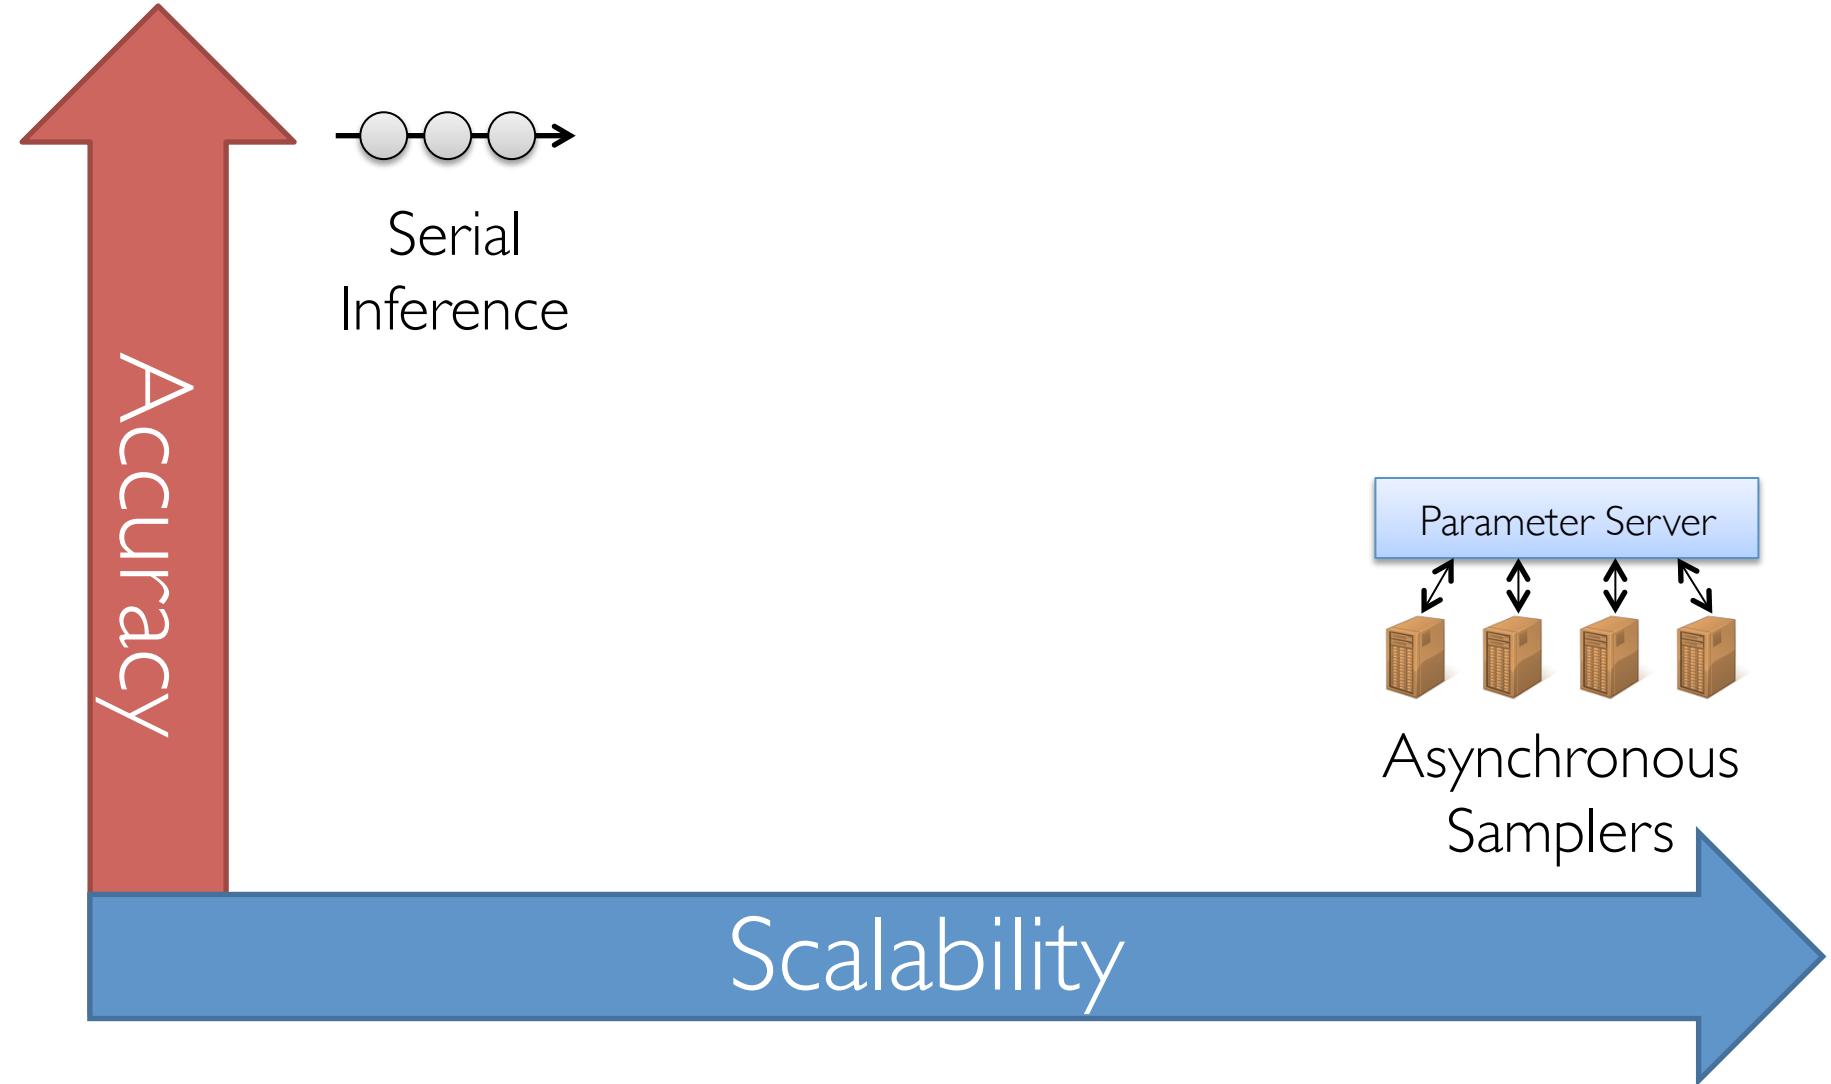
Opposing Forces



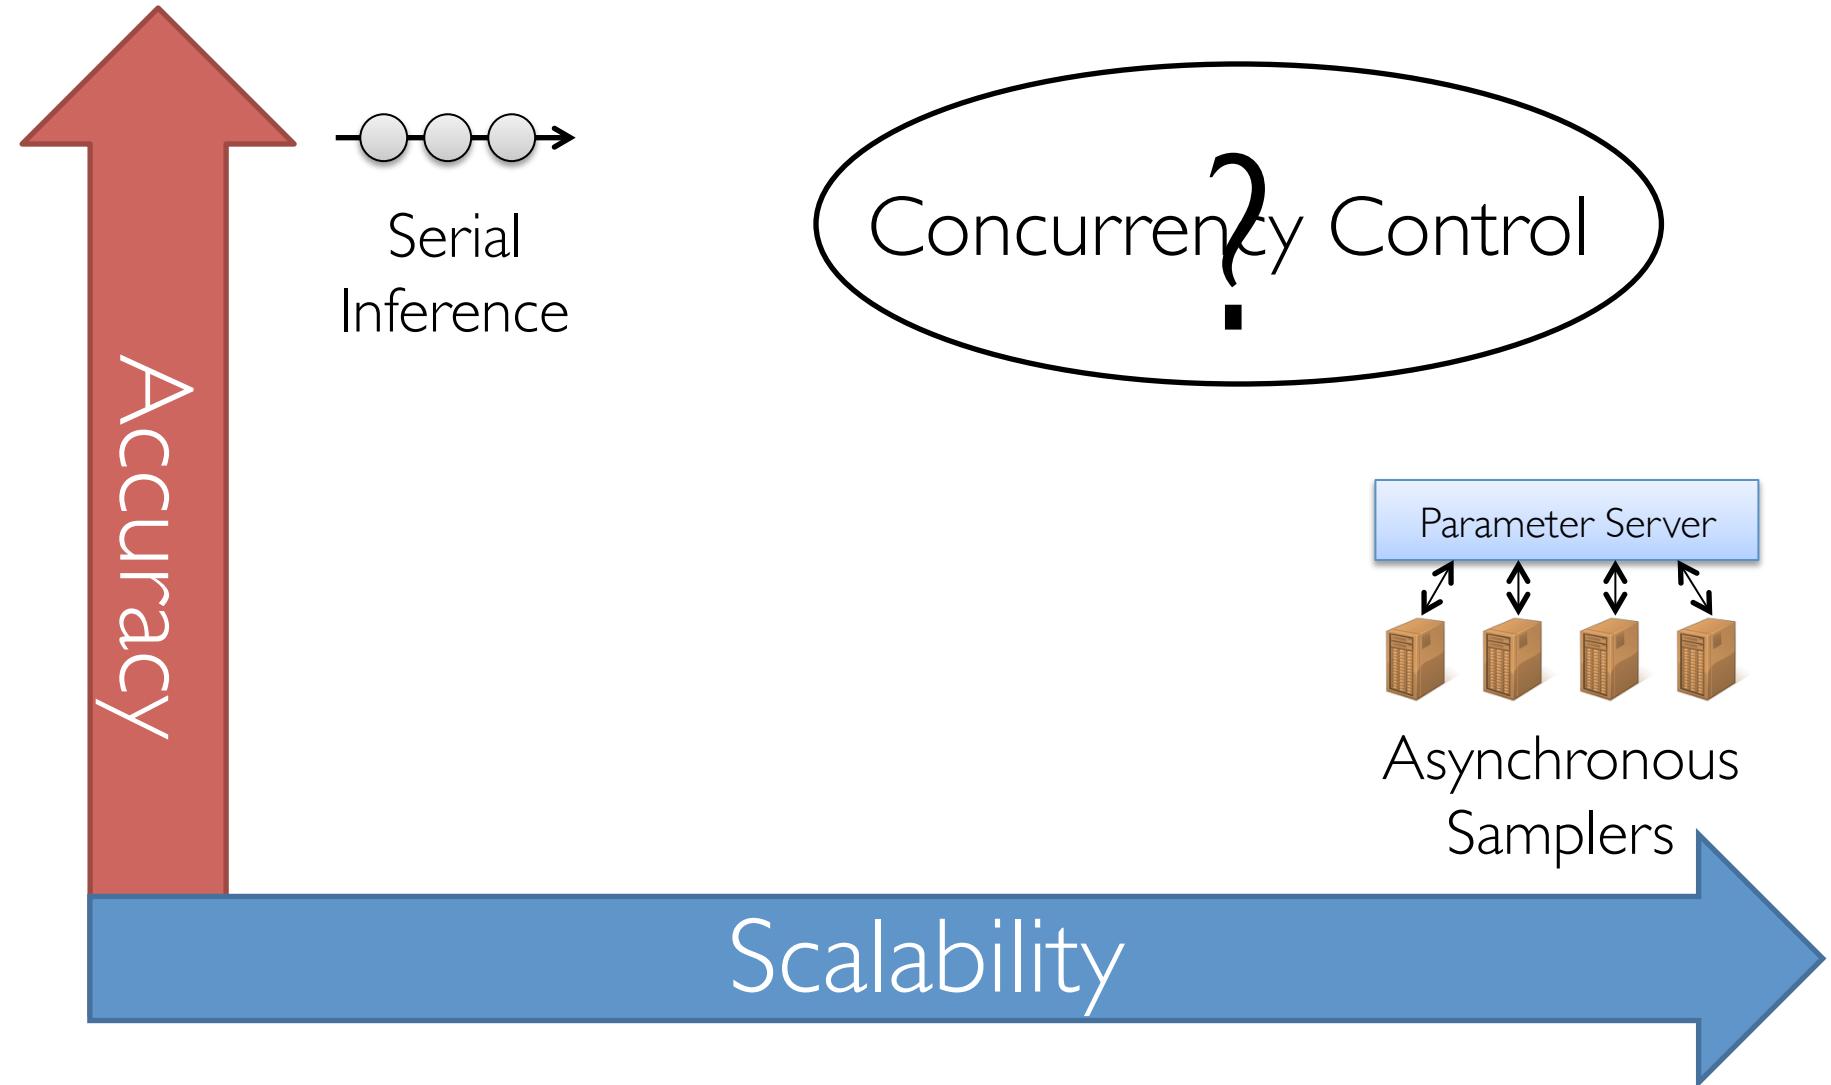
Opposing Forces



Opposing Forces



Opposing Forces



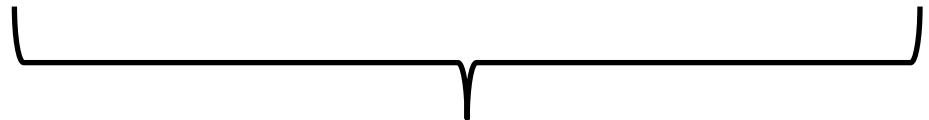
Concurrency Control

Coordination Free (Parameter Server):

Provably fast and correct under key assumptions.

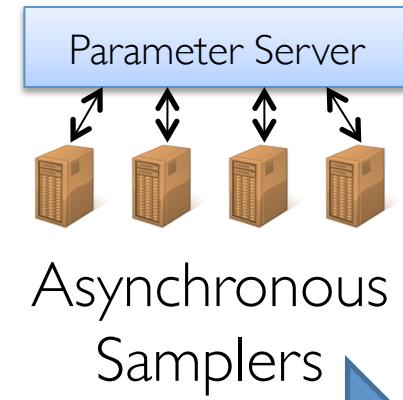
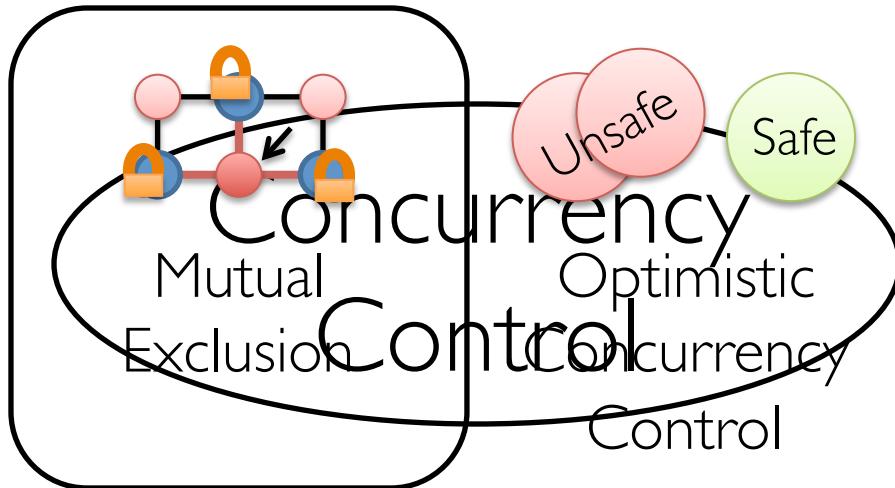
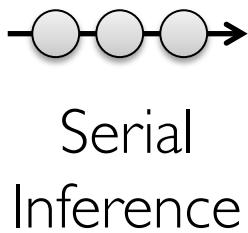
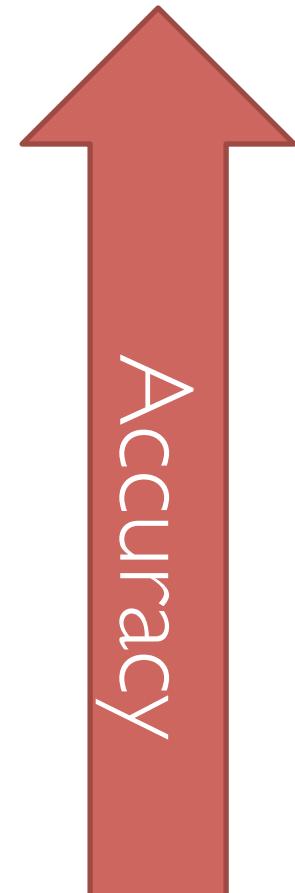
Concurrency Control:

Provably correct and fast under key assumptions.



Systems Ideas to
Improve Efficiency

Opposing Forces



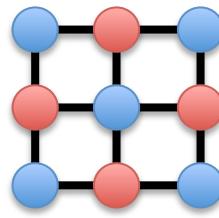
Scalability

Mutual Exclusion

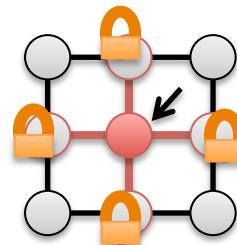
Conditional Independence

J. Gonzalez, Y. Low, A. Gretton, and C. Guestrin. *Parallel Gibbs Sampling: From Colored Fields to Thin Junction Trees*. AISTATS'11

Exploit the Markov random field for
Parallel Gibbs Sampling



Graph
Coloring



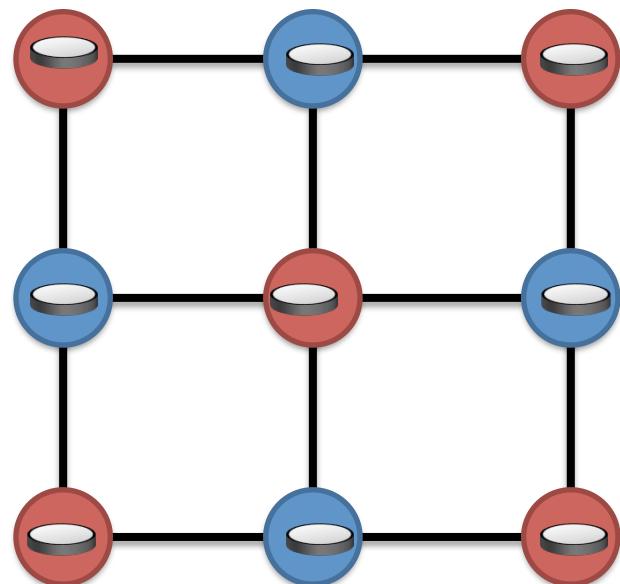
R/W Lock
Mutual Exclusion

Mutual Exclusion through Scheduling Chromatic Gibbs Sampler

Compute a k-coloring of the graphical model

Sample all variables with same color in parallel

Serial Equivalence:



Theorem: Chromatic Sampler

Ergodic: converges to the correct distribution

» Based on graph coloring of the Markov Random Field

Quantifiable acceleration in **mixing**

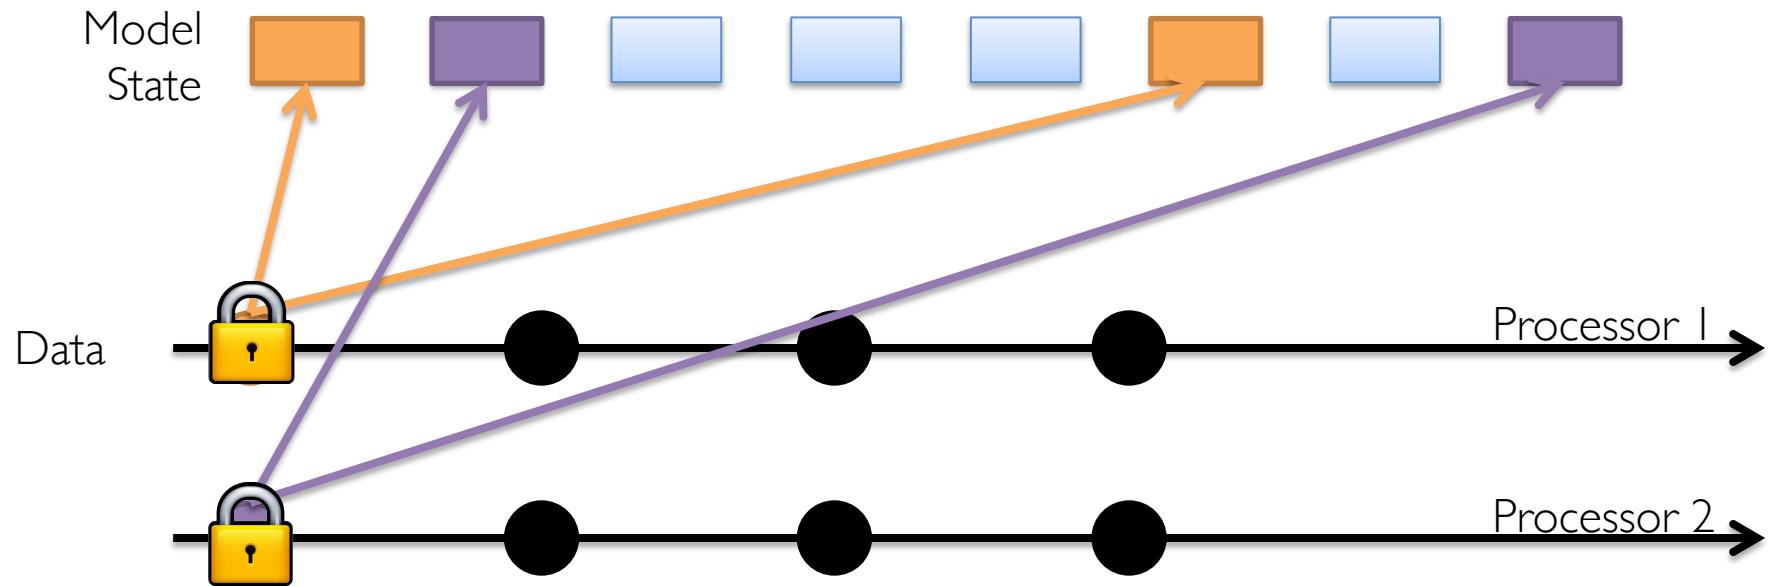
Time to update
all variables once

$$O\left(\frac{n}{p} + k\right)$$

Variables
Colors
Processors

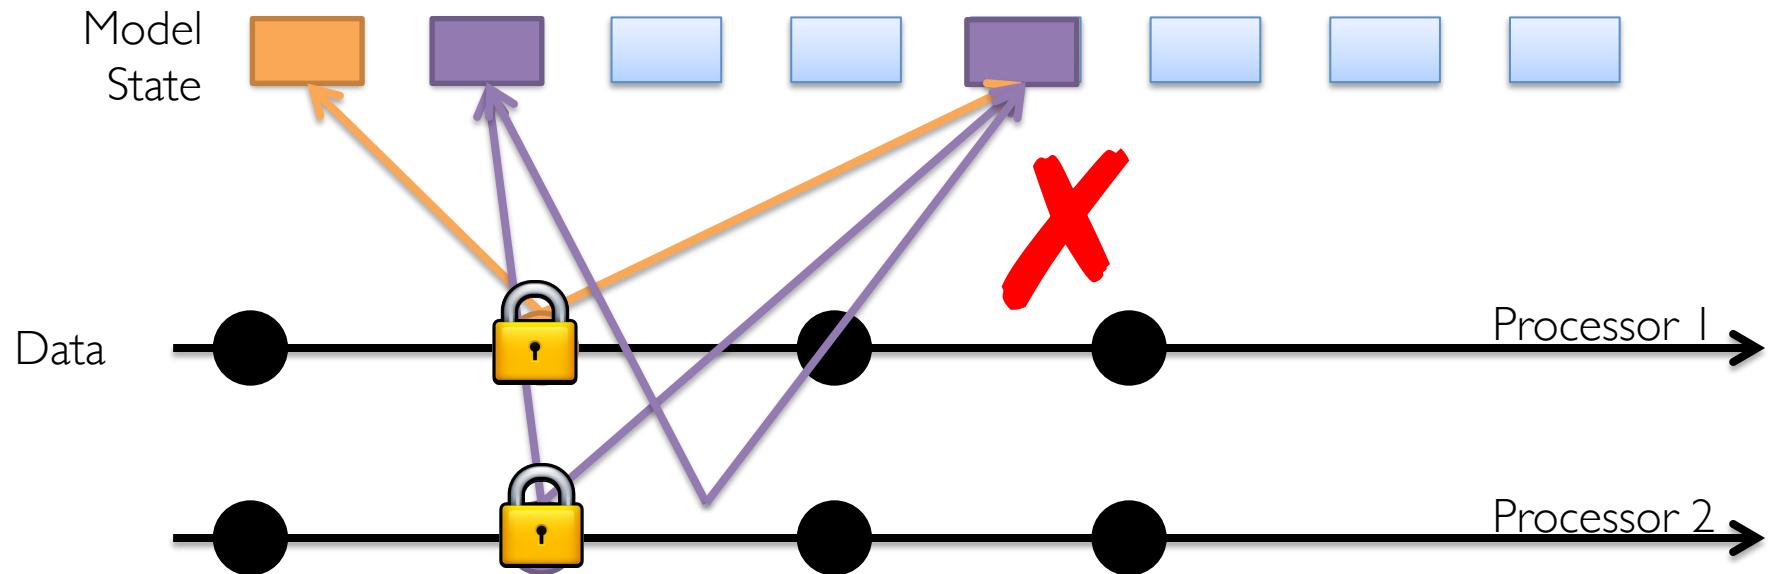
The diagram illustrates the time complexity of the Chromatic Sampler. The expression $O\left(\frac{n}{p} + k\right)$ is shown in a large circle. Three blue arrows point from the terms outside the circle to their corresponding labels: a top arrow points to "# Variables", a middle arrow points to "# Colors", and a bottom arrow points to "# Processors".

Mutual Exclusion Through Locking



Introducing locking (scheduling) protocols to identify potential conflicts.

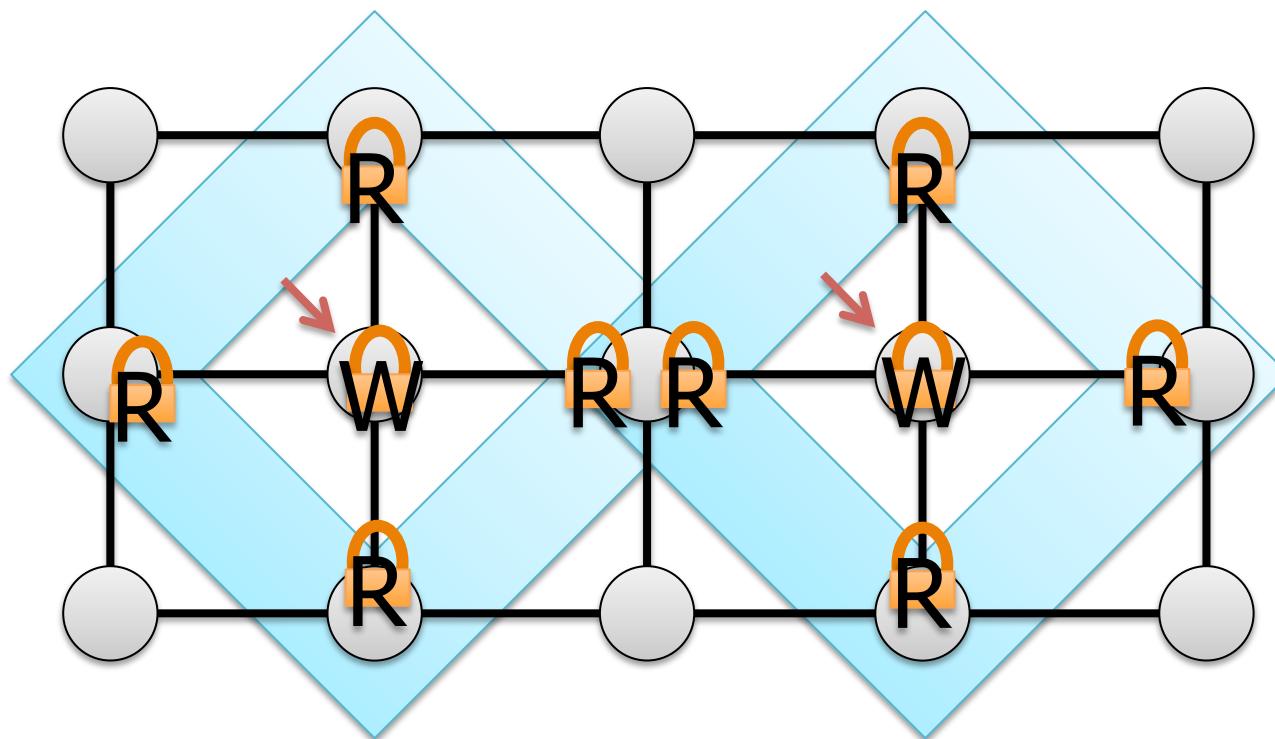
Mutual Exclusion Through Locking



Enforce serialization of computation that could conflict.

Markov Blanket Locks

Read/Write Locks:



Markov Blanket Locks

Eliminate fixed schedule and global coordination

Supports more advanced block sampling

Expected Parallelism:

$\mathbf{E}(\#\text{active processors})$

$$\geq 1 + (p - 1) \left(1 - (p - 1) \left(\frac{d + 1}{n} \right) \right)$$

Max Degree
Processors # Variables

A System for Mutual Exclusion on Markov Random Fields



GraphLab/PowerGraph [UAI'10, OSDI'12]:

- Chromatic Sampling
- Markov Blanket Locks + Block Sampling

Limitation

Densely Connected MRF

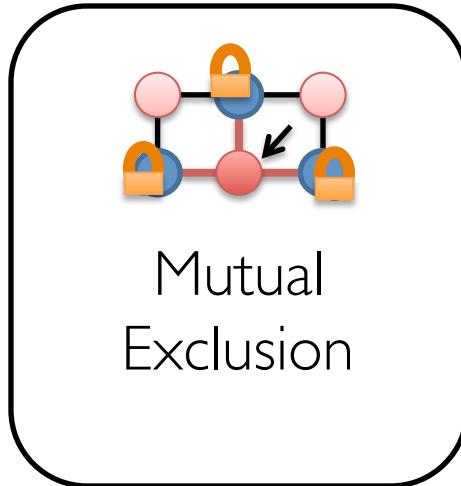
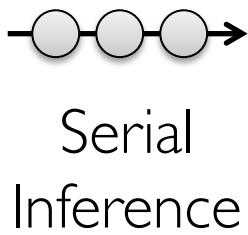
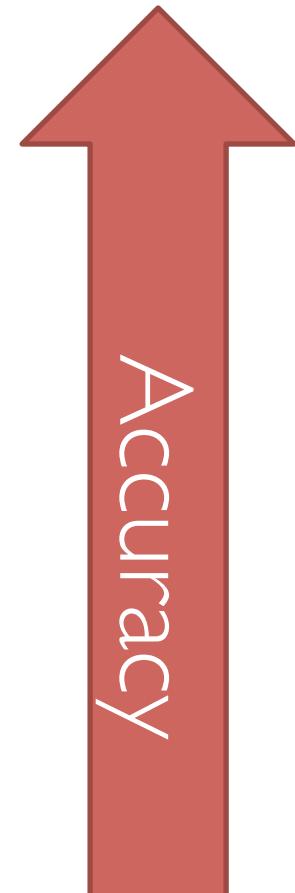
V-Structures: observations couple many variables

Collapsed models: clique-like MRFs

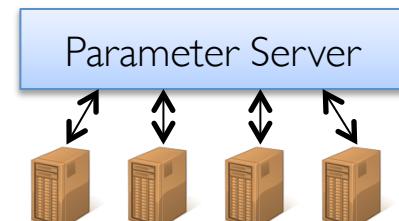
Mutual exclusion *pessimistically* serializes computation that *could* interfere.

Can we be *optimistic* and only serialize computation that *does* interfere?

Opposing Forces



Optimistic
Concurrency
Control



Asynchronous
Samplers

Scalability

Optimistic Concurrency Control

assume the best and correct

X. Pan, J. Gonzalez, S. Jegelka, T. Broderick, M. Jordan. Optimistic Concurrency Control for Distributed Unsupervised Learning. NIPS'13



Xinghao
Pan



Tamara
Broderick



Stefanie
Jegelka



Michael
Jordan

Optimistic Concurrency Control

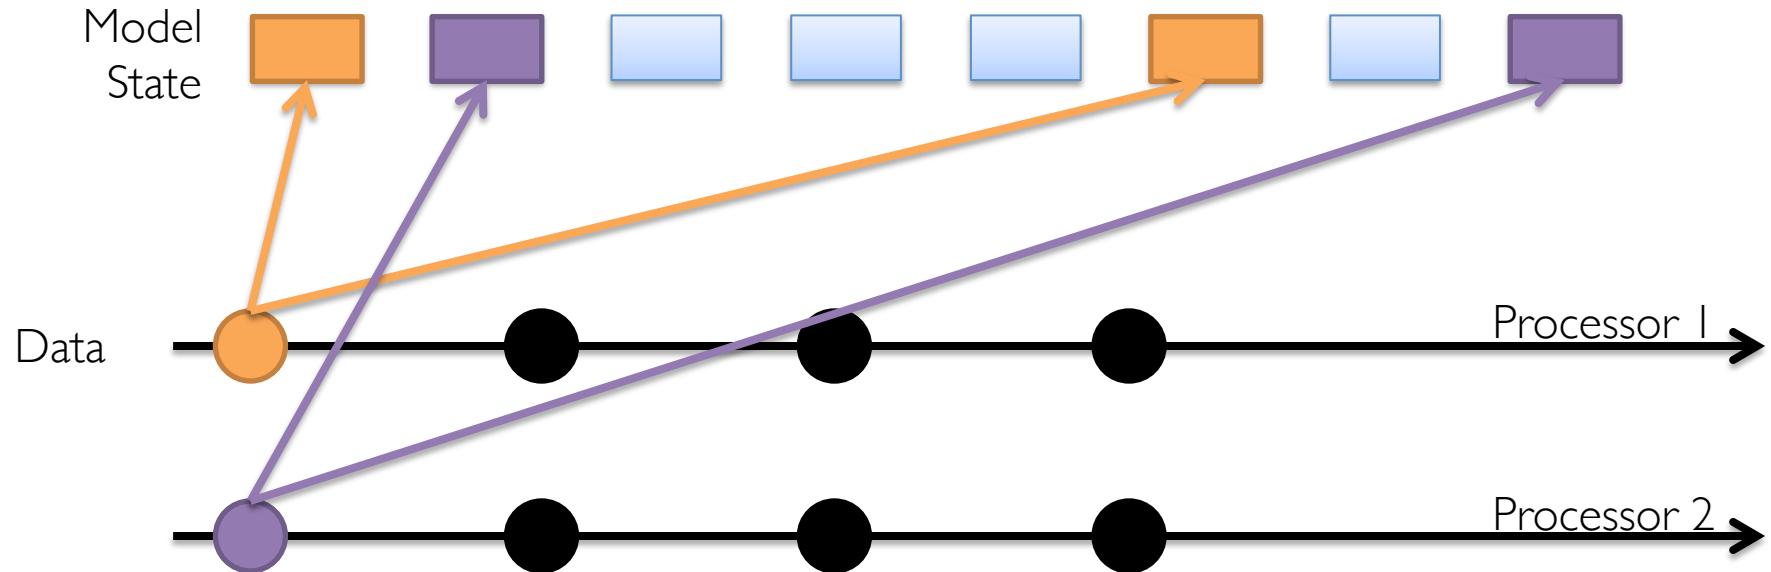
Classic idea from Database Systems:

Kung & Robinson. *On optimistic methods for concurrency control.*
ACM Transactions on Database Systems. 1981

Assume most operations won't conflict:

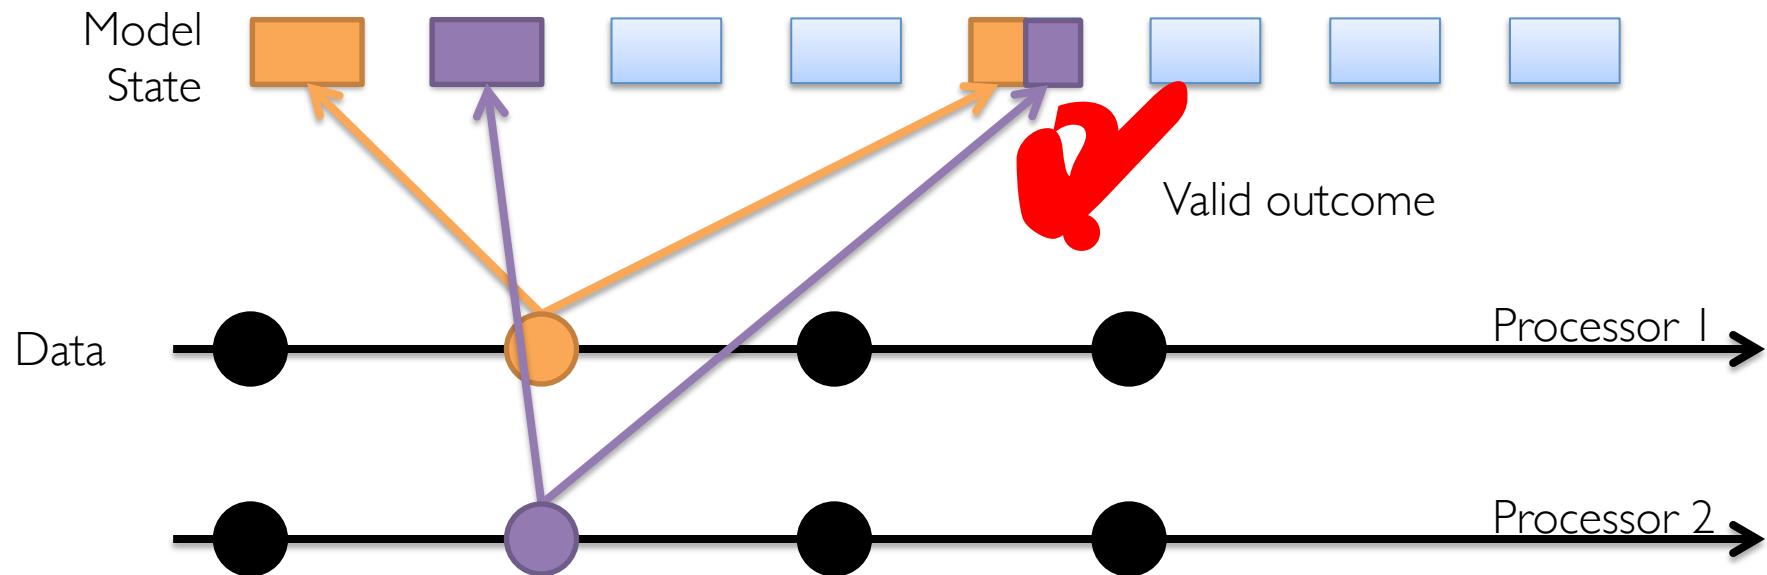
- Execute operations without blocking
Frequent case is fast
- Identify and resolve conflicts after they occur
Infrequent case with potentially costly resolution

Optimistic Concurrency Control



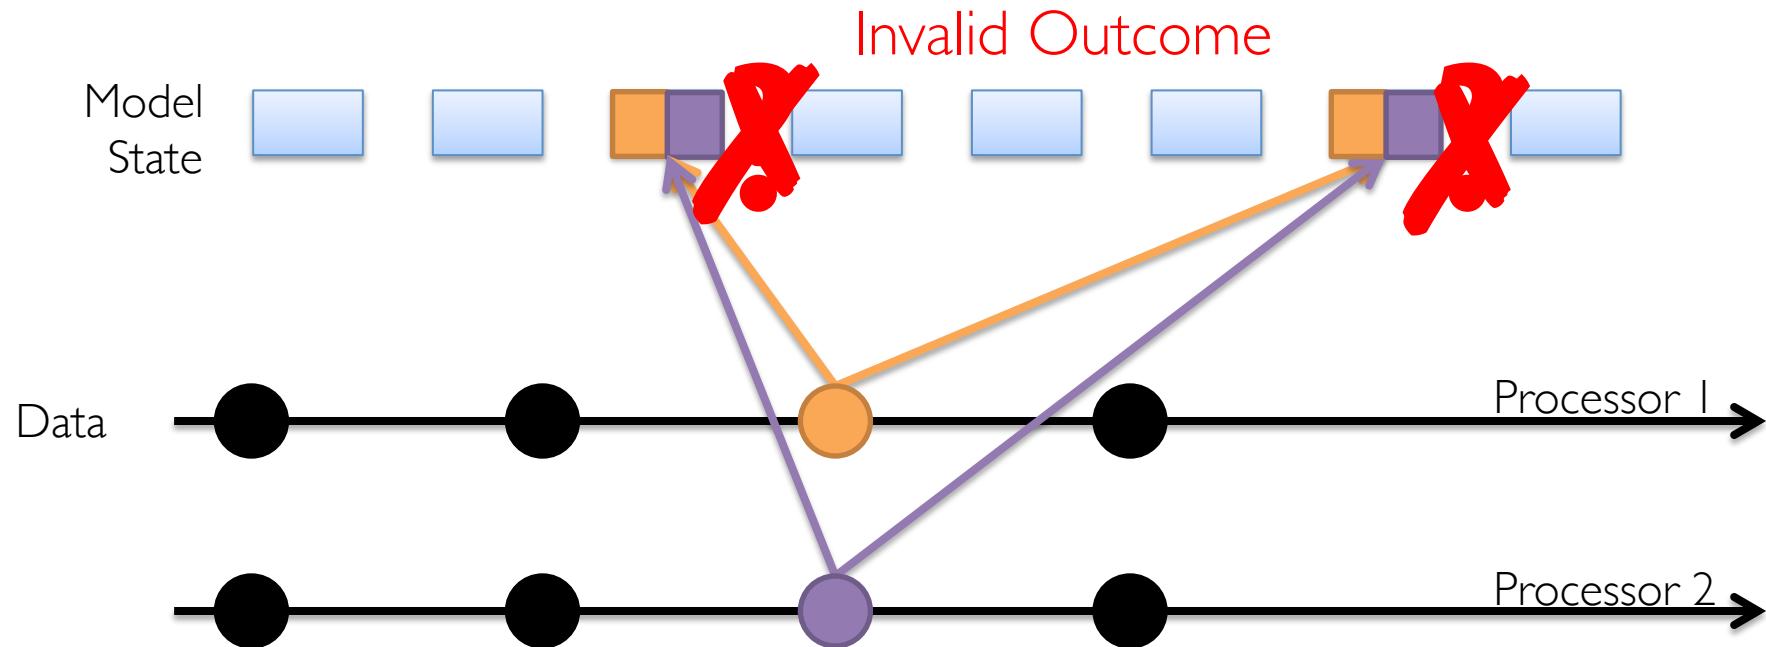
Allow computation to proceed without blocking.

Optimistic Concurrency Control



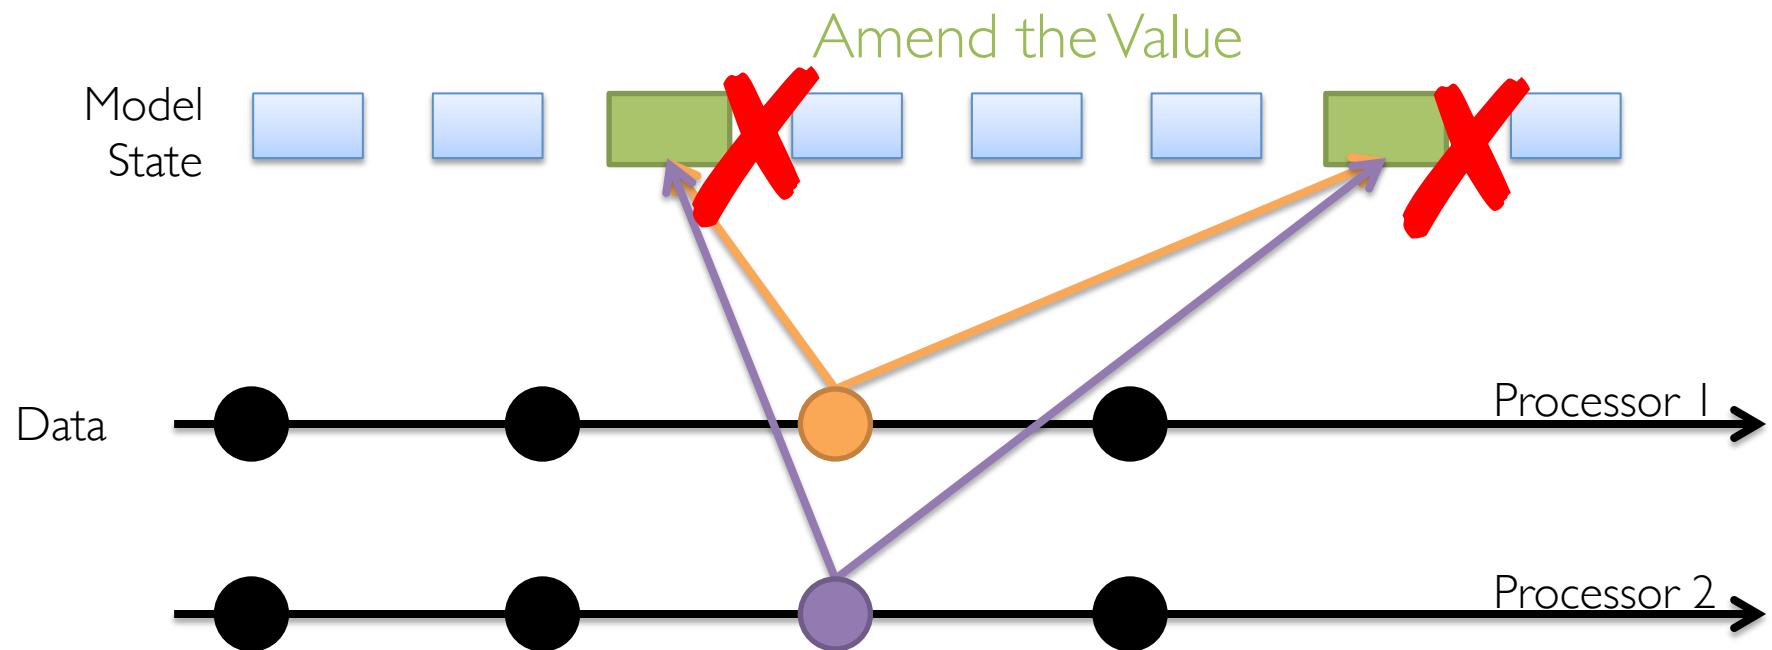
Validate potential conflicts.

Optimistic Concurrency Control



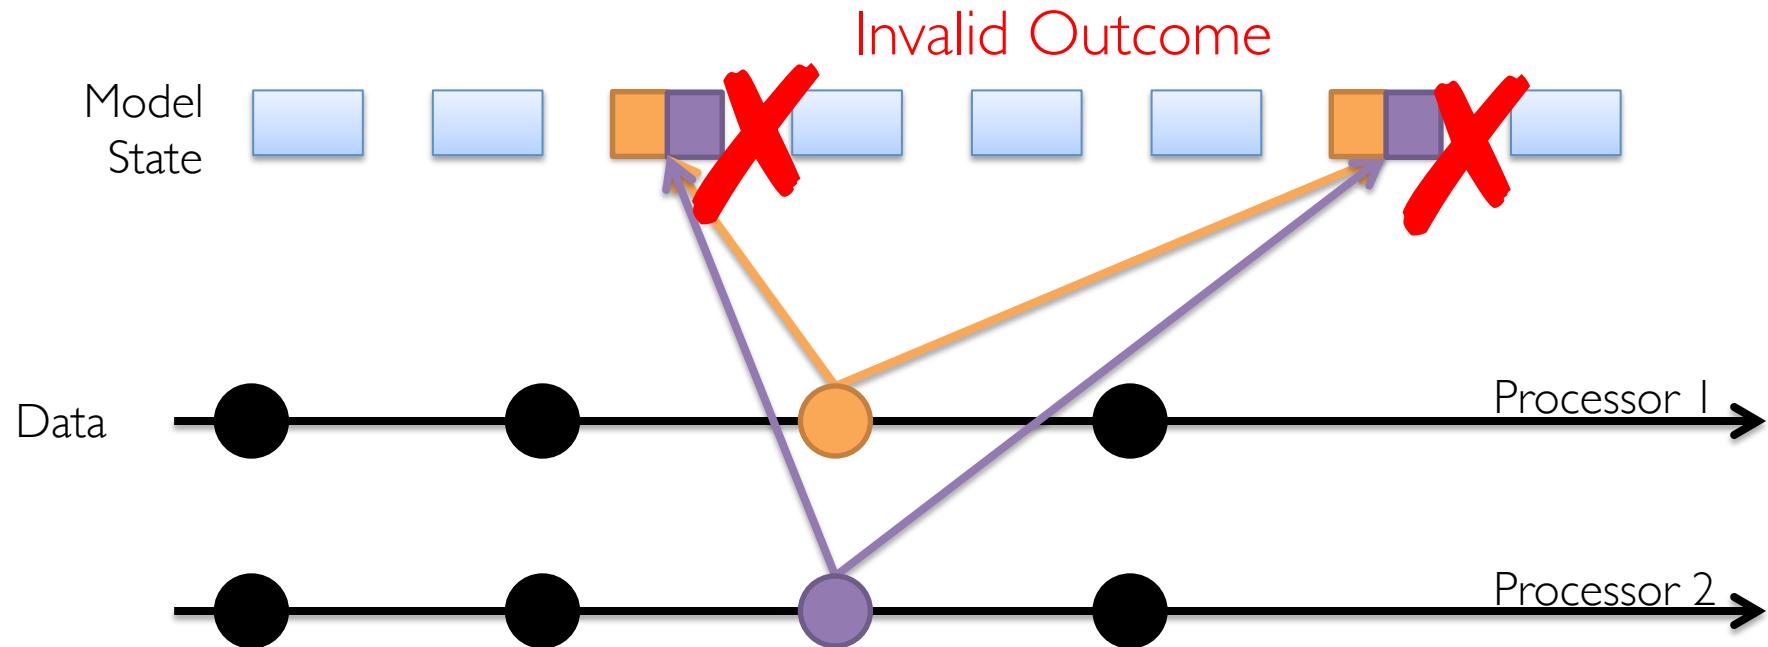
Validate potential conflicts.

Optimistic Concurrency Control



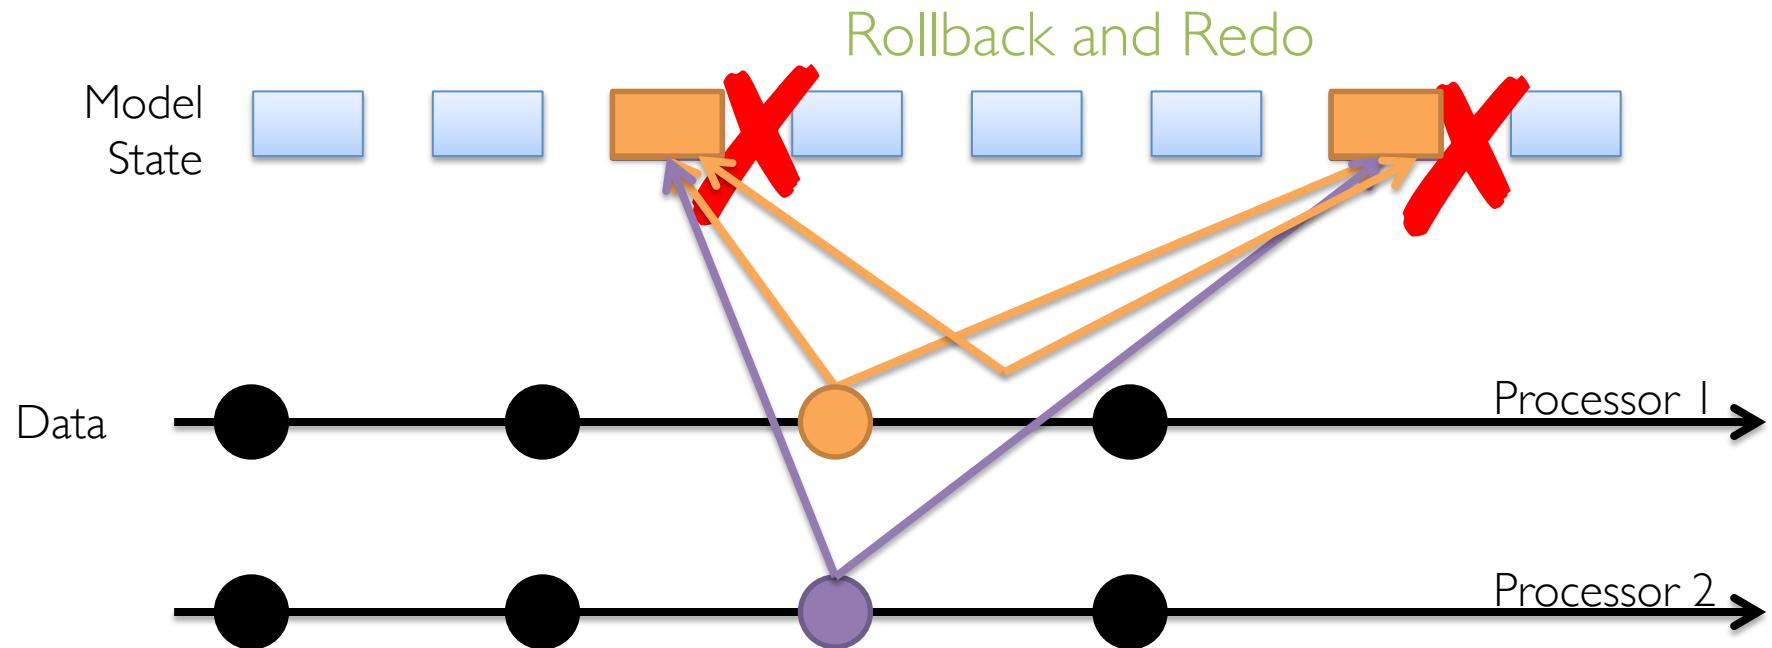
Take a compensating action.

Optimistic Concurrency Control



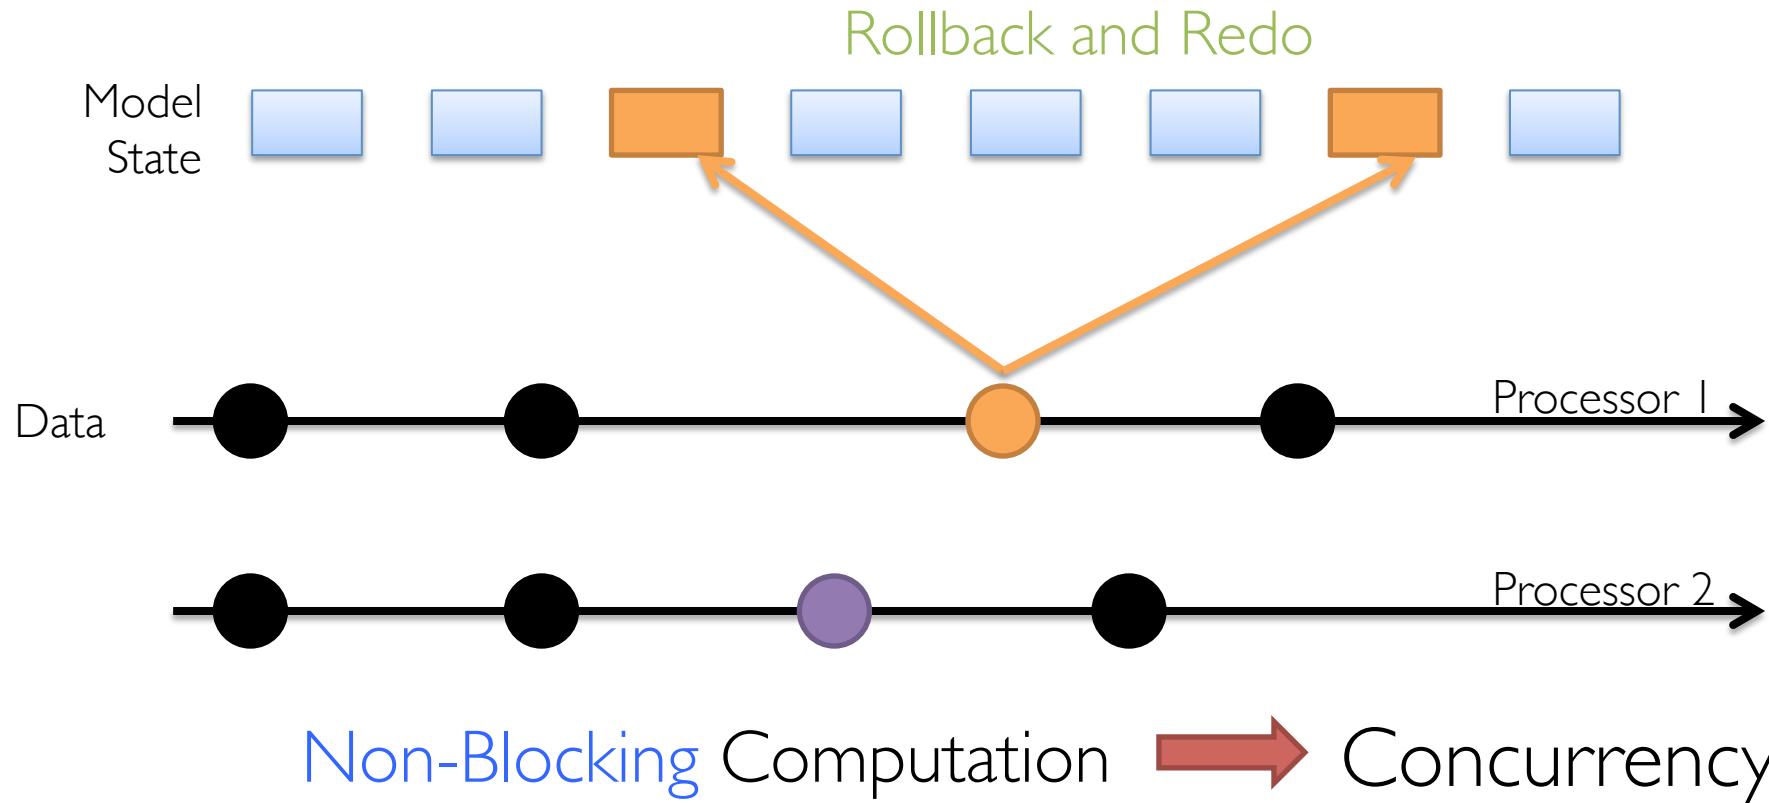
Validate potential conflicts.

Optimistic Concurrency Control



Take a compensating action.

Optimistic Concurrency Control



Requirements:

Fast → Validation: Identify Errors

Infrequent → Resolution: Correct Errors

Accuracy

Optimistic Concurrency Control for Bayesian Inference

Non-parametric Models [Pan et al., NIPS'13]:

- OCC DP-Means: Dirichlet Process Clustering
- OCC BP-Means: Beta Process Feature Learning

Conditional Sampling: (*In Progress*)

- Collapsed Gibbs LDA
- Retrospective Sampling for HDP

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

Start with DP Gaussian mixture model:

$$G \sim \text{DP}(\alpha, G_0 = \mathcal{N}(0, \rho I))$$

$$\phi_i \sim G$$

$$x_i \sim \mathcal{N}(\phi_i, \sigma I)$$

small variance limit $\sigma \rightarrow 0$

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

Start with DP Gaussian mixture model:

$$G \sim \text{DP}(\alpha, G_0 = \mathcal{N}(0, \rho I))$$

$$\phi_i \sim G$$

$$x_i \sim \mathcal{N}(\phi_i, \sigma I)$$

small variance limit $\sigma \rightarrow 0$ redefine α :

$$\alpha(\sigma) = \left(1 + \frac{\rho}{\sigma}\right)^{d/2} \exp\left(-\frac{\lambda}{2\sigma}\right)$$

Decreases
Rapidly

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

Corresponding Gibbs sampler conditionals:

$$P(\text{join } c) \propto n_{-i,c} \exp\left(-\frac{\|x_i - \mu_c\|^2}{2\sigma}\right)$$

$$P(\text{new}) \propto \exp\left(-\frac{1}{2\sigma} \left[\lambda + \frac{\sigma}{\rho + \sigma} \|x_i\|^2 \right]\right)$$

Taking the small variance limit $\sigma \rightarrow 0$

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

Gibbs updates become deterministic:

```
for  $i \in \{1, \dots, n\}$  do
     $c_i^{\min} = \arg \min_c \|x_i - \mu_c\|$ 
    if  $\|x_i - \mu_{c_i^{\min}}\| < \lambda$  then join  $c_i^{\min}$ 
    else create new cluster at  $\mu_{k+1} = x_i$ 
```

Taking the small variance limit $\sigma \rightarrow 0$

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

Gibbs updates become deterministic:

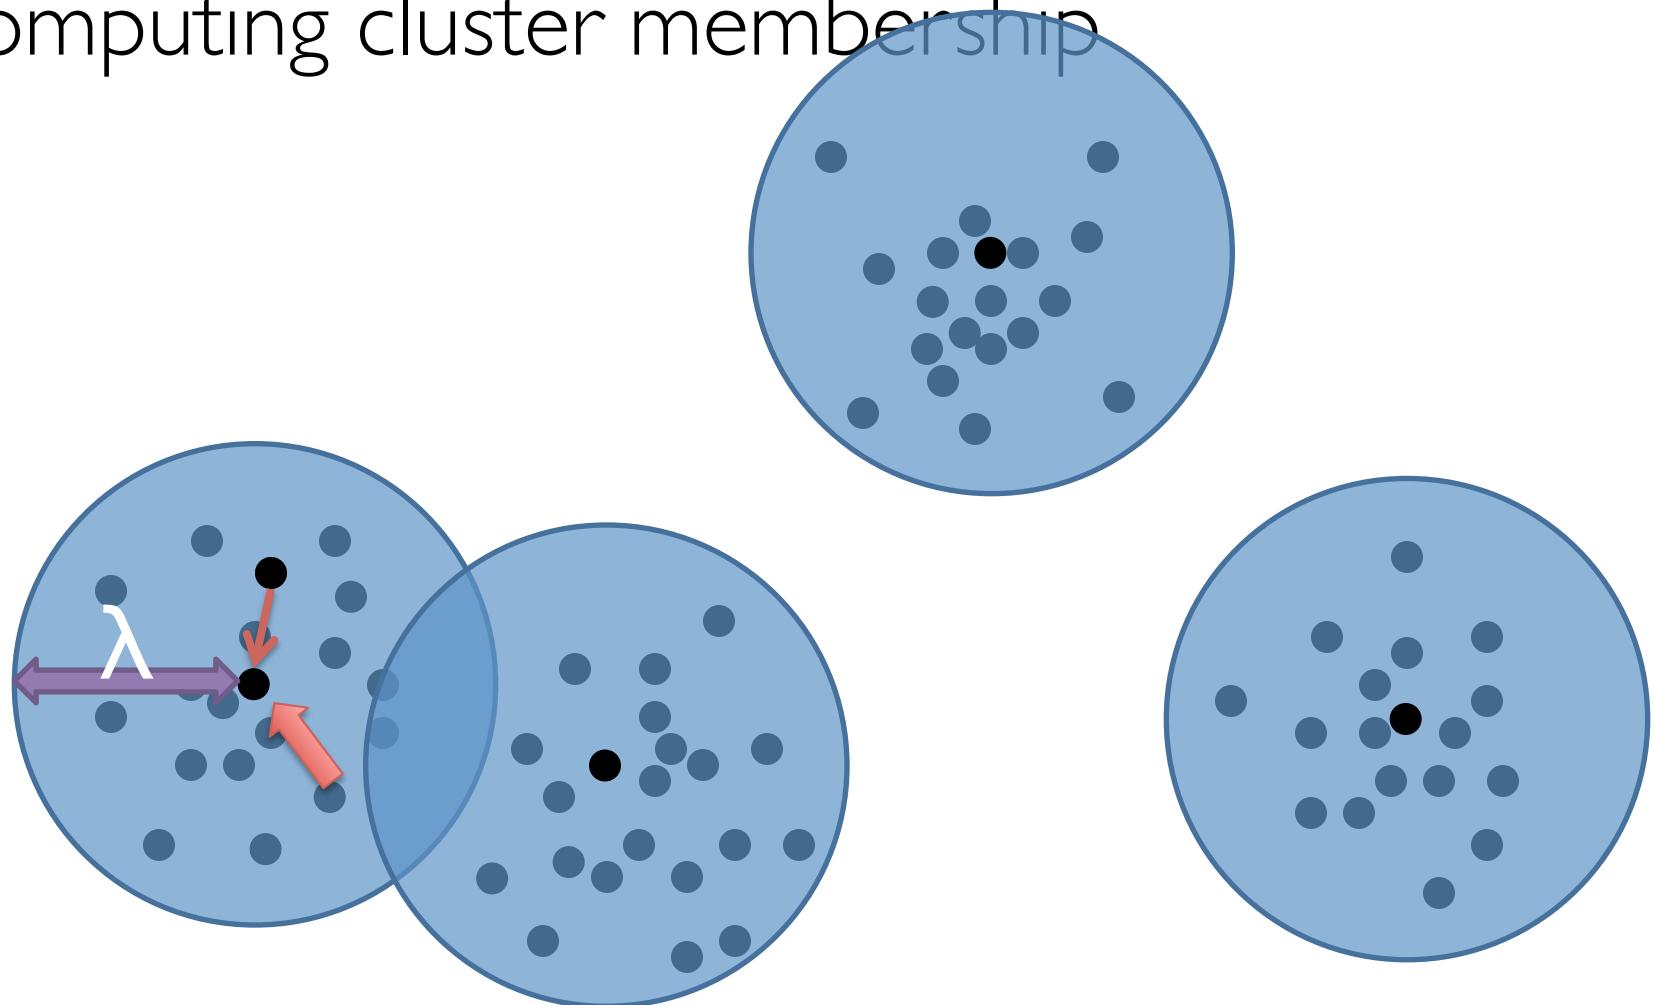
```
for  $i \in \{1, \dots, n\}$  do
     $c_i^{\min} = \arg \min_c \|x_i - \mu_c\|$ 
    if  $\|x_i - \mu_{c_i^{\min}}\| < \lambda$  then join  $c_i^{\min}$ 
    else create new cluster at  $\mu_{k+1} = x_i$ 
```

```
for  $c$  in clusters do  $\mu_c \leftarrow \frac{1}{n_c} \sum_{x \in c} x$ 
```

DP-Means Algorithm

[Kulis and Jordan, ICML'12]

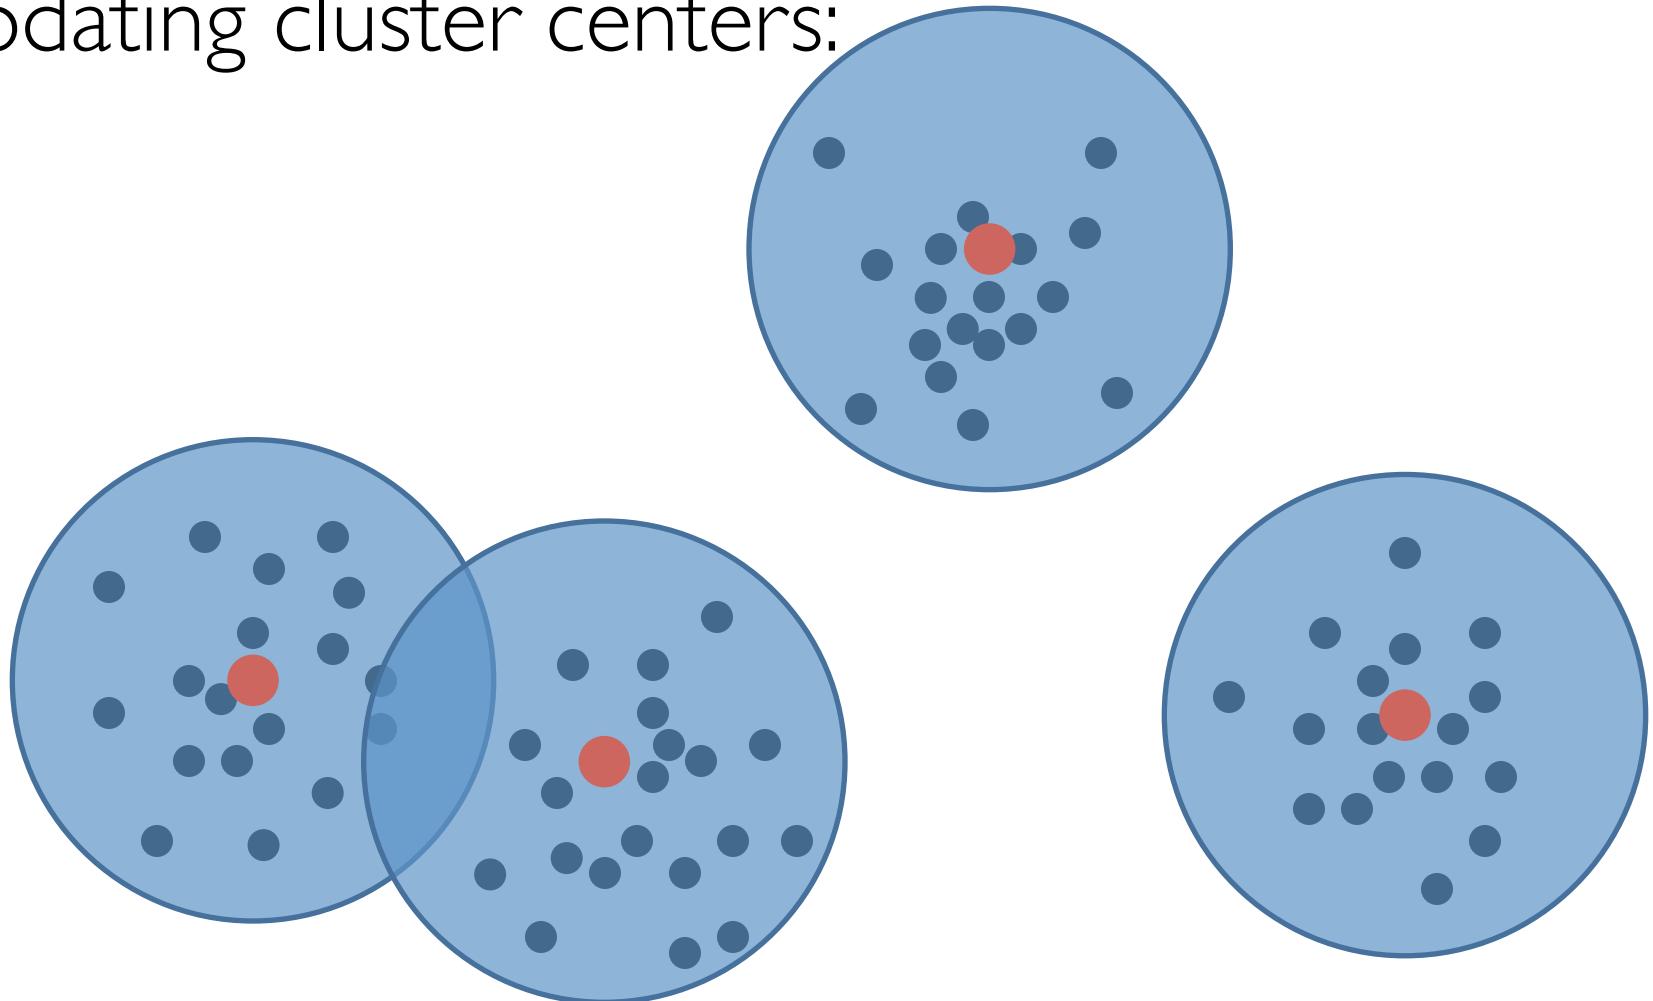
Computing cluster membership



DP-Means Algorithm

[Kulis and Jordan, ICML'12]

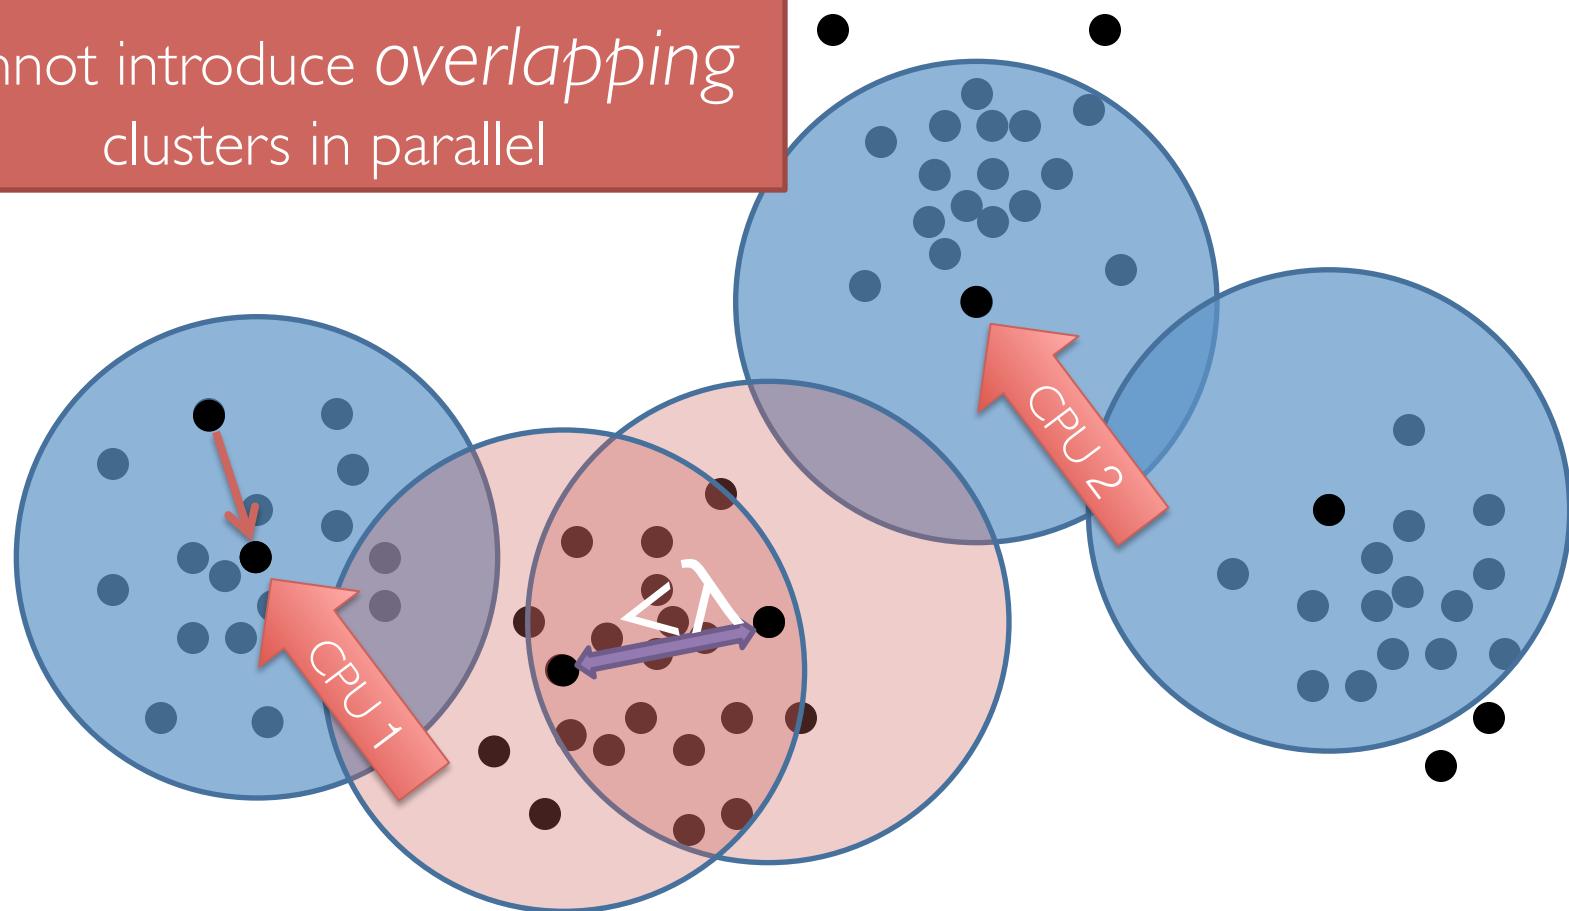
Updating cluster centers:



DP-Means Parallel Execution

Computing cluster membership in parallel:

Cannot introduce *overlapping* clusters in parallel



Optimistic Concurrency Control for Parallel DP-Means

Optimistic Assumption

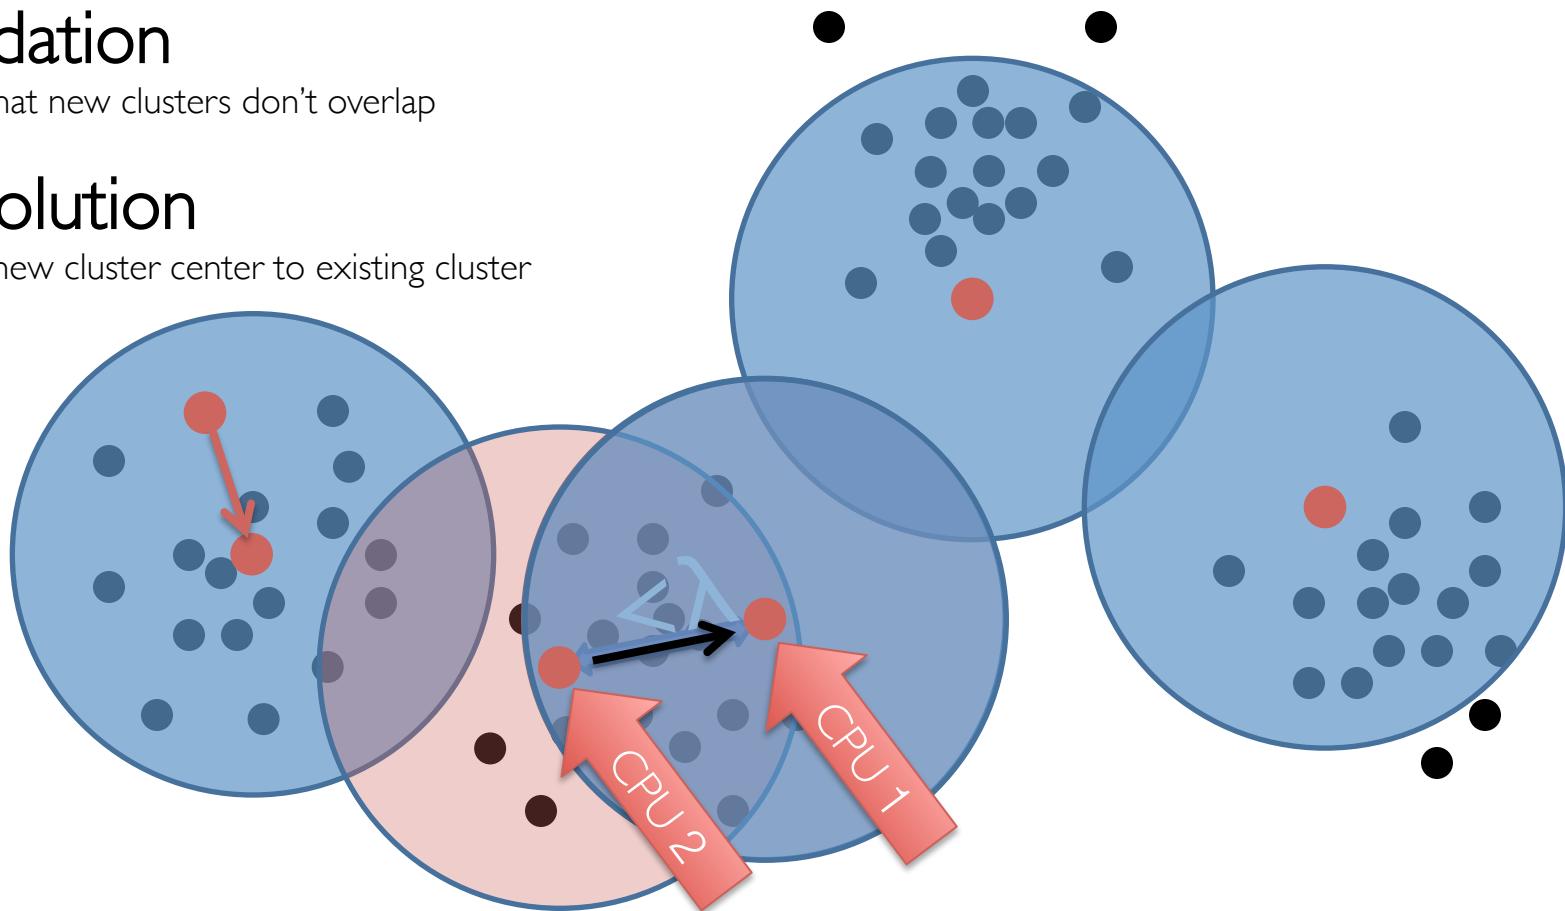
No new cluster created nearby

Validation

Verify that new clusters don't overlap

Resolution

Assign new cluster center to existing cluster



OCC DP-means

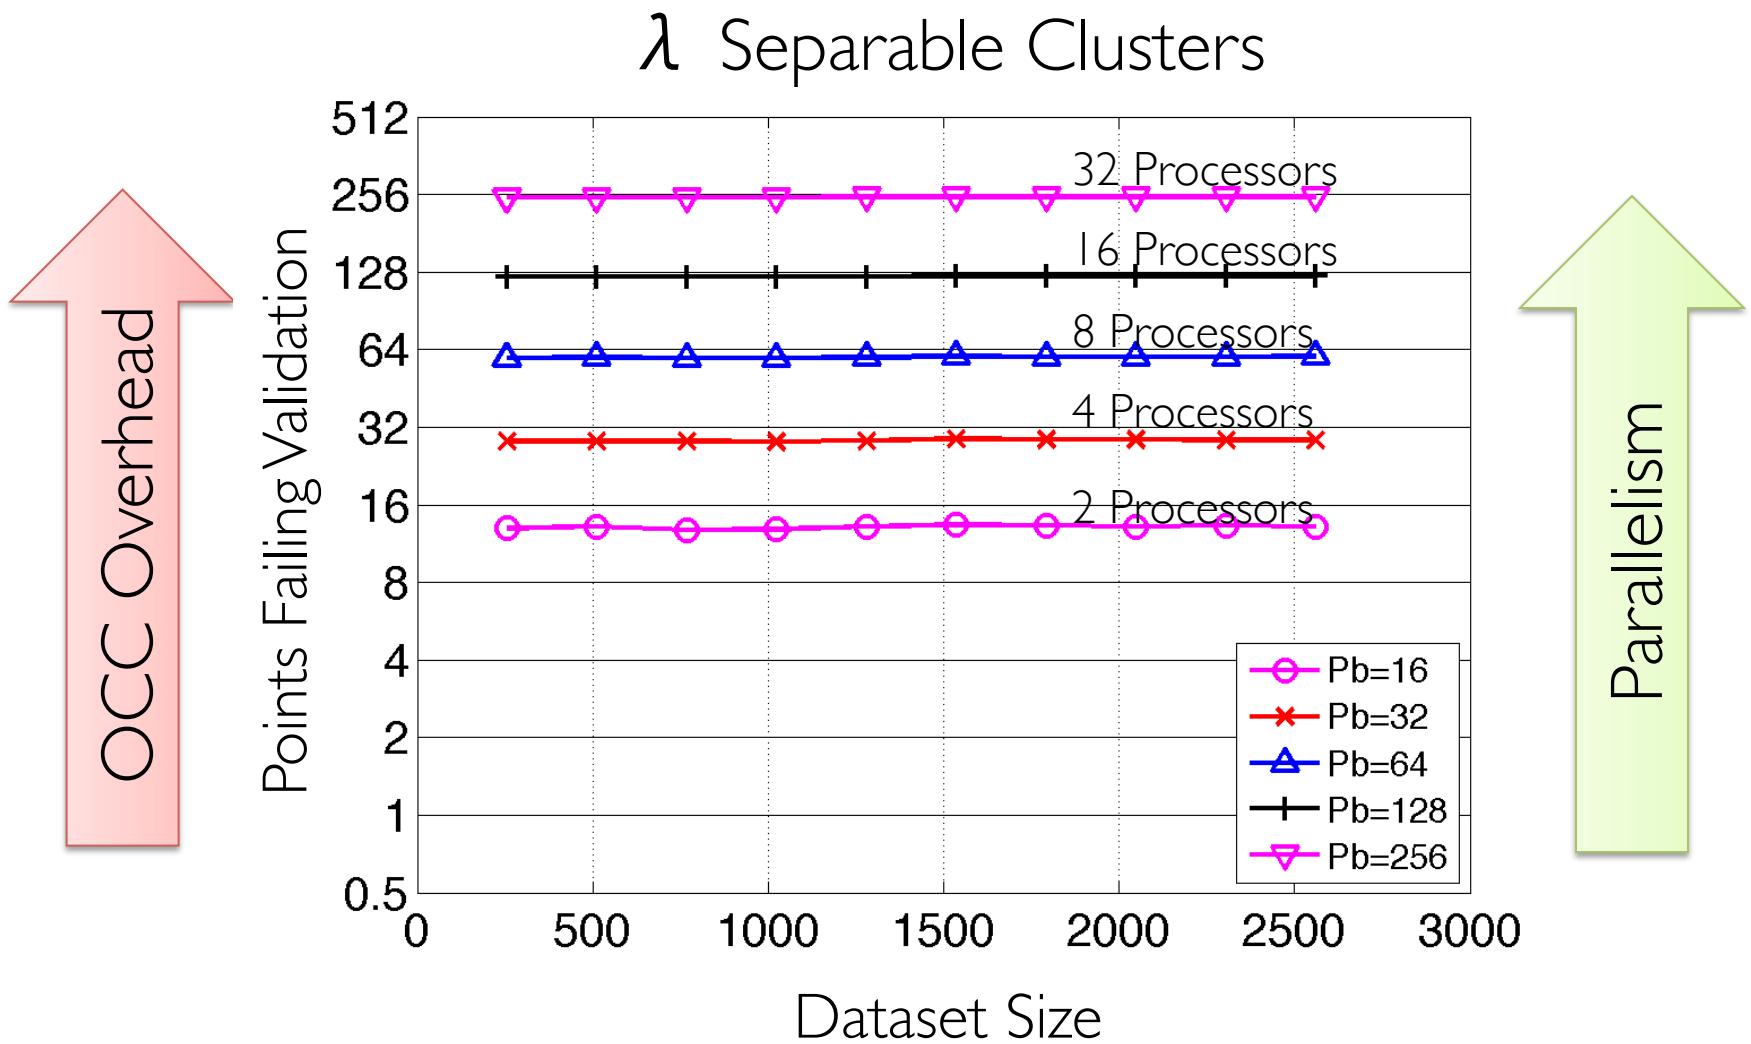
Theorem: OCC DP-means is serializable
and therefore preserves theoretical
properties of DP-means.

Correctness

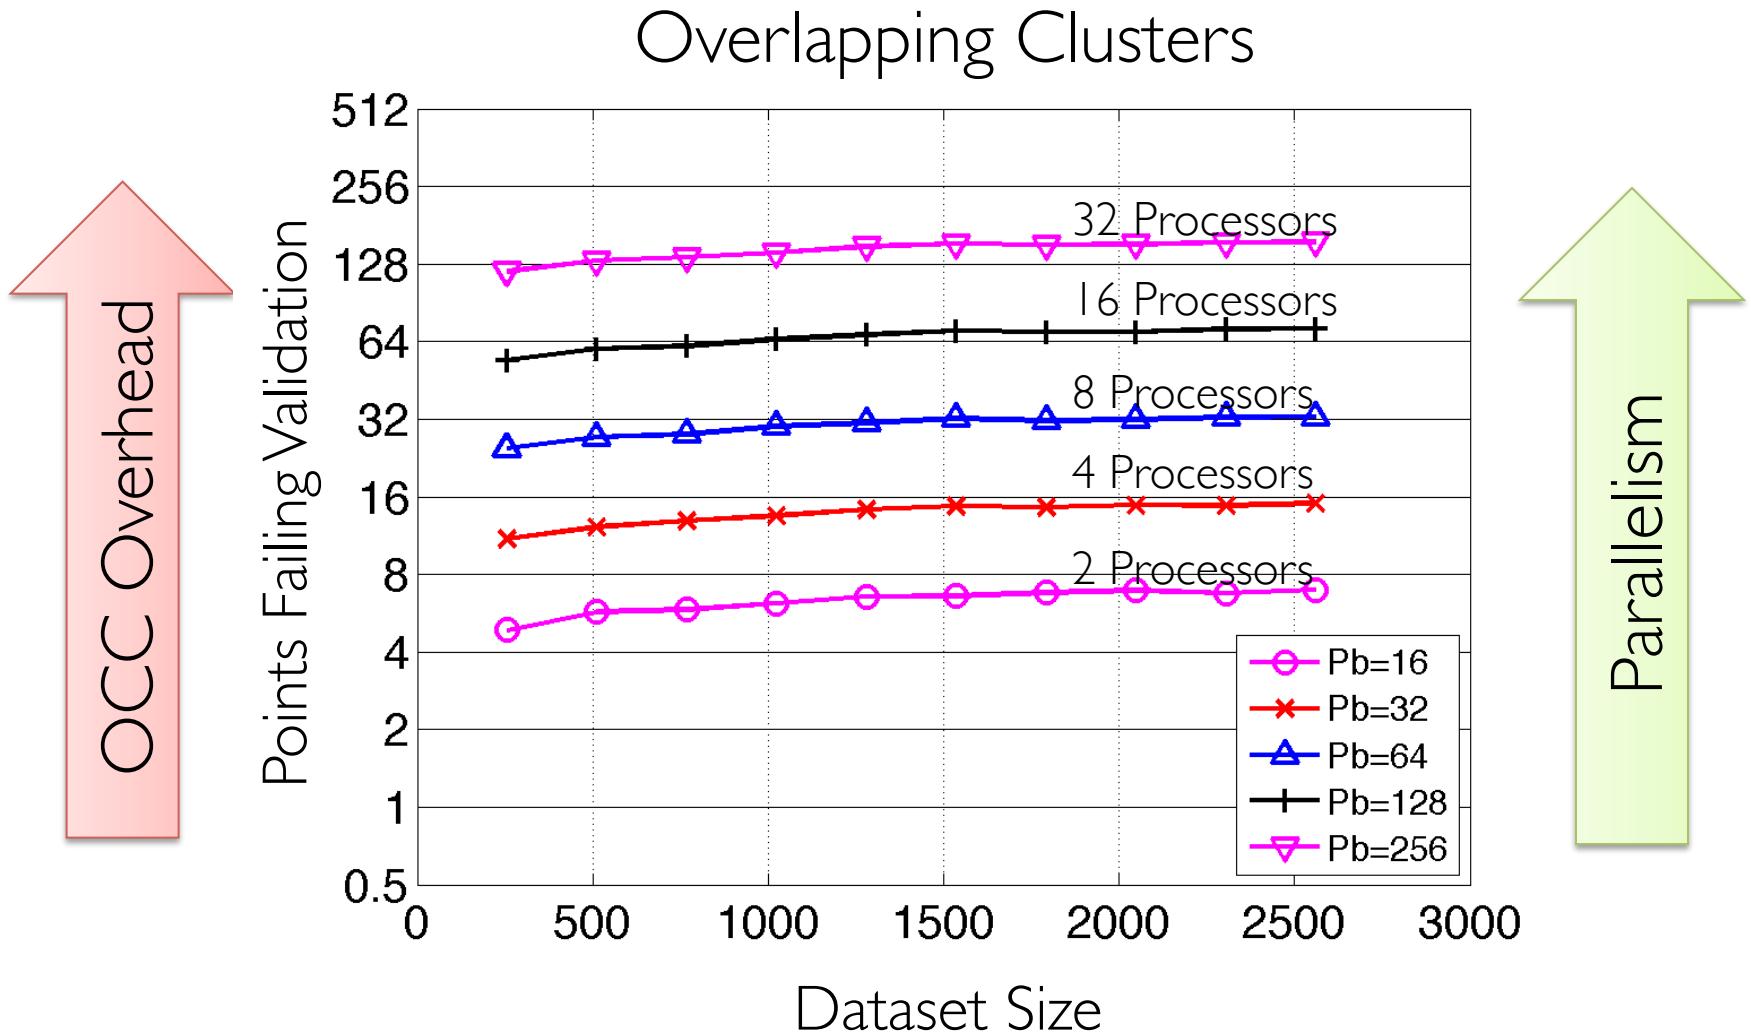
Theorem: Assuming well spaced clusters
the expected overhead of OCC DP-
means does not depend on data size.

Concurrency

Empirical Validation Failure Rate

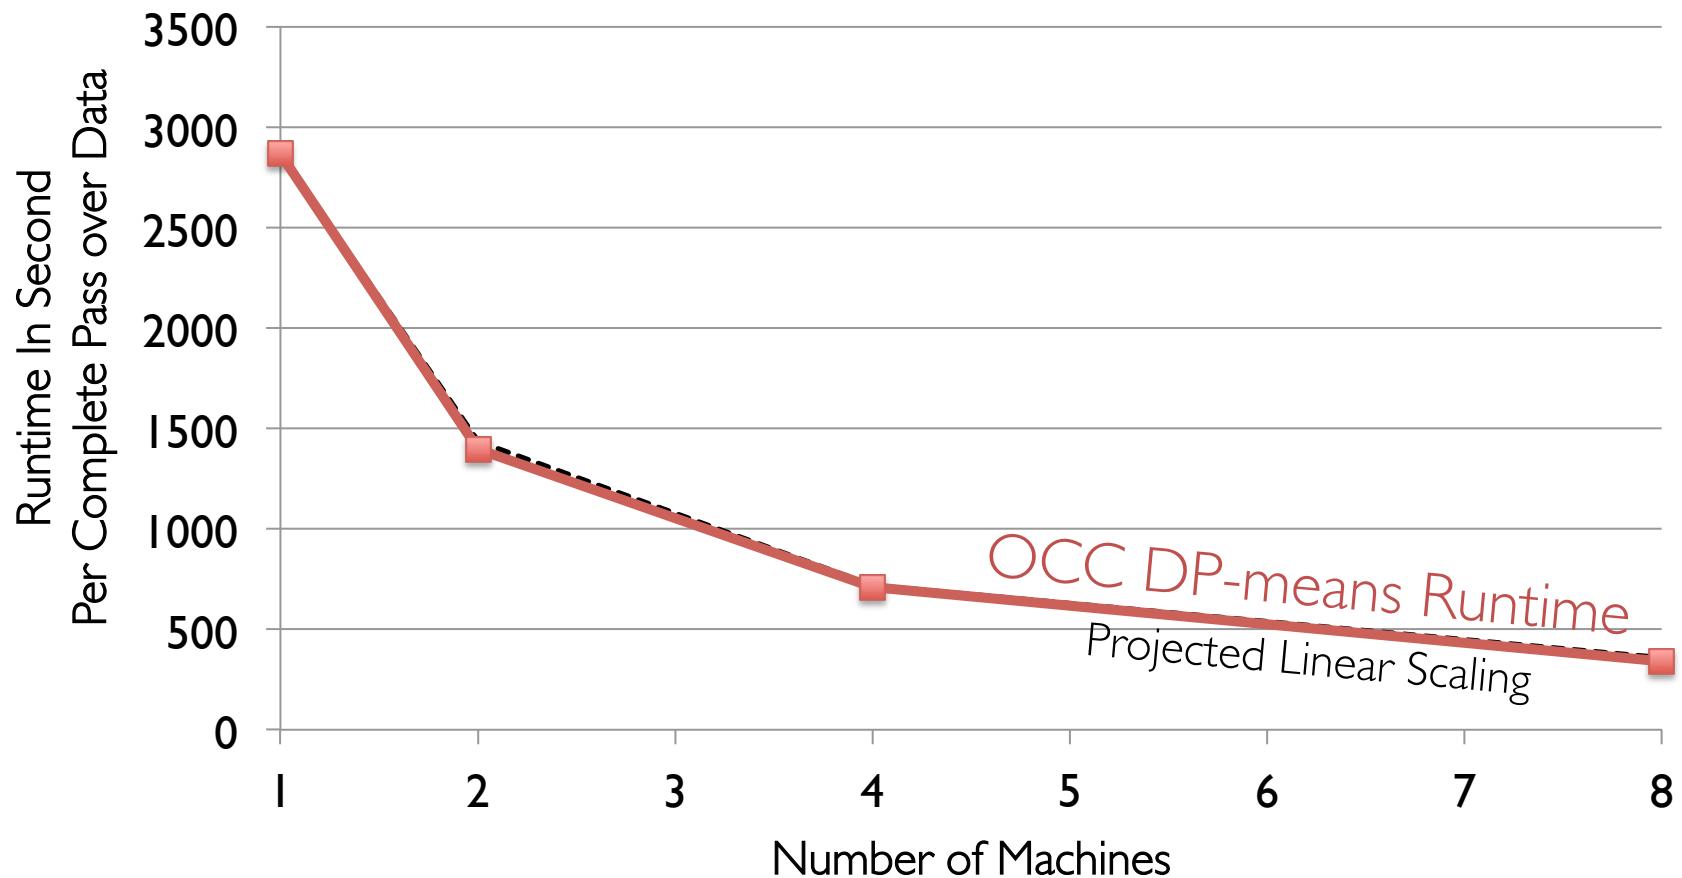


Empirical Validation Failure Rate



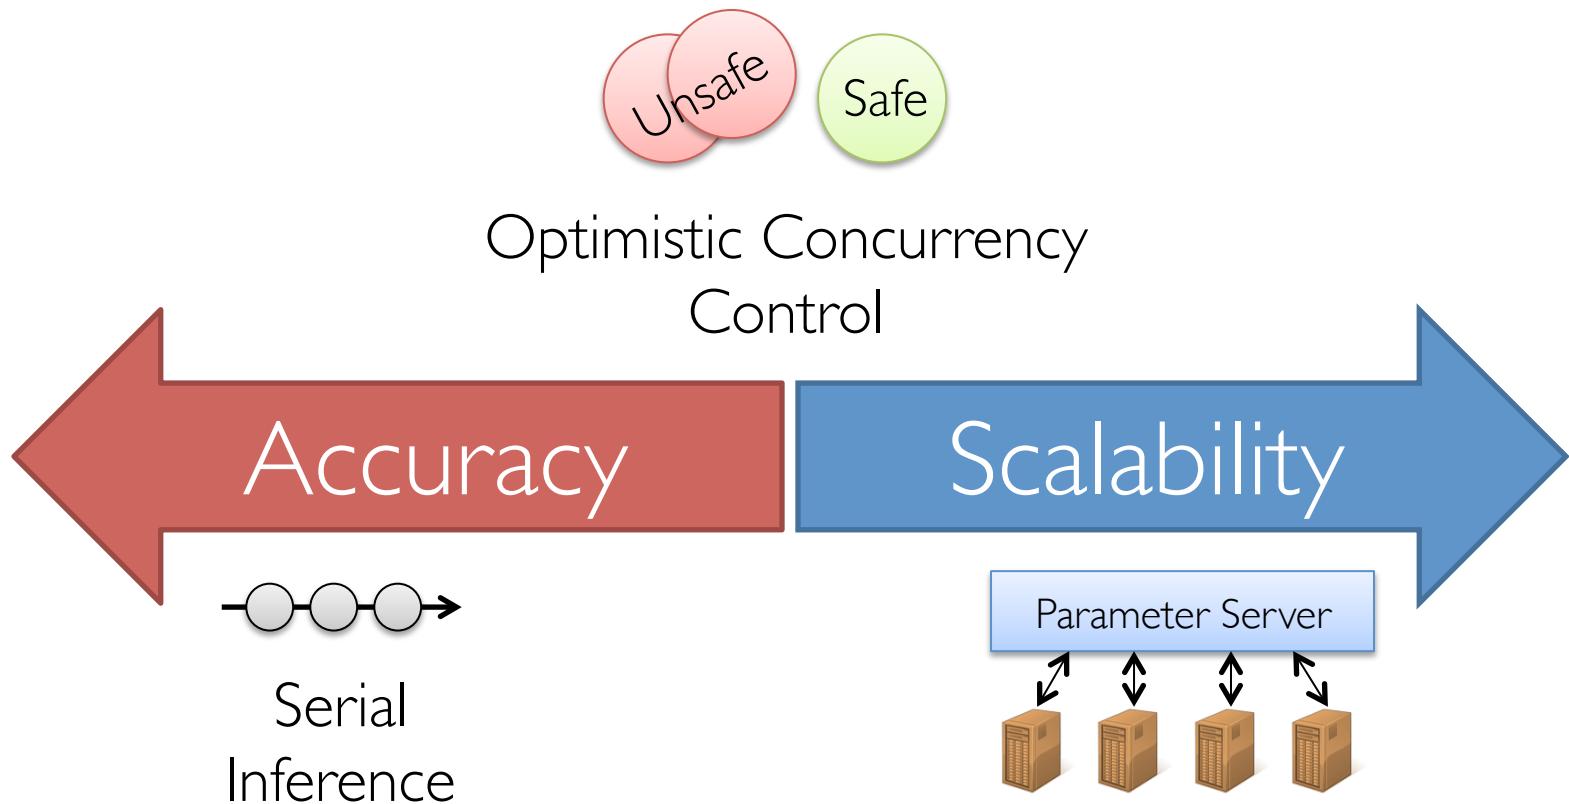
Distributed Evaluation Amazon EC2

~140 million data points; 1, 2, 4, 8 machines



2x #machines
≈ ½x runtime

Optimistic Future for Optimistic Concurrency Control

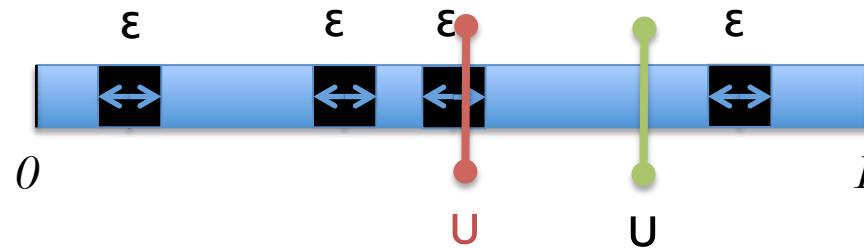


Thank You Questions?

Joseph E. Gonzalez
jegonzal@eecs.berkeley.edu

OCC Collapsed Gibbs for LDA

Maintain epsilon intervals on the conditional:

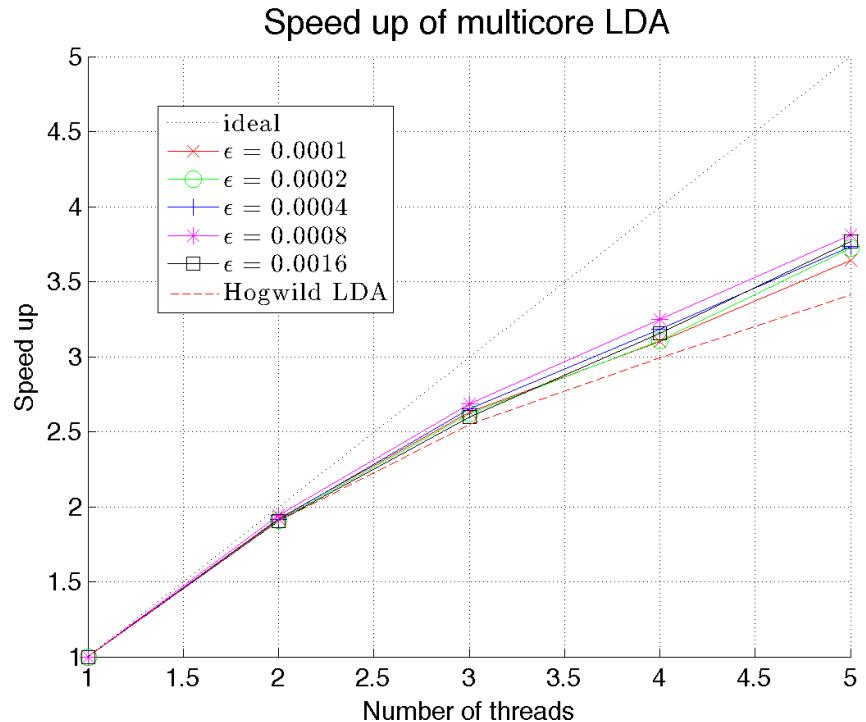
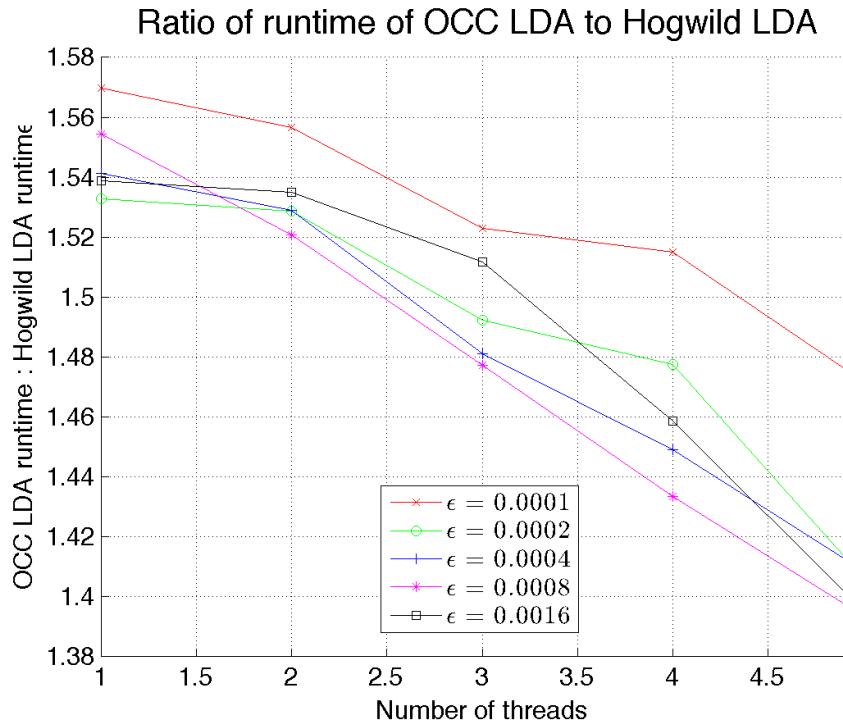


System ensures conditionals are ϵ -accurate

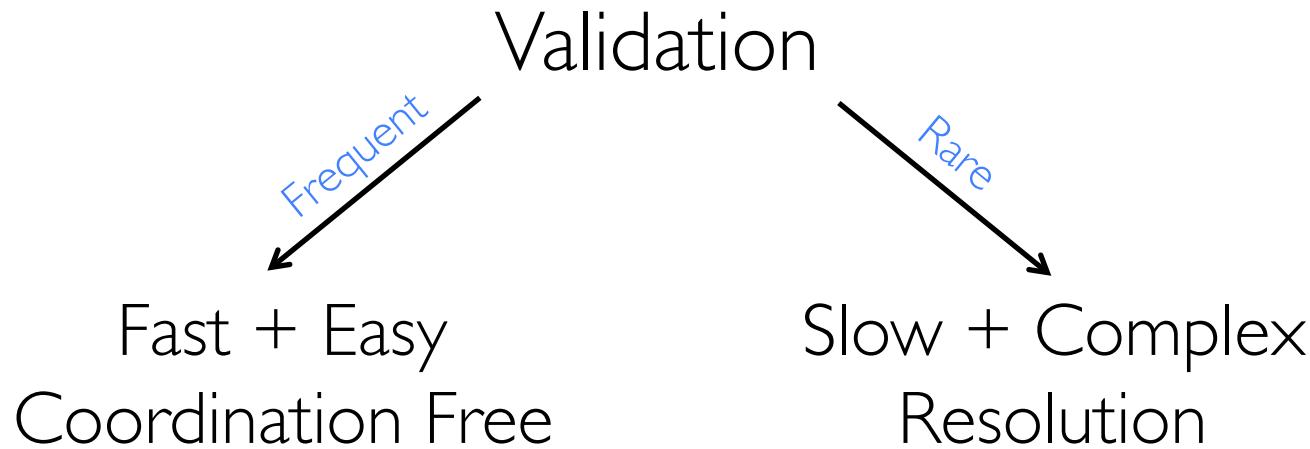
Validate: Accept draws that land outside interval

Resolution: Serially resample rejected tokens

OCC Collapsed Gibbs for LDA



Optimism for Optimistic Concurrency Control



Enable decades of work in *serial Bayesian inference* to be extended to the parallel setting.