

Thesis Defense

Parallel and Distributed Systems for Probabilistic Reasoning

Joseph E. Gonzalez



Thesis Committee:



Carlos Guestrin
University of
Washington & CMU



Guy Blelloch
CMU



David O'Hallaron
CMU



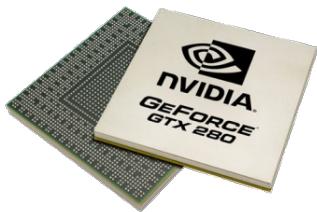
Alex Smola
CMU & Google



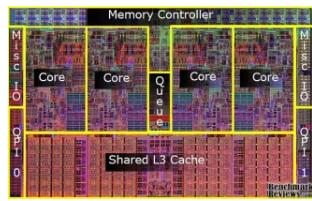
Jeff Bilmes
University of
Washington

*The foundations of computation
have changed ...*

New Parallel and Distributed Platforms



GPUs



Multicore



Clusters



Single Chip
Cloud Computers



Clouds

- New Opportunities
 - Increased processing and storage
- New Challenges
 - Parallel algorithm design and implementation

*The scale of
machine learning problems
is exploding ...*

The Age of Big Data



28 Million
Wikipedia Pages



1 Billion
Facebook Users



6 Billion
Flickr Photos



72 Hours a Minute
YouTube

The New York Times

SundayReview

WORLD U.S. N.Y. / REGION BUSINESS TEC

NEWS ANALYSIS

The Age of Big Data

By STEVE LOHR

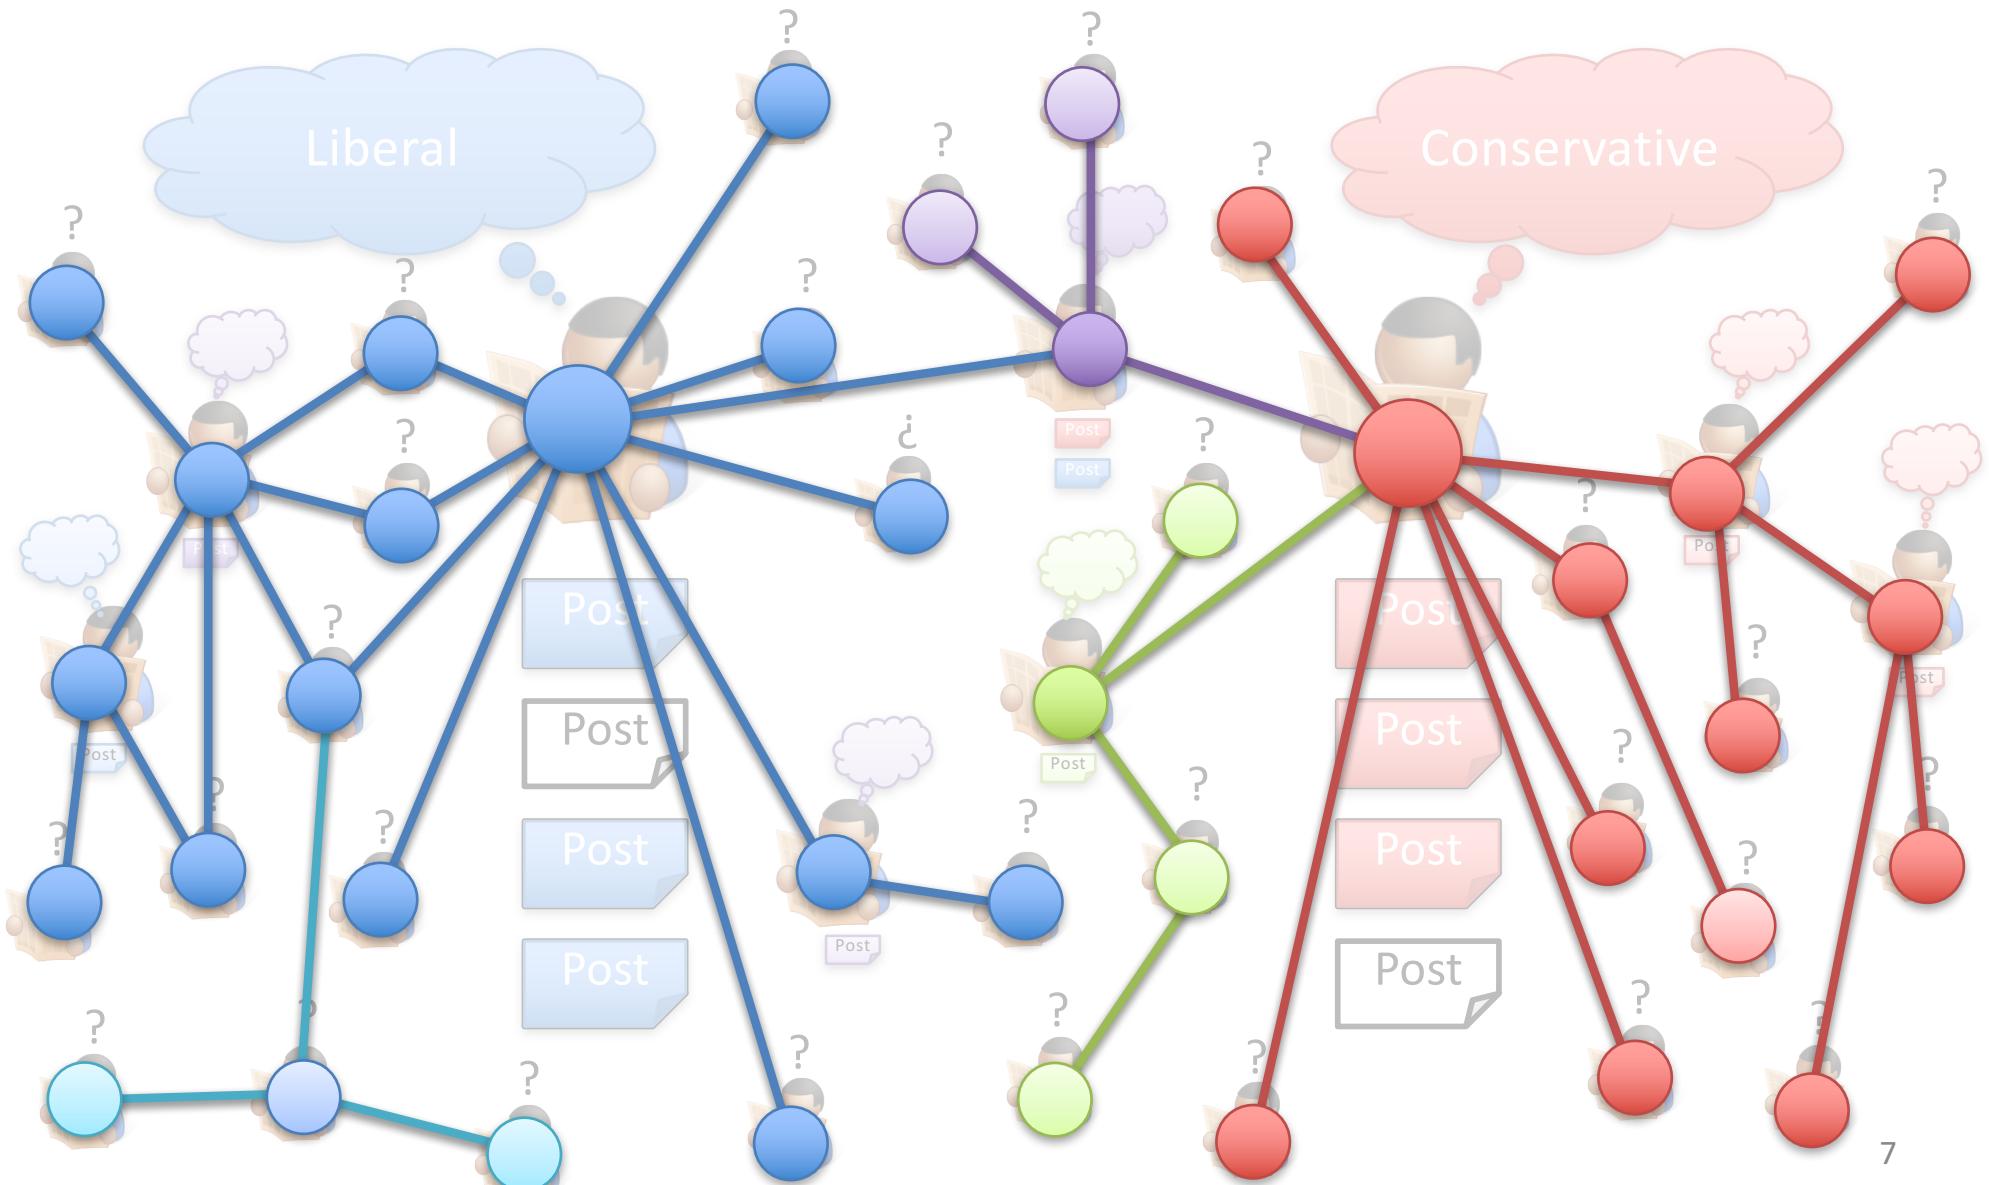
Published: February 11, 2012

“...growing at 50 percent a year...”

“... data a new class of economic asset,
like currency or gold.”

*Massive data provides
opportunities for
structured models...*

Example: Estimate Political Bias





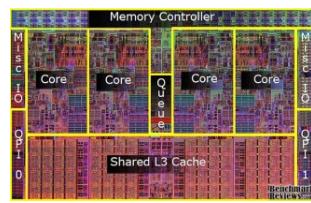
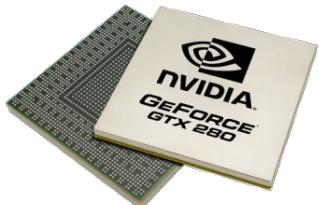
Massive Structured Problems

Thesis:

*Parallel and Distributed Systems
for Probabilistic Reasoning*



Advances Parallel Hardware



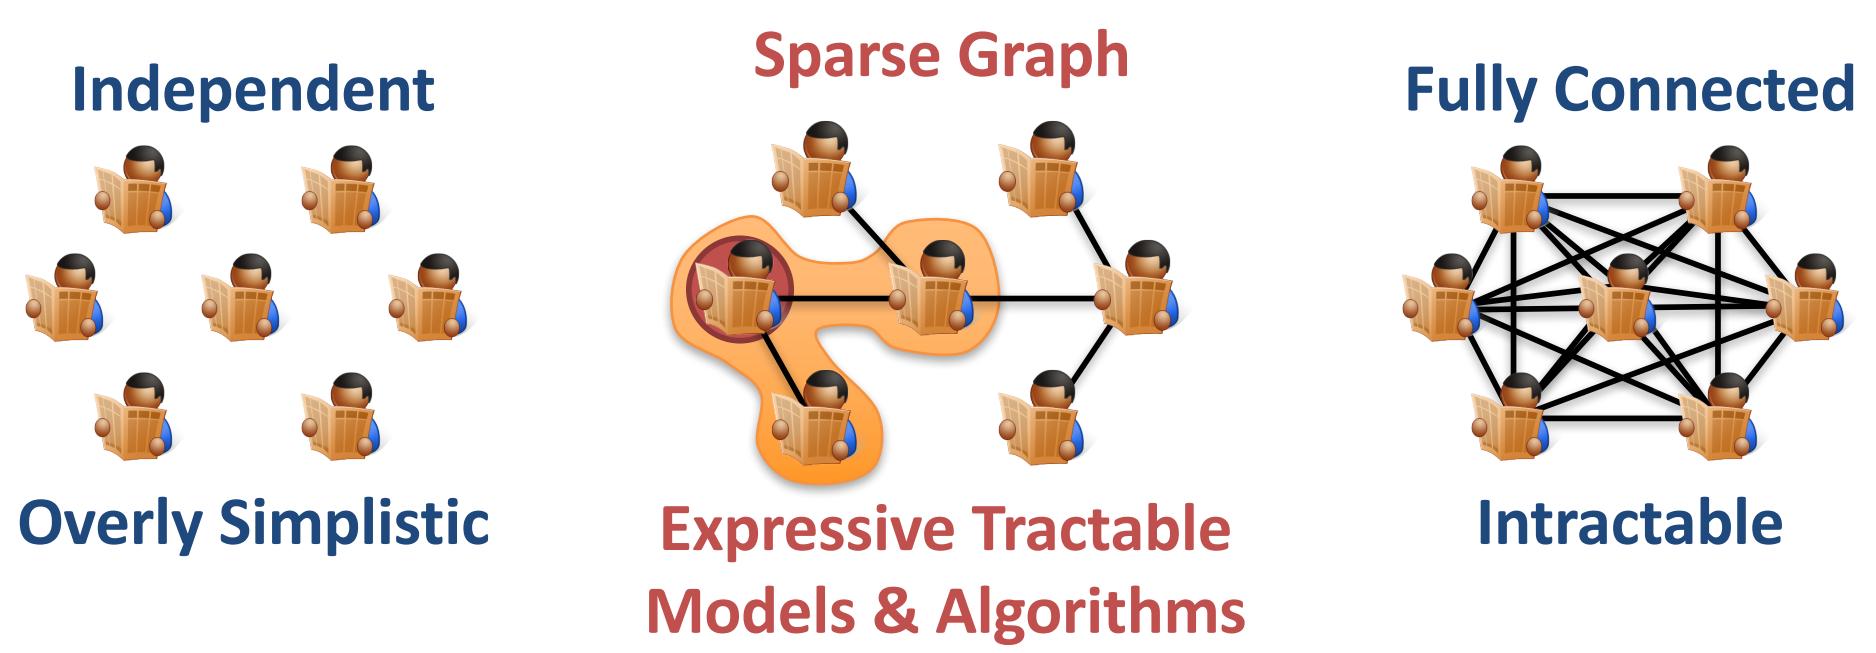
Thesis Statement: GrAD Methodology

Efficient parallel and distributed systems for probabilistic reasoning:

1. **Graphically** decompose *computational* and *statistical* dependencies
2. **Asynchronously** schedule computation
3. **Dynamically** identify and *prioritize* computation along critical paths

GrAD Methodology: Graphical

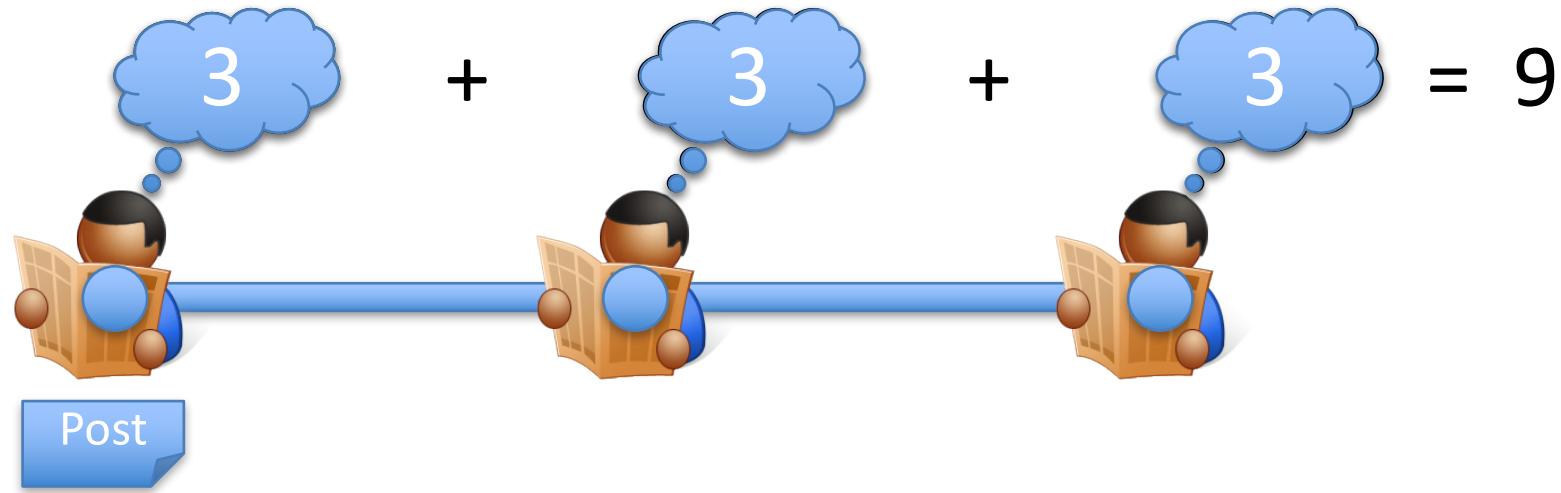
- Factor **statistical** and **computational** dependencies



- Improves **computational** and **statistical** efficiency
- Increases **parallelism**

Synchronous vs. Asynchronous

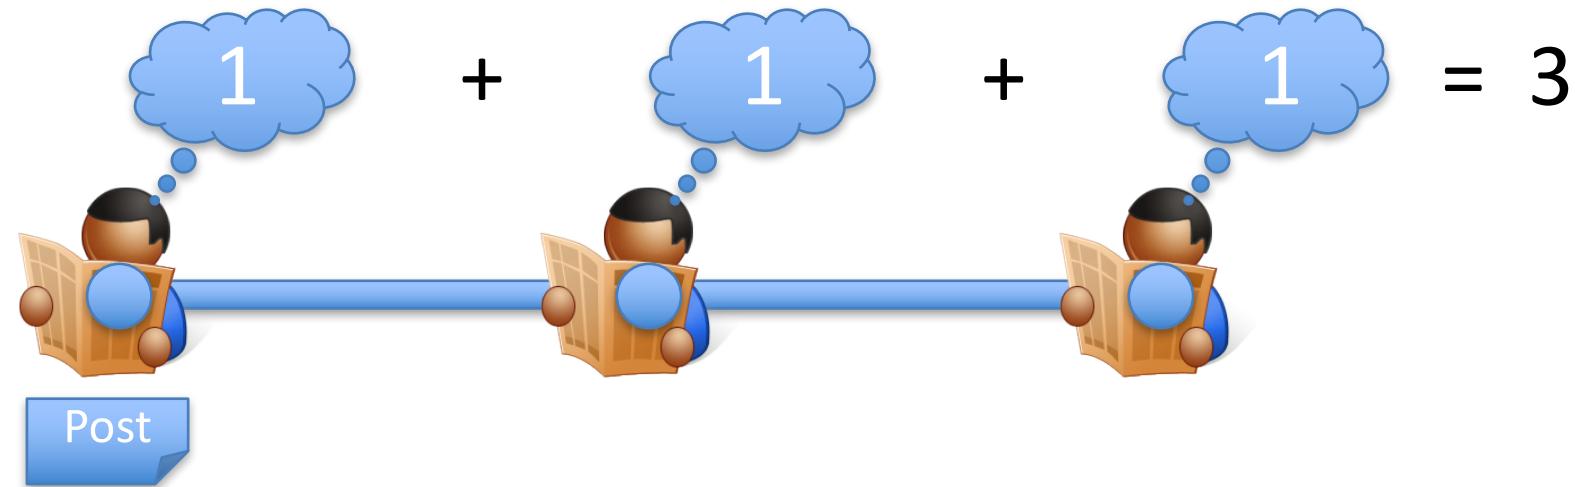
- **Synchronous:** *compute everything in parallel*



- Highly **parallel** – Maximum independent work
- Highly **inefficient** – Many wasted cycles

GrAD Methodology: Asynchronous

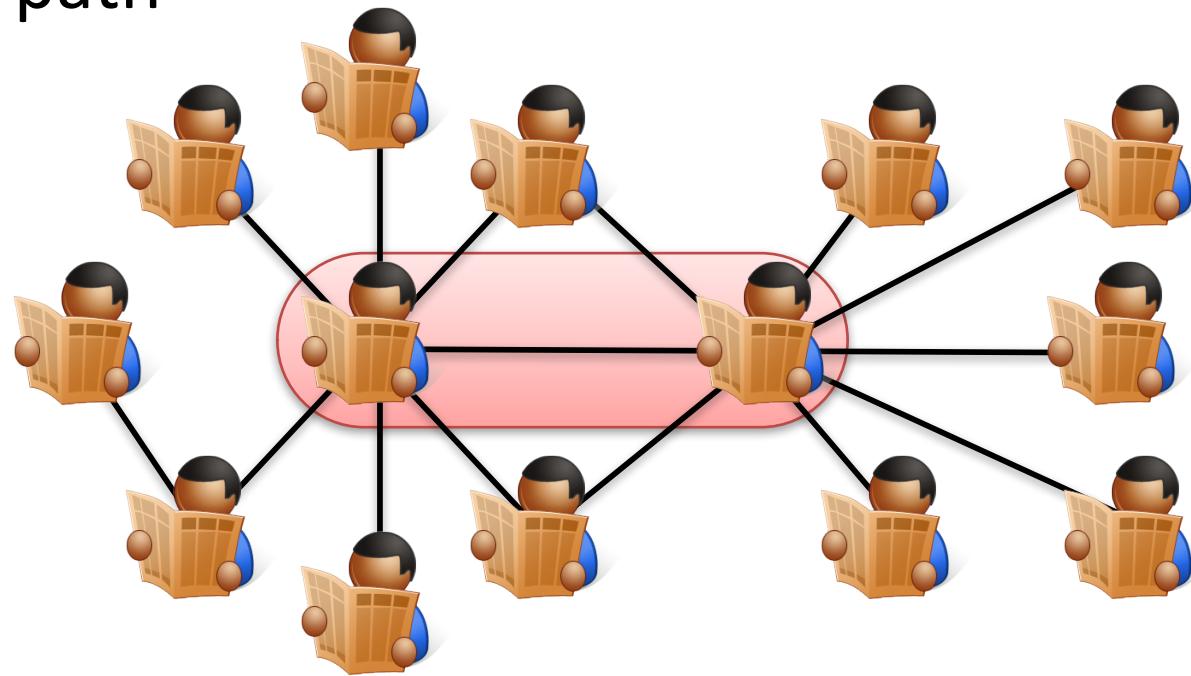
- Trigger computation as **new information arrives**



- Capture the flow of information:
 - More efficiently use network and processor resources
 - Guarantee algorithm correctness

GrAD Methodology: Dynamic

- Dynamically **identify** and **prioritize** computation along the critical path



- Focus computational resources where most effective:
 - Accelerated convergence
 - Increased work efficiency

We apply the GrAD methodology to

1. Probabilistic Graphical Models

2. **Parallel and Distributed Algorithms
for Probabilistic Inference**

3. **GraphLab & PowerGraph**



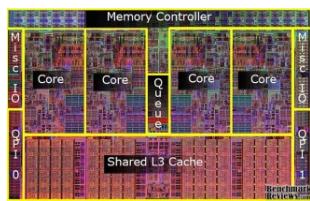
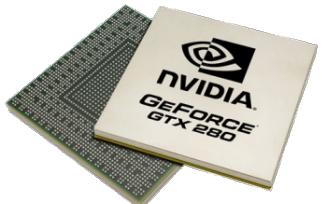
Massive Structured Problems

Probabilistic Graphical Models

**Parallel and Distributed Algorithms
for Probabilistic Inference**

GraphLab & PowerGraph

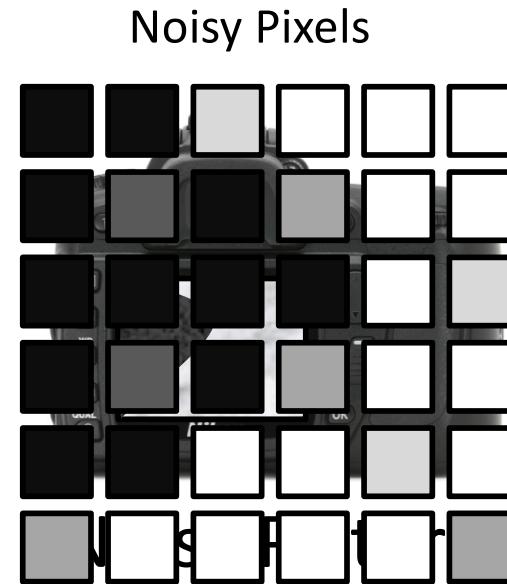
Advances Parallel Hardware



Encode Probabilistic Structure



True Image



Random Variables

True *unobserved* values

Dependency Graph:

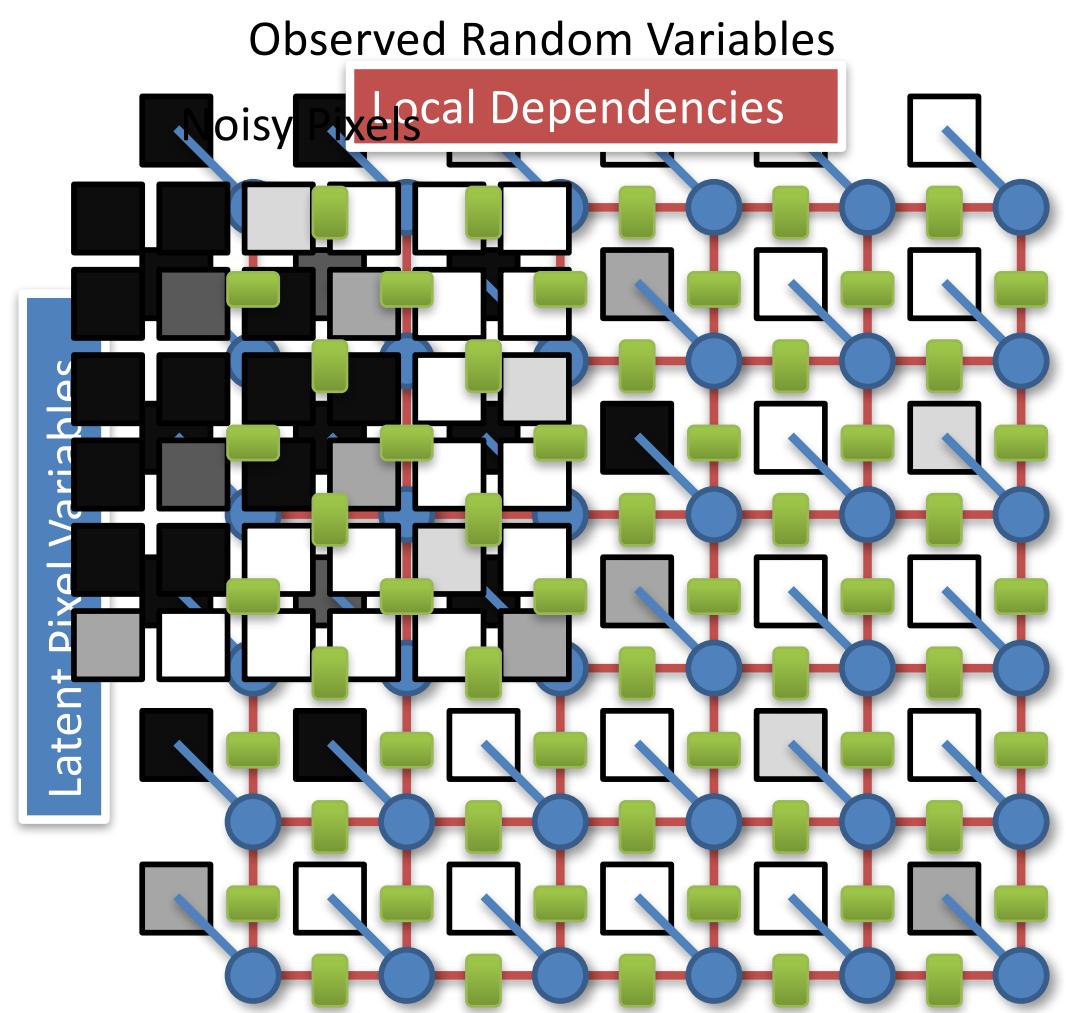
Represent dependencies

Parameters:

Characterize probabilities

$$P(X_1, \dots, X_n; \theta) \propto$$

Joint Probability

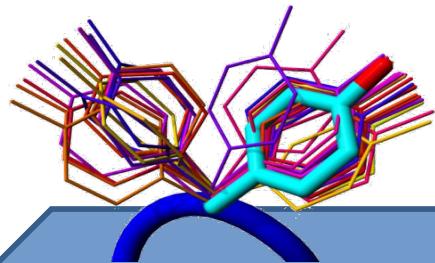


$$\prod_{(u,v) \in E} f(X_u, X_v; \theta_{u,v})$$

Graph Factors

Graphical models provide a common representation

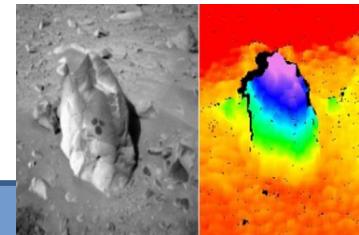
Protein Structure
Prediction



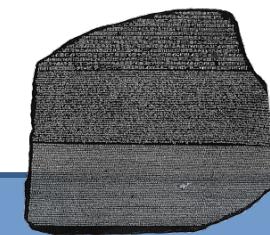
Movie
Recommendation



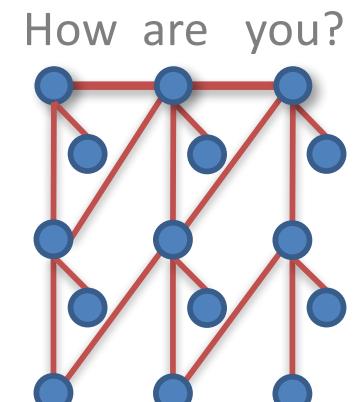
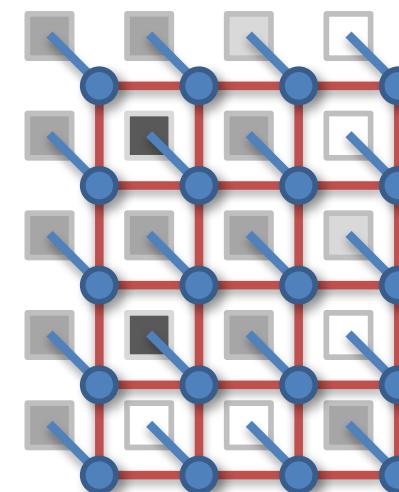
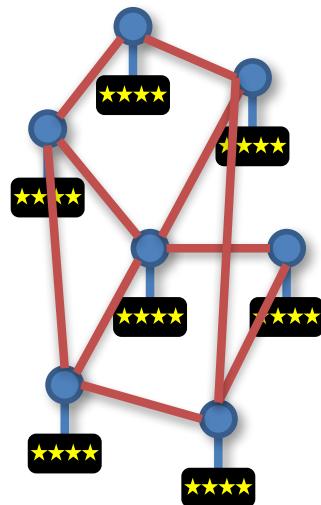
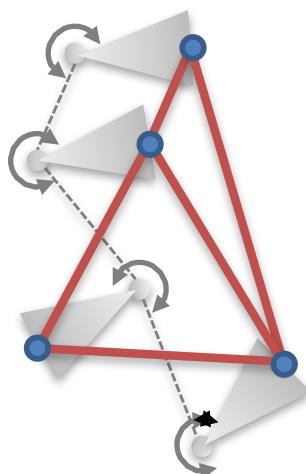
Computer
Vision



Machine
Translation

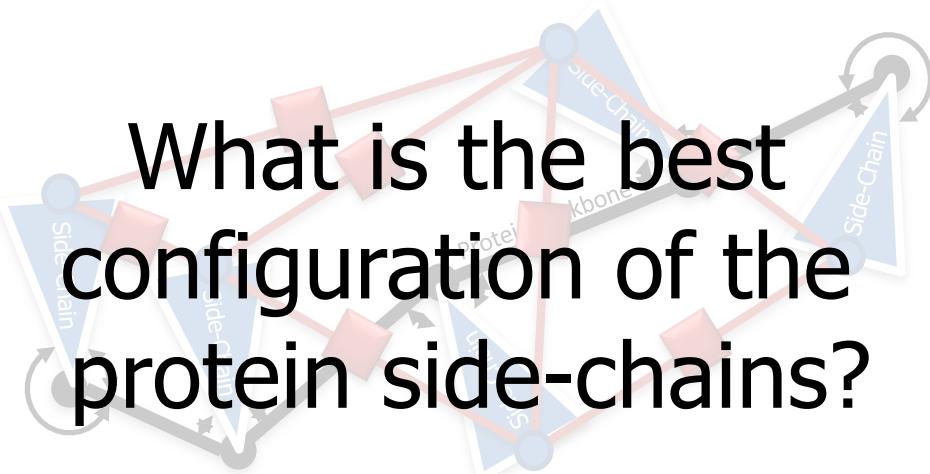


Probabilistic Graphical Models



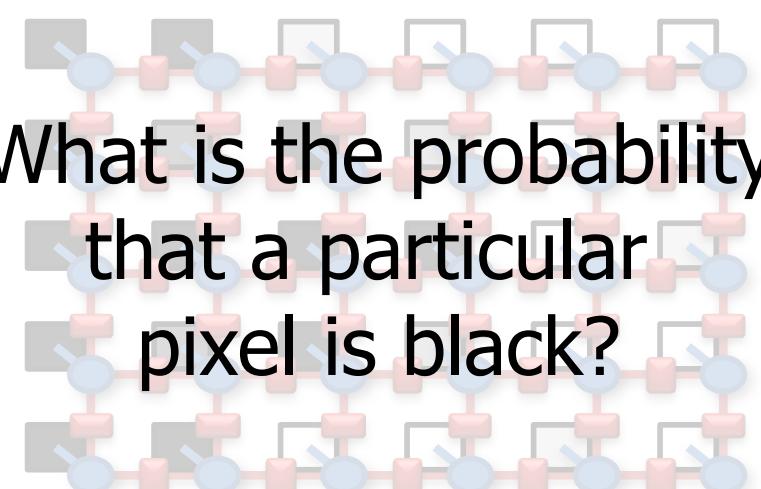
Probabilistic Inference

Making predictions given the model structure and parameters



What is the best configuration of the protein side-chains?

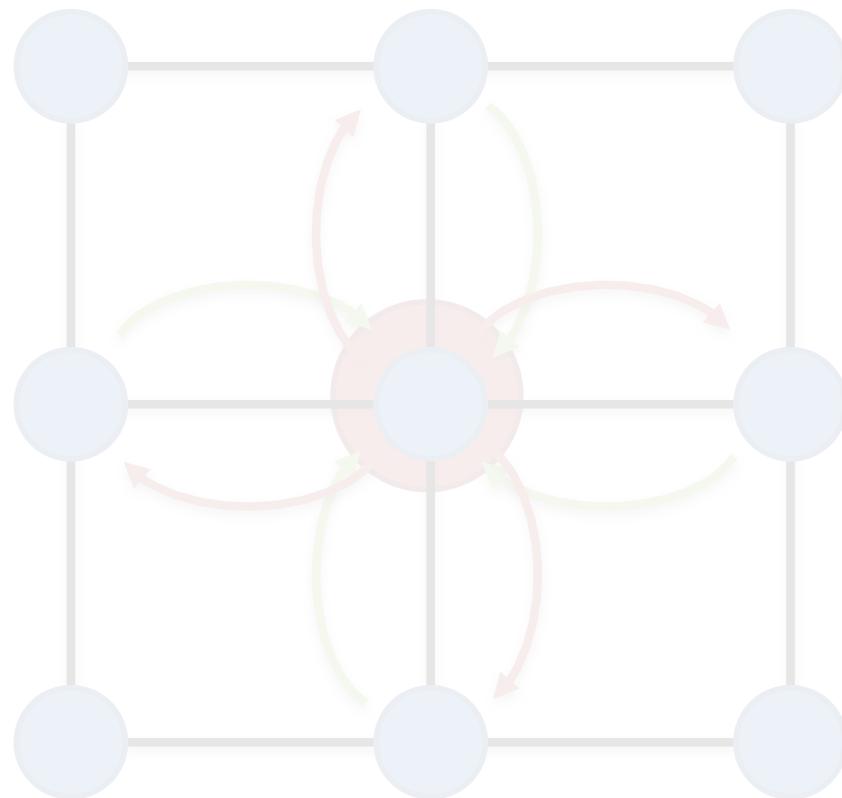
- **NP-complete** in general
 - Focus on *approximate* methods



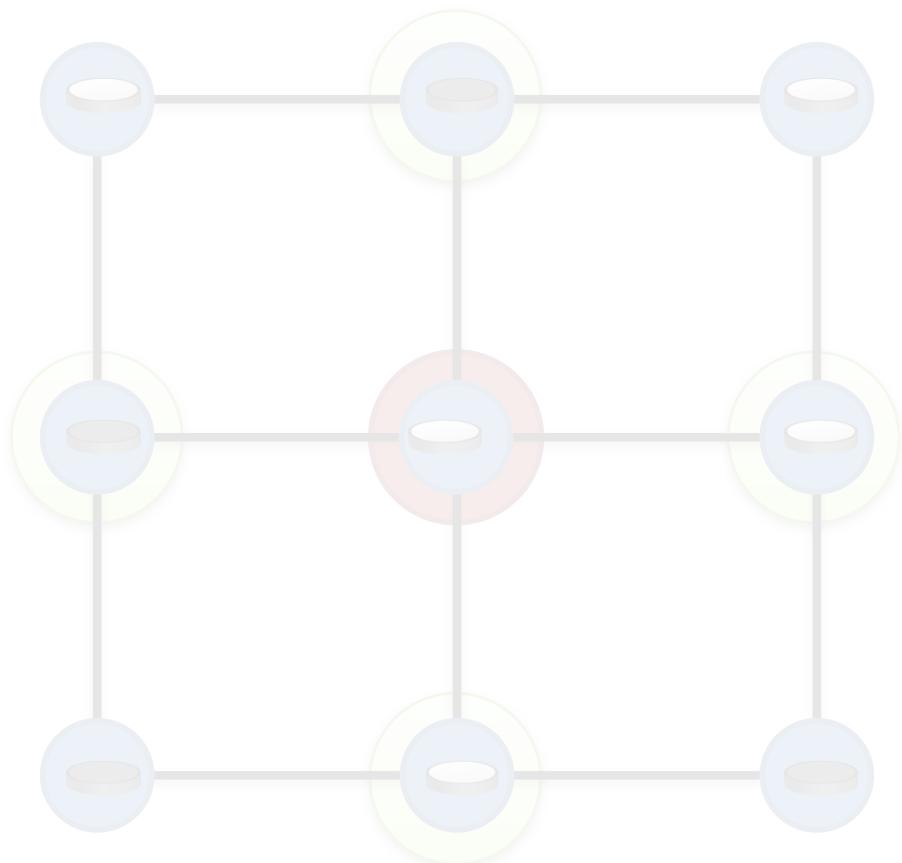
What is the probability that a particular pixel is black?

Parallel and Distributed Algorithms for Probabilistic Inference

Belief Propagation



Gibbs Sampling



Parallel Belief Propagation

Joint Work With:

Yucheng Low

Carlos Guestrin

David O'Hallaron

Published Results

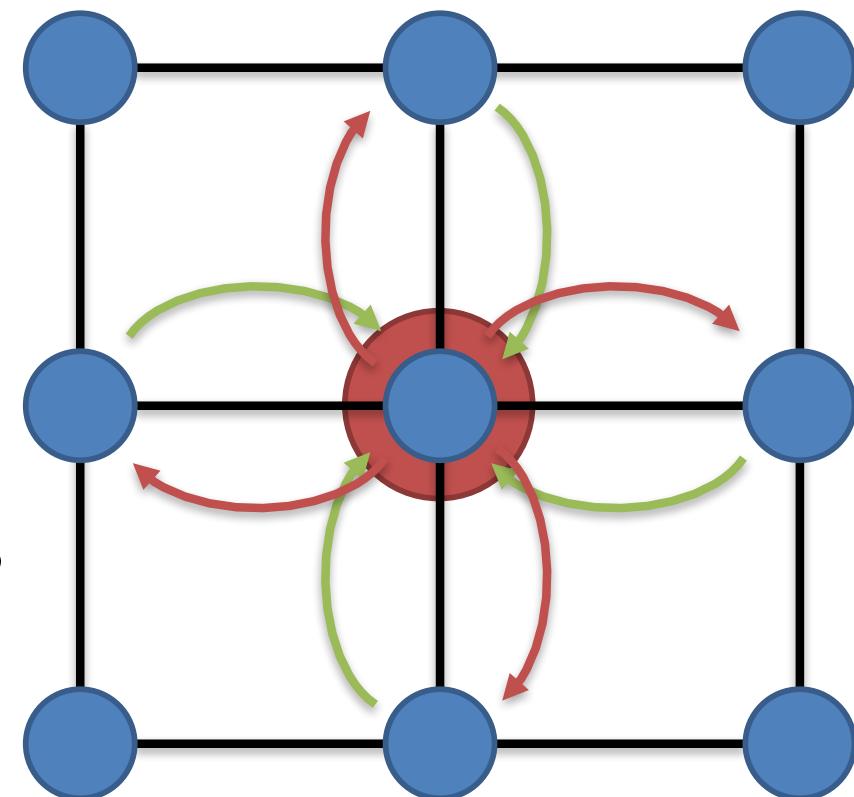
AISTATS'09

UAI'09

Chapter in SUML'10

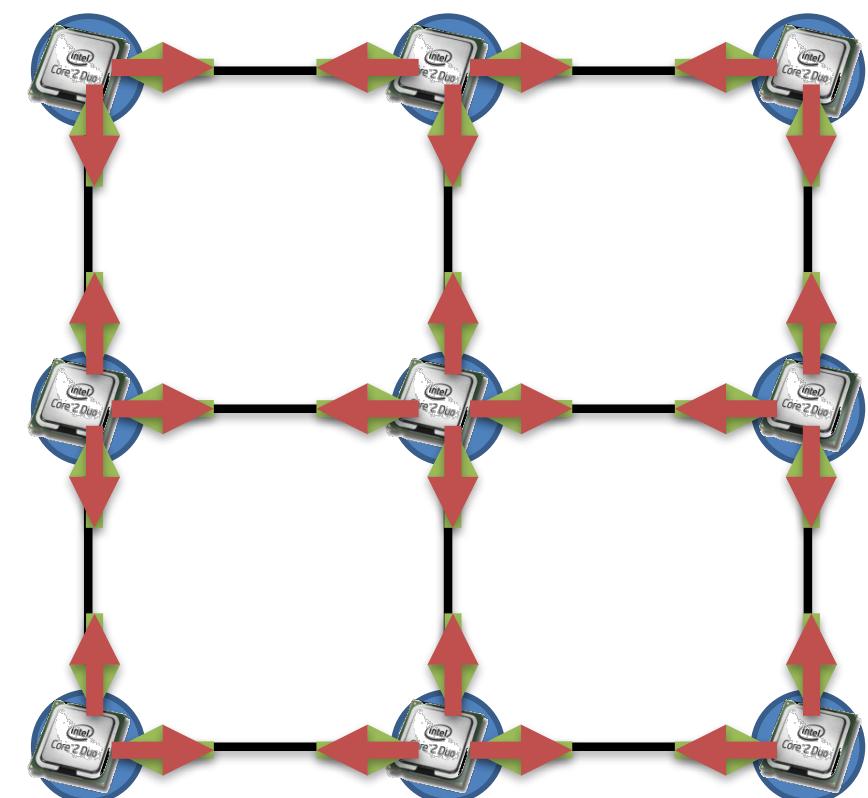
Loopy Belief Propagation (Loopy BP)

- Iteratively estimate the variable beliefs
 - Read **in messages**
 - Updates marginal estimate (**belief**)
 - Send updated **out messages**
- Repeat for all variables until convergence



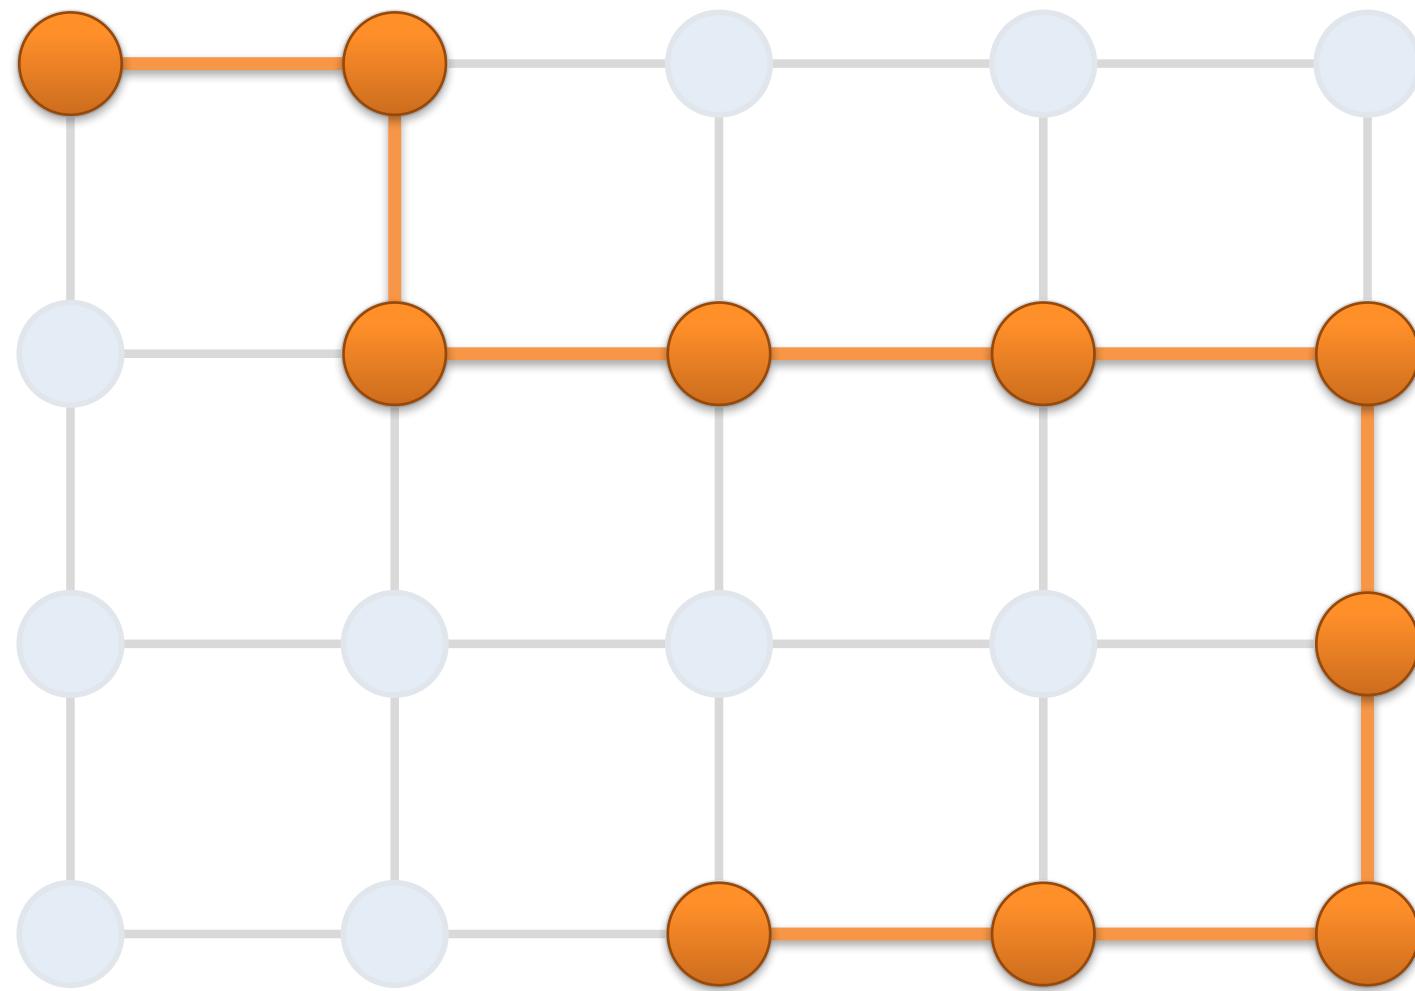
Synchronous Loopy BP

- Often considered embarrassingly parallel
 - Associate processor with each vertex
 - Receive all messages
 - Update all beliefs
 - Send all messages
- Proposed by:
 - Brunton et al. CRV'06
 - Mendiburu et al. GECC'07
 - Kang,et al. LDMTA'10
 - ...

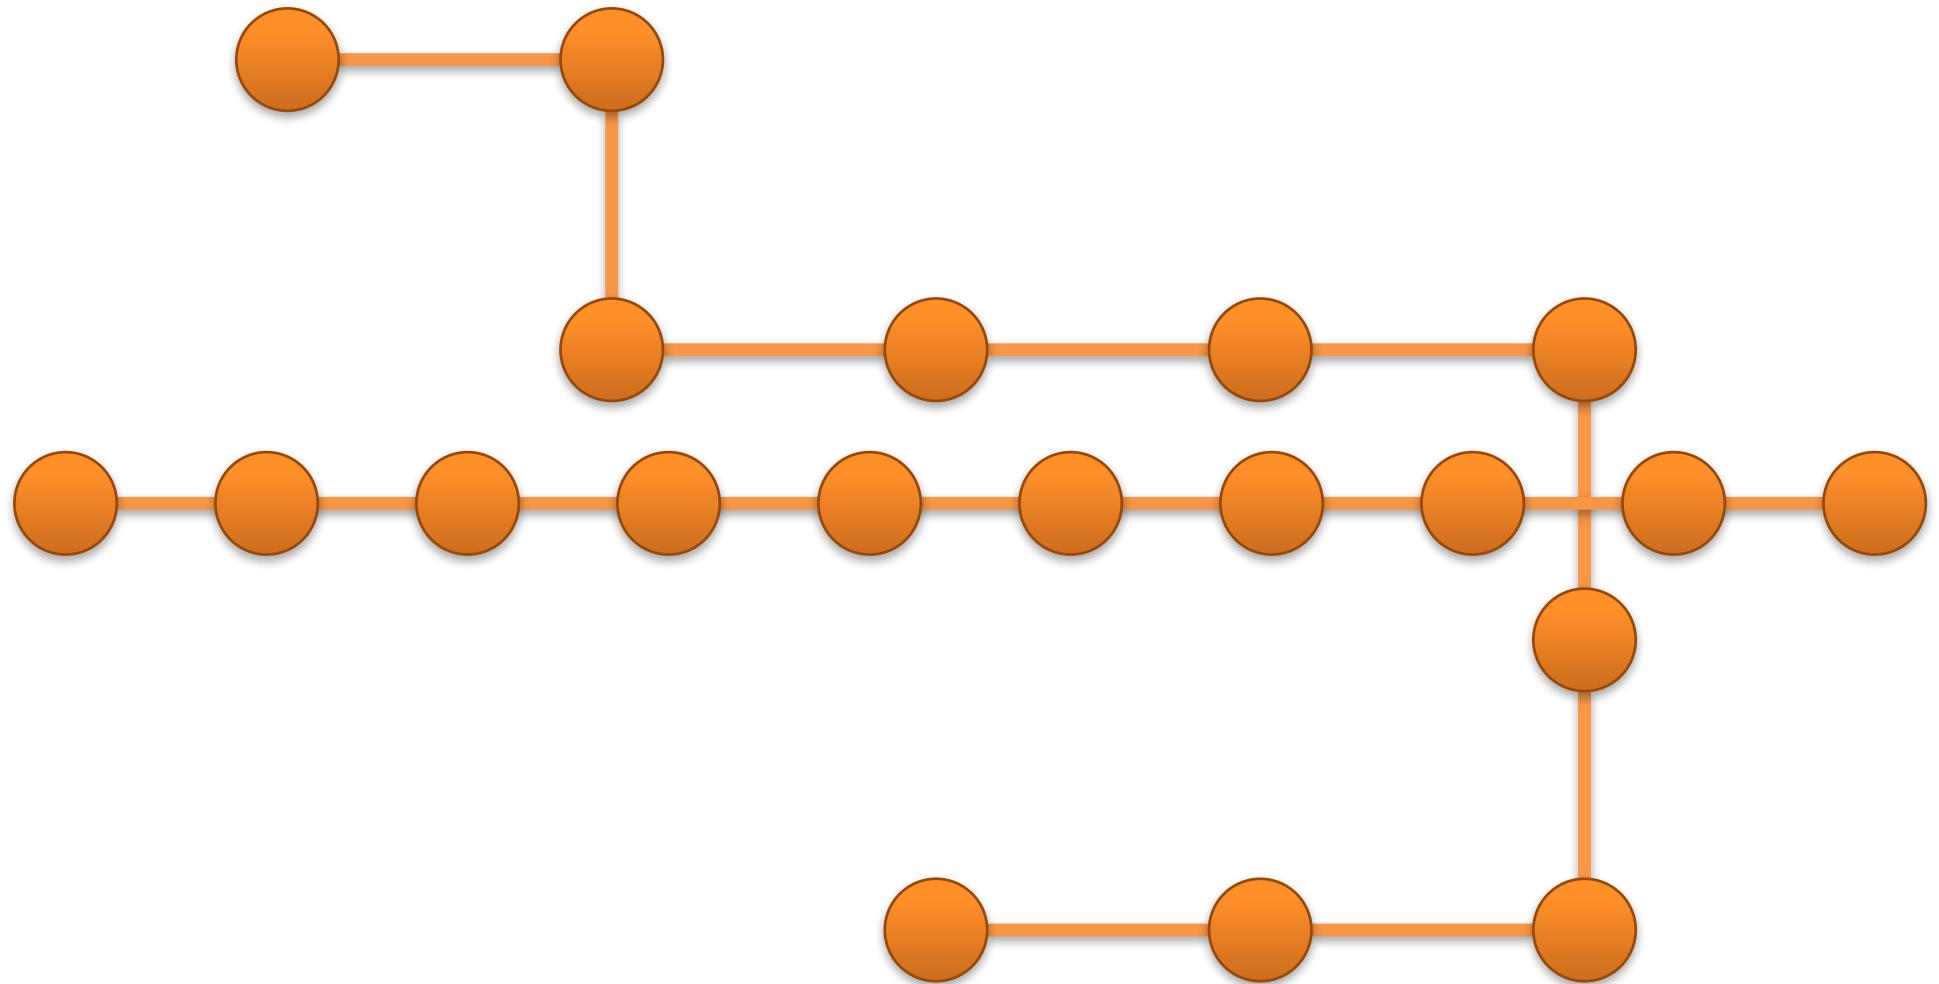


*Is Synchronous Loopy BP
an **efficient** parallel algorithm?*

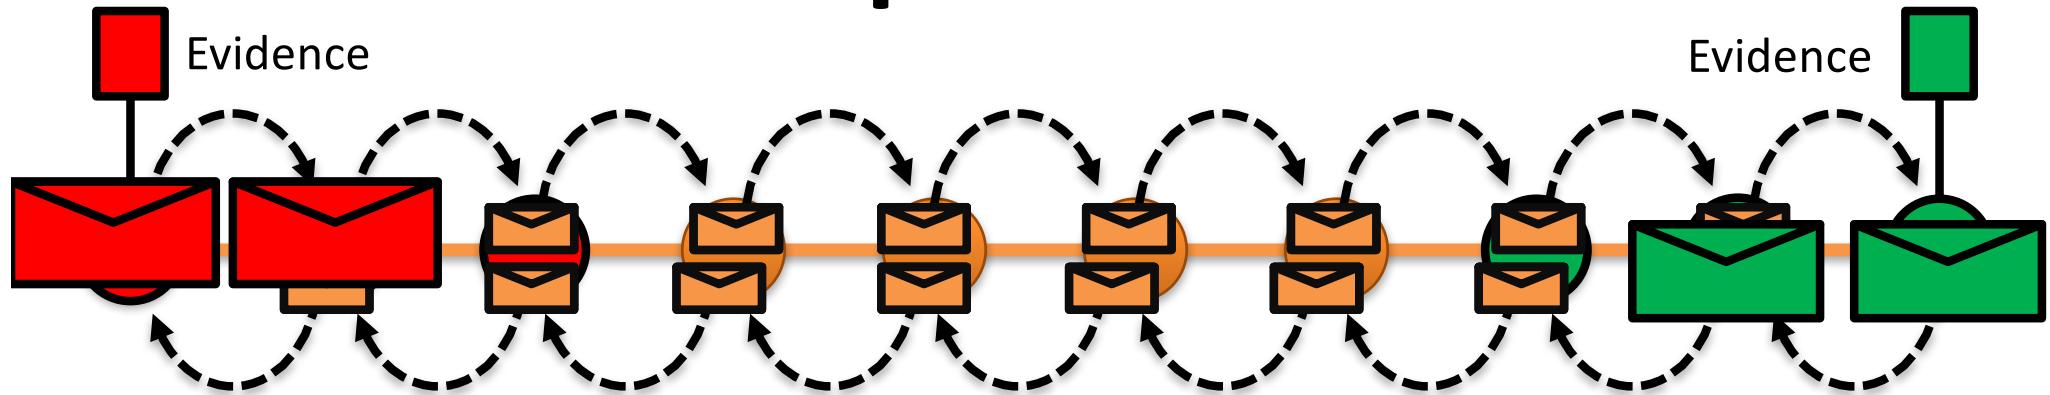
Sequential Computational Structure



Hidden Sequential Structure



Hidden Sequential Structure



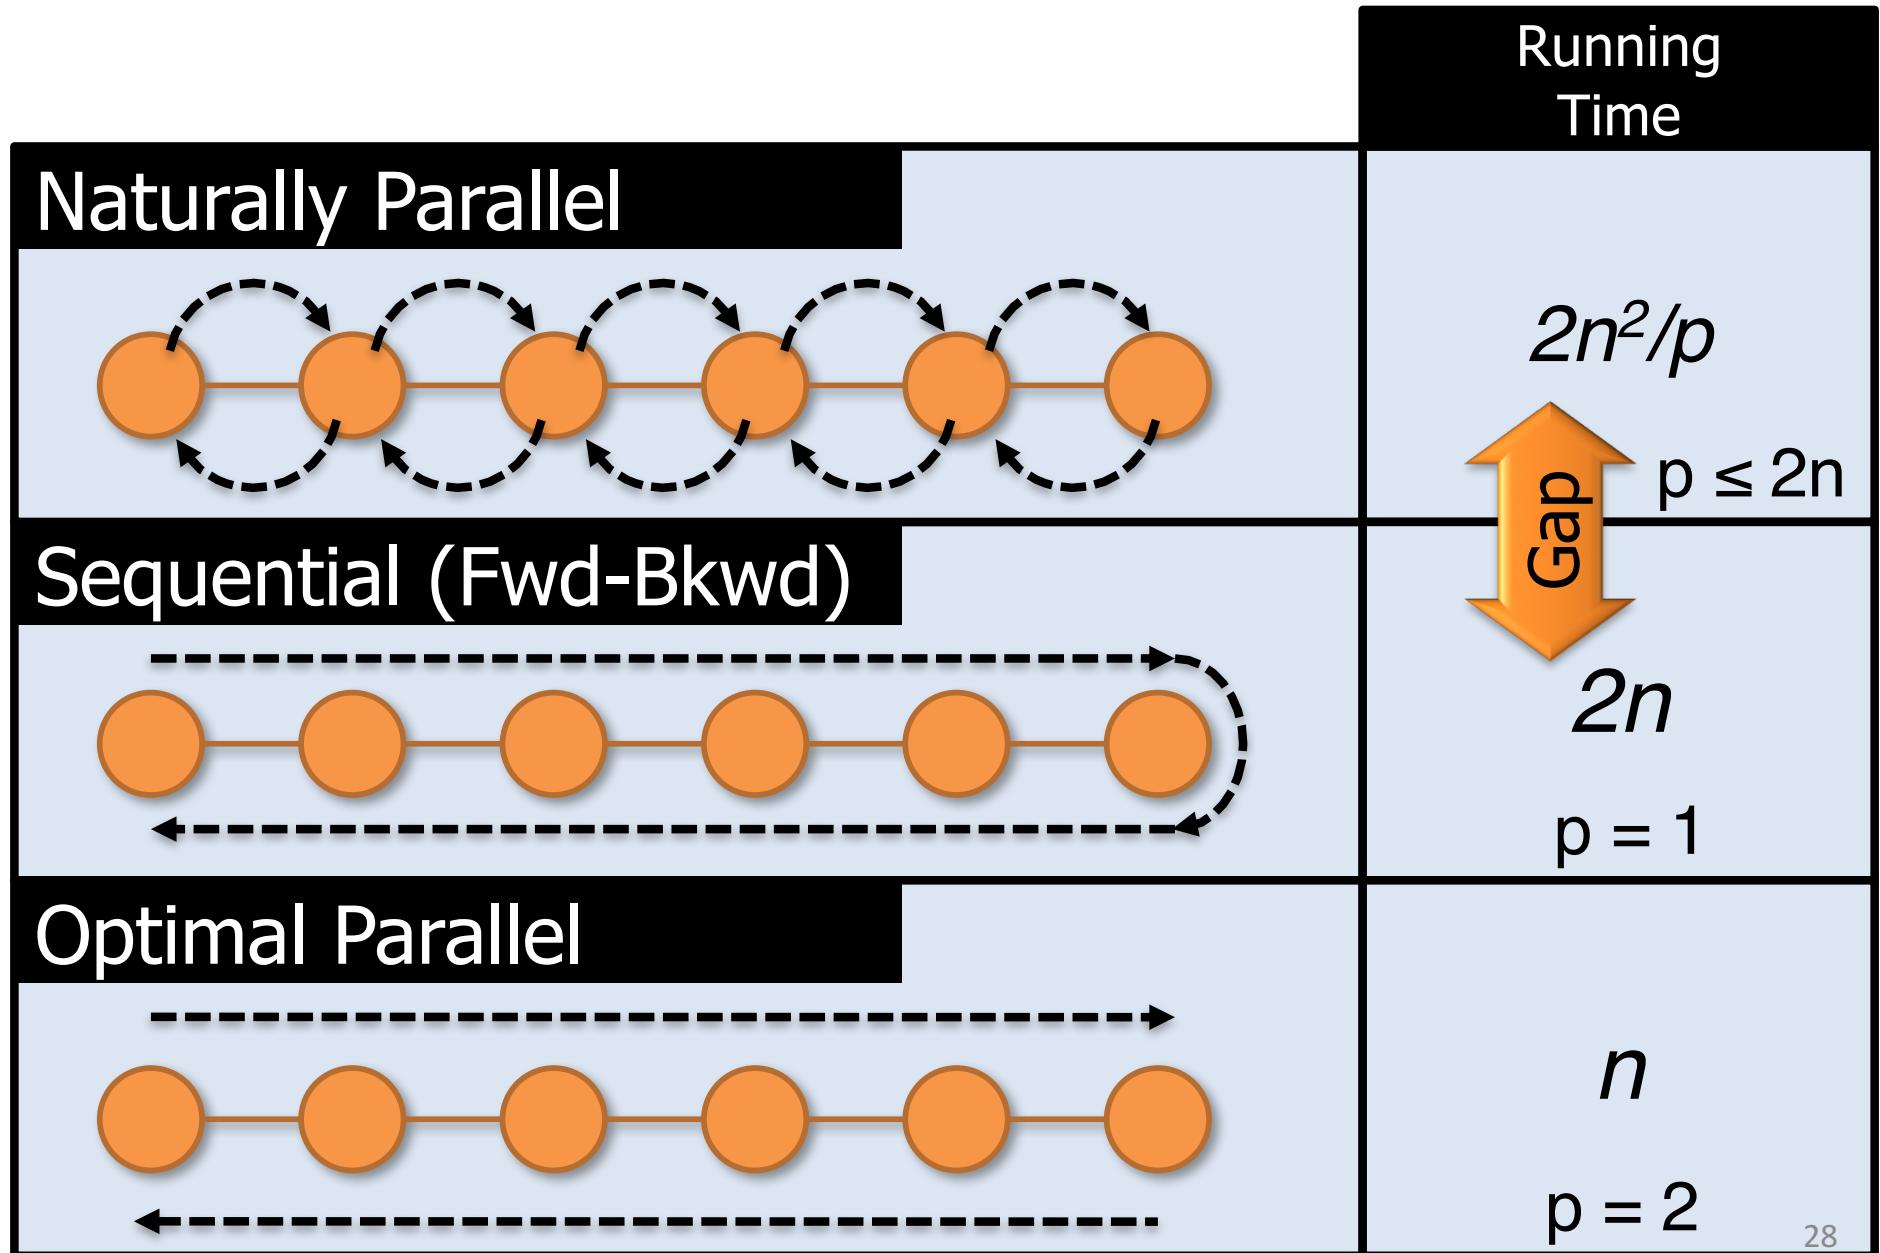
- Running Time:

$$\frac{2n \text{ Messages Calculations}}{p \text{ Processors}} \times (n \text{ Iterations to Converge}) = \frac{2n^2}{p}$$

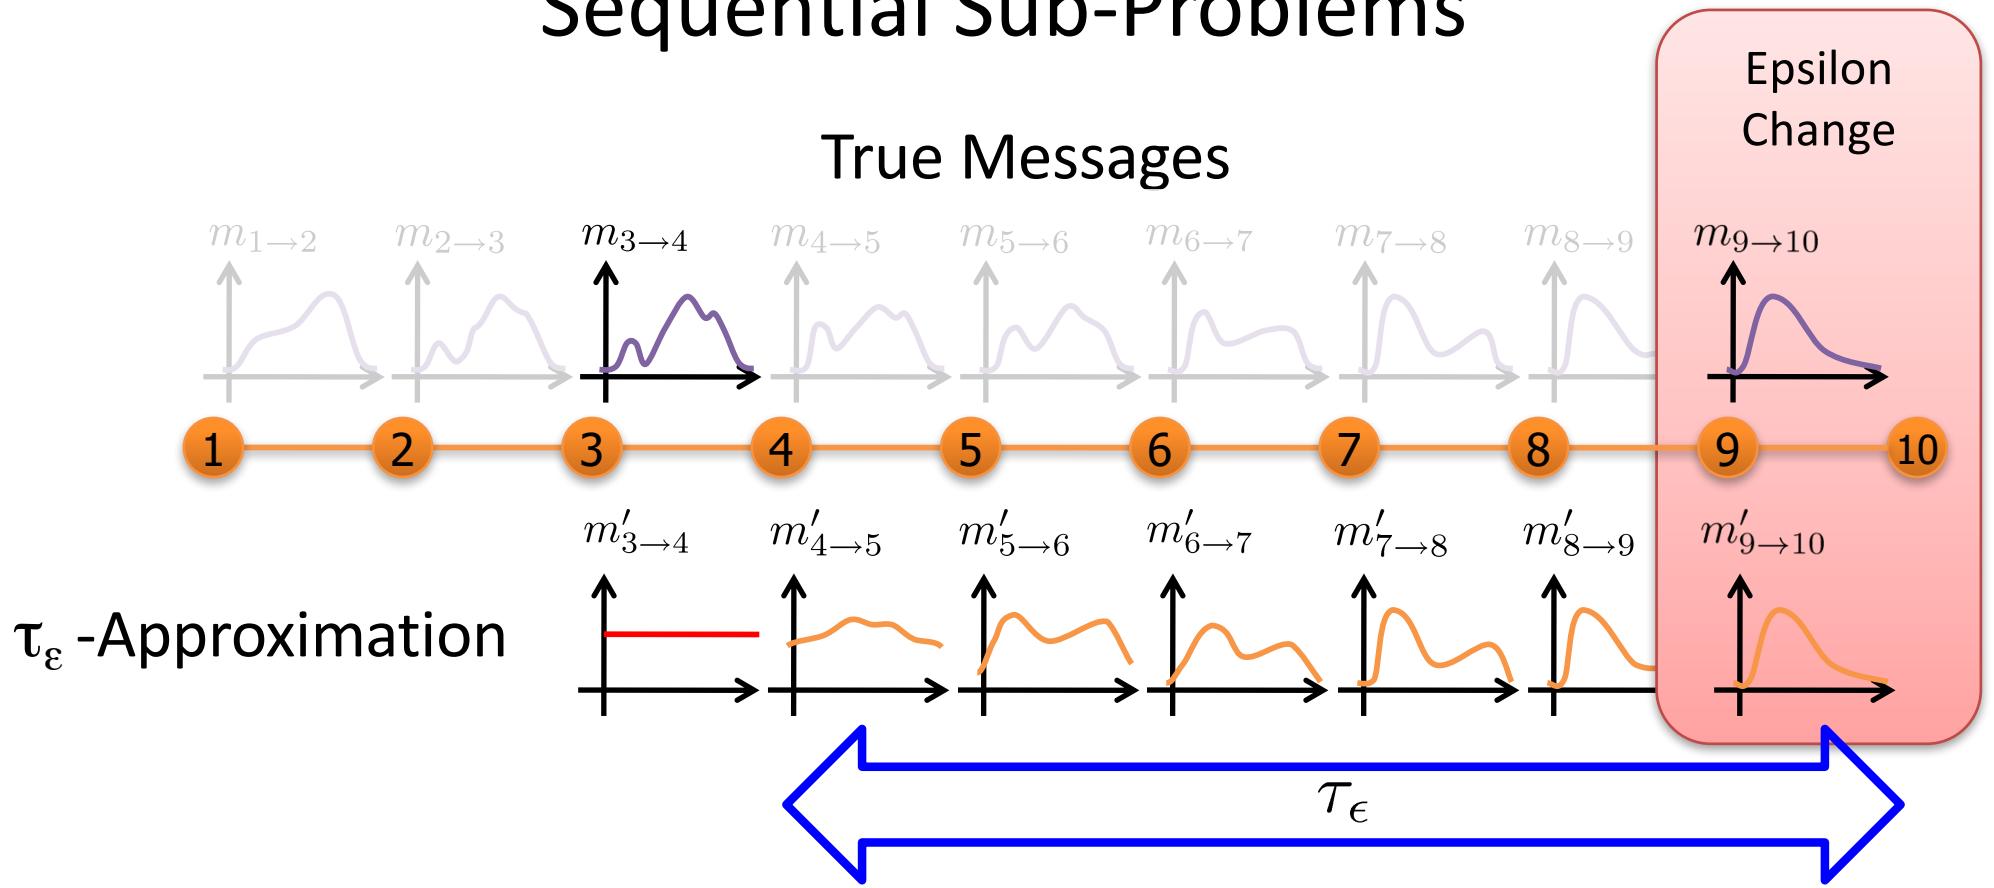
Time for a single parallel iteration

Number of Iterations

Optimal Sequential Algorithm

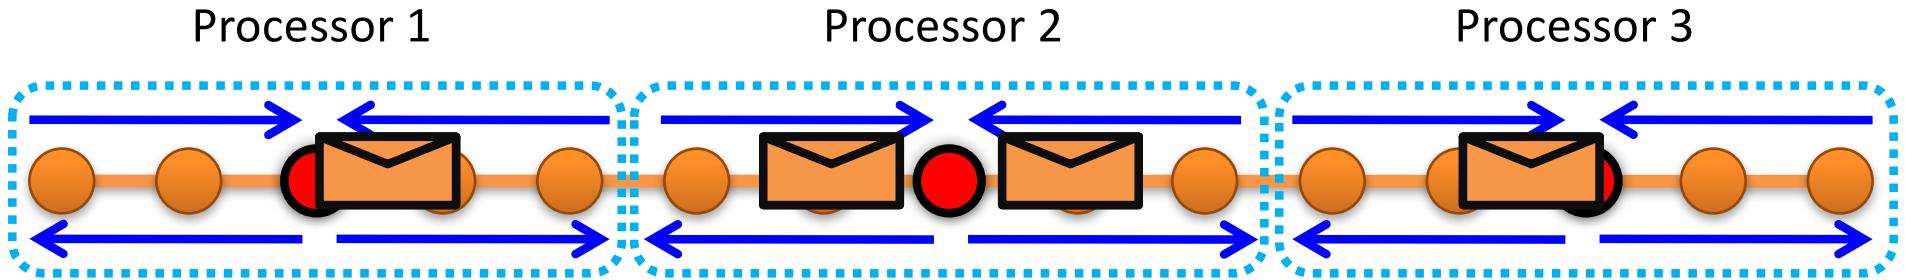


Role of model Parameters on Sequential Sub-Problems



- τ_ϵ represents the minimal sequential sub-problem
- Captures dependence on **model parameters**

Optimal Parallel Scheduling



Theorem:

Using p processors this algorithm achieves a τ_ϵ approximation in time:

Parallel Component

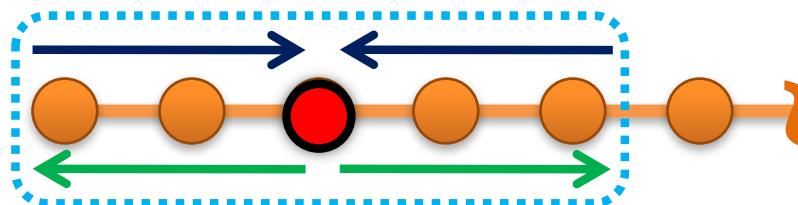
$$O\left(\frac{n}{p} + \tau_\epsilon\right)$$

Sequential Component

and is **optimal** for chain graphical models.

The Splash Operation

- Generalize the optimal chain algorithm:

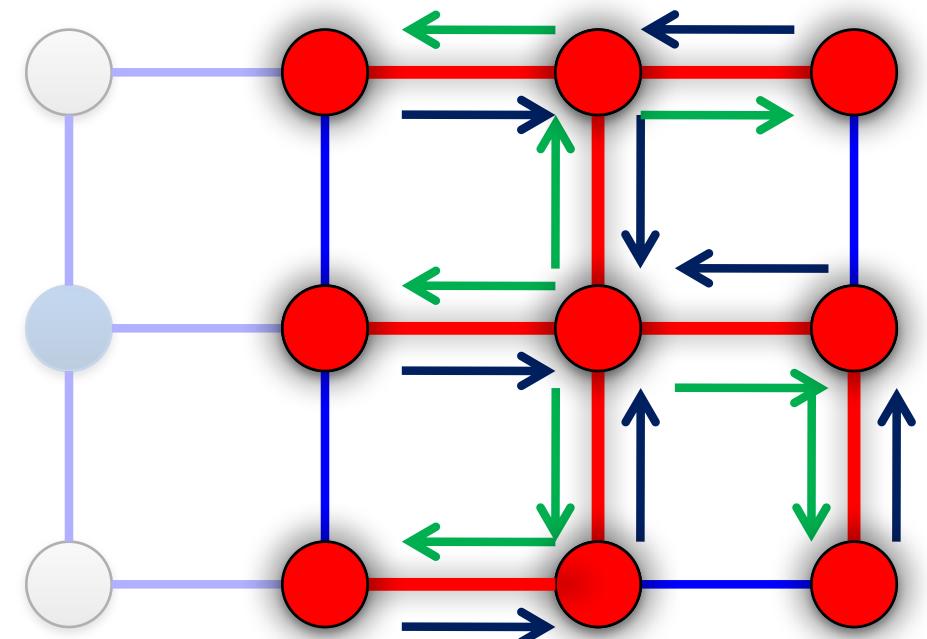


to arbitrary cyclic graphs:

1) Grow a BFS Spanning tree with fixed size

2) Forward Pass computing all messages at each vertex

3) Backward Pass computing all messages at each vertex



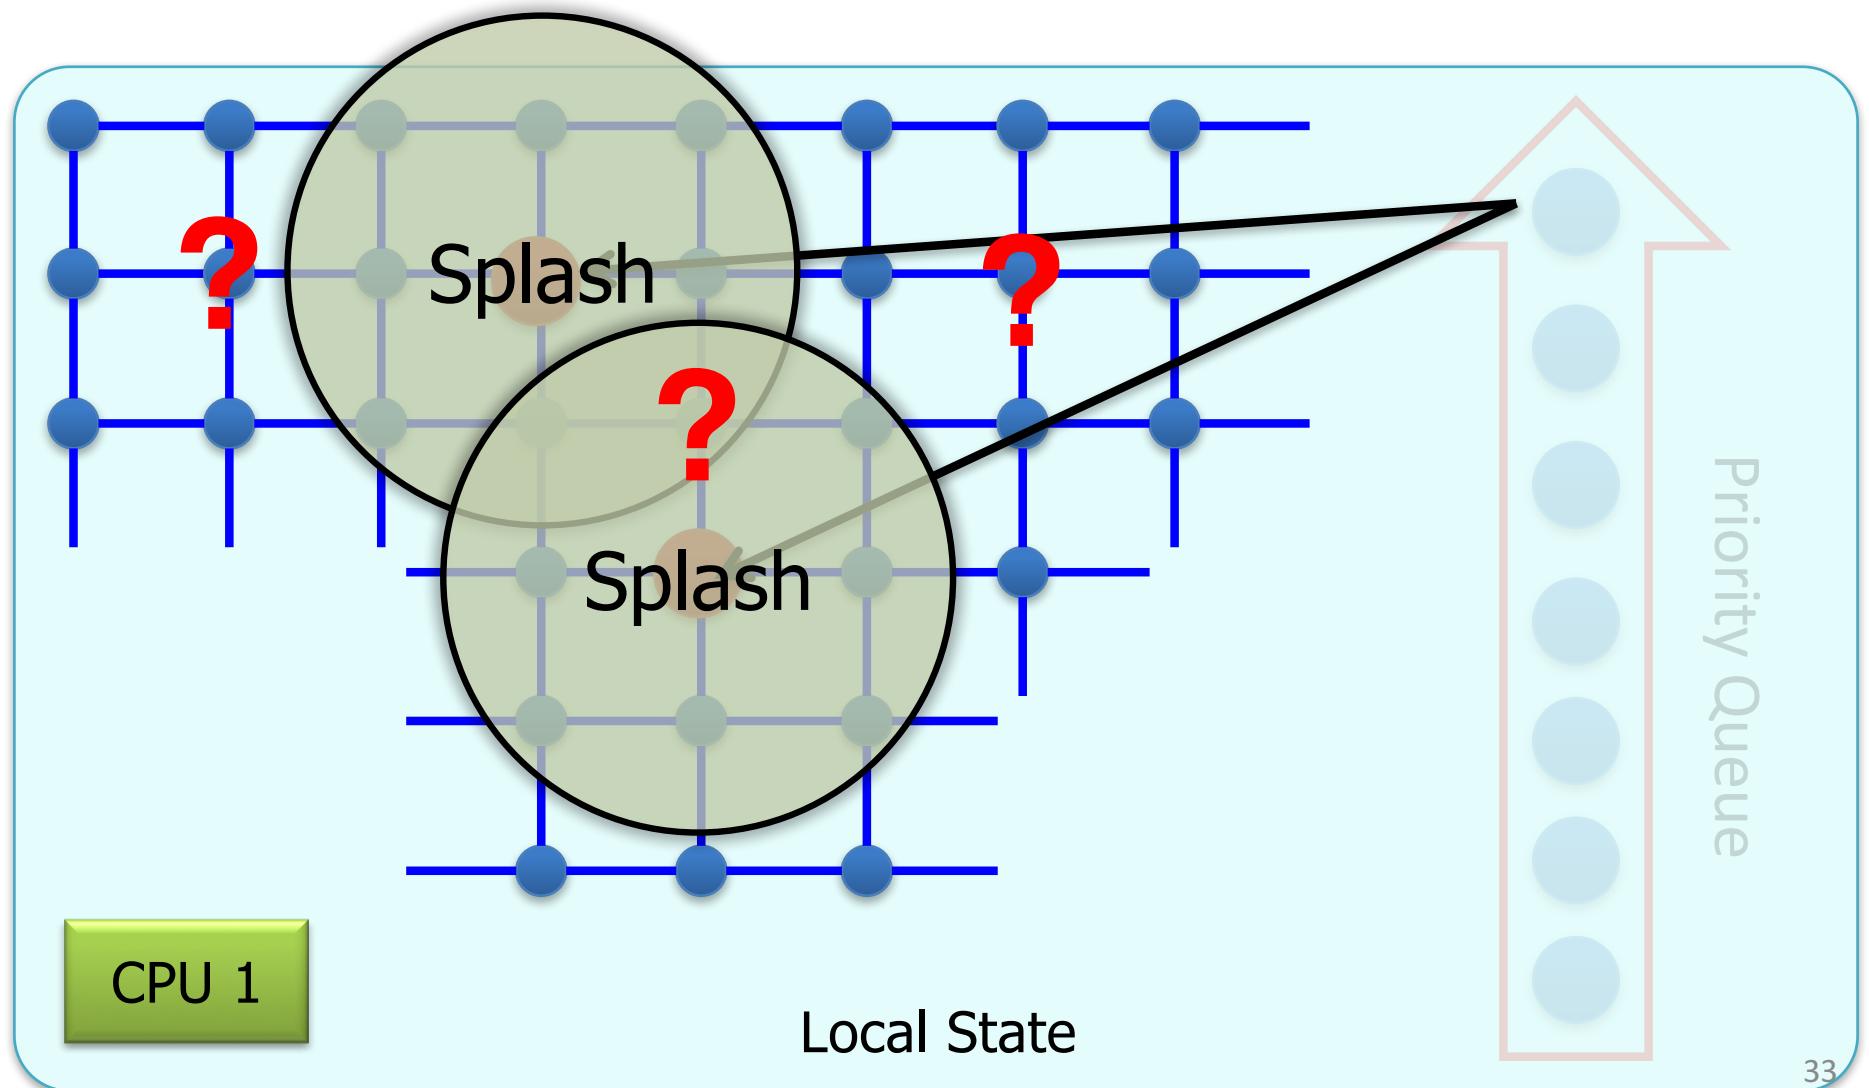
Running Parallel Splashes



- Partition the graph
- Schedule Splashes locally
- Transmit the messages along the boundary of the partition

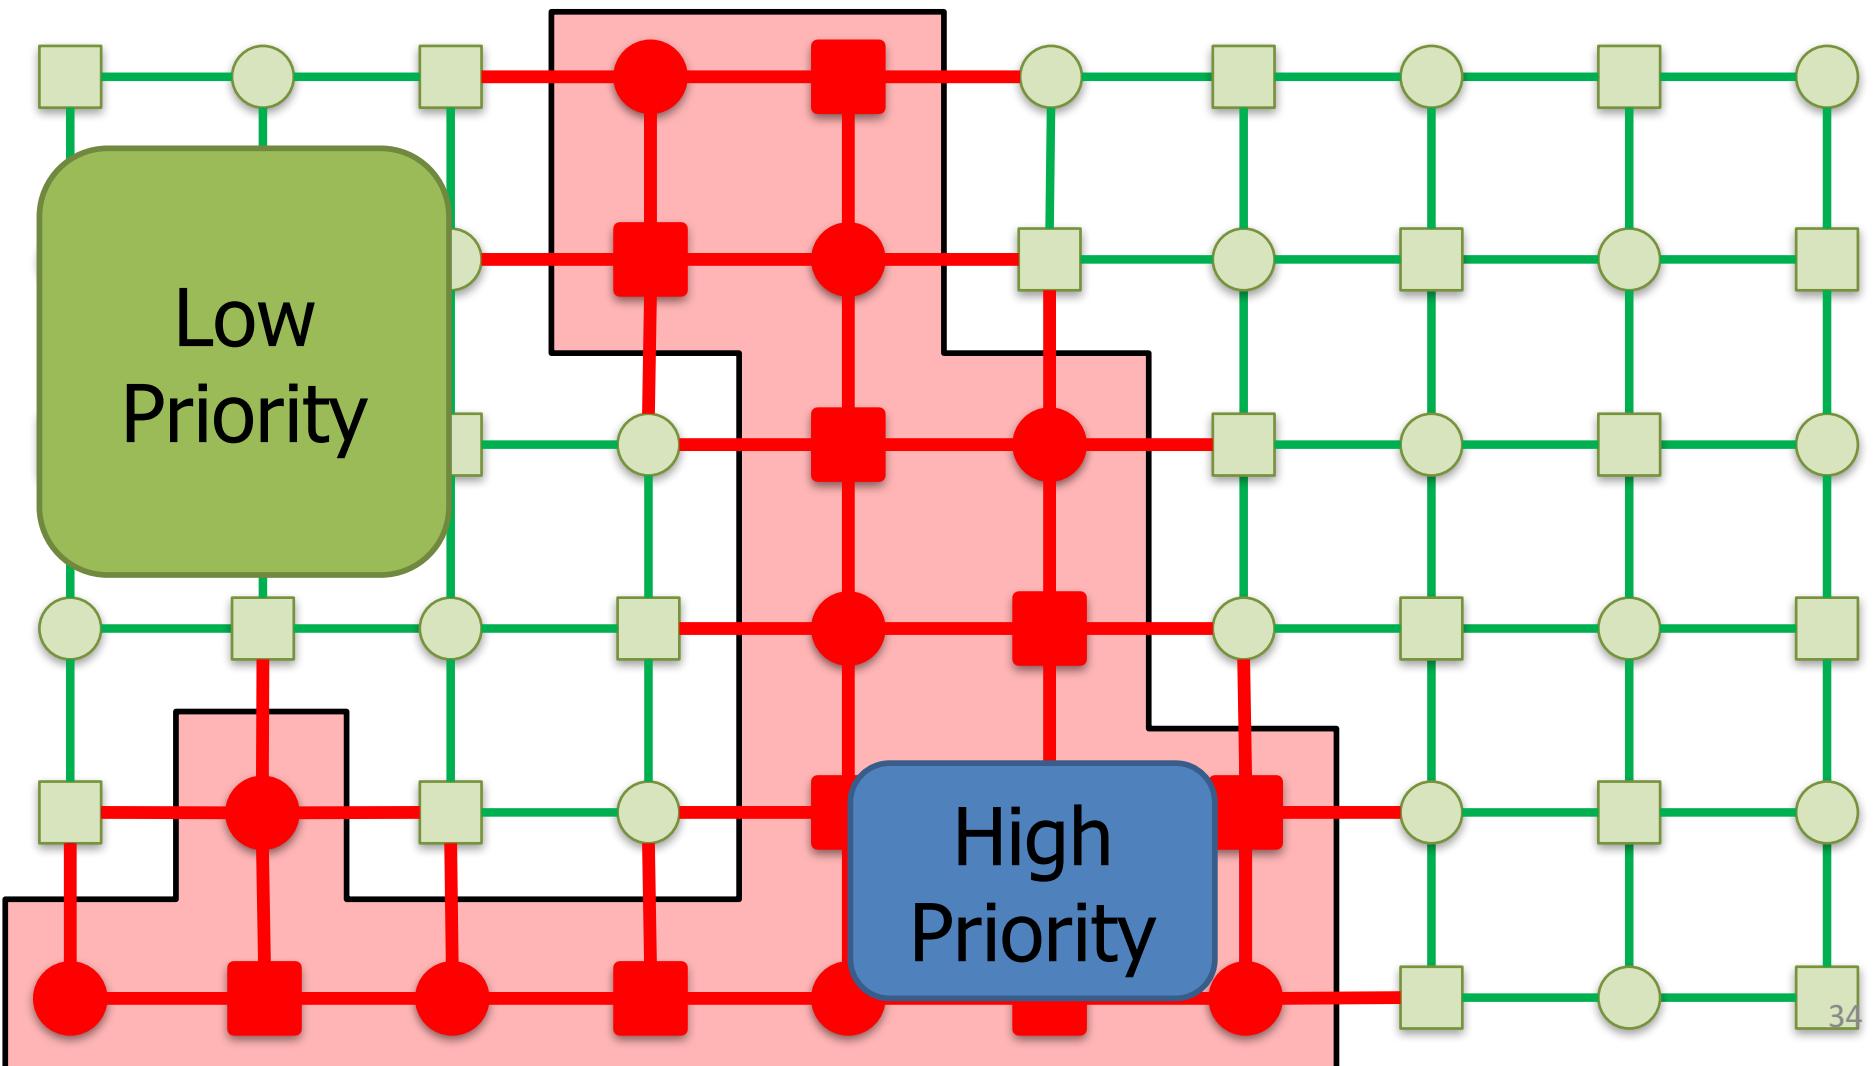
Priorities Determine the Roots

- Use a residual priority queue to select roots:

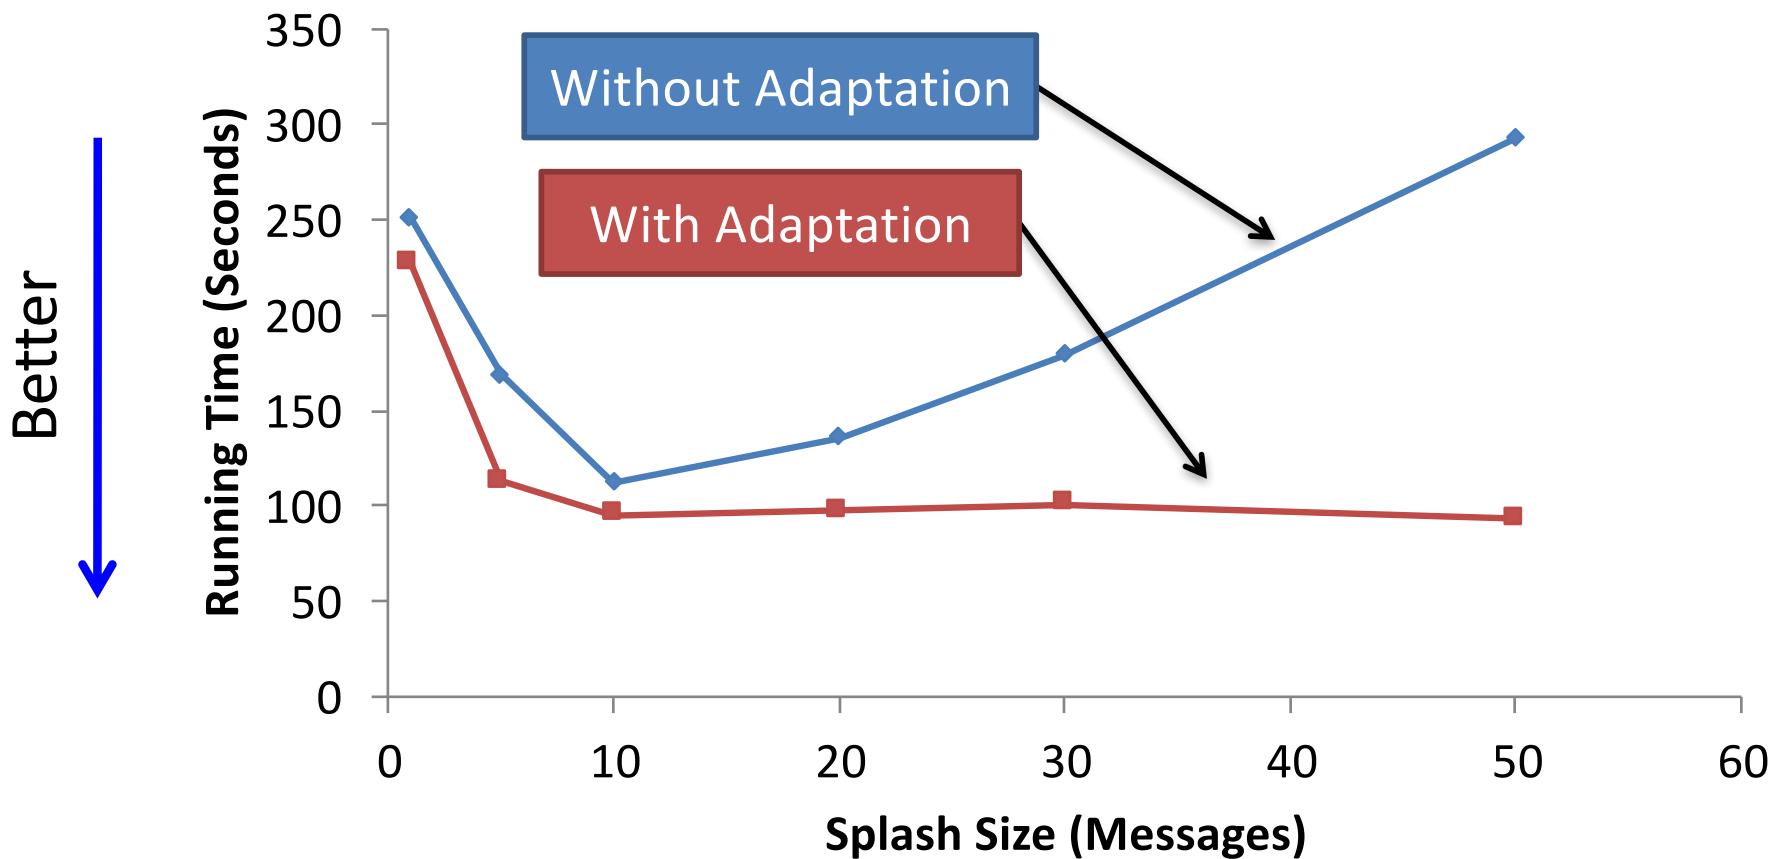


Dynamic Splashes

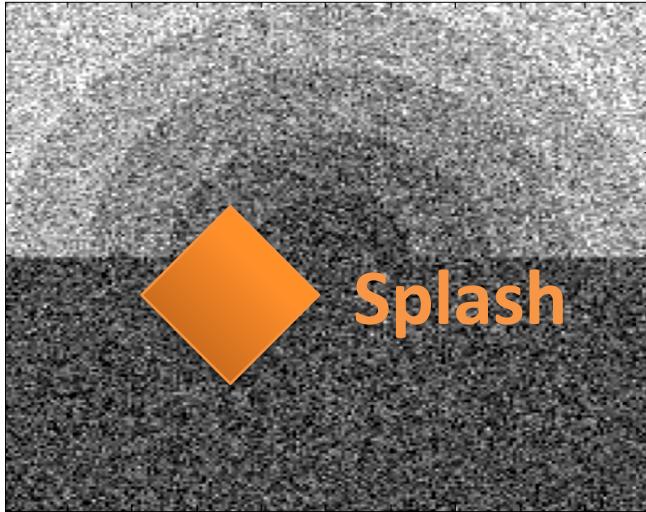
Priorities **adaptively** focus computation by determining the **shape** and **size** of each Splash



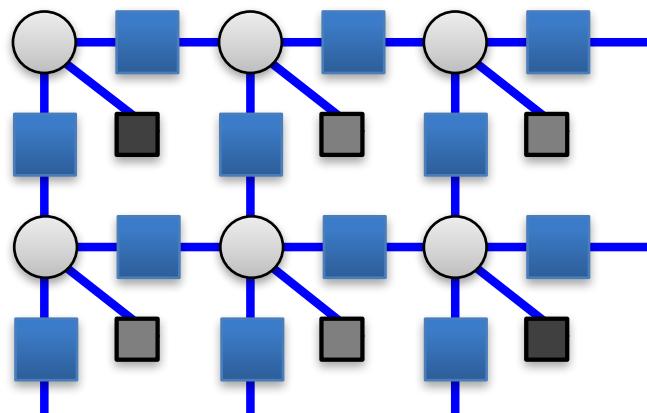
Dynamic Splashes automatically identify the optimal splash size



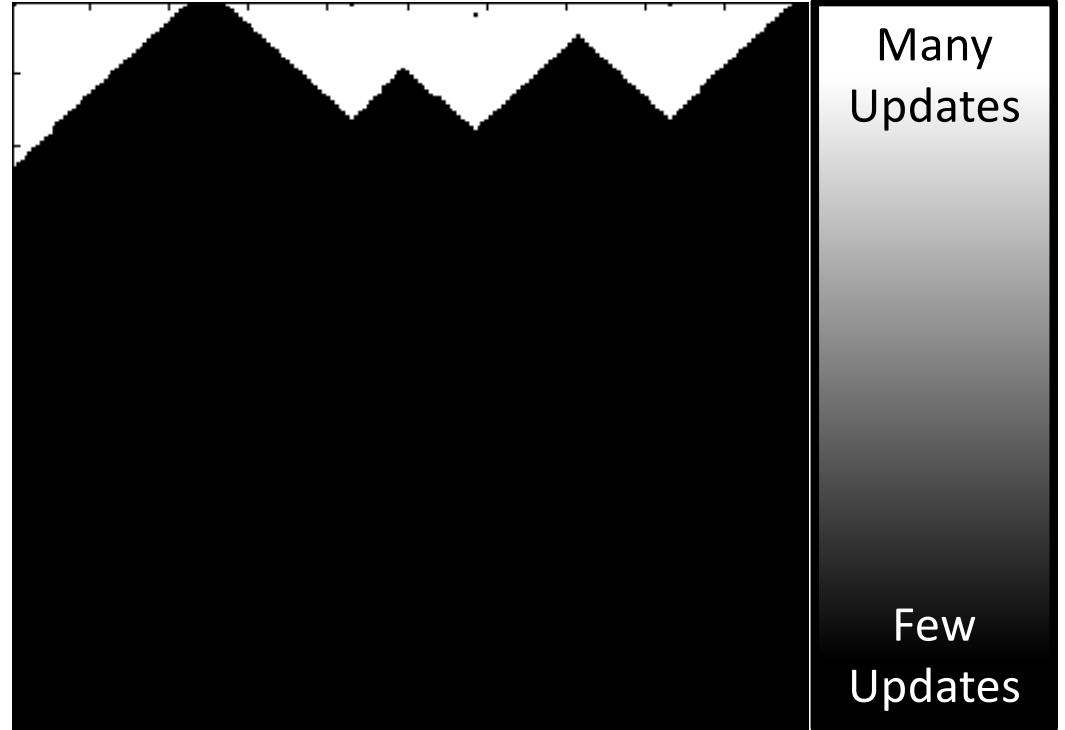
Splash Belief Propagation



Synthetic Noisy Image



Factor Graph



Vertex Updates

Algorithm identifies and focuses
on hidden sequential structure

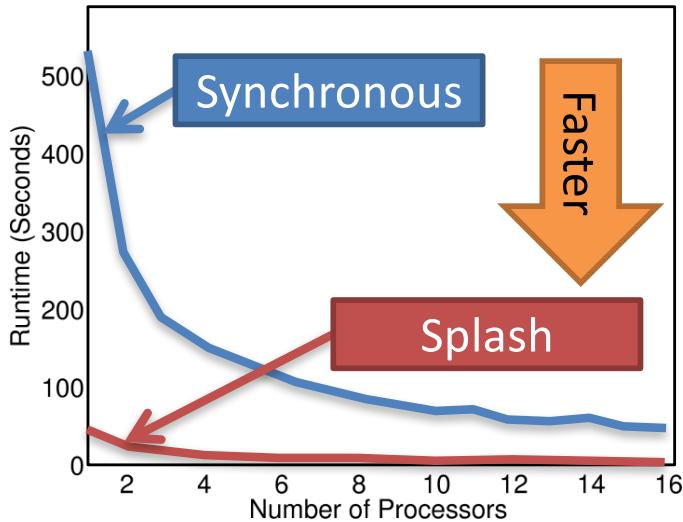
Evaluation

- System Design
 - Multicore and distributed implementations
 - Development was **time consuming**
- Evaluated on several real-world problems
 - Protein interaction and structure prediction
 - Markov Logic Networks
- Compared against several other variants
 - Faster, more efficient, more stable

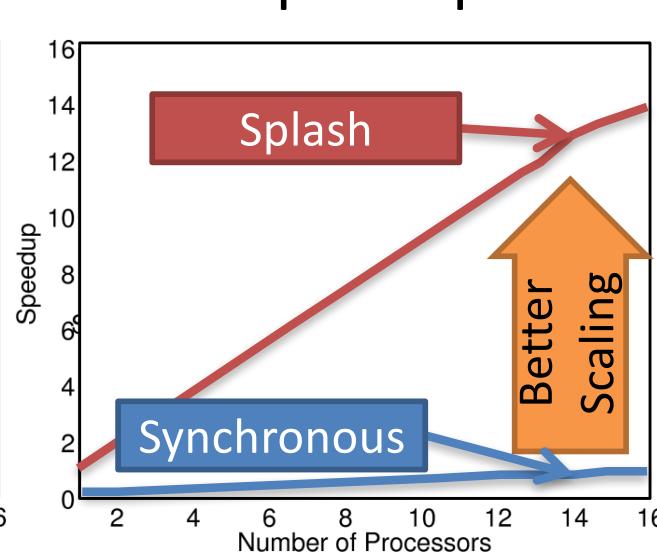
Representative Results

Protein Interaction Models: 14K Vertices, 21K Factors

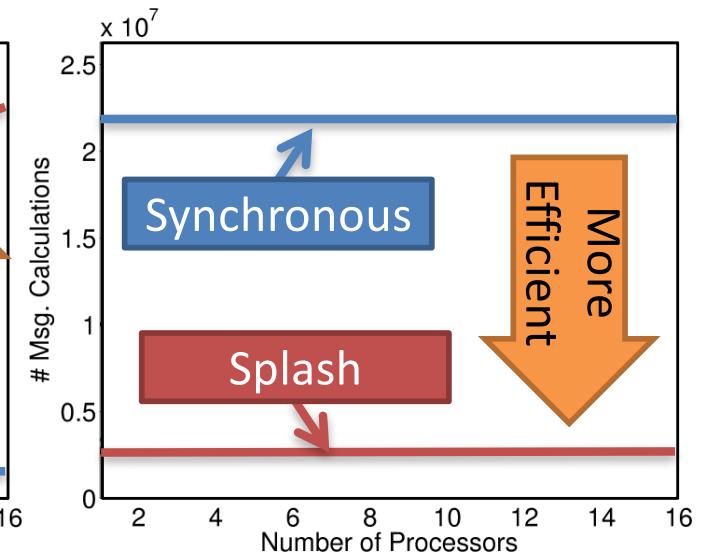
Runtime



Speedup



Total Work



- SplashBP converges more often
- Achieves better prediction accuracy

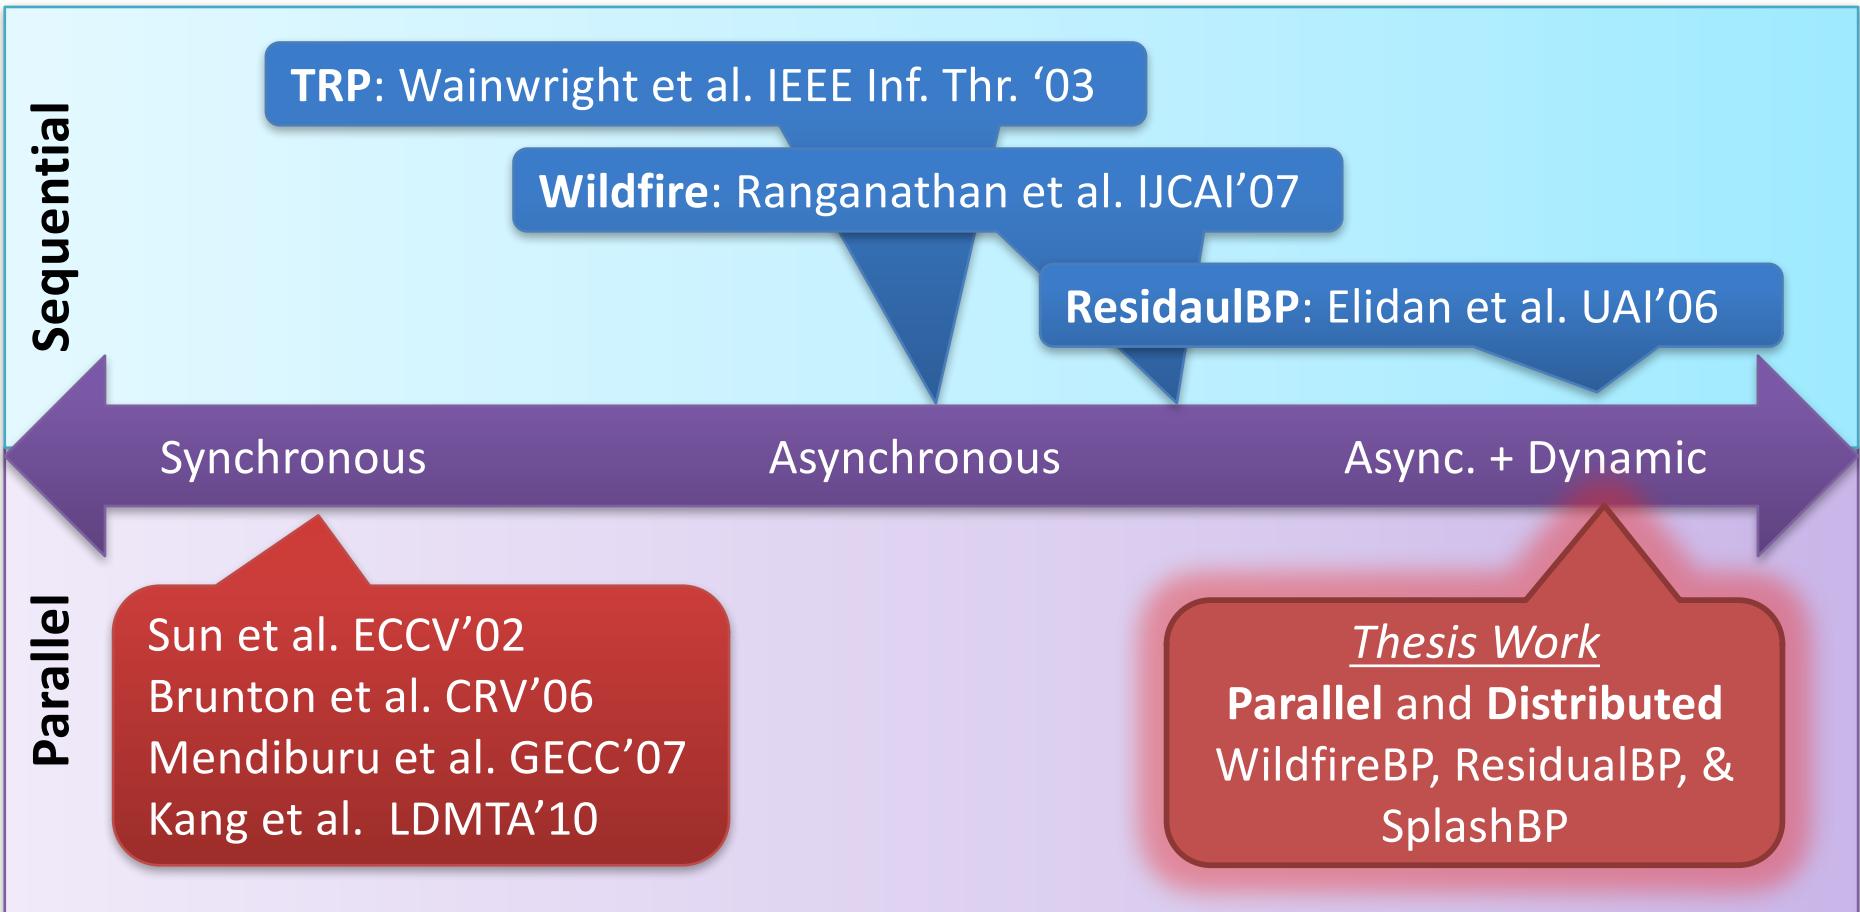
Summary: Belief Propagation

- **Asynchronous + Dynamic** → more efficient
 - *Theoretically and experimentally*
 - *Insight:* parallelize optimal **sequential** algorithm
 - *Tradeoff: Parallelism & Convergence*
- *Approximation* → Increased Parallelism
 - Exploit *weak* interactions (τ_ε – approximation)
- Key Contributions:
 - Demonstrate the importance of dynamic asynchronous scheduling in parallel inference
 - Theoretical analysis of work efficiency and relationship to model structure and parameters

GrAD Methodology

- **Graphical**
 - BP updates only depend on adjacent vertices
- **Asynchronous**
 - Compute messages sequentially within Splash
- **Dynamic**
 - Priority scheduling and adaptive Splashes

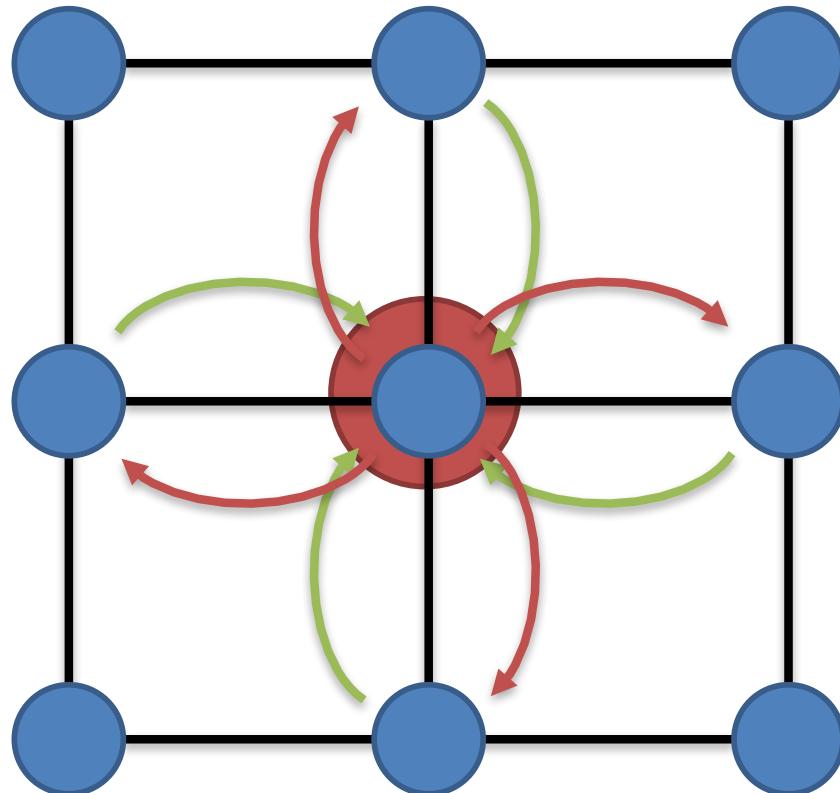
Additional Related Work



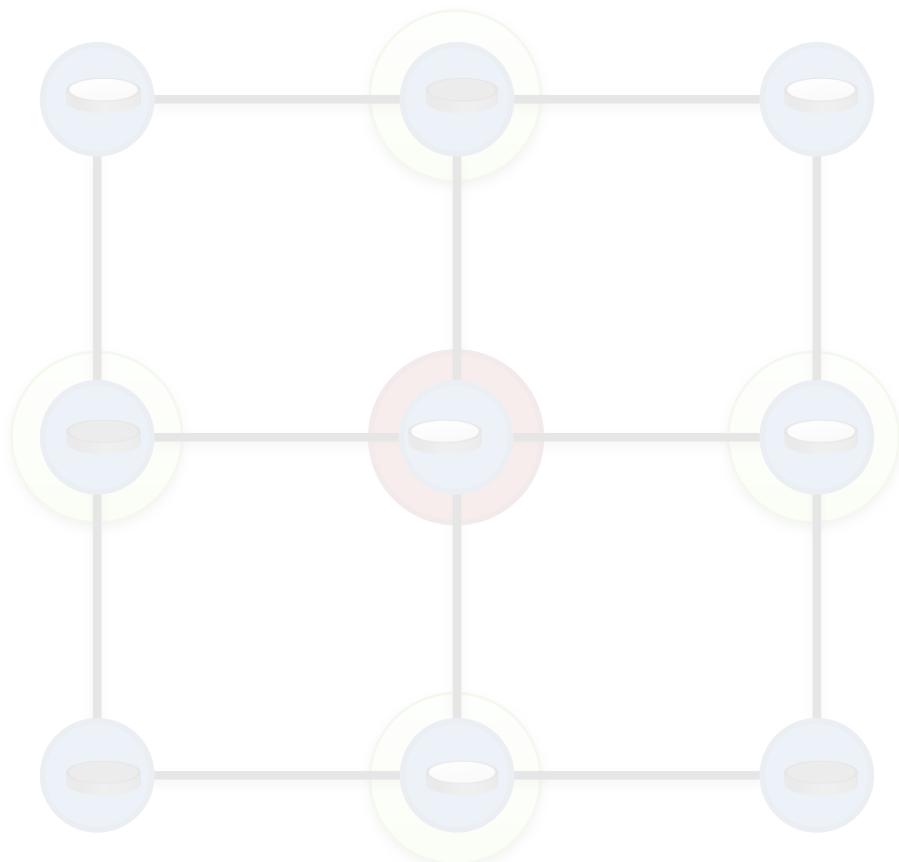
- **Parallel Exact Inference:** Pennock et al. UAI'98
- **Approximate Messages:** Ihler et al. JMLR'05

Parallel and Distributed Algorithms for Probabilistic Inference

Belief Propagation



Gibbs Sampling



Parallel Gibbs Sampling

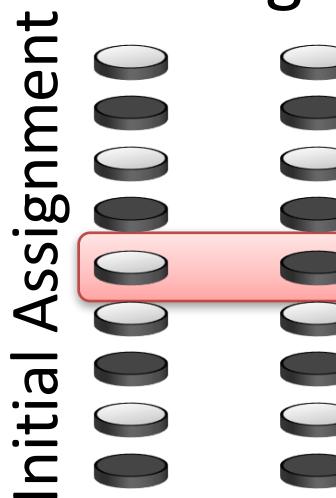
*An asynchronous Gibbs Sampler that
dynamically addresses strong dependencies.*

Joint Work With
Yucheng Low Arthur Gretton Carlos Guestrin

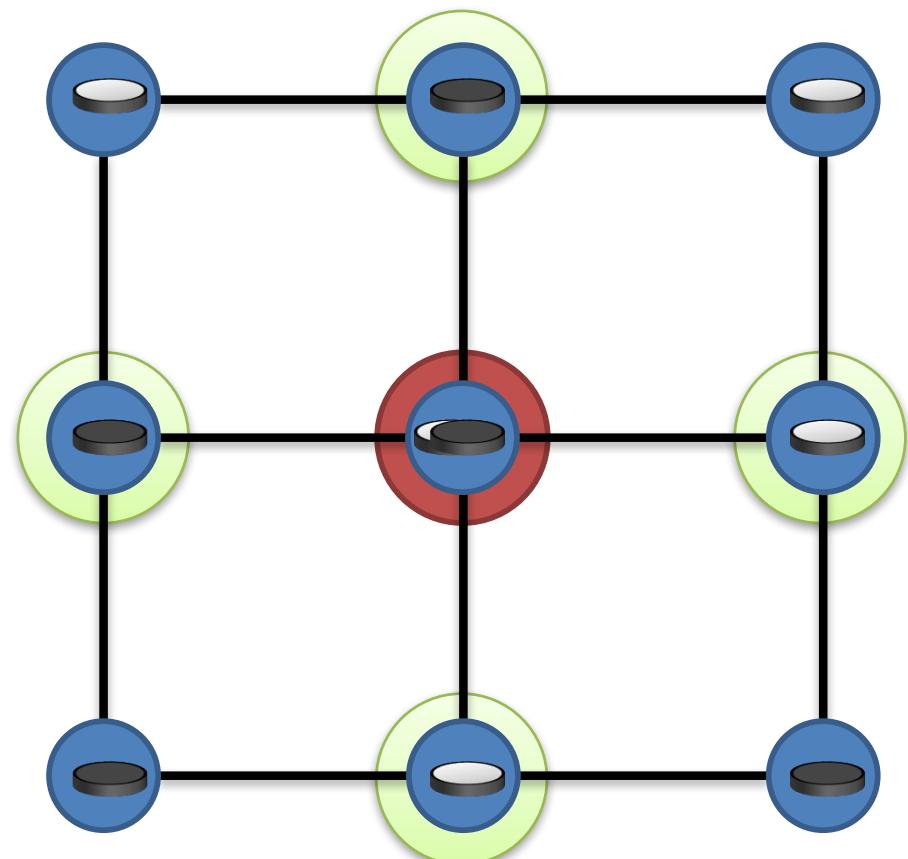
Published
AISTATS'11 (Related to work in WSDM'12)

Gibbs Sampling [Geman & Geman, 1984]

- **Sequentially** for each variable in the model
 - Select **variable**
 - Use **adjacent assignments** to construct a biased coin
 - Flip coin and update assignment to **variable**



...

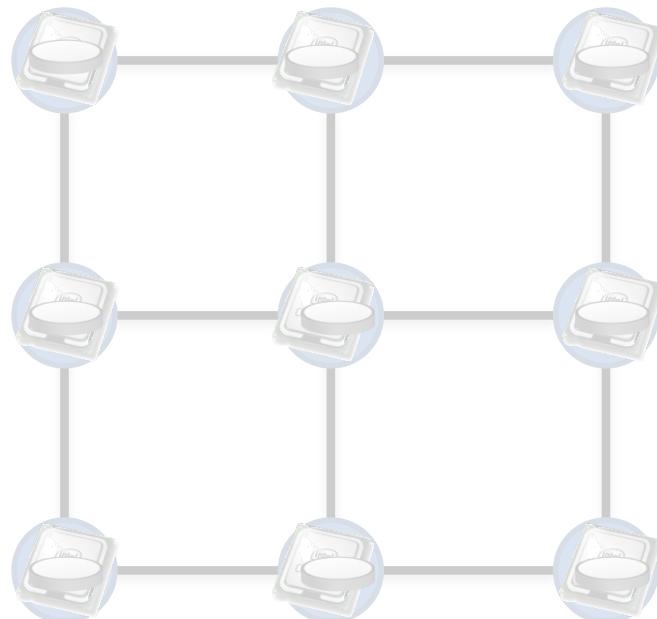


*Can we sample multiple
variables in parallel?*

From the original paper on Gibbs Sampling:

*“...the MRF can be divided into collections of [variables] with each collection assigned to an **independently** running **asynchronous processor**.”*

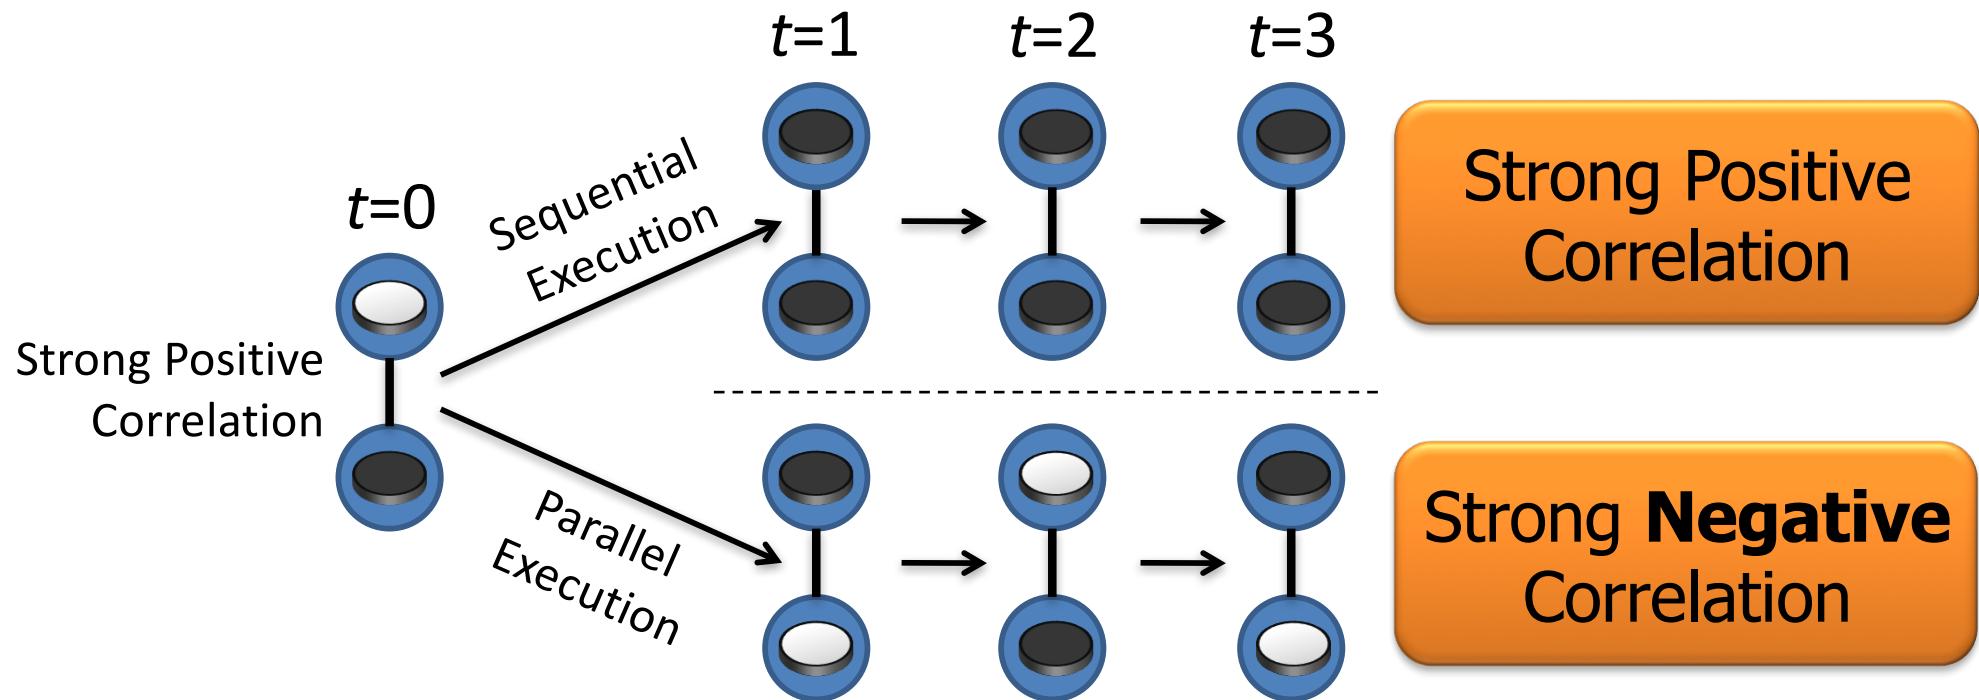
-- Stuart and Donald Geman, 1984.



Embarrassingly
Parallel!

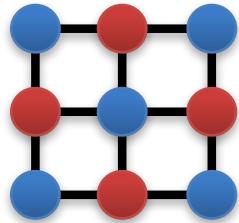
Converges to the
wrong distribution!

The problem with Synchronous Gibbs sampling

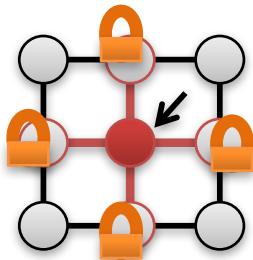


- *Adjacent variables **cannot** be sampled simultaneously.*

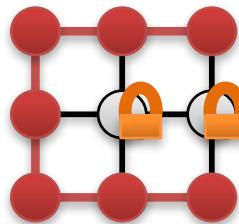
Introduced Three Convergent Samplers



Chromatic: Use graph coloring to synchronously sample independent sets



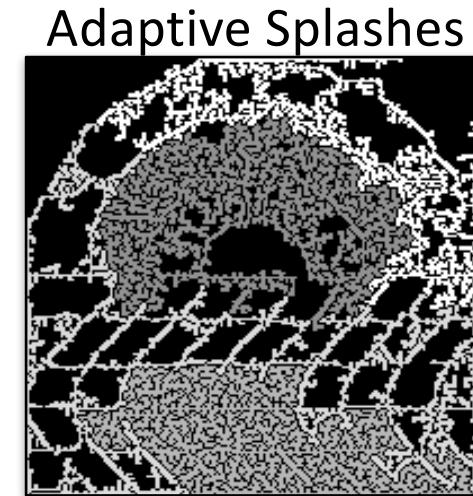
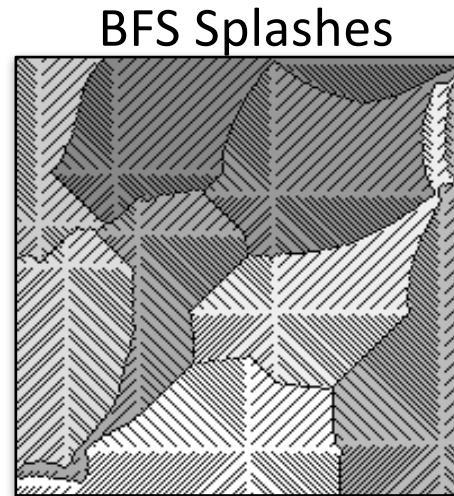
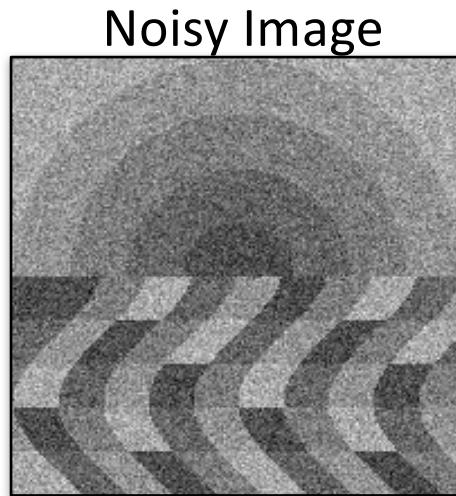
Asynchronous: Markov Blanket Locks ensure serializable execution



Splash: Adaptively constructs thin junction tree blocks

Dynamically Prioritized Sampling

- Prioritize Gibbs updates
- Adapt the **shape** of the Splash to span strongly coupled variables:



Theorem: Chromatic Sampler

- **Ergodic:** converges to the correct distribution
 - Based on graph coloring of the Markov Random Field
- **Quantifiable acceleration in mixing**

Time to update
all variables once

$$O\left(\frac{n}{p} + k\right)$$

The diagram shows the time complexity expression $O\left(\frac{n}{p} + k\right)$. Three blue arrows point from the terms to the right: one from $\frac{n}{p}$ to "# Variables", one from k to "# Colors", and one from p to "# Processors".

Variables
Colors
Processors

Theorem

Asynchronous and *Splash Gibbs Sampler*

- **Ergodic:** converges to the correct distribution
 - Requires vanishing adaptation
 - Corrected an error in a result by Levin & Casella *J. Multivar. Anal.* '06
- **Expected Parallelism:**

$$\mathbf{E}(\#\text{active processors})$$

$$\geq 1 + (p - 1) \left(1 - (p - 1) \left(\frac{d + 1}{n} \right) \right)$$

Processors Max Degree
 # Variables

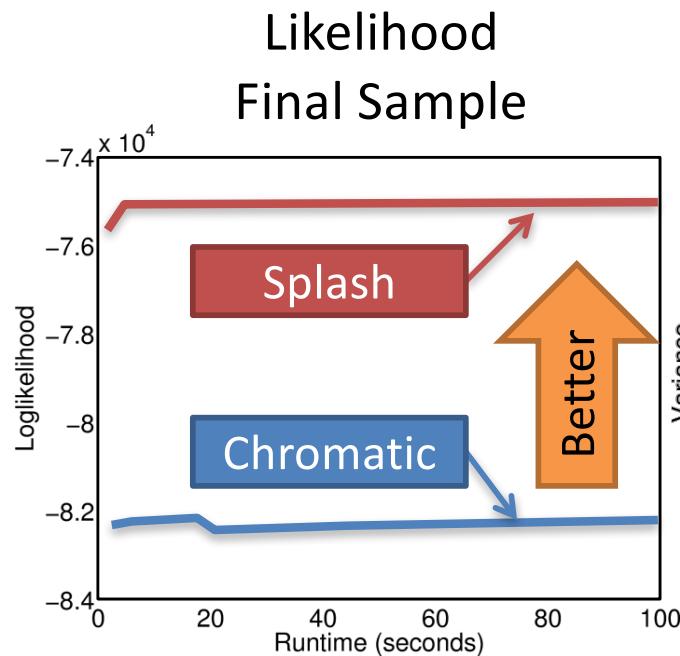
Evaluation

- Implemented multicore version:
 - Built using a GraphLab prototype
 - Substantially shorter development time
 - Novel **junction tree** construction algorithm
 - Markov blanket locking protocol
- Evaluated on large real-world problems

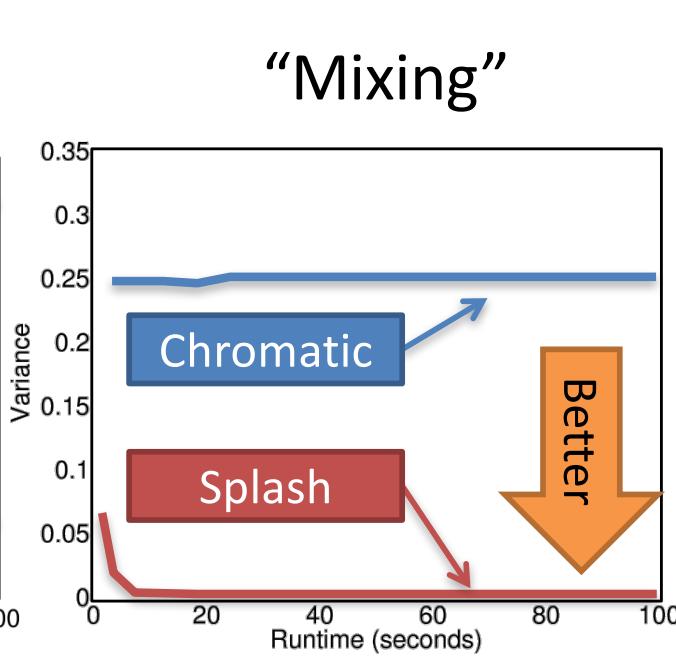
Experimental Results

- Markov logic network with **strong dependencies**

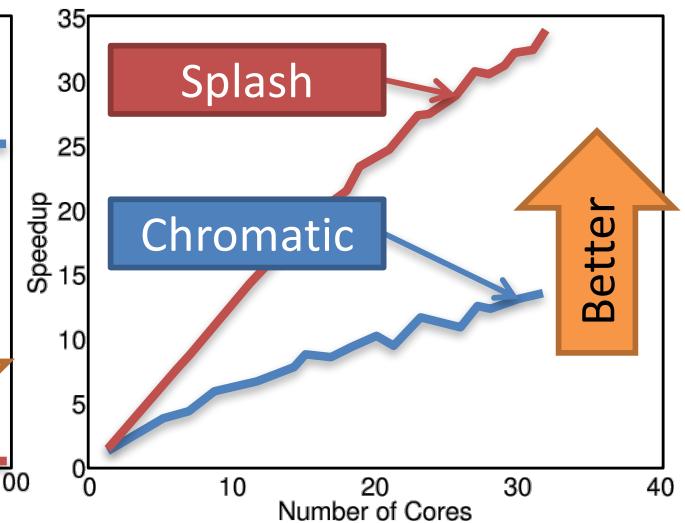
10K Variables



28K Factors



Speedup in Sample Generation



- The *Splash* sampler outperforms the *Chromatic* sampler on models with **strong dependencies**

Contributions: Gibbs Sampling

- Proposed **three convergent** Gibbs samplers
 - Chromatic, Asynchronous, Splash
 - Spectrum partially synchronous to asynchronous
 - New algorithms for junction tree construction
- Theoretical analysis of parallel Gibbs sampling
 - Convergence of asynchronous blocking
 - Relate parallelism to model structure
 - Stationary distribution of synchronous sampler
- Experimental analysis on real-world problems and systems

GrAD Methodology

- **Graphical**
 - Gibbs updates depend only on neighbors in MRF
- **Asynchronous**
 - Graph *Coloring* and *Markov Blanket Locks*
- **Dynamic**
 - Prioritized updates and adaptive Splash

Related Work

Ergodic (Convergent)

- Geman & Geman. Pami '84
- **Trees:** Hamze et al. UAI'04
- **Dynamic Blocking:** Barbu et al. IEEE Trans Pattern Analysis '05

Thesis
Chromatic,
Asynchronous,
and
Splash Gibbs

Parallel & Distributed

- LDA & Bayesian Networks
- Newman et al. NIPS'07
 - Asuncion et al. NIPS'08
 - Yan et al. NIPS'09

Amr et al. WSDM'12

- Asynchronous approximations empirically perform well



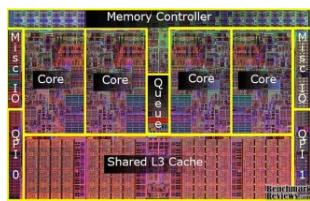
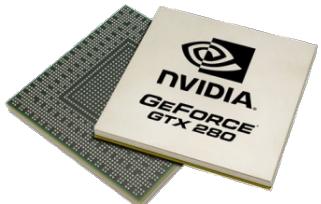
Massive Structured Problems

Probabilistic Graphical Models

**Parallel and Distributed Algorithms
for Probabilistic Inference**

GraphLab & PowerGraph

Advances Parallel Hardware



Parallel Algorithms for Probabilistic Inference

GraphLab & PowerGraph

Parallel Hardware

Joint Work With

*Yucheng Low Aapo Kyrola Haijie Gu Danny Bickson
Carlos Guestrin Joe Hellerstein Guy Blelloch David O'Hallaron*

Published Results

UAI'10 VLDB'12

How do we design and implement GrAD Algorithms

We could:

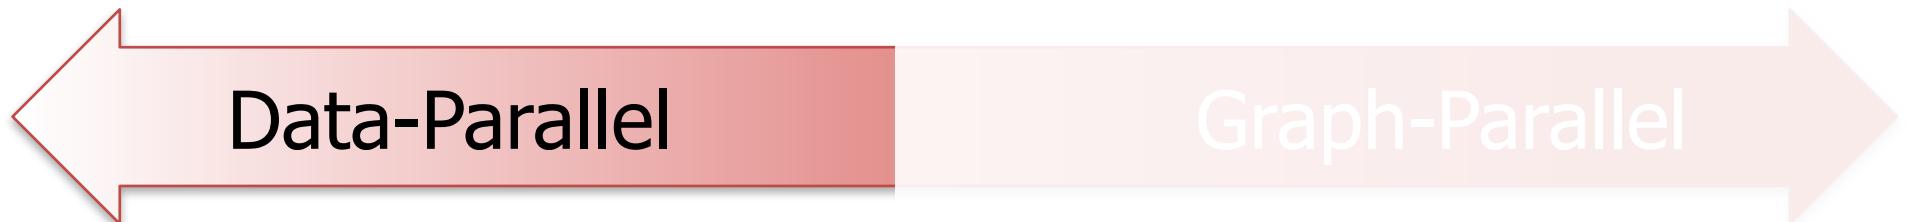
- *design and implement* for each architecture?
 - **Time consuming**
 - Repeatedly solving the same system problems
- use high-level abstractions like **MapReduce**?
 - Unable to express:

- Graphical
- Asynchronous
- Dynamic

GrAD Methodology

Solution: GraphLab

GraphLab is a Graph-Parallel Abstraction



Map Reduce

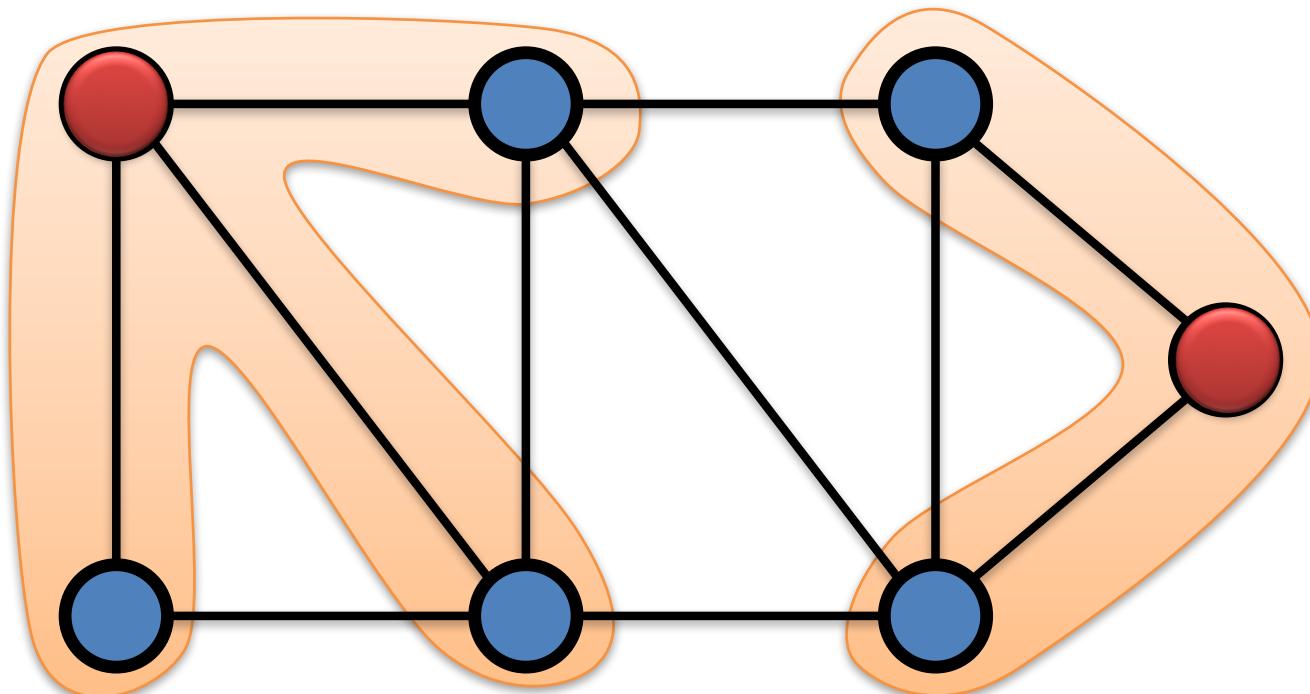
- *Independent Data*
- *Single Pass*
- *Synchronous*

GraphLab

- Graph Structured Data
- Iterative Computation
- Dynamic + Asynchronous

The GraphLab Abstraction

- A user-defined **Vertex Program** runs on each vertex
- **Graph** constrains **interaction** along edges
 - Directly **read** and **modify** the state of adjacent vertices and edges
- **Parallelism:** run multiple vertex programs simultaneously



The GraphLab Vertex Program

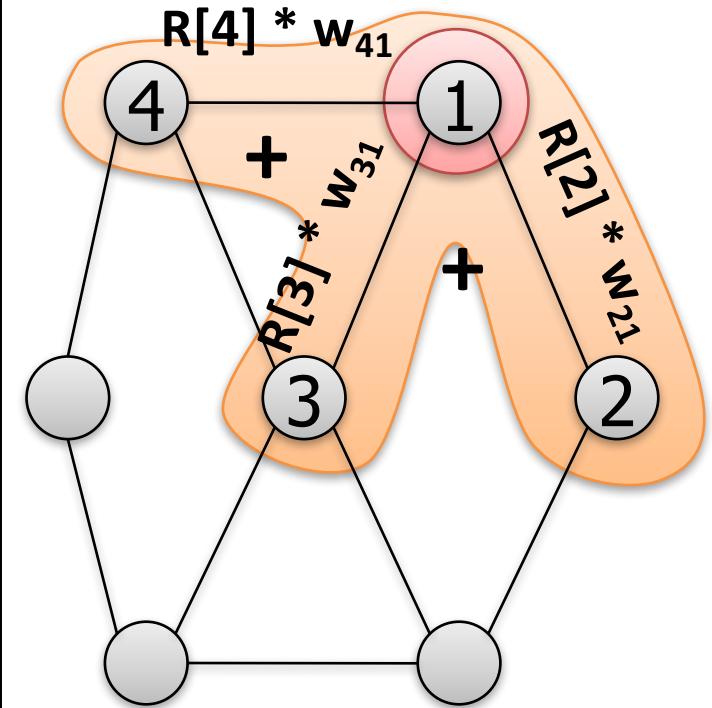
Vertex Programs directly **access** adjacent vertices and edges

```
GraphLab_PageRank(i)
```

```
// Compute sum over neighbors
total = 0
foreach( j in neighbors(i)):
    total = total + R[j] * wji
```

```
// Update the PageRank
R[i] = total
```

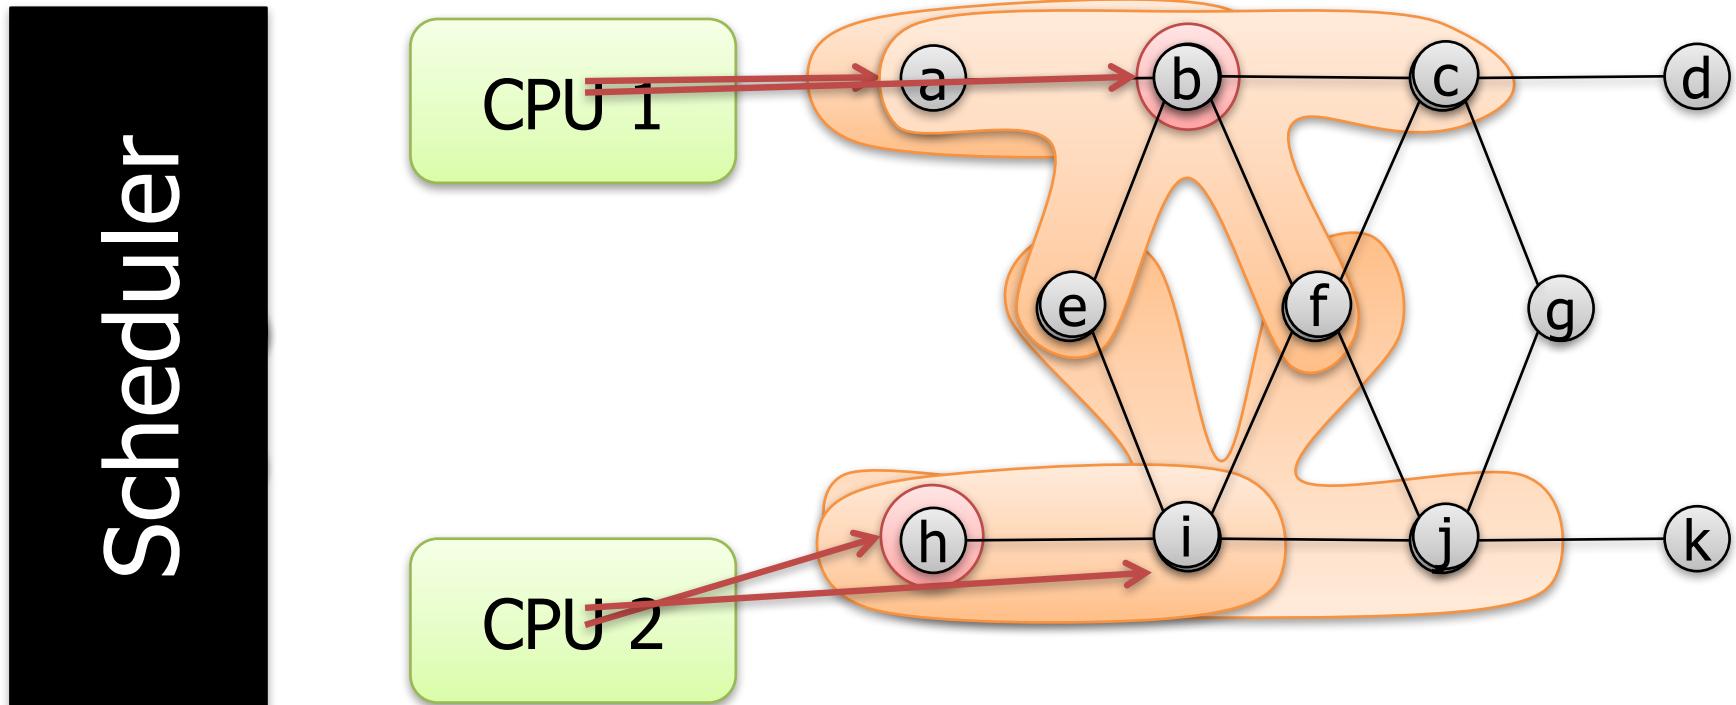
```
// Trigger neighbors to run again
priority = |R[i] - oldR[i]|
if R[i] not converged then
    signal neighbors(i) with priority
```



Dynamics

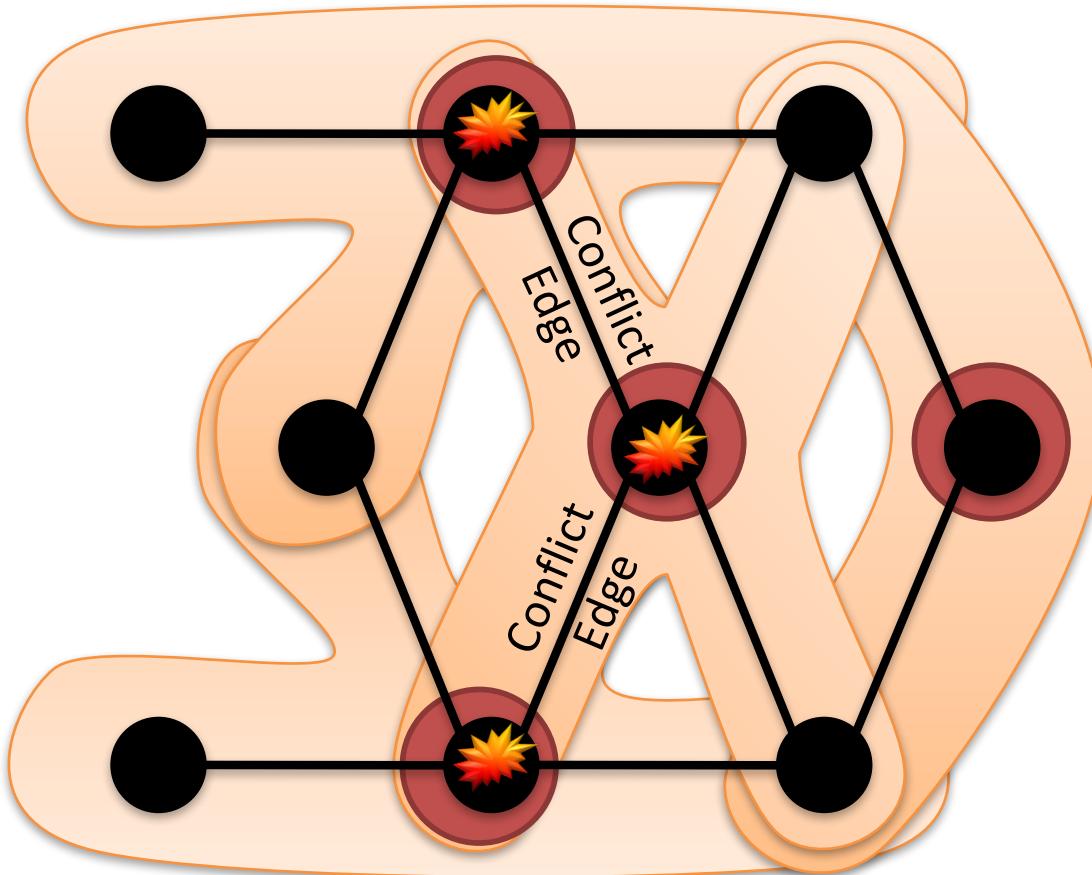
GraphLab is Asynchronous

The **scheduler** determines the order that vertices are executed



Scheduler can **prioritize** vertices.

GraphLab is Serializable



- Automatically ensures **serializable** executions

Serializable Execution

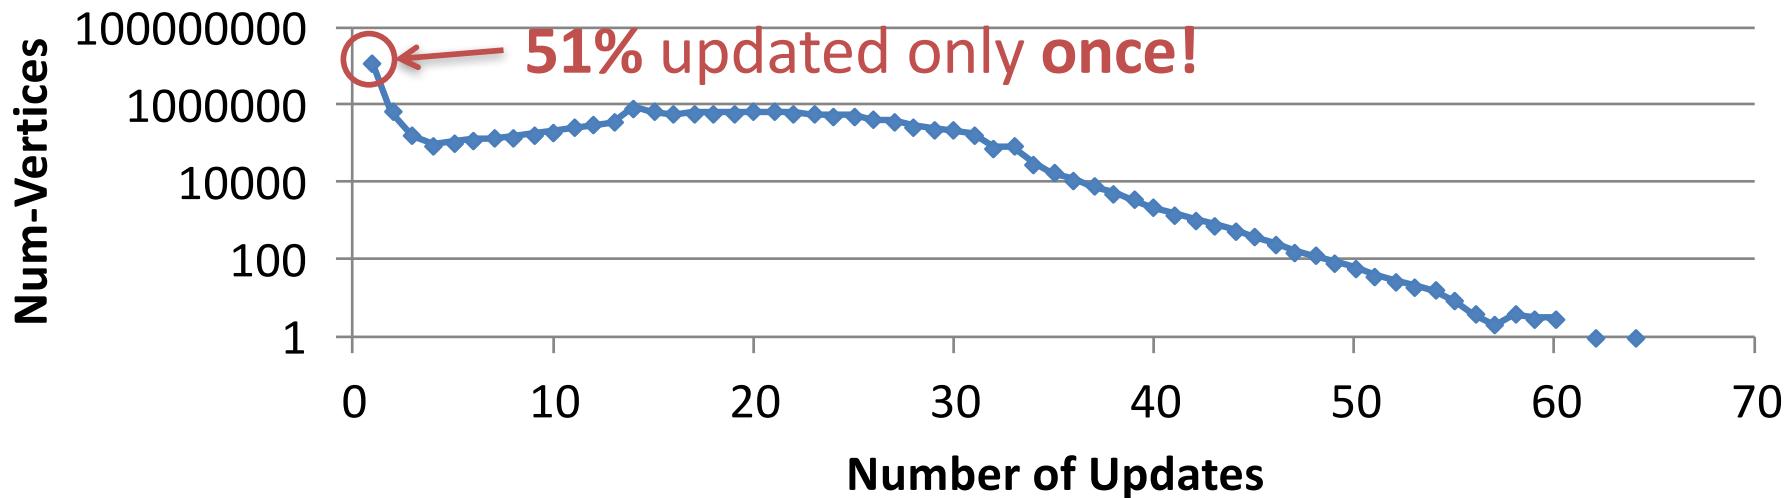
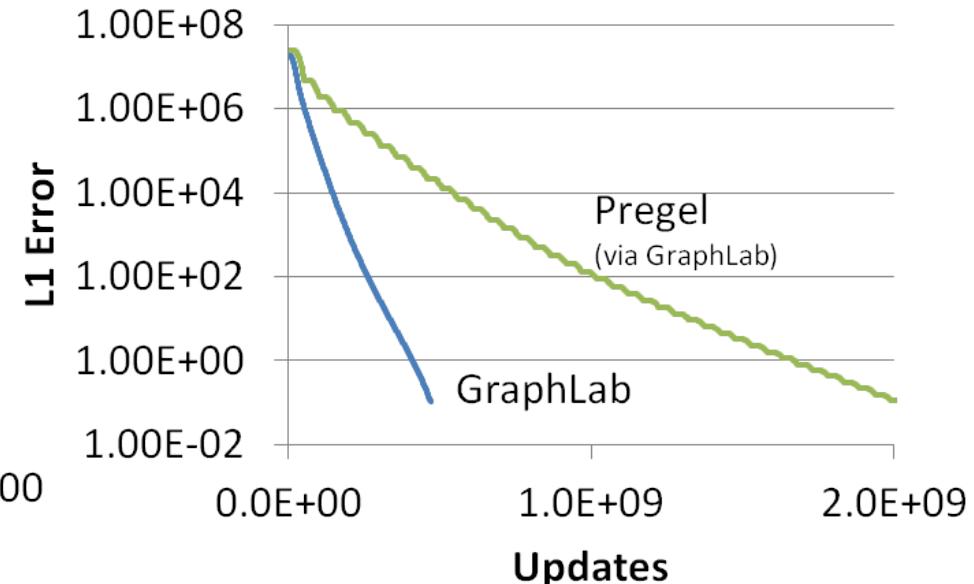
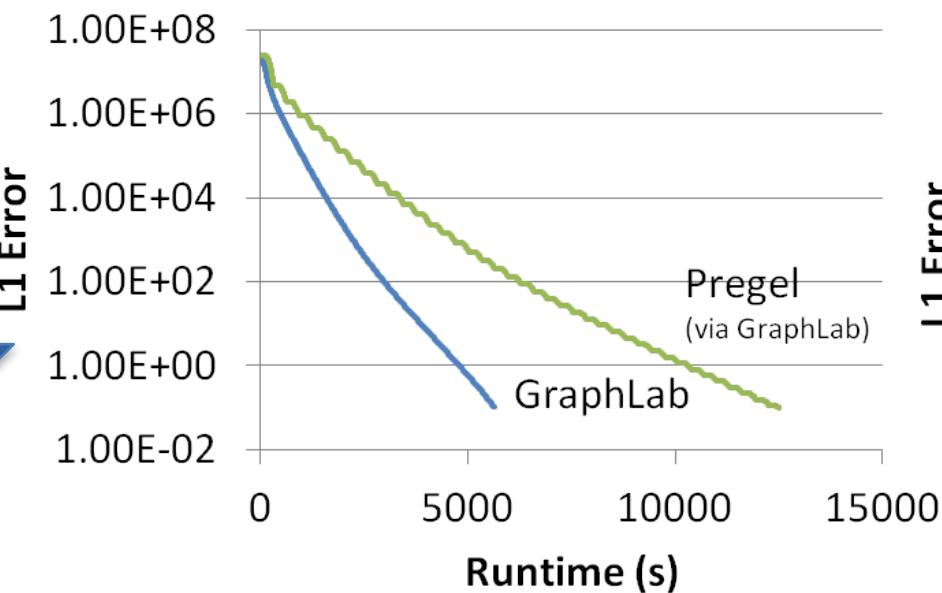
For **each parallel execution**, there exists a **sequential execution** of vertex-programs which produces the same result.



The GraphLab System

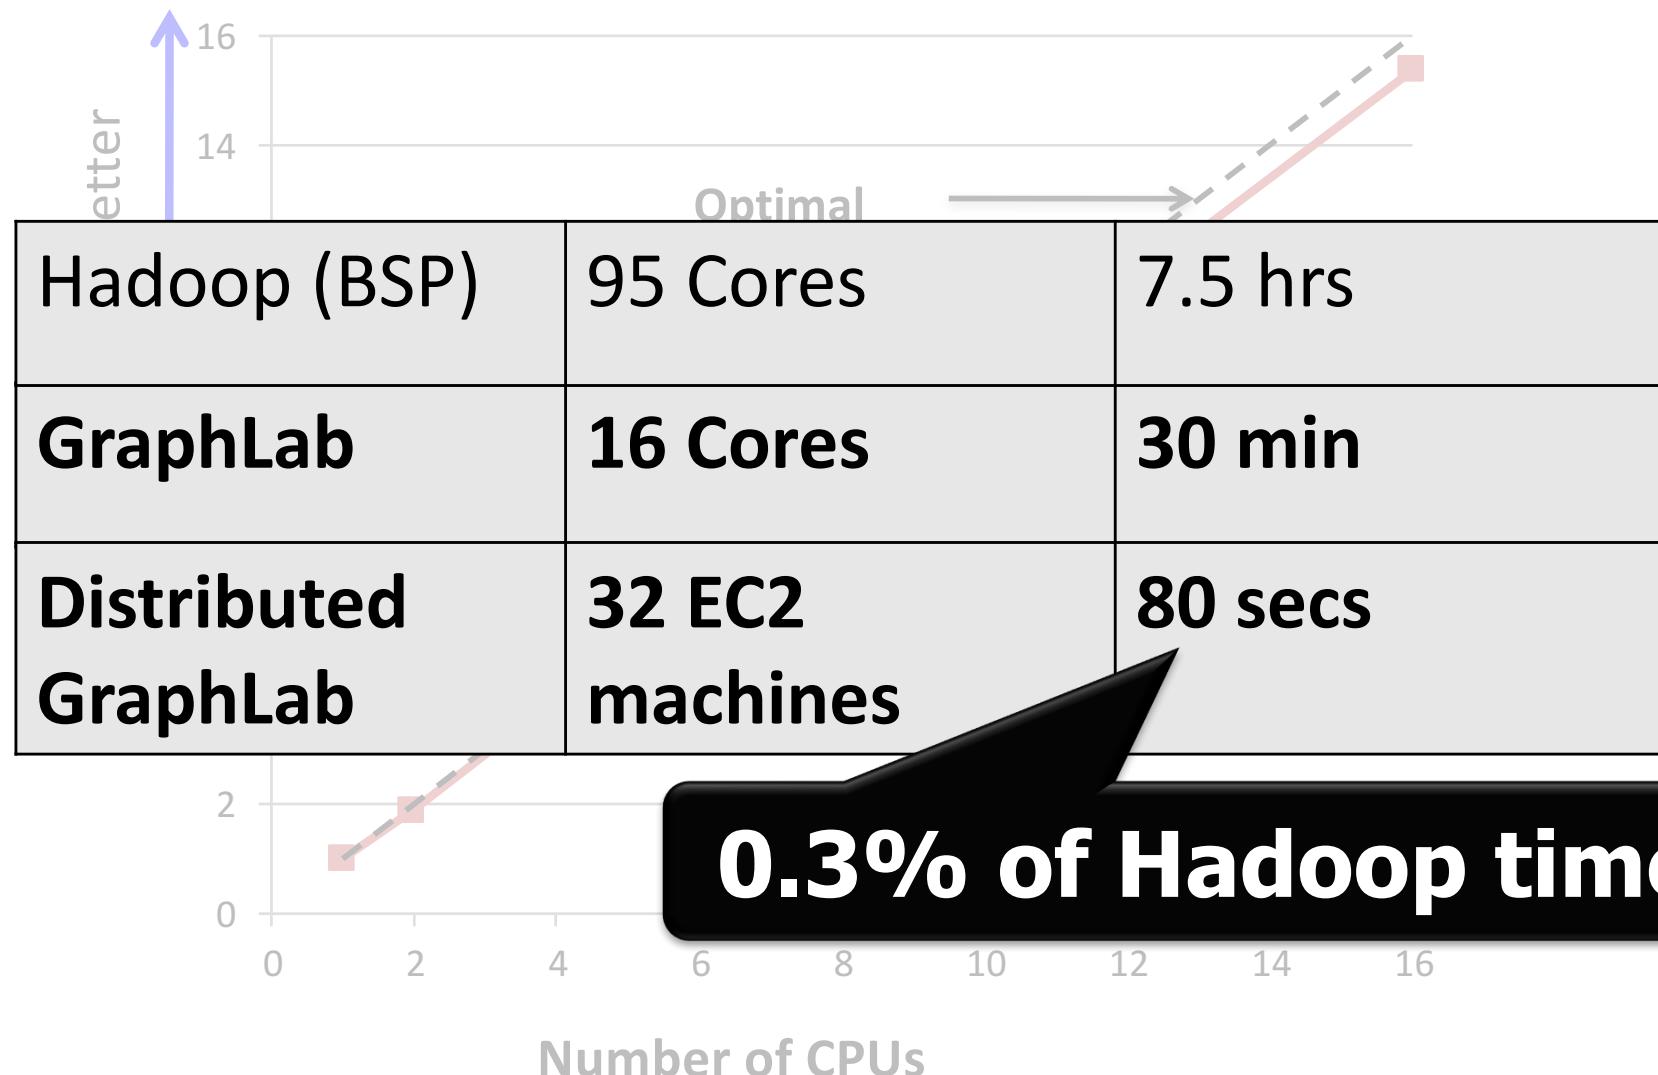
- Implemented as a C++ API
 - Widely downloaded open-source project
- **Multicore and distributed** versions:
 - **Hide Latency:** Pipelined locking
 - **Fault Tolerance:** Chandy-Lamport Snapshot
- Tested on a wide range of ML algorithms
 - ALS, BP, Gibbs, Lasso, CoEM, SVM, LDA, ...

GraphLab vs. Pregel (BSP)



PageRank (25M Vertices, 355M Edges)

Never Ending Learner Project (CoEM)



Summary: GraphLab

- **Generalizes** the GrAD Methodology
 - ALS, BP, Gibbs, Lasso, CoEM, SVM, PageRank, LDA, ...
- **Simplifies** the *design* and *implementation* of GrAD Algorithms
- Substantially outperforms existing systems
- Key Contributions:
 - Formalized the graph-parallel setting
 - Isolates computation from movement of data
 - Strong serializability guarantees
 - Evaluation on a wide range of algorithms

Thus far...

GraphLab provided exciting scaling performance

But...

We couldn't scale up to
Altavista Webgraph from 2002
1.4B vertices, 6.6B edges

Parallel Algorithms for Probabilistic Inference

GraphLab & PowerGraph

Parallel Hardware

Joint Work With

*Yucheng Low Aapo Kyrola Haijie Gu Danny Bickson
Carlos Guestrin Joe Hellerstein Guy Blelloch David O'Hallaron*

Published Results

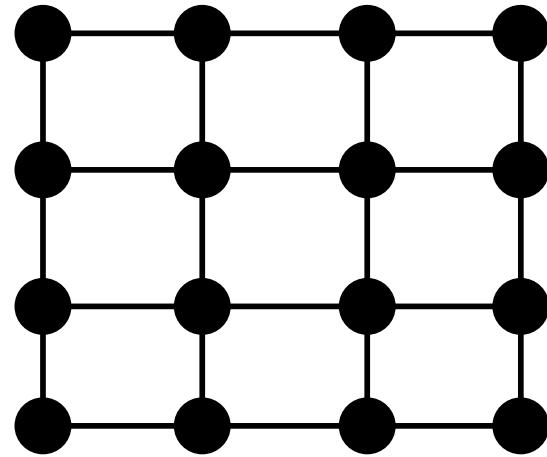
OSDI'12

Natural Graphs

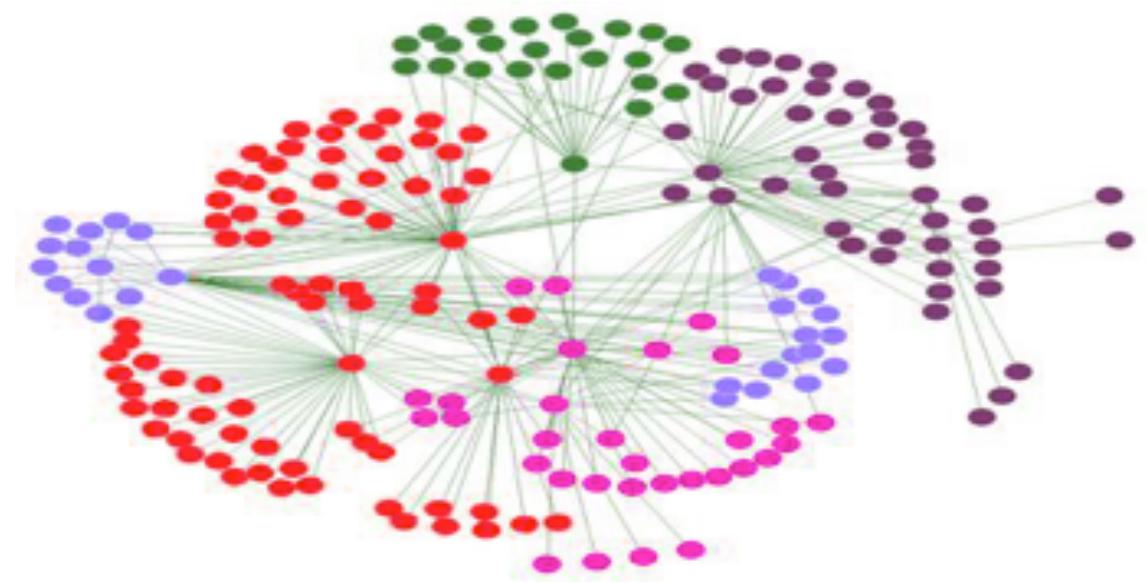
Graphs derived from natural
phenomena



Properties of Natural Graphs



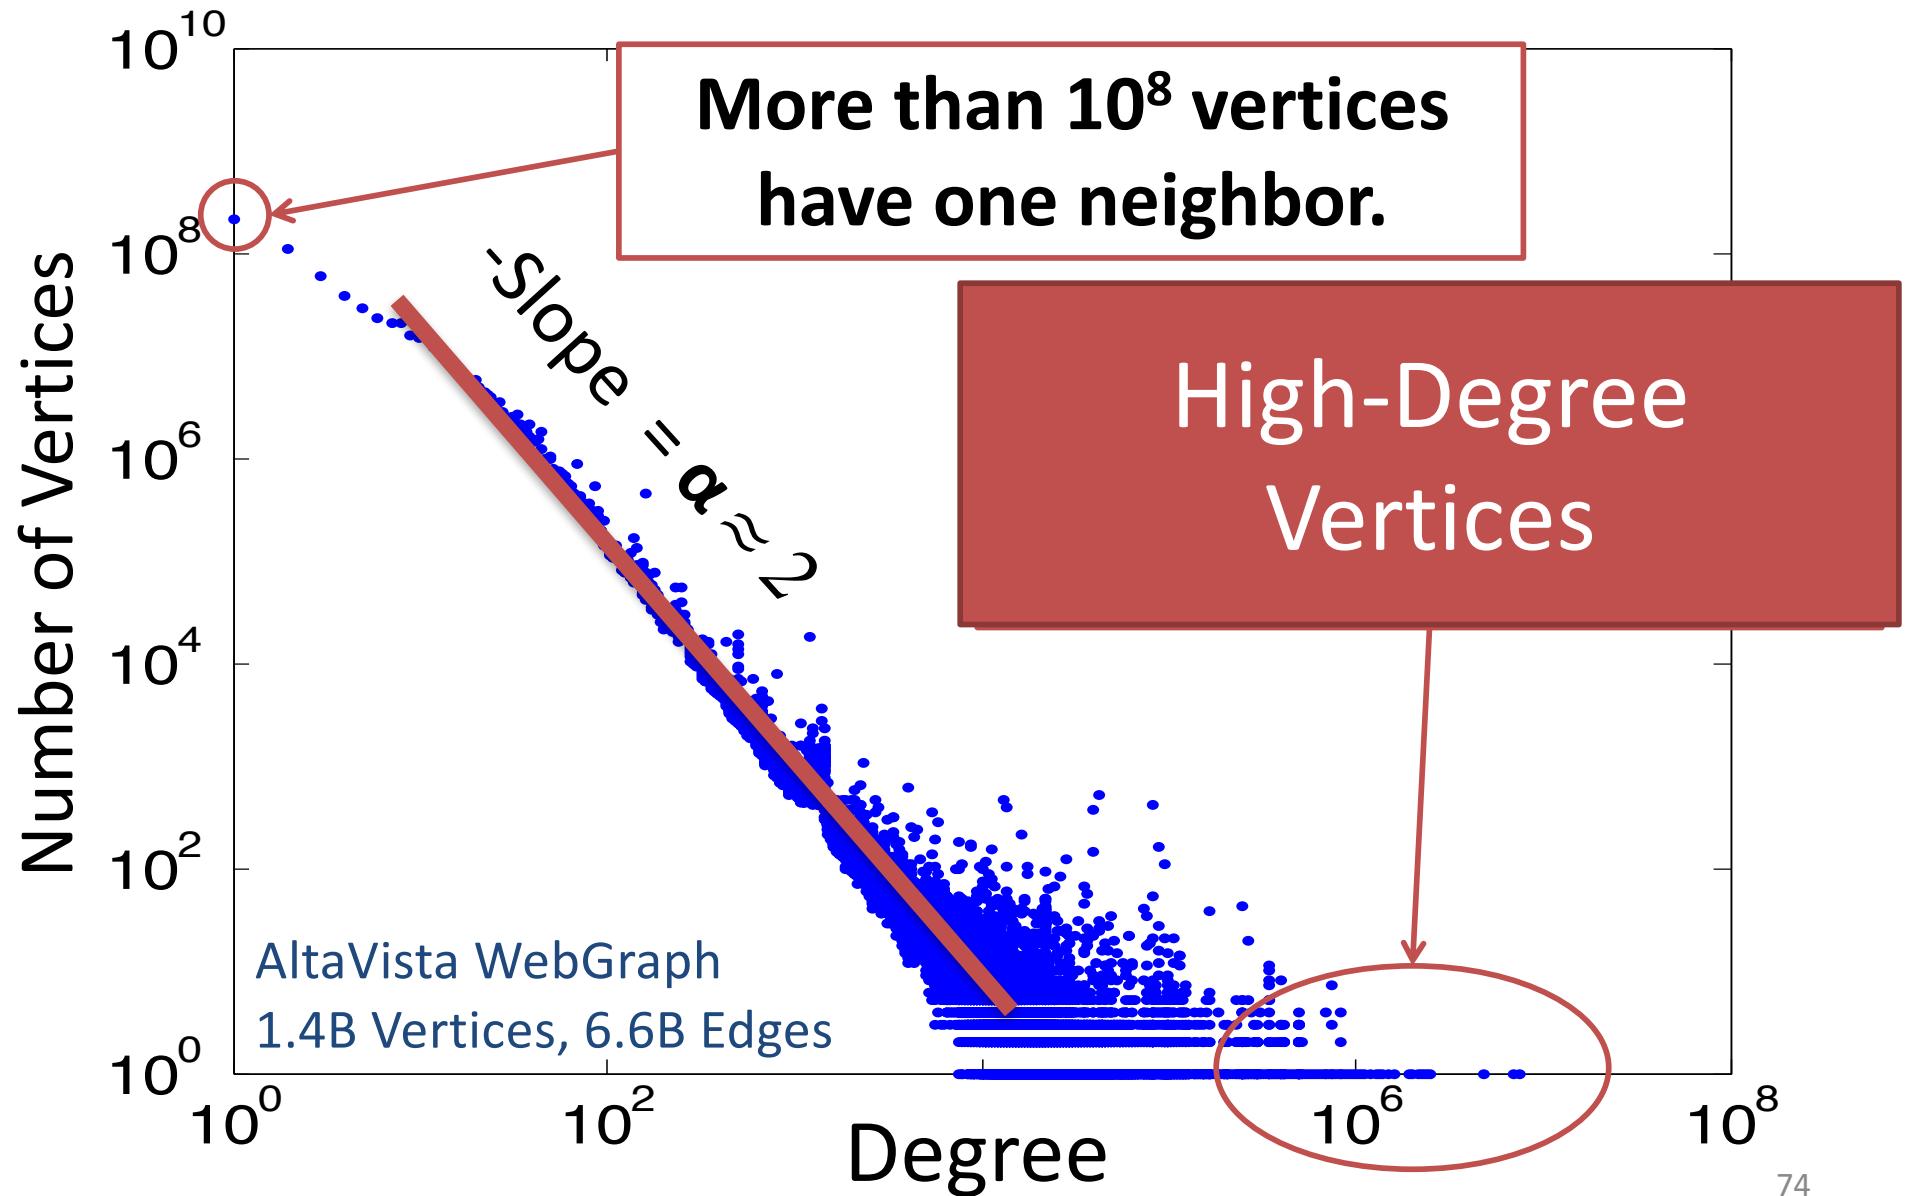
Regular Mesh



Natural Graph

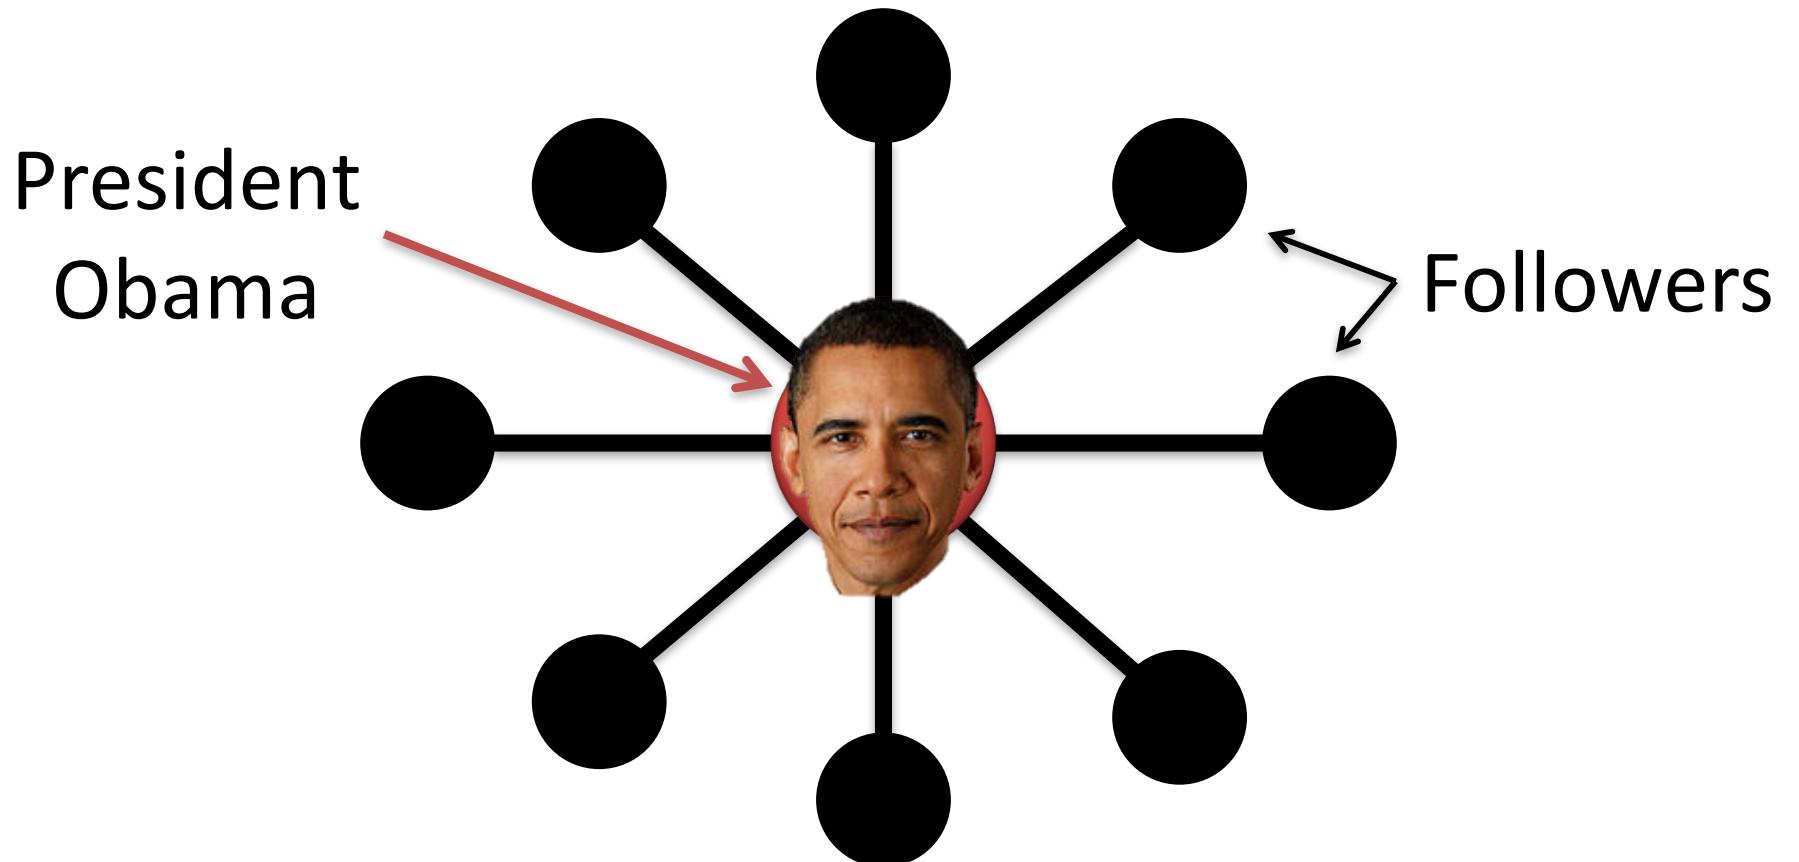
Power-Law Degree Distribution

Power-Law Degree Distribution

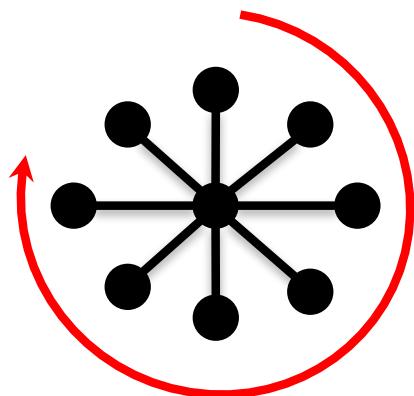


Power-Law Degree Distribution

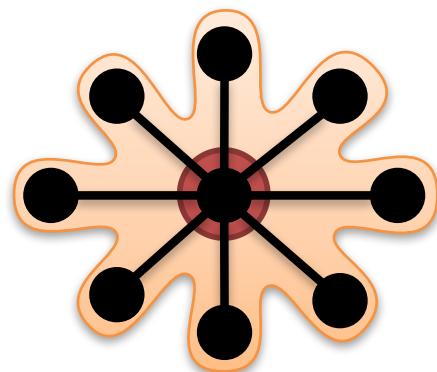
“Star Like” Motif



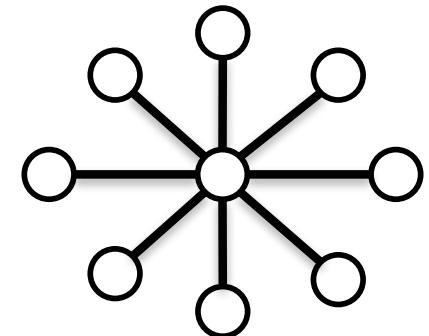
Challenges of High-Degree Vertices



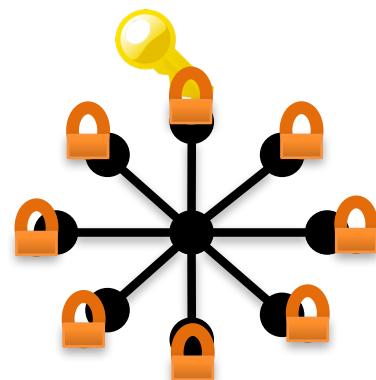
Sequentially process edges



Touches a large fraction of graph



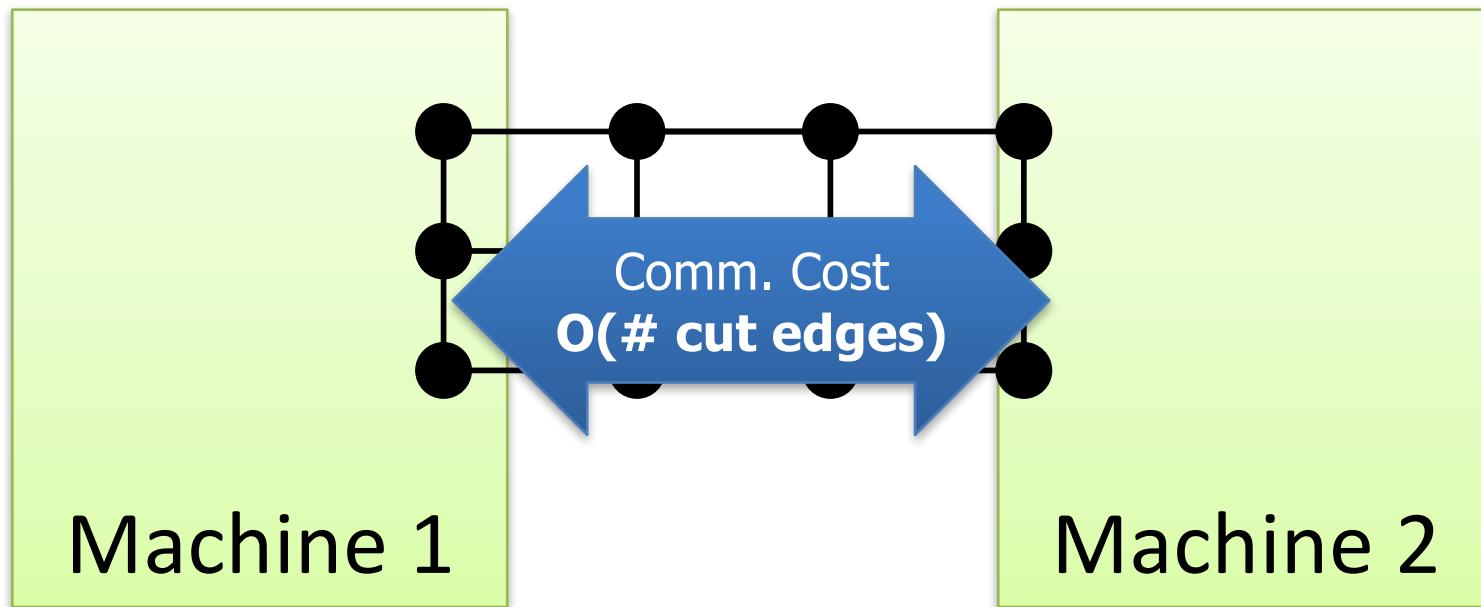
Edge meta-data too large for single machine



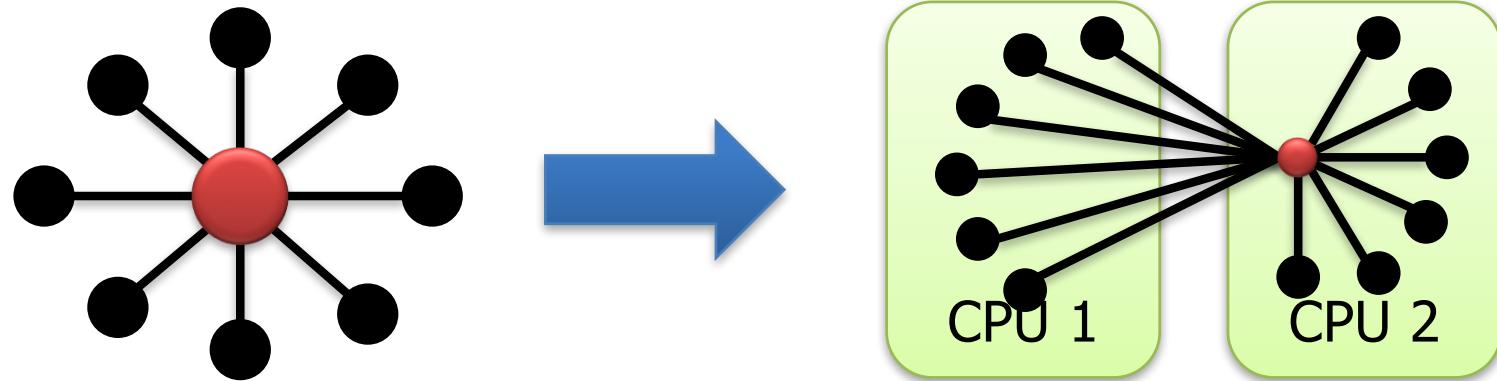
Serializability Requires
Heavy Locking

Graph Partitioning

- Graph parallel abstractions rely on partitioning:
 - Minimize communication
 - Balance computation and storage



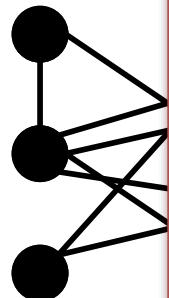
Power-Law Graphs are Difficult to Partition



- Power-Law graphs do not have **low-cost** balanced cuts [*Leskovec et al. 08, Lang 04*]
- Traditional graph-partitioning algorithms perform poorly on Power-Law Graphs.
[*Abou-Rjeili et al. 06*]

Random Partitioning

- GraphLab resorts to **random** (hashed) partitioning on **natural graphs**



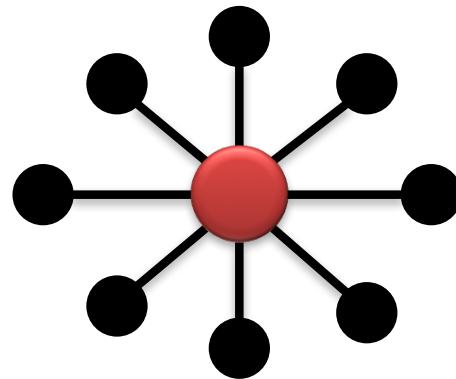
$$\mathbb{E} \left[\frac{|Edges\ Cut|}{|E|} \right] = 1 - \frac{1}{p}$$

10 Machines → 90% of edges cut

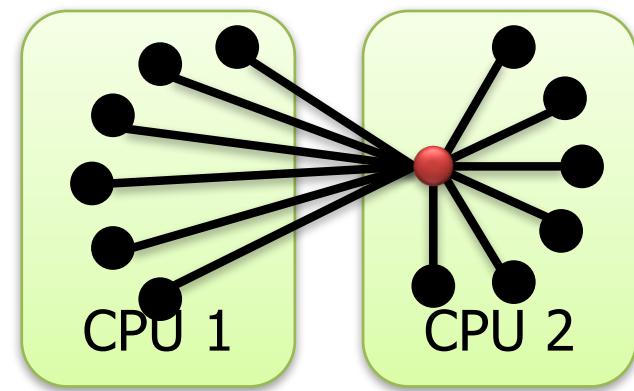
100 Machines → 99% of edges cut!

In Summary

GraphLab is not well suited for
natural graphs



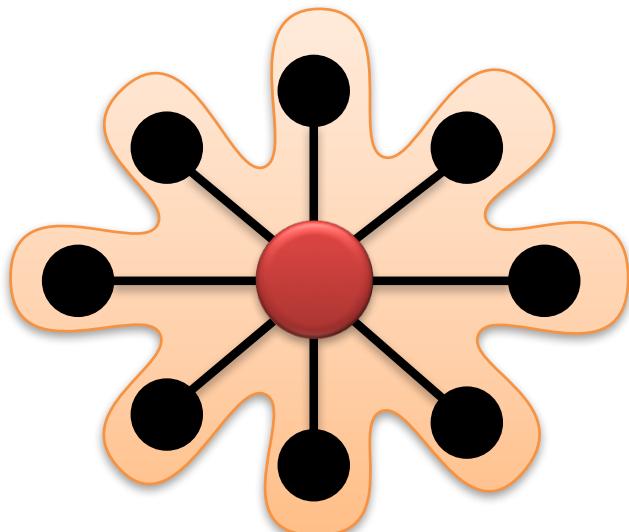
Challenges of **high-degree**
vertices



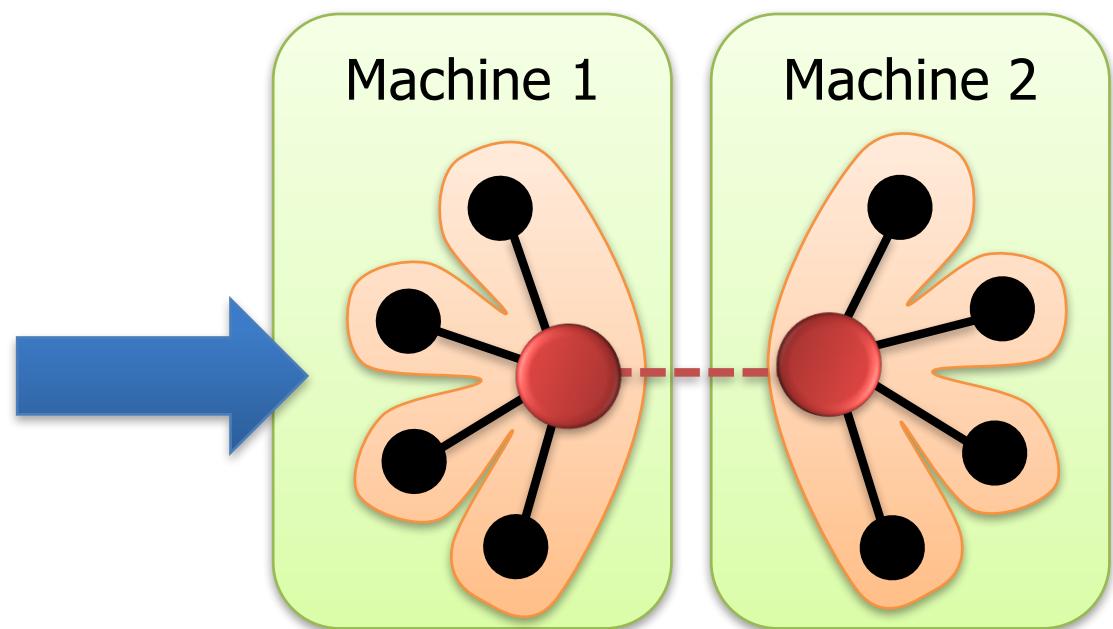
Low quality
partitioning

PowerGraph

Program
For This



Run on This



- Split **High-Degree** vertices
- New Abstraction → *Equivalence on Split Vertices*

A Common Pattern for Vertex-Programs

```
GraphLab_PageRank(i)
```

```
// Compute sum over neighbors  
total = 0  
foreach( j in neighbors(i)):  
    total = total + R[j] * wji
```

**Gather Information
About Neighborhood**

```
// Update the PageRank  
R[i] = total
```

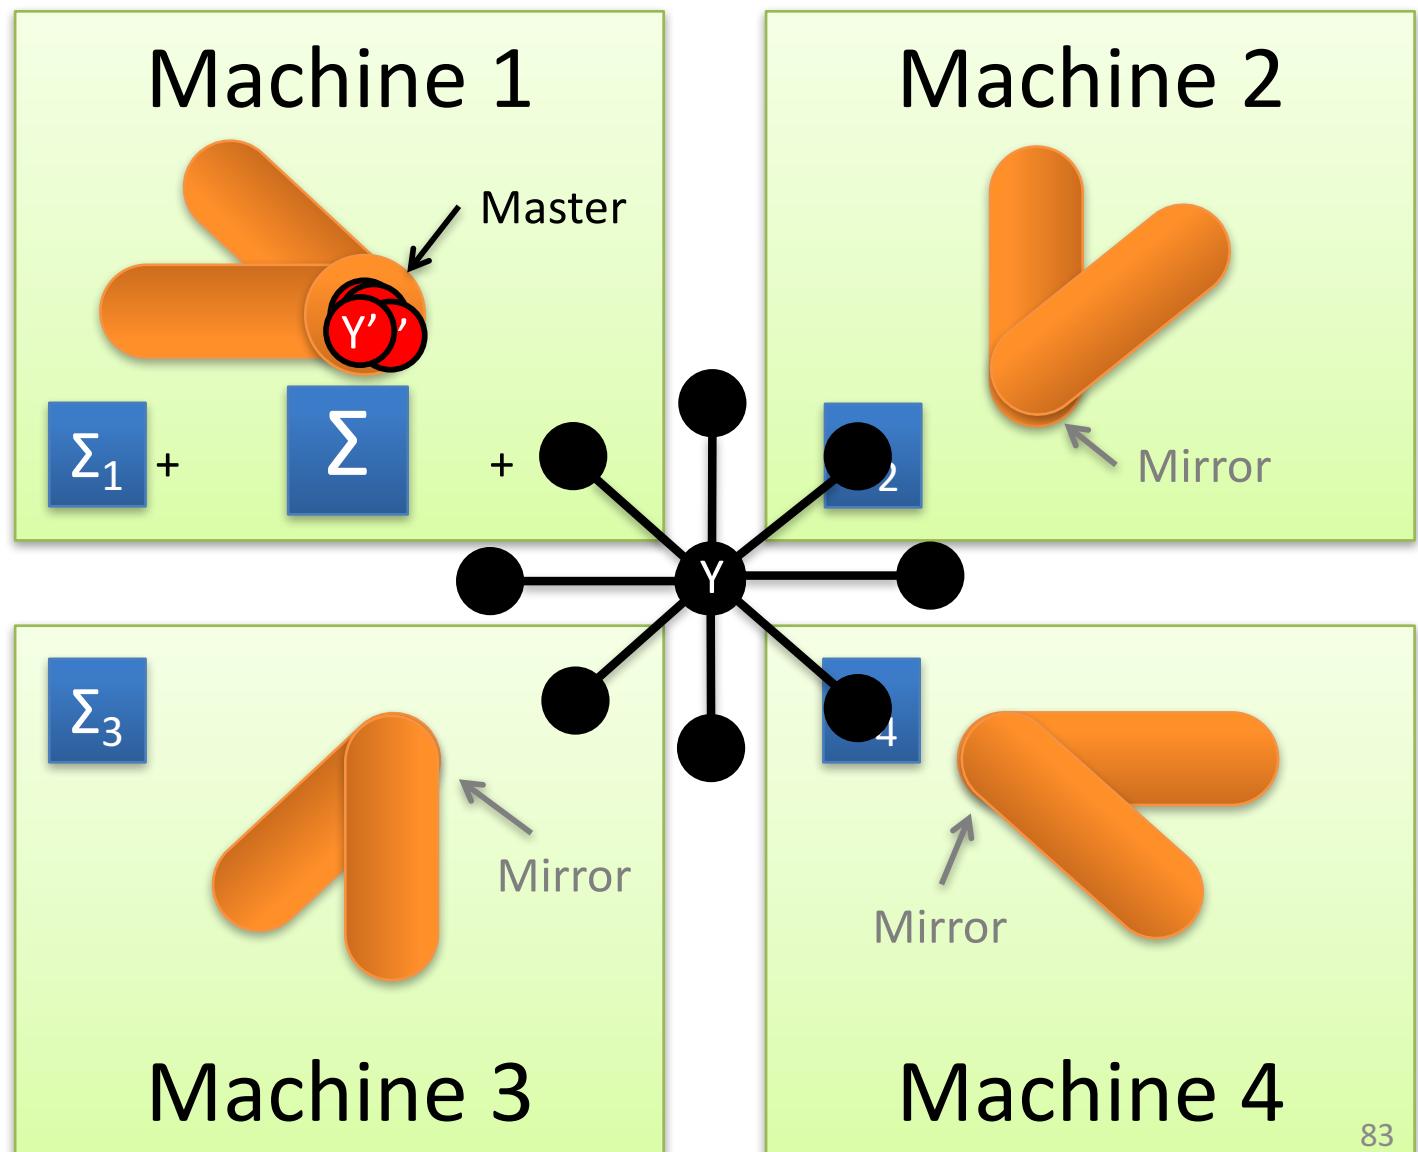
Update Vertex

```
// Trigger neighbors to run again  
priority = |R[i] - oldR[i]|  
if R[i] not converged then  
    signal neighbors(i) with priority
```

**Signal Neighbors &
Modify Edge Data**

GAS Decomposition

Gather
Apply
Scatter



Minimizing Communication in PowerGraph

New Theorem:

For any edge-cut we can directly construct a vertex-cut which requires strictly less communication and storage.

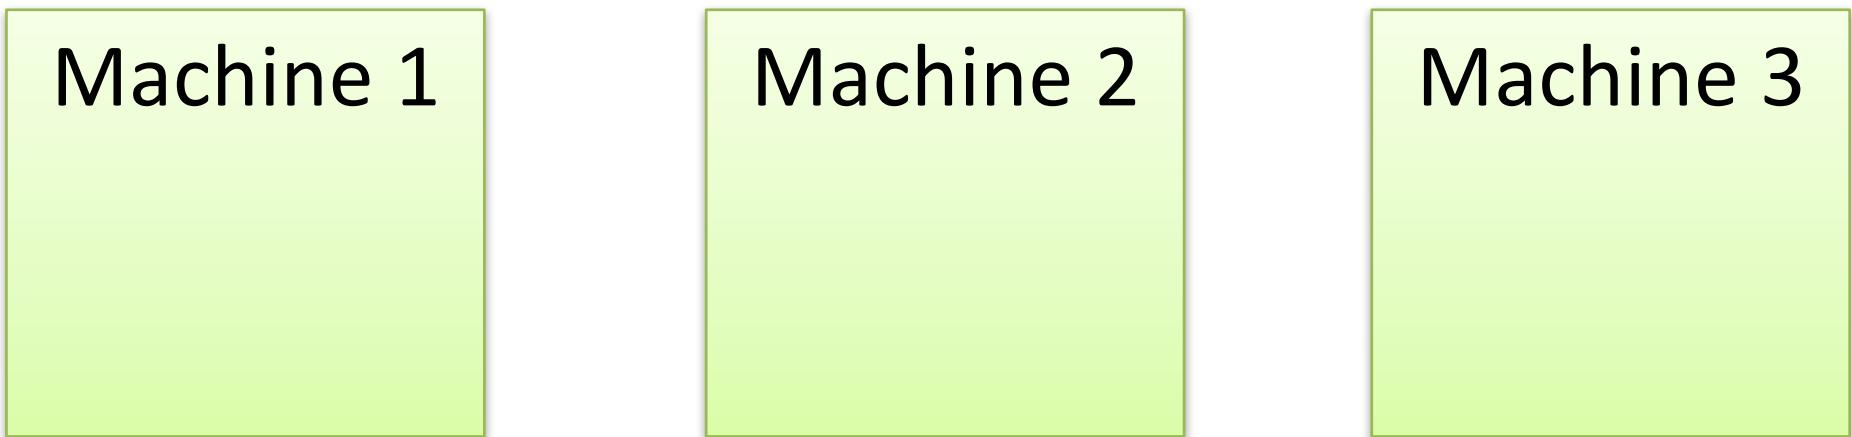
Percolation theory suggests that power law graphs have good vertex cuts. [Albert et al. 2000]

Constructing Vertex-Cuts

- **Evenly assign edges to machines**
 - Minimize machines spanned by each vertex
- Assign each edge **as it is loaded**
 - Touch each edge only once
- Propose two **distributed** approaches:
 - *Random Vertex Cut*
 - *Greedy Vertex Cut*

Random Vertex-Cut

- Randomly assign edges to machines

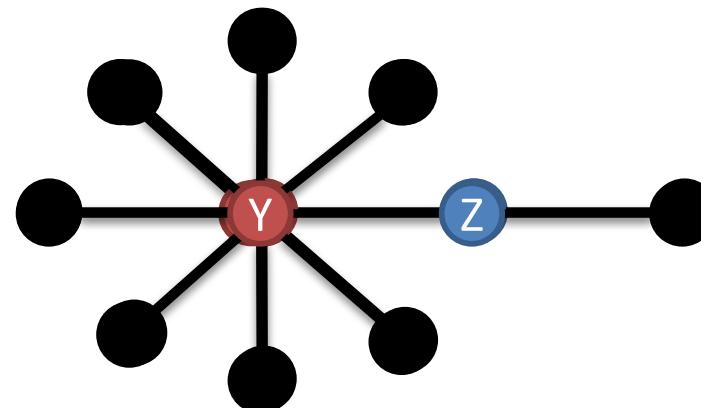


Balanced Vertex-Cut

 Y Spans 3 Machines

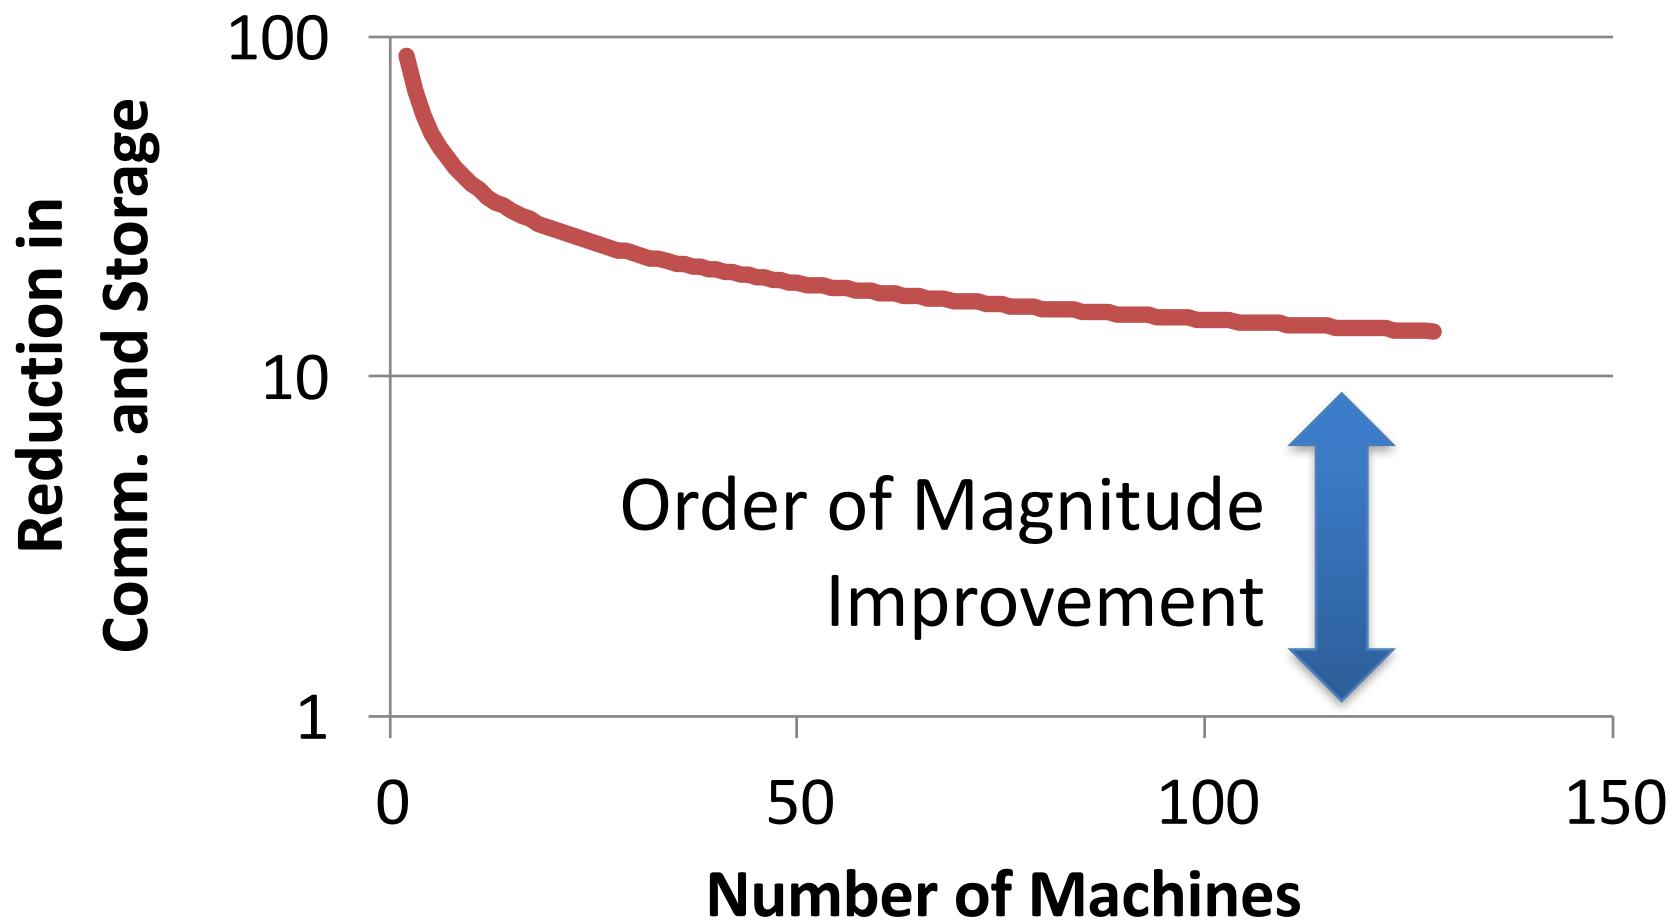
 Z Spans 2 Machines

 Not cut!



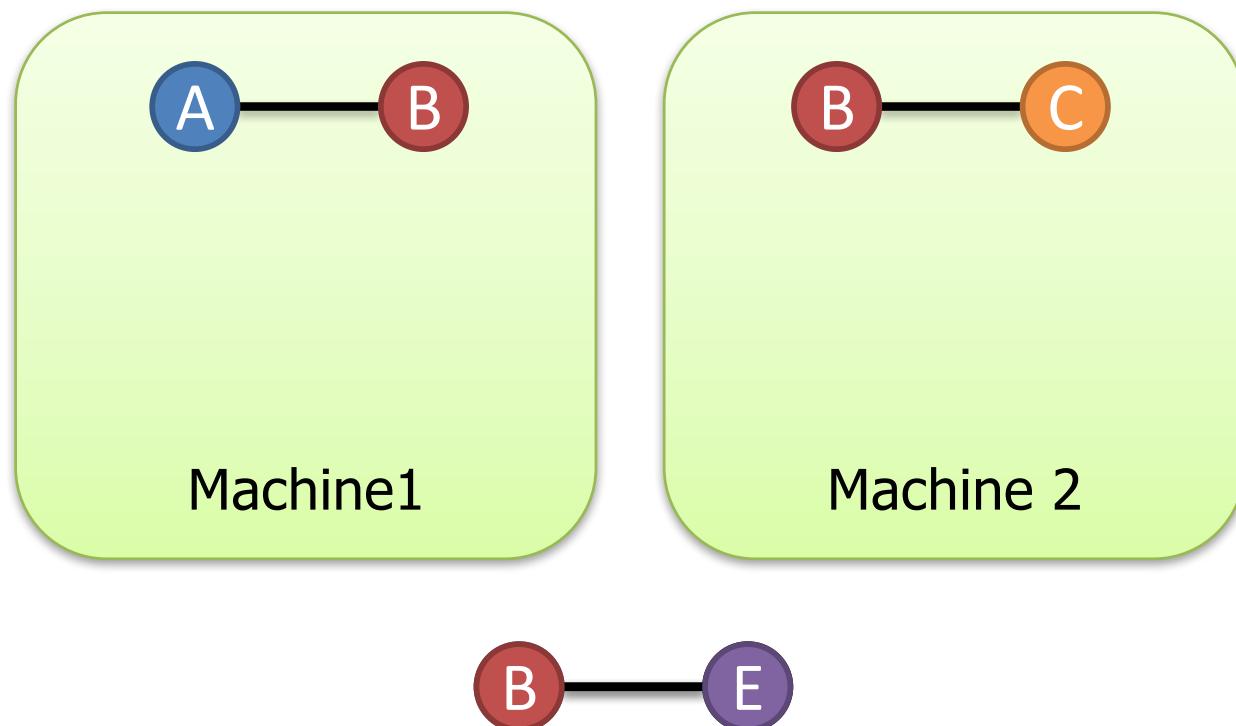
Random Vertex-Cuts vs. Edge-Cuts

- Expected improvement from vertex-cuts:

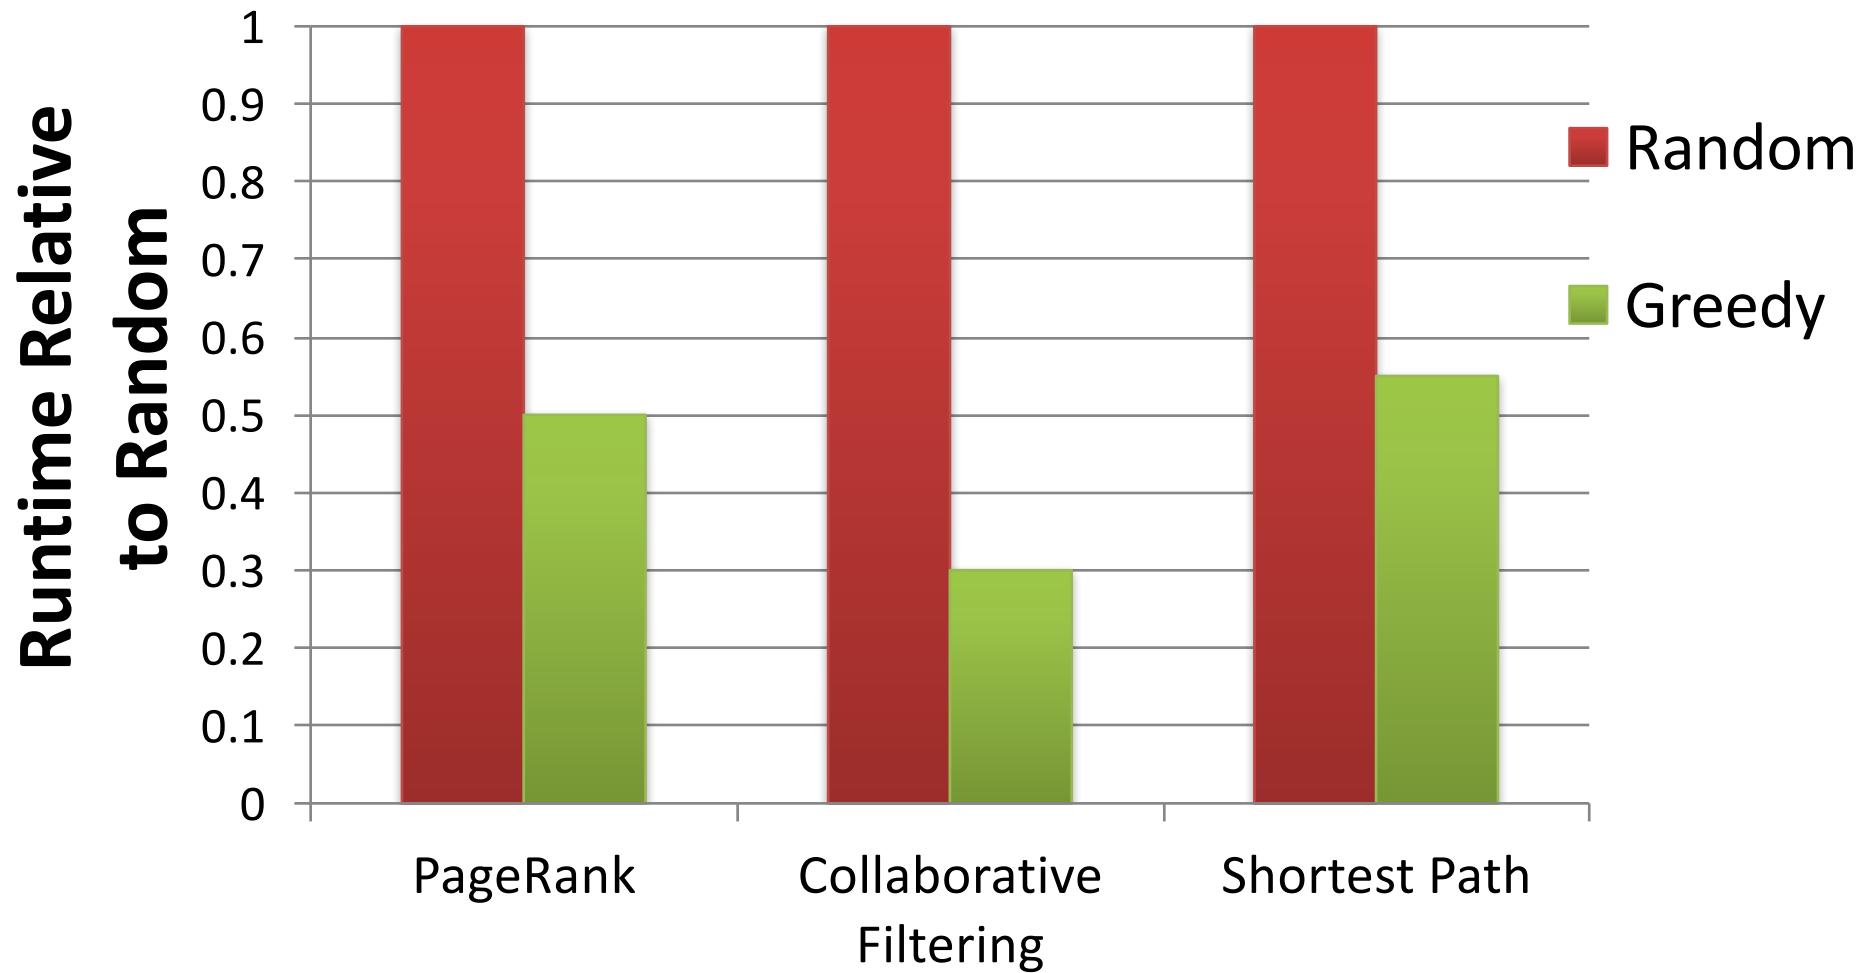


Greedy Vertex-Cuts

- Place edges on machines which already have the vertices in that edge.

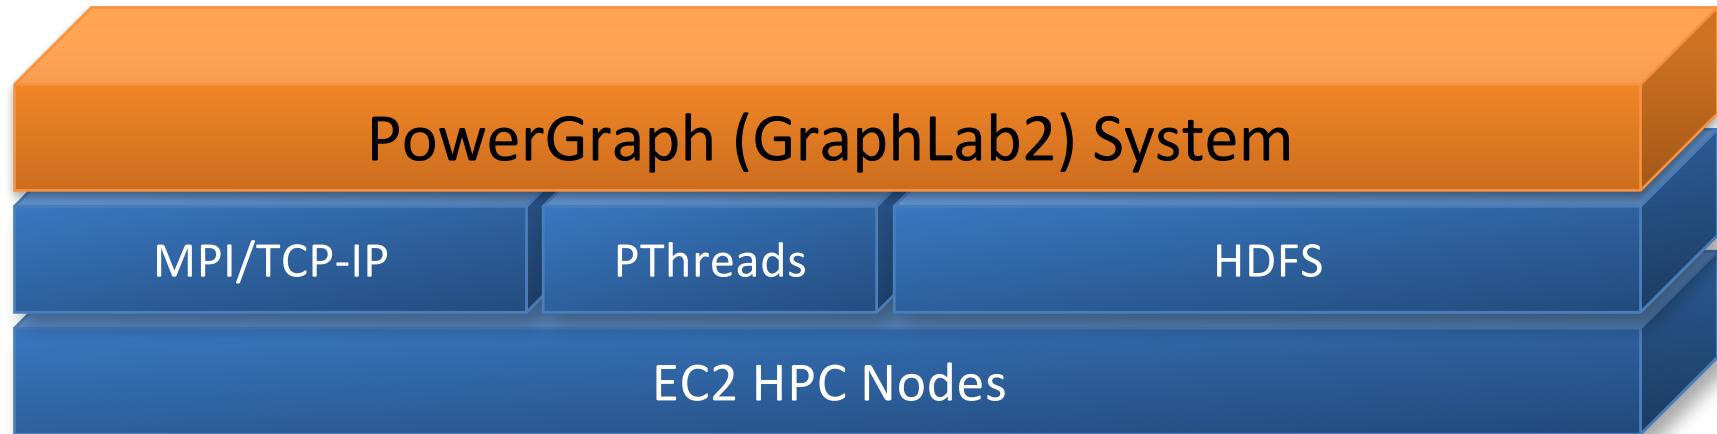


Greedy Vertex-Cuts Improve Performance



Greedy partitioning improves computation performance.

System Design

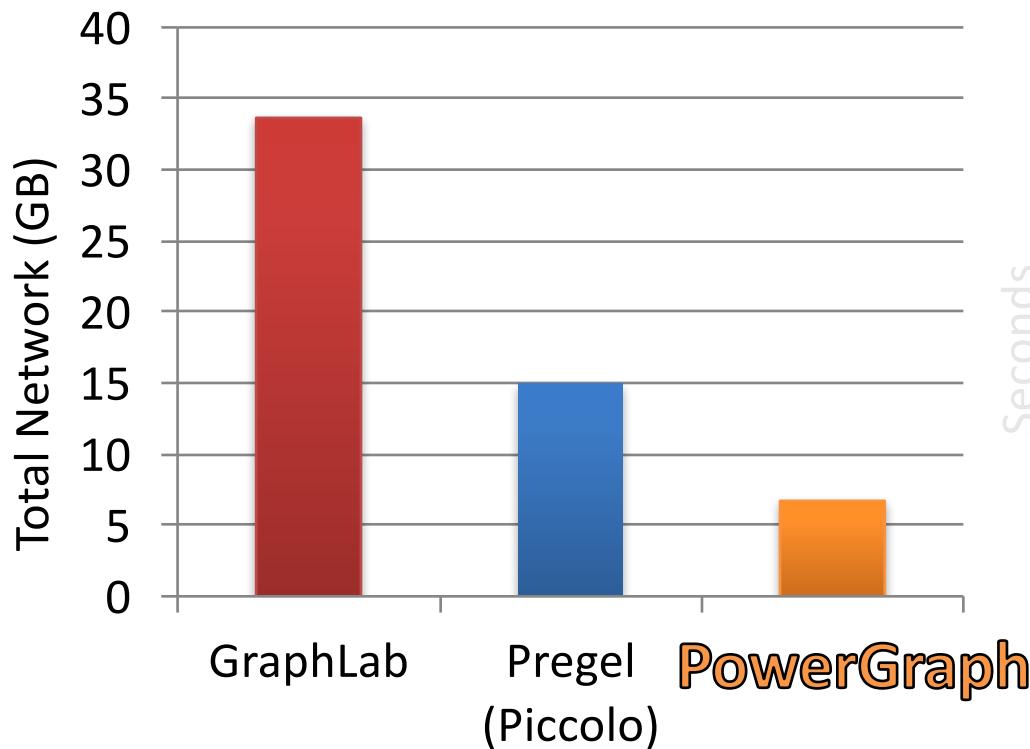


- Implemented as C++ API
- Uses HDFS for Graph Input and Output
- Fault-tolerance is achieved by check-pointing
 - Snapshot time < 5 seconds for twitter network

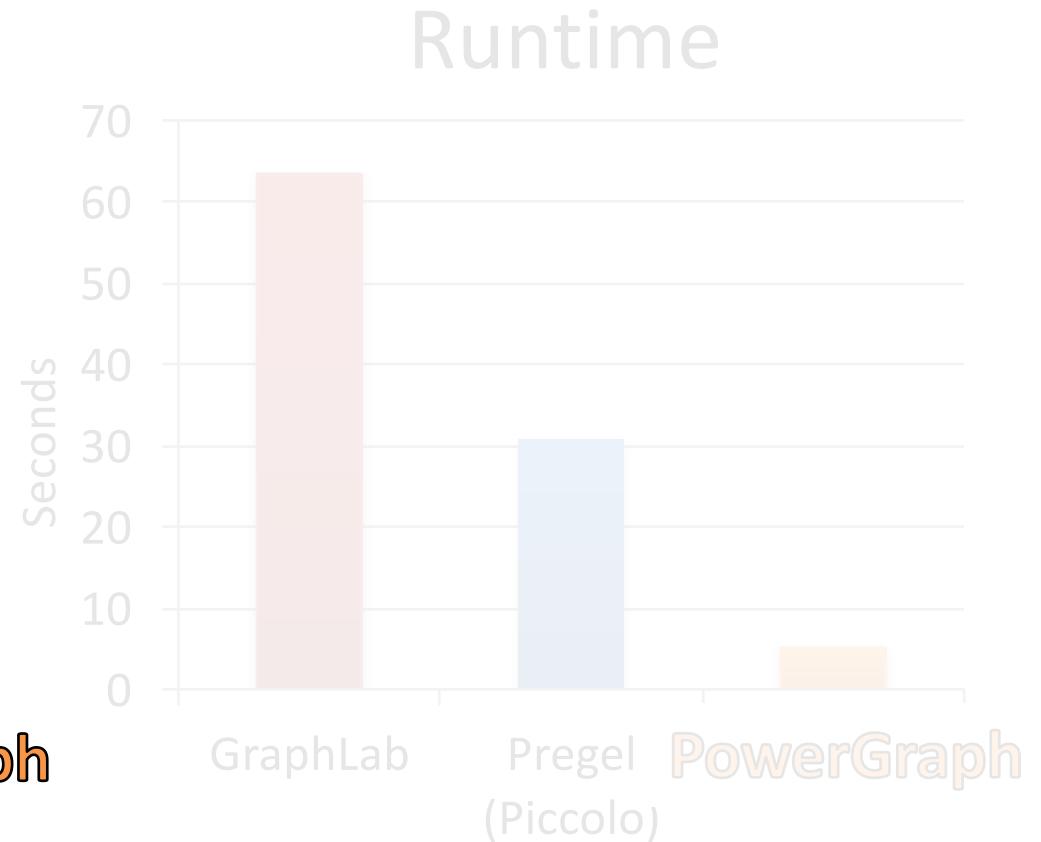
PageRank on the Twitter Follower Graph

Natural Graph with 40M Users, 1.4 Billion Links

Communication



Runtime



Reduces Communication

Runs Faster

32 Nodes x 8 Cores (EC2 HPC cc1.4x)

PowerGraph is Scalable

Yahoo Altavista Web Graph (2002):

One of the largest publicly available web graphs

1.4 Billion Webpages, 6.6 Billion Links

7 Seconds per Iter.

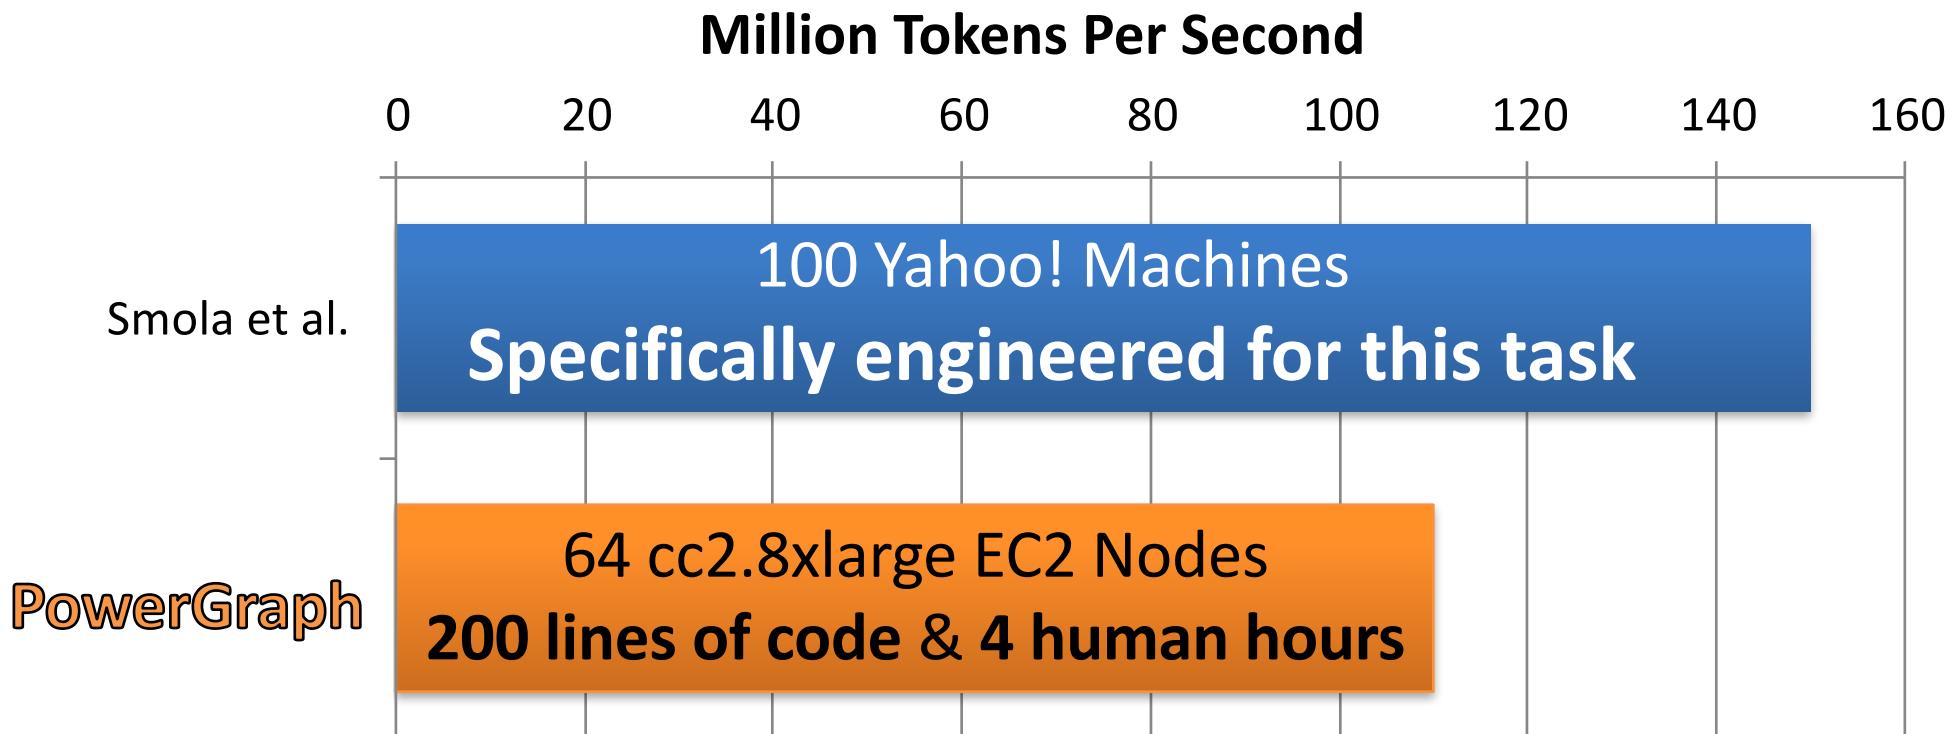
1B links processed per second

30 lines of user code

Topic Modeling



- English language Wikipedia
 - 2.6M Documents, 8.3M Words, 500M Tokens
 - Computationally intensive algorithm



Triangle Counting on The Twitter Graph

Identify individuals with **strong communities**.

Counted: 34.8 Billion Triangles

Hadoop
[WWW'11]

1536 Machines
423 Minutes

PowerGraph

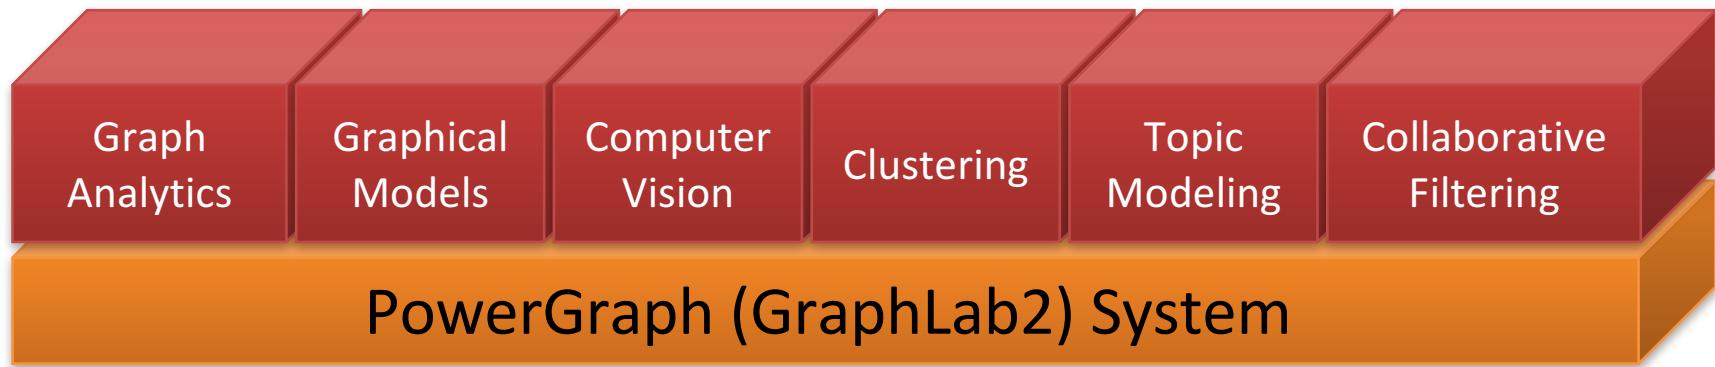
64 Machines
1.5 Minutes

282 x Faster

Why? Wrong Abstraction →

Broadcast $O(\text{degree}^2)$ messages per Vertex

Machine Learning and Data-Mining Toolkits

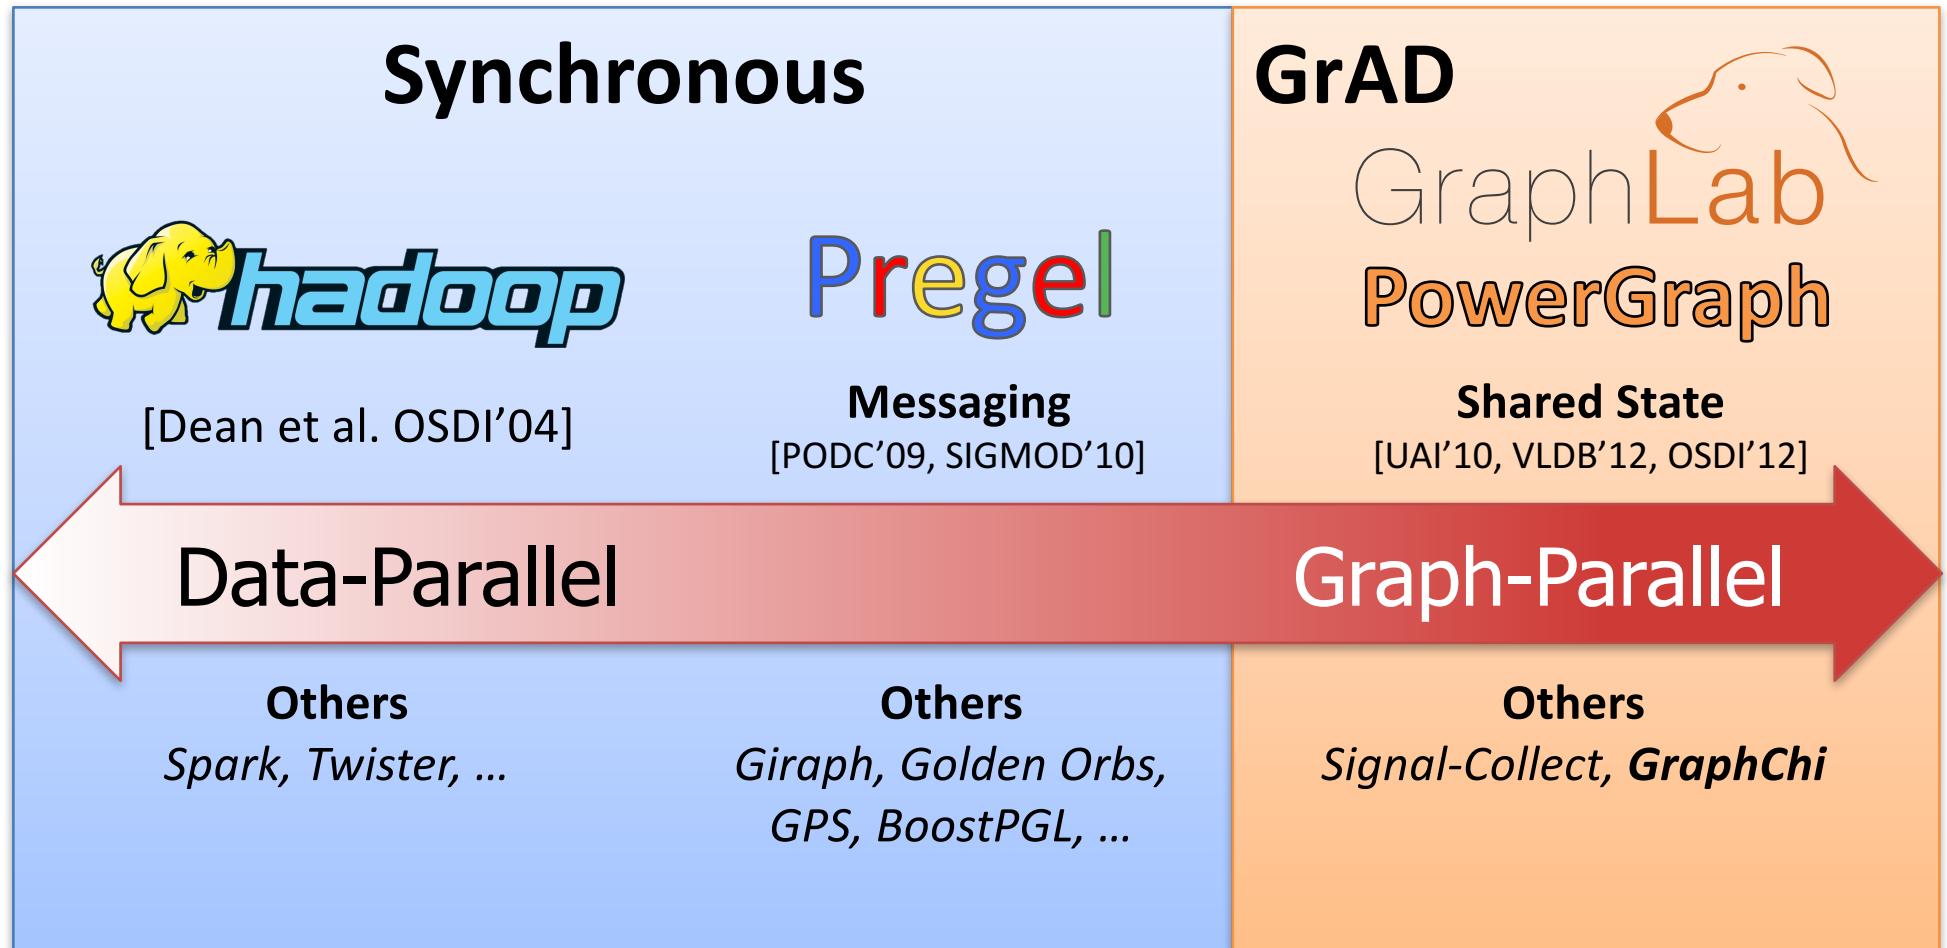


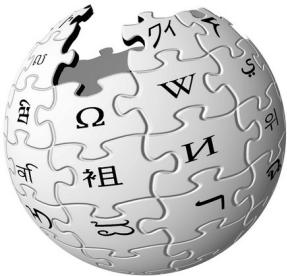
Demonstrates the Applicability
of the GrAD Methodology

Summary: PowerGraph

- Identify the **challenges** of Natural Graphs
 - High-degree vertices, Low-quality edge-cuts
- Solution **PowerGraph** System
 - **GAS Decomposition**: split vertex programs
 - **Vertex-partitioning**: distribute natural graphs
- PowerGraph **theoretically** and **experimentally** outperforms existing graph-parallel systems.

Related High-Level Abstractions





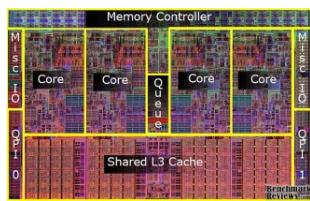
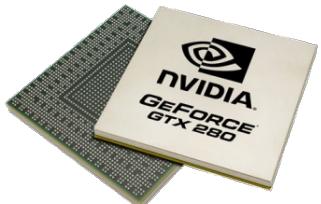
Massive Structured Problems

Probabilistic Graphical Models

**Parallel and Distributed Algorithms
for Probabilistic Inference**

GraphLab & PowerGraph

Advances Parallel Hardware



Thesis Statement

*Efficient parallel and distributed systems for probabilistic reasoning follow the **GrAD Methodology***

1. Graphically decomposition:

- Expose parallelism and distribute state

2. Asynchronous scheduling

- Improved convergence and correctness

3. Dynamic prioritization

- Eliminated wasted work

Observations

- *Graphical models* encode **statistical**, **computational**, and **parallel** structure
- **Tradeoff: Convergence and Parallelism**
 - Many things can be computed in parallel
 - Not all parallel computation is productive
- **Approximation → Increased Parallelism**
 - τ_ε -approximation, approximate sampling
- Power of high-level abstractions
 - Enables the exploration of GrAD methodology

Future: Declarative Models

- Models are *recursive relationships*
 - BP, Gibbs Sampling, PageRank, ...

My Interests Sum of my friends interests

$$A[x_i] = a \left(\sum_{j \in \mathcal{N}[i]} g(A[x_i], A[x_i, x_j], A[x_j]) \right)$$

“Closeness” number of overlapping posts

$$A[x_i, x_j] = s(A[x_i], A[x_i, x_j], A[x_j])$$

- System determines the optimal schedule

Future: Online Probabilistic Reasoning

- The world is rapidly evolving:
 - Make friends and rate movies in real-time
- How do we define and maintain models?
 - **Declarative specification:** *time invariant*
 - τ_ε -*approximation:* **small** change → **local** effect
- Exploit **Power-Law** structure in change
 - Popular items are rated more frequently
 - Exploit burstiness for better caching

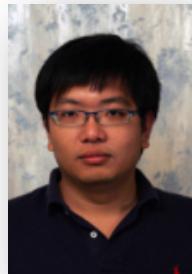
Contributions & Broader Impact

- Theoretically and experimentally characterized
 - Importance of **dynamic asynchronous** scheduling
 - Effect of model **structure** and **parameters** on parallelism
 - Effect of **approximation accuracy** on parallelism
 - Tradeoff between **parallelism** and **convergence**
- Developed two **graph-parallel** abstractions
 - **GraphLab**: vertex-centric view of computation
 - **PowerGraph**: *Distributed* vertex-centric view of computation
- Fostered a community around GraphLab/PowerGraph
 - Substantial industry and academic interest
- Built a foundation for the future design of scalable systems for probabilistic reasoning

Thank You!



Sue Ann
Hong



Yucheng
Low



Aapo
Kyrola



Haijie
Gu



Danny
Bickson



Arthur
Gretto



Andreas
Krause

n



Carlos
Guestrin



Alex
Smola



Jeff
Bilmes



David
O'Hallaron



Guy
Blelloch



Joe
Hellerstein

The Select Lab & My Family