

# From *Graphs* to *Tables*: The Design of Scalable Systems for *Graph Analytics*

Joseph E. Gonzalez

Post-doc, UC Berkeley AMPLab

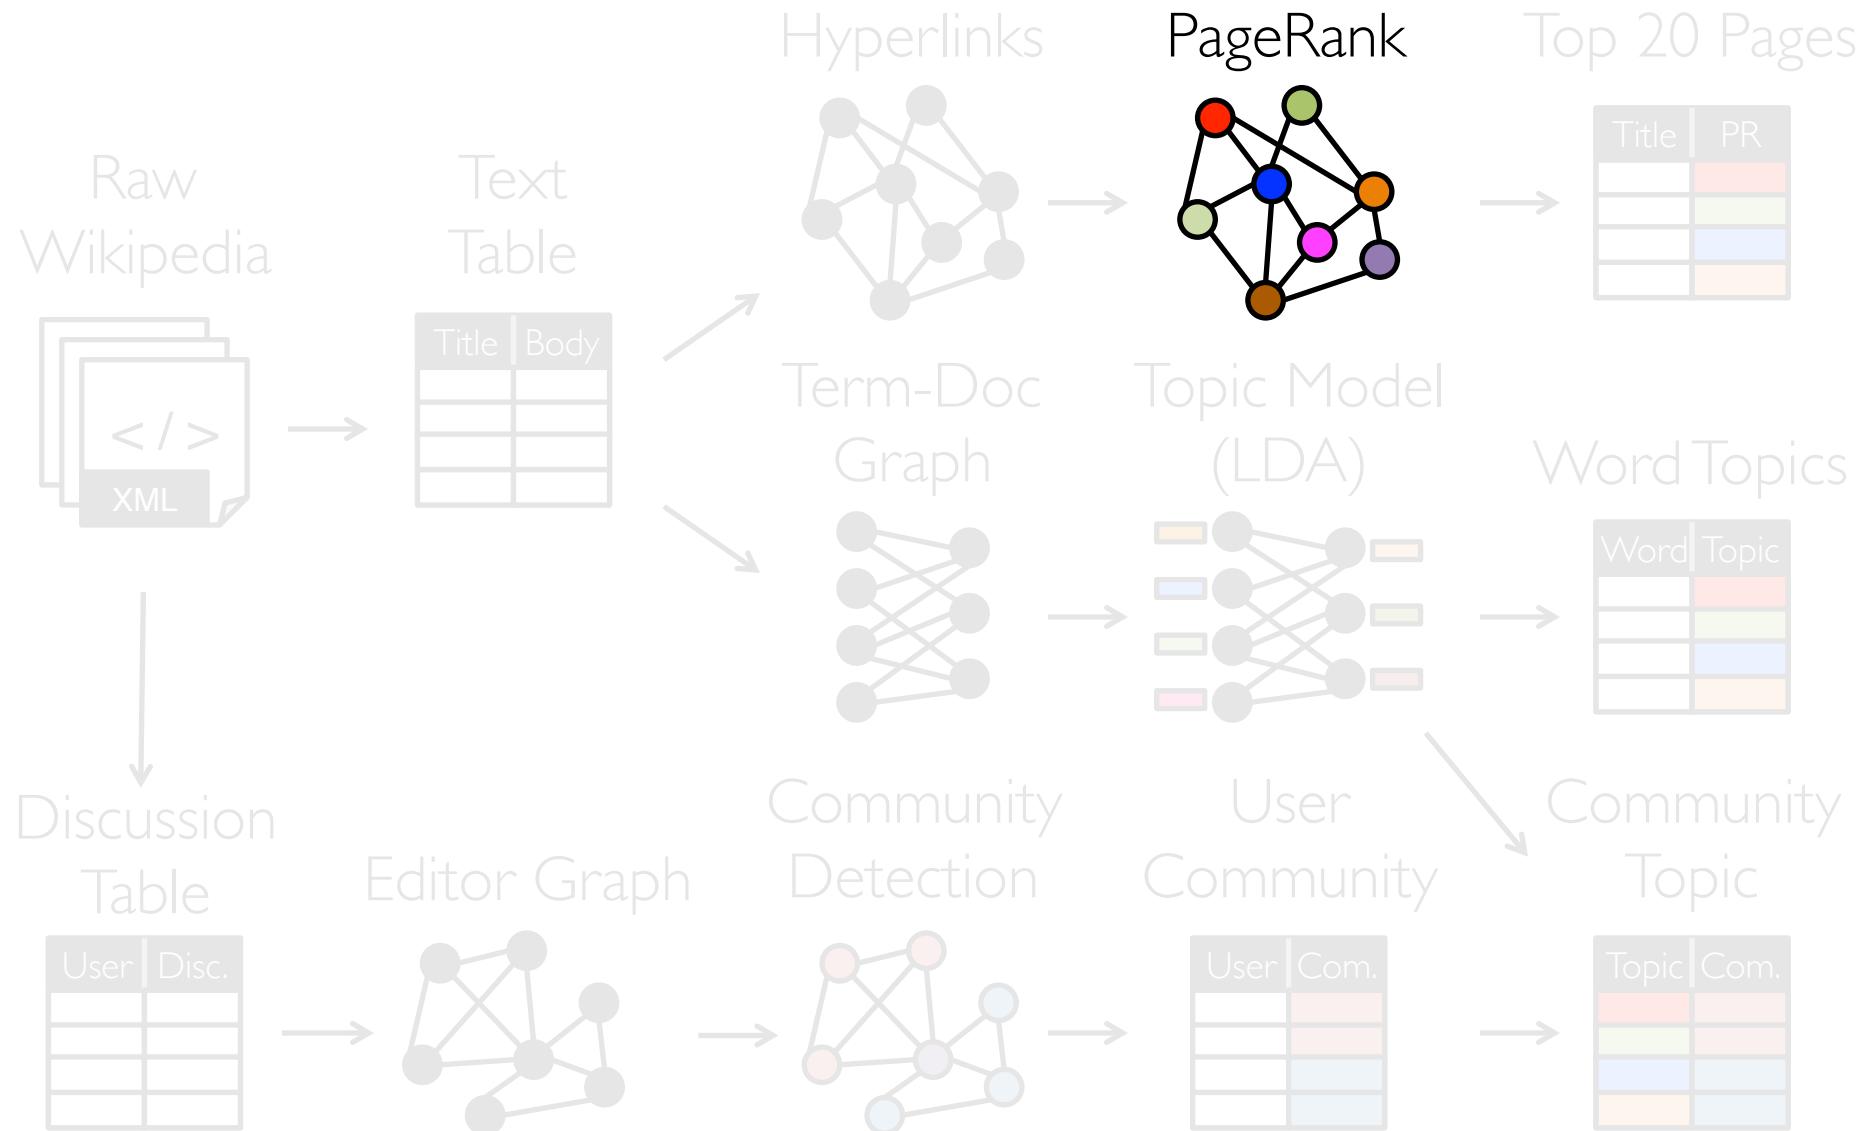
[jegonzal@eecs.berkeley.edu](mailto:jegonzal@eecs.berkeley.edu)

Co-founder, GraphLab Inc.

[joseph@graphlab.com](mailto:joseph@graphlab.com)

WWW'14 Workshop on Big Graph Mining

# Graphs are Central to Analytics



# PageRank: Identifying Leaders

$$R[i] = 0.15 + \sum_{j \in \text{Nbrs}(i)} w_{ji} R[j]$$

Rank of  
user  $i$

Sum of neighbors

Update ranks in parallel

Iterate until convergence

# Recommending Products

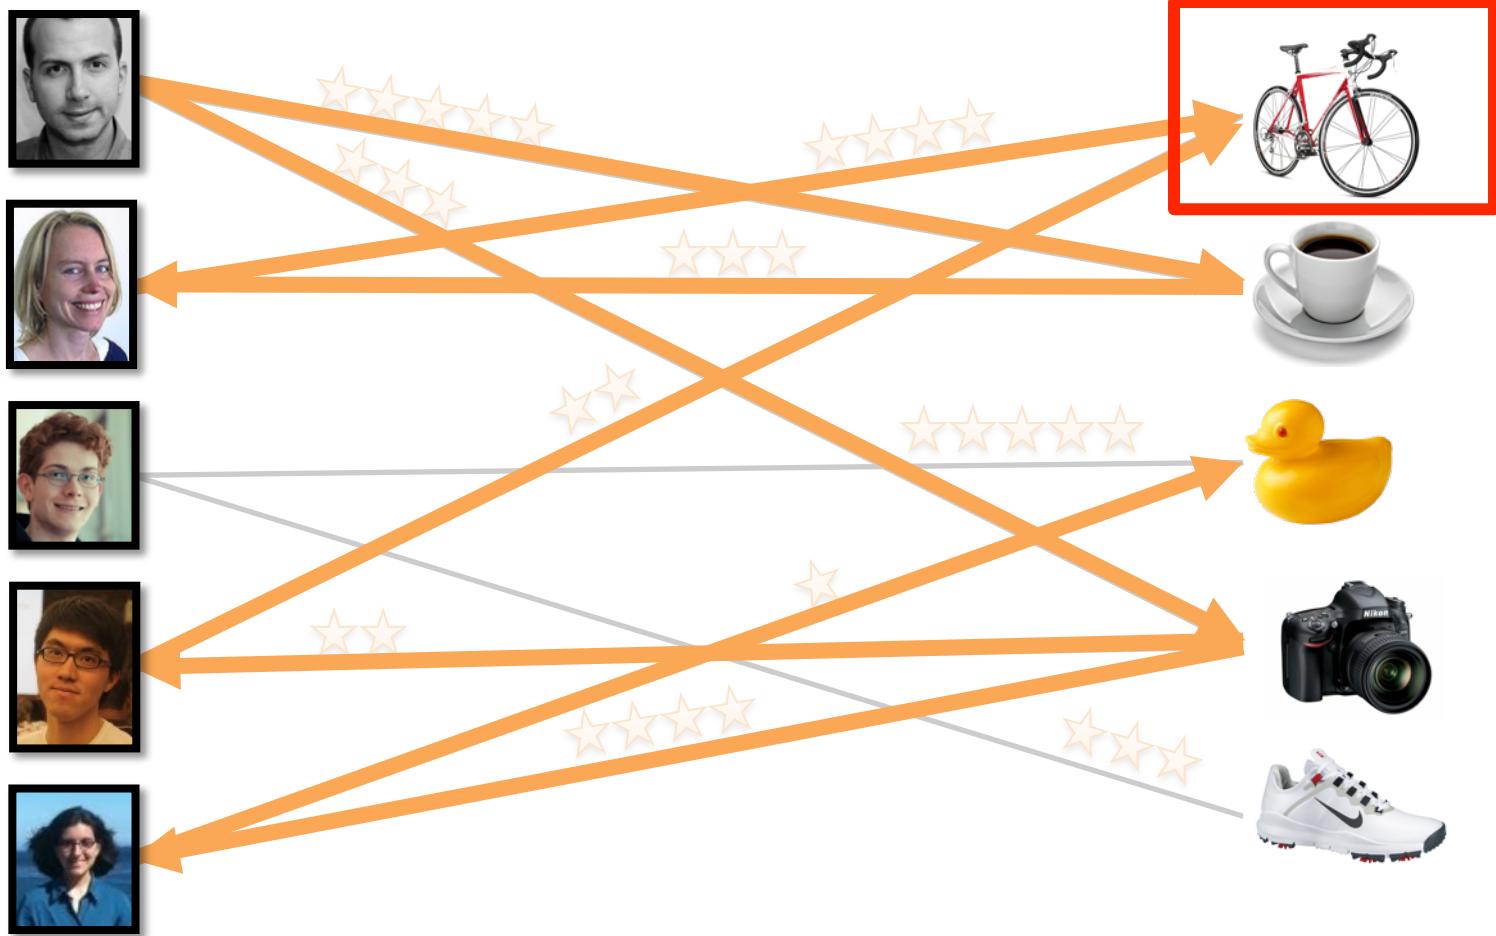
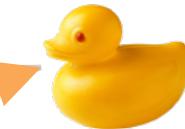
Users



Ratings

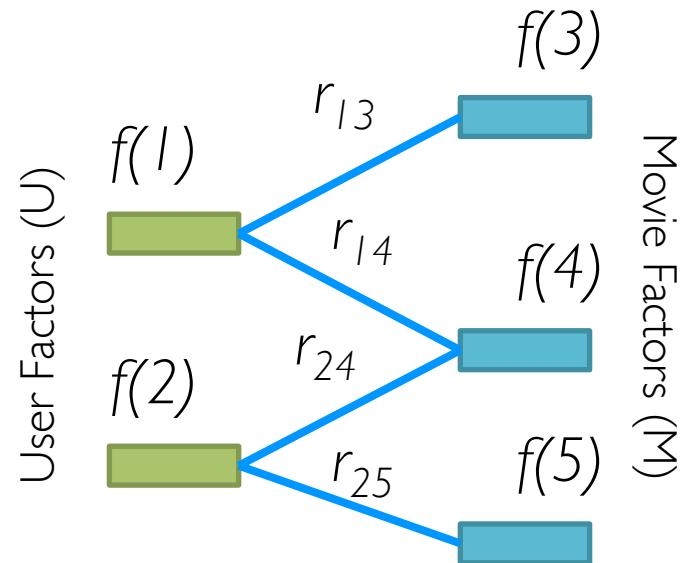
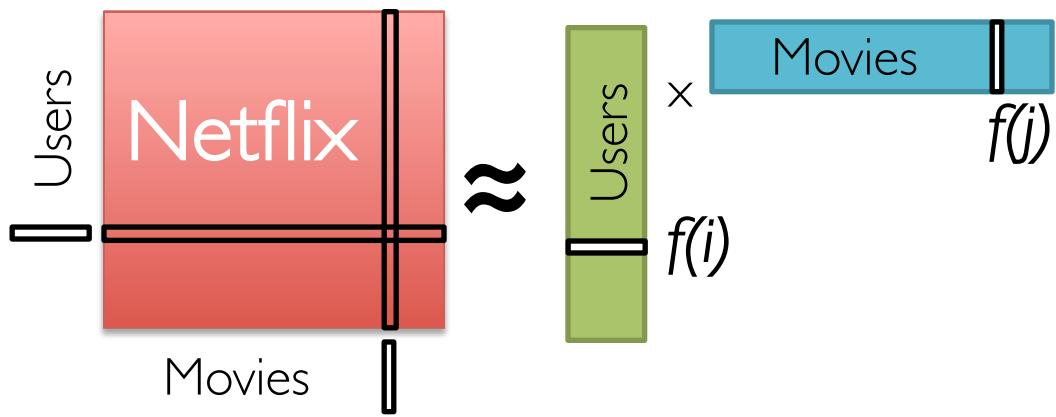


Items



# Recommending Products

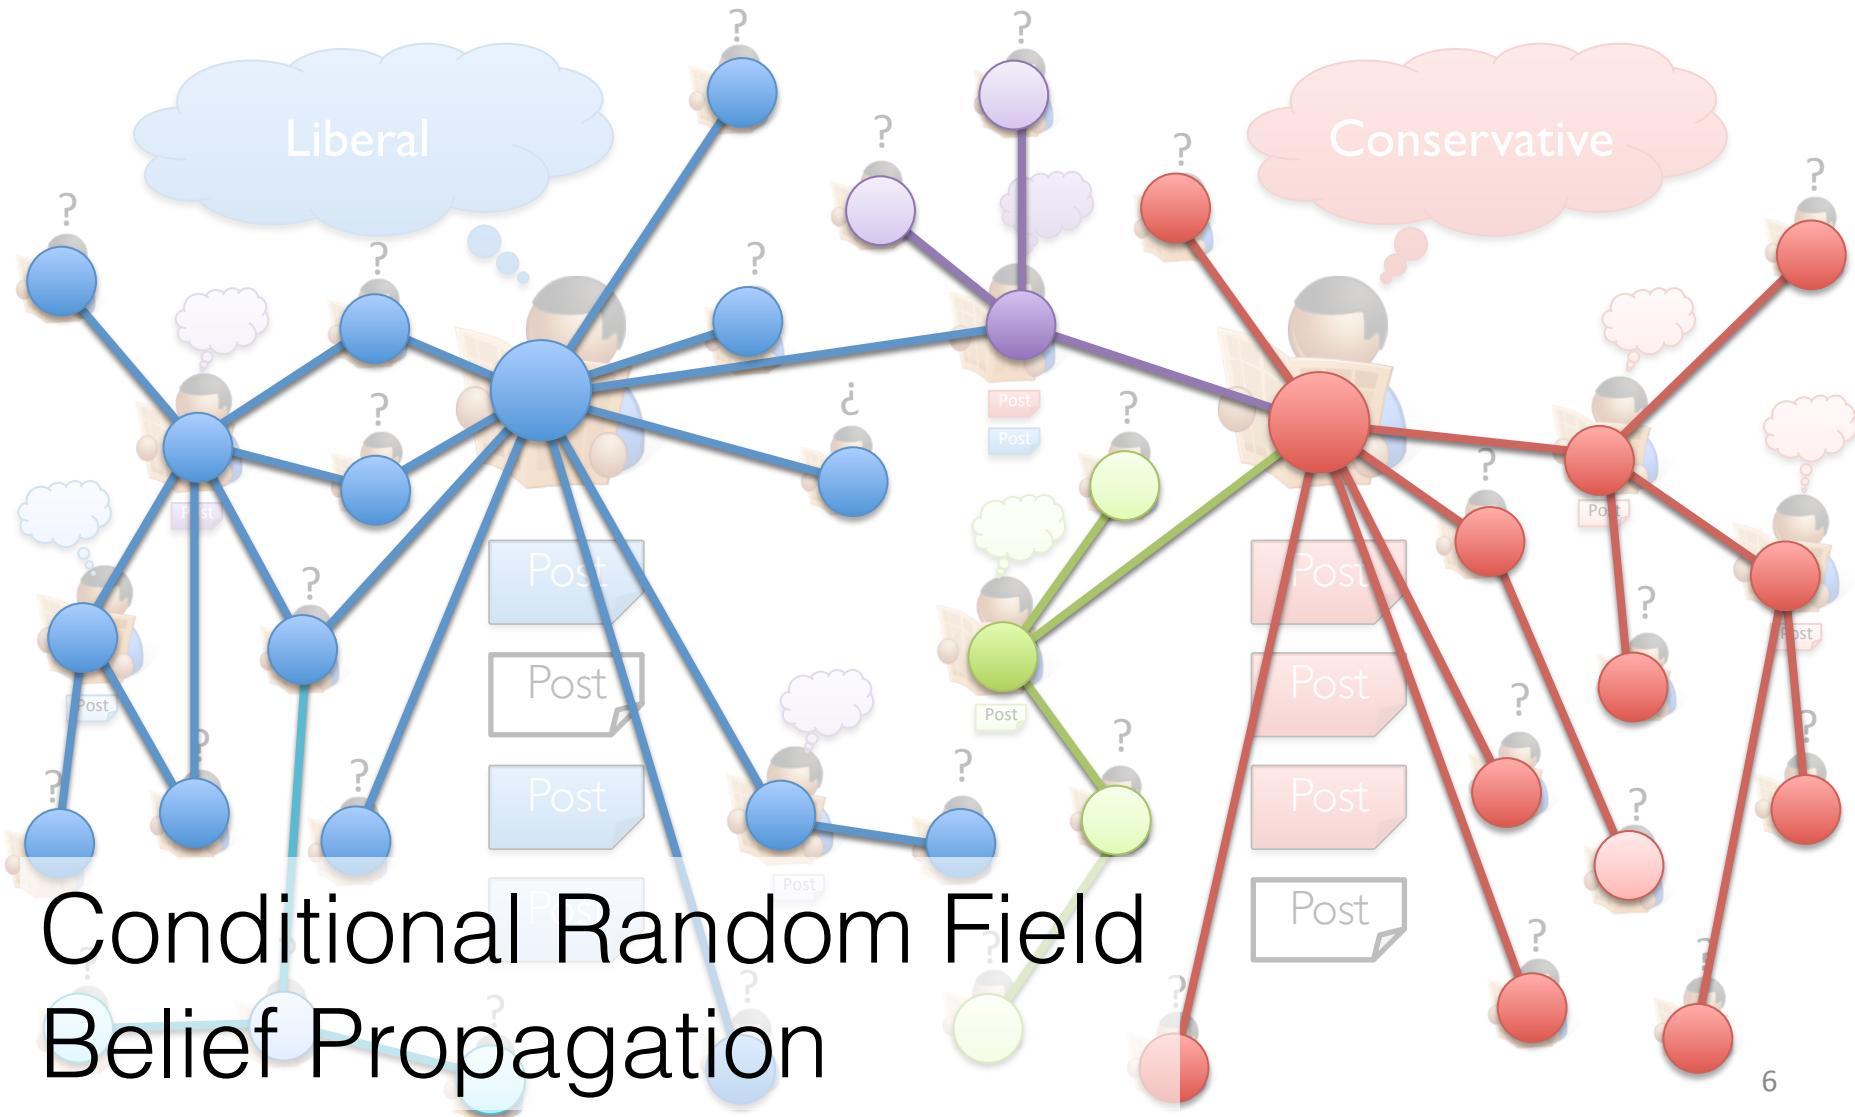
Low-Rank Matrix Factorization:



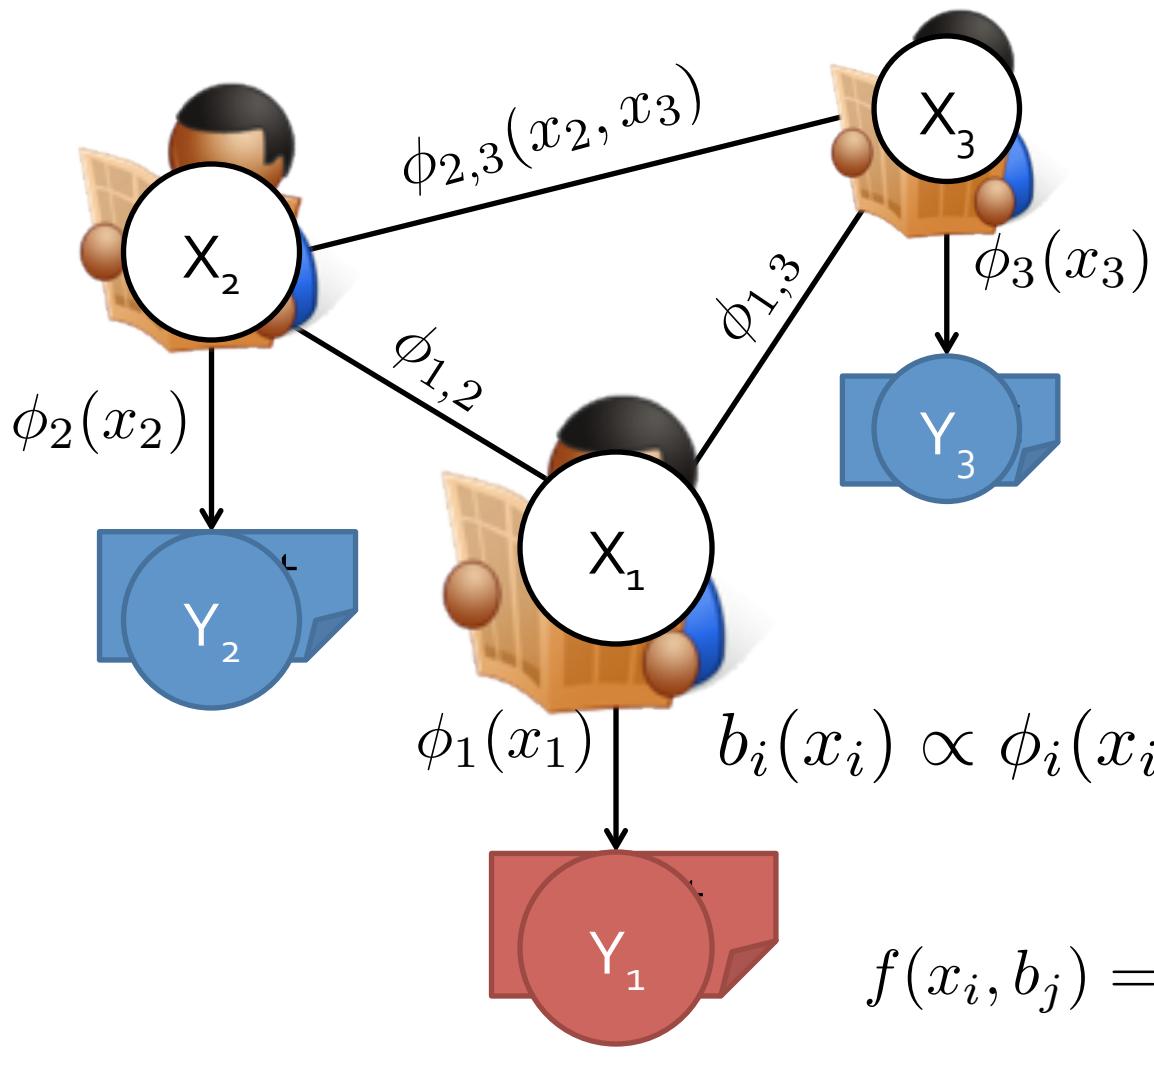
Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda \|w\|_2^2$$

# Predicting User Behavior



# Mean Field Algorithm

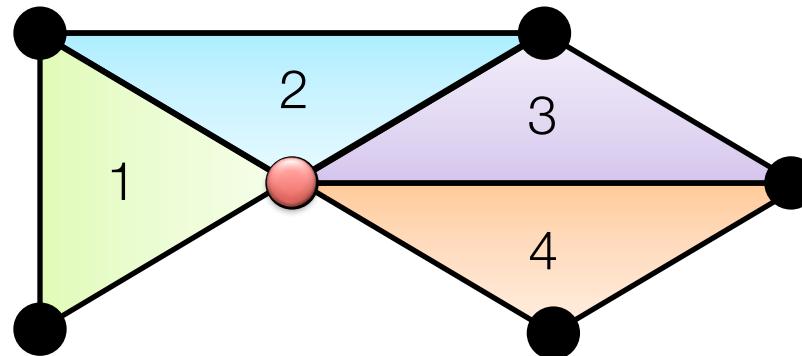


Sum over Neighbors

$$b_i(x_i) \propto \phi_i(x_i) \exp \left( \sum_{j \in N_i} f(x_i, b_j) \right)$$
$$f(x_i, b_j) = \sum_{x_j} b_j(x_j) \log \phi_{i,j}(x_i, x_j)$$

# Finding Communities

Count triangles passing through each vertex:



Measures “cohesiveness” of local community

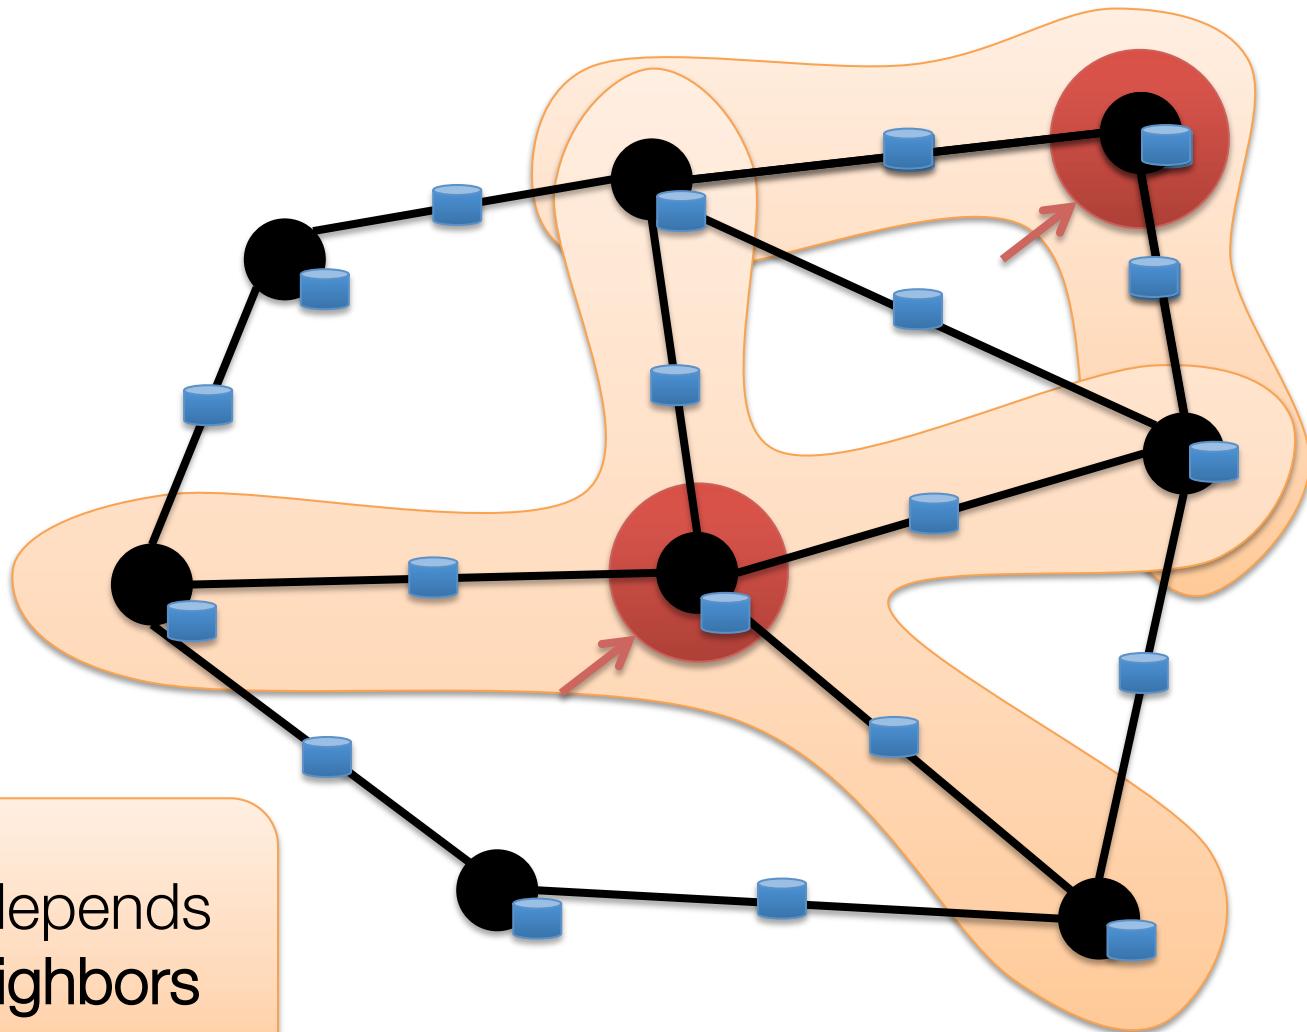
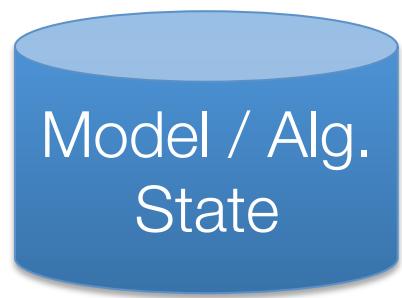


Fewer Triangles  
Weaker Community



More Triangles  
Stronger Community

# The Graph-Parallel Pattern



Computation depends  
only on the **neighbors**

# Many Graph-Parallel Algorithms

- Collaborative Filtering
  - Alternating Least Squares
  - Stochastic Gradient Descent
  - Tensor Factorization
- Structured Prediction
  - Loopy Belief Propagation
  - Max-Product Linear Programs
  - Gibbs Sampling
- Semi-supervised ML
  - Graph SSL
  - CoEM
- Community Detection
  - Triangle-Counting
  - K-core Decomposition
  - K-Truss
- Graph Analytics
  - PageRank
  - Personalized PageRank
  - Shortest Path
  - Graph Coloring
- Classification
  - Neural Networks

# Graph-Parallel Systems

Pregel

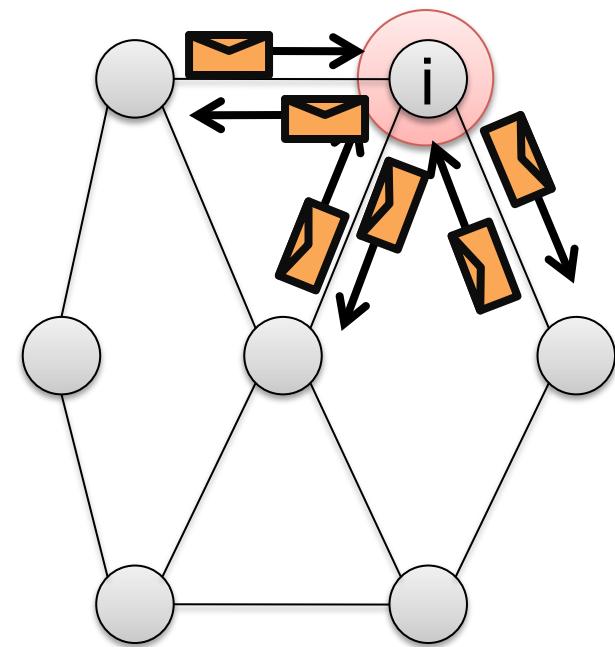


*Expose specialized APIs to simplify  
graph programming.*

# The Pregel (Push) Abstraction

Vertex-Programs interact by sending messages.

```
Pregel_PageRank(i, messages) :  
    // Receive all the messages  
    total = 0  
    foreach( msg in messages ) :  
        total = total + msg  
  
    // Update the rank of this vertex  
    R[i] = 0.15 + total  
  
    // Send new messages to neighbors  
    foreach(j in out_neighbors[i]) :  
        Send msg(R[i]) to vertex j
```



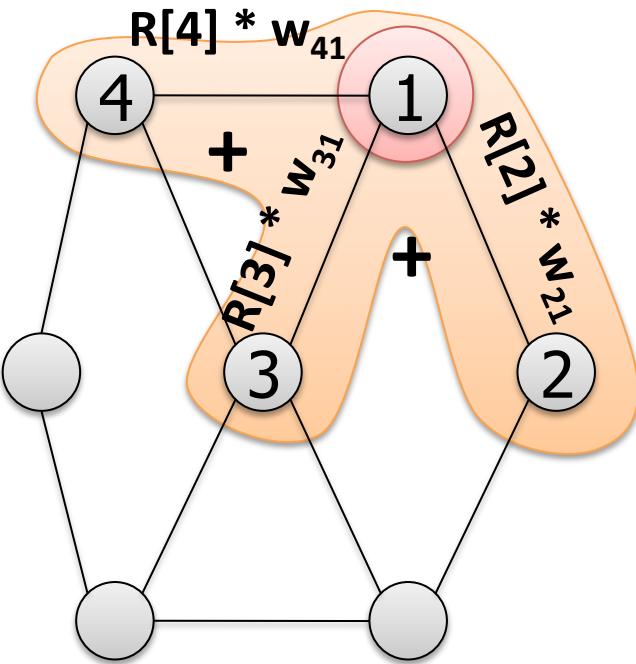
# The GraphLab (Pull) Abstraction

Vertex Programs directly access adjacent vertices and edges

```
GraphLab_PageRank(i)
```

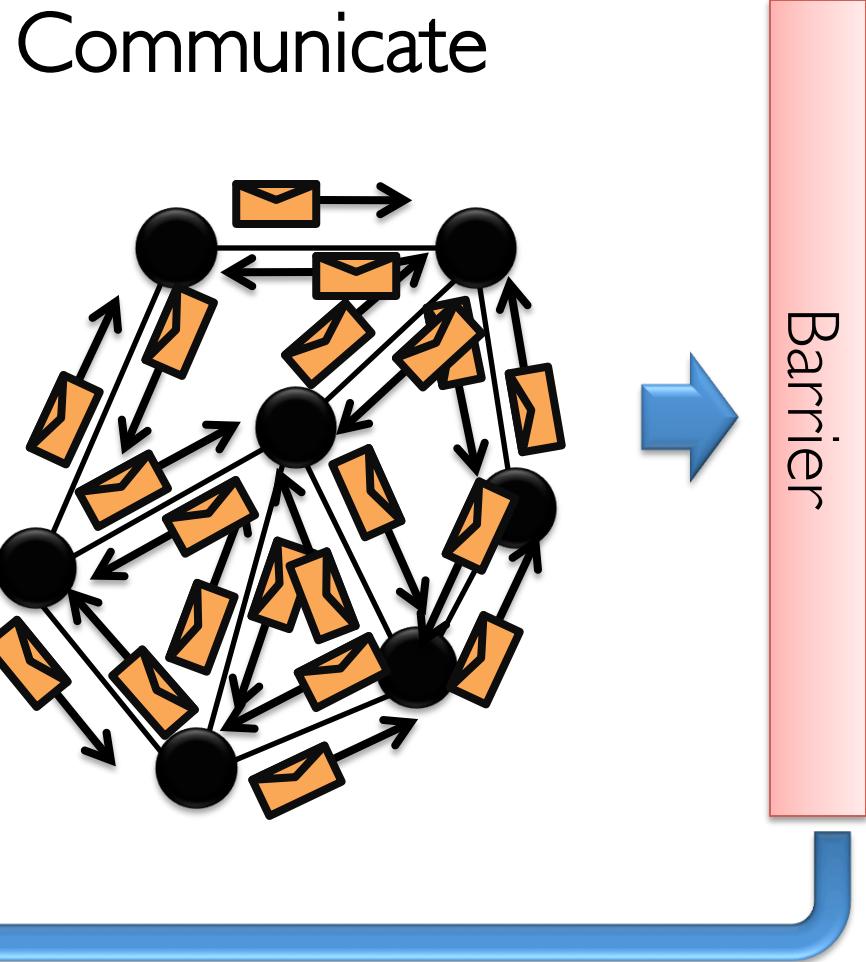
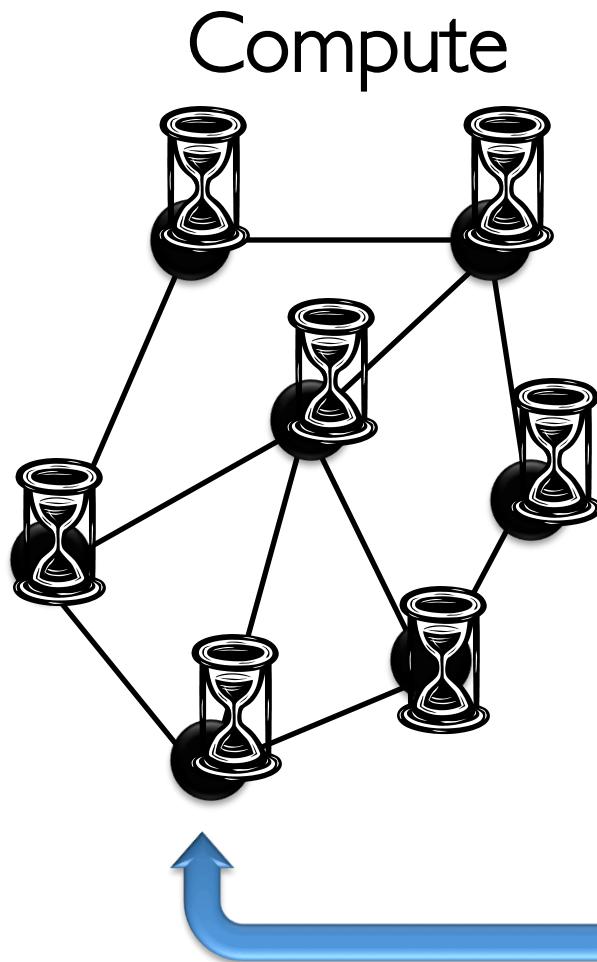
```
// Compute sum over neighbors
total = 0
foreach( j in neighbors(i)):
    total = total + R[j] * wji
```

```
// Update the PageRank
R[i] = 0.15 + total
```



Data movement is managed by the system  
and not the user.

# Iterative Bulk Synchronous Execution



# Graph-Parallel Systems

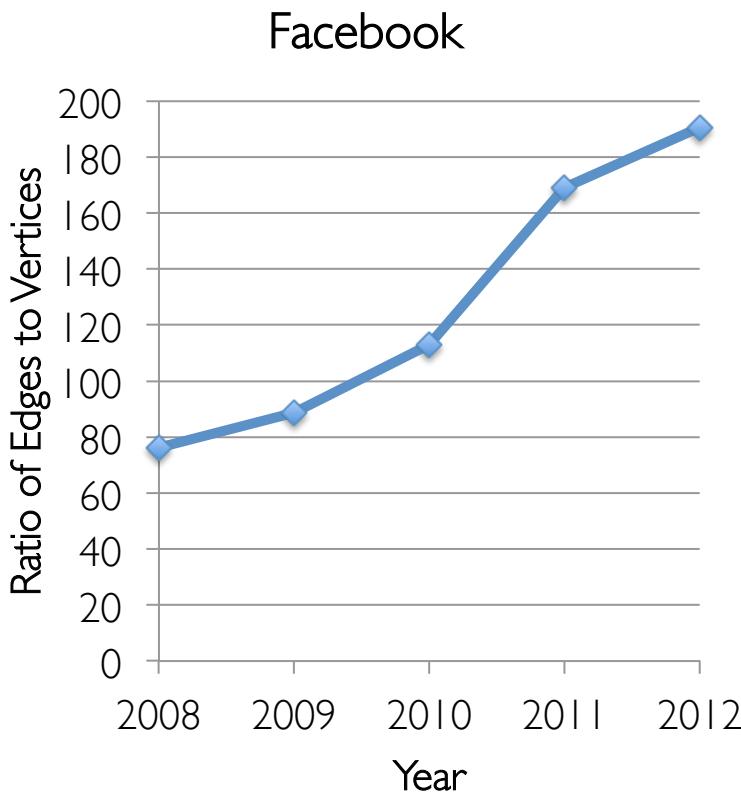
Pregel



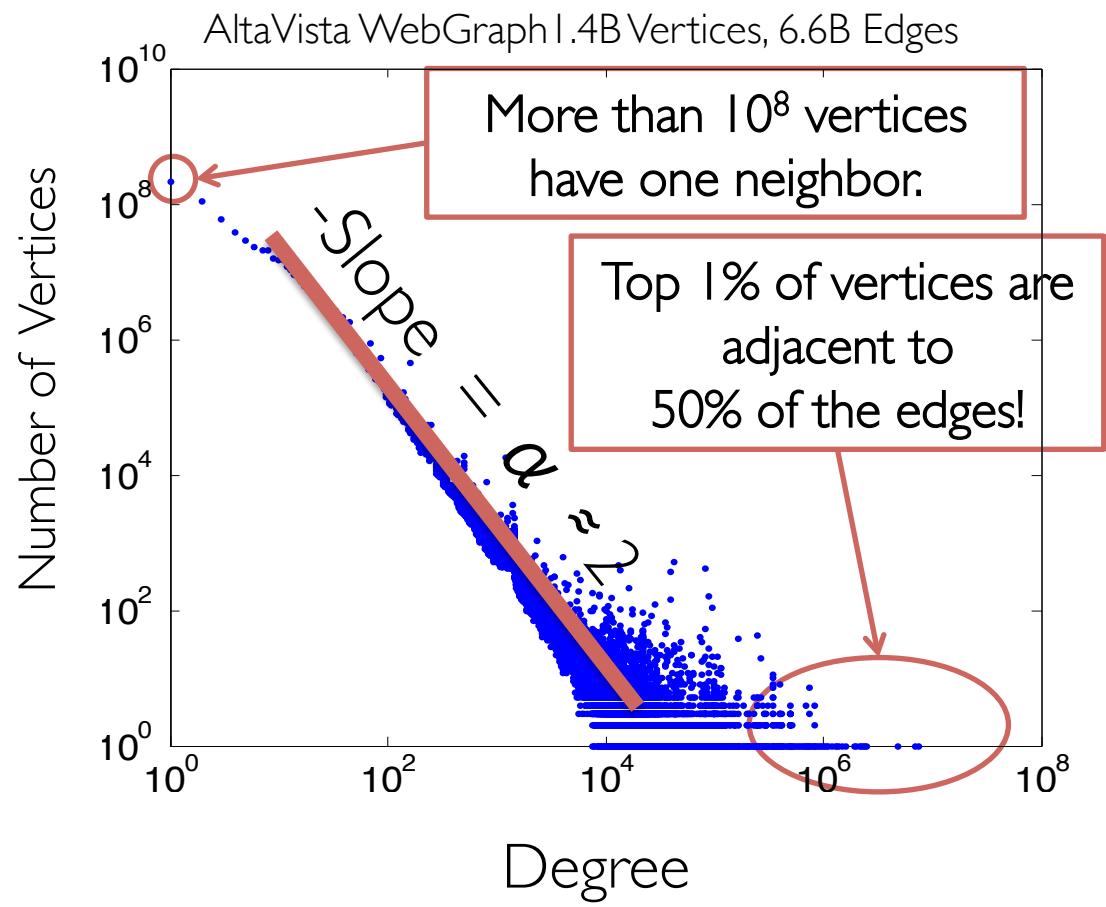
*Exploit graph structure to achieve  
orders-of-magnitude performance gains  
over more general data-parallel systems.*

# Real-World Graphs

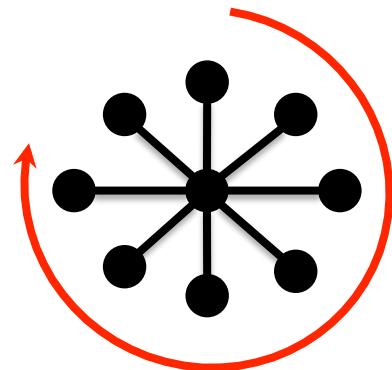
Edges >> Vertices



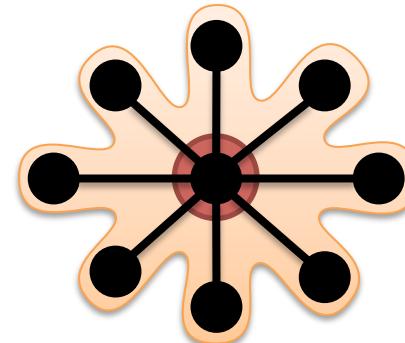
Power-Law Degree Distribution



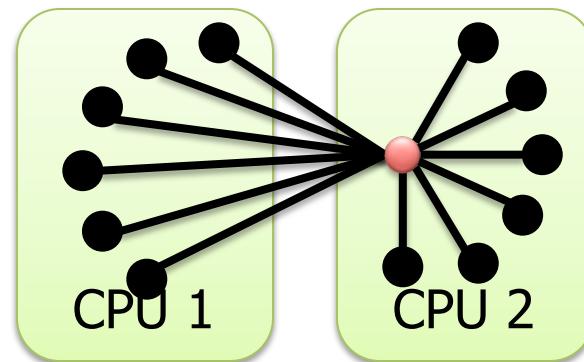
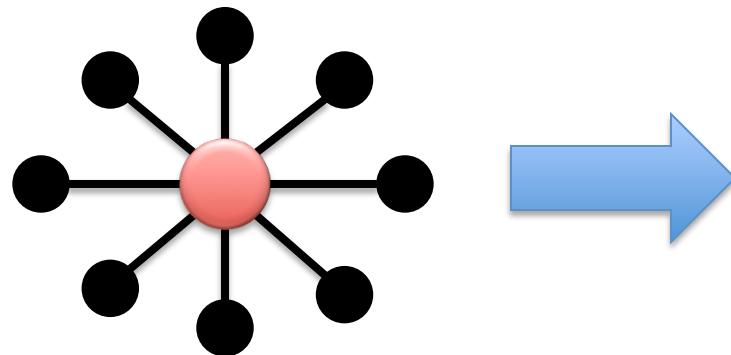
# Challenges of High-Degree Vertices



Sequentially process edges



Touches a large fraction of graph

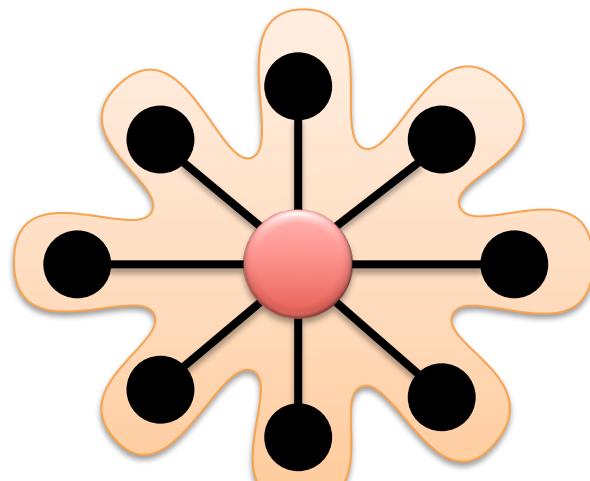


Provably Difficult to Partition

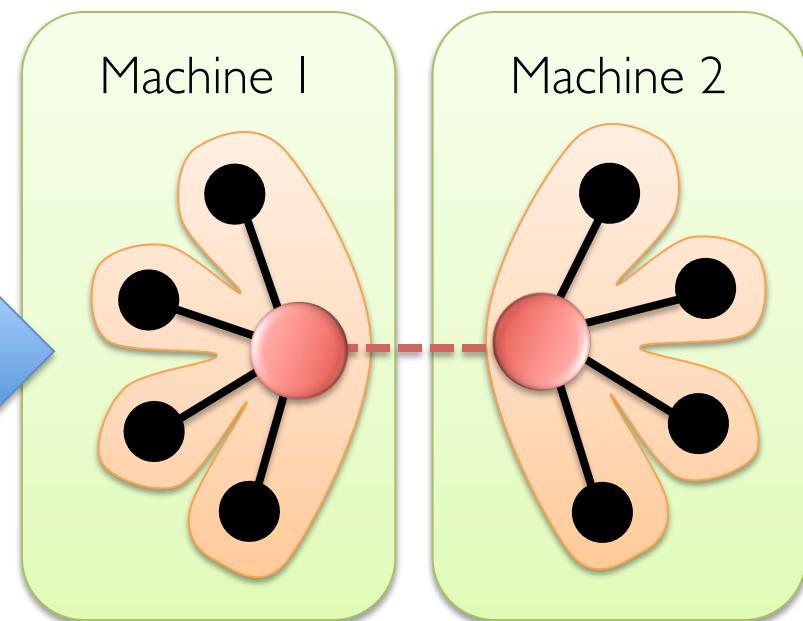
# GraphLab

(PowerGraph, OSDI'12)

Program This



Run on This

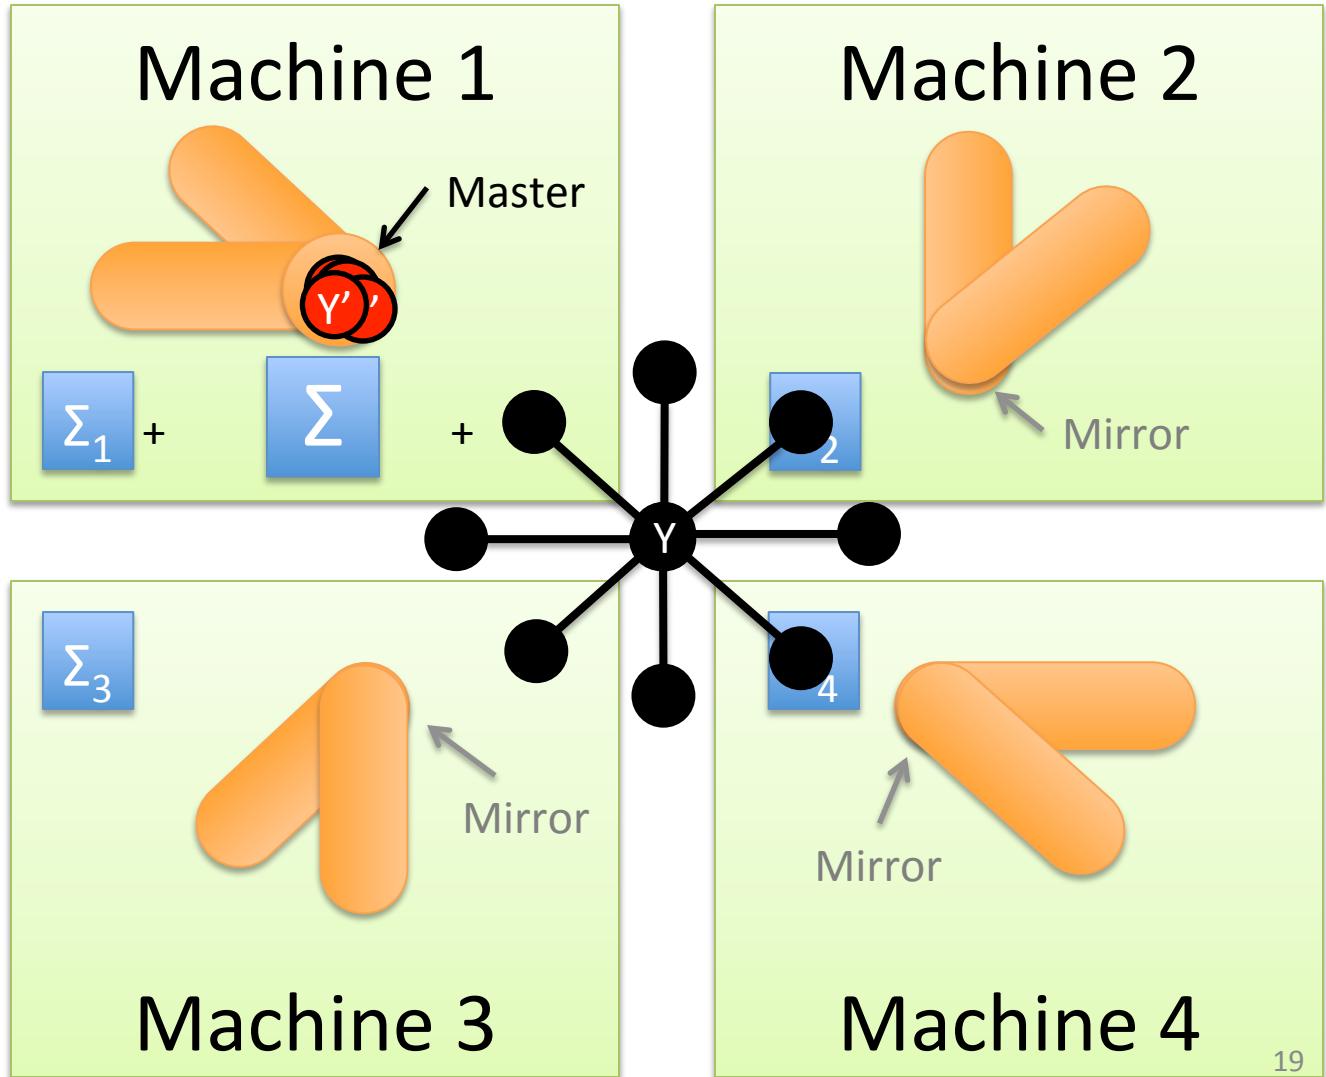


Split High-Degree vertices

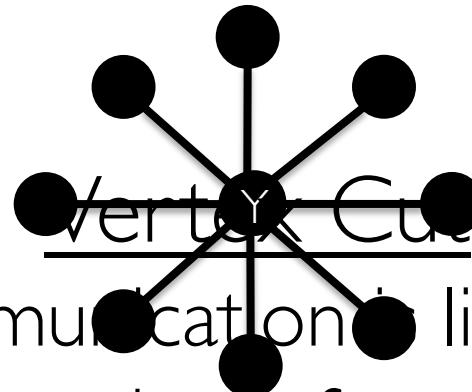
New Abstraction → Equivalence on Split Vertices

# GAS Decomposition

Gather  
Apply  
Scatter



# Minimizing Communication in PowerGraph

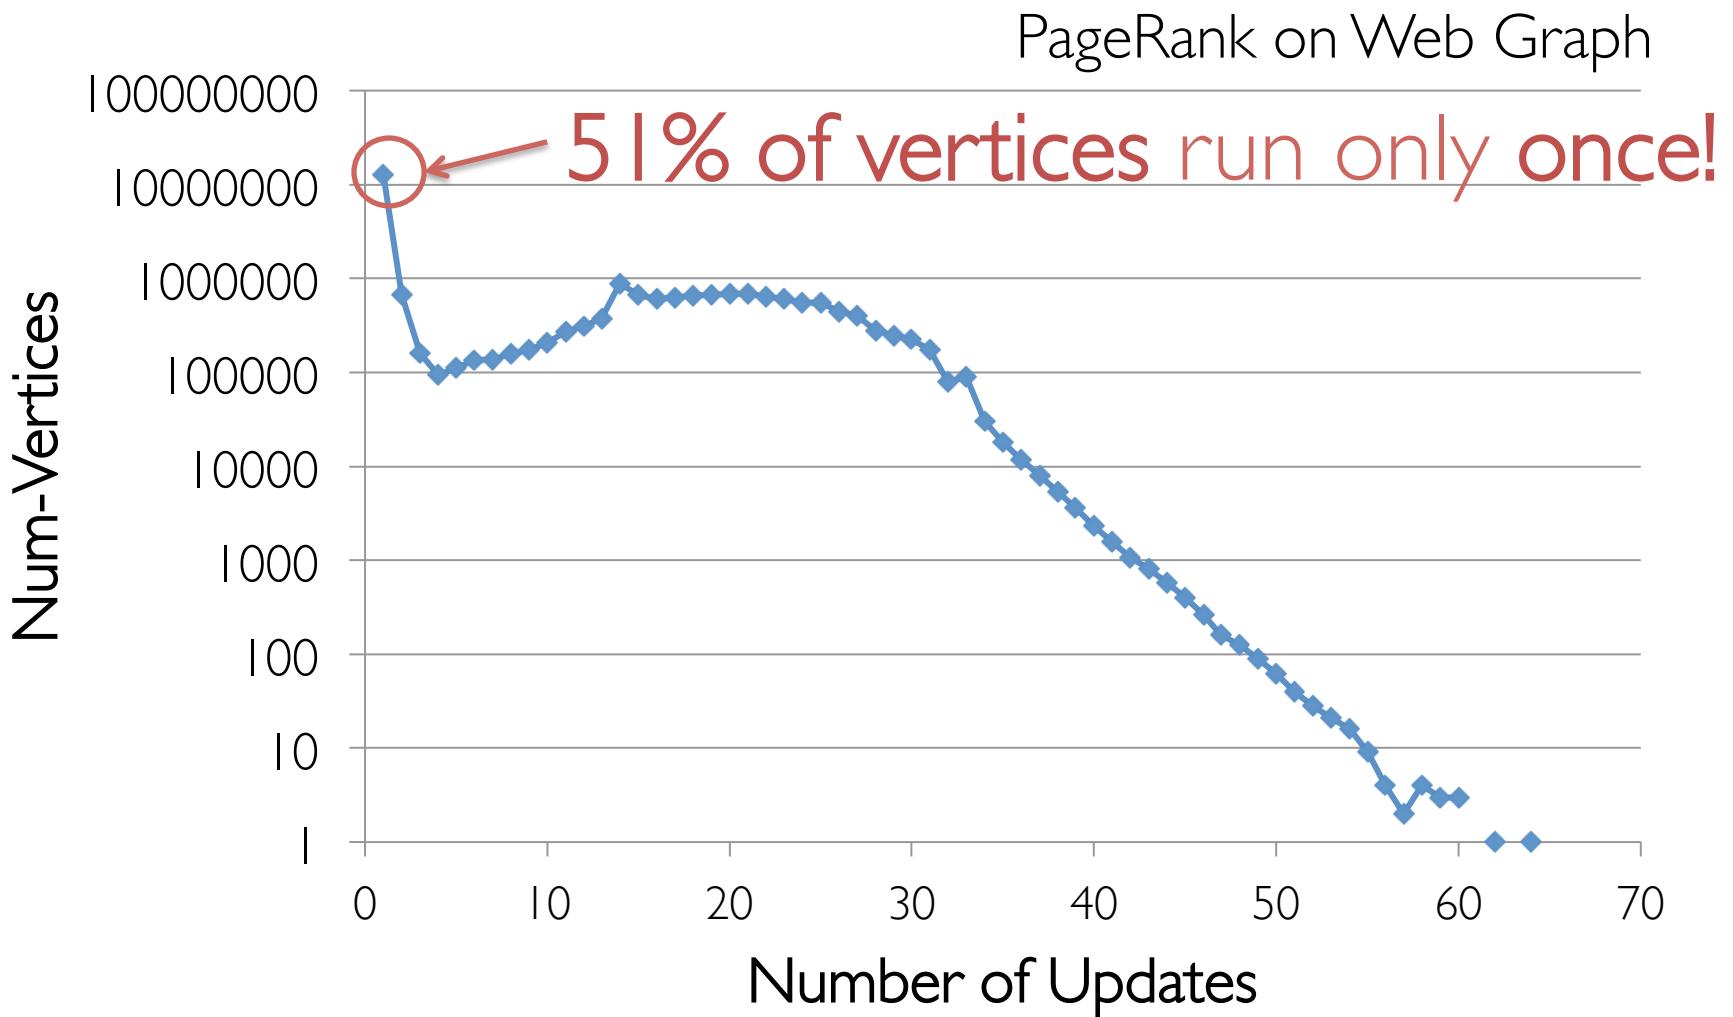


Communication is linear in  
the number of machines  
each vertex spans.

Total communication upper bound:

$$O\left(\#\text{vertices} \sqrt{\#\text{machines}}\right)$$

# Shrinking Working Sets



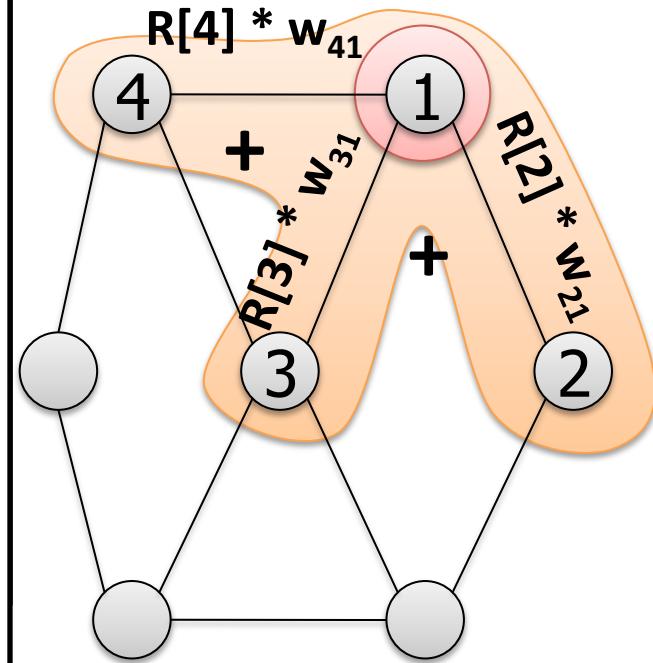
# The GraphLab (Pull) Abstraction

Vertex Programs directly access adjacent vertices and edges

```
GraphLab_PageRank(i)
    // Compute sum over neighbors
    total = 0
    foreach( j in neighbors(i)):
        total = total + R[j] * wji
```

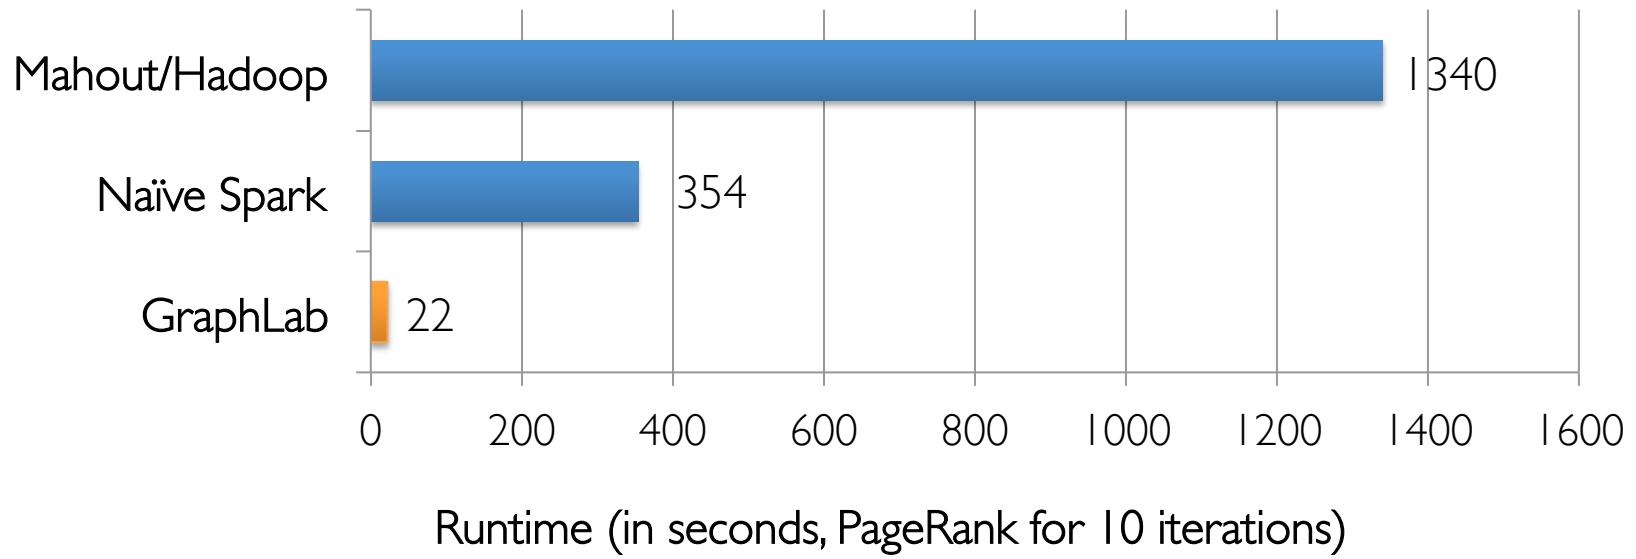
```
// Update the PageRank
R[i] = 0.15 + total
```

```
// Trigger neighbors to run again
if R[i] not converged then
    signal nbrsOf(i) to be recomputed
```



Trigger computation *only* when necessary.

# PageRank on the Live-Journal Graph



GraphLab is *60x faster* than Hadoop  
GraphLab is *16x faster* than Spark

# Triangle Counting on Twitter

40M Users, 1.4 Billion Links

Counted: 34.8 Billion Triangles

Hadoop  
[WWW'11]

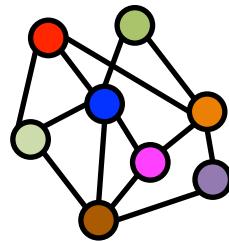
1536 Machines  
423 Minutes

GraphLab

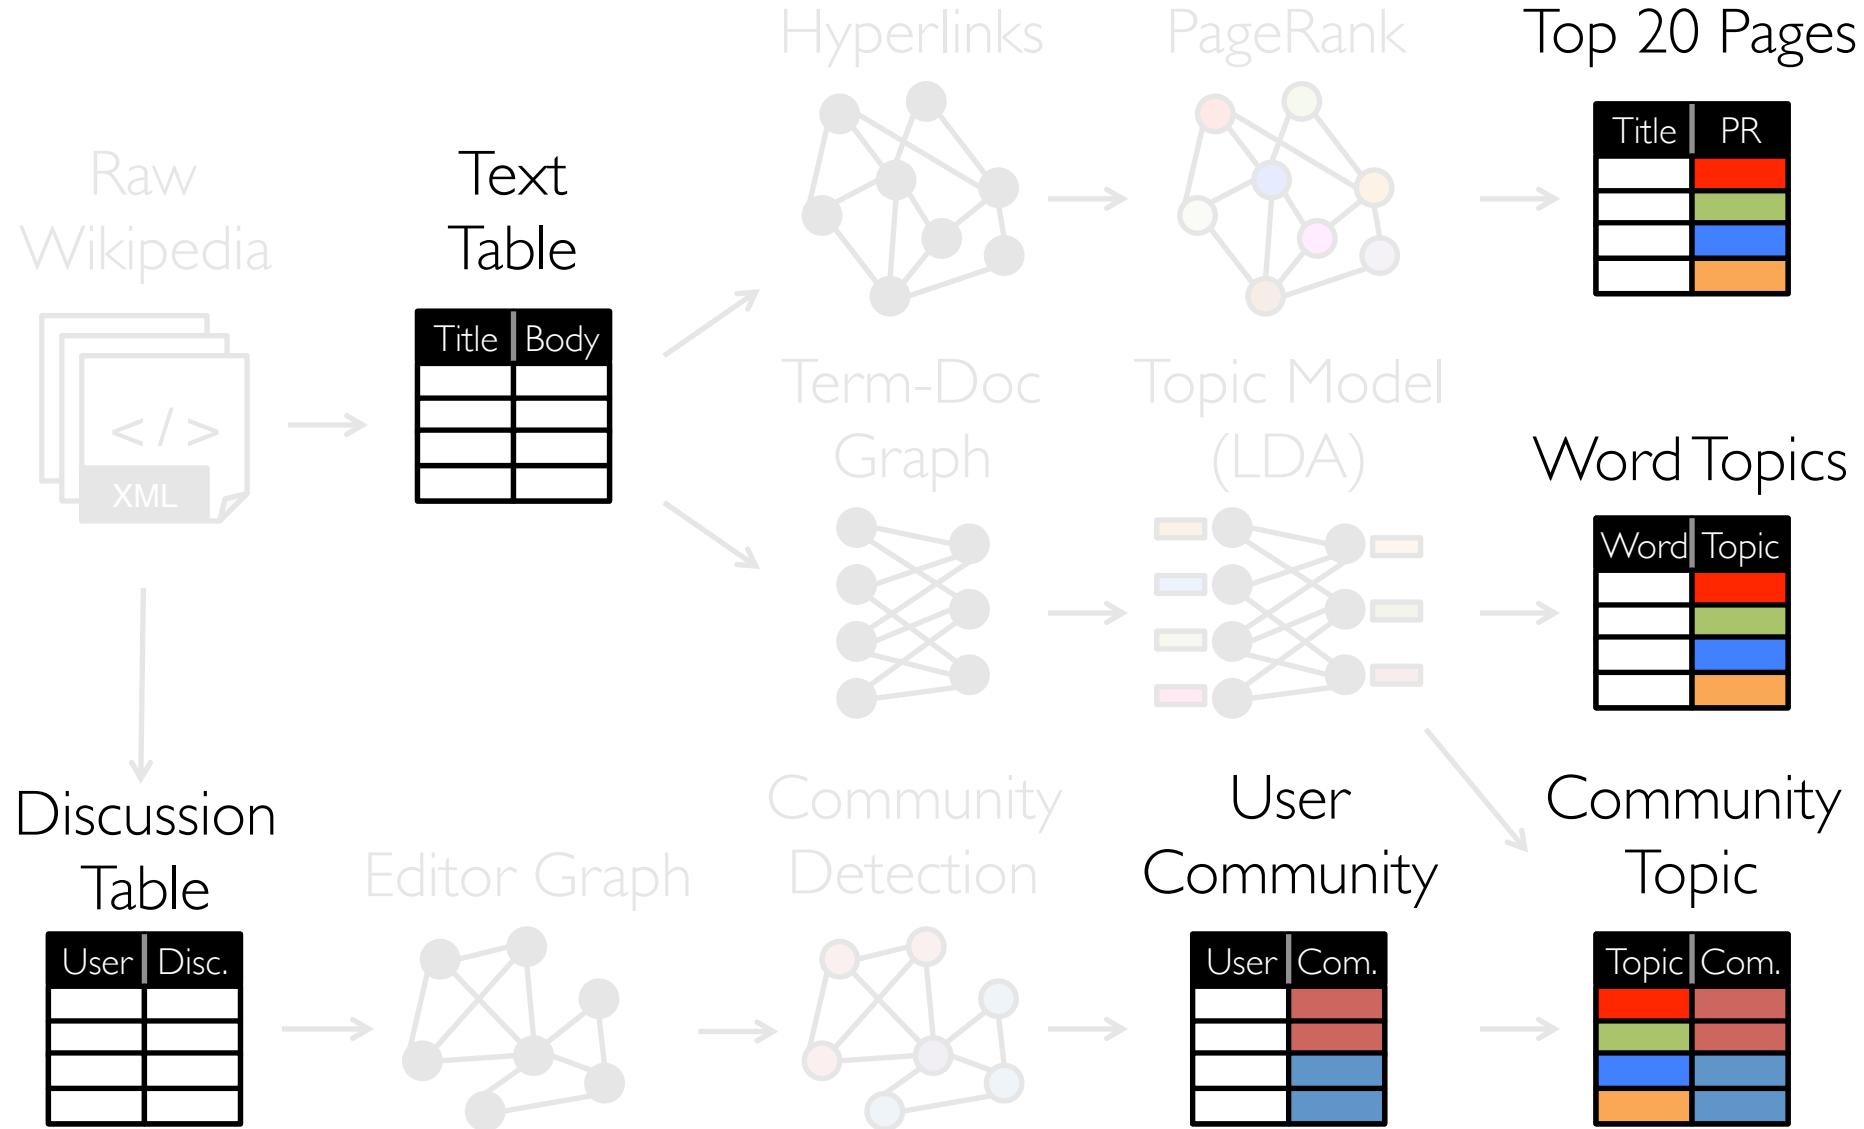
64 Machines  
15 Seconds

1000 × Faster

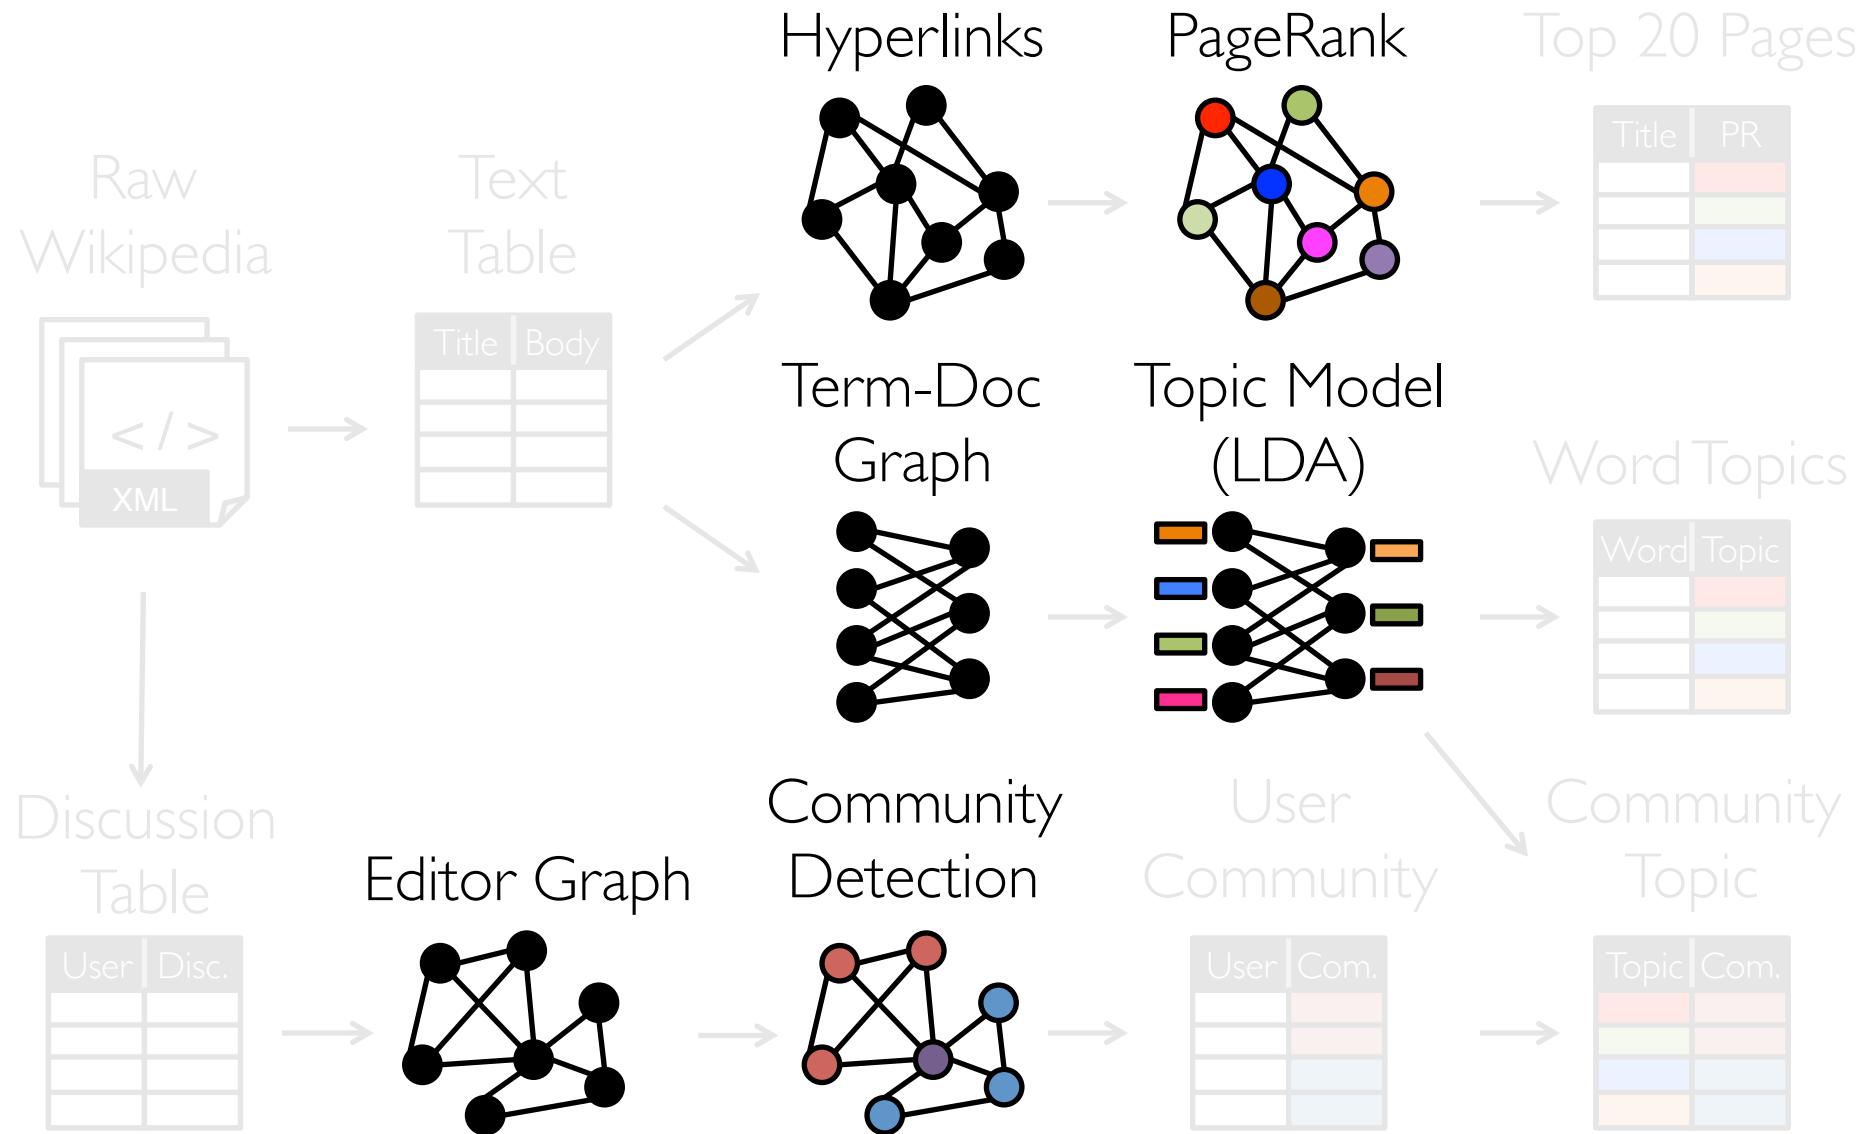
## PageRank



# Tables

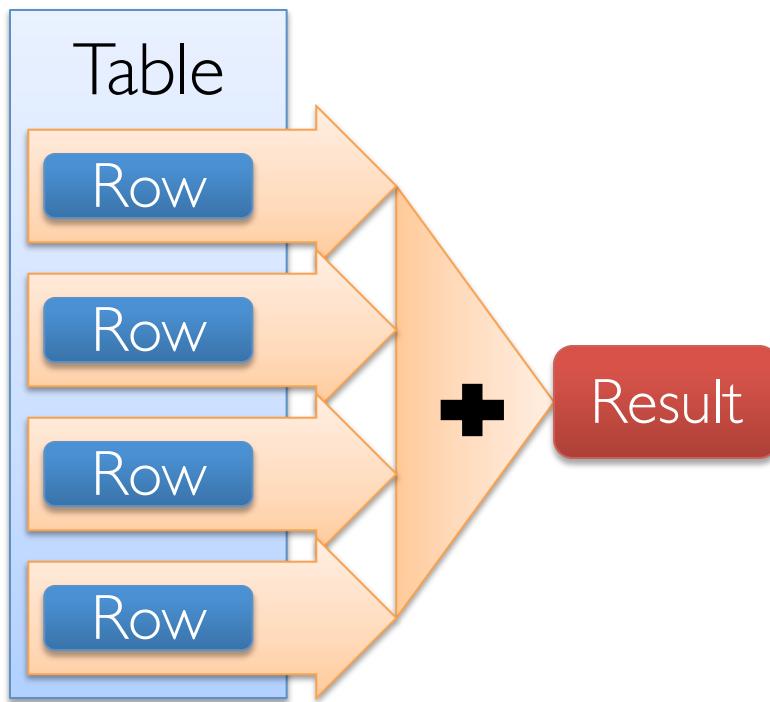


# Graphs



# Separate Systems to Support Each View

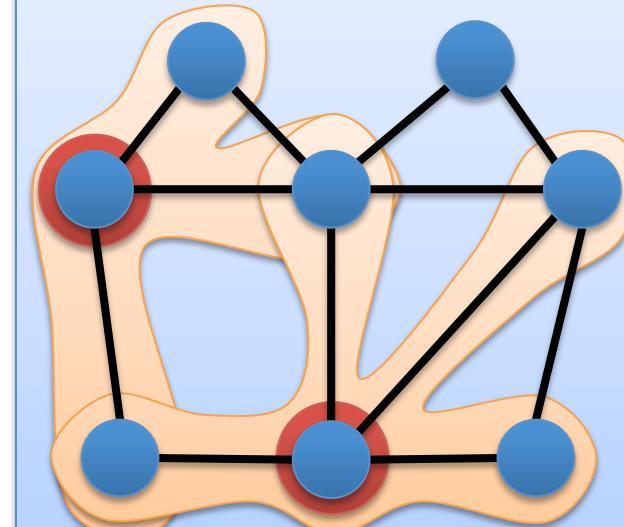
## Table View



## Graph View



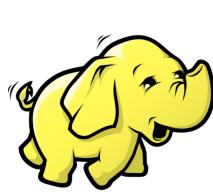
## Dependency Graph



*Separate systems  
for each view can be  
difficult to use and inefficient*

# Difficult to Program and Use

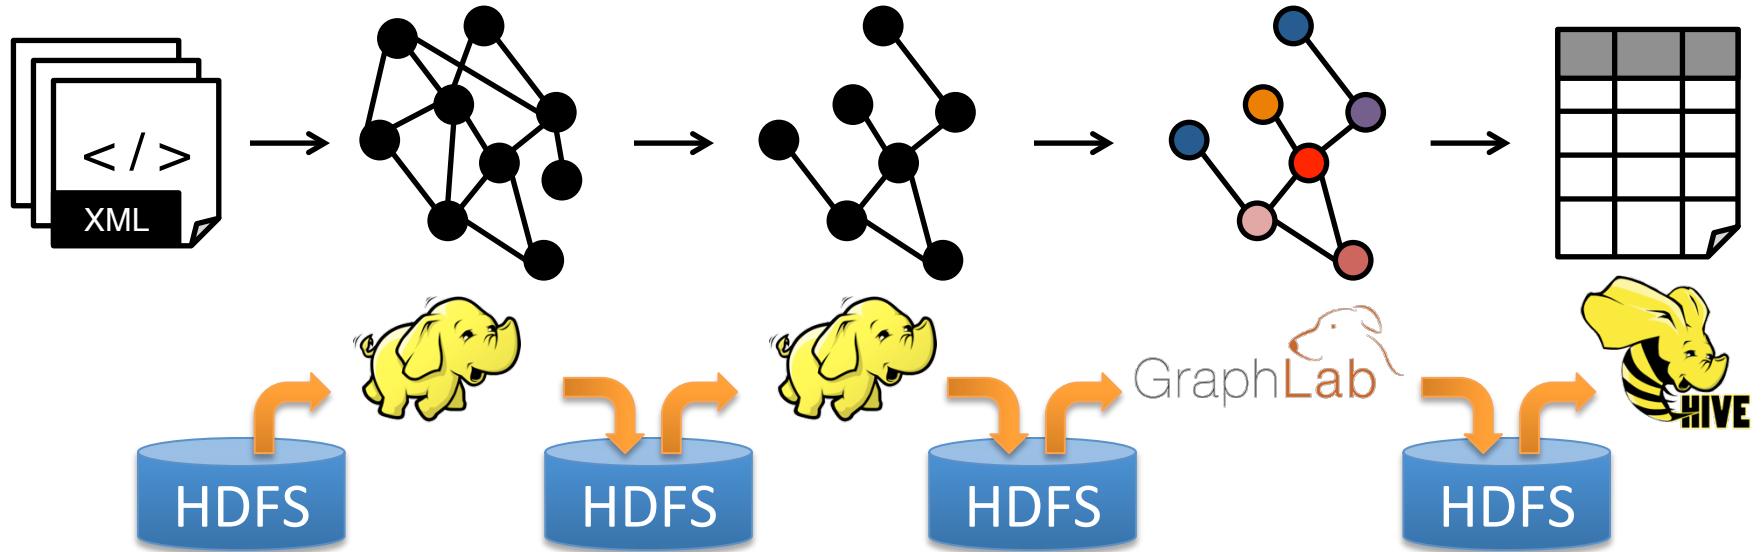
Users must *Learn, Deploy, and Manage* multiple systems



Leads to brittle and often complex interfaces

# Inefficient

Extensive **data movement** and **duplication** across  
the network and file system

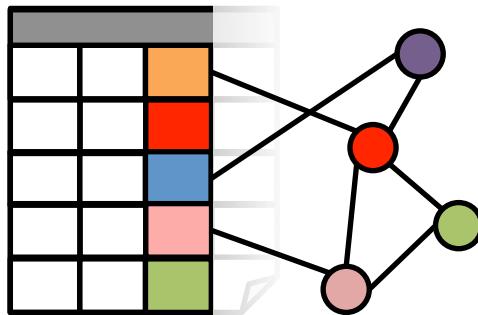


Limited reuse internal data-structures  
across stages

# Solution: The GraphX Unified Approach

## New API

*Blurs the distinction between  
Tables and Graphs*



## New System

*Combines Data-Parallel  
Graph-Parallel Systems*



Enabling users to **easily** and **efficiently**  
express the entire graph analytics pipeline

Tables and Graphs are **composable views** of the same *physical* data

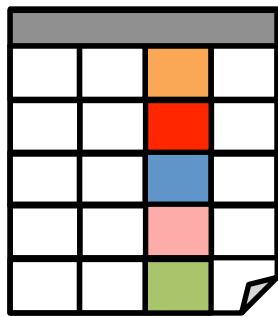
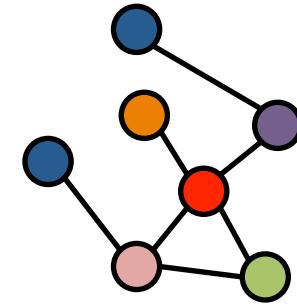
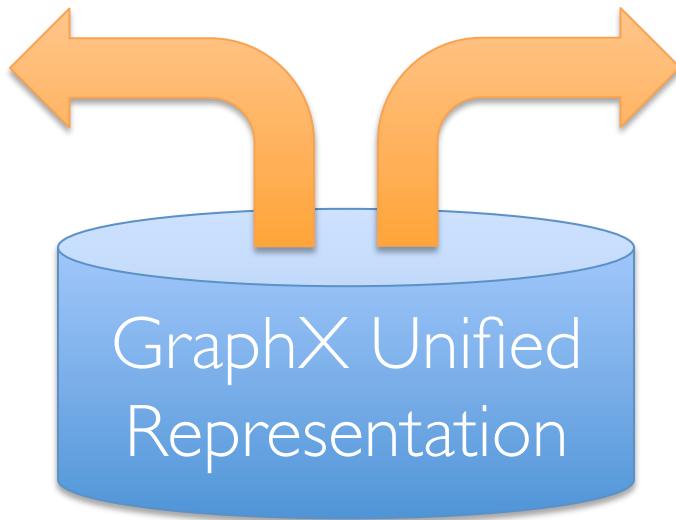


Table View

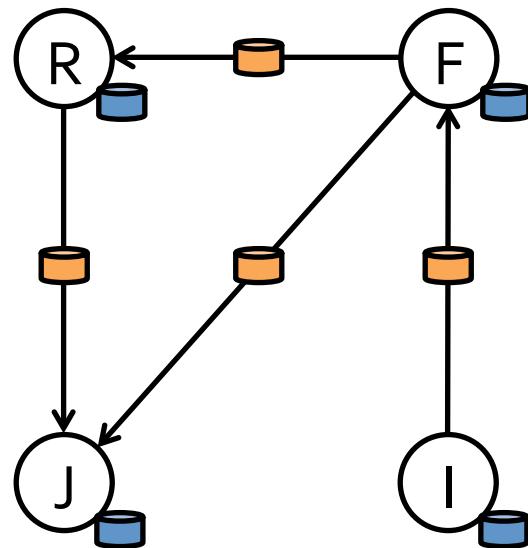


Graph View

Each view has its own **operators** that  
exploit the semantics of the view  
to achieve efficient execution

# View a Graph as a Table

## Property Graph



Vertex Property Table

<b>Id</b>	<b>Property (V)</b>
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

Edge Property Table

<b>SrcId</b>	<b>DstId</b>	<b>Property (E)</b>
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

# Table Operators

Table (RDD) operators are inherited from Spark:

map	reduce	sample
filter	count	take
groupBy	fold	first
sort	reduceByKey	partitionBy
union	groupByKey	mapwith
join	cogroup	pipe
leftOuterJoin	cross	save
rightOuterJoin	zip	...

# Graph Operators

```
class Graph [ V, E ] {  
    def Graph(vertices: Table[ (Id, V) ],  
              edges: Table[ (Id, Id, E) ])  
        // Table views -----  
        def vertices: Table[ (Id, V) ]  
        def edges: Table[ (Id, Id, E) ]  
        def triplets: Table [ ((Id, V), (Id, V), E) ]  
        // Transformations -----  
        def reverse: Graph[V, E]  
        def subgraph(pV: (Id, V) => Boolean,  
                    pE: Edge[V, E] => Boolean): Graph[V, E]  
        def mapV(m: (Id, V) => T ): Graph[T, E]  
        def mapE(m: Edge[V, E] => T ): Graph[V, T]  
        // Joins -----  
        def joinV(tbl: Table [(Id, T)]): Graph[(V, T), E ]  
        def joinE(tbl: Table [(Id, Id, T)]): Graph[V, (E, T)]  
        // Computation -----  
        def mrTriplets(mapF: (Edge[V, E]) => List[(Id, T)],  
                      reduceF: (T, T) => T): Graph[T, E]  
}
```

# Triplets Join Vertices and Edges

The *triplets* operator joins vertices and edges:

**SELECT** *s.Id*, *d.Id*, *s.P*, *e.P*, *d.P*

**FROM** *M* Edges AS *e*

**JOIN** *V* Vertices AS *s*, *V* Vertices AS *d*

**ON** *e.srcId* = *s.Id* **AND** *e.dstId* = *d.Id*

The *mrtriplets* operator sums adjacent triplets.

**SELECT** *t.dstId*, *reduce( map(t) ) AS sum*

**FROM** triplets AS *t* **GROUPBY** *t.dstId*

We express *enhanced* Pregel and GraphLab abstractions using the GraphX operators in less than 50 lines of code!

# Enhanced to Pregel in GraphX

```
pregelPR(i, messageSum) :  
    // Receive all the messages  
    total = 0  
    foreach( msg in messageList ) :  
        total = total + msg  
  
    // Update the rank of this vertex  
    R[i] = 0.15 + total  
combineMsg(a, b) :  
    // Compute sum of two messages  
    sendMsg(i → newR[i], R[i], E[i, j]) :  
    foreach(j ← out_neighbors[i]) :  
        // Compute single message  
        Send msg(R[i]/E[i, j]) to vertex
```

Require Message Combiners

Remove Message Computation from the Vertex Program

# Implementing PageRank in GraphX

```
// Load and initialize the graph
val graph = GraphBuilder.text("hdfs://web.txt")
val prGraph = graph.joinVertices(graph.outDegrees)

// Implement and Run PageRank
val pageRank =
  prGraph.pregel(initialMessage = 0.0, iter = 10) (
    (oldV, msgSum) => 0.15 + 0.85 * msgSum,
    triplet => triplet.src.pr / triplet.src.deg,
    (msgA, msgB) => msgA + msgB)
```

We express the Pregel and GraphLab *like* abstractions using the GraphX operators in less than 50 lines of code!

By composing these operators we can construct entire graph-analytics pipelines.

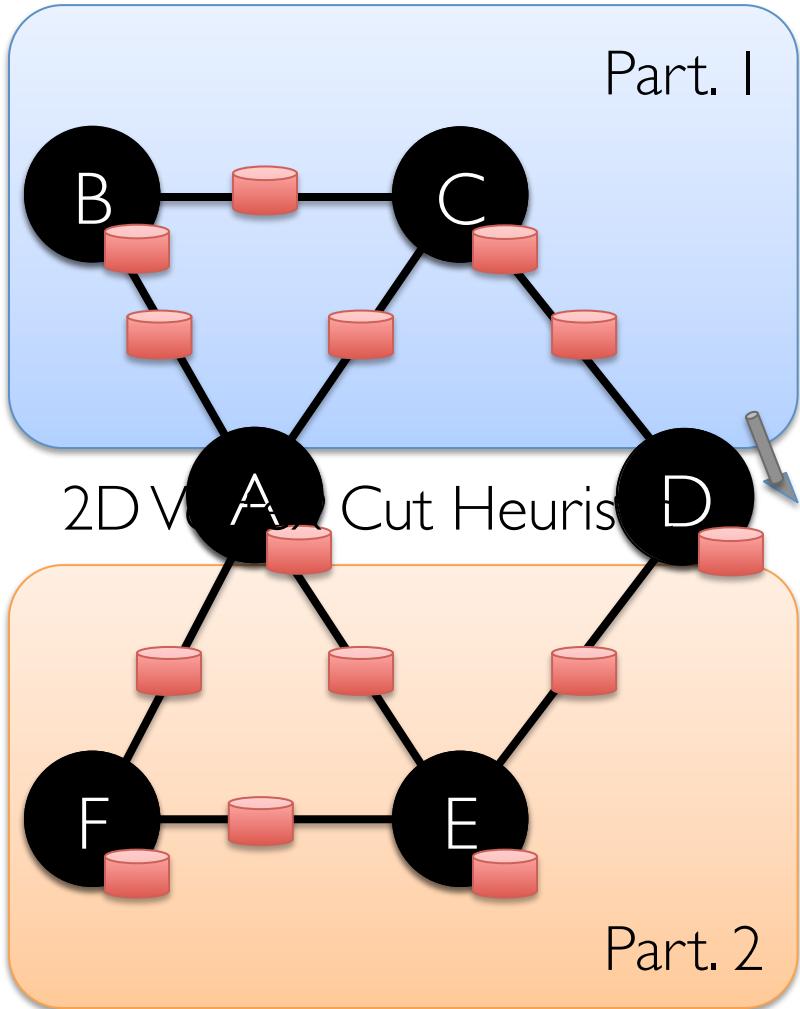
# Example Analytics Pipeline

```
// Load raw data tables  
val verts = sc.textFile("hdfs://users.txt").map(parserV)  
val edges = sc.textFile("hdfs://follow.txt").map(parserE)  
// Build the graph from tables and restrict to recent links  
val graph = new Graph(verts, edges)  
val recent = graph.subgraph(edge => edge.date > LAST_MONTH)  
// Run PageRank Algorithm  
val pr = graph.PageRank(tol = 1.0e-5)  
// Extract and print the top 25 users  
val topUsers = verts.join(pr).top(25).collect  
topUsers.foreach(u => println(u.name + '\t' + u.pr))
```

# GraphX System Design

# Distributed Graphs as Tables (RDDs)

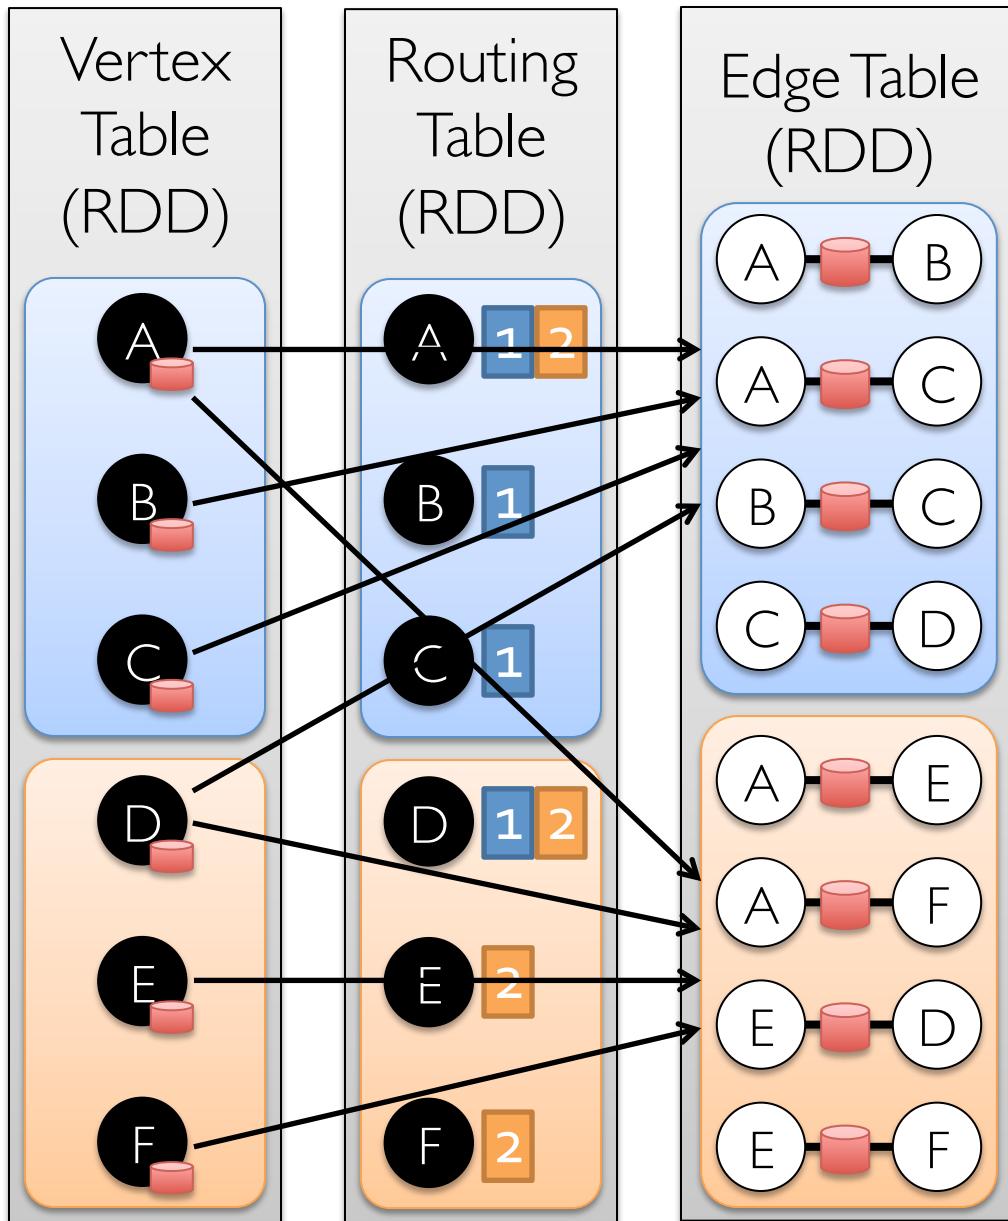
Property Graph



Part. 1

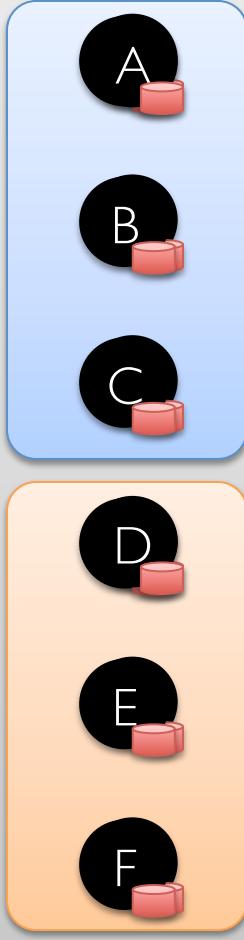
Cut Heuris

Part. 2

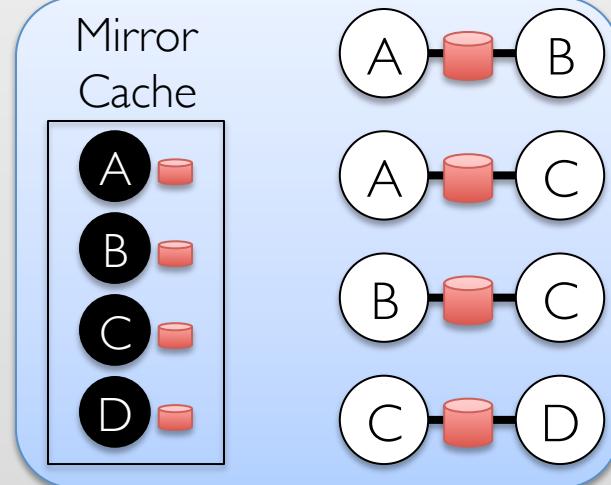


# Caching for Iterative mrTriplets

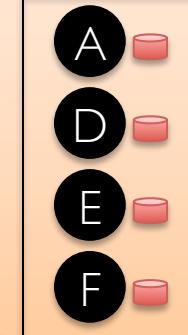
Vertex  
Table  
(RDD)



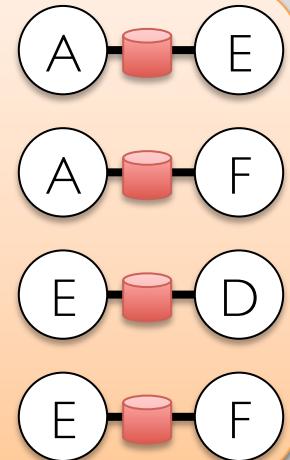
Edge Table  
(RDD)



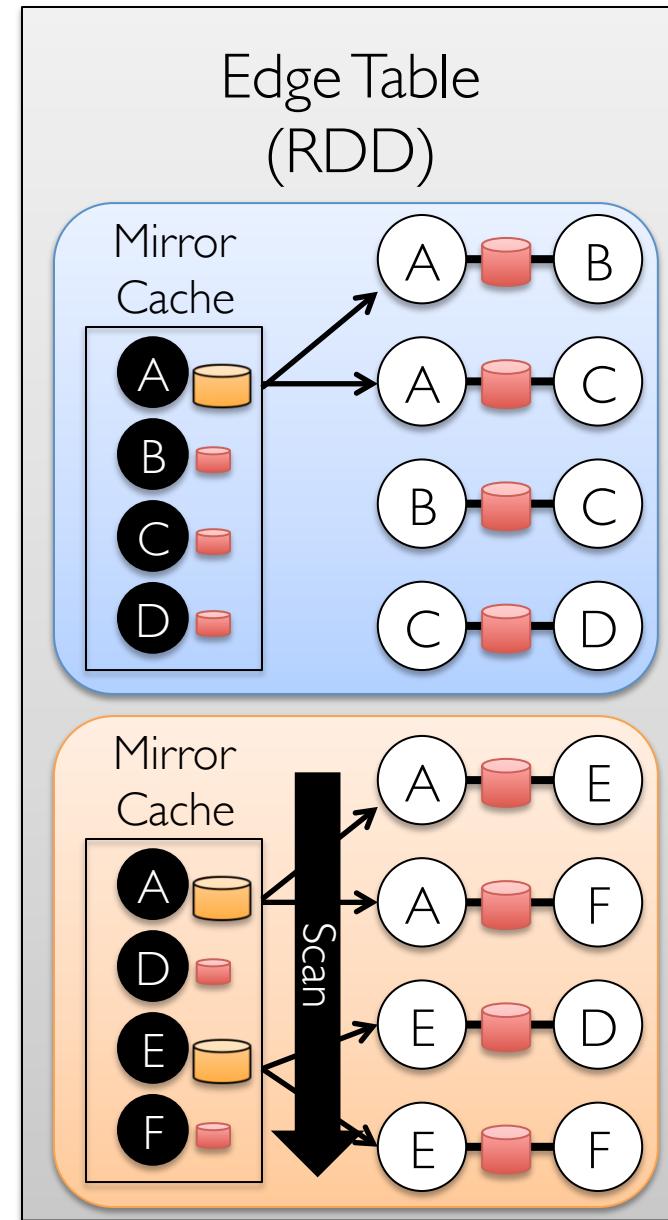
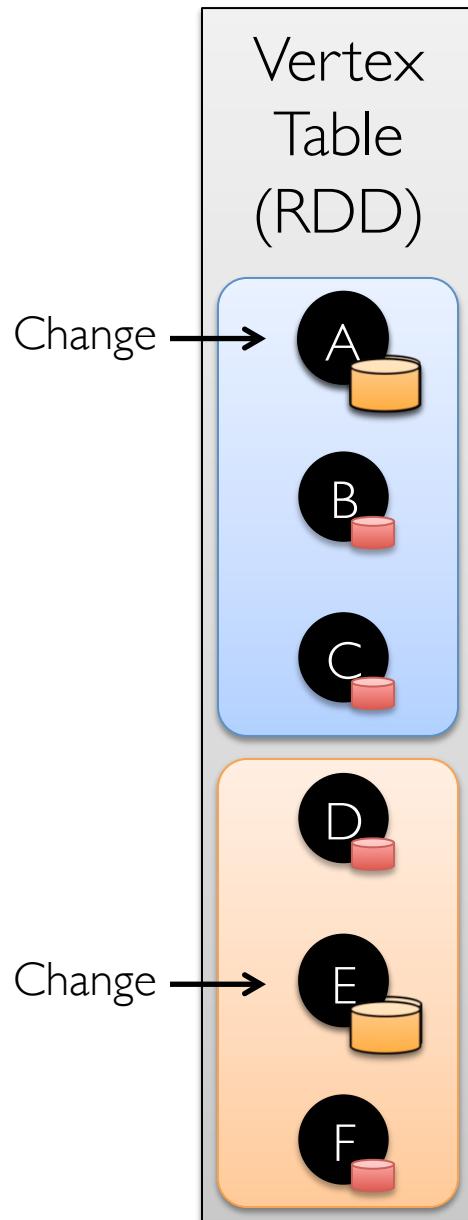
Mirror  
Cache



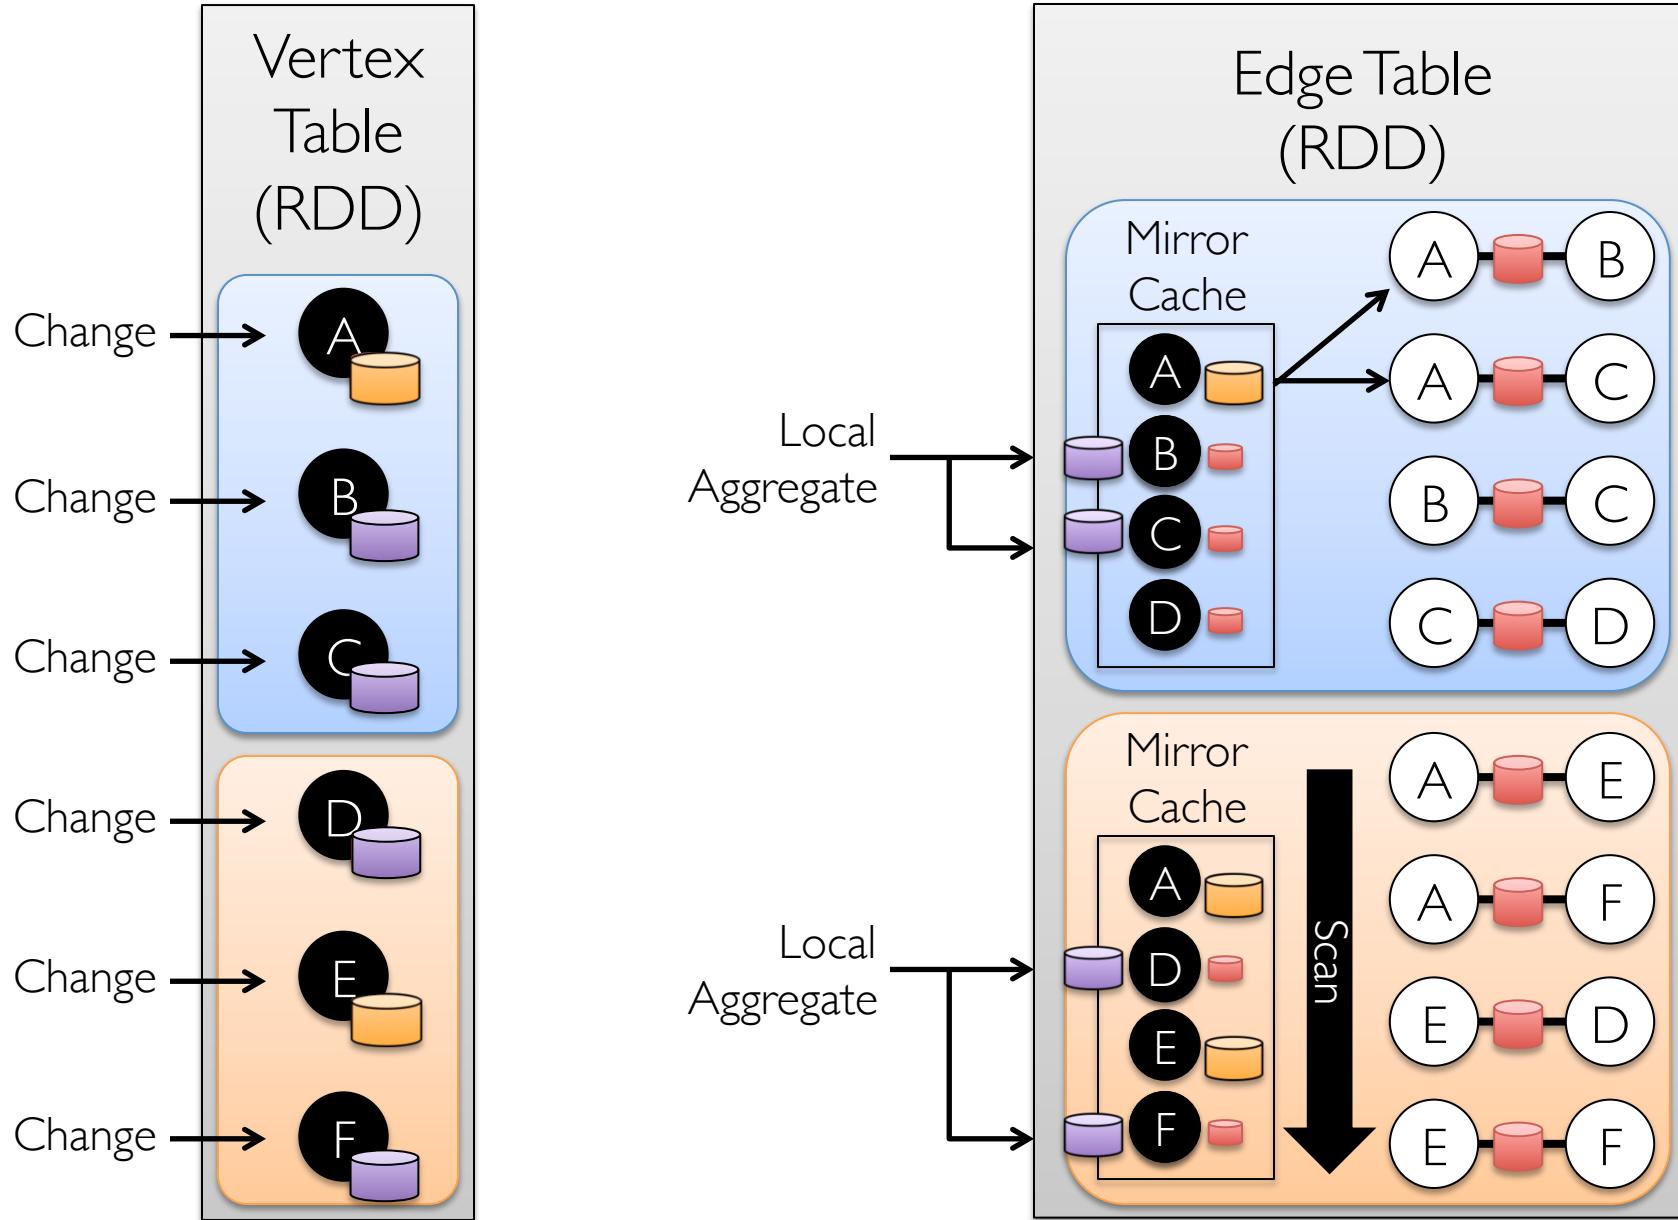
Mirror  
Cache



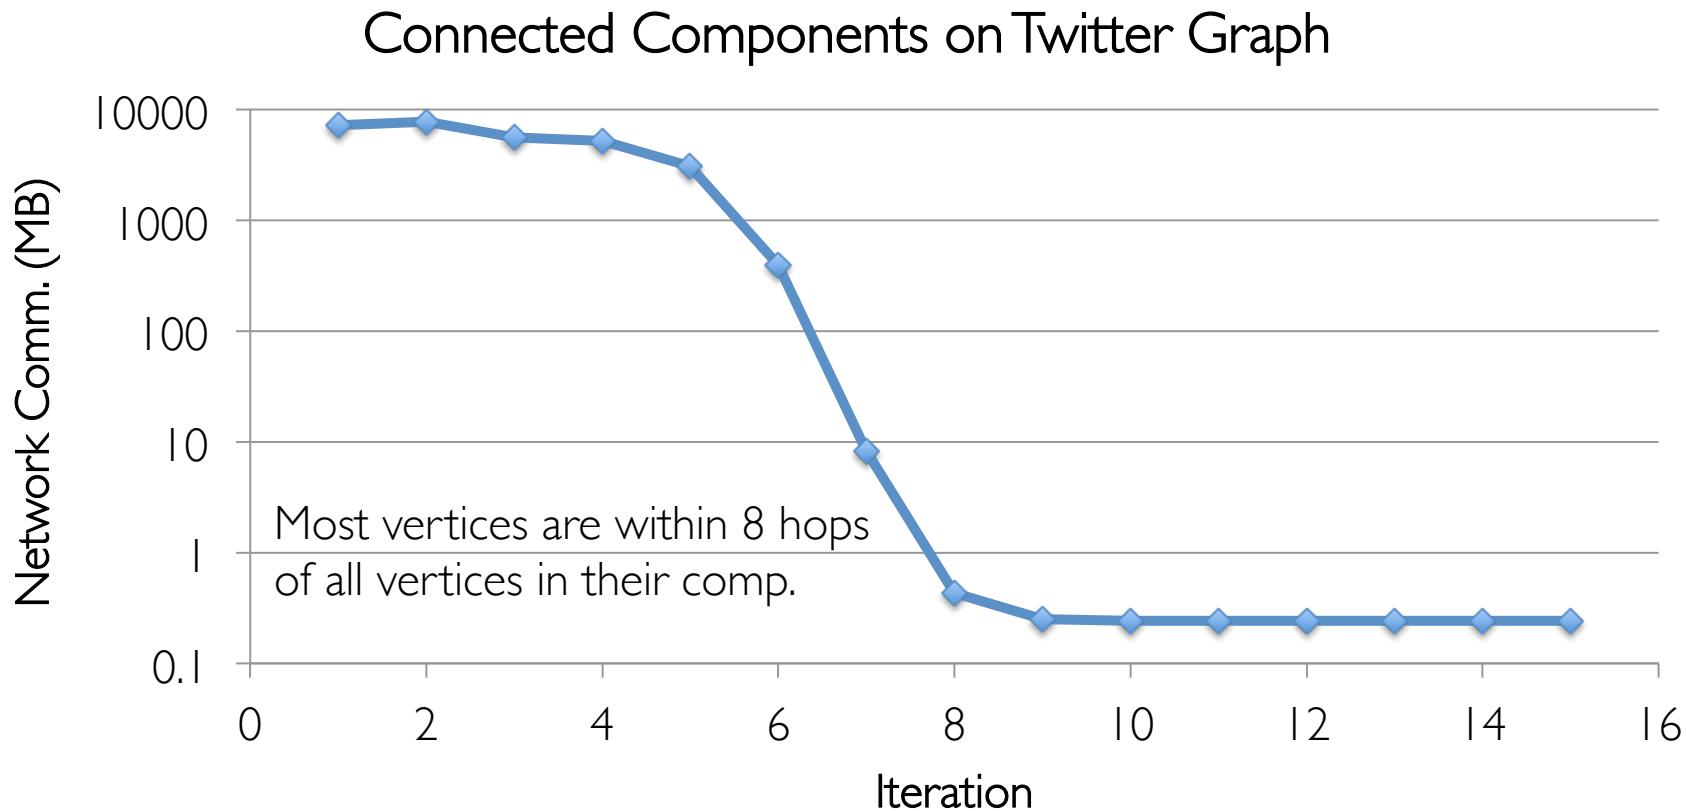
# Incremental Updates for Iterative mrTriplets



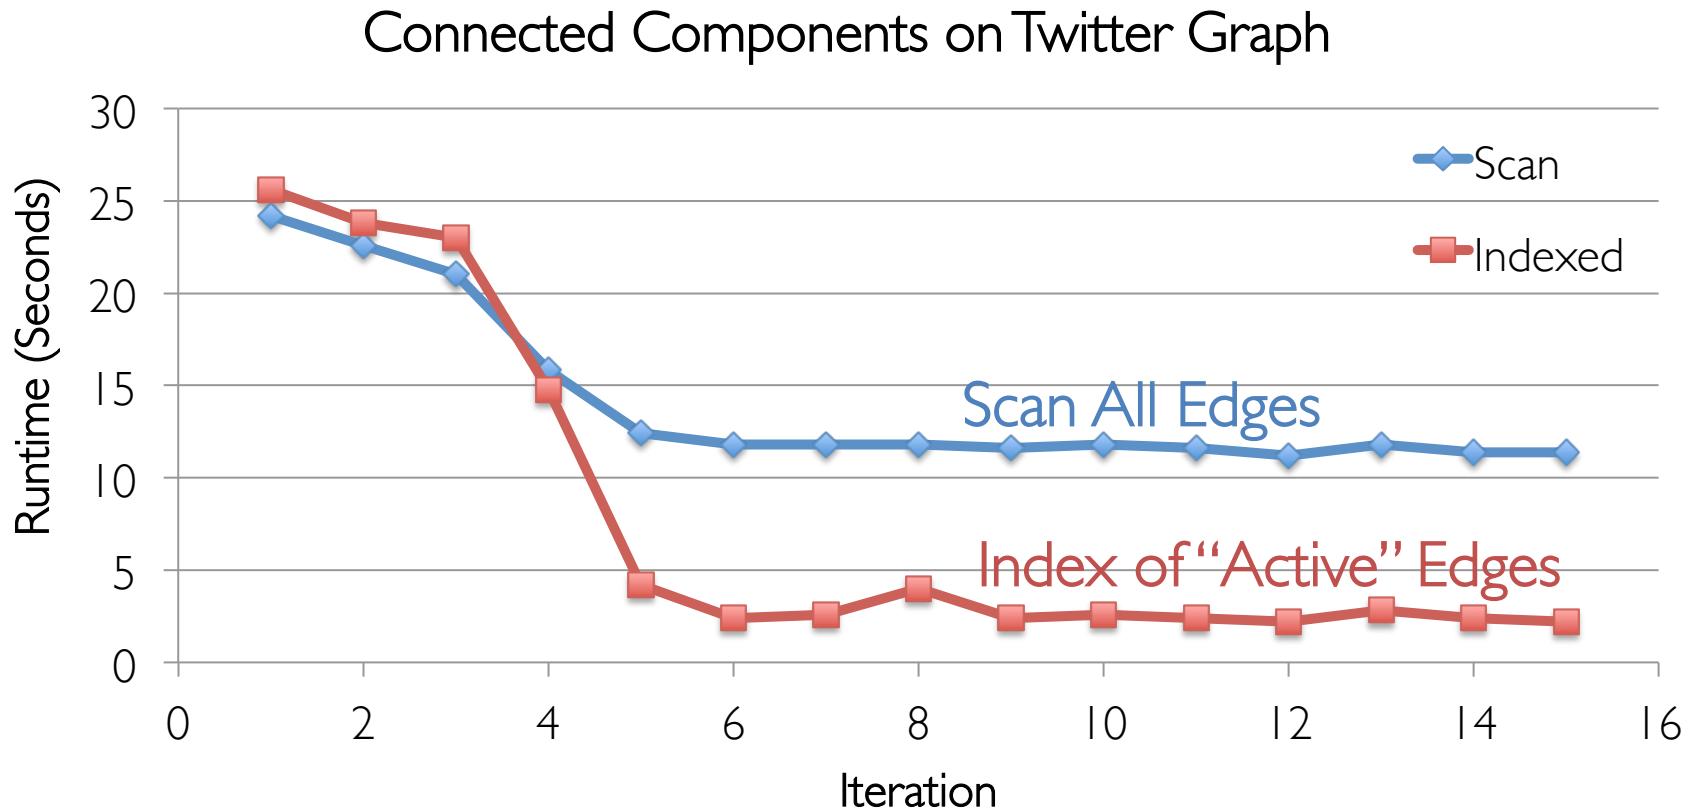
# Aggregation for Iterative mrTriplets



# Reduction in Communication Due to Cached Updates



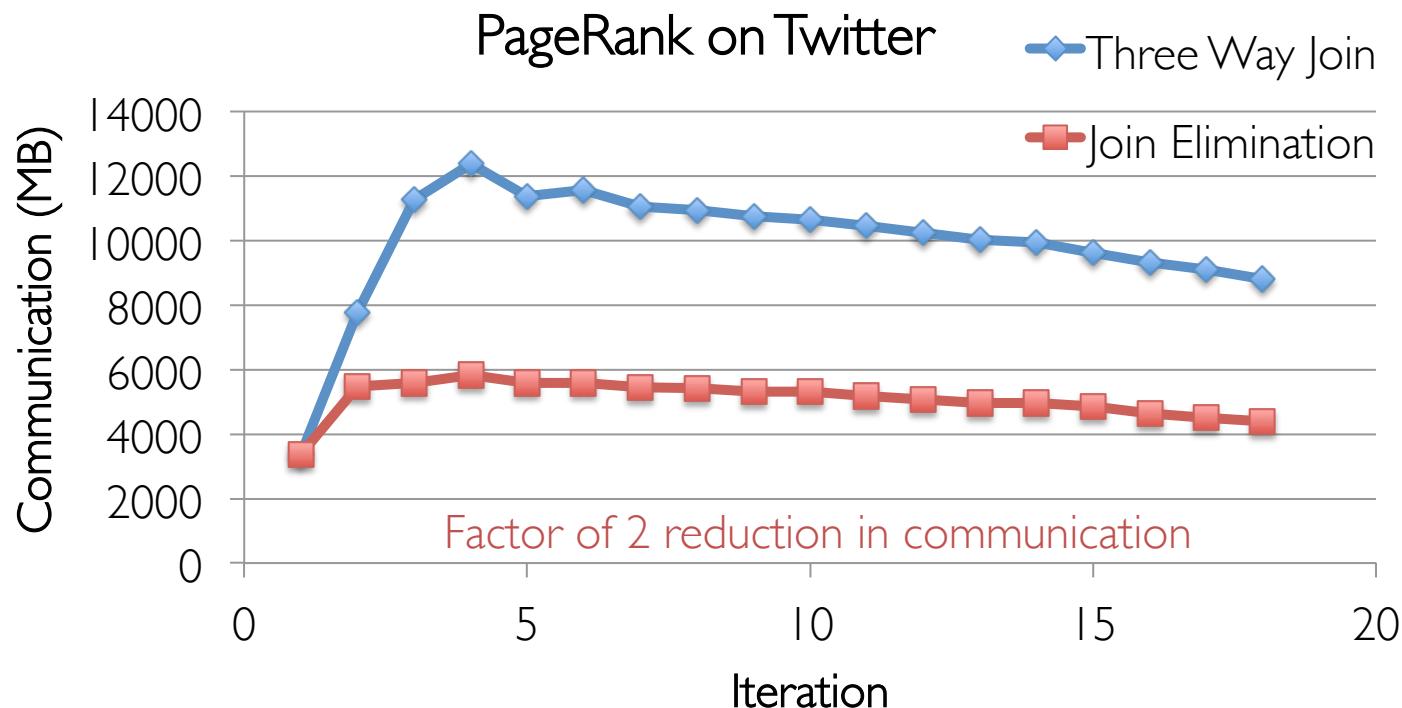
# Benefit of Indexing Active Edges



# Join Elimination

Identify and bypass joins for unused triplet fields

```
sendMsg(i→j, R[i], R[j], E[i,j]):  
    // Compute single message  
    return msg(R[i]/E[i,j])
```



# Additional Query Optimizations

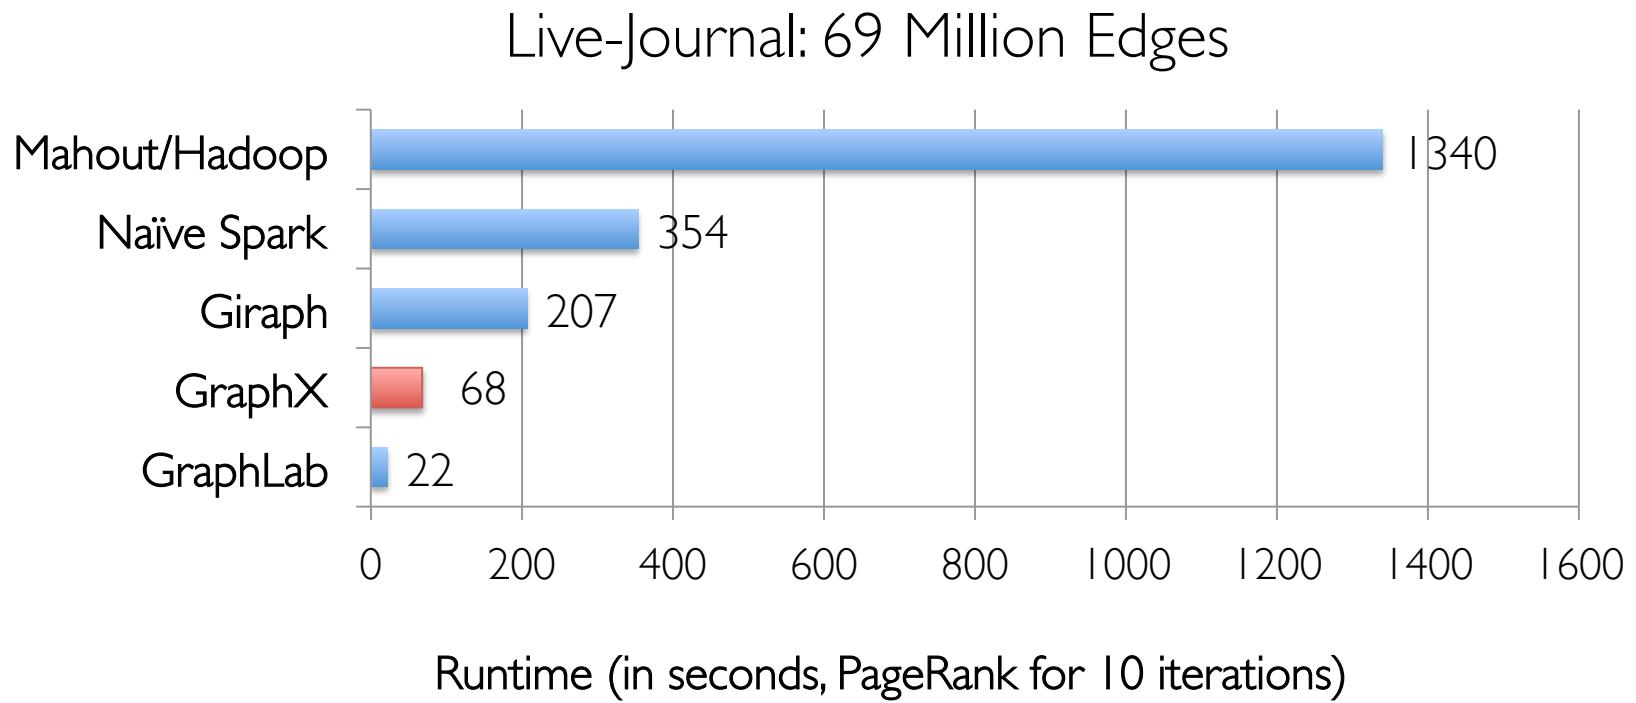
Indexing and Bitmaps:

- » To accelerate joins across graphs
- » To efficiently construct sub-graphs

Substantial Index and Data Reuse:

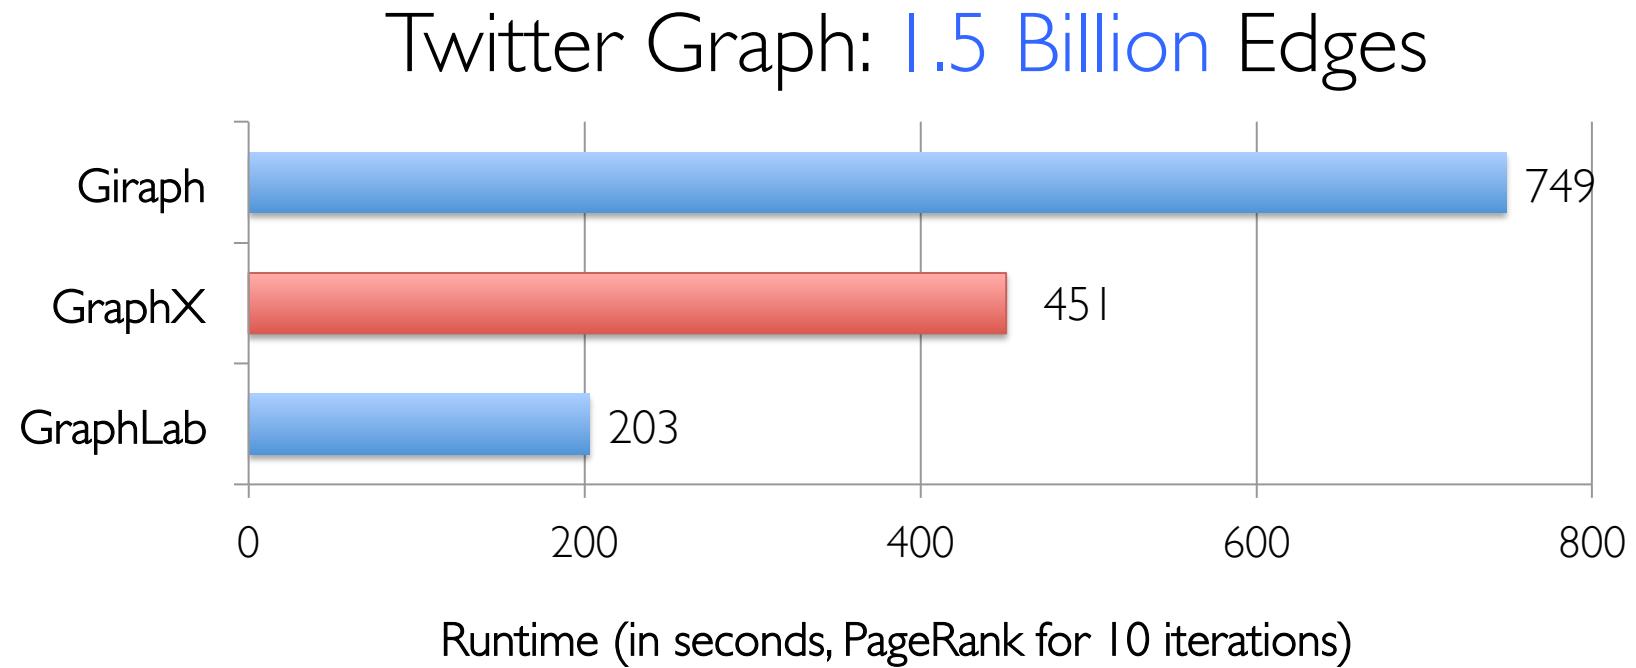
- » Reuse routing tables across graphs and sub-graphs
- » Reuse edge adjacency information and indices

# Performance Comparisons



GraphX is roughly *3x slower* than GraphLab

# GraphX scales to larger graphs



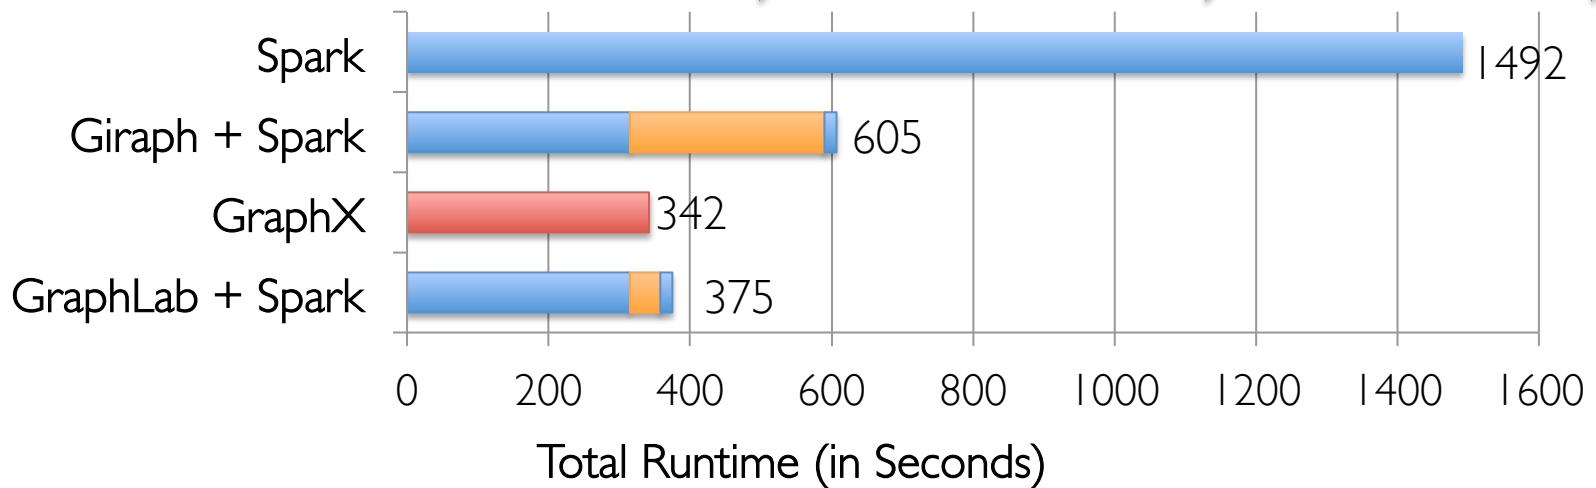
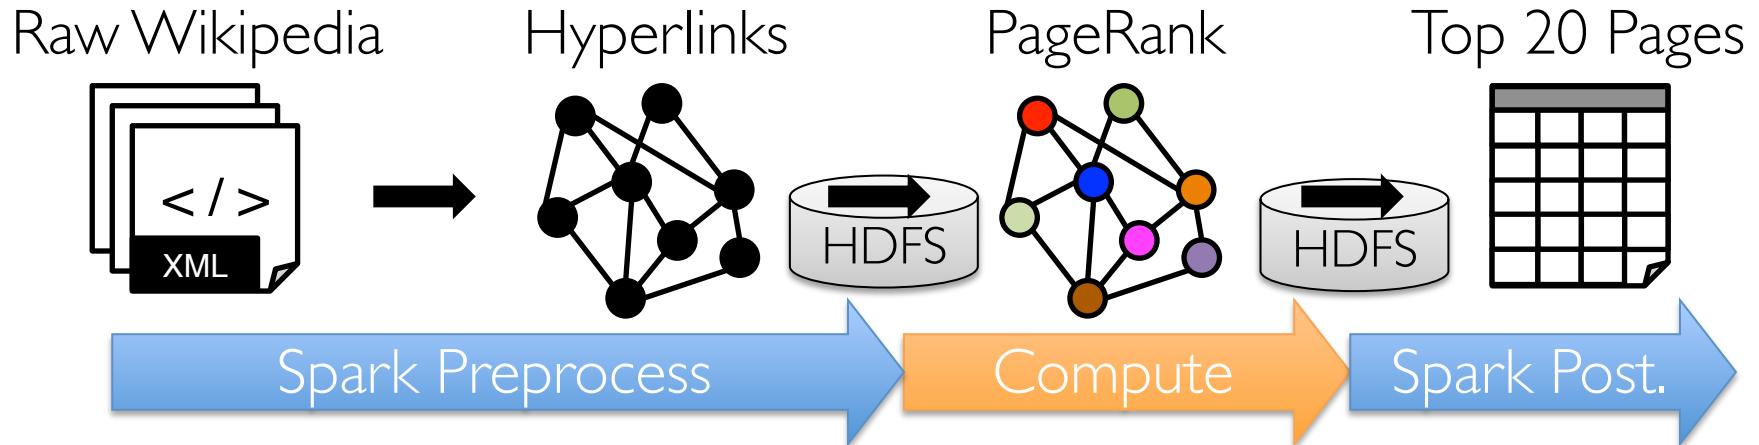
GraphX is roughly *2x slower* than GraphLab

- » Scala + Java overhead: Lambdas, GC time, ...
- » No shared memory parallelism: *2x increase* in comm.

PageRank is just one stage....

What about a pipeline?

# A Small Pipeline in GraphX



Timed end-to-end GraphX is *faster* than GraphLab

# Conclusion and Observations

Domain specific views: *Tables and Graphs*

- » tables and graphs are first-class composable objects
- » specialized operators which exploit view semantics

Single system that efficiently spans the pipeline

- » minimize data movement and duplication
- » eliminates need to learn and manage multiple systems

Graphs through the lens of database systems

- » Graph-Parallel Pattern → Triplet joins in relational alg.
- » Graph Systems → Distributed join optimizations

# Open Source Project

Alpha release as part of Spark 0.9

The screenshot shows a web browser displaying the "GraphX Programming Guide - Spark 0.9.0 Documentation" at [spark.incubator.apache.org/docs/latest/graphx-programming-guide.html](http://spark.incubator.apache.org/docs/latest/graphx-programming-guide.html). The page features the Spark 0.9.0 logo and navigation links for Overview, Programming Guides, API Docs, Deploying, and More. A large graphic on the left illustrates a graph structure connected to a matrix, with the word "GraphX" written in a large, stylized font. Below this, the "Overview" section is visible, followed by a detailed description of GraphX's capabilities and its relationship to other graph-parallel systems like Hadoop and Pregel.

**GraphX**

## Overview

GraphX is the new (alpha) Spark API for graphs and graph-parallel computation. At a high-level, GraphX extends the Spark [RDD](#) by introducing the [Resilient Distributed Property Graph](#): a directed multigraph with properties attached to each vertex and edge. To support graph computation, GraphX exposes a set of fundamental operators (e.g., [subgraph](#), [joinVertices](#), and [mapReduceTriplets](#)) as well as an optimized variant of the [Pregel](#) API. In addition, GraphX includes a growing collection of graph [algorithms](#) and [builders](#) to simplify graph analytics tasks.

## Background on Graph-Parallel Computation

From social networks to language modeling, the growing scale and importance of graph data has driven the development of numerous new *graph-parallel* systems (e.g., [Giraph](#) and [GraphLab](#)). By restricting the types of computation that can be expressed and introducing new techniques to partition and distribute graphs, these systems can efficiently execute sophisticated graph algorithms orders of magnitude faster than more general *data-parallel* systems.

**Data-Parallel**

**Graph-Parallel**

# Active Research

## Static Data → Dynamic Data

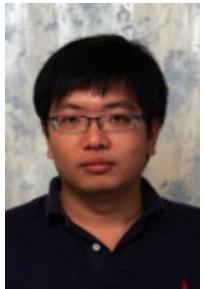
- » Apply GraphX unified approach to time evolving data
- » Materialized view maintenance for graphs

## Serving Graph Structured Data

- » Allow external systems to interact with GraphX
- » Unify distributed graph databases with relational database technology

# Collaborators

GraphLab:



Yucheng  
Low

Haijie  
Gu

Aapo  
Kyrola

Danny  
Bickson

Carlos  
Guestrin

Alex  
Smola

Guy  
Blelloch

GraphX:



Reynold  
Xin

Ankur  
Dave

Daniel  
Crankshaw

Michael  
Franklin

Ion  
Stoica

# Thanks!

<http://tinyurl.com/ampgraphx>

[jegonzal@eecs.berkeley.edu](mailto:jegonzal@eecs.berkeley.edu)