

Prediction Serving

what happens after learning?

Joseph E. Gonzalez

Asst. Professor, UC Berkeley


jegonzal@cs.berkeley.edu

Co-founder, Dato Inc.

joseph@dato.com




Outline



Daniel Crankshaw, Xin Wang
Michael Franklin, & Ion Stoica

Learning



Big Model

Timescale: minutes to days


Systems: offline and batch optimized

*Heavily studied ... major focus of the **AMPLab***

Learning




Training



Inference

Query

Decision





Application

Learning



Training

Inference



Application

Timescale: ~10 milliseconds

Systems: *online* and *latency* optimized
Less studied ...

Learning




Training

Feedback


Big Model

Inference



Query

Decision




Application

Learning




Training



Inference

Decision



Application

Timescale: hours to weeks

Systems: combination of systems

Less studied ...

Feedback

Learning

Inference

Adaptive
(~1 seconds)

Responsive
(~10ms)

Feedback

Big Model

Training Data


Query

Decision

Application

VELOX Model Serving System [CIDR'15]

Daniel Crankshaw, Peter Bailis, Haoyuan Li, Zhao Zhang,
Joseph Gonzalez, Michael J. Franklin, Ali Ghodsi, and Michael I. Jordan




Key Insight:

Decompose models into fast and slow changing components


Learning



Training



Inference



Application


Feedback

Learning



Training

Slow Changing Model




Feedback

Slow


Inference

Fast Changing Model



Query

Decision



Application

Fast Feedback

Hybrid Offline + Online Learning

Update feature functions **offline** using batch solvers

- Leverage high-throughput systems (Tensor Flow)
- Exploit slow change in population statistics

$$f(x; \theta)^T$$

$$w_u$$

Update the user weights **online**:

- Simple to train + more robust model
- Address rapidly changing user statistics


Common modeling structure

$$f(x; \theta)^T w_u$$


Matrix
Factorization



Deep
Learning



Ensemble
Methods




Learning



Training

Slow Changing Model




Feedback

Slow


Inference

Fast Changing Model



Query

Decision



Application


Fast Feedback

Learning



Training


Slow Changing Model



Feedback


Slow

Fast Changing Model per user



Query


Decision



Application

Fast Feedback


Velox Online Learning for Recommendations (20-News Groups)



Online Updates: 0.4 ms
Retraining: 7.1 seconds

*>4 orders-of-magnitude
faster adaptation
given sufficient offline
training data*

Velox Online Learning for Recommendations (20-News Groups)



Partial Updates: 0.4 ms
Retraining: 7.1 seconds


*>4 orders-of-magnitude
faster adaptation*

Learning



Training


Slow Changing Model



Feedback

Slow

Fast Changing Model per user



Query

Decision




Application

Fast Feedback

Learning




Training



Inference

Query



Decision




Feedback

Slow

VELOX: the Missing Piece of BDAS




VELOX: the Missing Piece of BDAS



— 
Berkeley
Data
Analytics
Stack

VELOX: the Missing Piece of BDAS



— 
Berkeley
Data
Analytics
Stack


VELOX Architecture

Fraud
Detection



Content
Rec.

NETFLIX



VELOX Architecture

Fraud
Detection



Content
Rec.



Personal
Asst.



Robotic
Control



Machine
Translation



Keystone ML

MLLib

Spark

Velox

Single JVM Instance



Caffe



TensorFlow



OKALDI
theano

VELOX as a Middle Layer Arch?

Fraud
Detection



Content
Rec.



Personal
Asst.



Robotic
Control



Machine
Translation



Generalize Velox?

theano

KeystoneML

Dato



Caffe

TensorFlow

scikit
learn

dmlc
mxnet

VW

KALDI

Clipper

A Low-Latency Online Prediction Serving System

Daniel Crankshaw

Xin Wang

Michael Franklin

Joseph E. Gonzalez

Ion Stoica



Clipper Generalizes Velox Across ML Frameworks

Fraud
Detection



Content
Rec.



Personal
Asst.



Robotic
Control



Machine
Translation



Clipper

theano

KeystoneML

Dato



Create

Caffe

TensorFlow

scikit
learn

dmlc
mxnet

VW

KALDI



Clipper



Key Insight:

The challenges of prediction serving can be addressed between end-user applications and machine learning frameworks

As a result, Clipper is able to:

- **hide complexity**
 - by providing a *common prediction interface*
- **bound latency** and **maximize throughput**
 - through *approximate caching* and *adaptive batching*
- enable *robust online learning* and **personalization**
 - through *generalized split-model correction policies*

without modifying machine learning frameworks or end-user applications

Clipper Design Goals

Low and **bounded** latency predictions

- interactive applications need reliable latency objectives

Up-to-date and personalized predictions **across models** and **frameworks**

- generalize the split model decomposition

Optimize **throughput** for performance under heavy load

- single query can trigger many predictions

Simplify deployment

- serve models using the original code and systems

Clipper Architecture

Fraud
Detection



Content
Rec.



Personal
Asst.



Robotic
Control



Machine
Translation



Clipper

theano

KeystoneML

Dato



Caffe

TensorFlow

scikit
learn

dmlc
mxnet

VW

KALDI

Clipper Architecture

Applications



Predict

RPC/REST Interface

Observe

Clipper

theano

Dato



KeystoneML

Caffe

TensorFlow


scikit
learn

dmlc
mxnet



VW
KALDI

Clipper Architecture



Clipper Architecture

Applications



NETFLIX



Predict

RPC/REST Interface

Observe

Clipper

*Improve accuracy through **ensembles**,
online learning and **personalization***

Correction Layer

*Provide a **common interface** to models
while **bounding latency** and
maximizing throughput.*

Model Abstraction Layer

RPC

RPC

RPC

RPC

Model Wrapper (MW)



MW

Caffe

MW




MW



...


Clipper Architecture



Approximate Caching

Adaptive Batching

Model Abstraction Layer




Provides a unified generic prediction API across **frameworks**

- **Reduce Latency** → Approximate Caching
- **Increase Throughput** → Adaptive Batching
- **Simplify Deployment** → RPC + Model Wrapper

Approximate Caching

Adaptive Batching


Model Abstraction Layer



Approximate Caching

Adaptive Batching

Model Abstraction Layer




Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes
 - Resource isolation

Approximate Caching

Adaptive Batching

Model Abstraction Layer




Common Interface → Simplifies Deployment:

- Evaluate models using original code & systems
- Models run in separate processes
 - Resource isolation
 - Scale-out

Problem: frameworks optimized for **batch processing** not **latency**


Adaptive Batching to Improve Throughput

- Why batching helps:



A single page load may generate many queries

Hardware Acceleration



Helps amortize system overhead

- Optimal batch depends on:


- hardware configuration
- model and framework
- system load

Clipper Solution:



be as **slow** as **allowed**...

- Inc. batch size *until the latency objective is exceeded* (**Additive Increase**)
- If latency exceeds SLO cut batch size by a fraction (**Multiplicative Decrease**)

Tensor Flow Conv. Net (GPU)




Comparison to TensorFlow Serving



Takeaway: Clipper is able to **match the average latency** of TensorFlow Serving while reducing **tail latency (2x)** and **improving throughput (2x)**

Approximate Caching to Reduce Latency

- Opportunity for caching



Popular items may be evaluated frequently

- Need for **approximation**




Bag-of-Words Model




High Dimensional and continuous valued queries have low cache hit rate.

Clipper Solution: Approximate Caching

apply *locality sensitive hash functions*



Clipper Architecture



Goal:

*Maximize **accuracy** through **ensembles**, **online learning**, and **personalization***

Generalize the **split-model** insight from Velox to achieve:

- **robust predictions** by combining multiple models & frameworks
- **online learning** and **personalization** by correcting and personalizing **predictions** in response to feedback


Learning

Inference


Velox



Slow Changing
Model



Fast Changing
User Model




Application

Feedback


Fast Feedback

Slow

Learning



**Slow Changing
Model**



Feedback

Slow

Inference


Clipper

**Fast Changing
User Model**



Application


Fast Feedback



Correction Policy

Improves prediction **accuracy** by:

- Incorporating real-time **feedback**
- Managing **personalization**
- **Combine** models & **frameworks**
 - enables frameworks to **compete**



Improved Prediction Accuracy (ImageNet)

| System | Model | Error Rate | #Errors |
|------------|--------------|------------|---------|
| Caffe | VGG | 13.05% | 6525 |
| Caffe | LeNet | 11.52% | 5760 |
| Caffe | ResNet | 9.02% | 4512 |
| TensorFlow | Inception v3 | 6.18% | 3088 |

sequence of pre-trained state-of-the-art models

Improved Prediction Accuracy

| System | Model | Accuracy (%) | Errors |
|------------|--------------|--------------|--------|
| Caffe | ResNet | 9.02% | 6525 |
| Caffe | Inception v3 | 6.18% | 5760 |
| Caffe | Ensemble | 5.86% | 4512 |
| TensorFlow | | | 3088 |
| Clipper | | | 2930 |

**5.2% relative improvement
in prediction accuracy!**


Cost of Ensembles

Increased Load


- *Solutions:*
 - **Caching and Batching**
 - **Load-shedding** correction policy can prioritize frameworks

Stragglers

- e.g., framework fails to meet SLO
- *Solution: Anytime* predictions
 - Correction policy must render predictions with missing inputs
 - e.g., built-in correction policies **substitute expected value**



Evaluation of Throughput Under Heavy Load



Takeaway: Clipper is able to *gracefully degrade accuracy* to maintain availability under heavy load.

Conclusion




Clipper sits between applications and ML frameworks to



Clipper



- to **simplifying deployment**
 - **bound latency** and **increase throughput**
 - and enable **real-time learning** and **personalization** across **machine learning frameworks**



Clipper

theano Dato Caffe scikit
 KeystoneML CreateTensorFlow mxnet KALDI
 vw

Ongoing & Future Research Directions

- Serving and updating RL models
- Bandit techniques in correction policies
- Splitting inference across the cloud and the client to reduce latency and bandwidth requirements
- Secure model evaluation on the client (model DRM)

Coarsening + Anytime Predictions

$$f_i(x; \theta) \approx f_i(z; \theta)$$

$$f_i(x; \theta) \approx \mathbb{E}[f_i(x; \theta)]$$

