

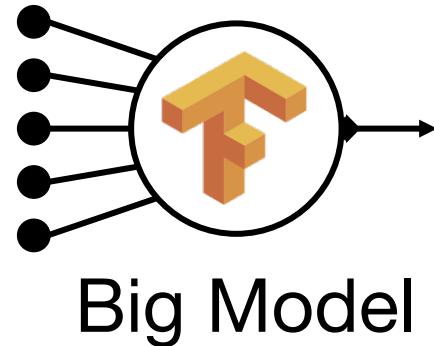


# RISE to the Challenges of AI Systems

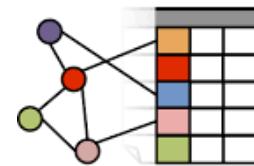
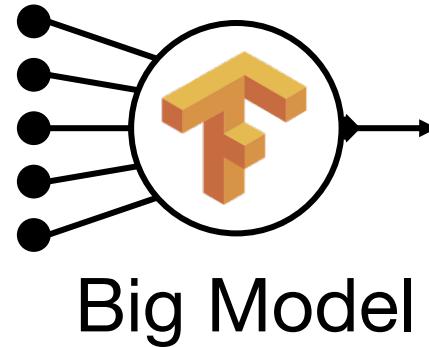
Joseph E. Gonzalez

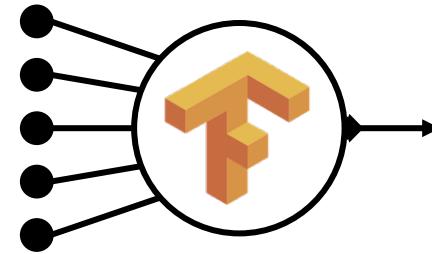
Assistant Professor, UC Berkeley

[jegonzal@cs.berkeley.edu](mailto:jegonzal@cs.berkeley.edu)

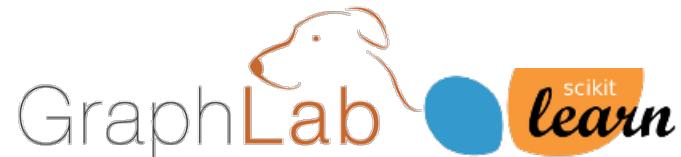


Large-Scale parallel and distributed systems





Big Model



# How to do Research in AI Systems

- Manage **Complexity**
  - seek parsimony in system design
  - great systems research is often about what features are taken away
  - Do a few things well and be composable
- Identify **Tradeoffs**
  - With each design decision what do you gain and lose?
  - What trade-offs are fundamental?
- **Evaluate** your System
  - **Positive:** How fast and scalable is it and *why*?
  - **Negative:** When does it fail and what are its *limitations*?

# Hemingway\*

## Modeling Throughput and Convergence for ML Workloads



Shivaram  
Venkataraman

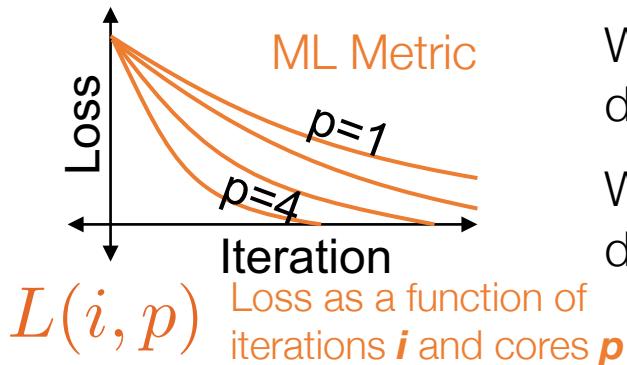
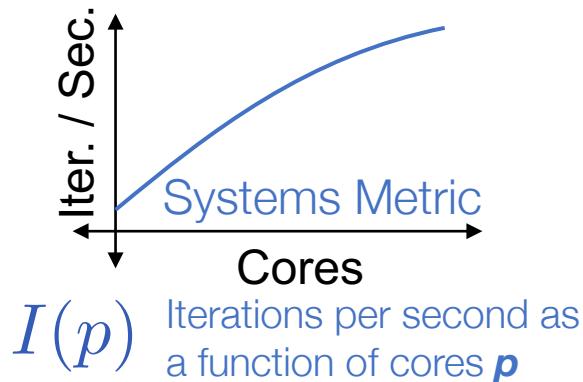


Xinghao  
Pan



Zi  
Zheng

- What is the best algorithm and level of parallelism for an ML task?
  - **Trade-off:** Parallelism, Coordination, & Convergence
- **Research challenge:** Can we model this trade-off explicitly?



We can estimate  $I$  from data on many systems

We can estimate  $L$  from data for our problem

# Hemingway\*

## Modeling Throughput and Convergence for ML Workloads



Shivaram  
Venkataraman



Xinghao  
Pan



Zi  
Zheng

- What is the best algorithm and level of parallelism for an ML task?
  - **Trade-off:** Parallelism, Coordination, & Convergence
- **Research challenge:** Can we model this trade-off explicitly?

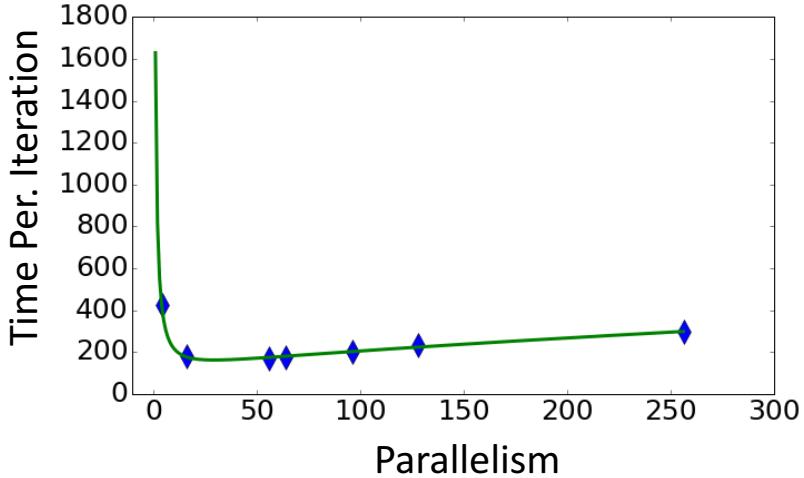
$$L(i, p) \text{ Loss as a function of iterations } i \text{ and cores } p$$

$$I(p) \text{ Iterations per second as a function of cores } p$$

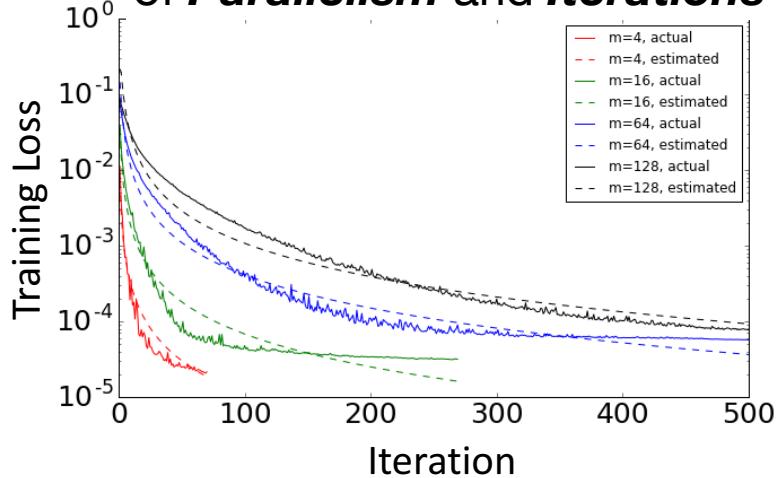
$$\text{loss}(t, p) = L(t * I(p), p)$$

- How long does it take to get to a given loss?
- Given a time budget and number of cores which algorithm will give the best result?

## System Performance as a function of *Parallelism*

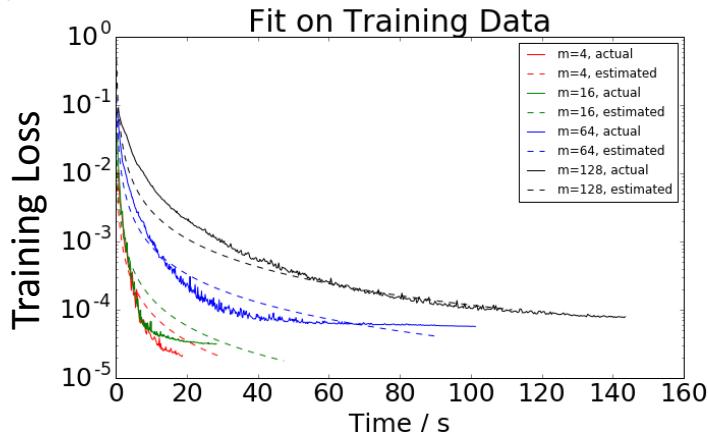


## Convergence as a function of *Parallelism and Iterations*



**Hemingway: Modeling Distributed Optimization Algorithms.**

Xinghao Pan, Shivaram Venkataraman,  
Zizheng Tai, Joseph Gonzalez.  
NIPS'16 ML-Sys Workshop.



**Convergence as a fn. of *Time* and *Parallelism***

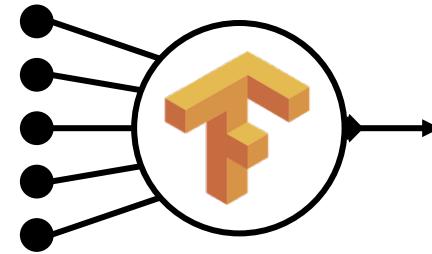
# Take away ...

try to decouple

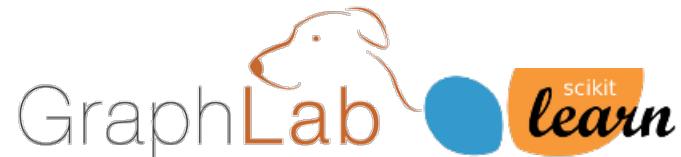
**System  
Improvements**

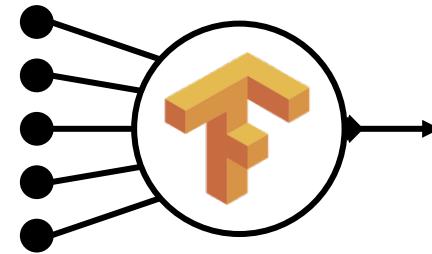
**Algorithm  
Improvements**

use data collection + sparse modeling  
to understand your system



Big Model





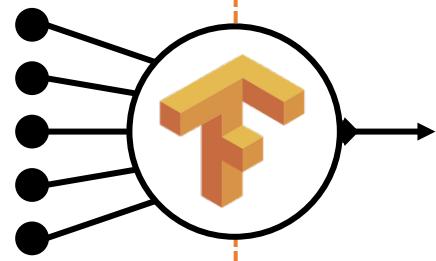
Big Model

-amplab

# Learning



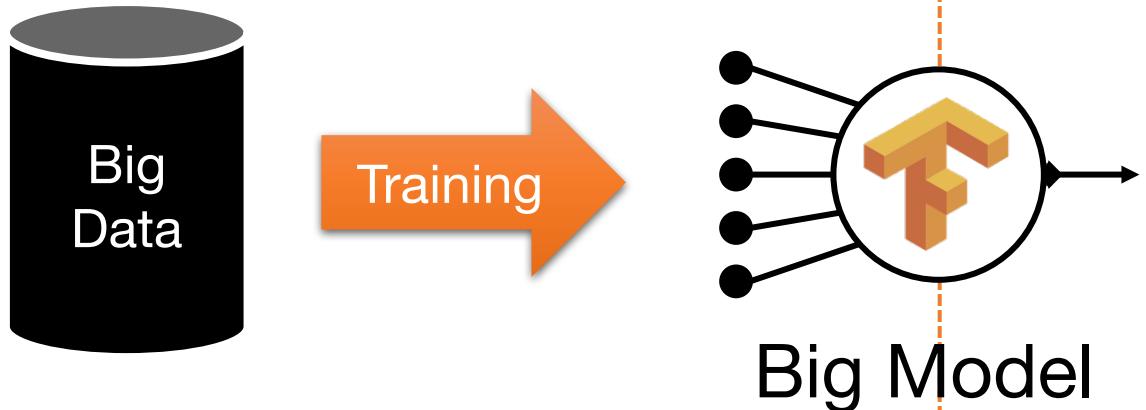
Training



Big Model



# Learning



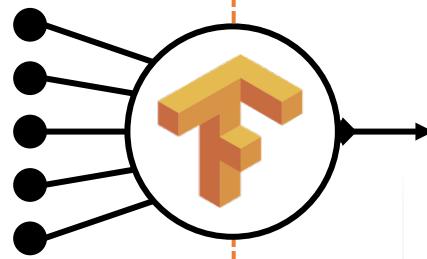
# Conference Papers



# Learning



Training



Big Model

# Conference Papers

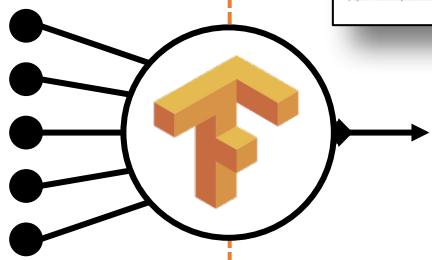


# Dashboards and Reports

# Learning



Training



Big Model



# Conference Papers

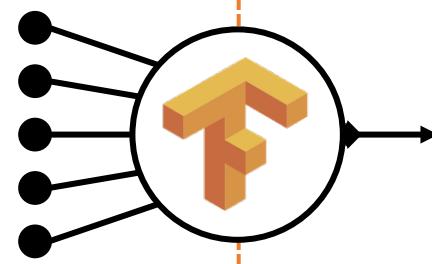


# Dashboards and Reports



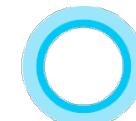
# Drive Actions

# Learning



Big Model

# Drive Actions



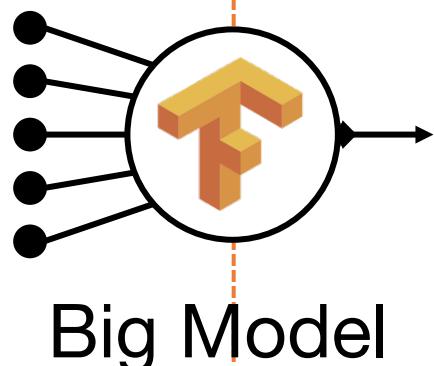
Hi, I'm Cortana.



# Learning



Training

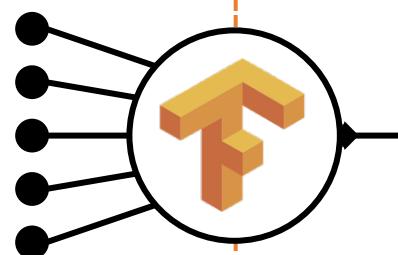


# Inference

# Learning



Training

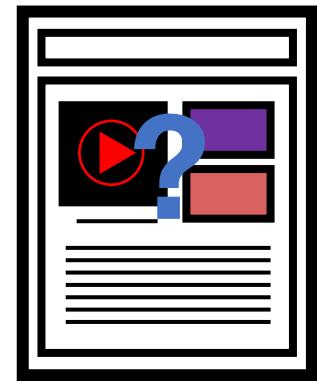


Big Model

# Inference

Query

Decision



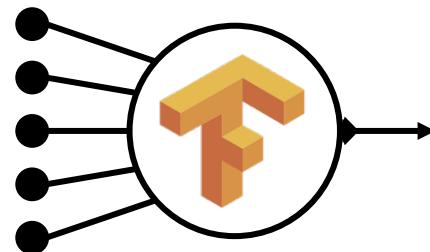
Application

# Learning

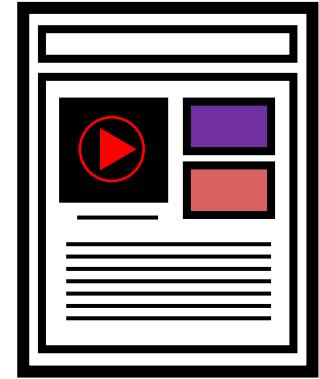


Training

# Inference



Big Model



Application

Often **overlooked**

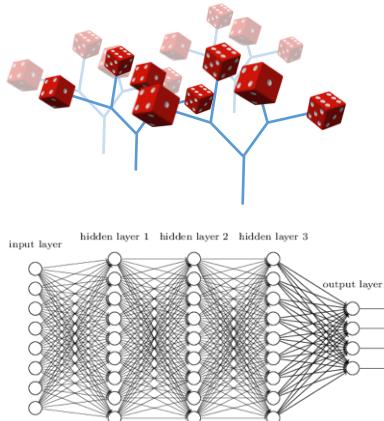
Timescale: ~10 milliseconds

**Billions of Queries a Day → Costly**

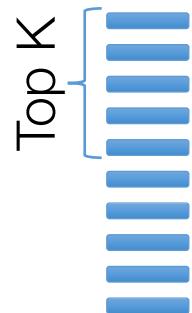
# why is **Inference** challenging?

Need to render **low latency** (< 10ms) predictions for **complex**

## Models



## Queries

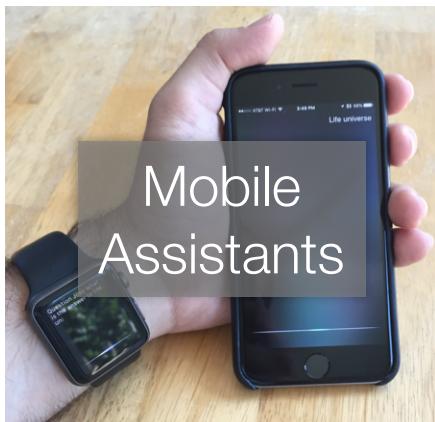


## Features

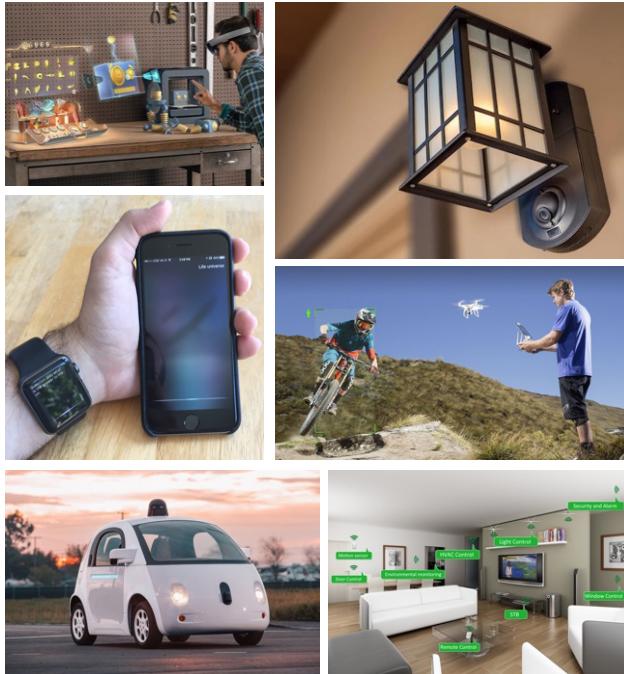
```
SELECT * FROM users JOIN items,  
click_logs, pages WHERE ...
```

under **heavy load** with system **failures**.

# Inference is moving beyond the cloud



# Inference is moving beyond the cloud



# Opportunities

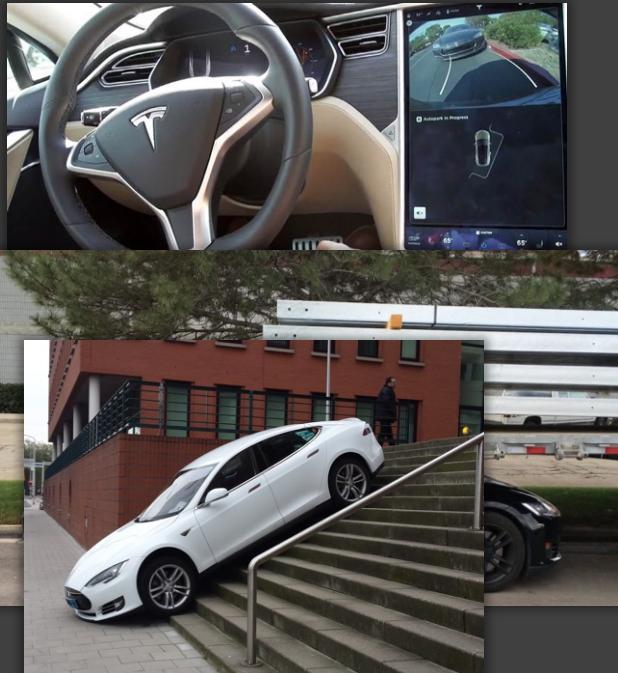
- Reduce latency and improve privacy
  - Address network partitions

# Research Challenges

- Minimize **power consumption**
  - **Limited hardware** & long life-cycles
  - Develop new **hybrid models** to leverage the cloud and edge devices

# Robust Inference is critical

Self “Parking” Cars



Self “Driving” Cars



Chat Als



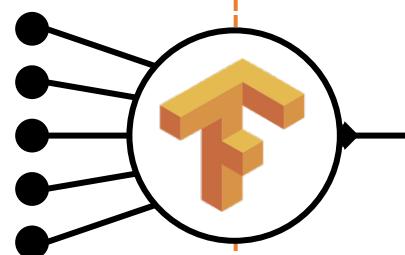
# Learning



Training

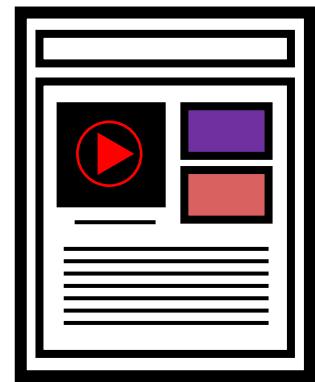
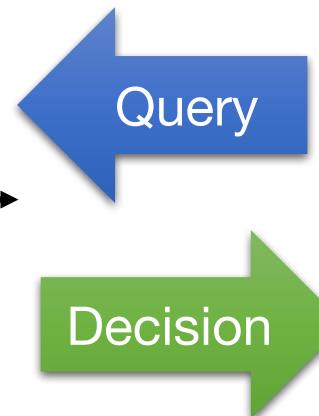


Feedback



Big Model

# Inference

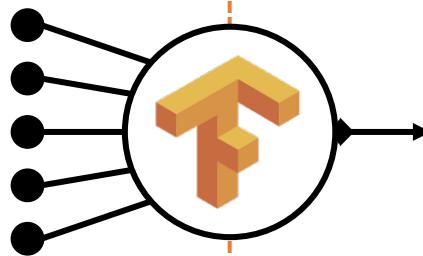


Application

# Learning

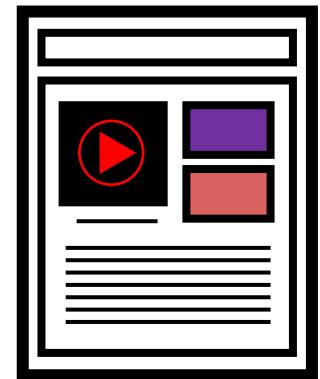


Training



# Inference

Decision



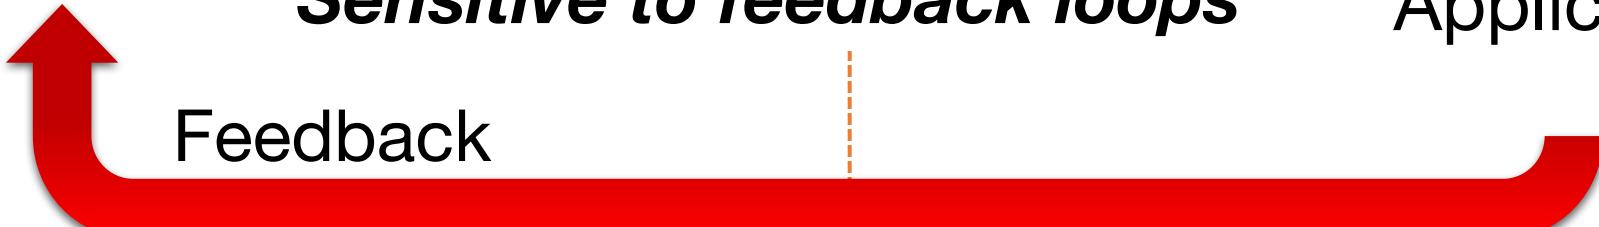
Application

**Timescale:** hours to weeks

***Often re-run training***

***Sensitive to feedback loops***

Feedback



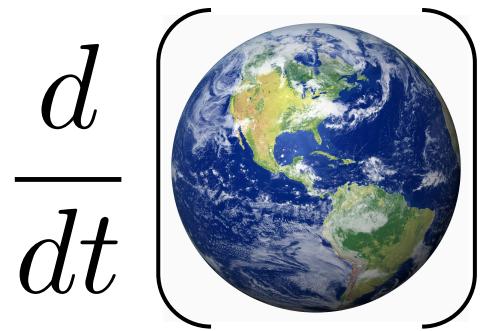
# Why is **Closing the Loop** challenging?



**Implicit and Delayed  
Feedback**



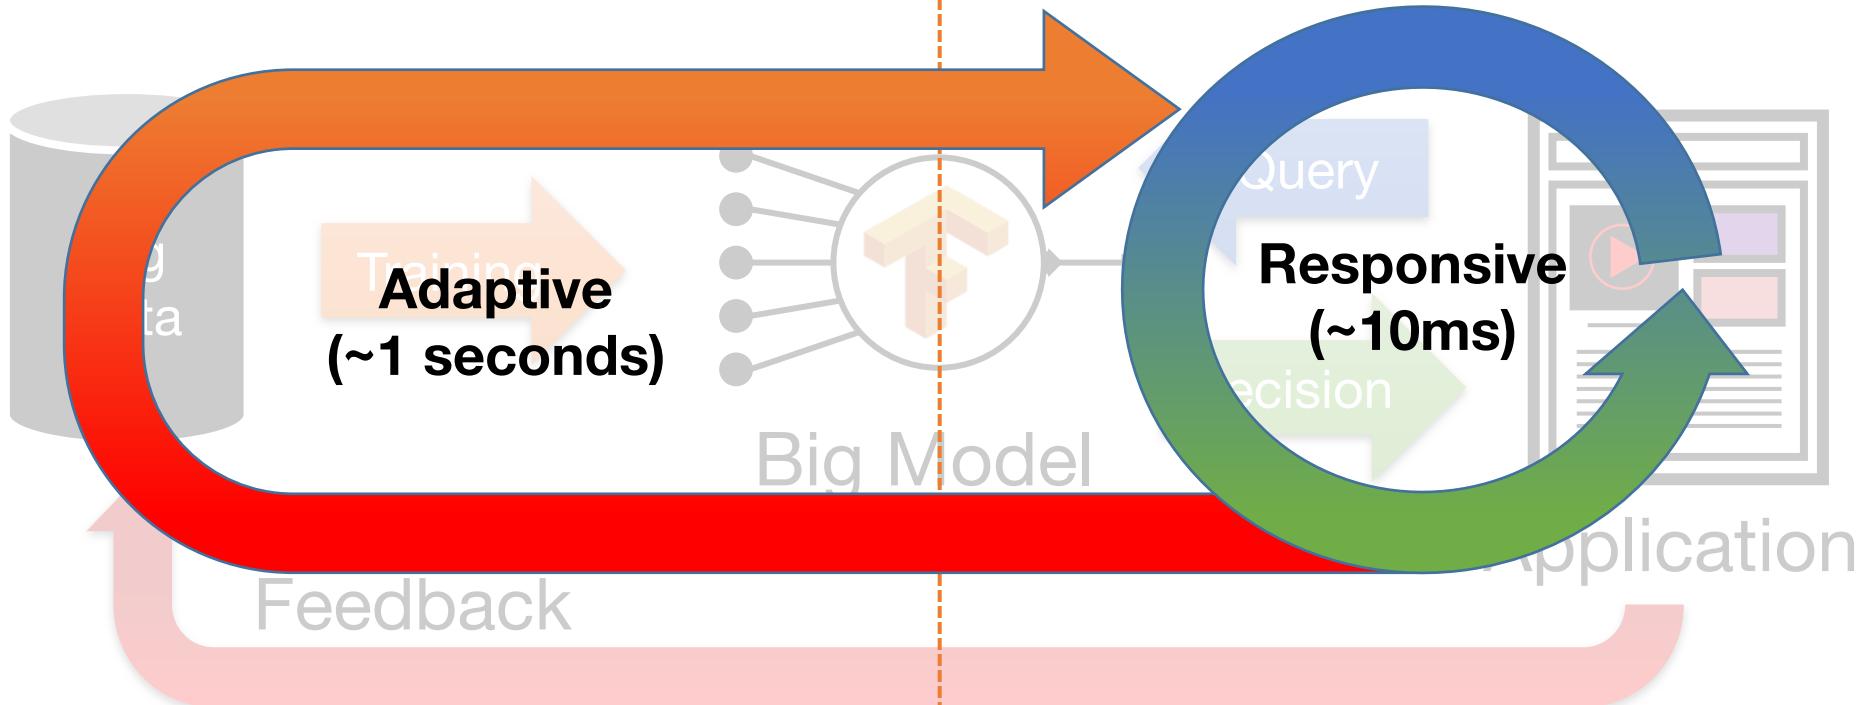
**Self Reinforcing  
Feedback Loops**

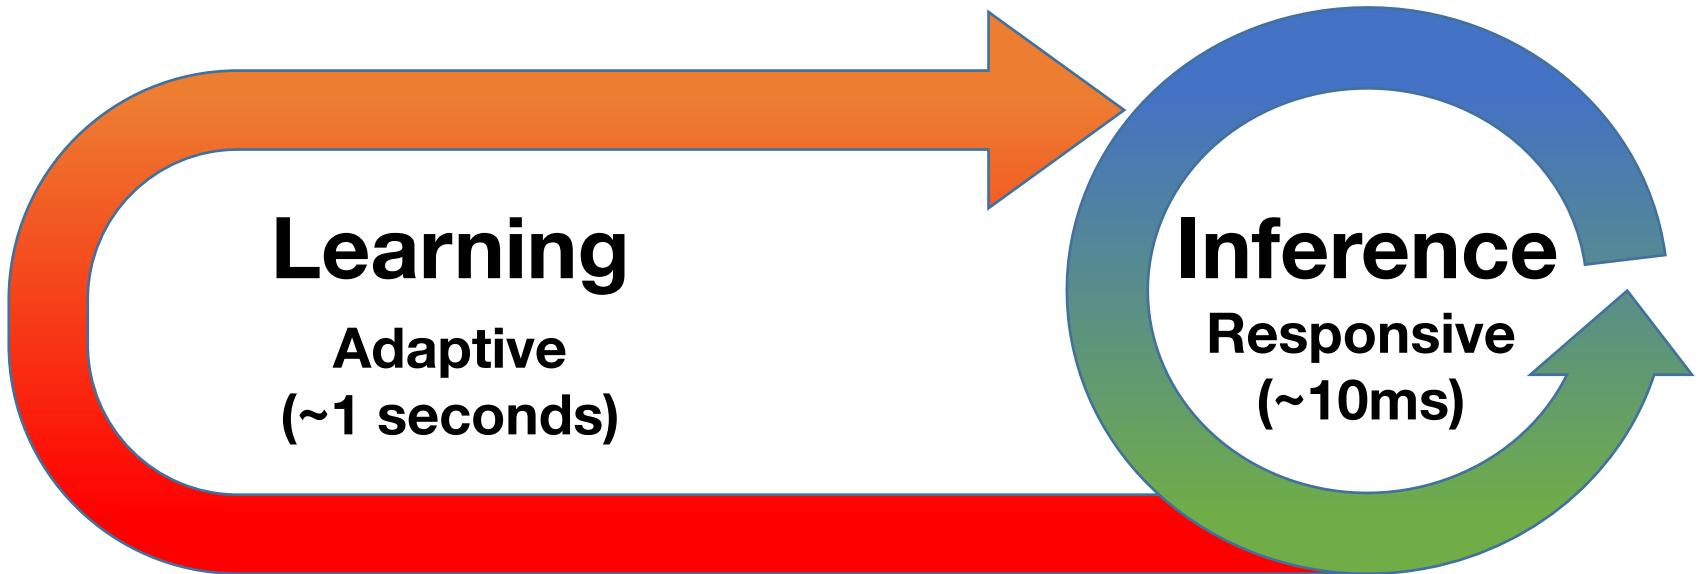


**World Changes  
at varying rates**

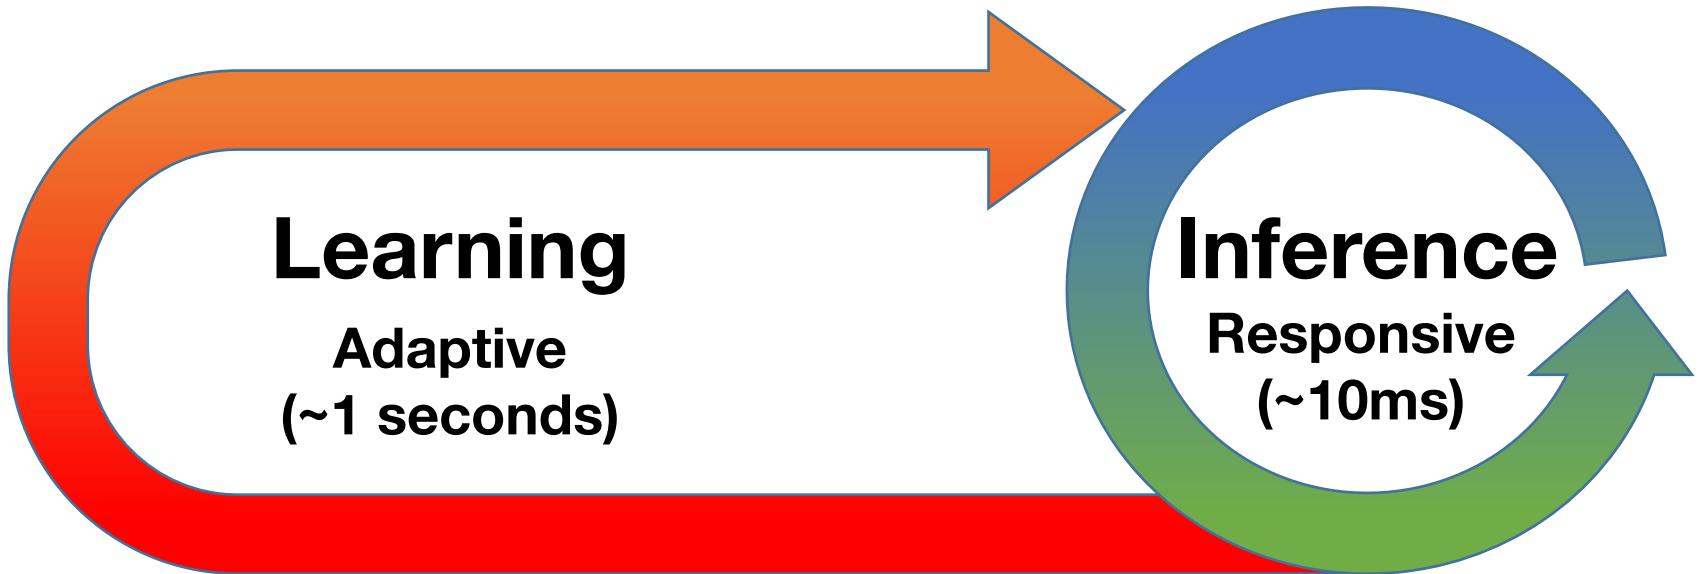
# Learning

# Inference





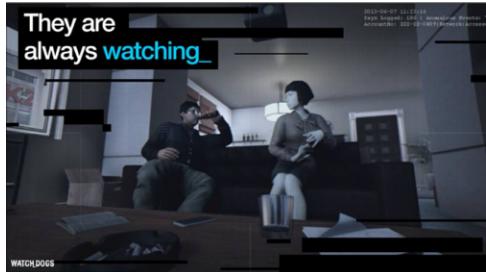
?



# Secure

# Intelligence in Sensitive Contexts

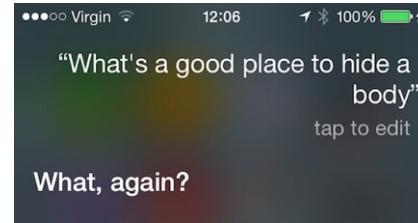
Augmented Reality



Home Monitoring



Voice Technologies



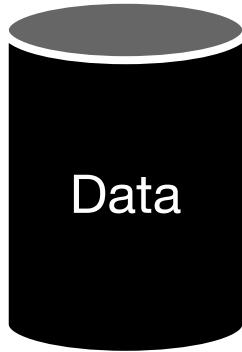
Medical Imaging



Protect the **data**, the **model**, and the **query**

# Protect the **data**, the **model**, and the **query**

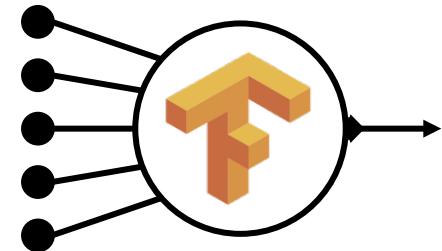
High-Value **Data is Sensitive**



- Medical Info.
- Home video
- Finance

**Models** capture **value** in data

- Core Asset
- Sensitive



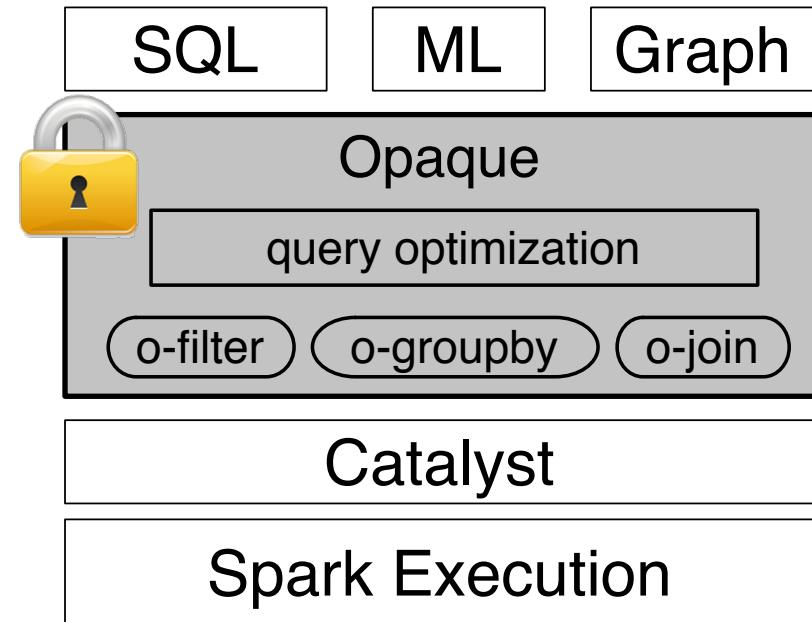
**Queries** can be as sensitive as the data



# Opaque: Analytics on Secure Enclaves

*Exploit hardware support to enable computing on encrypted data*

- **Today:** prototype system running in Apache Spark
  - support SQL queries in untrusted cloud
  - ~50% reduction in perf.
- **Future:** enable prediction serving on enc. queries

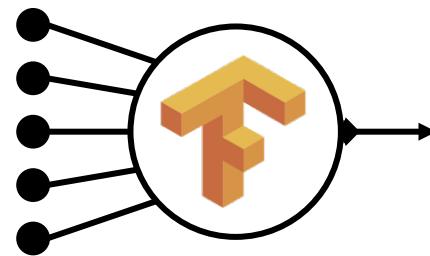
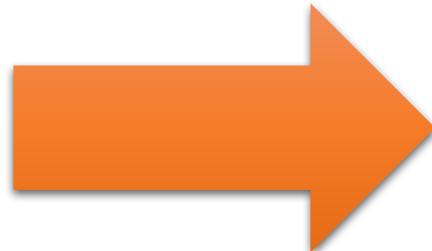
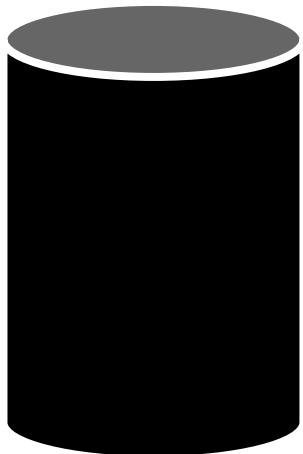


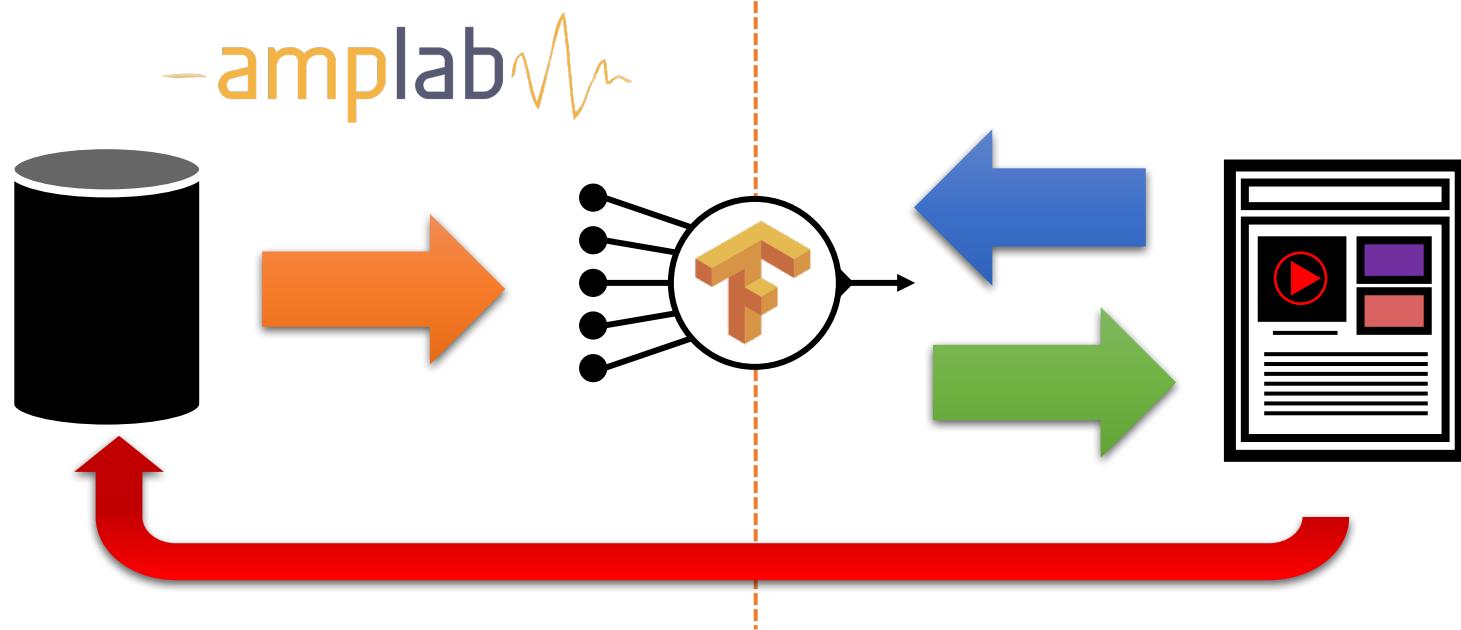
**Adaptive**

**Responsive**

**Secure**

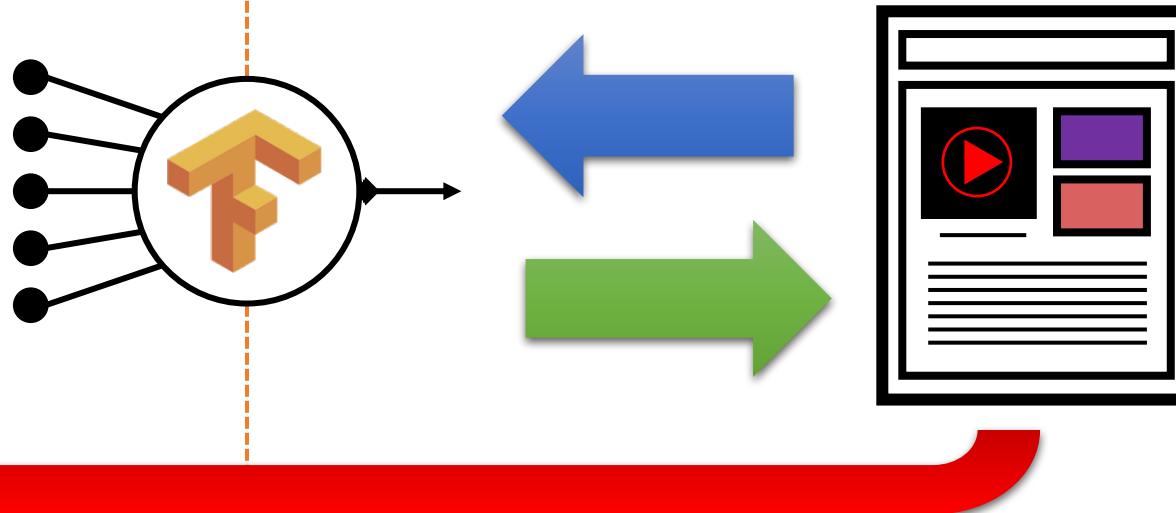
-amplab





# Clipper

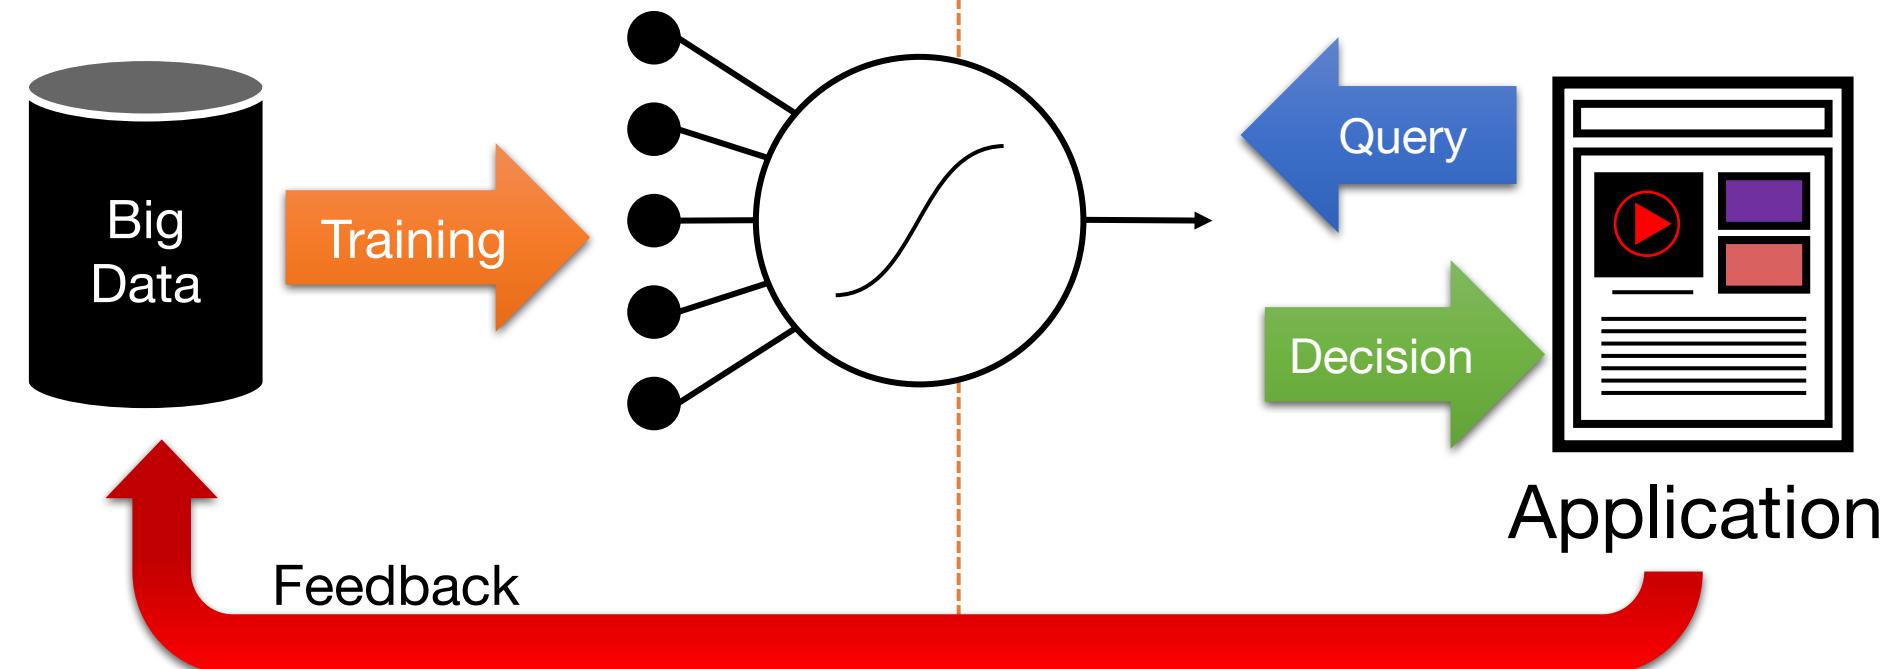
A Low-Latency Online Prediction  
Serving System



**NSDI'17**  
**Daniel Crankshaw**  
Xin Wang  
Giulio Zhou  
Michael J. Franklin  
Joseph E. Gonzalez  
Ion Stoica

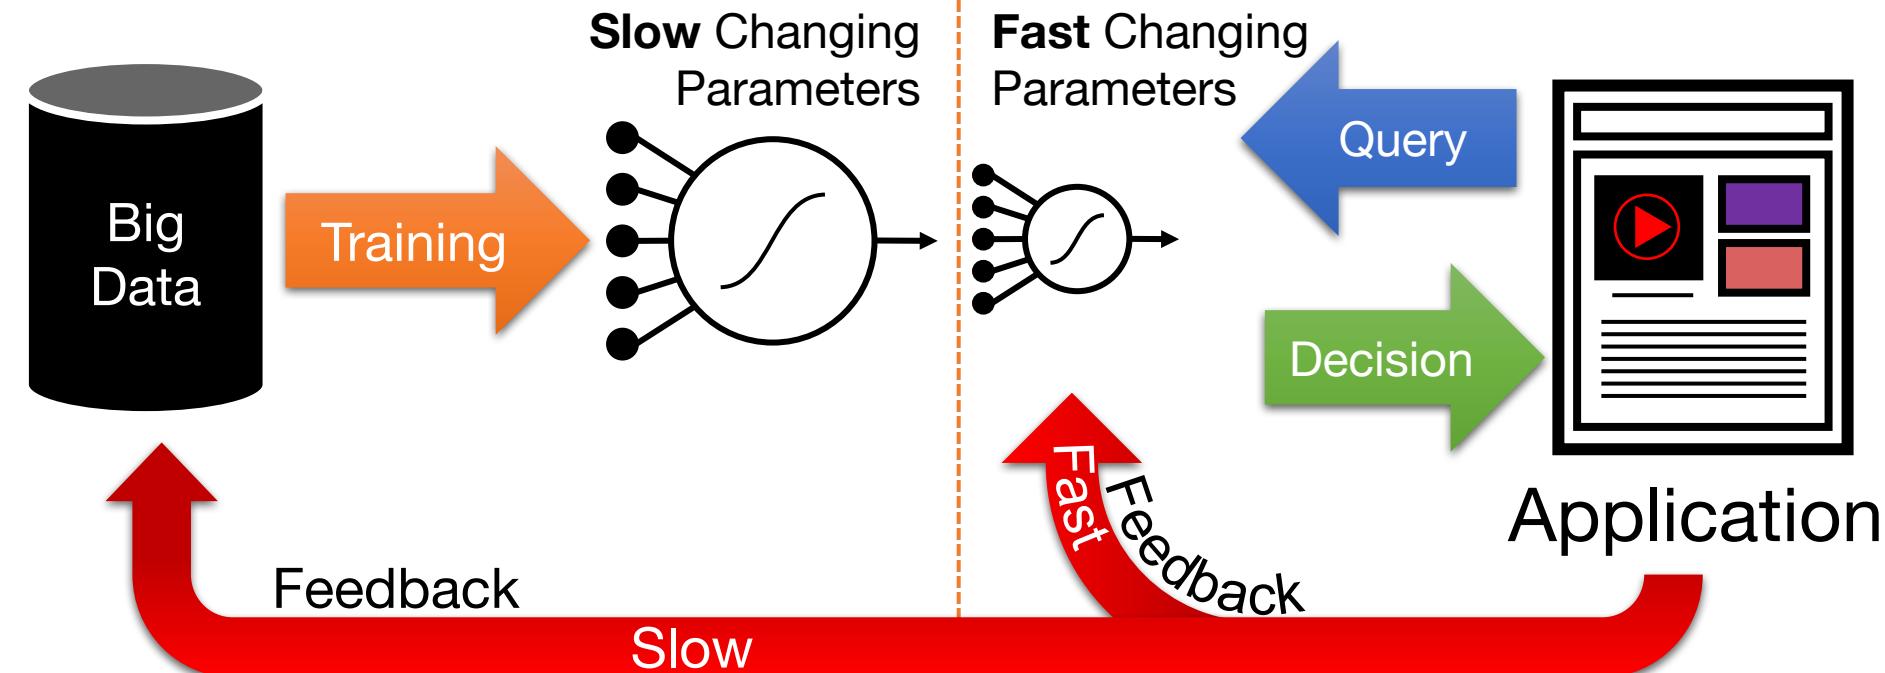
# Learning

# Inference



# Learning

# Inference



# Hybrid Offline + Online Learning

Update “feature” functions **offline** using batch solvers

- Leverage high-throughput systems (Tensor Flow)
- Exploit slow change in population statistics

$$f(x; \theta)^T$$

$$w_u$$

Update the user weights **online**:

- Simple to train + more robust model
- Address rapidly changing user statistics

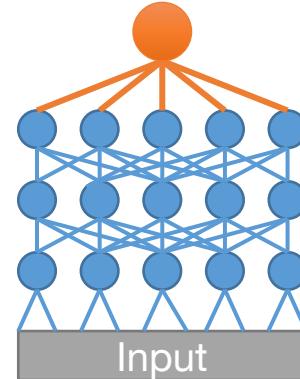
# Common modeling structure

$$f(x; \theta)^T w_u$$

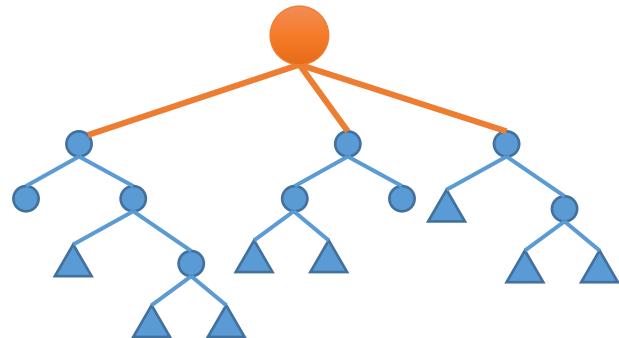
Matrix  
Factorization



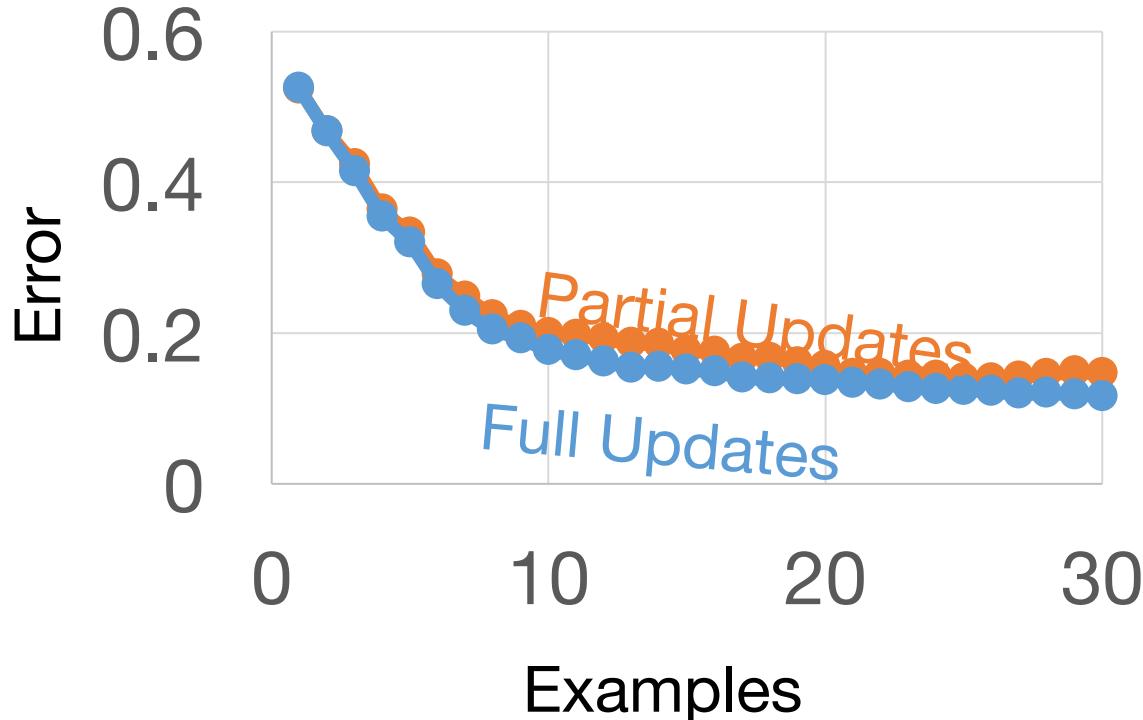
Deep  
Learning



Ensemble  
Methods



# Clipper Online Learning for Recommendations (Simulated News Rec.)

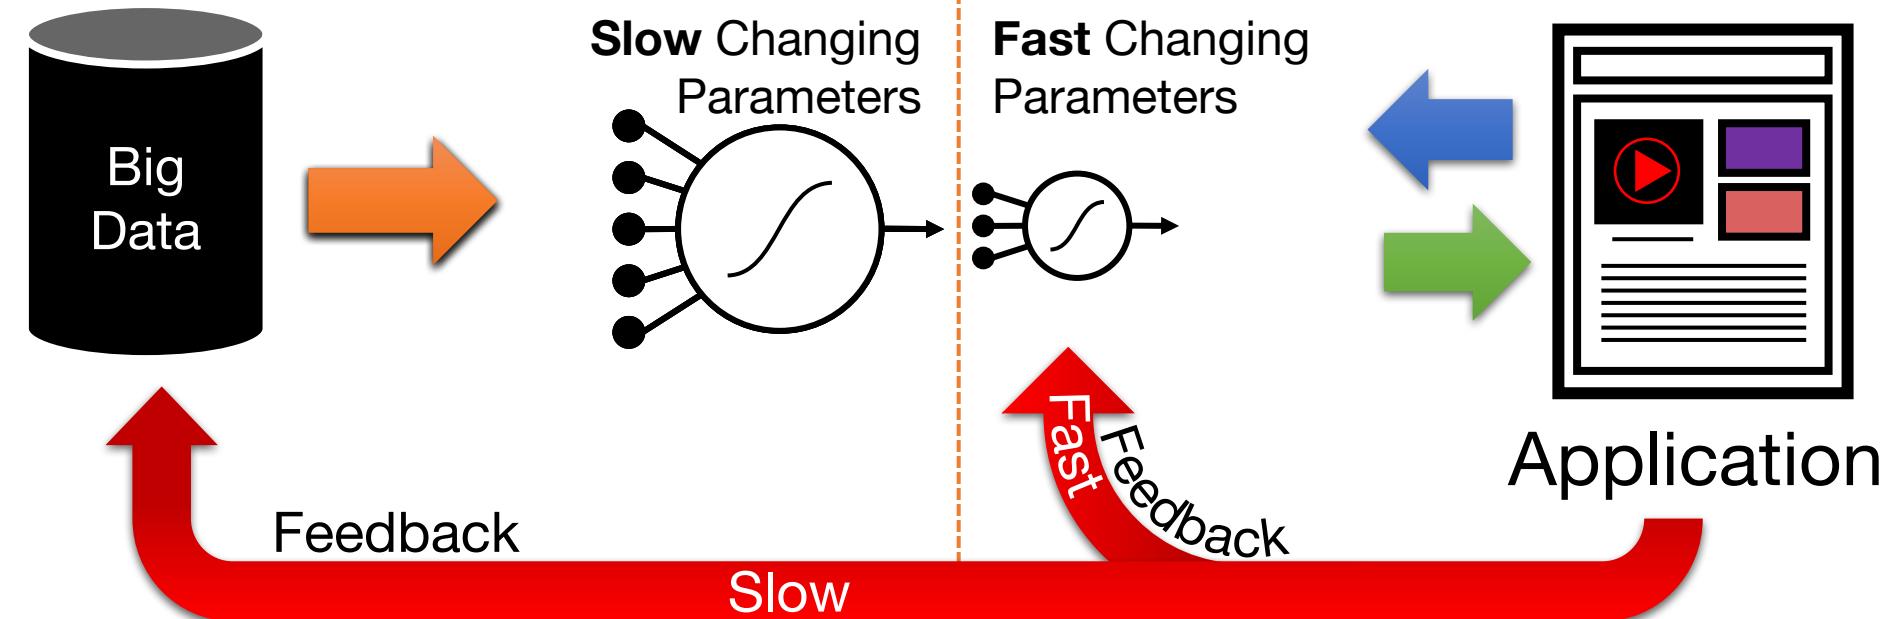


**Partial Updates:** 0.4 ms  
**Retraining:** 7.1 seconds

>4 orders-of-magnitude **faster adaptation**

# Learning

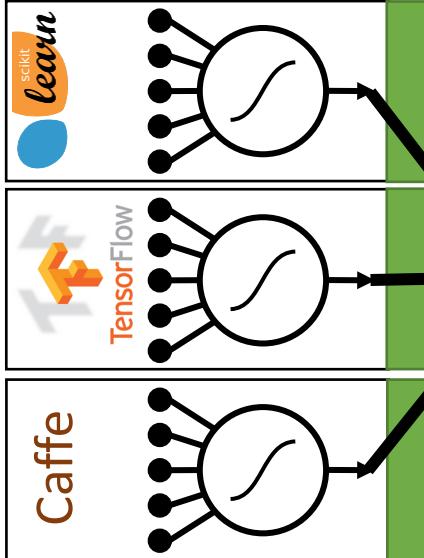
# Inference



# Learning



**Slow Changing  
Parameters**



Slow

# Inference

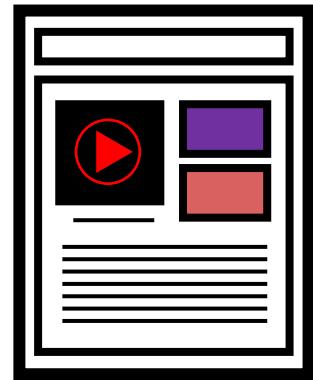
**Clipper**

**Fast Changing  
Parameters**



Fast Feedback

**Application**



# Clipper Serves Predictions across ML Frameworks

Fraud  
Detection



Content  
Rec.



Personal  
Asst.



Robotic  
Control



Machine  
Translation



## Clipper

theano

Dato



KeystoneML



Caffe

TensorFlow

scikit  
learn

dmlc

mxnet



VW

KALDI

# Clipper

## Key Insight:

*The challenges of prediction serving can be addressed between end-user applications and machine learning frameworks*

As a result, Clipper is able to:

- **hide complexity** by
  - providing a common interface to applications
- **bound latency** and **maximize throughput**
  - through caching, adaptive batching, model replication
- enable *robust online learning* and **personalization**
  - through model selection and ensemble algorithms

**without modifying** machine learning **frameworks** or front-end **applications**

# Clipper Architecture

Fraud  
Detection



Content  
Rec.



Personal  
Asst.



Robotic  
Control



Machine  
Translation



## Clipper

theano

KeystoneML

Dato



Create  
Caffe TensorFlow



scikit  
learn

dmlc  
mxnet



VW  
KALDI

# Clipper Architecture

Applications



Predict

RPC/REST Interface

Observe

Clipper

theano

Dato



Keystone

ML



Caffe

TensorFlow

scikit  
learn

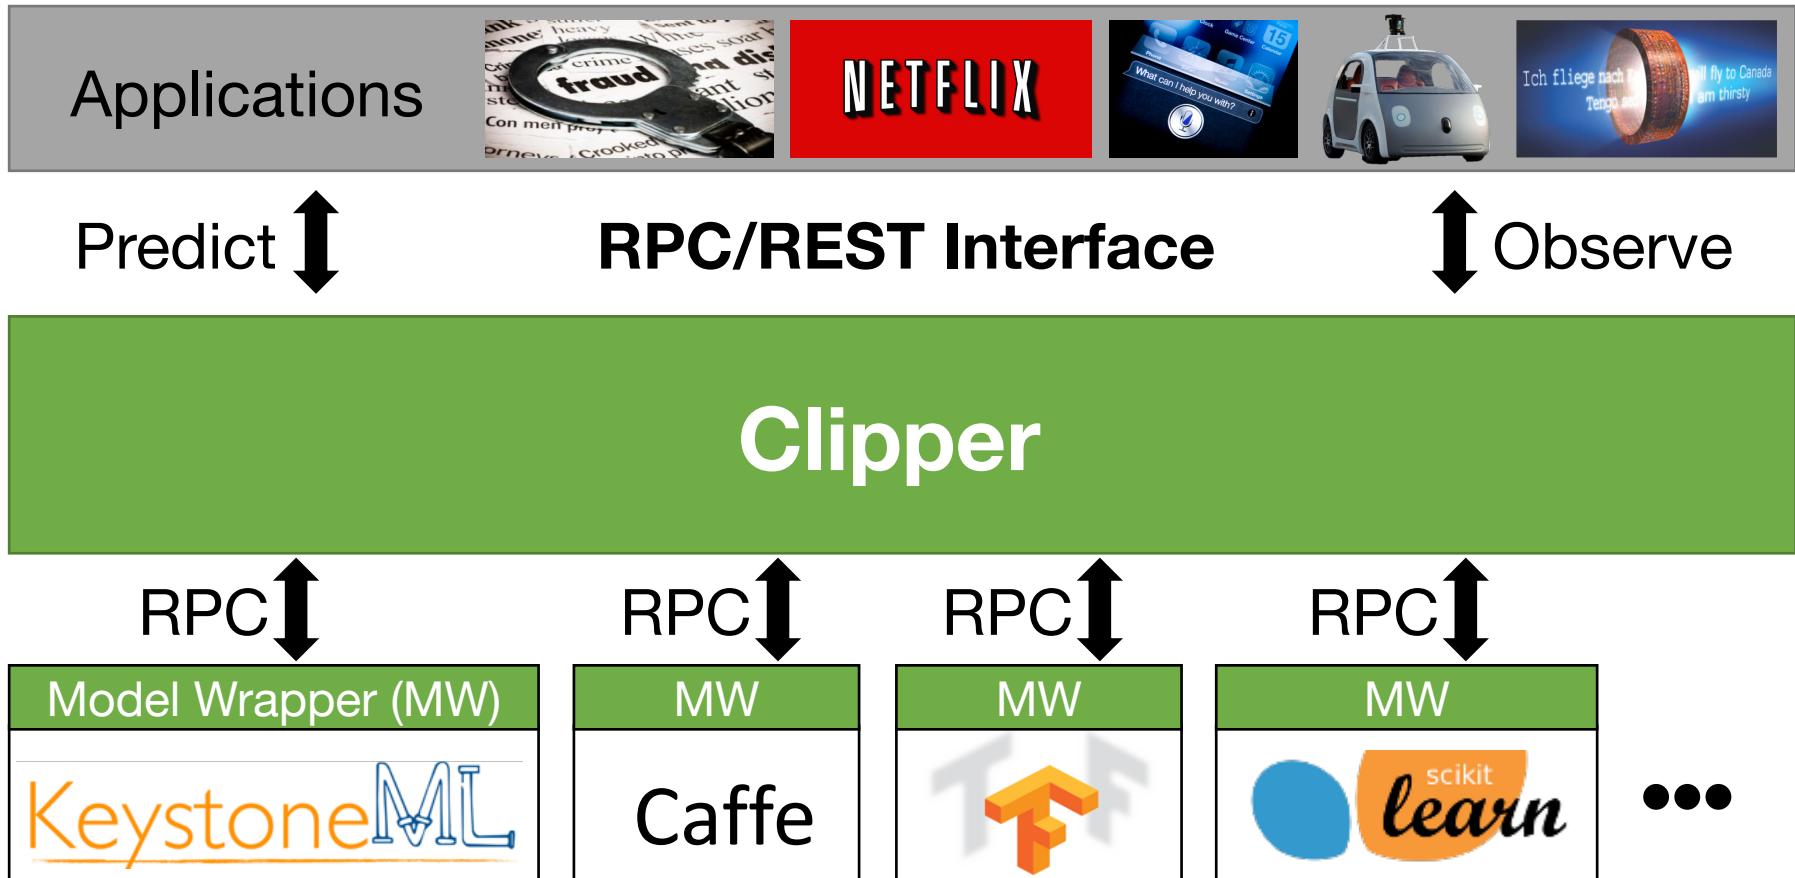
dmlc

mxnet

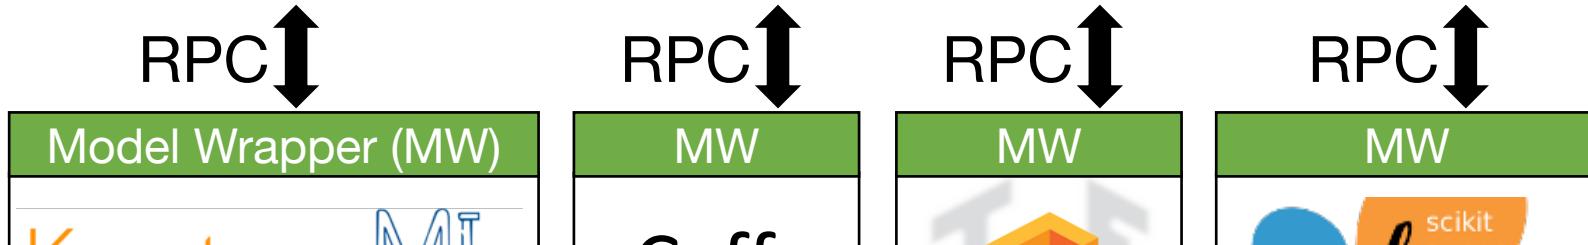
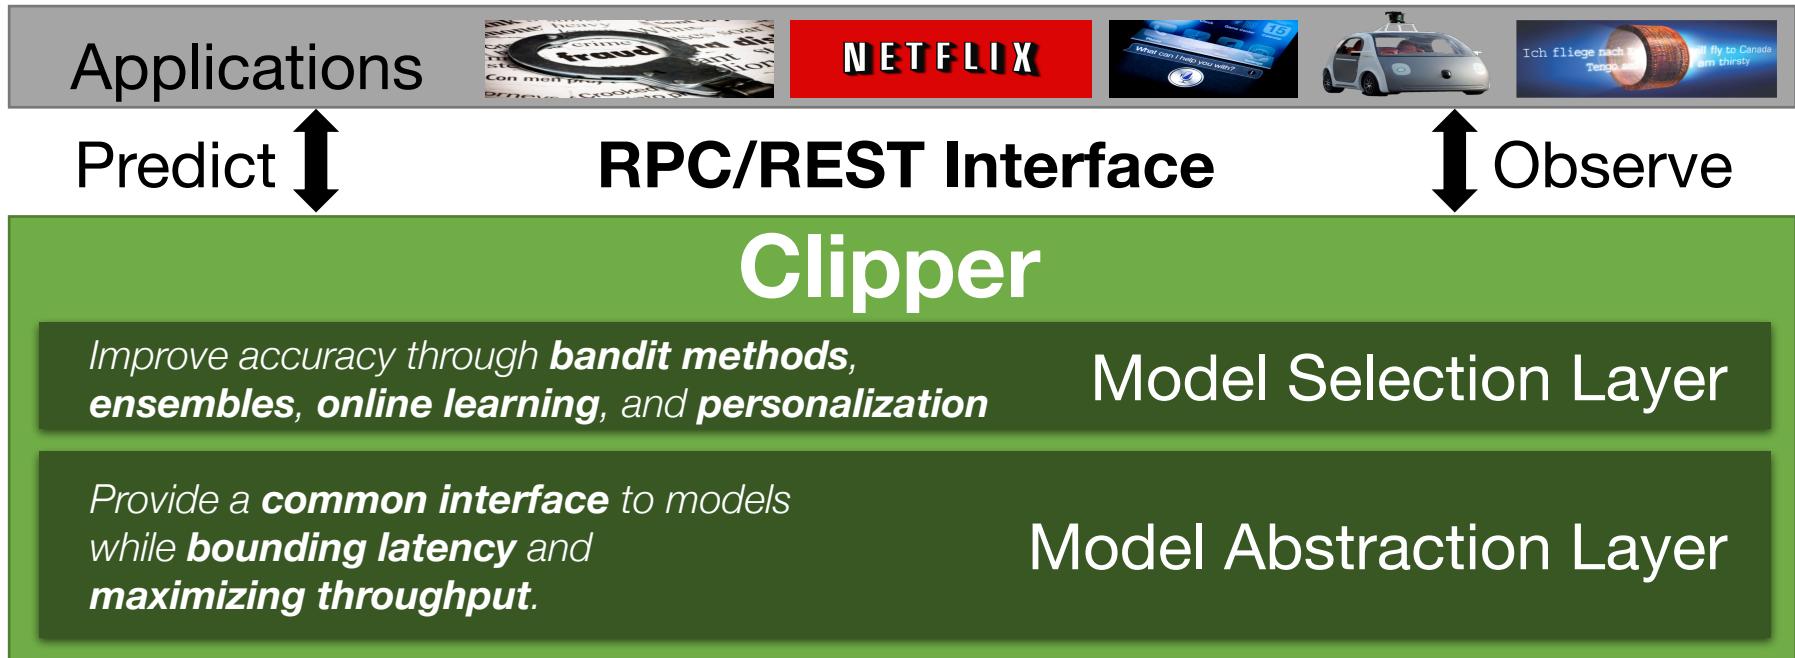


KALDI

# Clipper Architecture

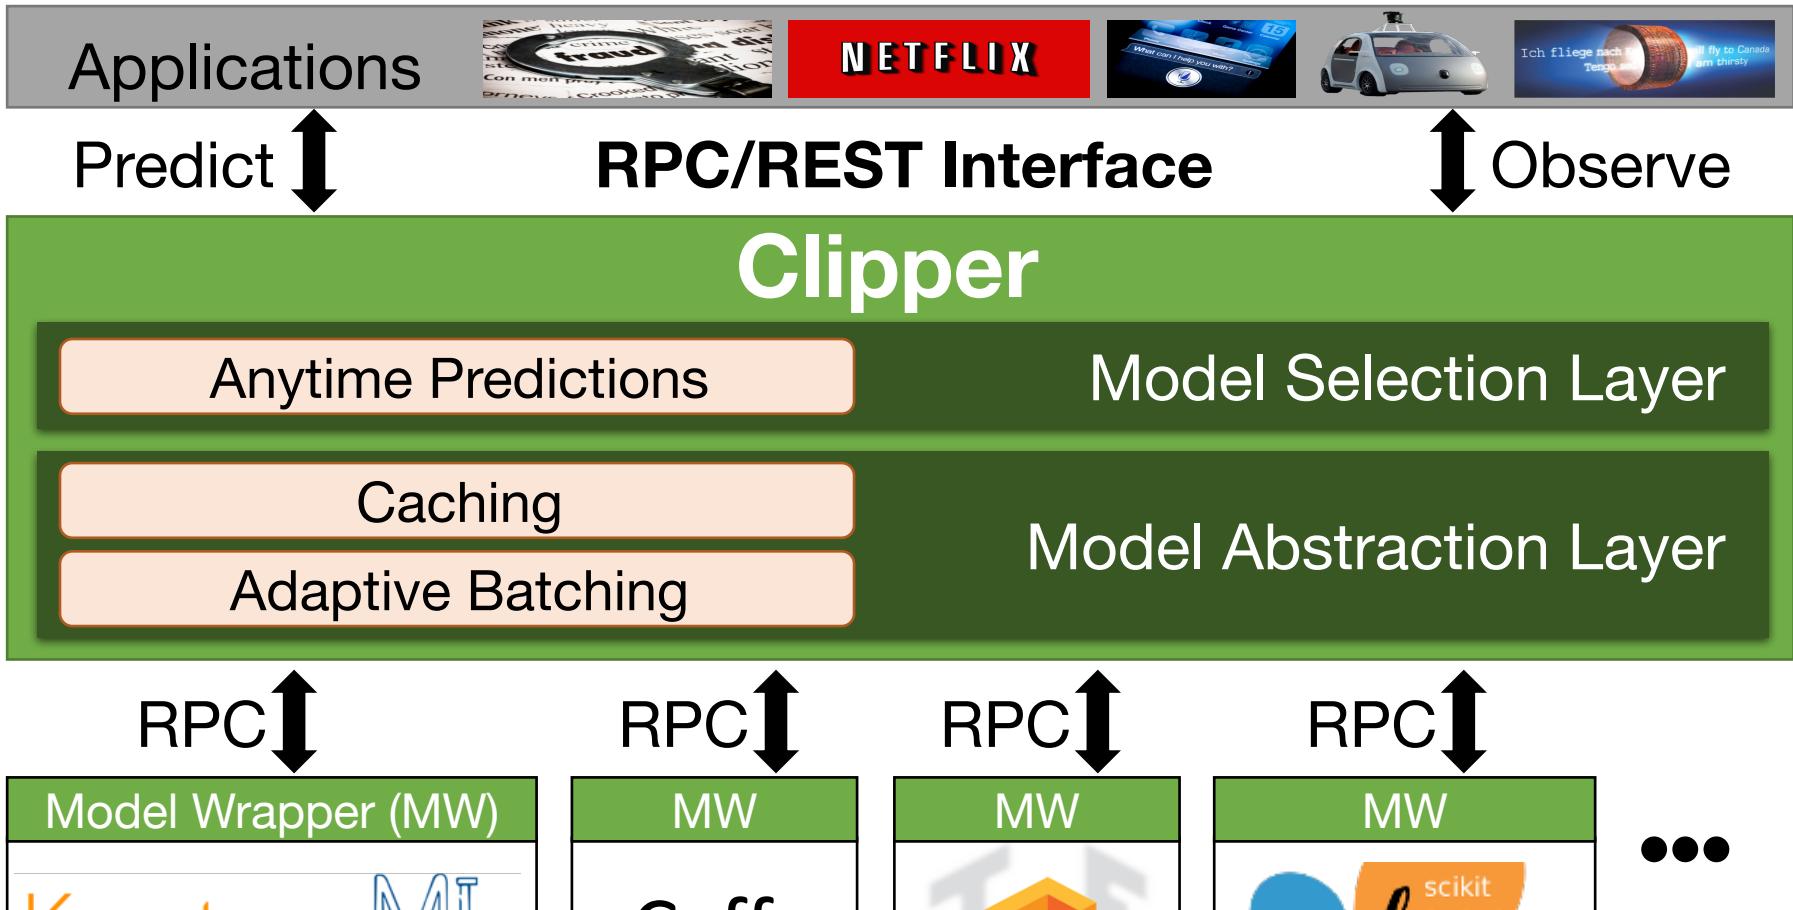


# Clipper Architecture



...

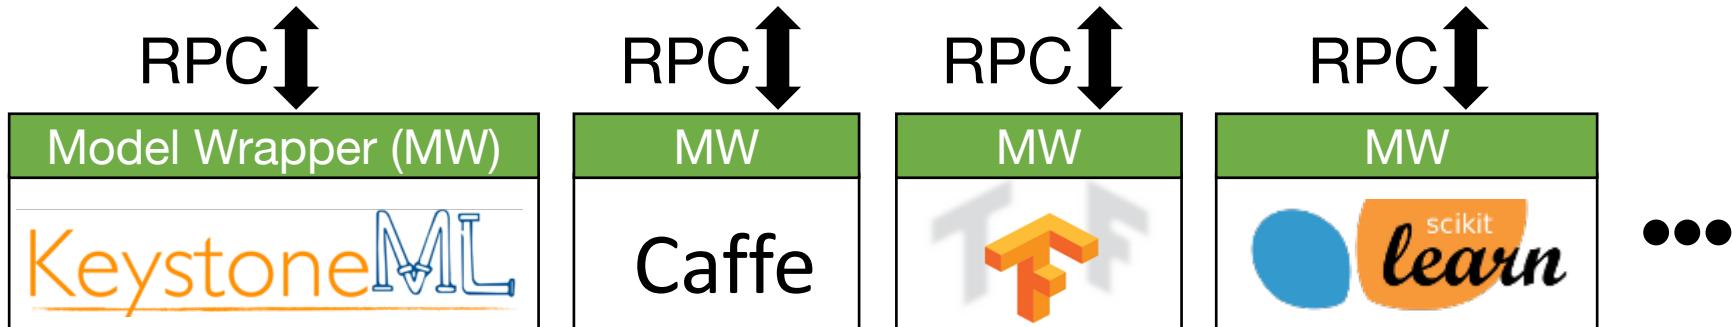
# Clipper Architecture



Caching

Adaptive Batching

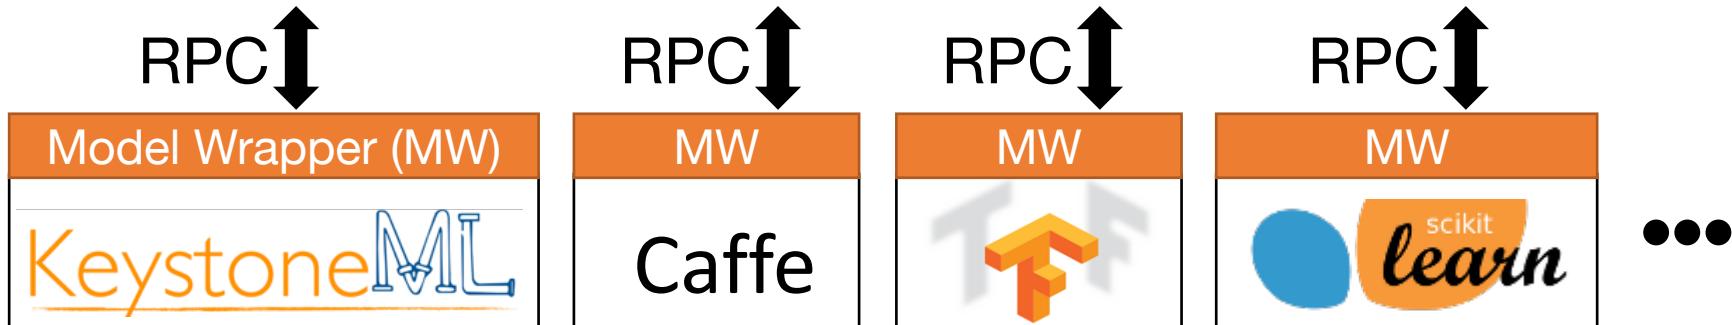
## Model Abstraction Layer



Caching

Adaptive Batching

## Model Abstraction Layer



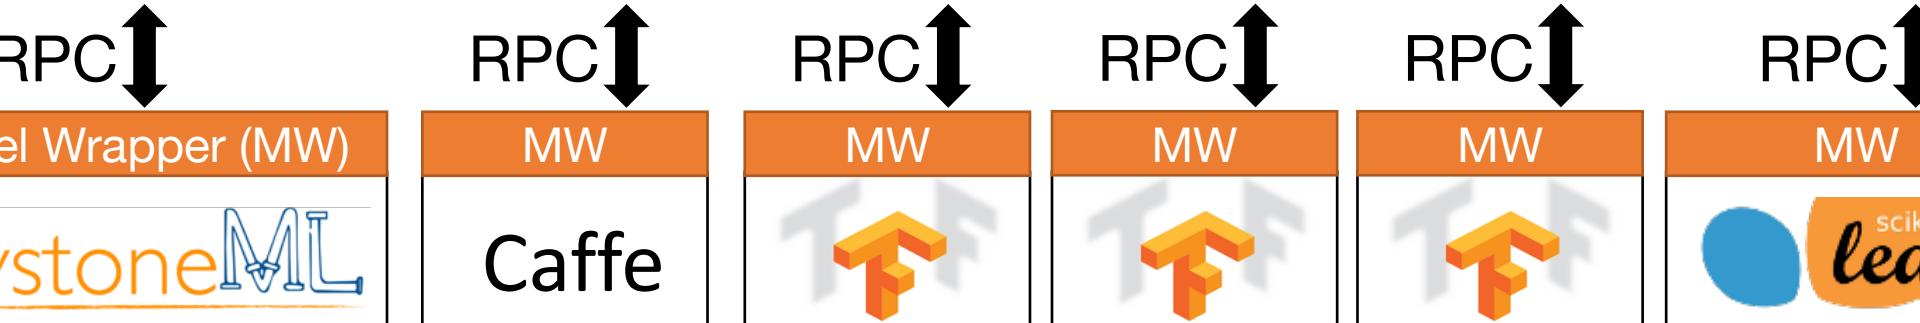
Provide a **common interface** to models while  
**bounding latency** and **maximizing throughput**.

- Models run in separate processes as Docker containers
  - Resource isolation

Caching

Adaptive Batching

Model Abstraction Layer



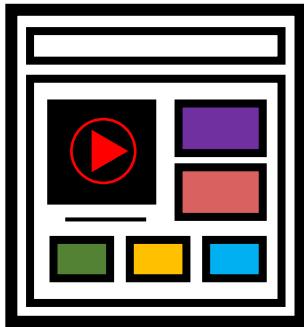
Provide a **common interface** to models while  
**bounding latency** and **maximizing throughput**.

- Models run in separate processes as Docker containers
  - Resource isolation
- Scaling under heavy load

**Problem:** frameworks optimized for **batch processing** not **latency**

# *Adaptive Batching* to Improve Throughput

- Why batching helps:



A single page load may generate many queries

Hardware Acceleration



Helps amortize system overhead

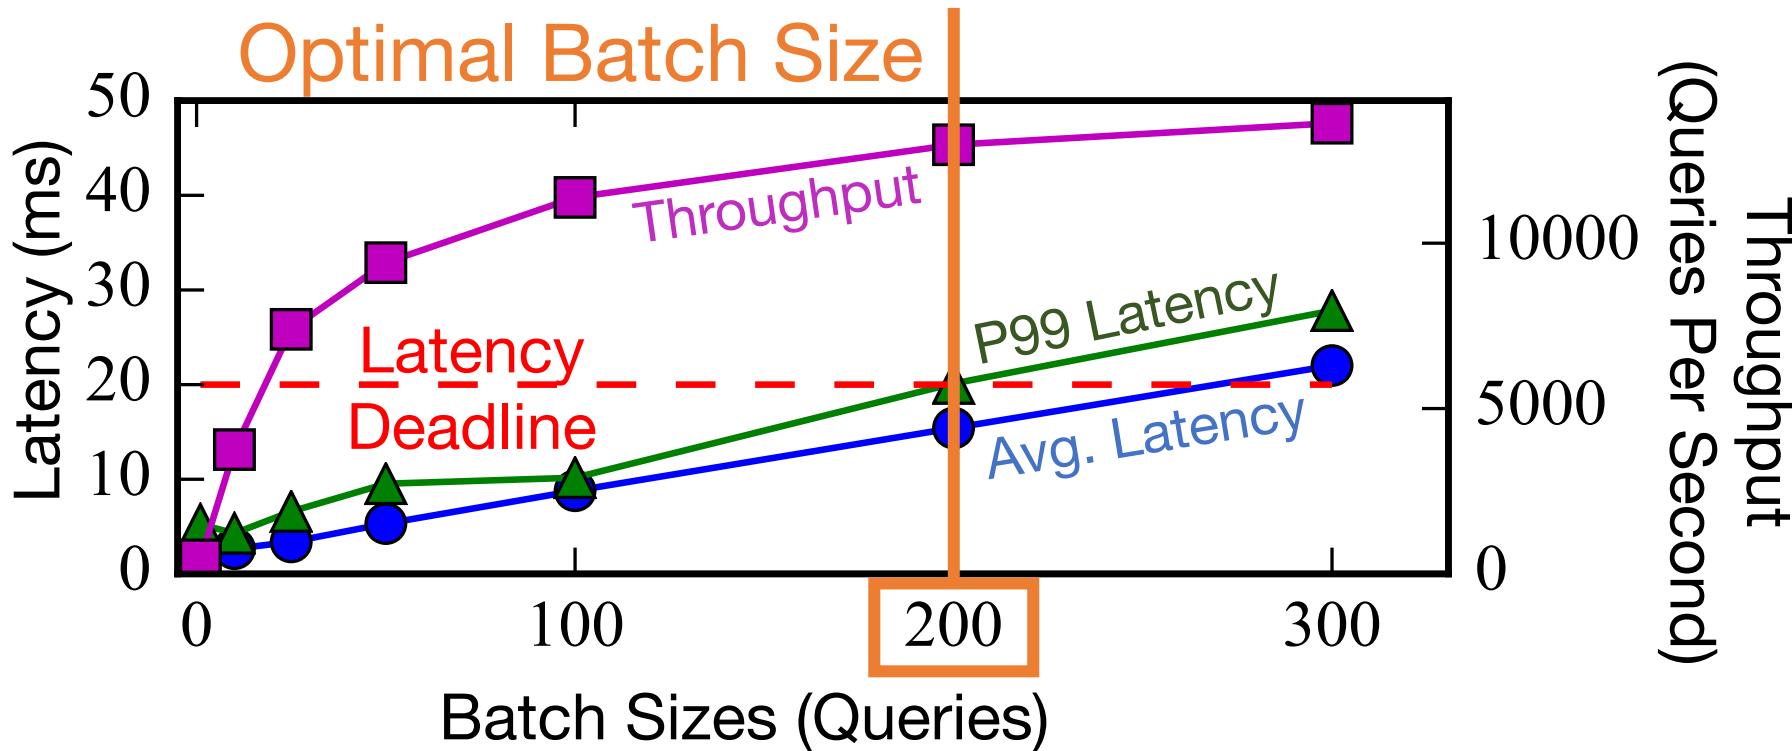
- Optimal batch depends on:
  - hardware configuration
  - model and framework
  - system load

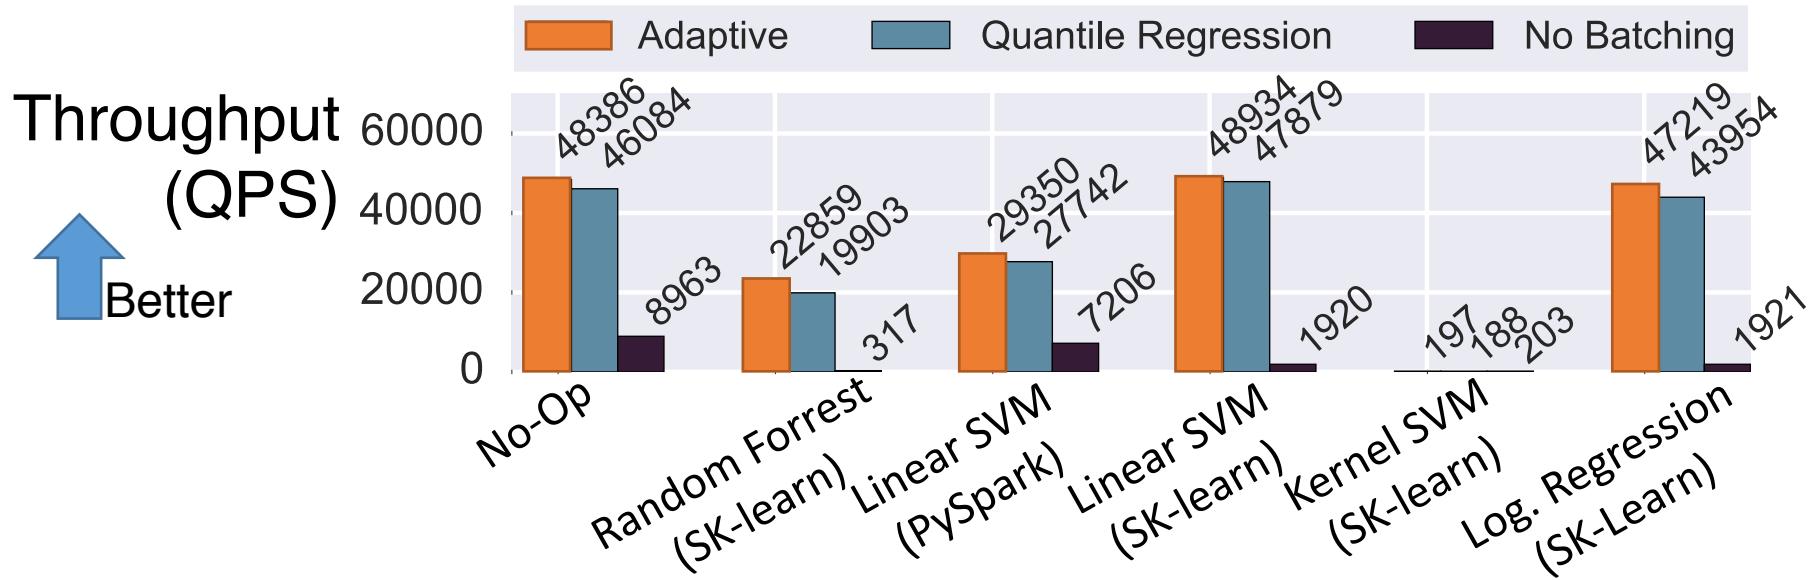
## Clipper Solution:

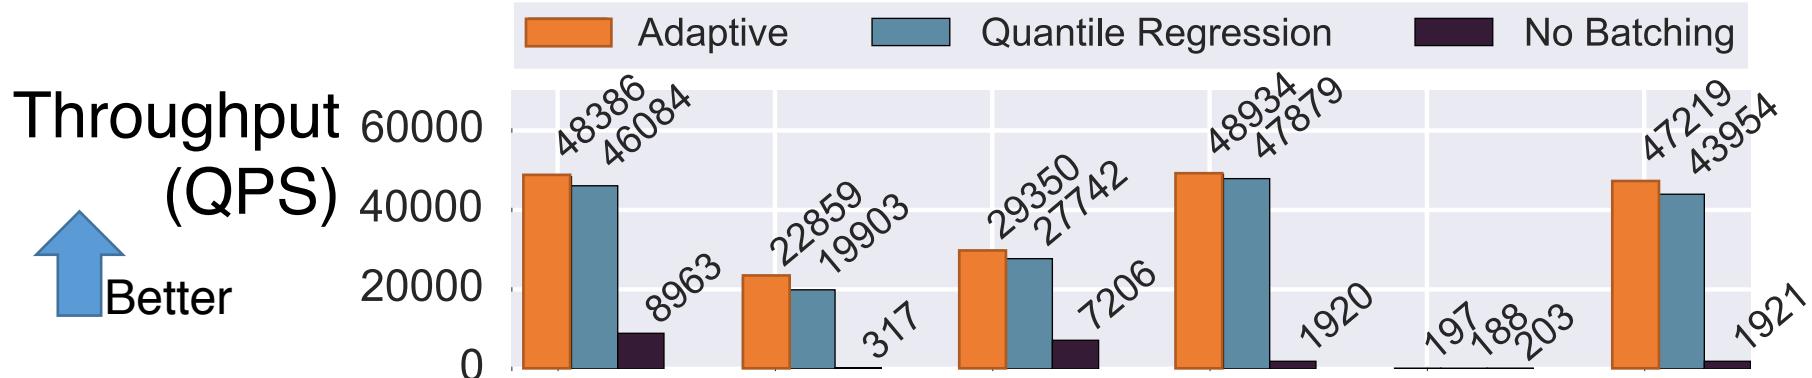
be as **slow** as **allowed**...

- Application specifies latency objective
- Clipper uses TCP-like tuning algorithm to **increase latency** up to the objective

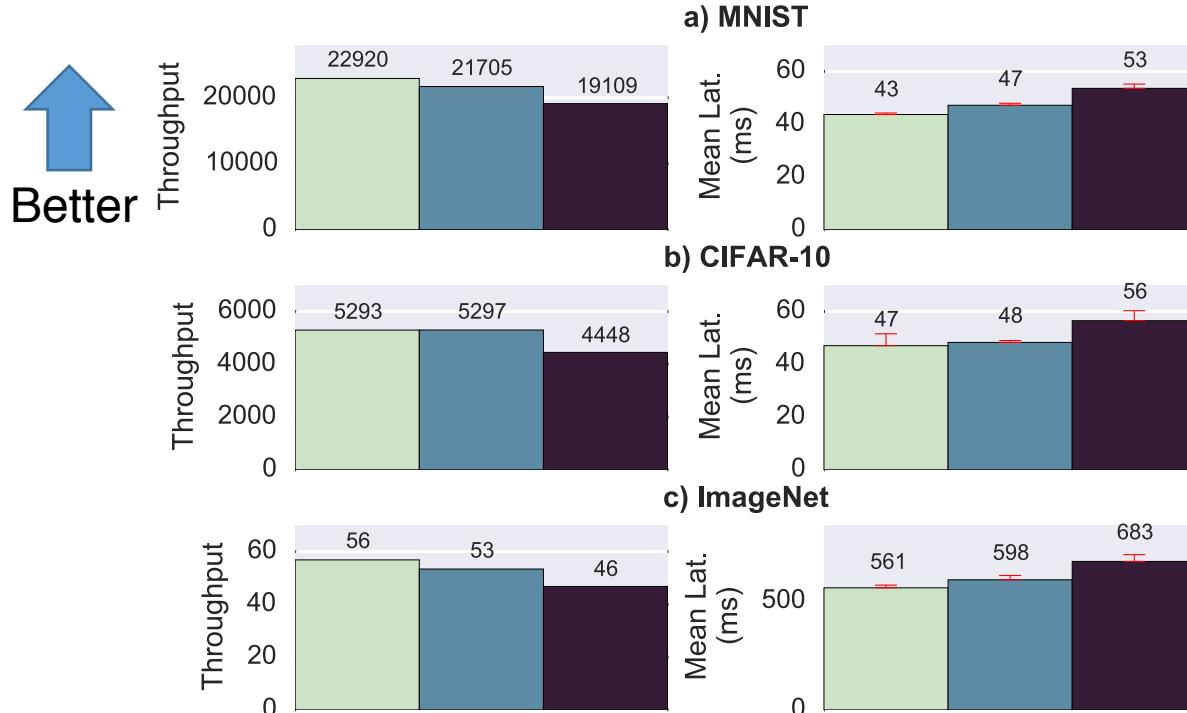
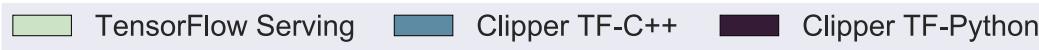
# Tensor Flow Conv. Net (GPU)







# Overhead of modularity?

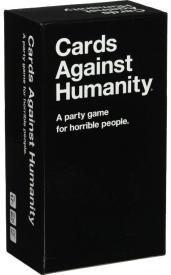


Better  
↓  
40000 is  
Good Enough

*The decoupled Clipper architecture  
can be as fast as the  
in-process approach adopted by  
TensorFlow-Serving*

# Approximate Caching to Reduce Latency

- Opportunity for caching



Popular items may be evaluated frequently

- Need for **approximation**

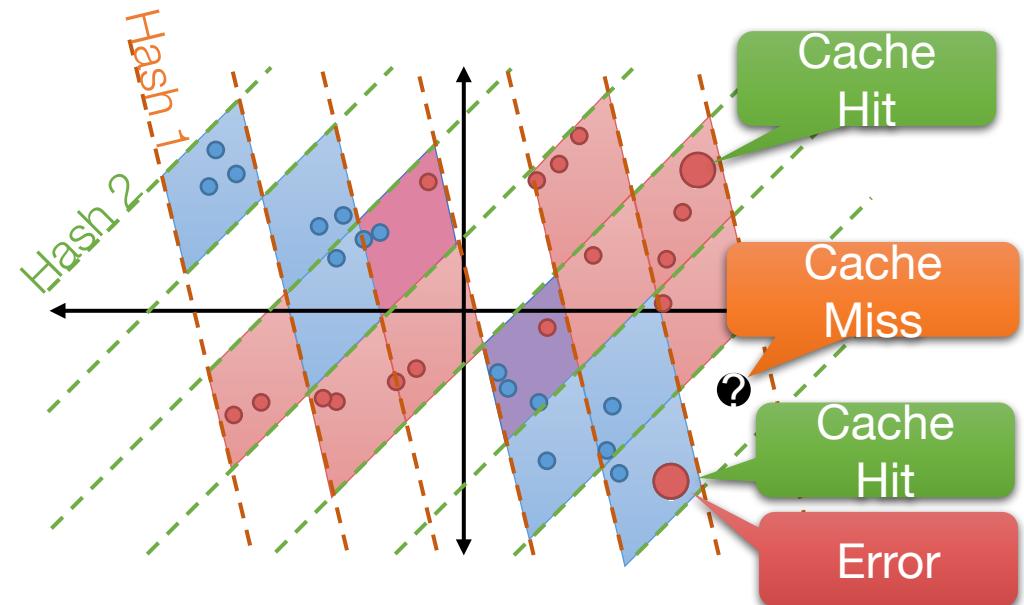


Bag-of-Words Model

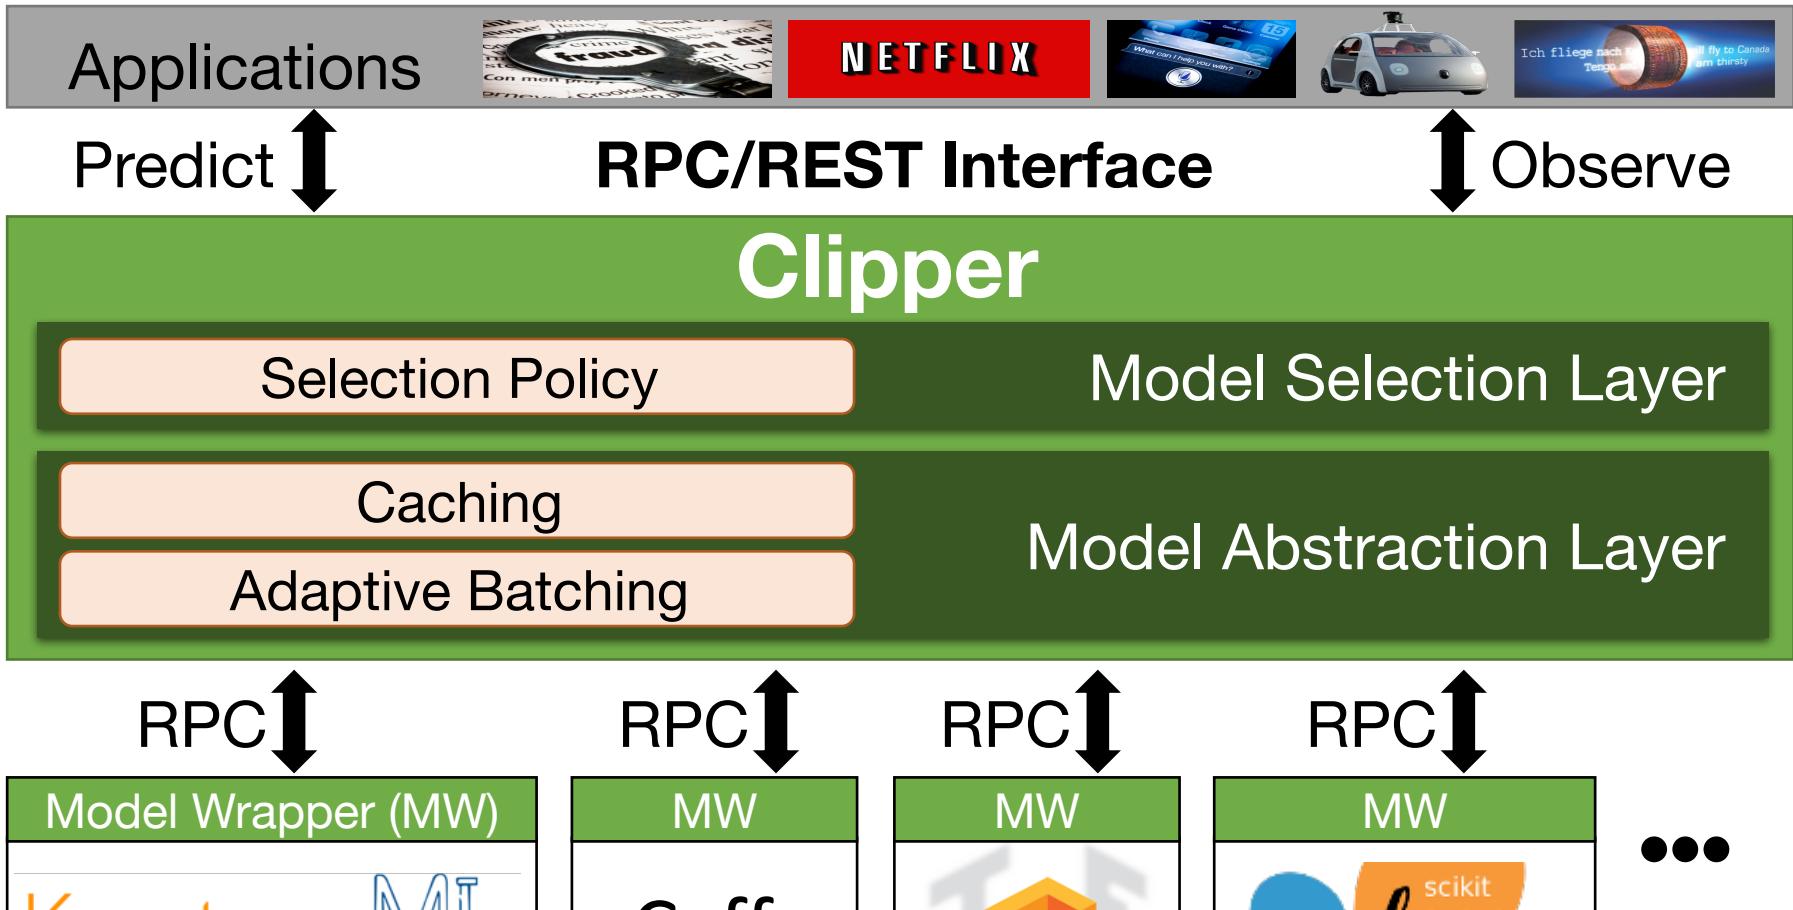
High Dimensional and continuous valued queries have low cache hit rate.

## Clipper Solution: Approximate Caching

apply *locality sensitive hash functions*



# Clipper Architecture



## Goal:

Maximize **accuracy** through **bandits**, **ensembles**, **online learning**, and **personalization**

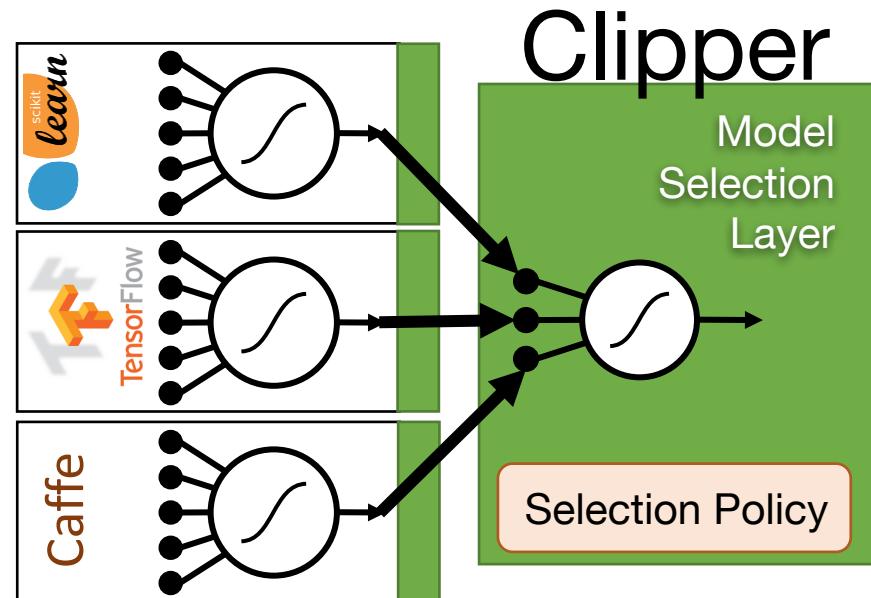
Incorporate feedback in real-time to achieve:

- **robust predictions** by adaptively combining predictions from multiple models and frameworks
- **online learning** and **personalization** by selecting and personalizing **predictions** in response to feedback

# Model Selection Policy

Improves prediction **accuracy** by:

- **Combining predictions** from multiple frameworks
  - Ensemble methods
- Incorporate real-time **feedback**
  - Personalized ensembles
  - Bandit algorithms
- Estimates **confidence** of predictions
  - Agreement between models



# Ensemble Prediction Accuracy (ImageNet)

System	Model	Error Rate	#Errors
Caffe	VGG	13.05%	6525
Caffe	LeNet	11.52%	5760
Caffe	ResNet	9.02%	4512
TensorFlow	Inception v3	6.18%	3088

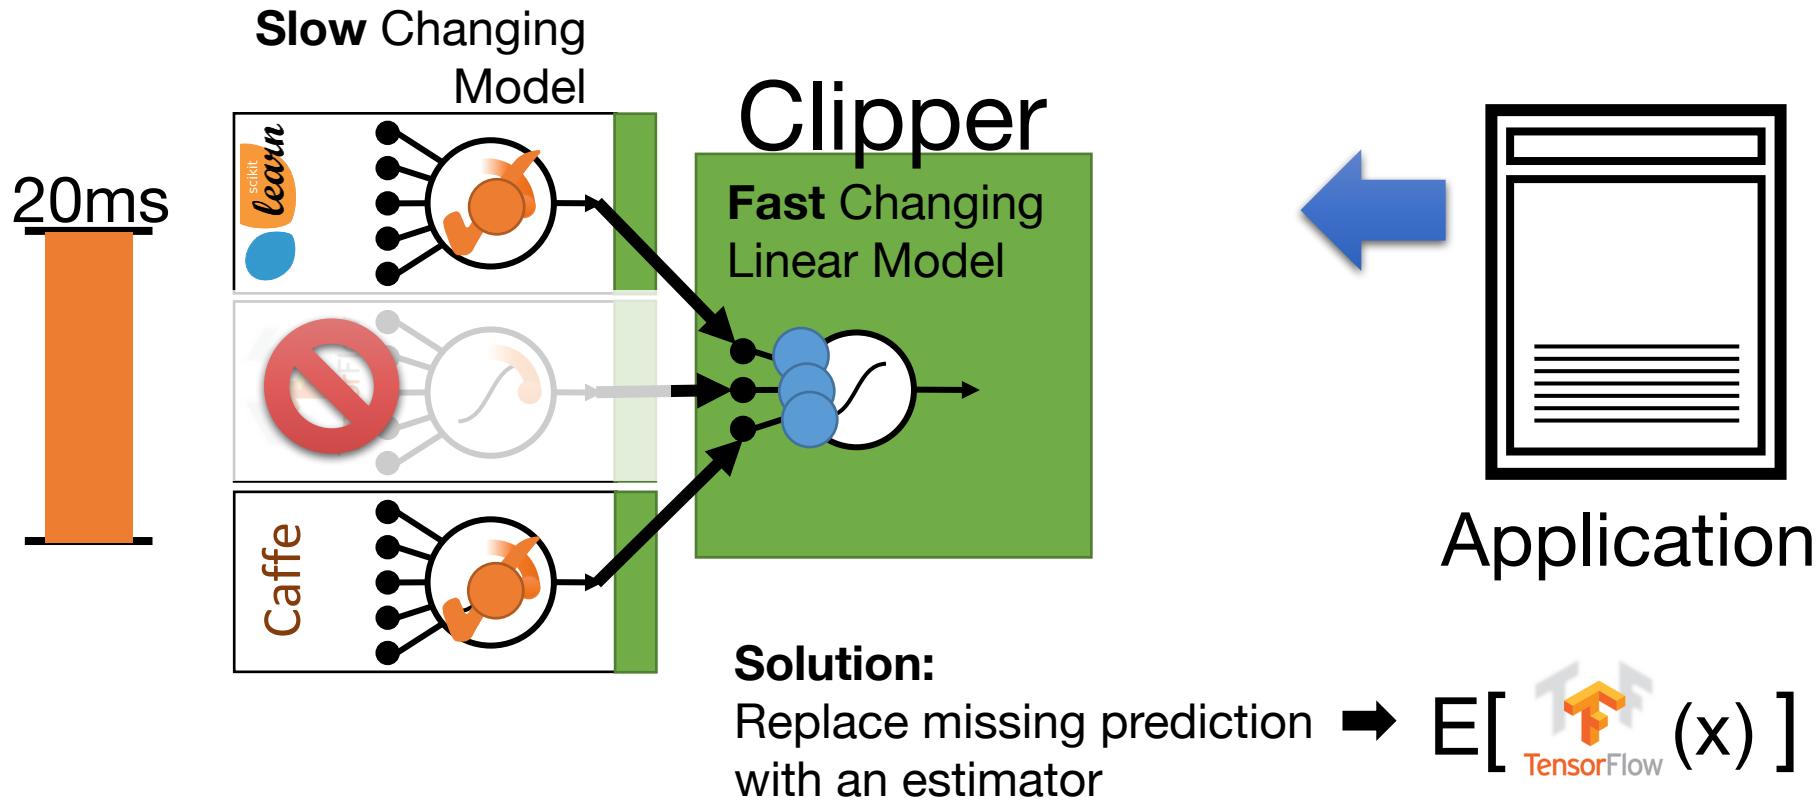
sequence of pre-trained models

# Ensemble Prediction Accuracy (ImageNet)

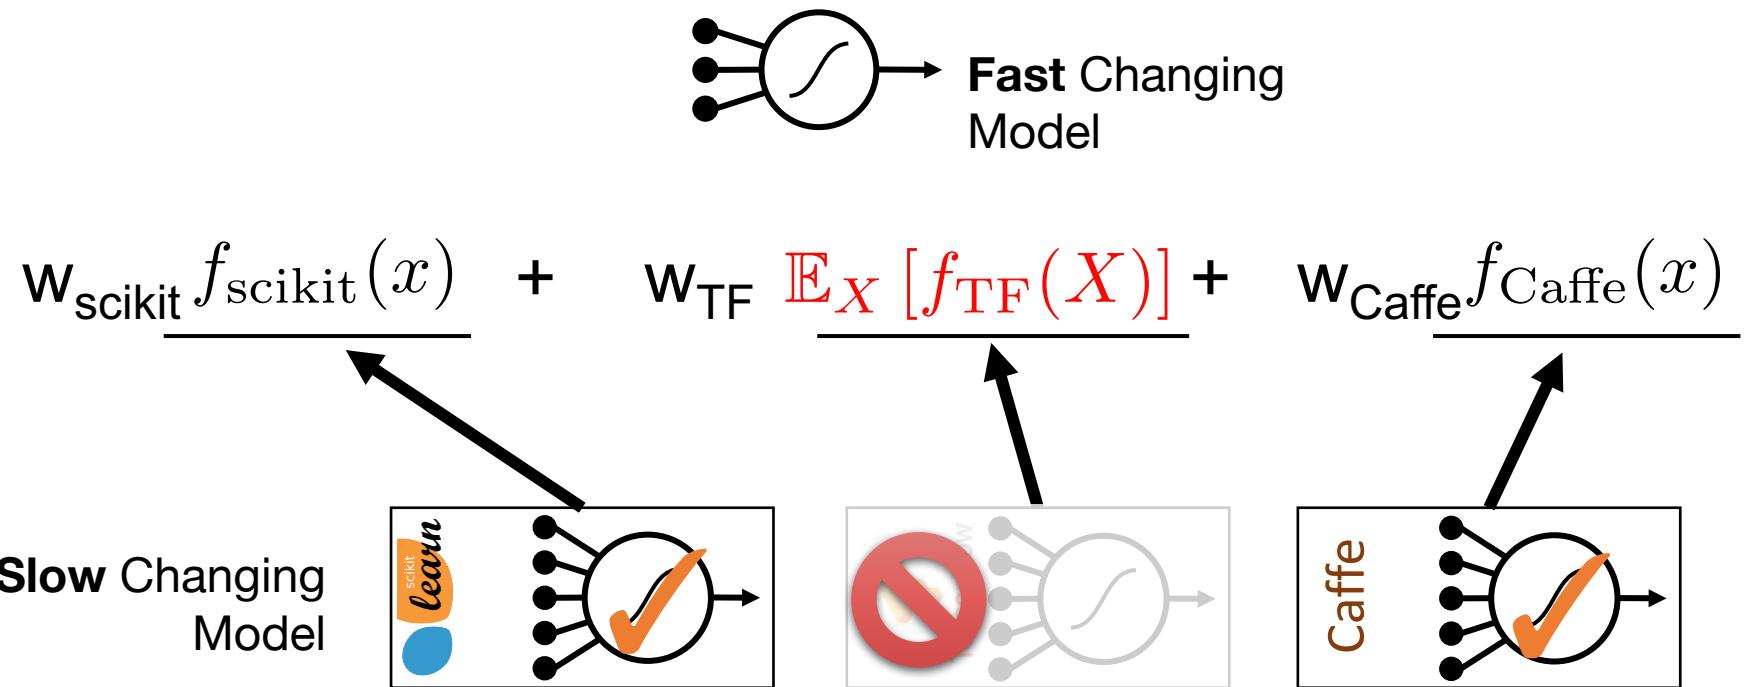
System	Model	Accuracy (%)	Errors
Caffe	ResNet	9.02%	6525
Caffe	Inception v3	6.18%	5760
Caffe	ResNet	9.02%	4512
TensorFlow	Inception v3	6.18%	3088
Clipper	<b>Ensemble</b>	<b>5.86%</b>	<b>2930</b>

5.2% relative improvement  
in prediction accuracy!

# Ensemble Methods Create Stragglers

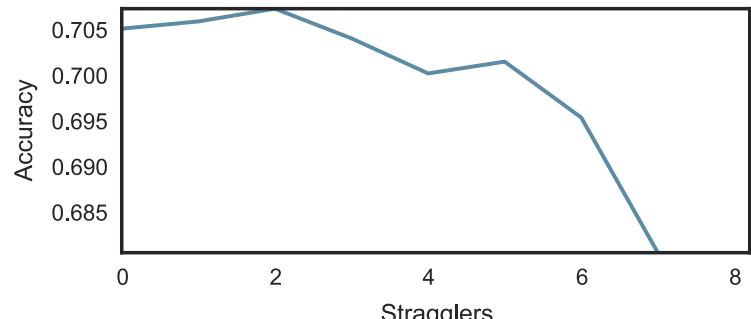
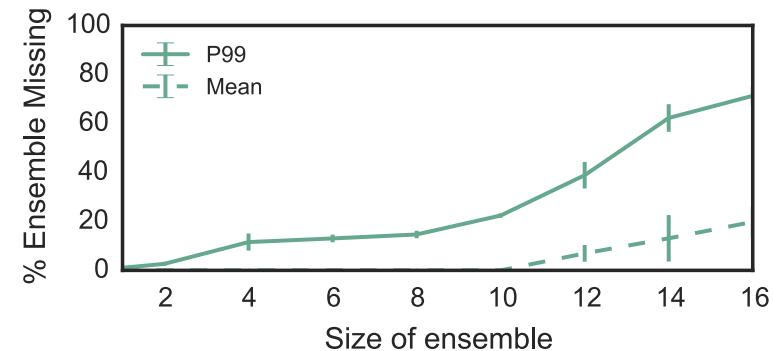
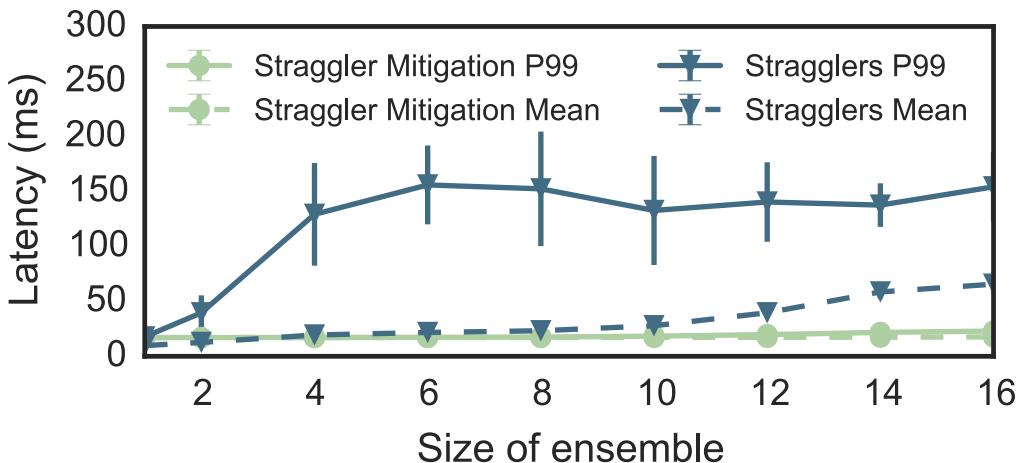


# Anytime Predictions

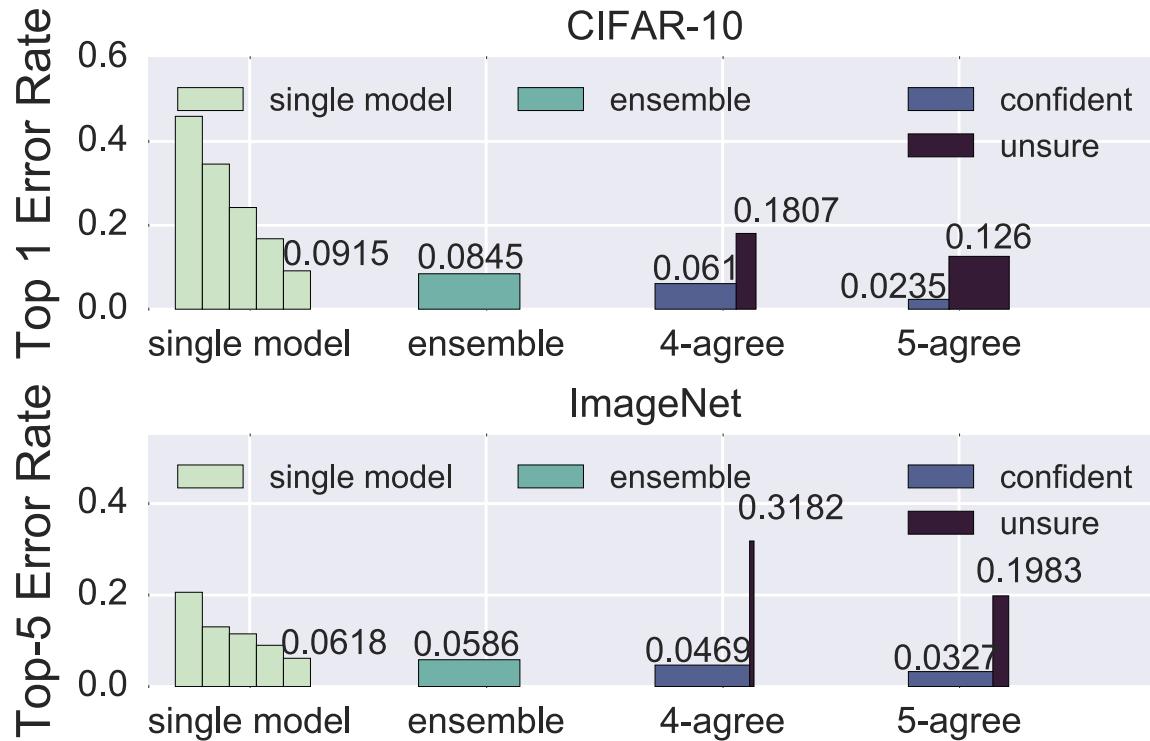


# Anytime Predictions

- Tolerates some loss of models
  - Depends heavily on ensemble



# Ensemble's to Estimate Confidence





# Clipper

theano

Dato

Caffe

F

dmlc

scikit  
learn

VW

Clipper is a prediction serving system that spans multiple ML Frameworks and is designed to

- to **simplifying** model serving
- **bound latency** and **increase throughput**
- and enable **real-time learning** and **personalization across** machine learning frameworks

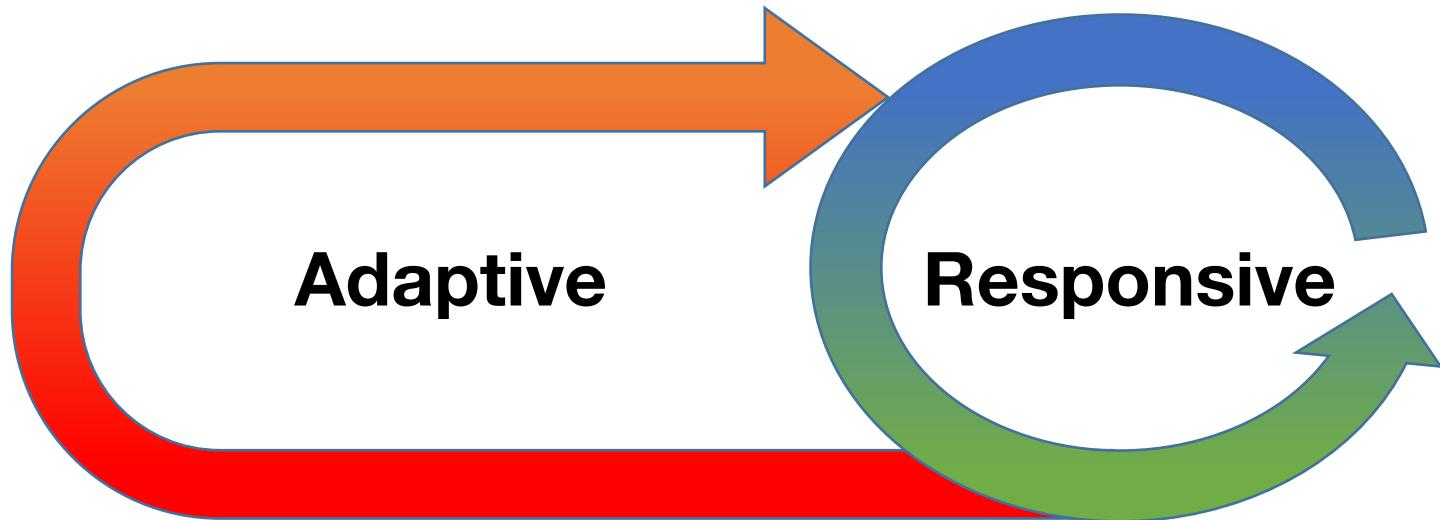
“*Clipper: A Low-Latency Online Prediction Serving System*”  
<https://github.com/ucbrise/clipper> (open source)

# Ongoing Clipper Subprojects

- **Adaptive Batching for Prediction**
  - Leverage internal data-parallelism and hardware acceleration
- **Approximate Caching**
  - Detect “similar” queries and re-use cached predictions
- **Prediction Cascades**
  - Automatically deriving cascades of increasingly GPU intensive models
- **RL/Control**
  - Serving and updating RL policies based on feedback
- **Scheduling and resource allocation**
  - Reduce the need to over-provision for bursty workloads



We are developing new technologies that will enable applications to make low-latency intelligent decisions on live data with strong security guarantees.



**Secure**